



# Operating Systems

David Hay

# SYNCHRONIZATION

# Motivation: Two processes (or threads) want to print

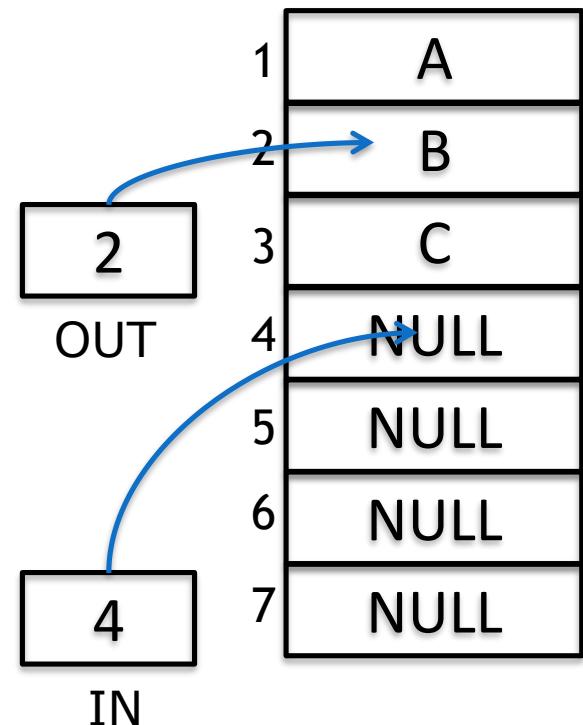
- Both processors write their job to the *spooler* - the (**shared**) queue from which the printer extracts job to print.
  - NULL: No job to print
  - OUT: Next job to print
  - IN: End of queue, write new job here
- Code for each process (adding jobs):  
Spooler[IN] = job  
IN++

IN and Spooler[] are shared among all processes

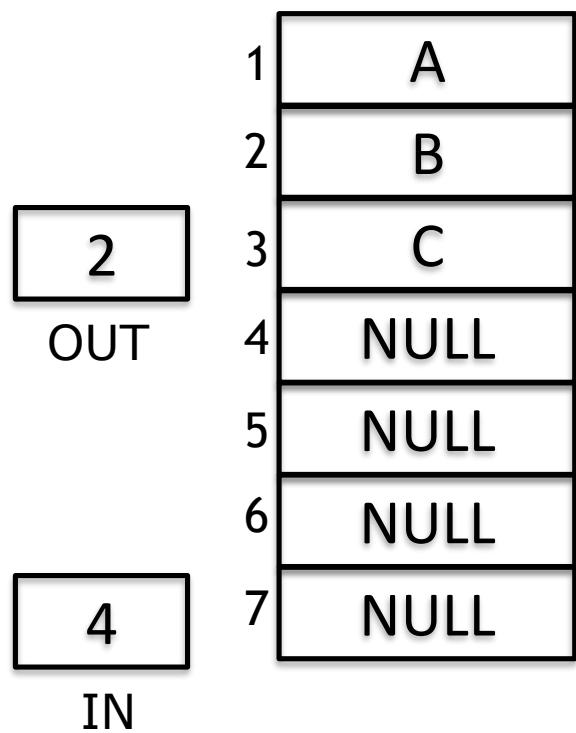
# Motivation: Two processes (or threads) want to print

- Both processors write their job to the *spooler* - the (**shared**) queue from which the printer extracts job to print.
  - NULL: No job to print
  - OUT: Next job to print
  - IN: End of queue, write new job here
- Code for each process (adding jobs):  
`Spooler[IN] = job  
IN++`

IN and Spooler[] are shared among all processes

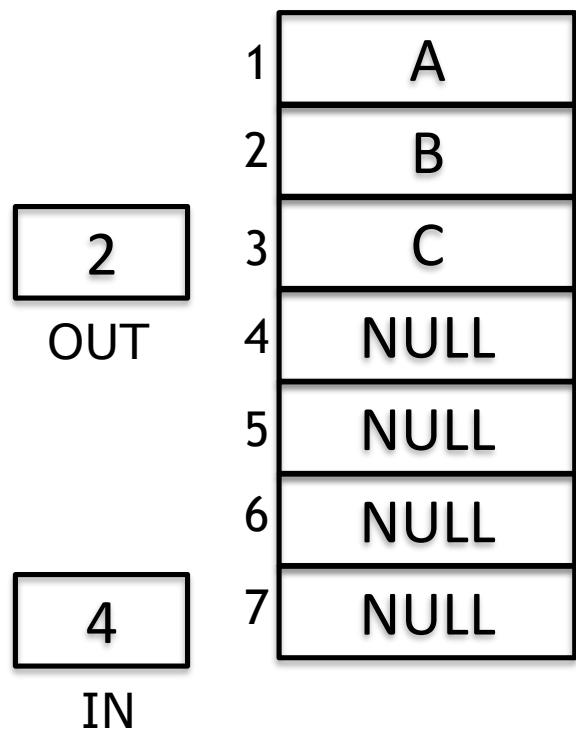


# Possible Outcomes 1/2



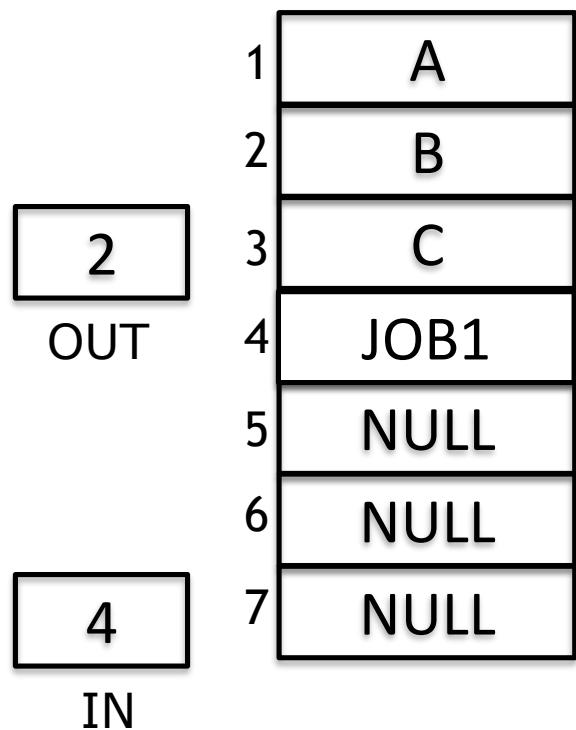
# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1



# Possible Outcomes 1/2

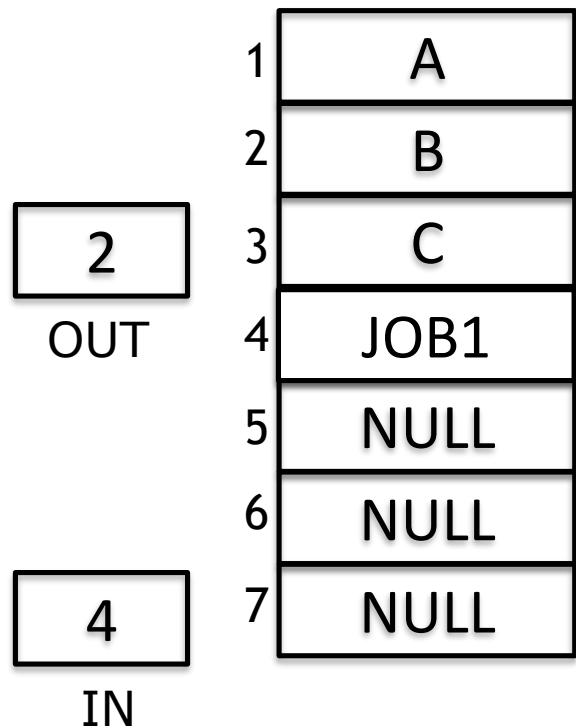
Process 1: Spooler[IN] = JOB1



# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1

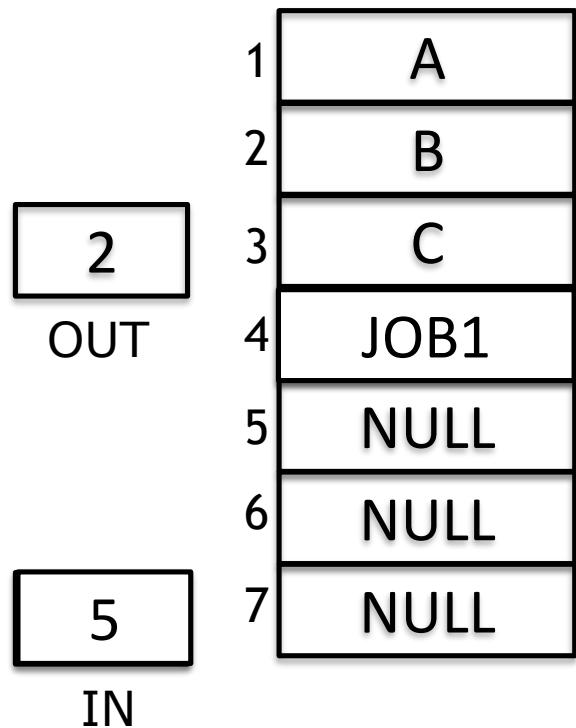
Process 1: IN++



# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1

Process 1: IN++

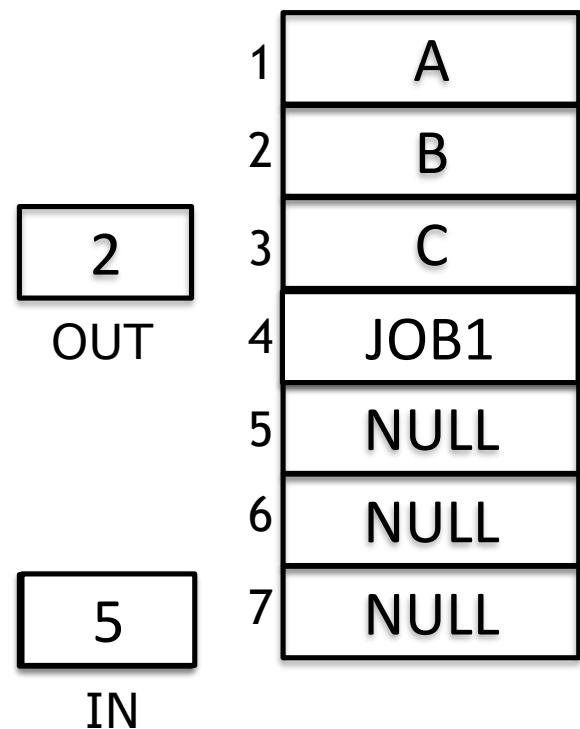


# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1

Process 1: IN++

<Context Switch by OS>



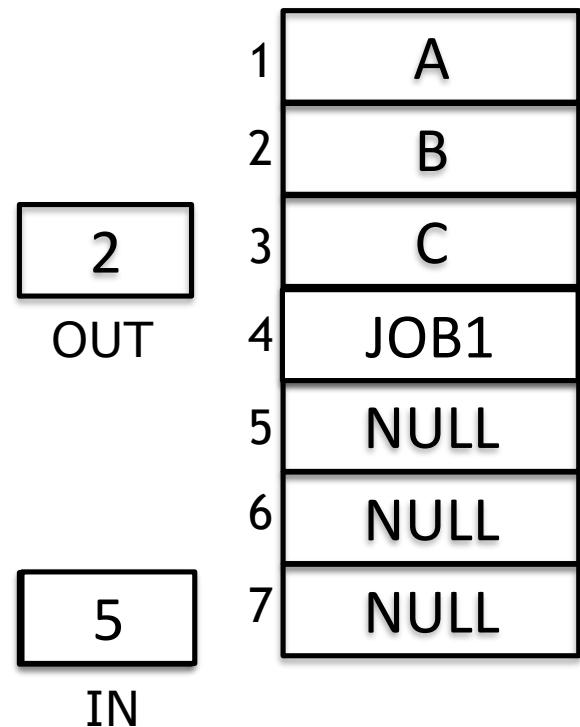
# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1

Process 1: IN++

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2



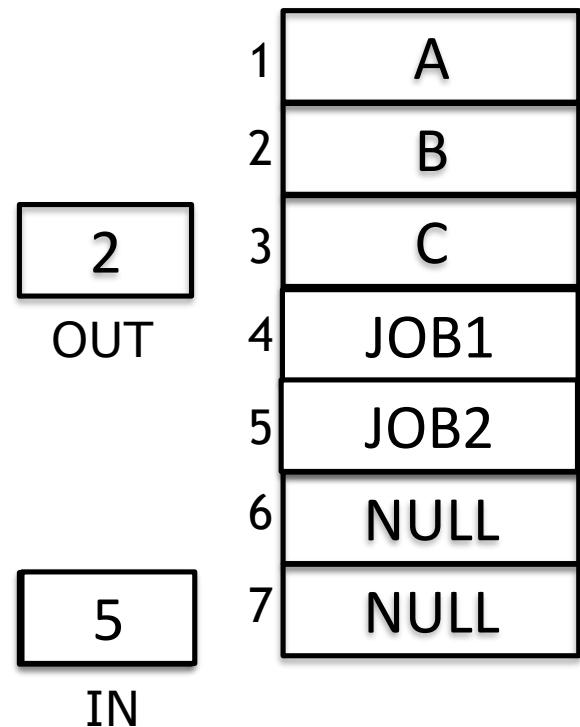
# Possible Outcomes 1/2

Process 1: Spooler[IN] = JOB1

Process 1: IN++

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2



# Possible Outcomes 1/2

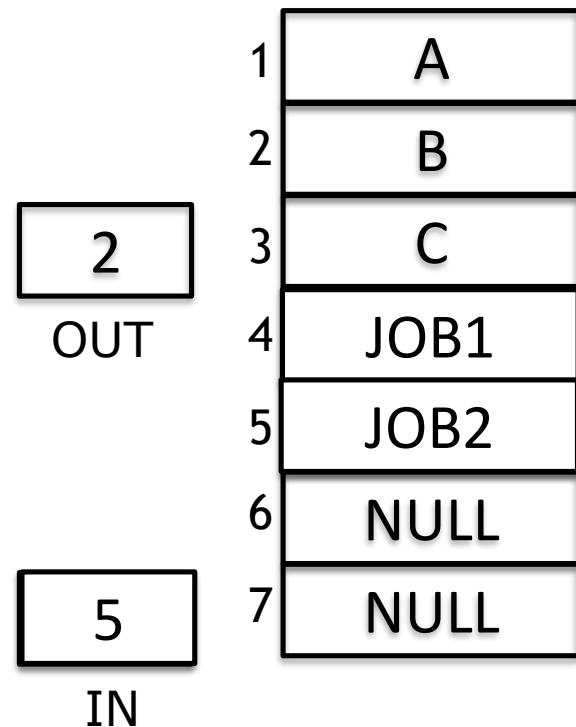
Process 1: Spooler[IN] = JOB1

Process 1: IN++

**<Context Switch by OS>**

Process 2: Spooler[IN] = JOB2

Process 2: IN++



# Possible Outcomes 1/2

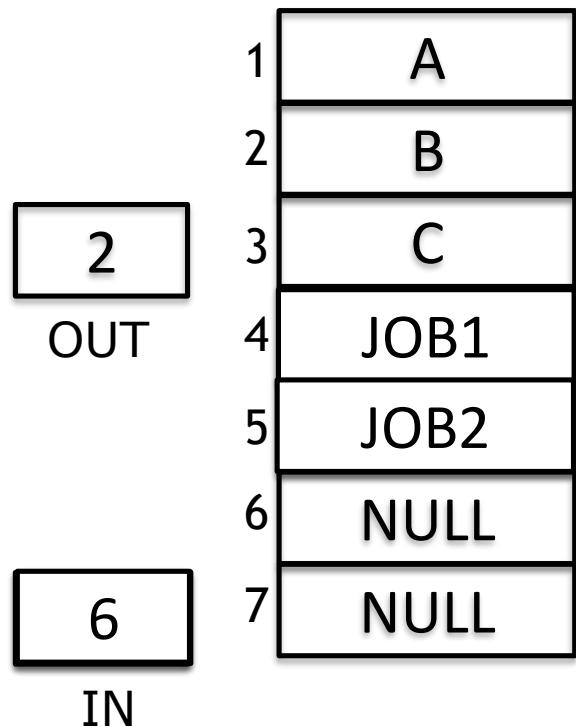
Process 1: Spooler[IN] = JOB1

Process 1: IN++

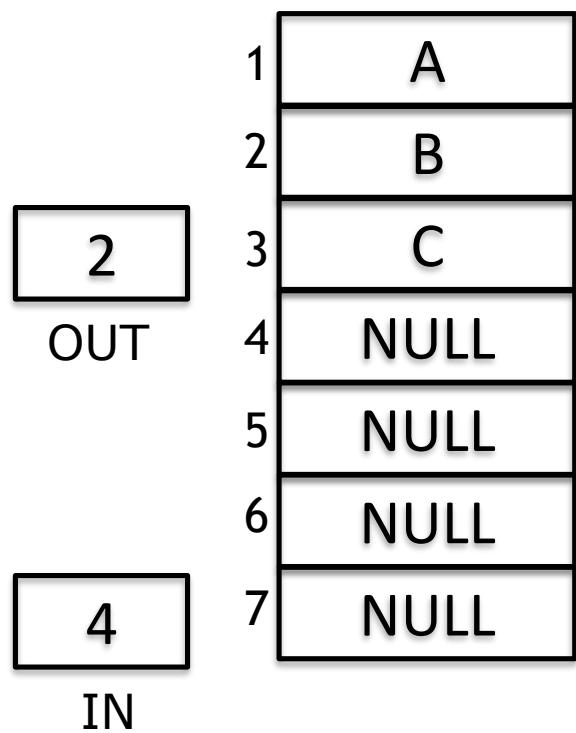
<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

Process 2: IN++

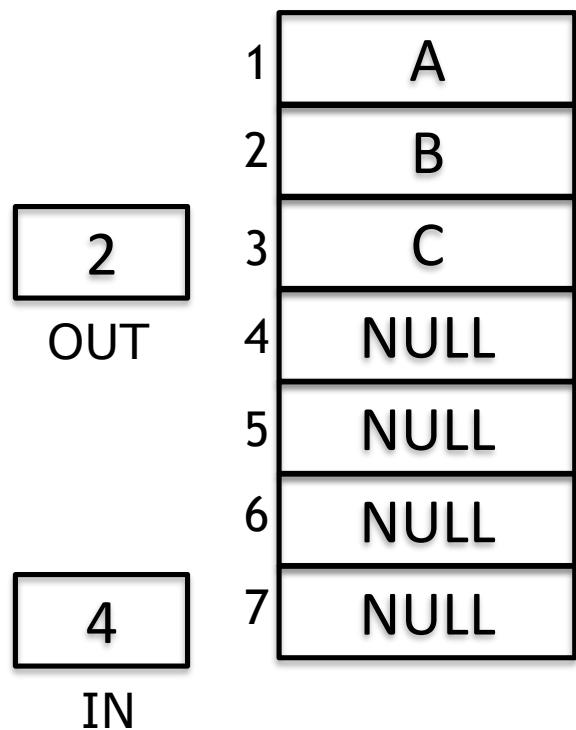


# Possible Outcomes 2/2



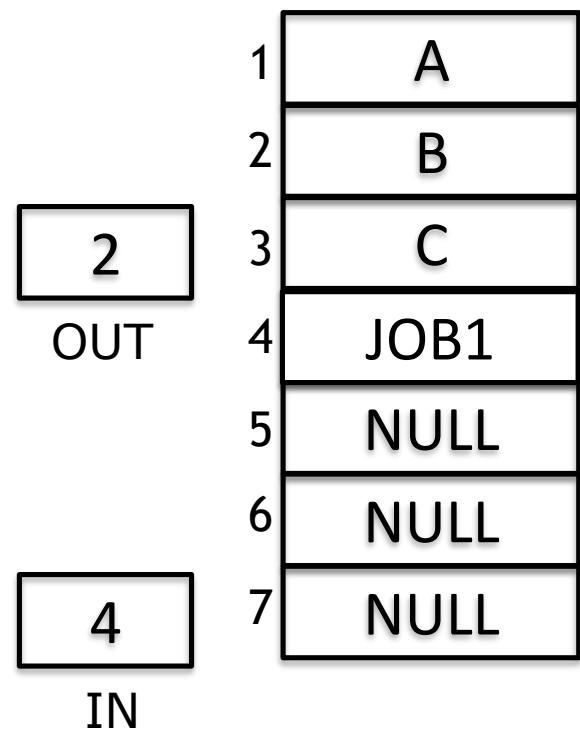
# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1



# Possible Outcomes 2/2

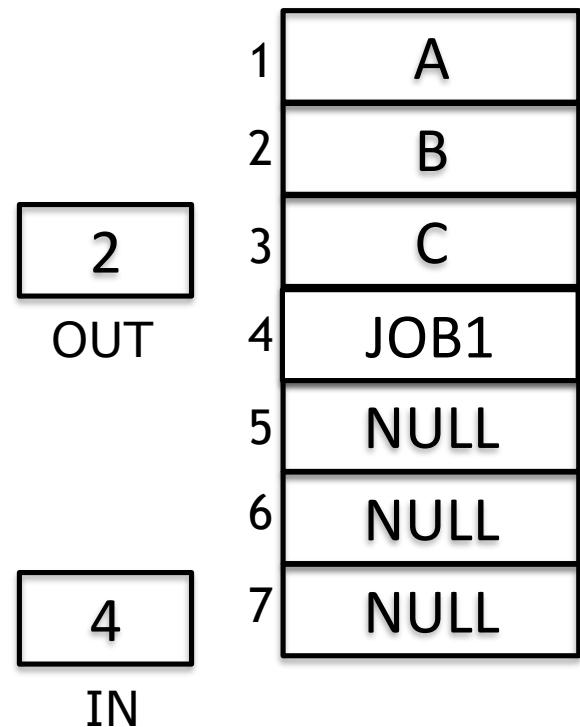
Process 1: Spooler[IN] = JOB1



# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

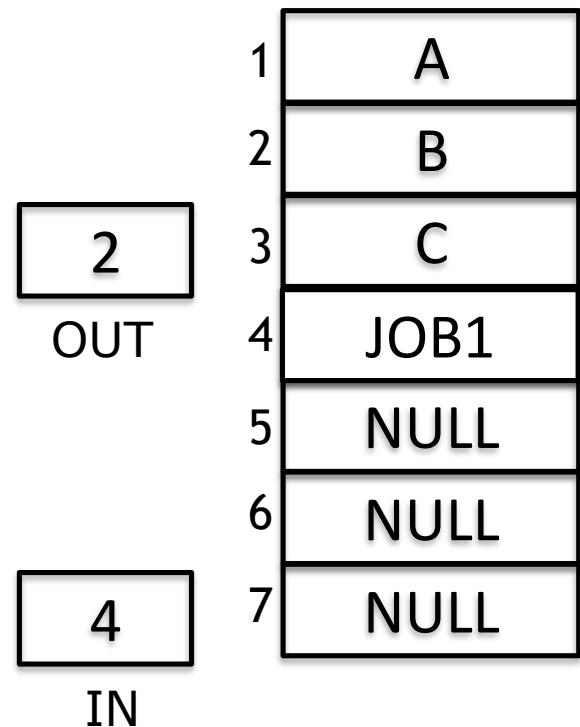


# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

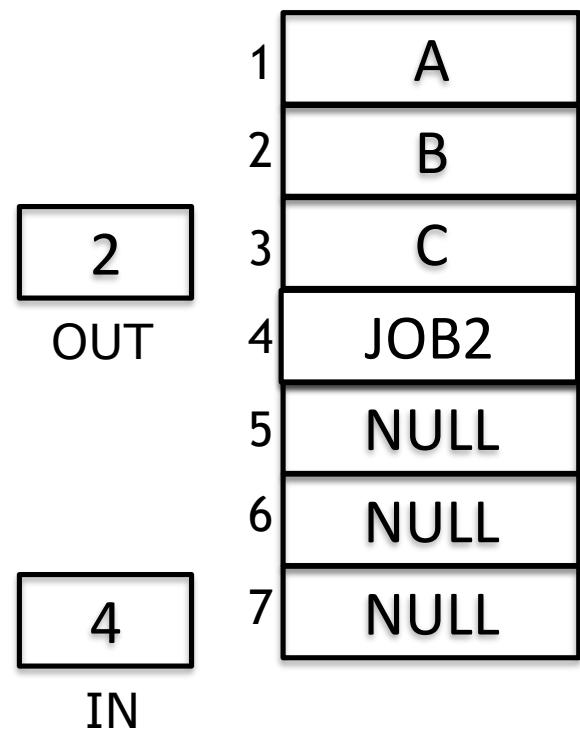


# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2



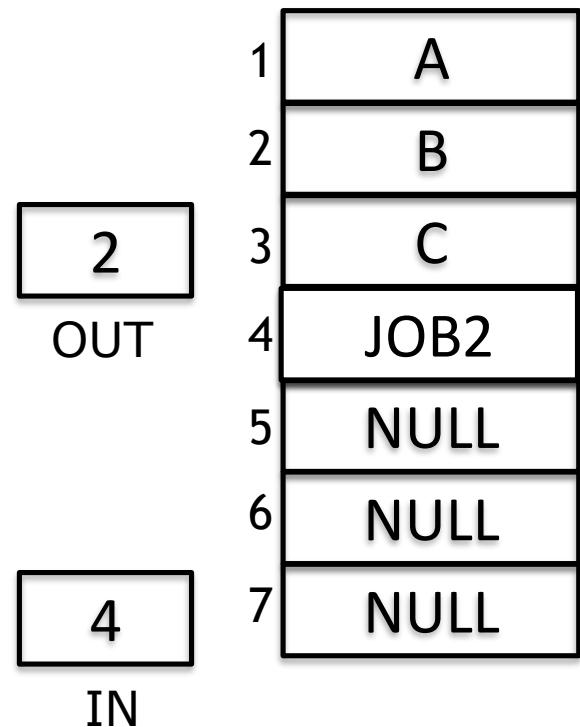
# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

<Context Switch by OS>



# Possible Outcomes 2/2

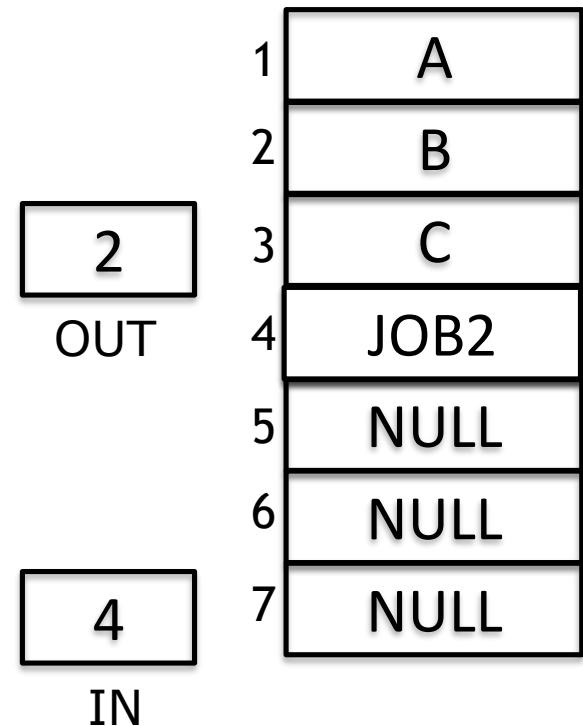
Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

<Context Switch by OS>

Process 1: IN++



# Possible Outcomes 2/2

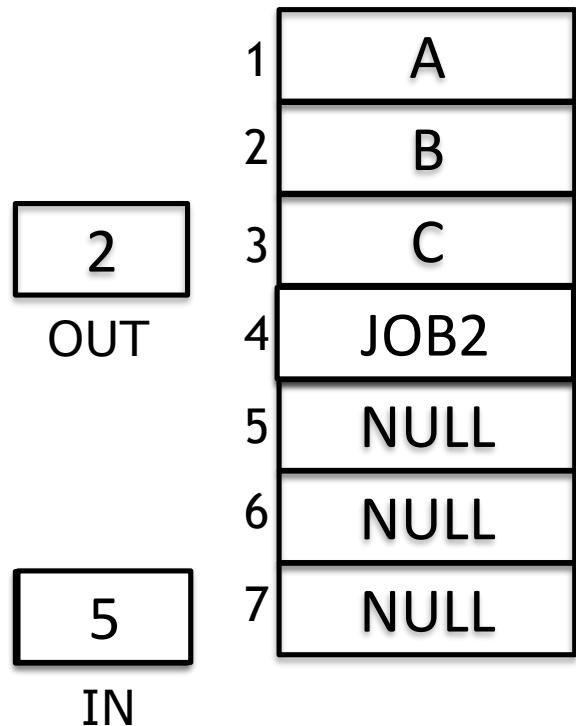
Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

<Context Switch by OS>

Process 1: IN++



# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

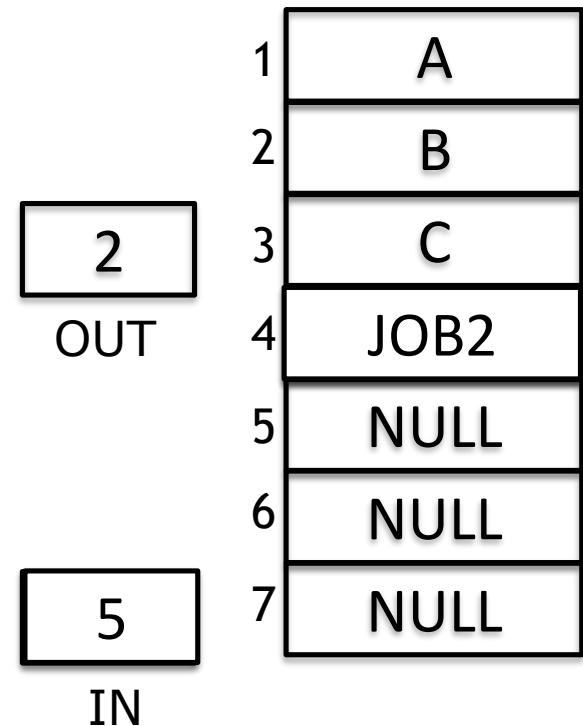
<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

<Context Switch by OS>

Process 1: IN++

<Context Switch by OS>



# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

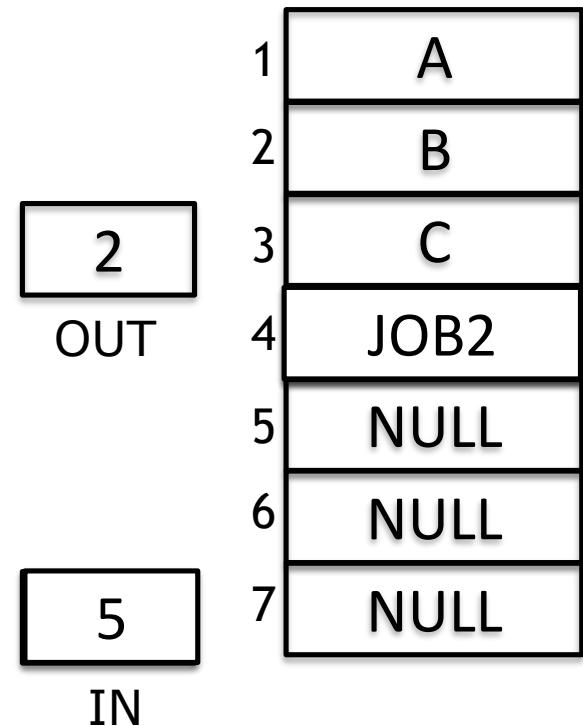
Process 2: Spooler[IN] = JOB2

<Context Switch by OS>

Process 1: IN++

<Context Switch by OS>

Process 2: IN++



# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

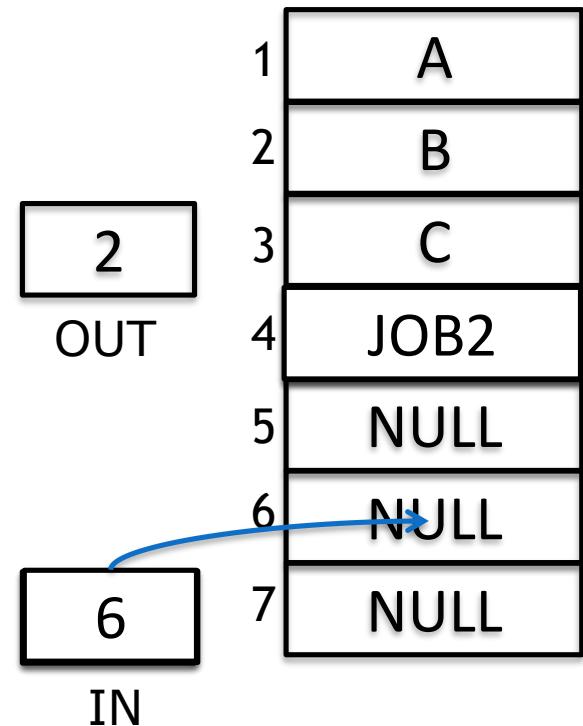
Process 2: Spooler[IN] = JOB2

<Context Switch by OS>

Process 1: IN++

<Context Switch by OS>

Process 2: IN++



# Possible Outcomes 2/2

Process 1: Spooler[IN] = JOB1

<Context Switch by OS>

Process 2: Spooler[IN] = JOB2

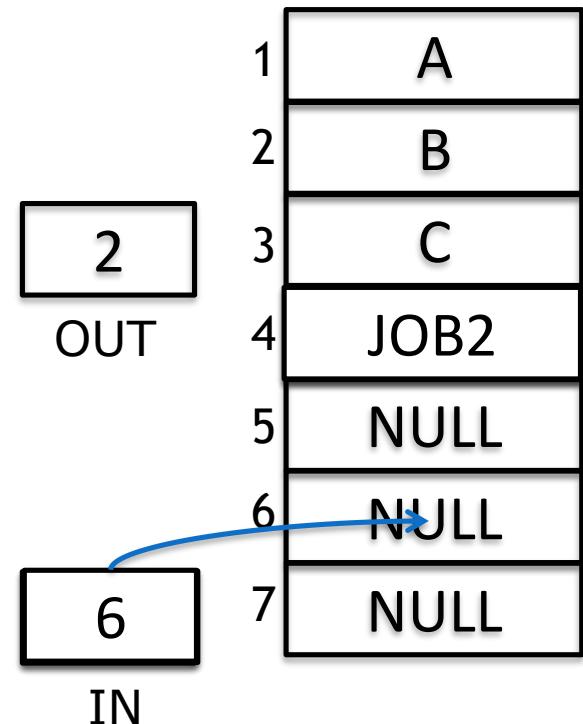
<Context Switch by OS>

Process 1: IN++

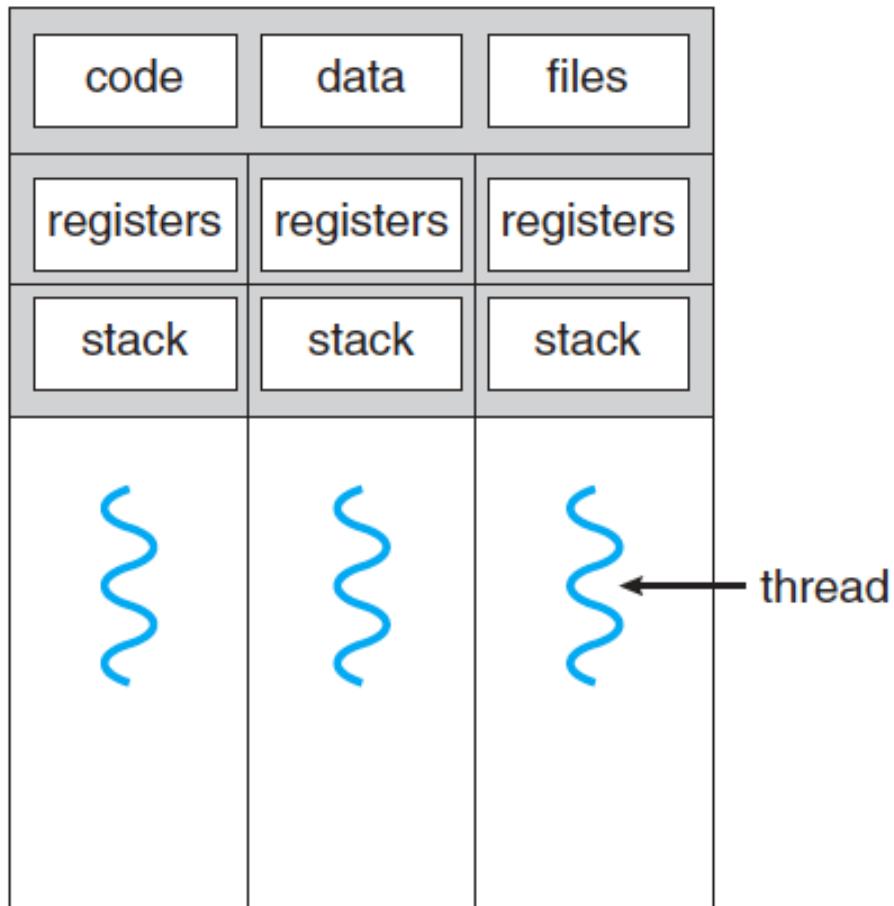
<Context Switch by OS>

Process 2: IN++

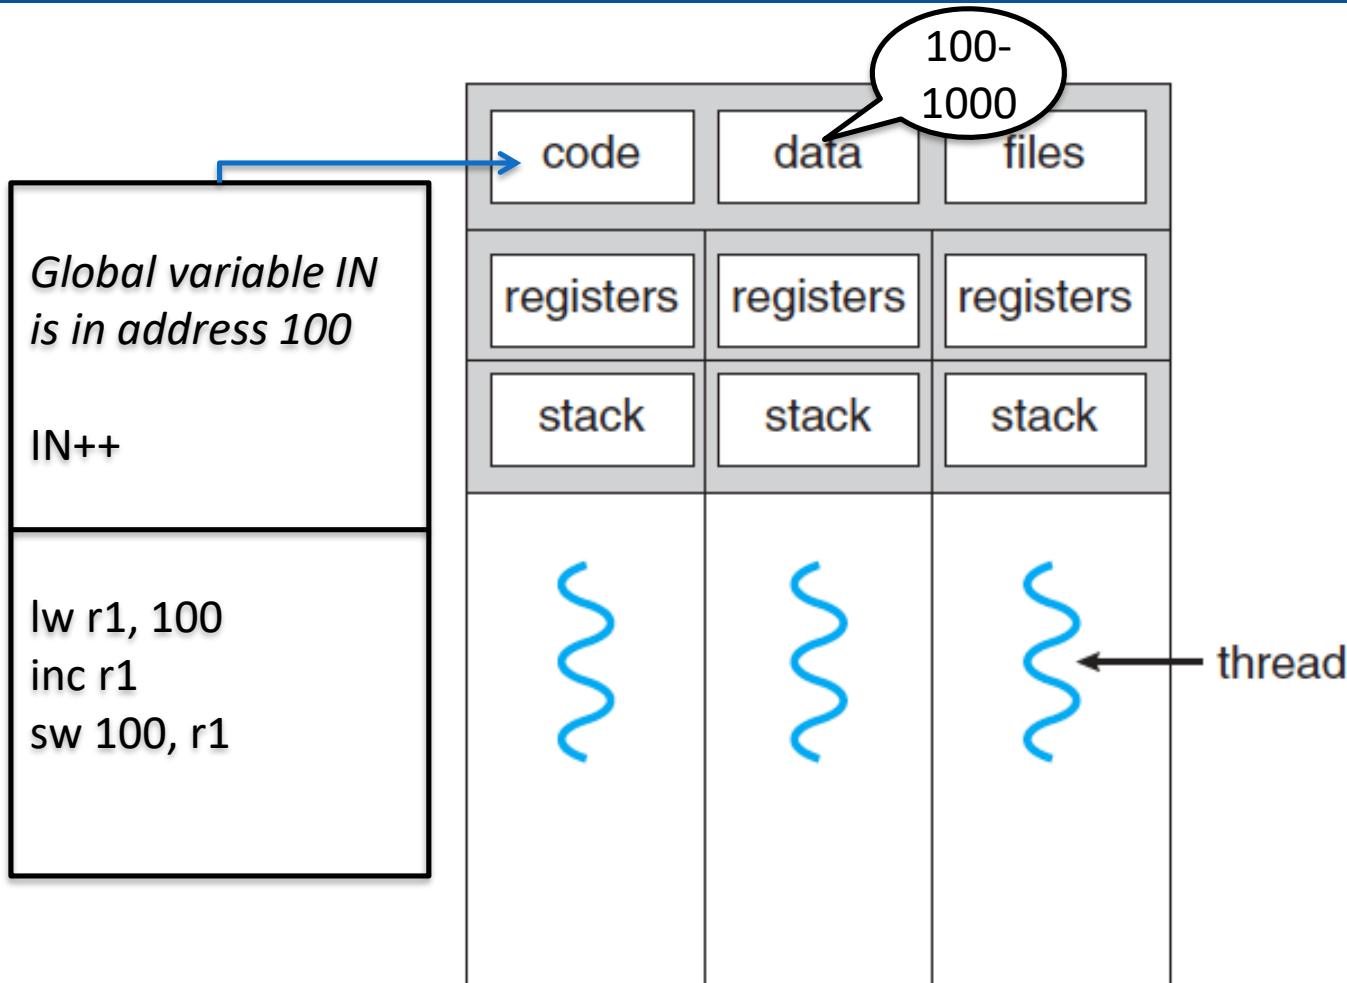
→ JOB1 is lost; all jobs after JOB2 may not be printed as well, because of the gap



# Further Motivation

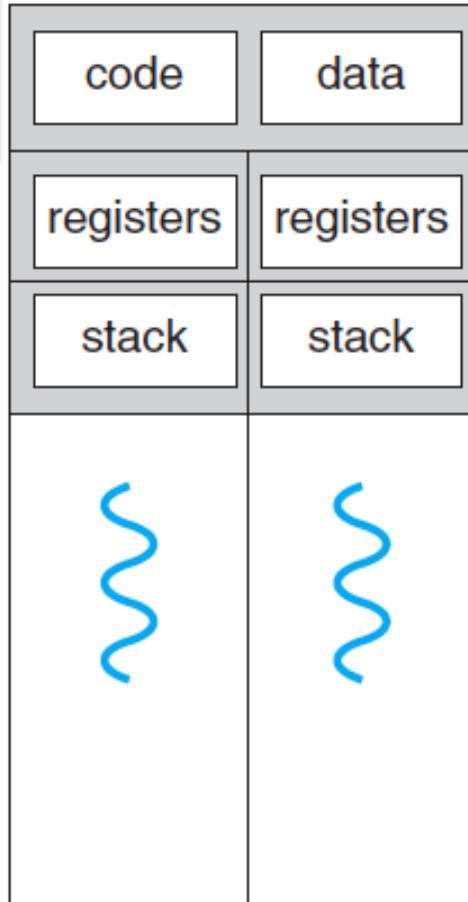


# Further Motivation



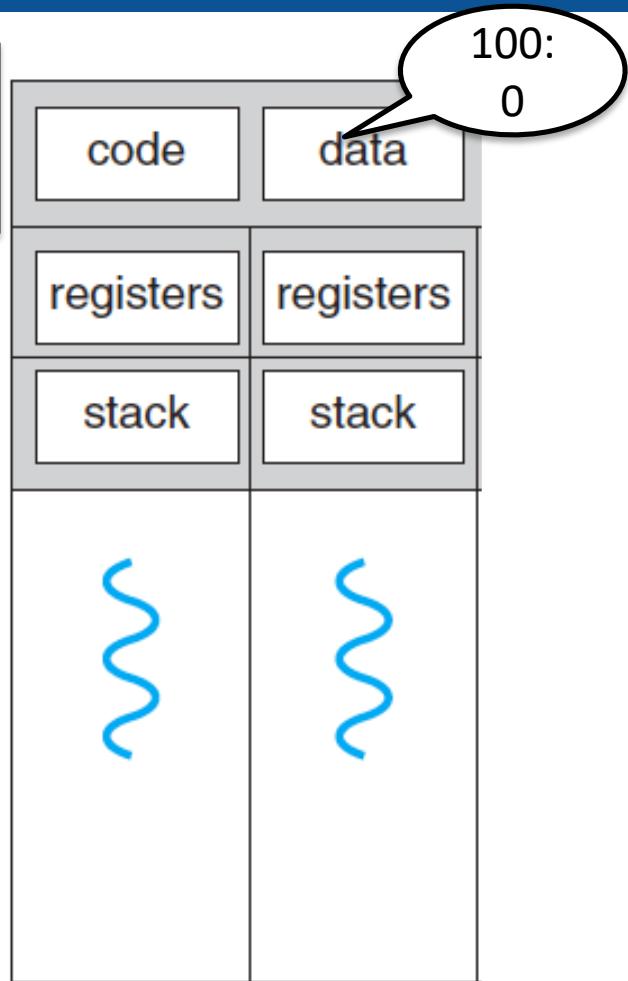
# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```



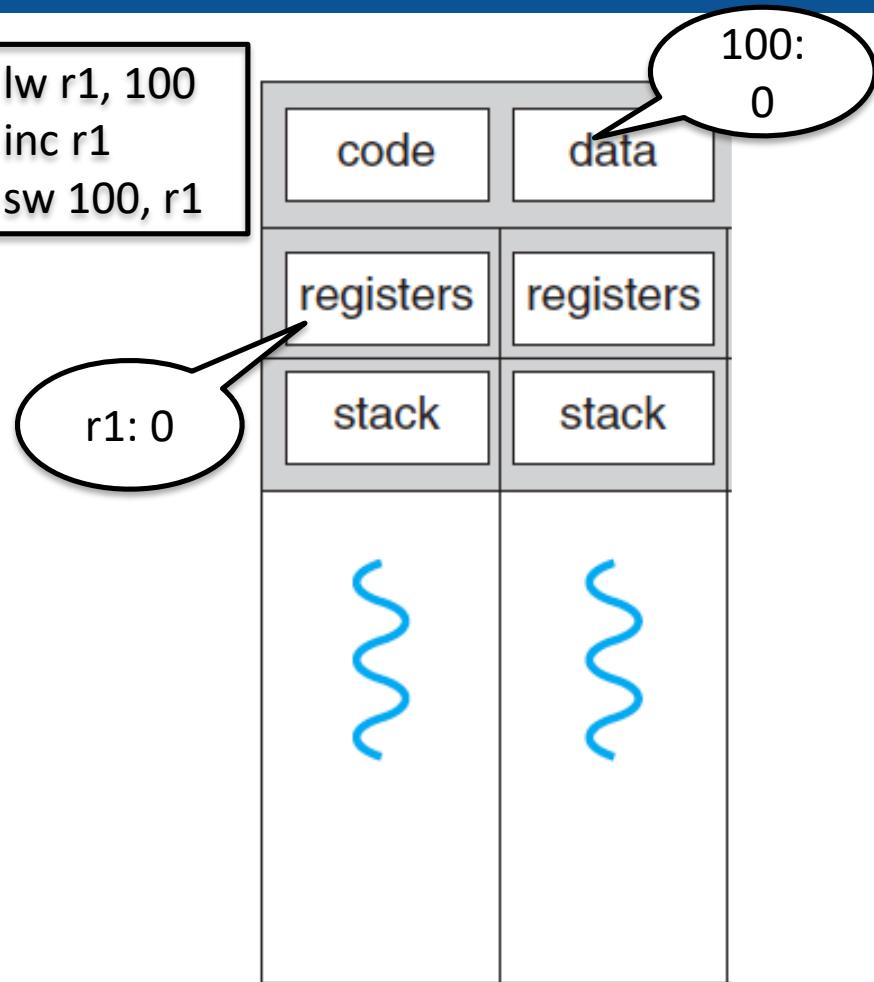
# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```



# Further Motivation

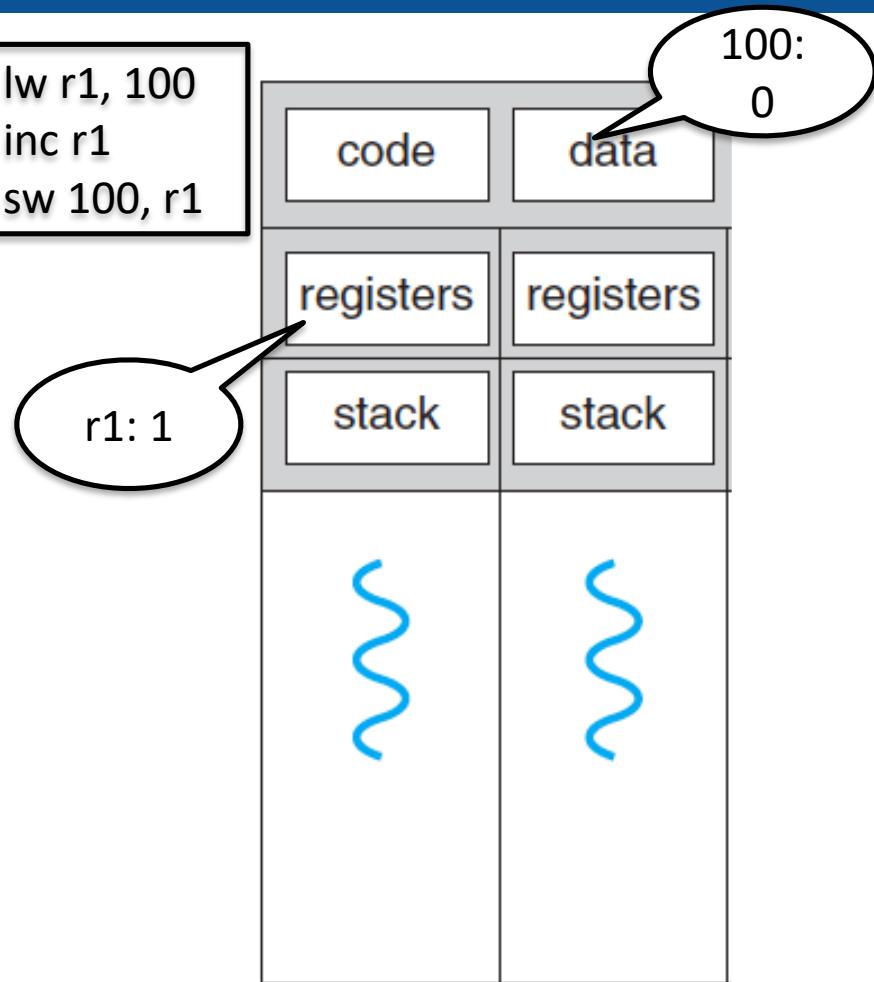
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: lw r1,100

# Further Motivation

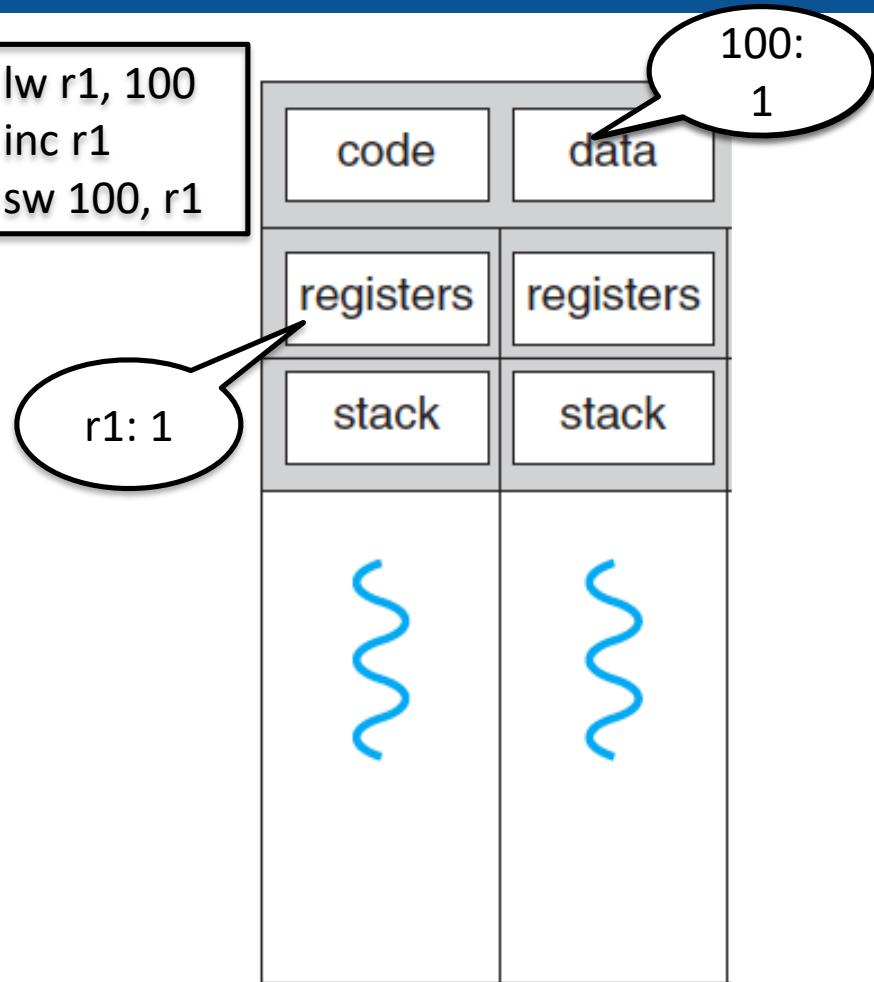
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: lw r1,100
- Thread 1: inc r1

# Further Motivation

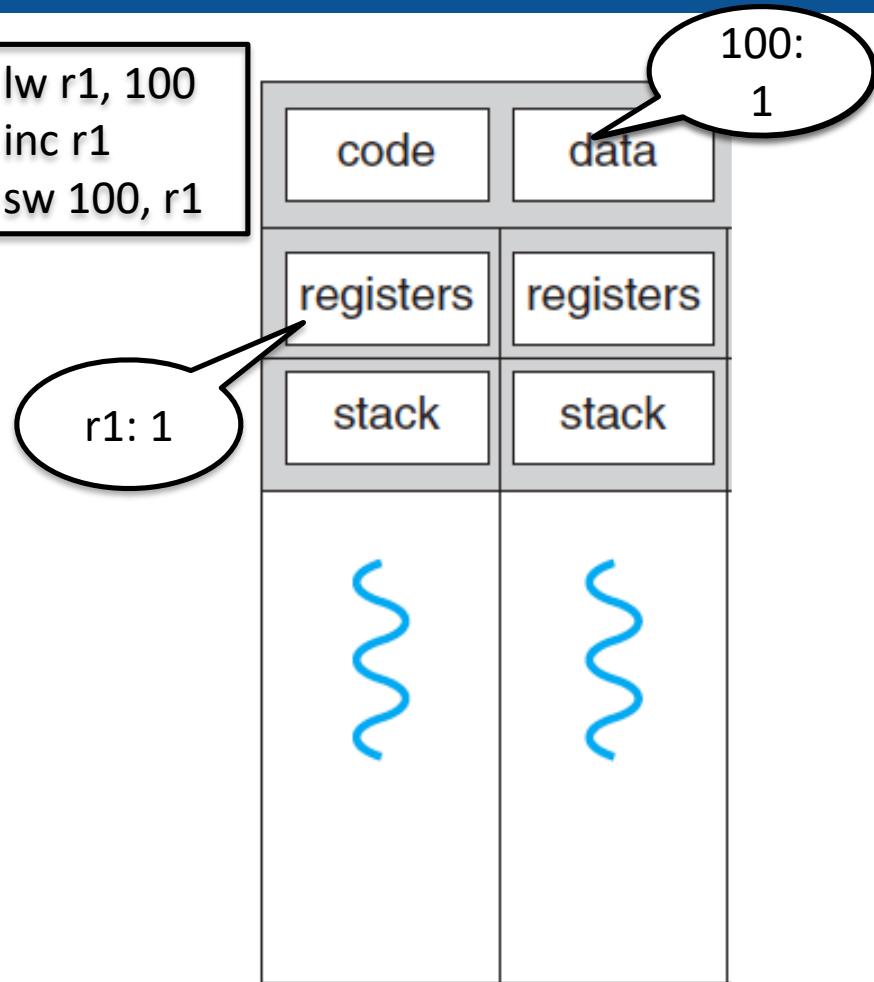
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: `lw r1,100`
- Thread 1: `inc r1`
- Thread 1: `sw 100,r1`

# Further Motivation

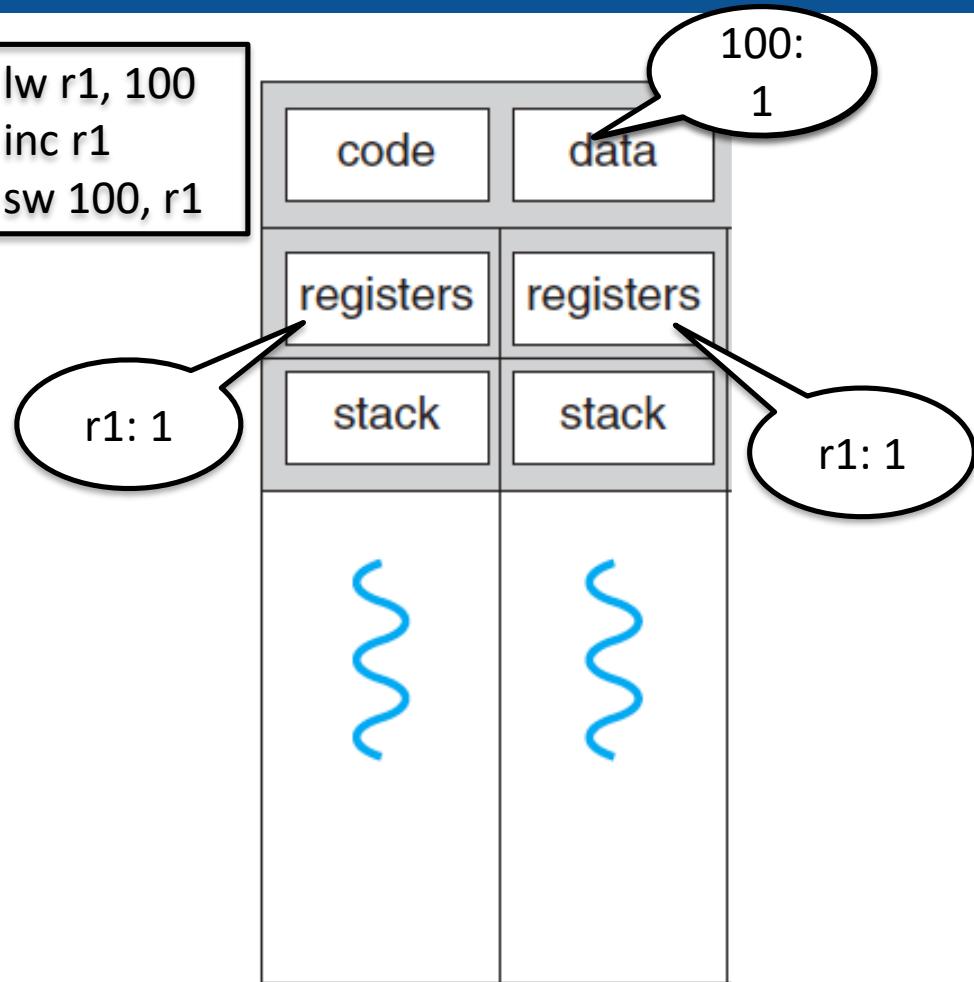
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: lw r1,100
- Thread 1: inc r1
- Thread 1: sw 100,r1
- <Context Switch>

# Further Motivation

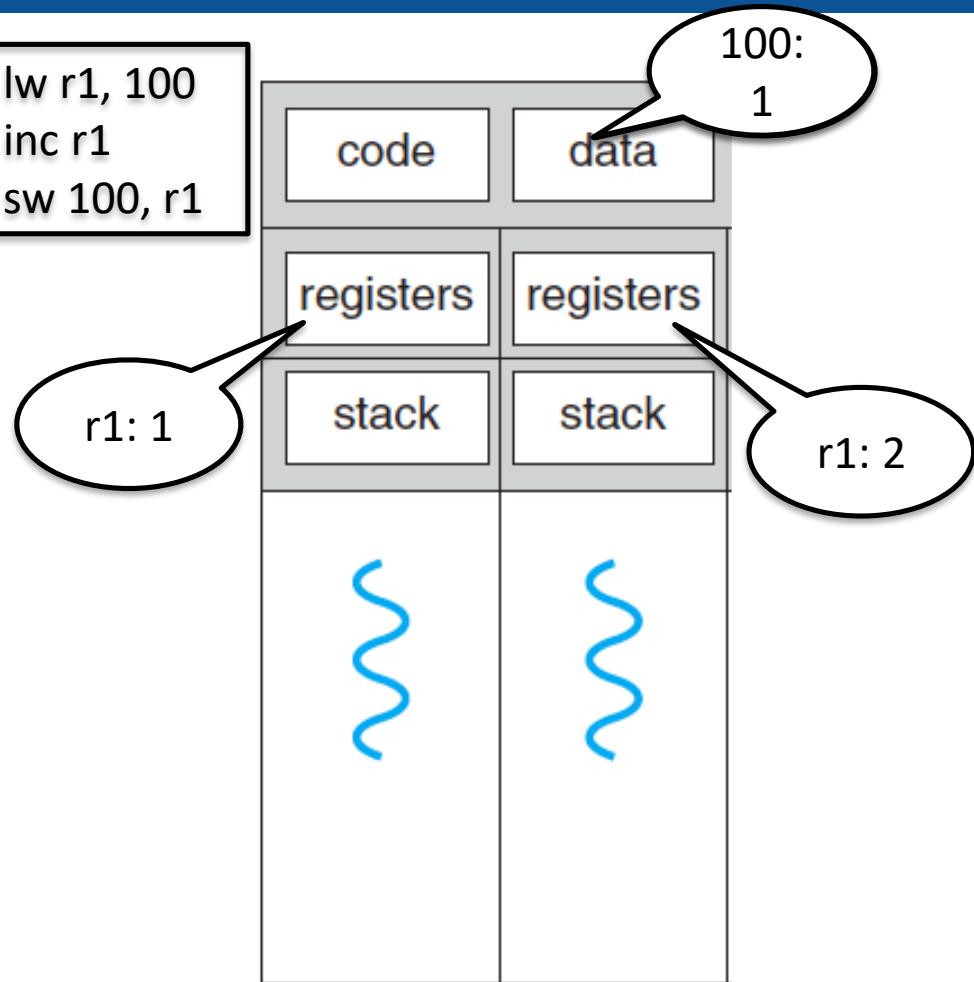
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- Thread 1: `sw 100, r1`
- **<Context Switch>**
- Thread 2: `lw r1, 100`

# Further Motivation

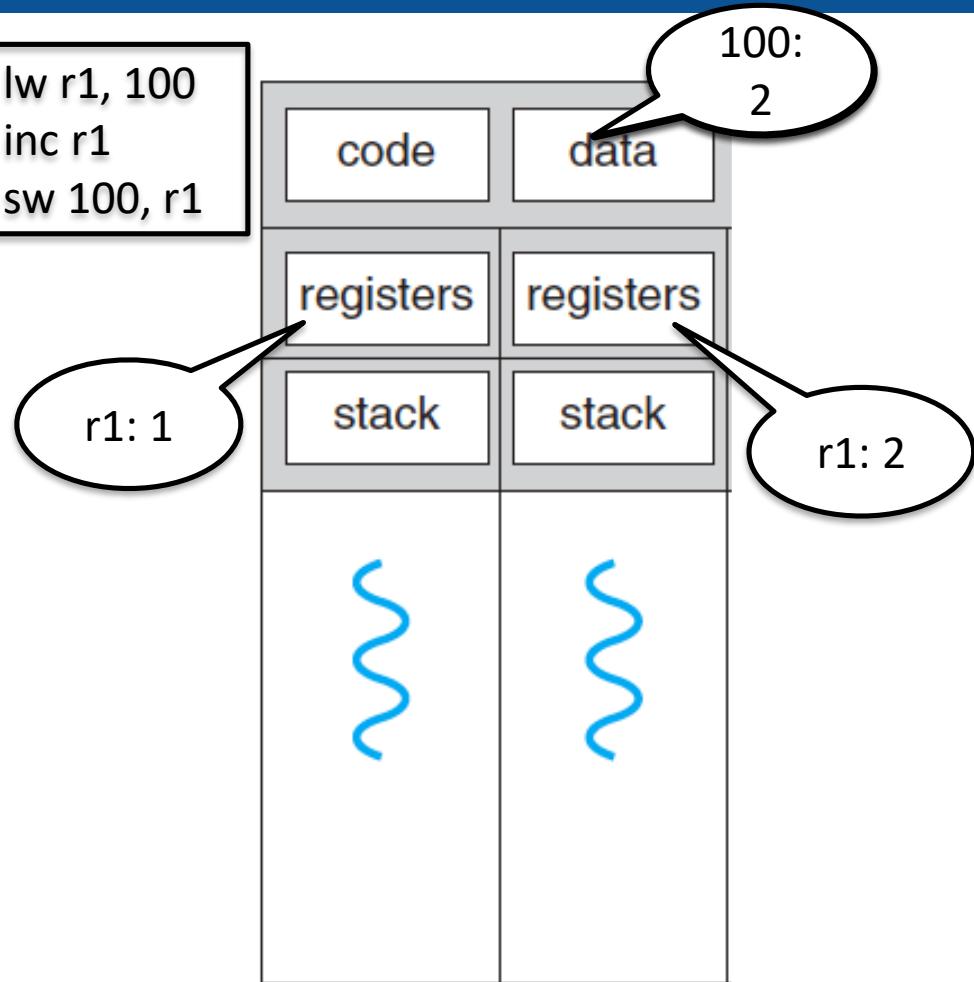
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- Thread 1: `sw 100, r1`
- **<Context Switch>**
- Thread 2: `lw r1, 100`
- Thread 2: `inc r1`

# Further Motivation

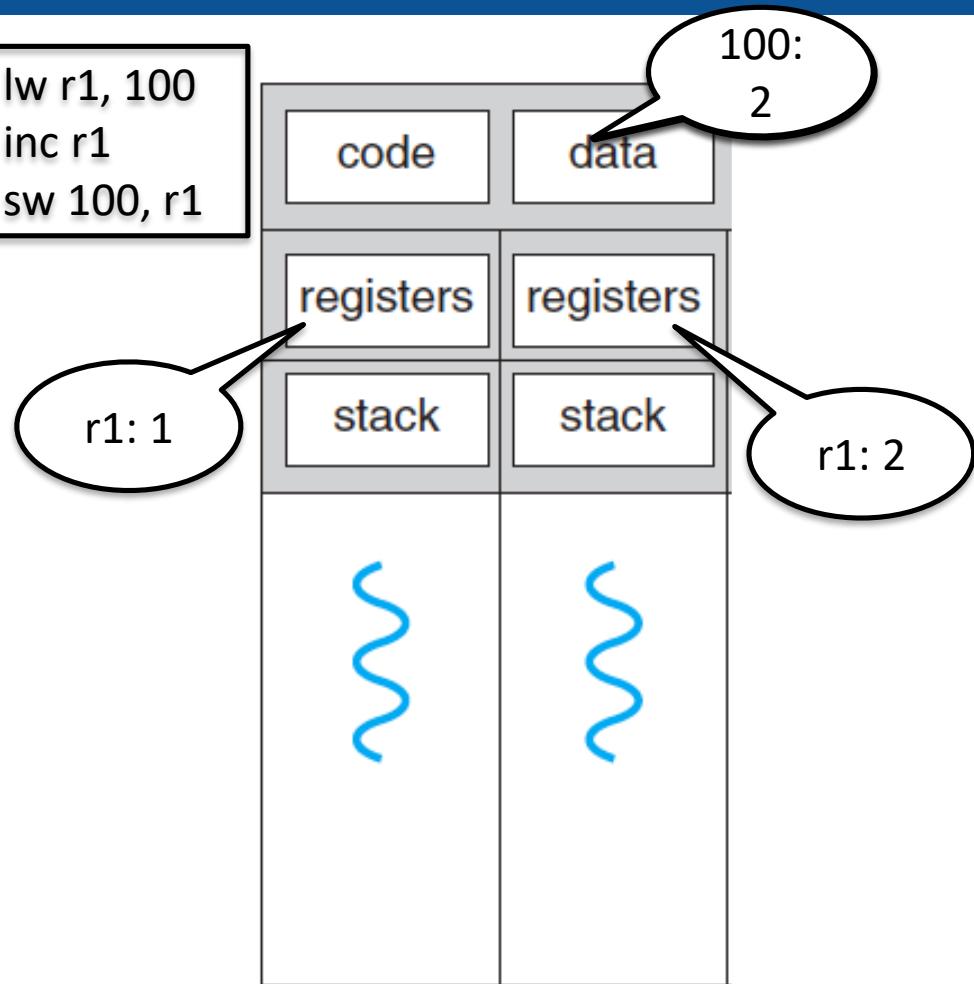
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: lw r1, 100
- Thread 1: inc r1
- Thread 1: sw 100, r1
- <Context Switch>
- Thread 2: lw r1, 100
- Thread 2: inc r1
- Thread 2: sw 100, r1

# Further Motivation

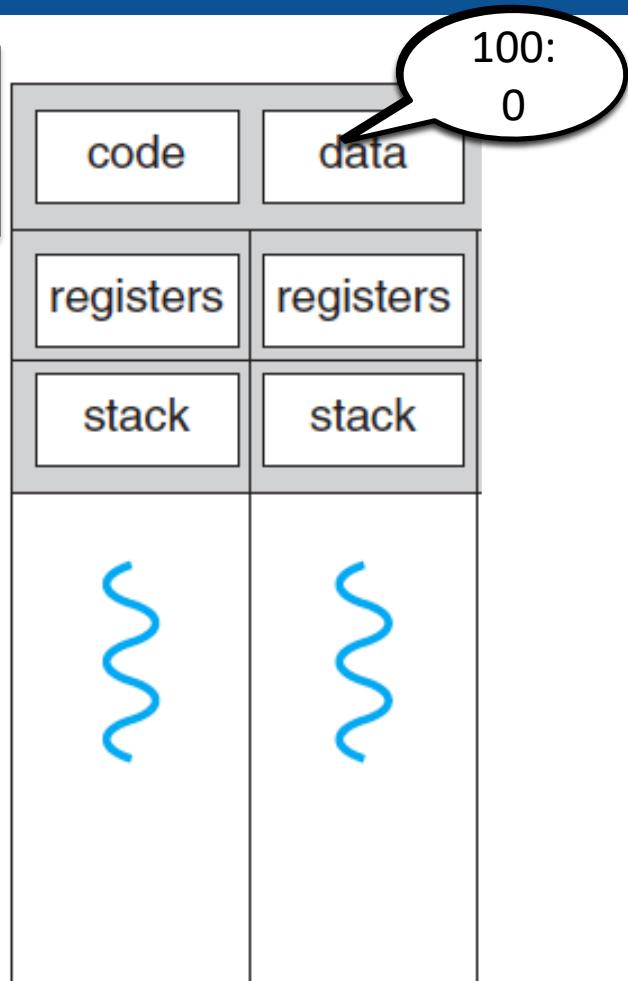
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- IN = 2

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

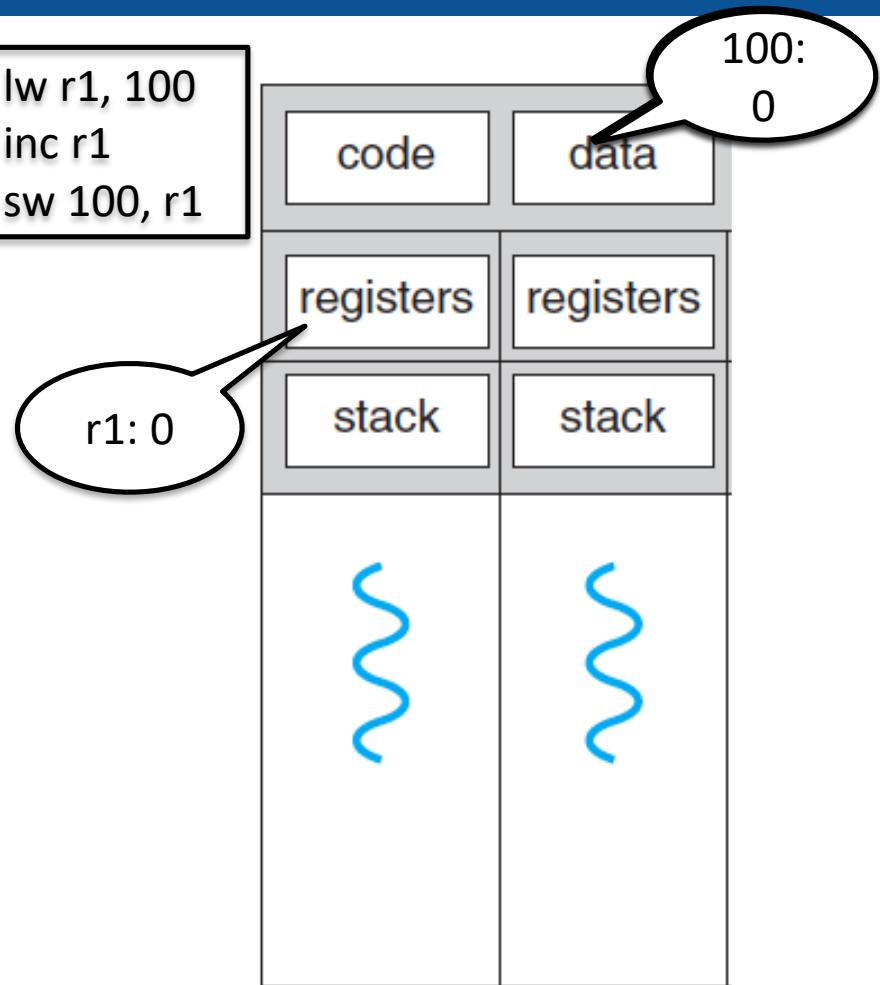


- Thread 1: lw r1,100
- Thread 1: inc r1
- Thread 1: sw 100,r1
- <Context Switch>
- Thread 2: lw r1, 100
- Thread 2: inc r1
- Thread 2: sw 100, r1

→ IN = 2

# Further Motivation

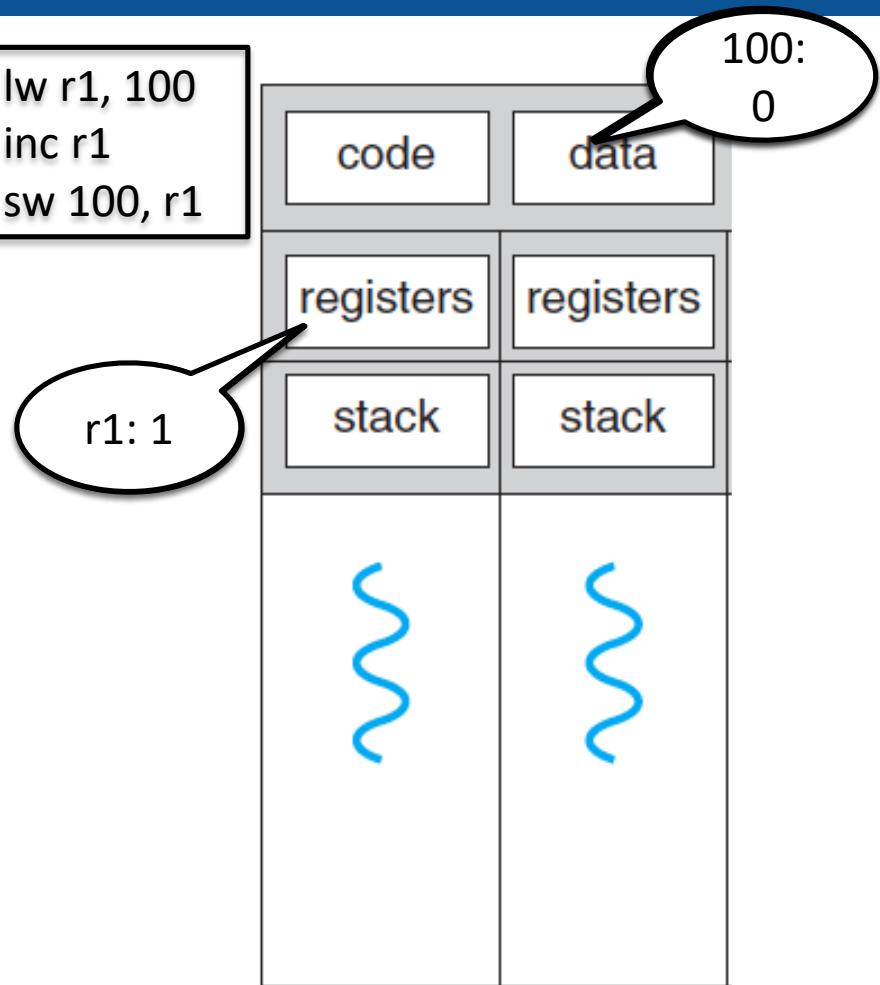
```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: lw r1,100
  - Thread 1: inc r1
  - Thread 1: sw 100,r1
  - <Context Switch>
  - Thread 2: lw r1, 100
  - Thread 2: inc r1
  - Thread 2: sw 100, r1
- IN = 2
- Thread 1: lw r1,100

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

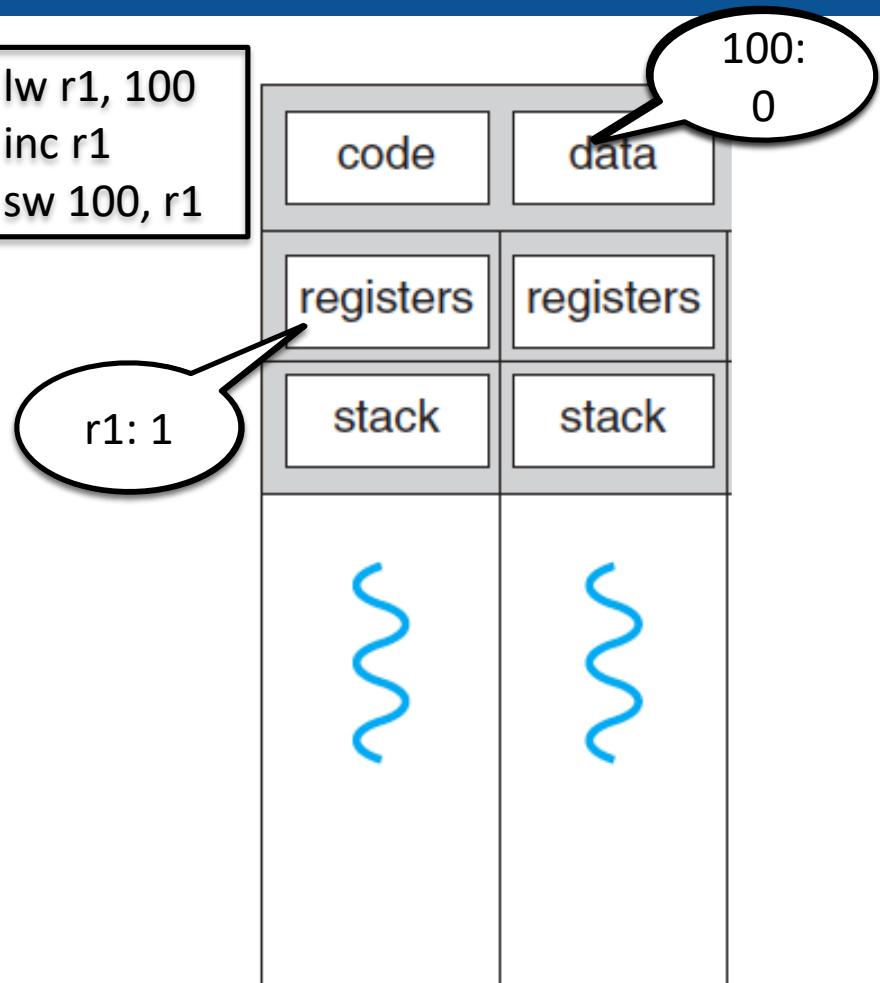


- Thread 1: lw r1,100
  - Thread 1: inc r1
  - Thread 1: sw 100,r1
  - <Context Switch>
  - Thread 2: lw r1, 100
  - Thread 2: inc r1
  - Thread 2: sw 100, r1
- IN = 2

- Thread 1: lw r1,100
- Thread 1: inc r1

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

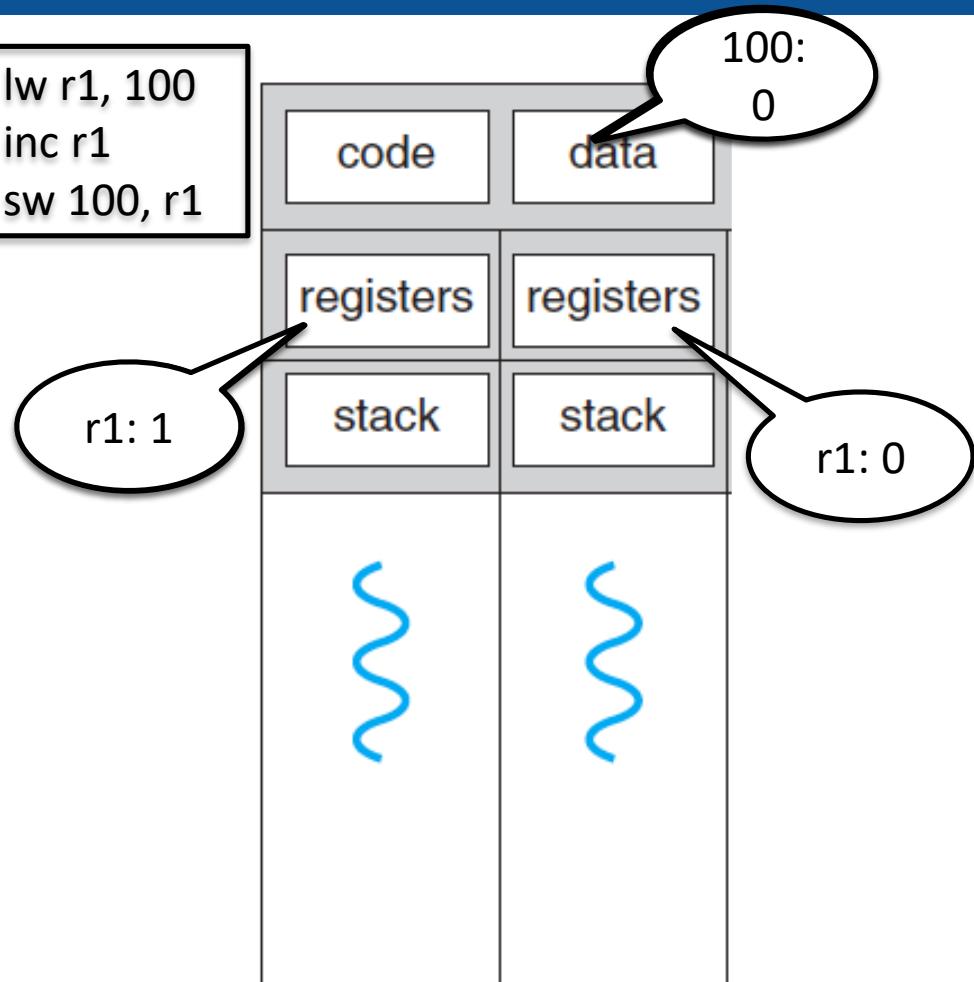


- Thread 1: lw r1,100
  - Thread 1: inc r1
  - Thread 1: sw 100,r1
  - <Context Switch>
  - Thread 2: lw r1, 100
  - Thread 2: inc r1
  - Thread 2: sw 100, r1
- IN = 2

- Thread 1: lw r1,100
- Thread 1: inc r1
- <Context Switch>

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

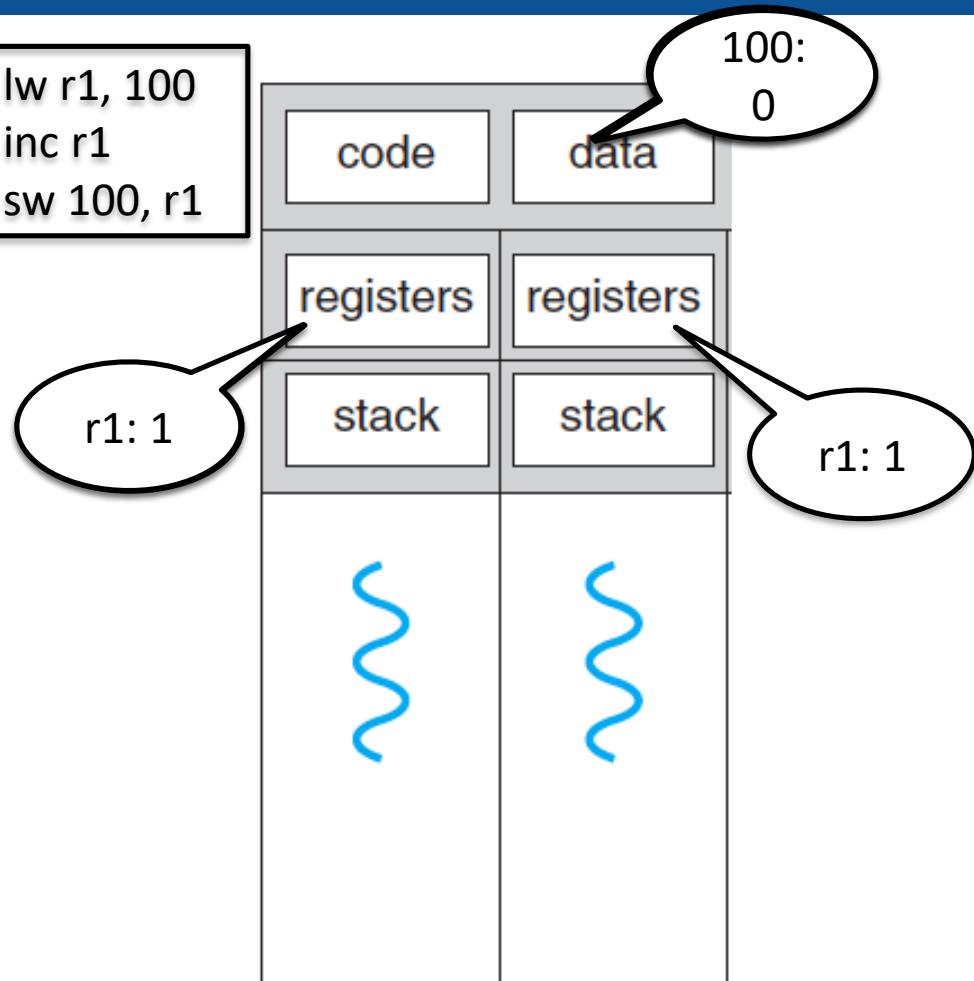


- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- IN = 2

- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- **<Context Switch>**
- Thread 2: `lw r1, 100`

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

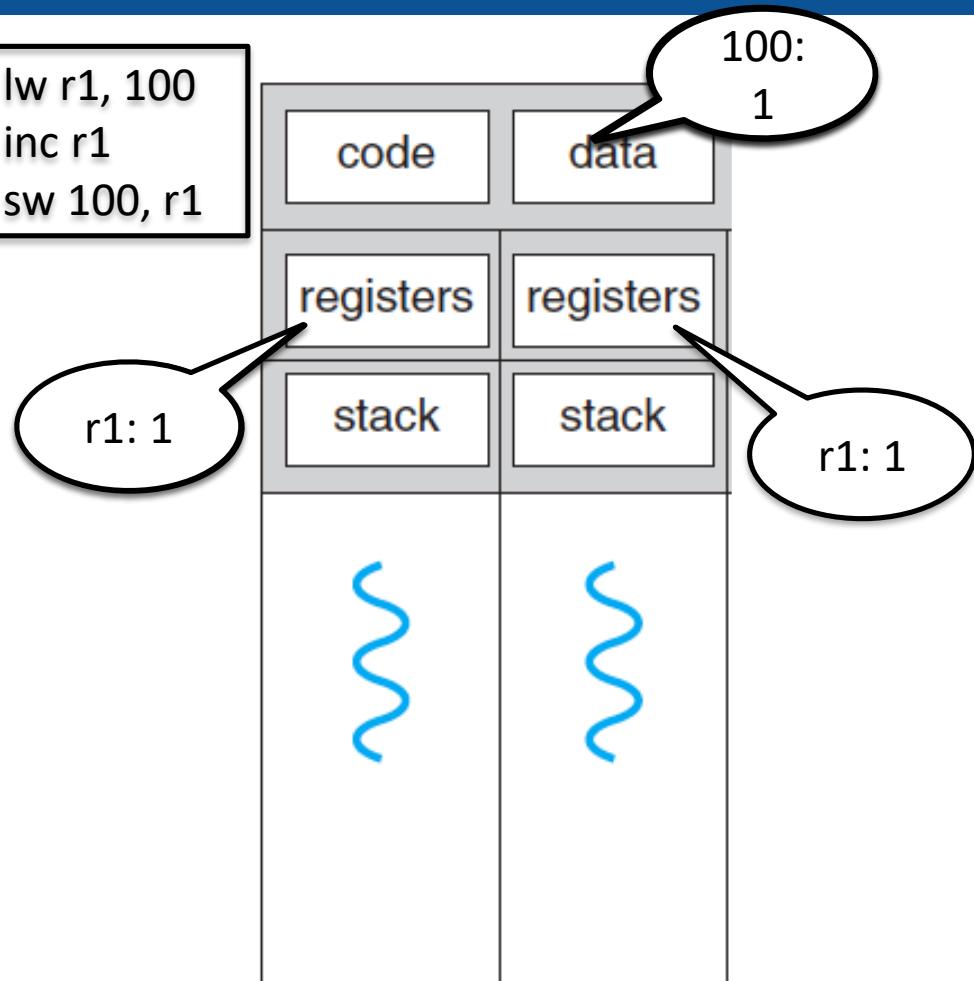


- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - <Context Switch>
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- IN = 2

- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- <Context Switch>
- Thread 2: `lw r1, 100`
- Thread 2: `inc r1`

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

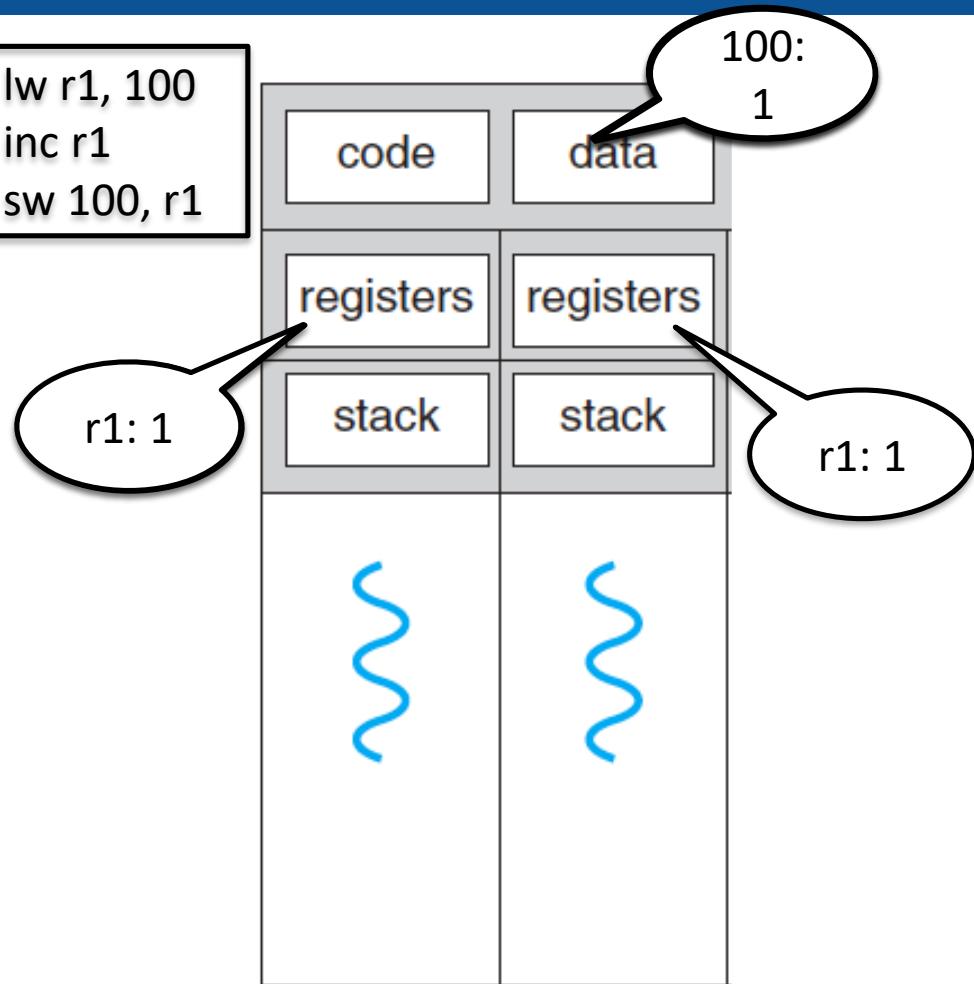


- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- $\rightarrow \text{IN} = 2$

- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- **<Context Switch>**
- Thread 2: `lw r1, 100`
- Thread 2: `inc r1`
- Thread 2: `sw 100, r1`

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

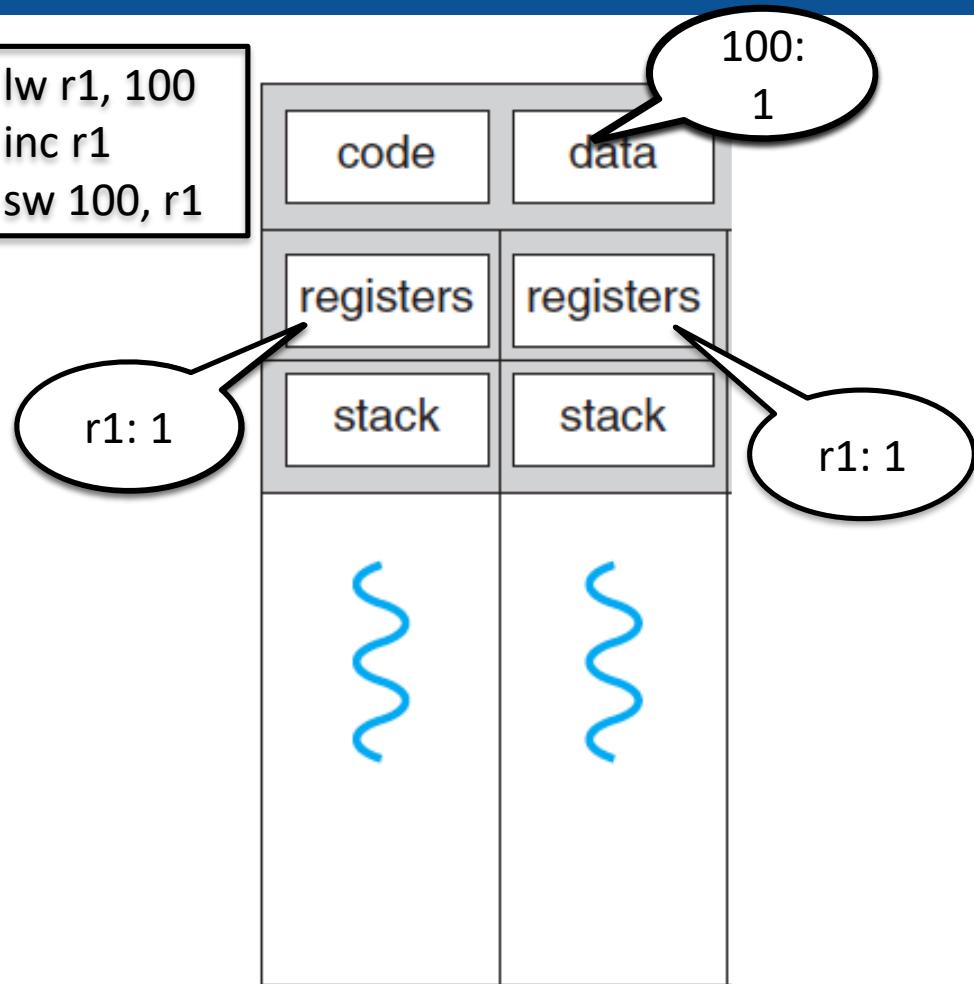


- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- IN = 2

- Thread 1: `lw r1, 100`
- Thread 1: `inc r1`
- **<Context Switch>**
- Thread 2: `lw r1, 100`
- Thread 2: `inc r1`
- Thread 2: `sw 100, r1`
- **<Context Switch>**

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```

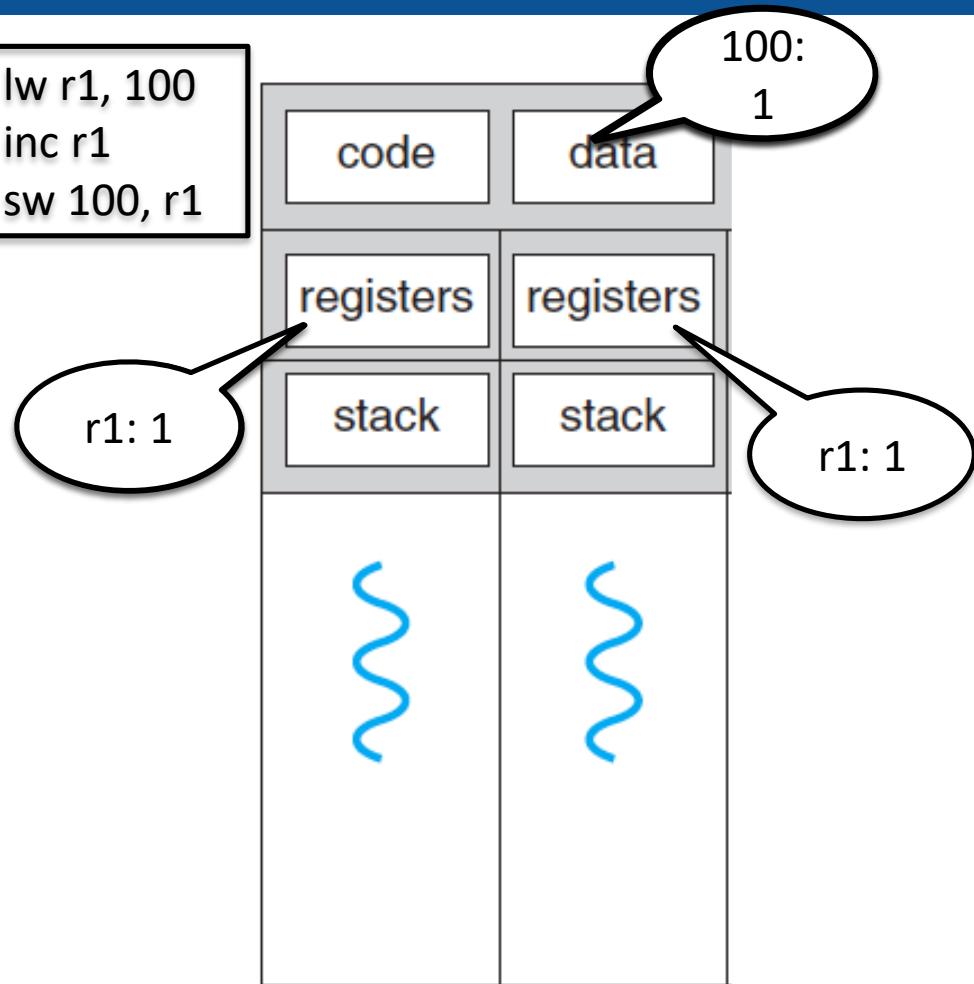


- Thread 1: lw r1,100
  - Thread 1: inc r1
  - Thread 1: sw 100,r1
  - <Context Switch>
  - Thread 2: lw r1, 100
  - Thread 2: inc r1
  - Thread 2: sw 100, r1
- IN = 2

- Thread 1: lw r1,100
- Thread 1: inc r1
- <Context Switch>
- Thread 2: lw r1, 100
- Thread 2: inc r1
- Thread 2: sw 100, r1
- <Context Switch>
- Thread 1: sw 100,r1

# Further Motivation

```
lw r1, 100  
inc r1  
sw 100, r1
```



- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - Thread 1: `sw 100, r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
- IN = 2

- Thread 1: `lw r1, 100`
  - Thread 1: `inc r1`
  - **<Context Switch>**
  - Thread 2: `lw r1, 100`
  - Thread 2: `inc r1`
  - Thread 2: `sw 100, r1`
  - **<Context Switch>**
  - Thread 1: `sw 100, r1`
- IN = 1

# “Too Much Milk”...

Time	You	Your Roommate
3:00	Arrive Home	
3:05	Look in fridge; no milk	
3:10	Leave for grocery	
3:15		Arrive home
3:20	Arrive at grocery	Look in fridge; no milk
3:25	Buy milk	Leave for grocery
3:35	Arrive home; put milk in fridge	
3:45		Buy milk
3:50		Arrive home; oh no!

# “Too Much Milk”...

Time	You	Your Roommate
3:00	Arrive Home	
3:05	Look in fridge; no milk	
3:10	Leave for grocery	
3:15		Arrive home
3:20	Arrive at grocery	Look in fridge; no milk
3:25	Buy milk	Leave for grocery
3:35	Arrive home; put milk in fridge	
3:45		Buy milk
3:50		Arrive home; oh no!

Hi Roommate,  
I went to buy  
milk from the  
grocery

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource  
(data structure or device)

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource  
(data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource (data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

**Where:** Critical section is a piece of code that accesses the shared resource.

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource (data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

**Where:** Critical section is a piece of code that accesses the shared resource.

- Can be one instruction or very large code

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource (data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

**Where:** Critical section is a piece of code that accesses the shared resource.

- Can be one instruction or very large code
- Correctness: must not be concurrently executed by more than one thread process.

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource (data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

**Where:** Critical section is a piece of code that accesses the shared resource.

- Can be one instruction or very large code
- Correctness: must not be concurrently executed by more than one thread process.

**Solution:** Mutual exclusion algorithms are used to avoid the simultaneous execution of a critical section

# What's happening here?

**The Problem:** Uncoordinated access to a shared resource (data structure or device)

- By multiple threads, by multiple processes, by multiple cores/processors

**Where:** Critical section is a piece of code that accesses the shared resource.

- Can be one instruction or very large code
- Correctness: must not be concurrently executed by more than one thread process.

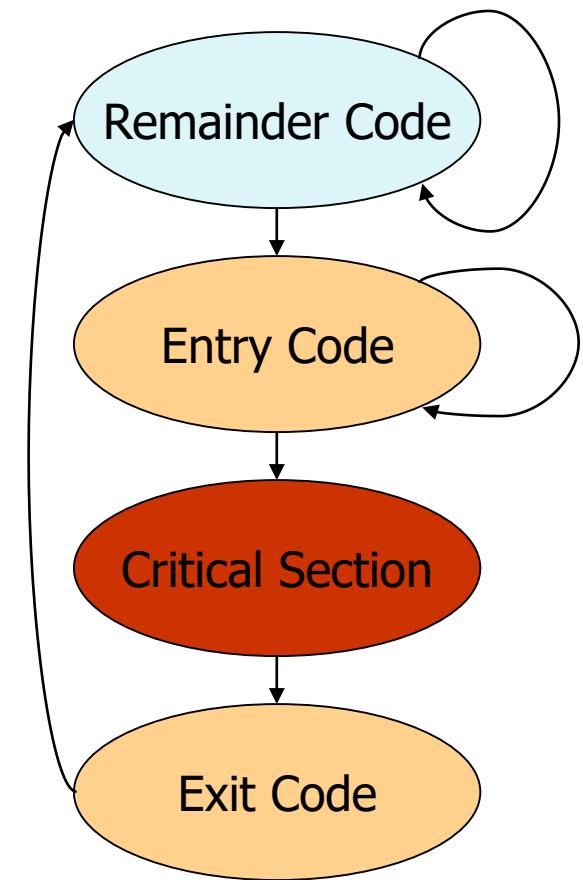
**Solution:** Mutual exclusion algorithms are used to avoid the simultaneous execution of a critical section

**Important:** Sometimes it is not enough to ensure mutual exclusion (e.g., when the order of execution matters)

- More on this later

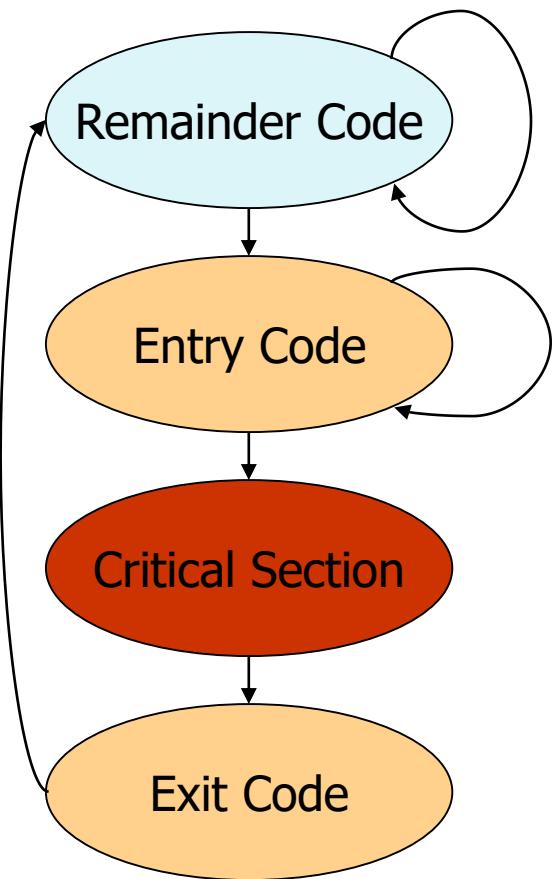
# Modeling the Mutual Exclusion Problem

[First modeling by Dijkstra, 1965]



# Modeling the Mutual Exclusion Problem

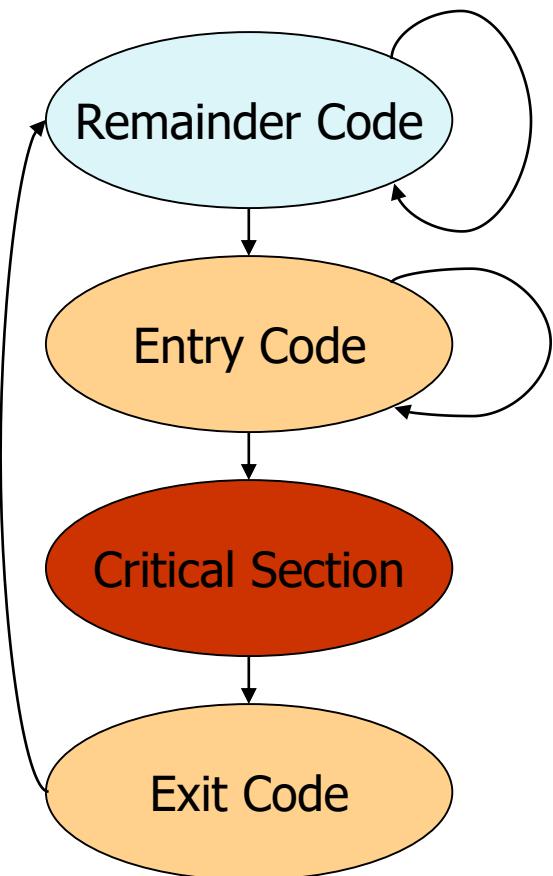
[First modeling by Dijkstra, 1965]



The Problem: Write entry and exit code. We will examine the code against the following properties (the goal is to ensure all of them)

# Modeling the Mutual Exclusion Problem

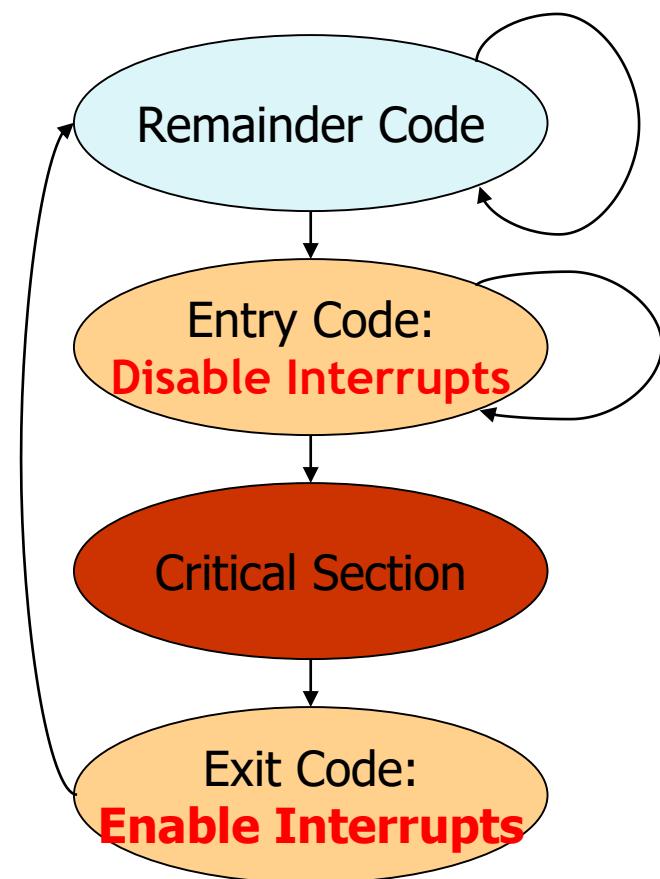
[First modeling by Dijkstra, 1965]



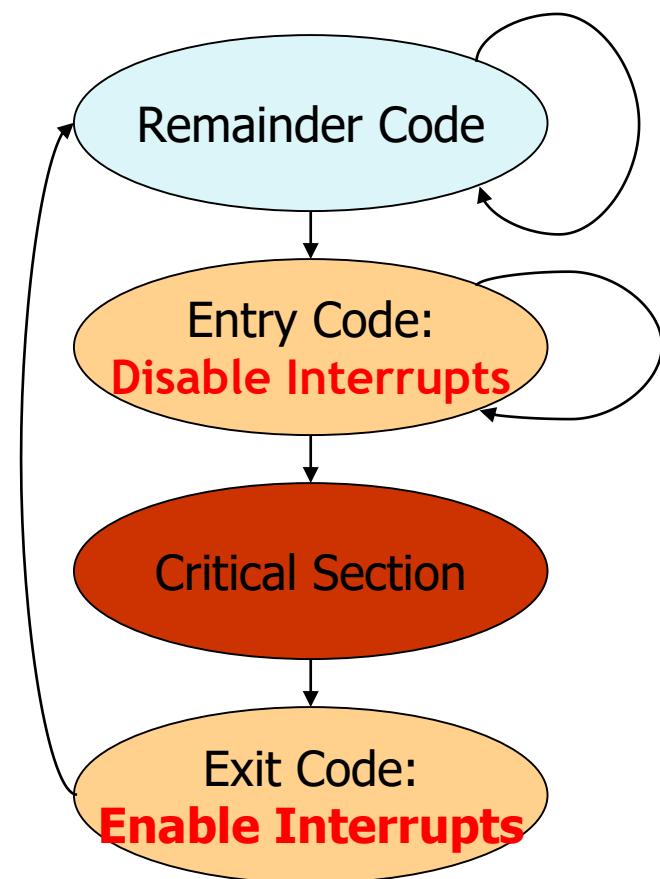
The Problem: Write entry and exit code. We will examine the code against the following properties (the goal is to ensure all of them)

1. **Mutual Exclusion:** No two processes are in their critical section at the same time.
2. **Progress:** If a process is trying to enter the critical section then **some** process must eventually enter the critical section.
3. **Starvation freedom:** If a process is trying to enter the critical section then **this** process must eventually enter the critical section.
4. **Generality:** no assumptions may be made about speeds or the number of participants (threads/processes/cores/CPUs)
5. **No blocking in the remainder** No process running outside its critical section may block other process

# First Try: Disabling interrupts

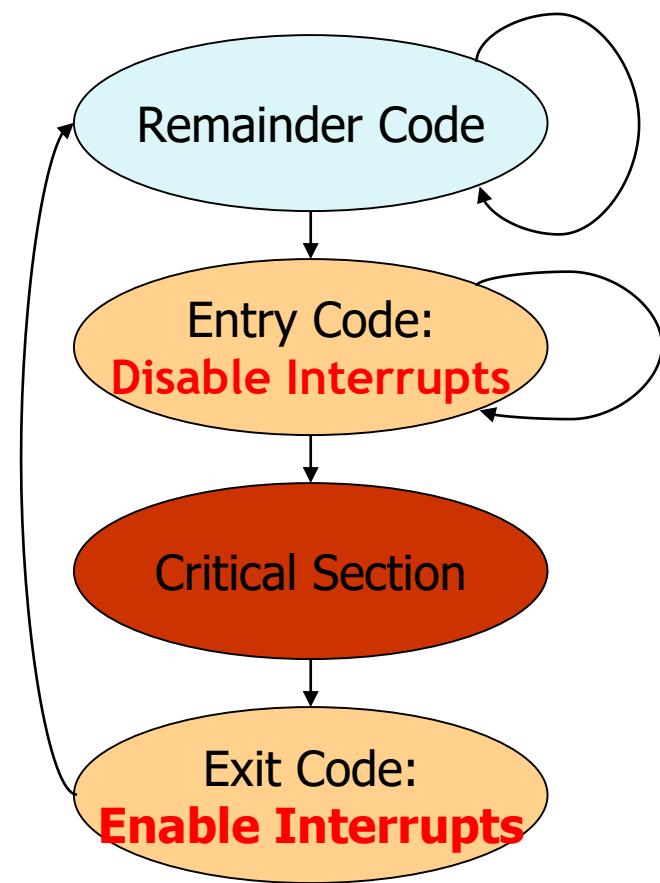


# First Try: Disabling interrupts



Satisfies Properties 1,2,3,5 but:

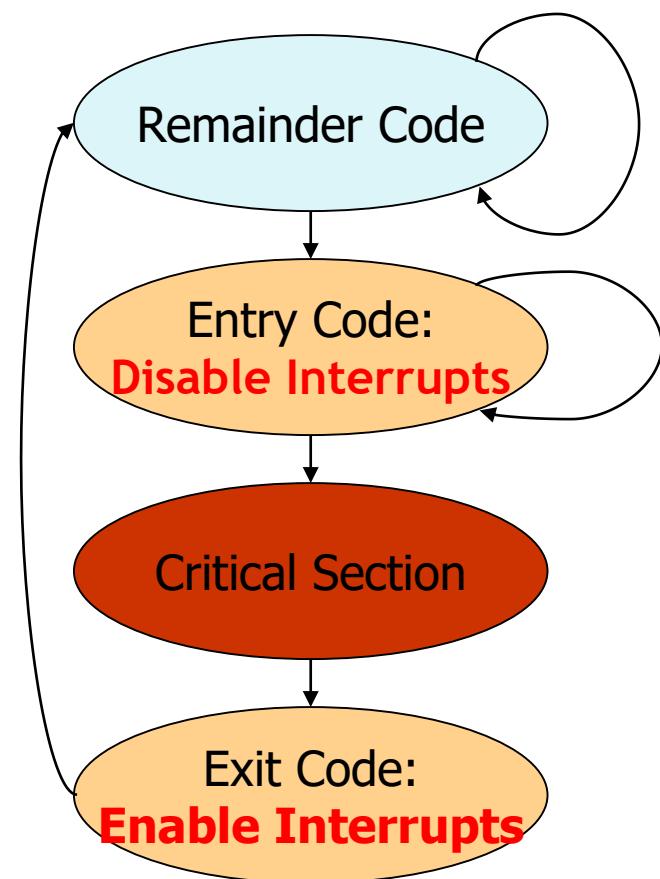
# First Try: Disabling interrupts



Satisfies Properties 1,2,3,5 but:

1. Does not solve the problem in a multi-processor system (where real parallelism exists) → **Property 4 is violated**

# First Try: Disabling interrupts



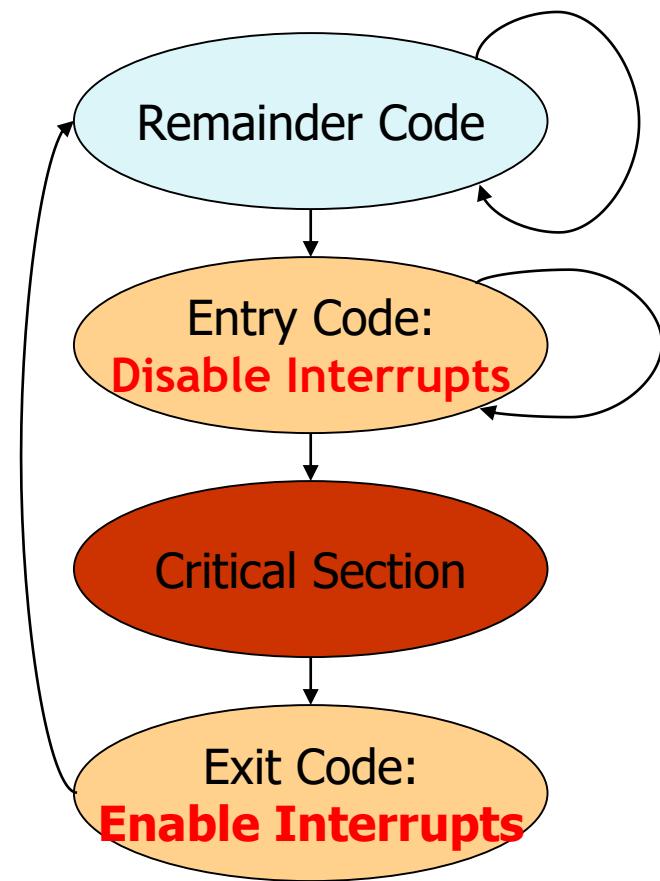
Satisfies Properties 1,2,3,5 but:

1. Does not solve the problem in a multi-processor system (where real parallelism exists) → **Property 4 is violated**

And:

2. User processes should not be allowed to disable interrupts

# First Try: Disabling interrupts



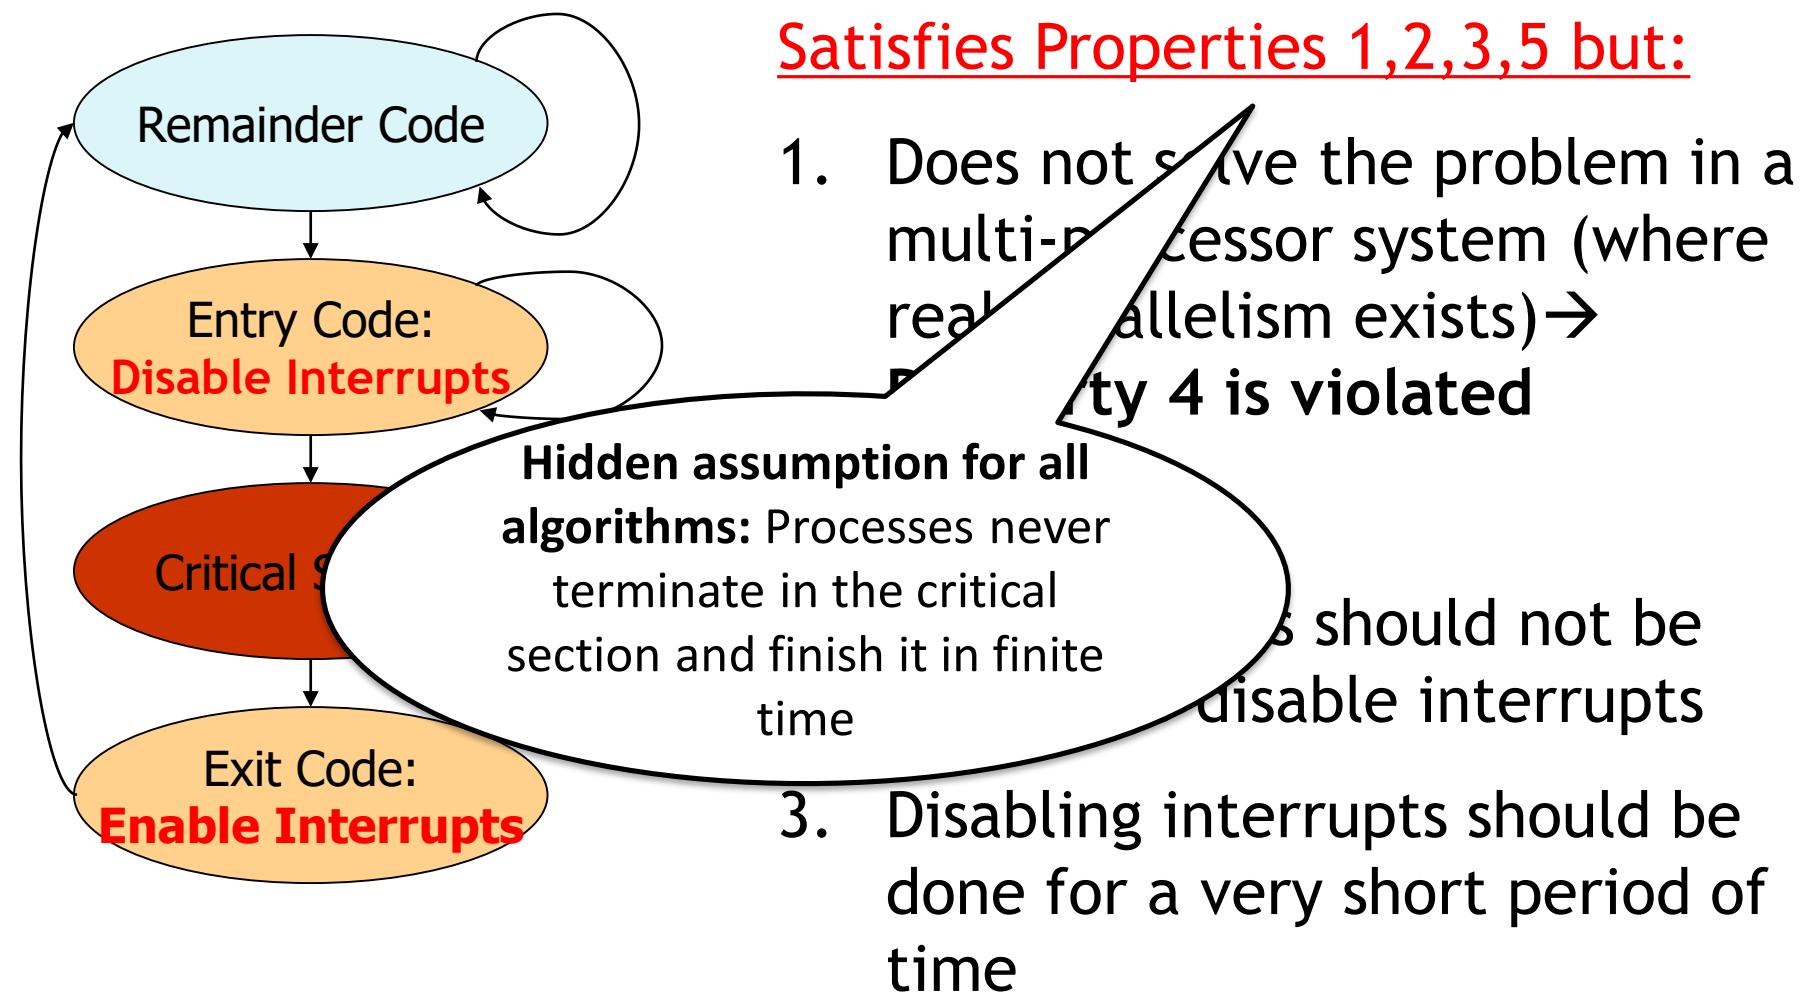
Satisfies Properties 1,2,3,5 but:

1. Does not solve the problem in a multi-processor system (where real parallelism exists) → **Property 4 is violated**

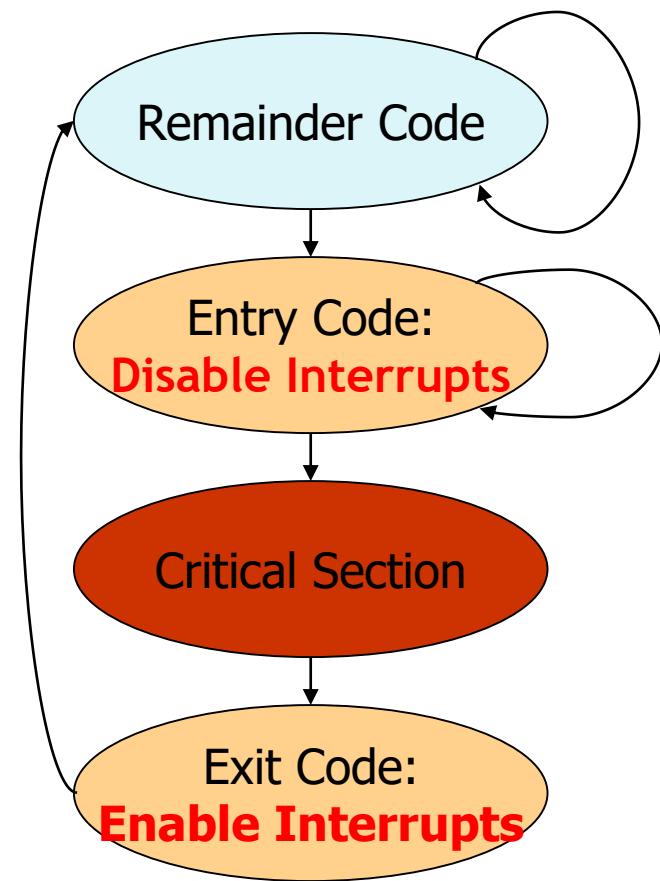
And:

2. User processes should not be allowed to disable interrupts
3. Disabling interrupts should be done for a very short period of time

# First Try: Disabling interrupts



# First Try: Disabling interrupts



Satisfies Properties 1,2,3,5 but:

1. Does not solve the problem in a multi-processor system (where real parallelism exists) → **Property 4 is violated**

And:

2. User processes should not be allowed to disable interrupts
3. Disabling interrupts should be done for a very short period of time

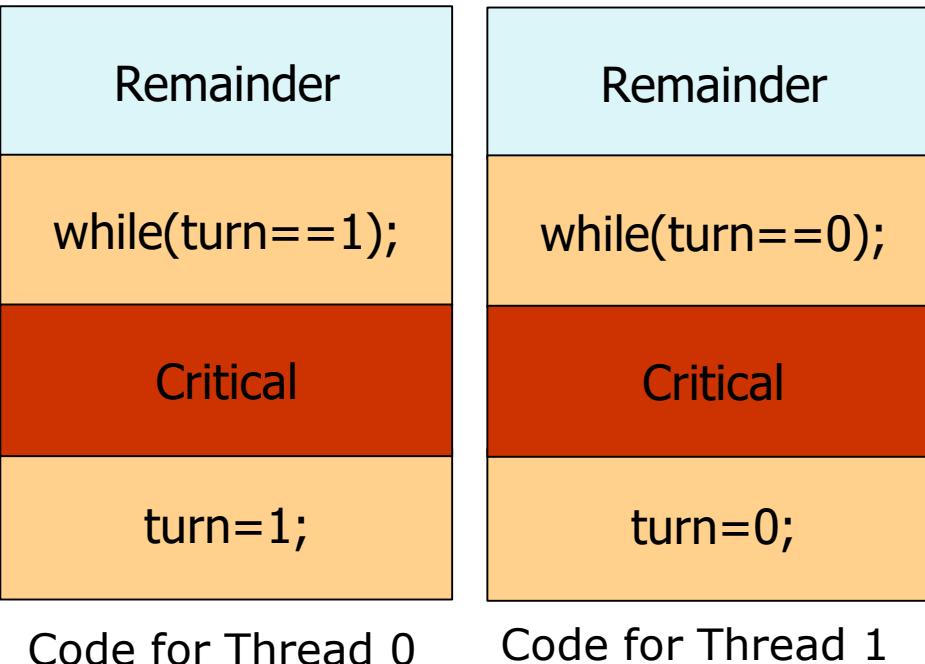
# No “Cheating” – First Try

Remainder	Remainder
while(turn==1);	while(turn==0);
Critical	Critical
turn=1;	turn=0;

Code for Thread 0      Code for Thread 1

- Algorithm for 2 threads/processes
- Uses global variable (1 bit) named **turn**. Cannot be used in the remainder/critical
- (Slightly) different code for each thread

# No “Cheating” – First Try



✓ **Mutual Exclusion:** Assume that the two processes are in the critical condition, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point,  $\text{turn} \neq 1 \rightarrow \text{turn} = 0$ . The only place turn can be changed to 1 is when thread 0 exits the critical section, thus  $\text{turn} = 0$  throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – First Try

Remainder	Remainder
while(turn==1);	while(turn==0);
Critical	Critical
turn=1;	turn=0;

Code for Thread 0      Code for Thread 1

Bad News:

**Progress / no blocking in the remainder** does not hold.

Thread 1 can simply stay in its remainder and never get to its entry (and therefore exit) code  
→ **Starvation Freedom** does not hold

✓ **Mutual Exclusion:** Assume that the two of the processes are in the critical condition, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point,  $\text{turn} \neq 1 \rightarrow \text{turn} = 0$ . The only place turn can be changed to 1 is when thread 0 exits the critical section, thus  $\text{turn} = 0$  throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

Remainder
flag[0] = true; while(flag[1]);
Critical
flag[0]=false;

Code for Thread 0

Remainder
flag[1]=true; while(flag[0]);
Critical
flag[1]=false

Code for Thread 1

Algorithm for 2 threads/processes

Uses global two global variables  
(1 bit), named **flag[0]**,  
**flag[1]**. Cannot be used in the  
remainder/critical section

(Slightly) different code for each  
thread

# No “Cheating” – Second Try

Remainder	Remainder
<pre>flag[0] = true; while(flag[1]);</pre>	<pre>flag[1]=true; while(flag[0]);</pre>
Critical	Critical
<pre>flag[0]=false;</pre>	<pre>flag[1]=false</pre>

Code for Thread 0      Code for Thread 1

# No “Cheating” – Second Try

Remainder	Remainder
<code>flag[0] = true; while(flag[1]);</code>	<code>flag[1]=true; while(flag[0]);</code>
Critical	Critical
<code>flag[0]=false;</code>	<code>flag[1]=false</code>

Code for Thread 0

Code for Thread 1

✓ **Mutual Exclusion:** Assume that the two processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
<code>flag[0] = true; while(flag[1]);</code>	<code>flag[1]=true; while(flag[0]);</code>
Critical	Critical
<code>flag[0]=false;</code>	<code>flag[1]=false</code>

Code for Thread 0

Code for Thread 1

**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

**✓ Mutual Exclusion:** Assume that the two processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
<code>flag[0] = true; while(flag[1]);</code>	<code>flag[1]=true; while(flag[0]);</code>
Critical	Critical
<code>flag[0]=false;</code>	<code>flag[1]=false</code>

Code for Thread 0      Code for Thread 1

**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

But still **Progress** does not hold  
→ **Starvation Freedom** does not hold

**✓ Mutual Exclusion:** Assume that the two of the processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
→ flag[0] = true; while(flag[1]);	flag[1]=true; while(flag[0]);
Critical	Critical
flag[0]=false;	flag[1]=false

Code for Thread 0

Code for Thread 1

**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

But still **Progress** does not hold  
→ **Starvation Freedom** does not hold

**✓ Mutual Exclusion:** Assume that the two of the processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
→ flag[0] = true; while(flag[1]);	→ flag[1]=true; while(flag[0]);
Critical	Critical
flag[0]=false;	flag[1]=false

Code for Thread 0

Code for Thread 1

**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

But still **Progress** does not hold  
→ **Starvation Freedom** does not hold

**✓ Mutual Exclusion:** Assume that the two of the processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
flag[0] = true; → while(flag[1]);	→ flag[1]=true; while(flag[0]);
Critical	Critical
flag[0]=false;	flag[1]=false

Code for Thread 0

Code for Thread 1

**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

But still **Progress** does not hold  
→ **Starvation Freedom** does not hold

**✓ Mutual Exclusion:** Assume that the two of the processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# No “Cheating” – Second Try

Remainder	Remainder
flag[0] = true; → while(flag[1]);	flag[1]=true; → while(flag[0]);
Critical	Critical
flag[0]=false;	flag[1]=false

Code for Thread 0

Code for Thread 1

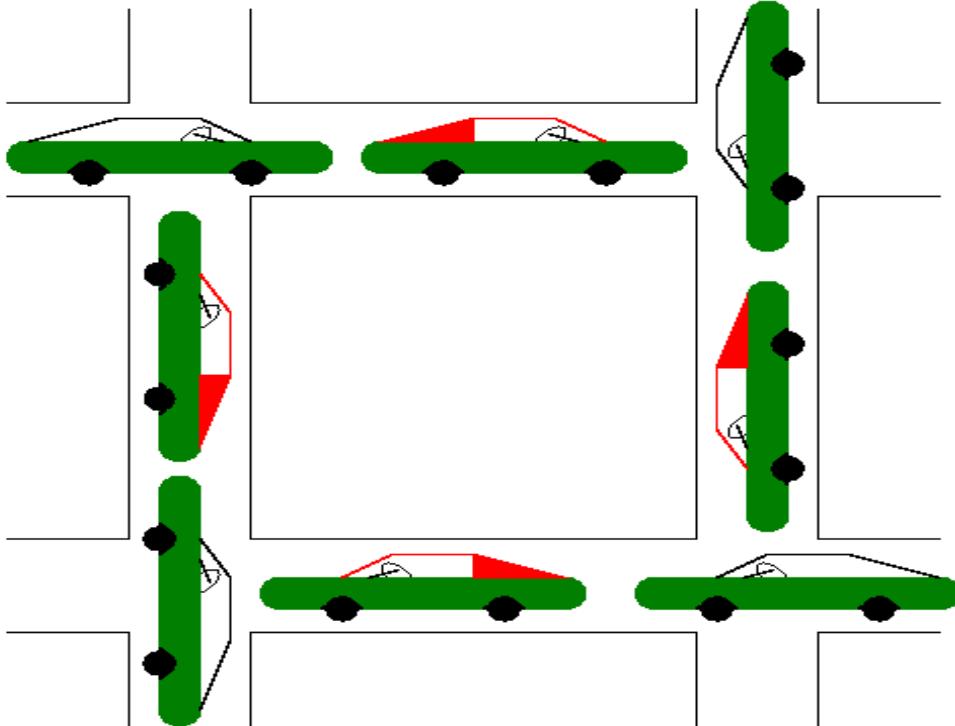
**No Blocking in the Remainder:**  
A thread can block another thread only if its flag is true, implying it is either in its entry, critical, or exit section.

But still **Progress** does not hold  
→ **Starvation Freedom** does not hold

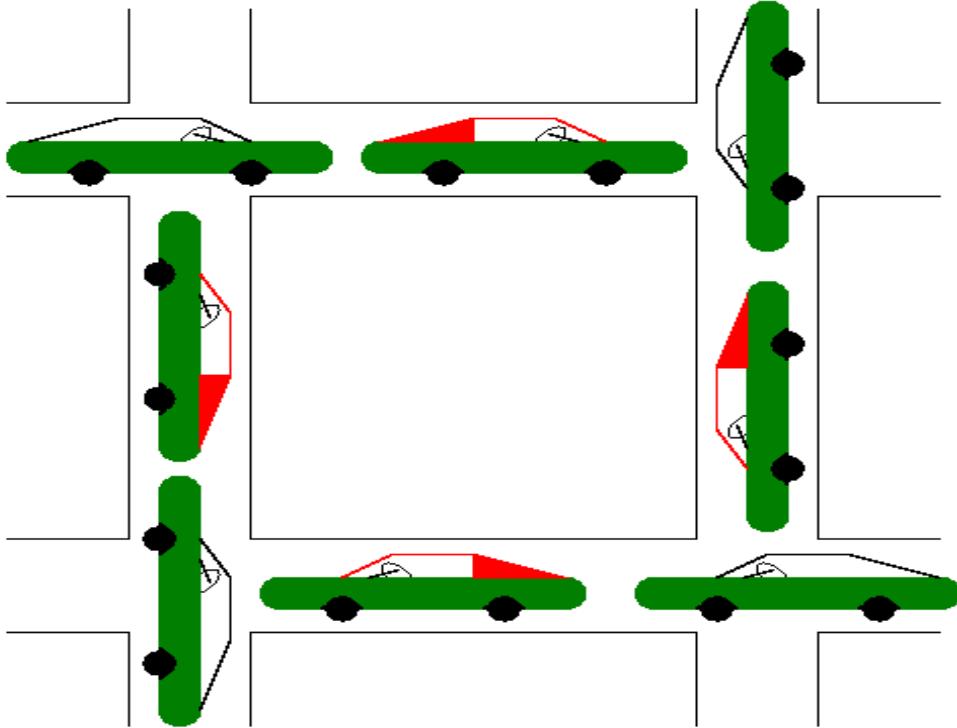
**Deadlock!**

**✓ Mutual Exclusion:** Assume that the two processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[0]=true** → The only place **flag[0]** can be changed to **false** is when thread 0 exits the critical section, thus **flag[0]=true** throughout the entire execution of the critical section by Thread 0 → Thread 1 is in its entry code when Thread 0 is the critical section → Contradiction.

# Illustration of Deadlocks



# Illustration of Deadlocks

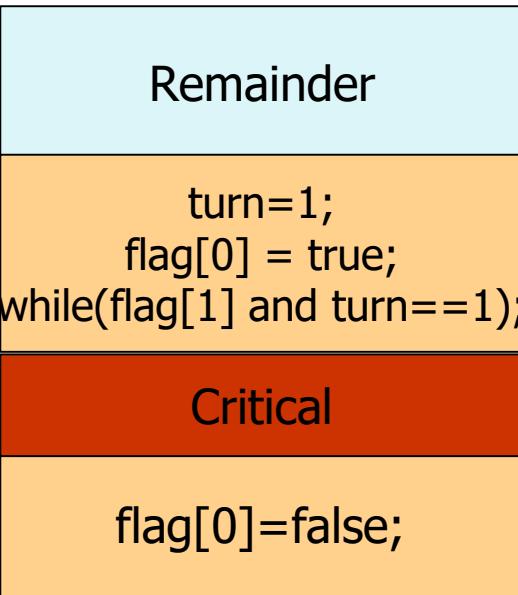


# Illustration of Deadlocks

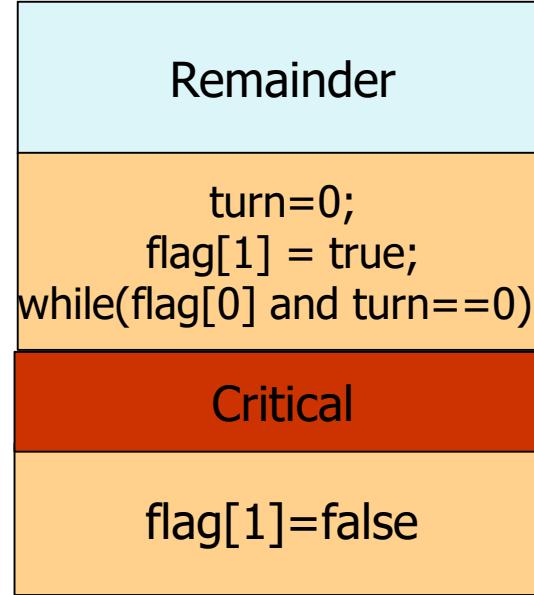


Tel Aviv, Winter 2011

# No “Cheating” – Third Try

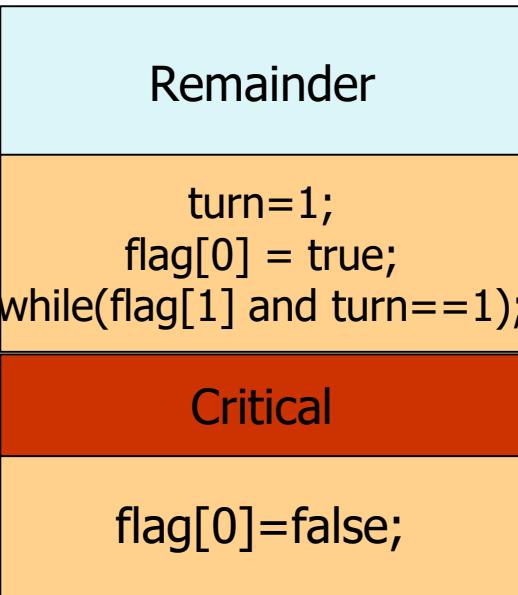


Code for Thread 0

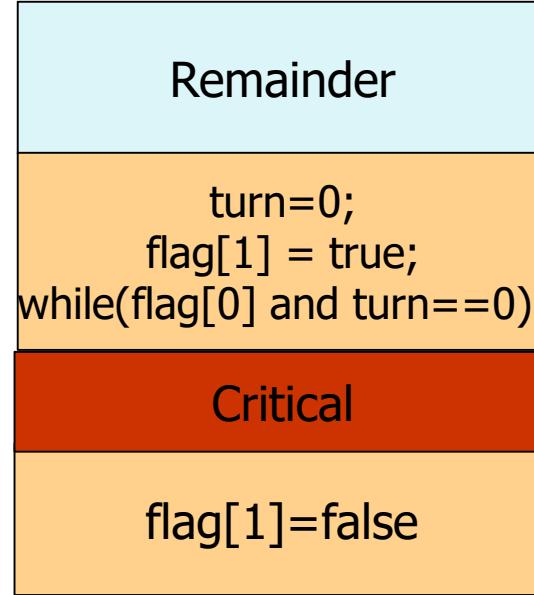


Code for Thread 1

# No “Cheating” – Third Try



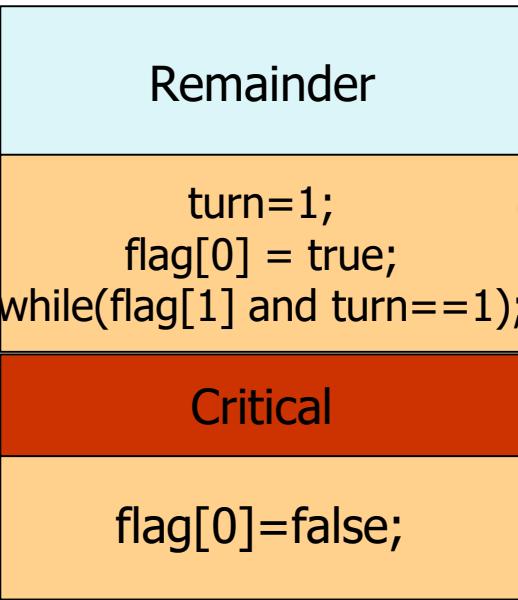
Code for Thread 0



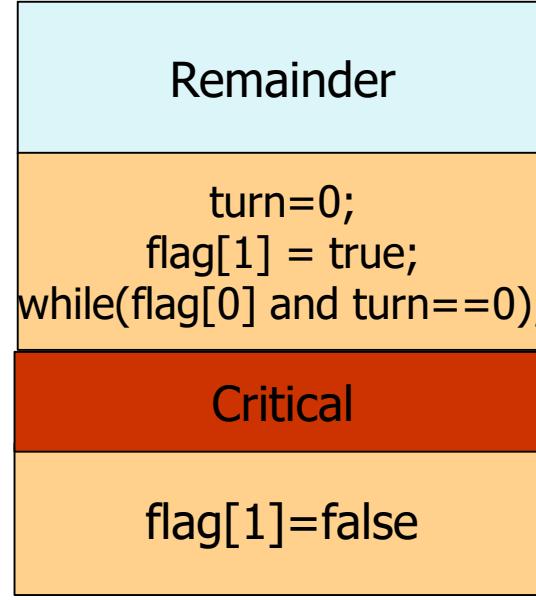
Code for Thread 1

**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



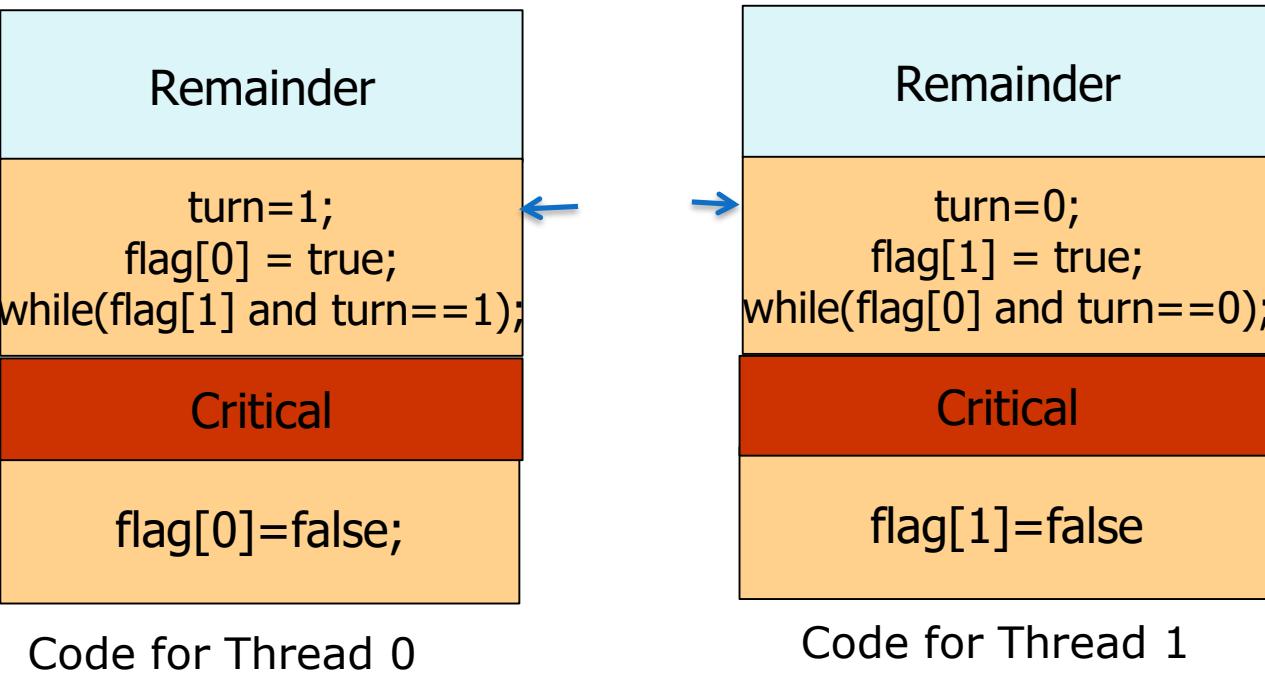
Code for Thread 0



Code for Thread 1

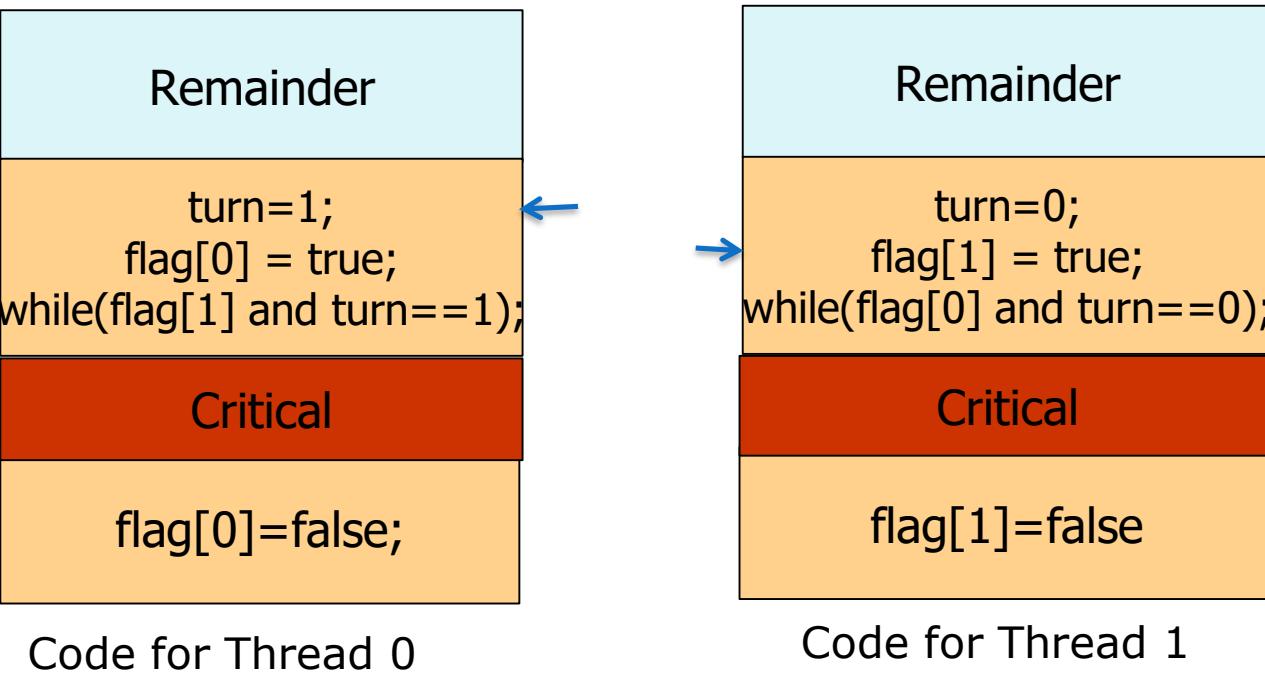
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



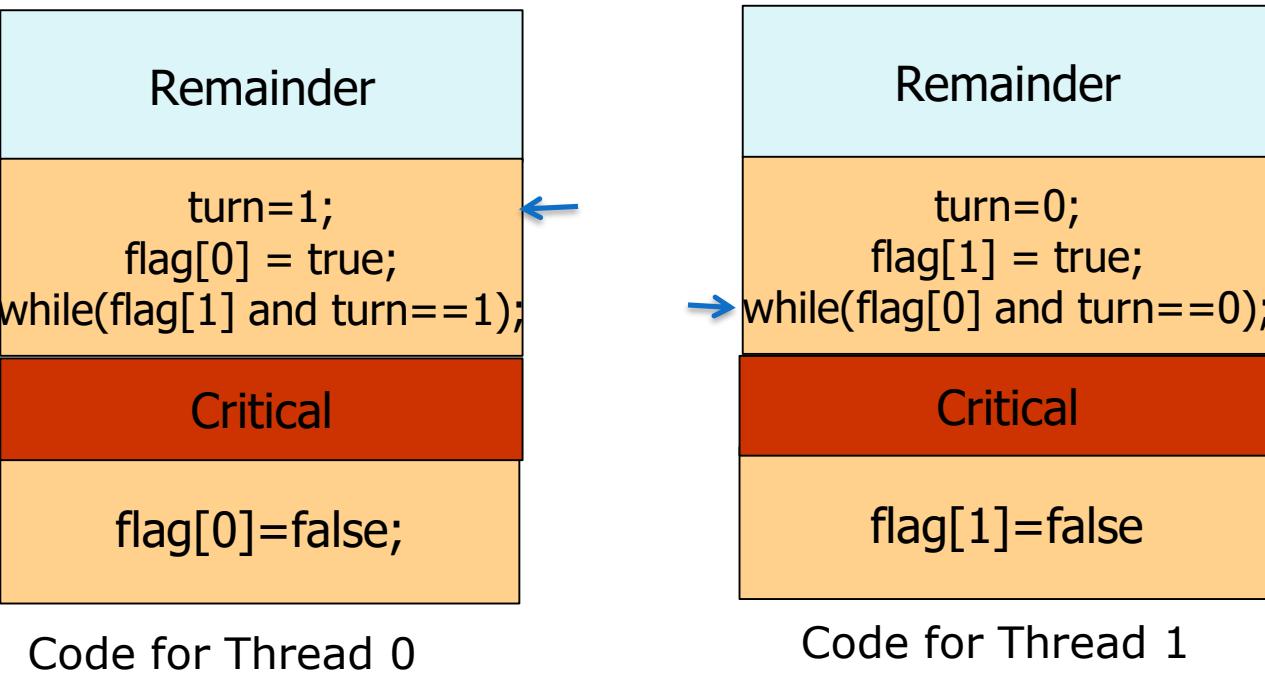
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



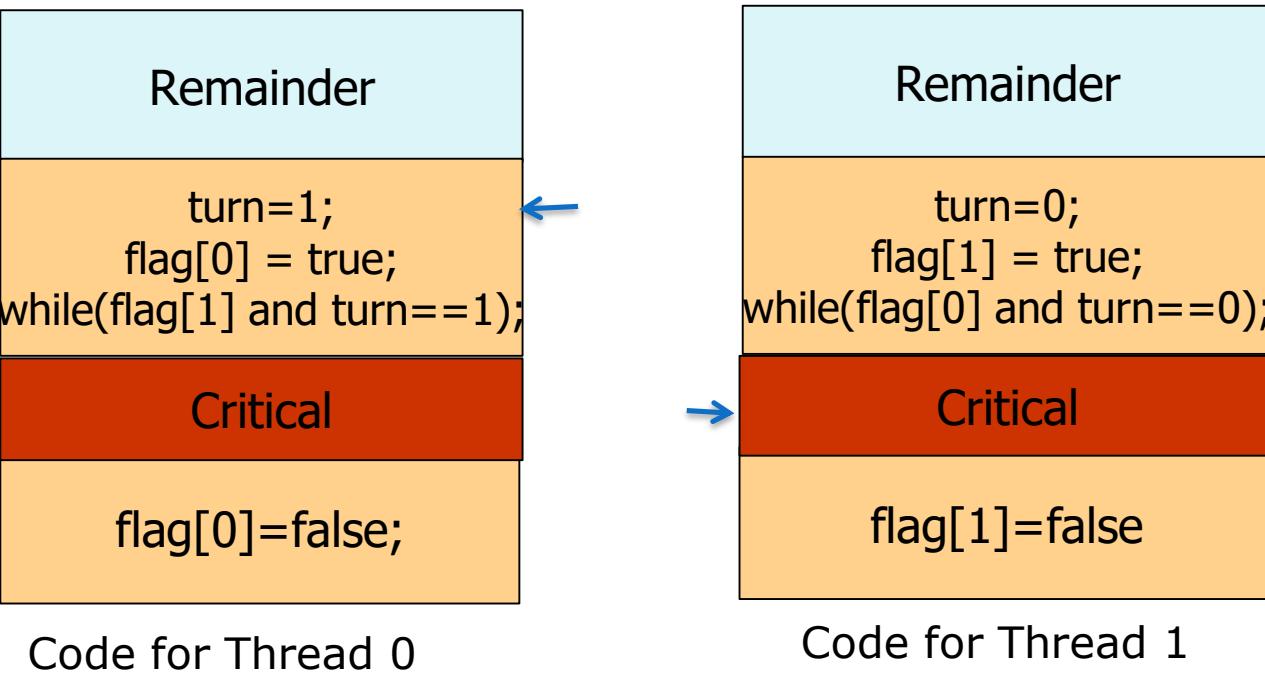
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



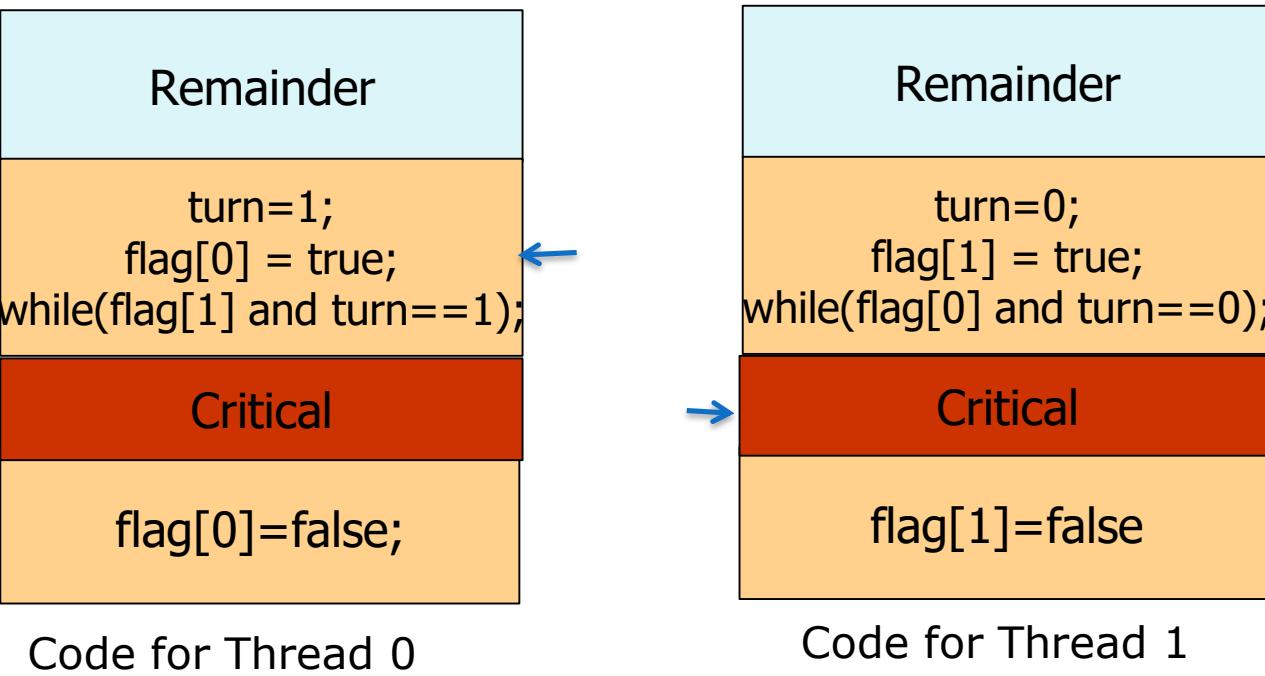
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



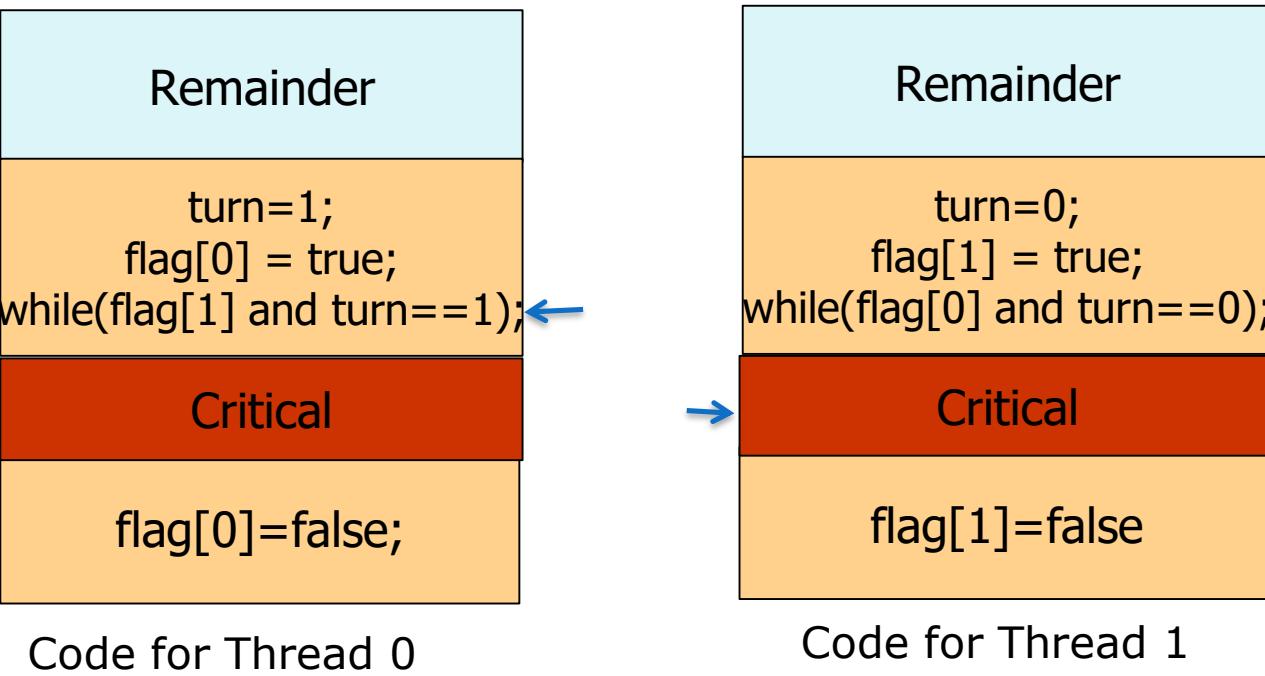
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



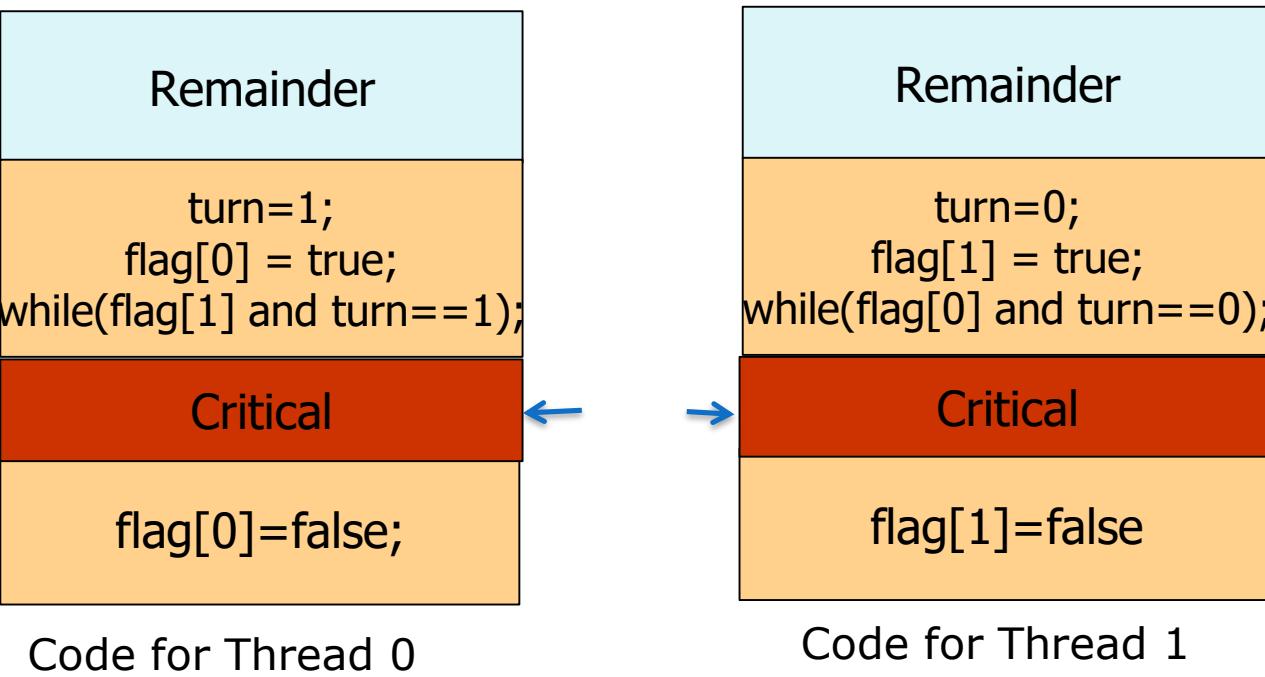
**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



**Mutual Exclusion** does not hold

# No “Cheating” – Third Try



**Mutual Exclusion** does not hold

# Peterson's Algorithm - 1981

Remainder

```
flag[0] = true;  
turn=1;  
while(flag[1] and turn==1);
```

Critical

```
flag[0]=false;
```

Code for Thread 0

Remainder

```
flag[1] = true;  
turn=0;  
while(flag[0] and turn==0);
```

Critical

```
flag[1]=false
```

Code for Thread 1

# Peterson's Algorithm - 1981

Remainder

```
flag[i] = true;  
turn=1-i;  
while(flag[1-i] and turn==1-i);
```

Critical

```
Flag[i]=false;
```

Code for Thread *i*

# Peterson's Algorithm - 1981

Remainder

```
flag[i] = true;  
turn=1-i;  
while(flag[1-i] and turn==1-i);
```

Critical

```
Flag[i]=false;
```

Code for Thread *i*

✓ **Mutual Exclusion:** Assume that the two processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[1]=false or turn=0**.

**Case 1: flag[1]=false** → Thread 1 has not executed Line 1 of it's entry code → When it will eventually get to Line 3, it will set turn to 0, and wait. The only place **turn** is set to 1, is in Thread's 0 entry code, after it left the critical section → Contradiction.

**Case 2: turn=0** → Thread 1 executed Line 2 between Thread 0 Line 2 and 3 →

# Peterson's Algorithm - 1981

Remainder
flag[ <i>i</i> ] = true; turn=1- <i>i</i> ; while(flag[1- <i>i</i> ] and turn==1- <i>i</i> );
Critical
Flag[ <i>i</i> ]=false;

✓ **No blocking in the Remainder**

✓ **Starvation Free:** Assume Thread 1 tries to get into the critical section but blocked (on Line 3) for infinite time → Thread 0 enters the critical section infinite number of times → The second time Thread 0 will execute the entry, **flag[0]=true, turn=1** and it will block → The next time Thread 1 gets the CPU, the **while** condition does not hold and Thread 1 enters the critical section → Contradiction.

Code for Thread *i* —

✓ **Mutual Exclusion:** Assume that the two processes are in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[1]=false or turn=0**.

**Case 1: flag[1]=false** → Thread 1 has not executed Line 1 of its entry code → When it will eventually get to Line 3, it will set turn to 0, and wait. The only place **turn** is set to 1, is in Thread's 0 entry code, after it left the critical section → Contradiction.

**Case 2: turn=0** → Thread 1 executed Line 2 between Thread 0 Line 2 and 3 →

# Peterson's Algorithm - 1981

Remainder
flag[ <i>i</i> ] = true; turn=1- <i>i</i> ; while(flag[1- <i>i</i> ] and turn==1- <i>i</i> );
Critical
Flag[ <i>i</i> ]=false;

Code for Thread *i* —

✓ **No blocking in the Remainder**

✓ **Starvation Free:** Assume Thread 1 tries to get into the critical section but blocked (on Line 3) for infinite time → Thread 0 enters the critical section infinite number of times → The second time Thread 0 will execute the entry, **flag[0]=true, turn=1** and it will block → The next time Thread 1 gets the CPU, the **while**

**Two threads only**

cannot be easily extended  
to more

**Busy wait (spinning)**

Thread 1 enters the  
on.

✓ **Mutual Exclusion:** A thread can be in the critical section, and then w.l.o.g. thread 0 left its entry section (and entered the critical section) first → At that point, **flag[1]=false or turn=0**.

**Case 1: flag[1]=false** → Thread 1 has not executed Line 1 of its entry code → When it will eventually get to Line 3, it will set turn to 0, and wait. The only place **turn** is set to 1, is in Thread 0's entry code, after it left the critical section → Contradiction.

**Case 2: turn=0** → Thread 1 executed Line 2 between Thread 0 Line 2 and 3 →

# Three Concepts We've Covered

- **Race Condition:** A situation in which the scheduling of processes/threads (the order of who is getting the CPU) changes the final result
  - Different order → different results
  - Typical scenario: Most of the time everything is OK, but sometime not → very hard to debug
- **Atomic Instruction:** instruction that completes in a single step relative to other threads
  - $x=x+1$  was not atomic
  - Read (lw) and write (sw) were atomic
  - Sometimes this doesn't hold (wide words that spans more than one memory address)

# Three Concepts We've Covered

- **Busy-waiting (spinning, busy-looping):** technique in which a process repeatedly checks to see if a condition is true
  - “do nothing” loops: `while(flag[1-i] and turn==1-i);`

# Three Concepts We've Covered

- **Busy-waiting (spinning, busy-looping):** technique in which a process repeatedly checks to see if a condition is true
  - “do nothing” loops: `while(flag[1-i] and turn==1-i);`
  - Usually, should be avoided as it wastes CPU cycles and yields large overhead

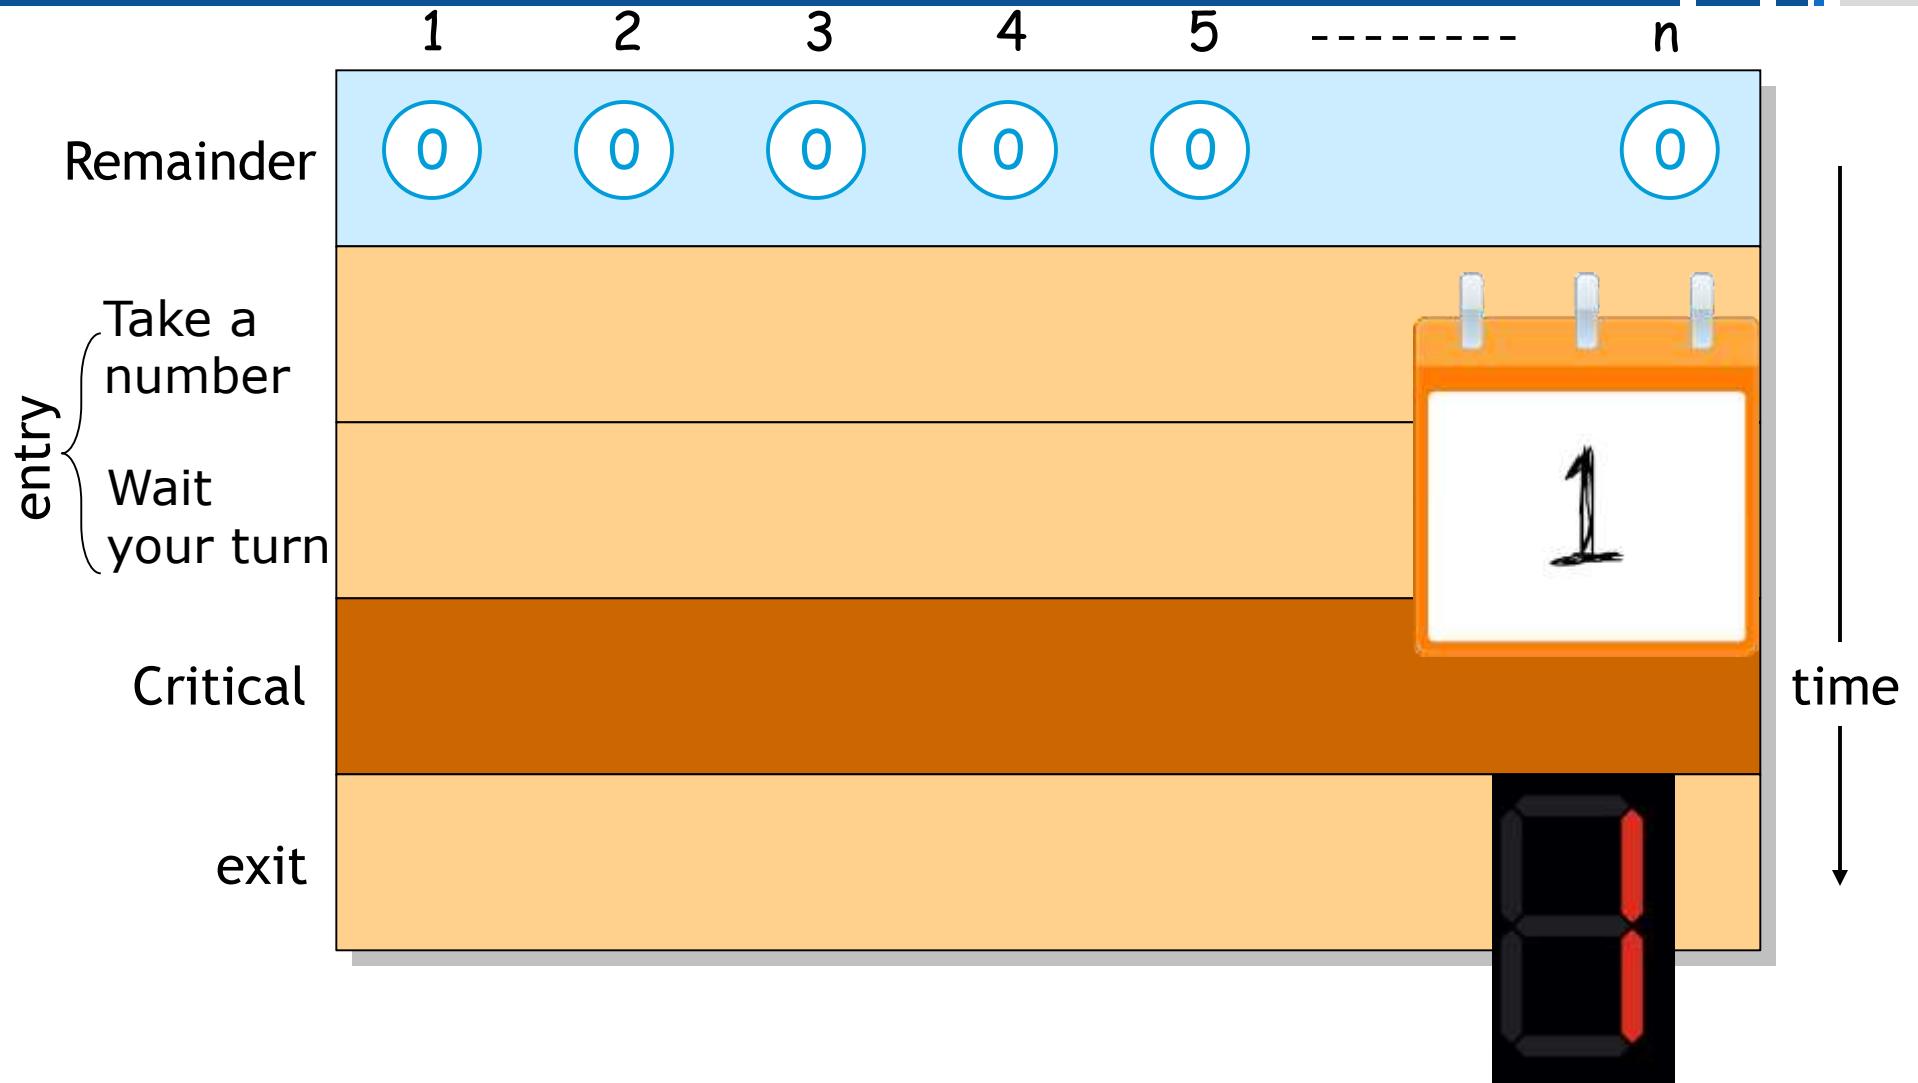
# Three Concepts We've Covered

- **Busy-waiting (spinning, busy-looping):** technique in which a process repeatedly checks to see if a condition is true
  - “do nothing” loops: `while(flag[1-i] and turn==1-i);`
  - Usually, should be avoided as it wastes CPU cycles and yields large overhead
  - Polling vs. interrupts

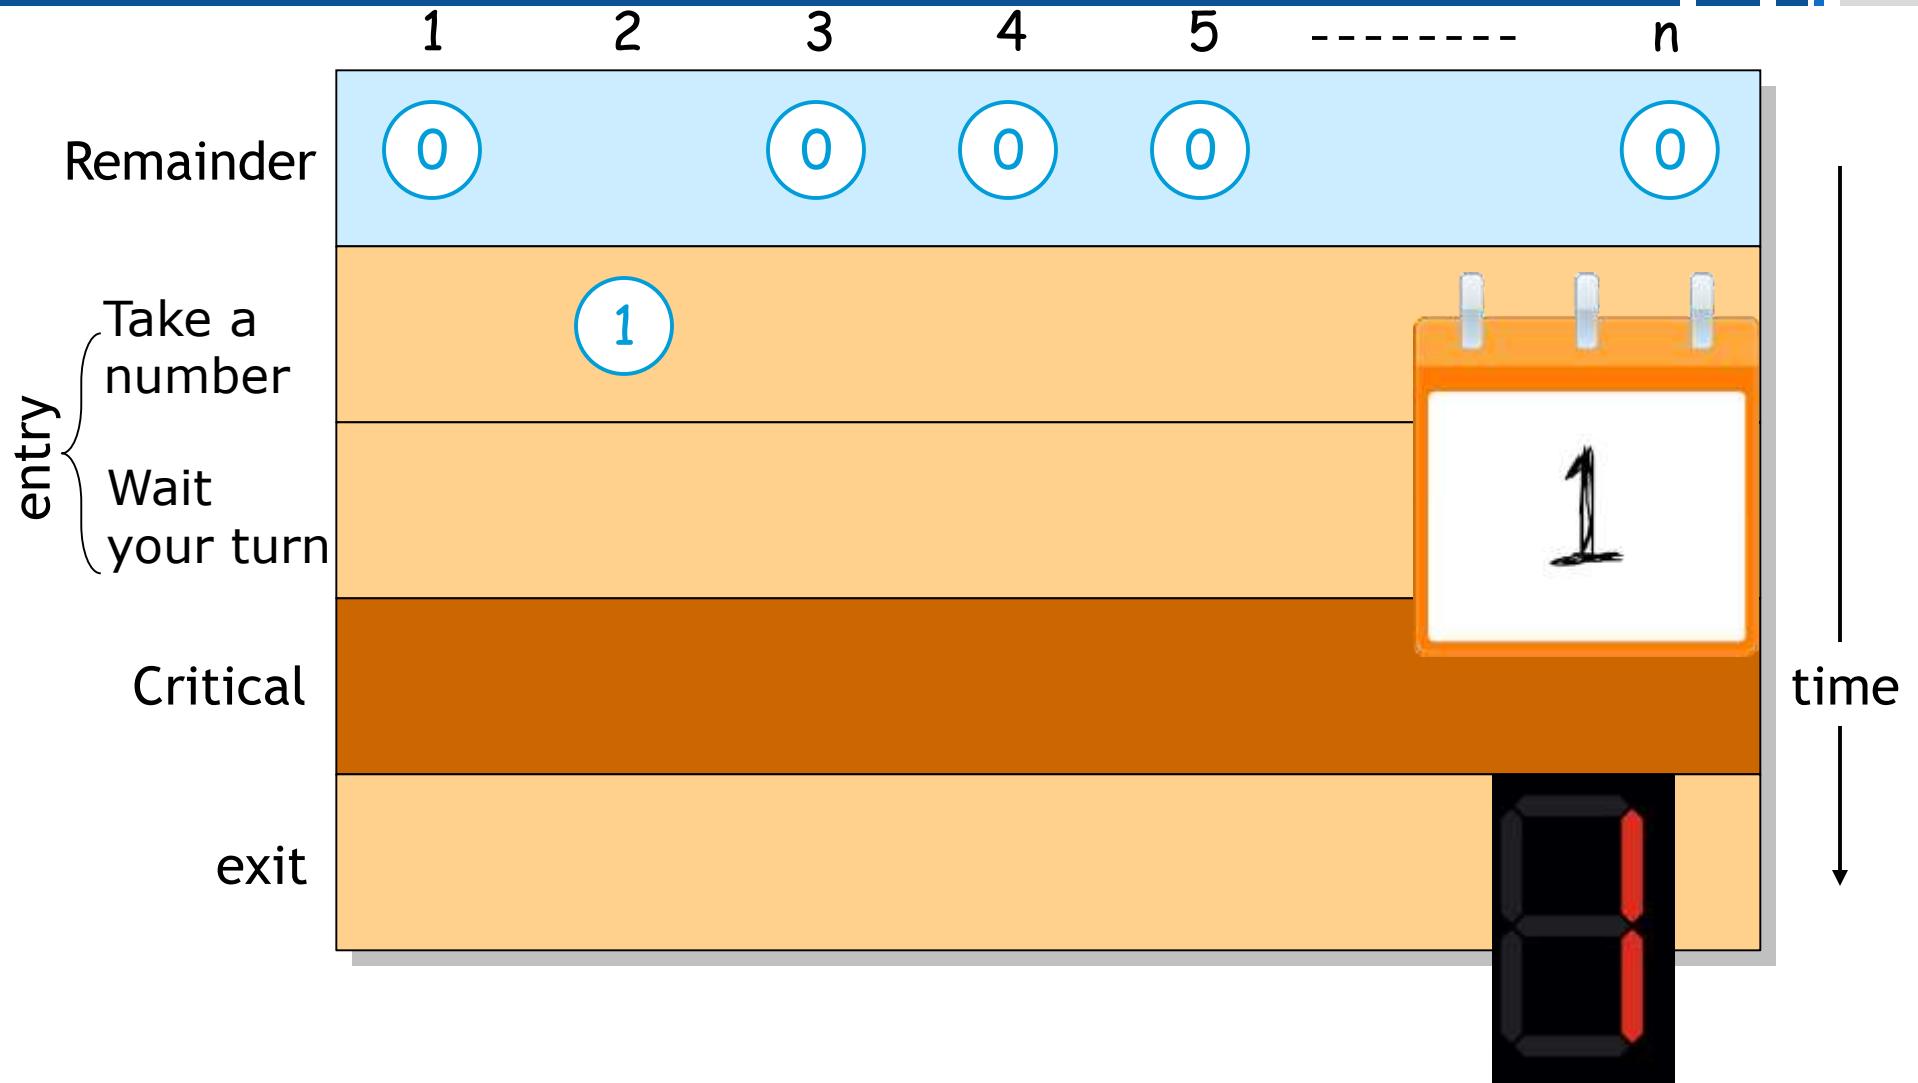
# Three Concepts We've Covered

- **Busy-waiting (spinning, busy-looping):** technique in which a process repeatedly checks to see if a condition is true
  - “do nothing” loops: `while(flag[1-i] and turn==1-i);`
  - Usually, should be avoided as it wastes CPU cycles and yields large overhead
  - Polling vs. interrupts
  - We will discuss alternative later

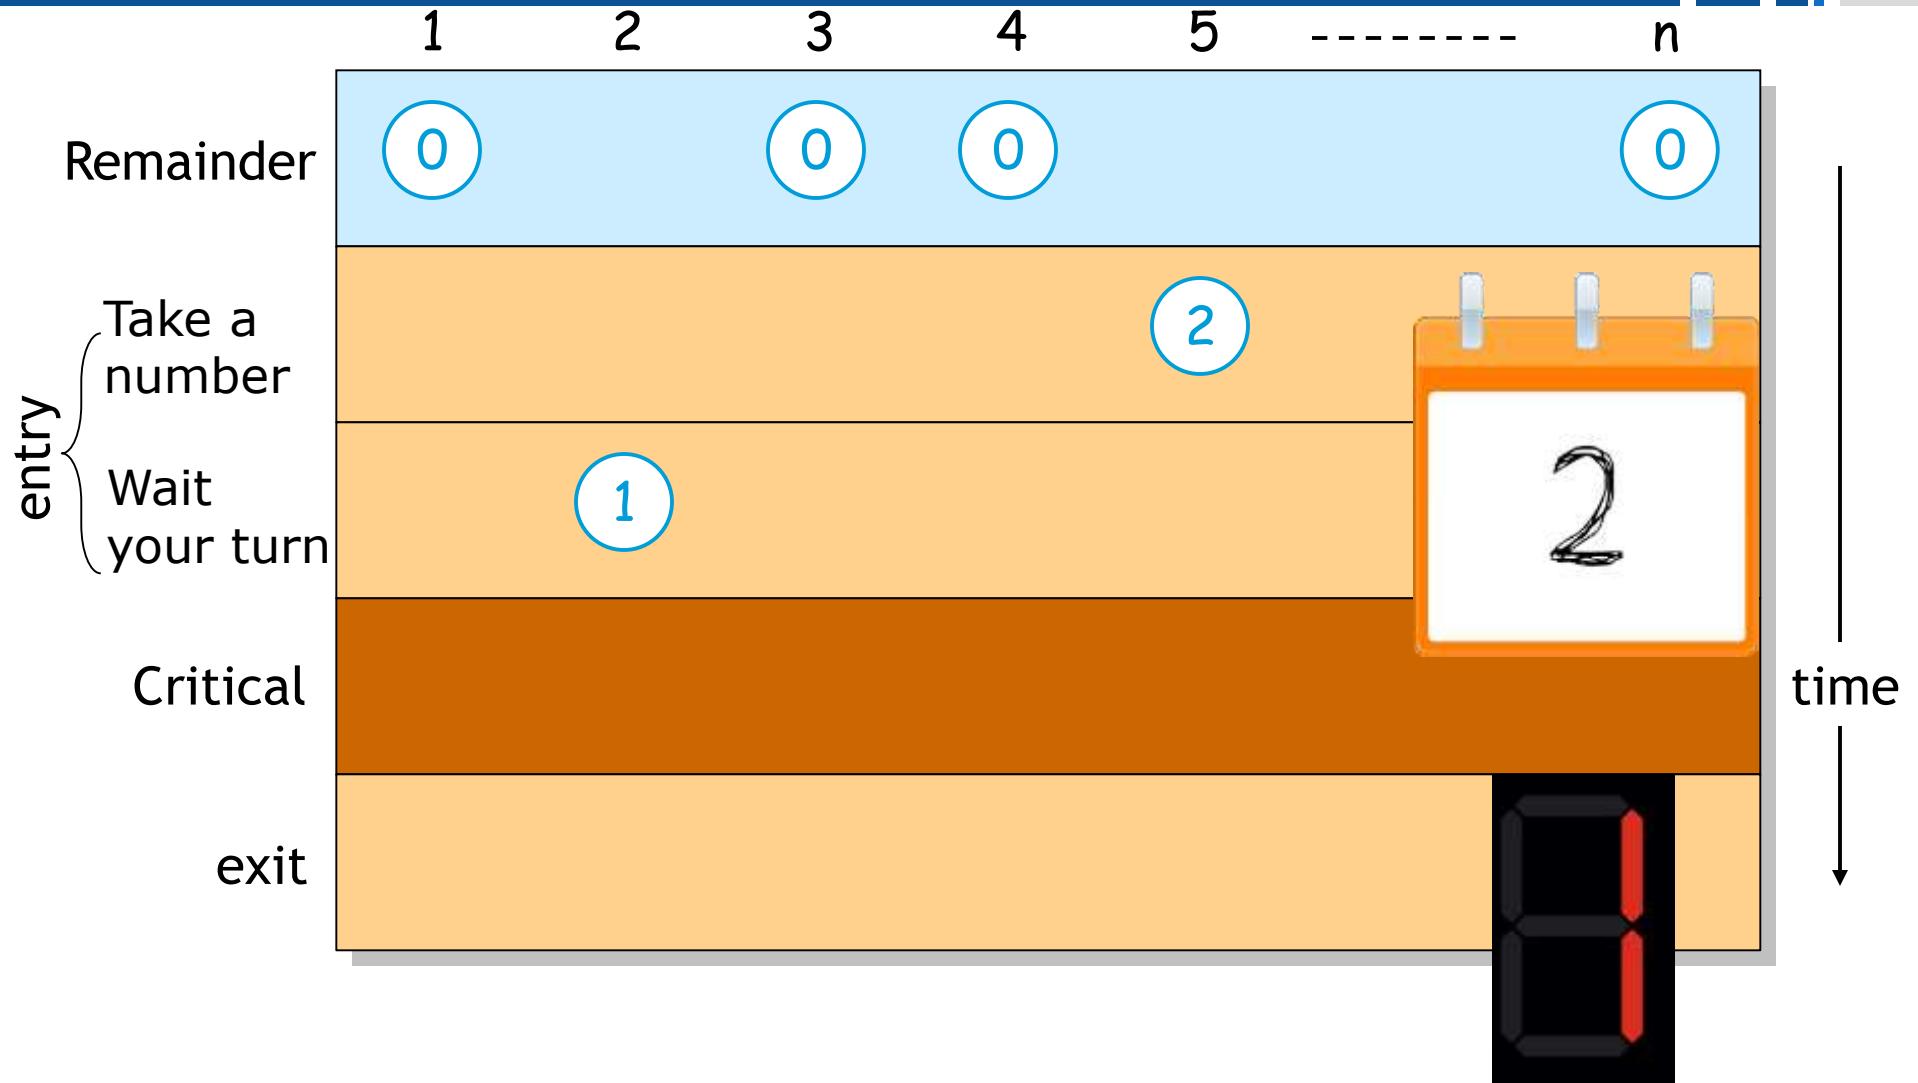
# Lamport's Bakery Algorithm - 1974



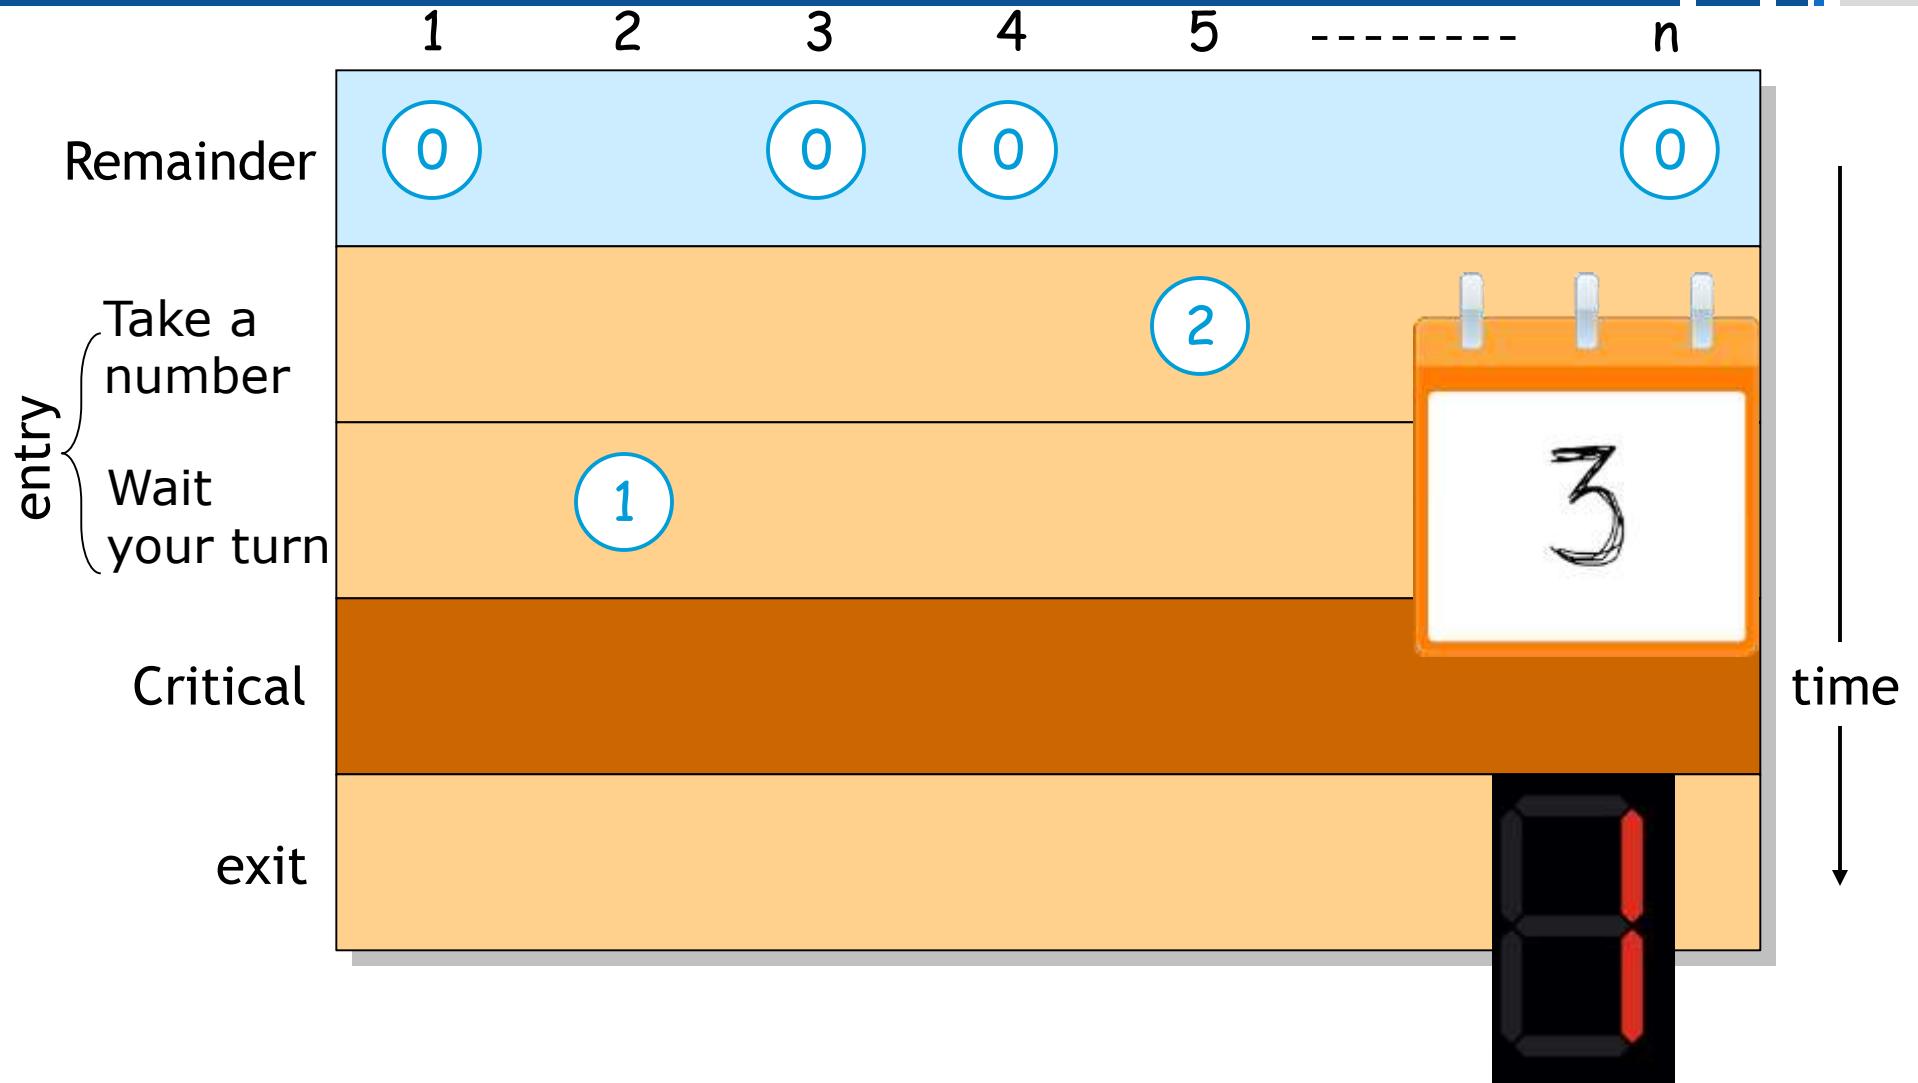
# Lamport's Bakery Algorithm - 1974



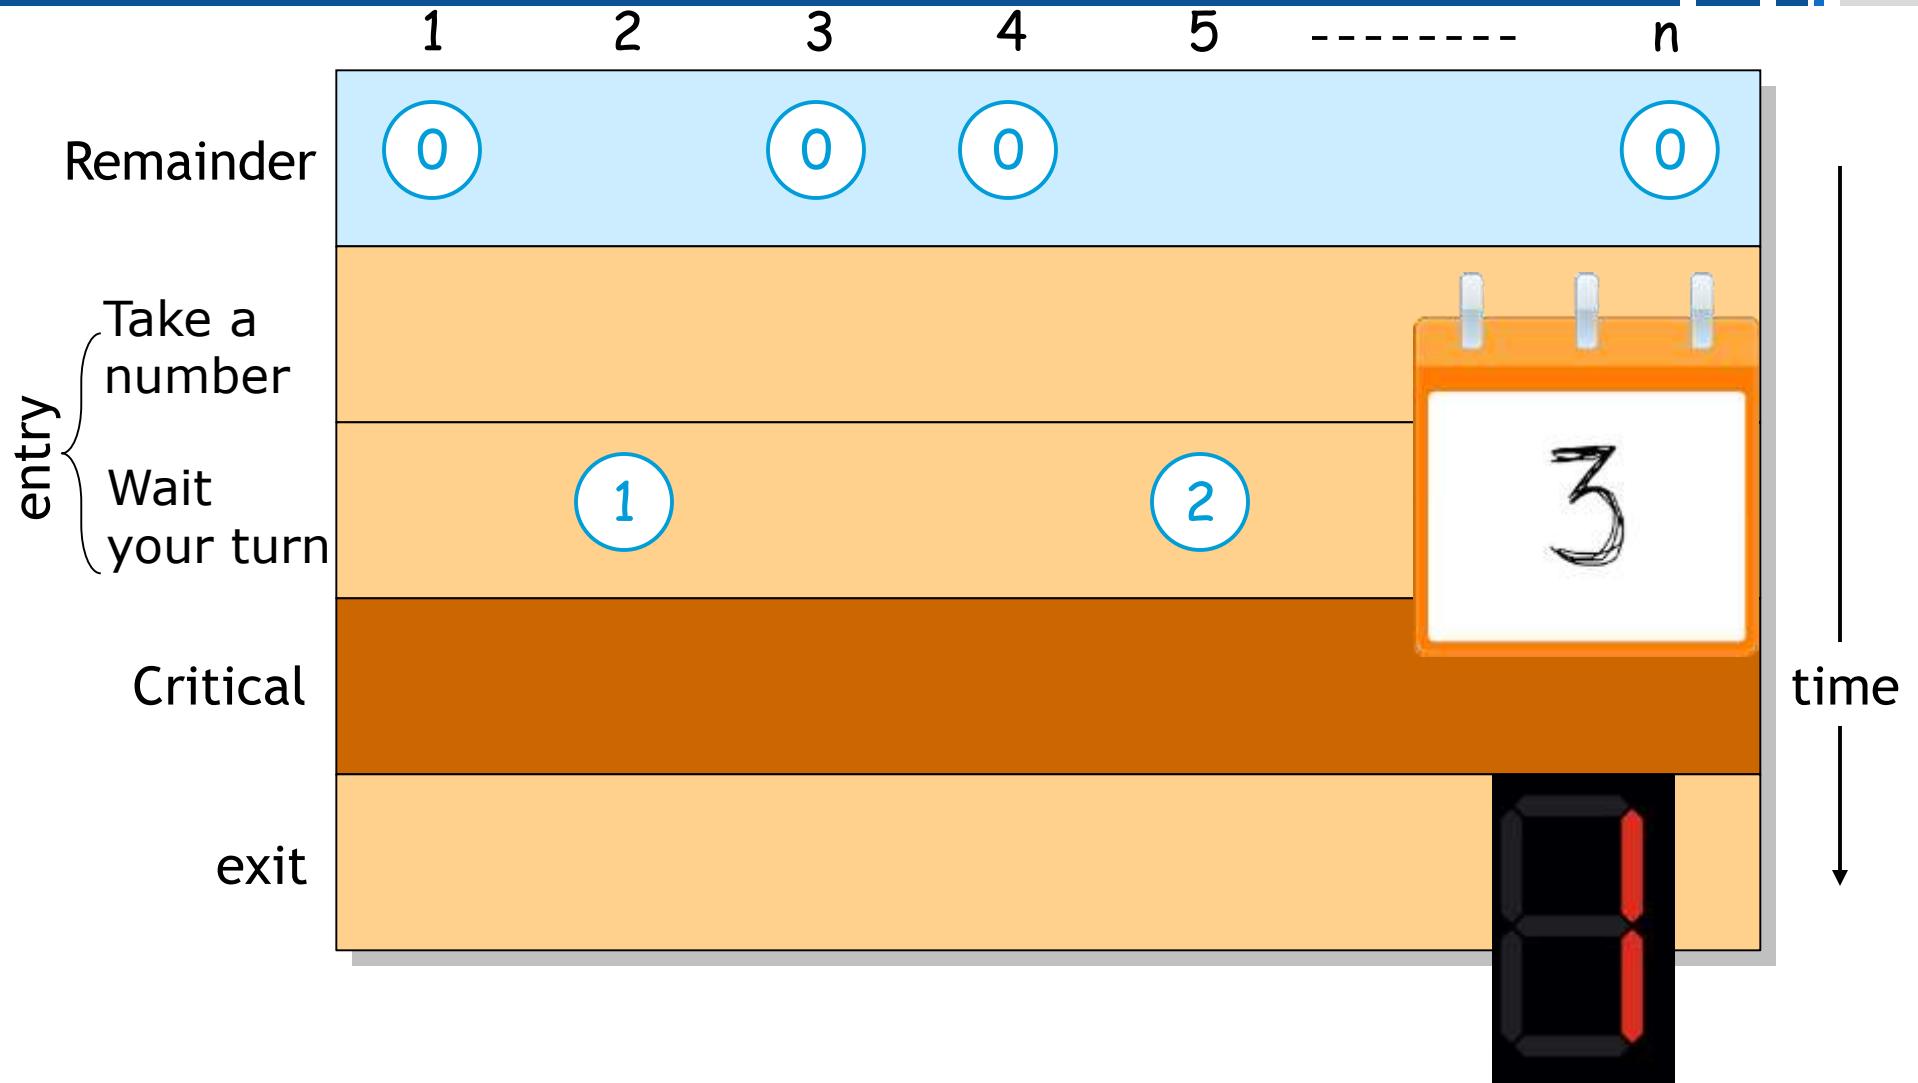
# Lamport's Bakery Algorithm - 1974



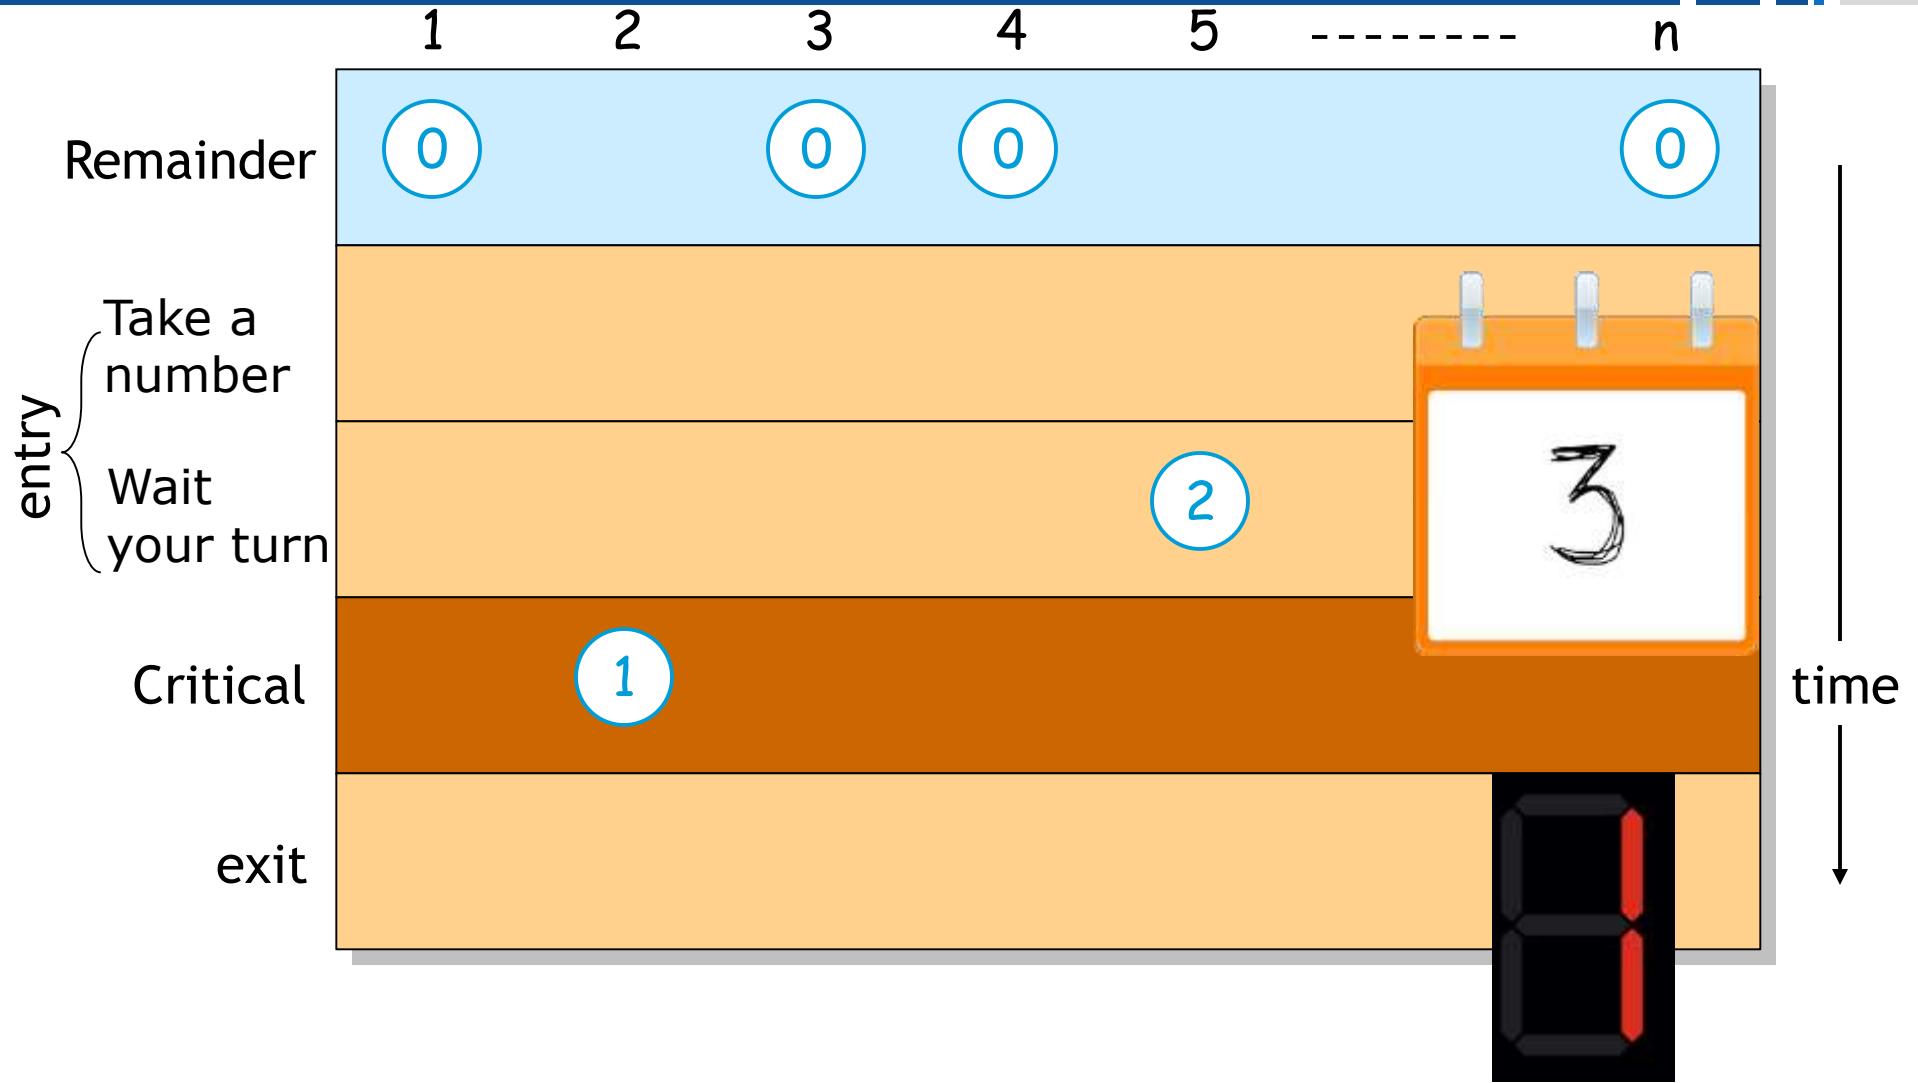
# Lamport's Bakery Algorithm - 1974



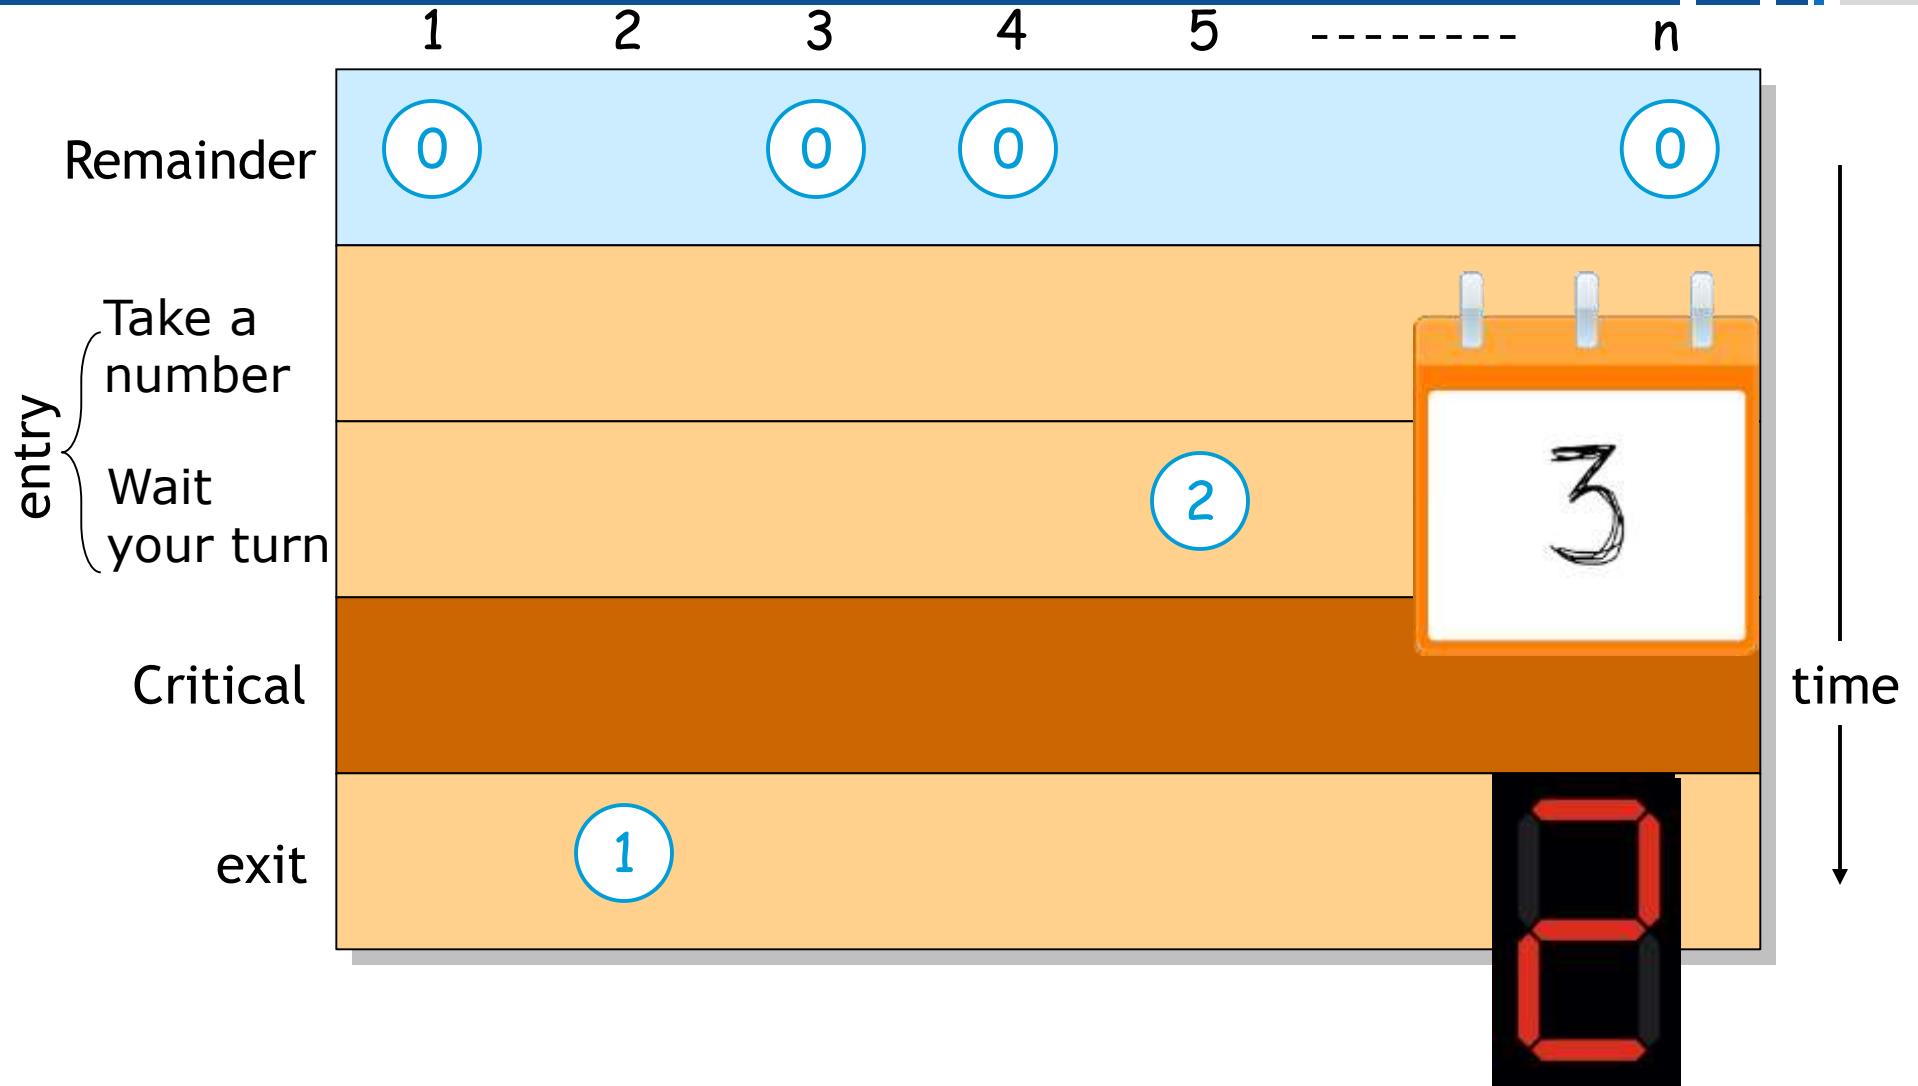
# Lamport's Bakery Algorithm - 1974



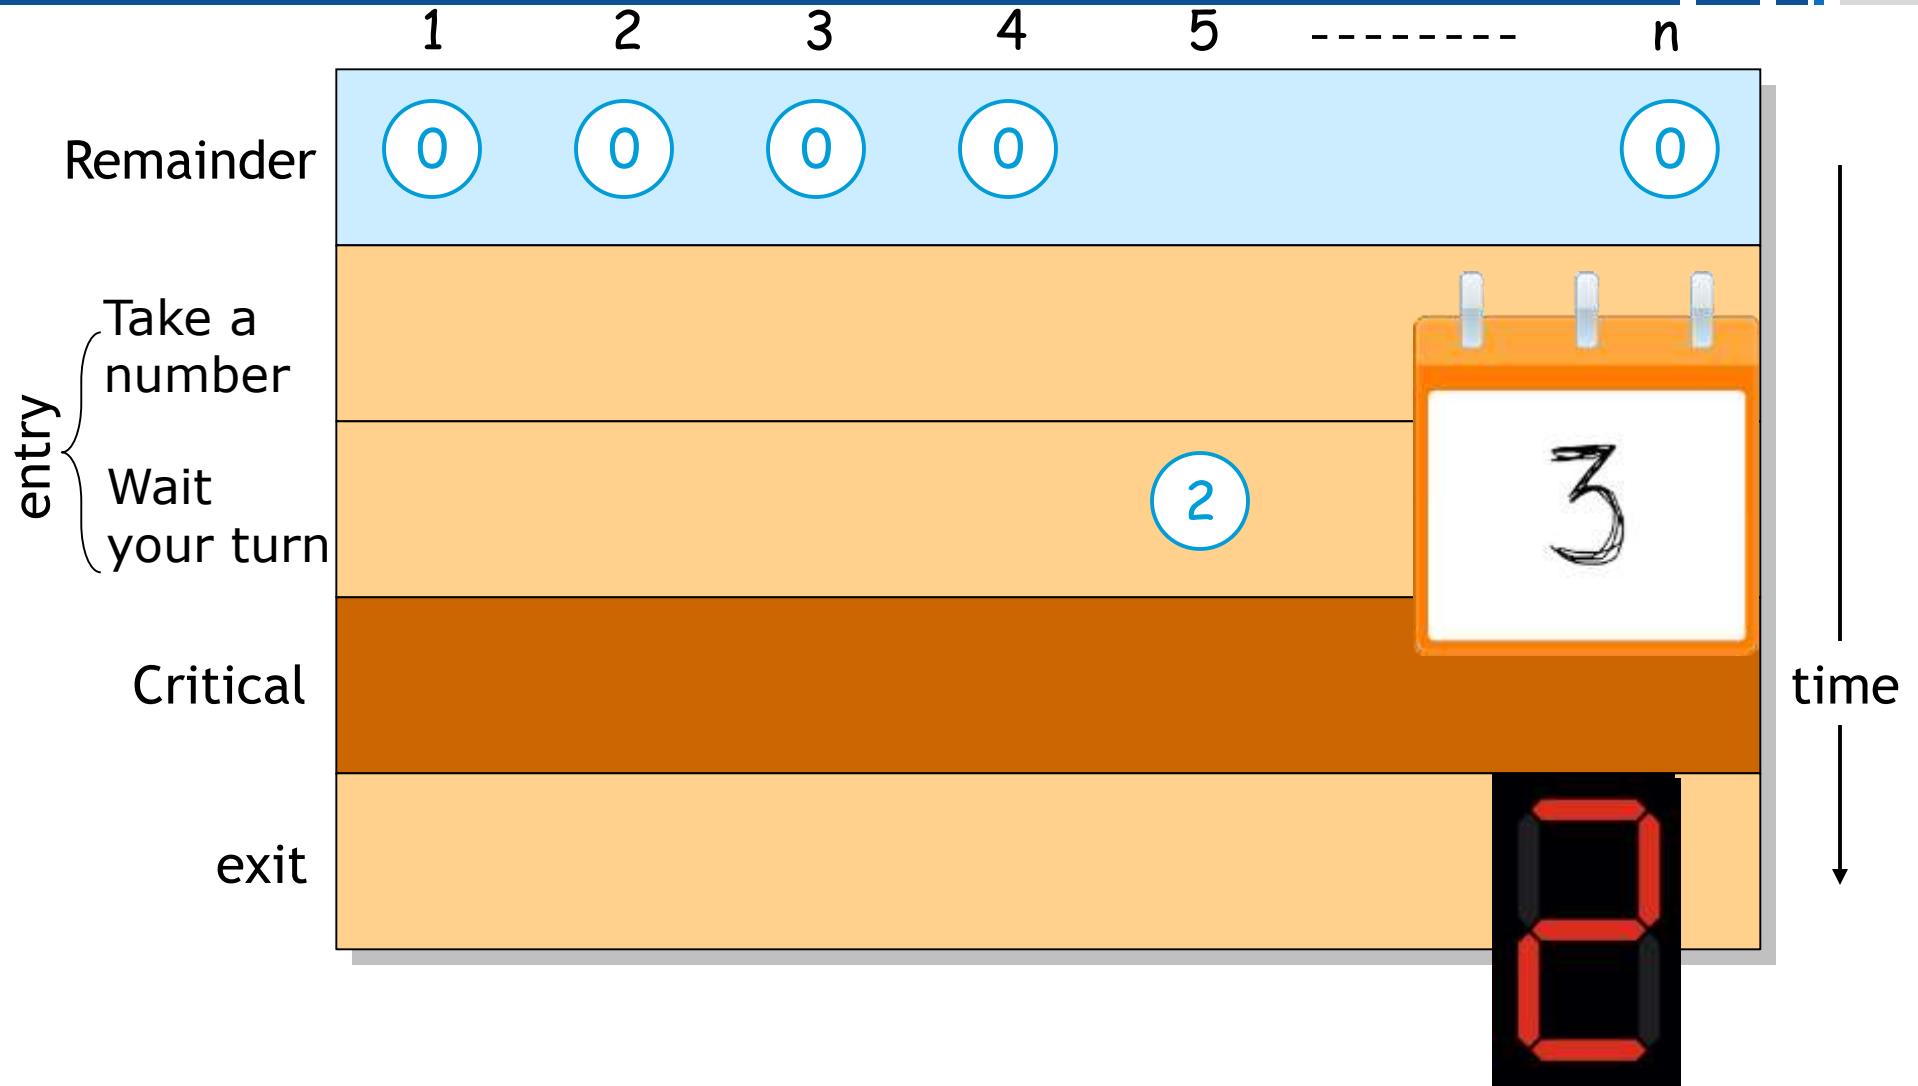
# Lamport's Bakery Algorithm - 1974



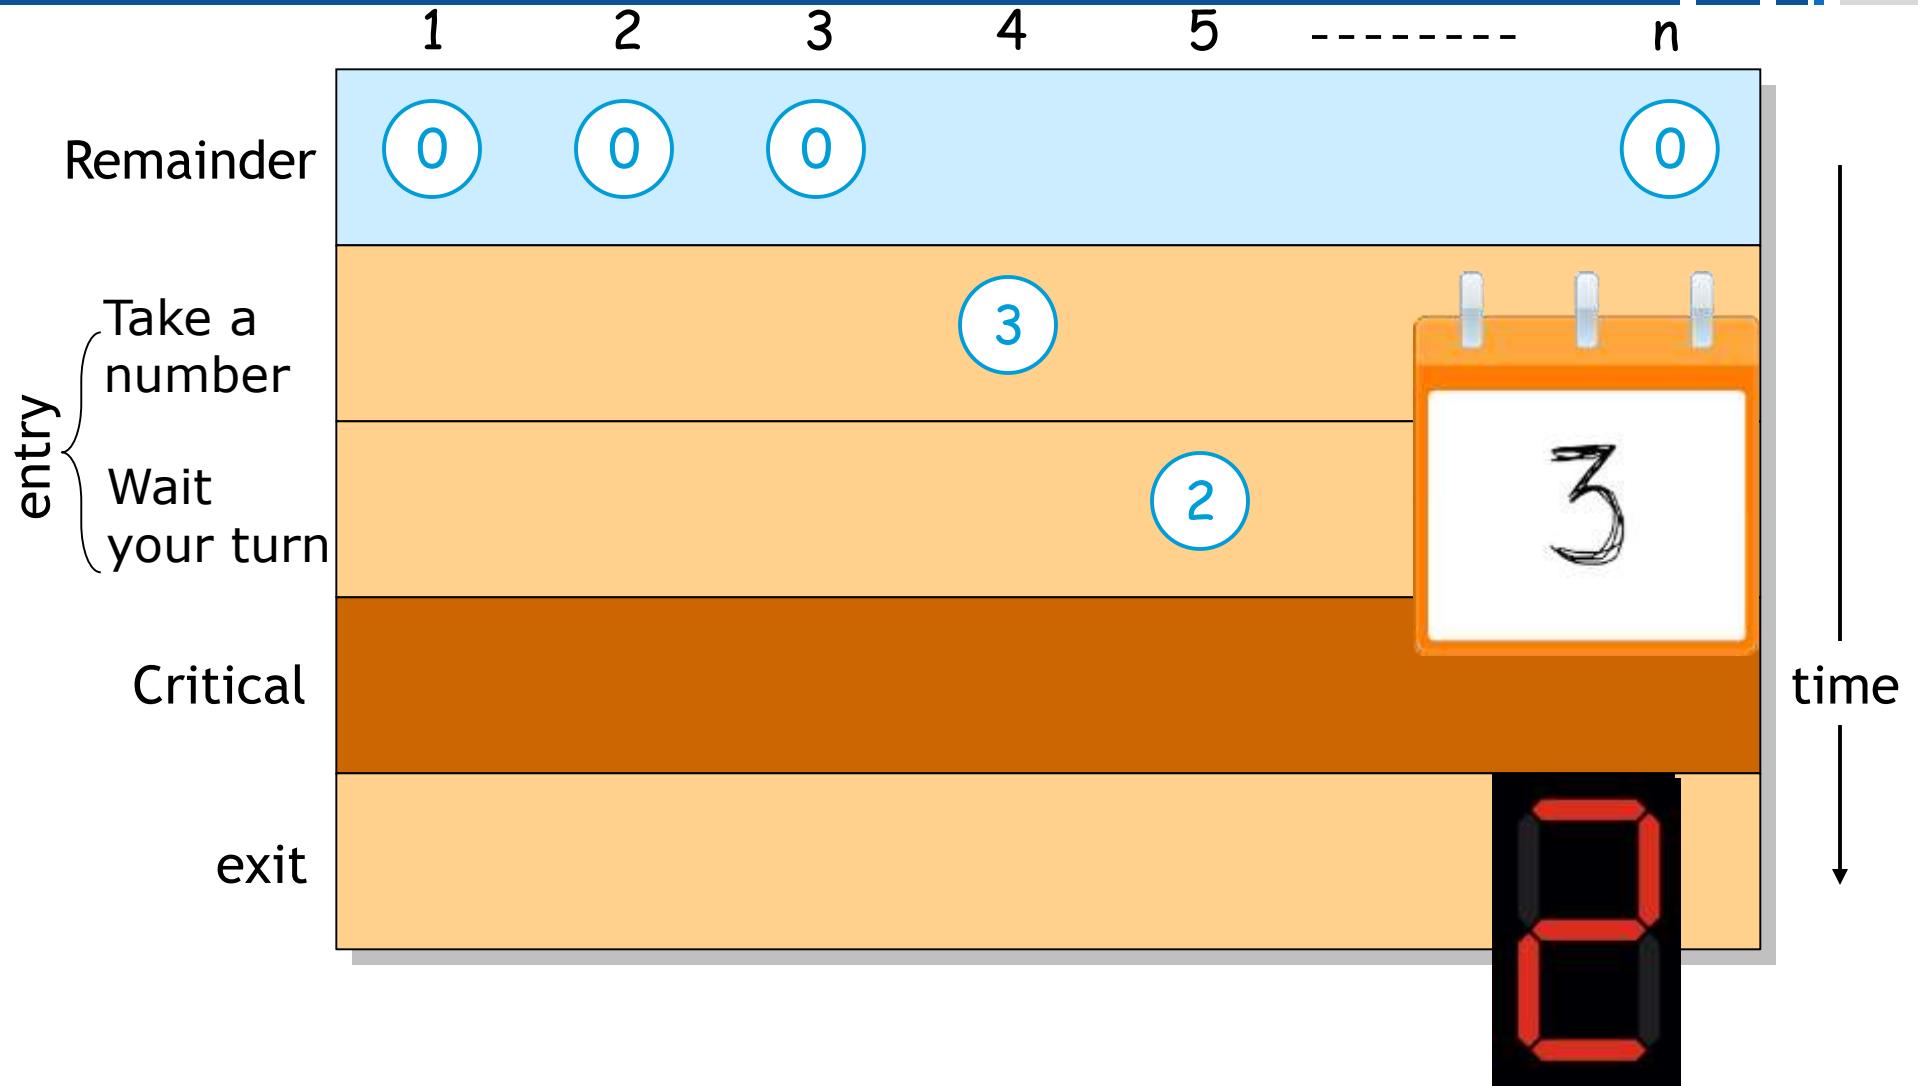
# Lamport's Bakery Algorithm - 1974



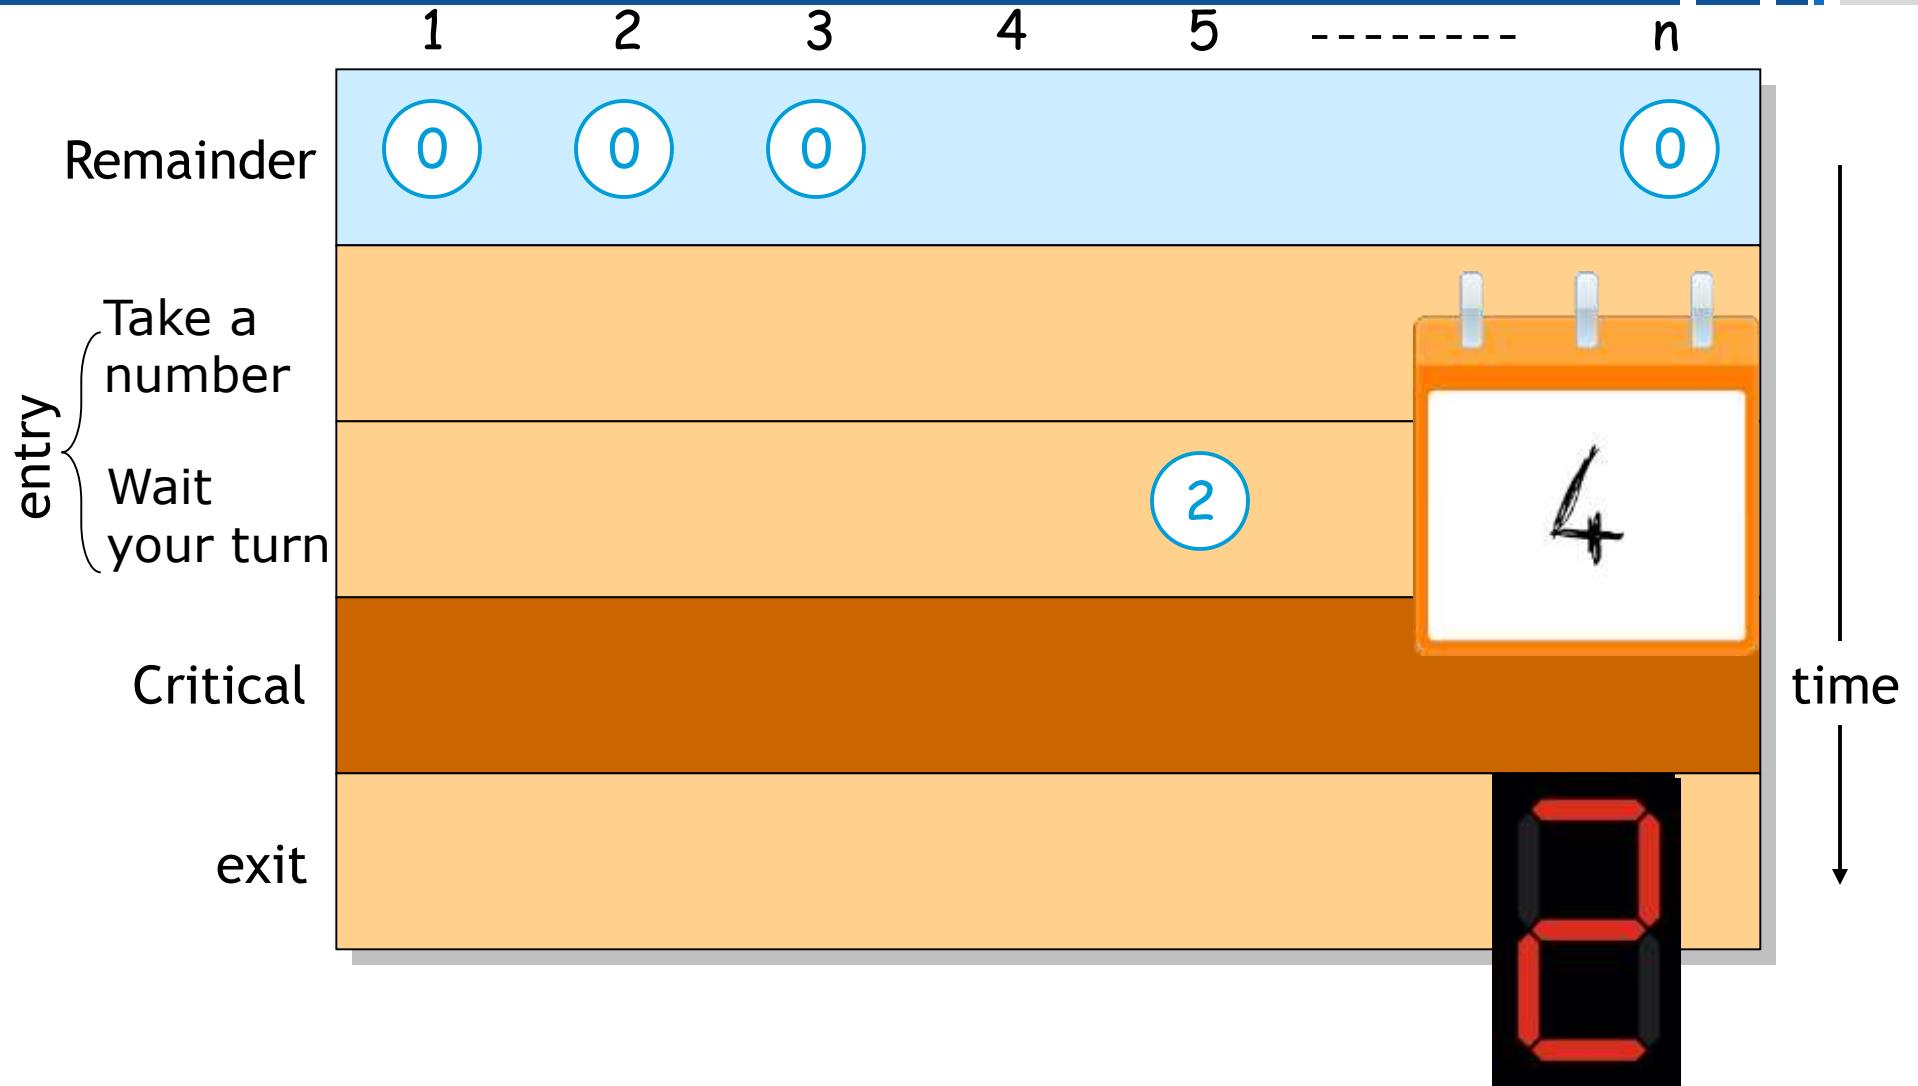
# Lamport's Bakery Algorithm - 1974



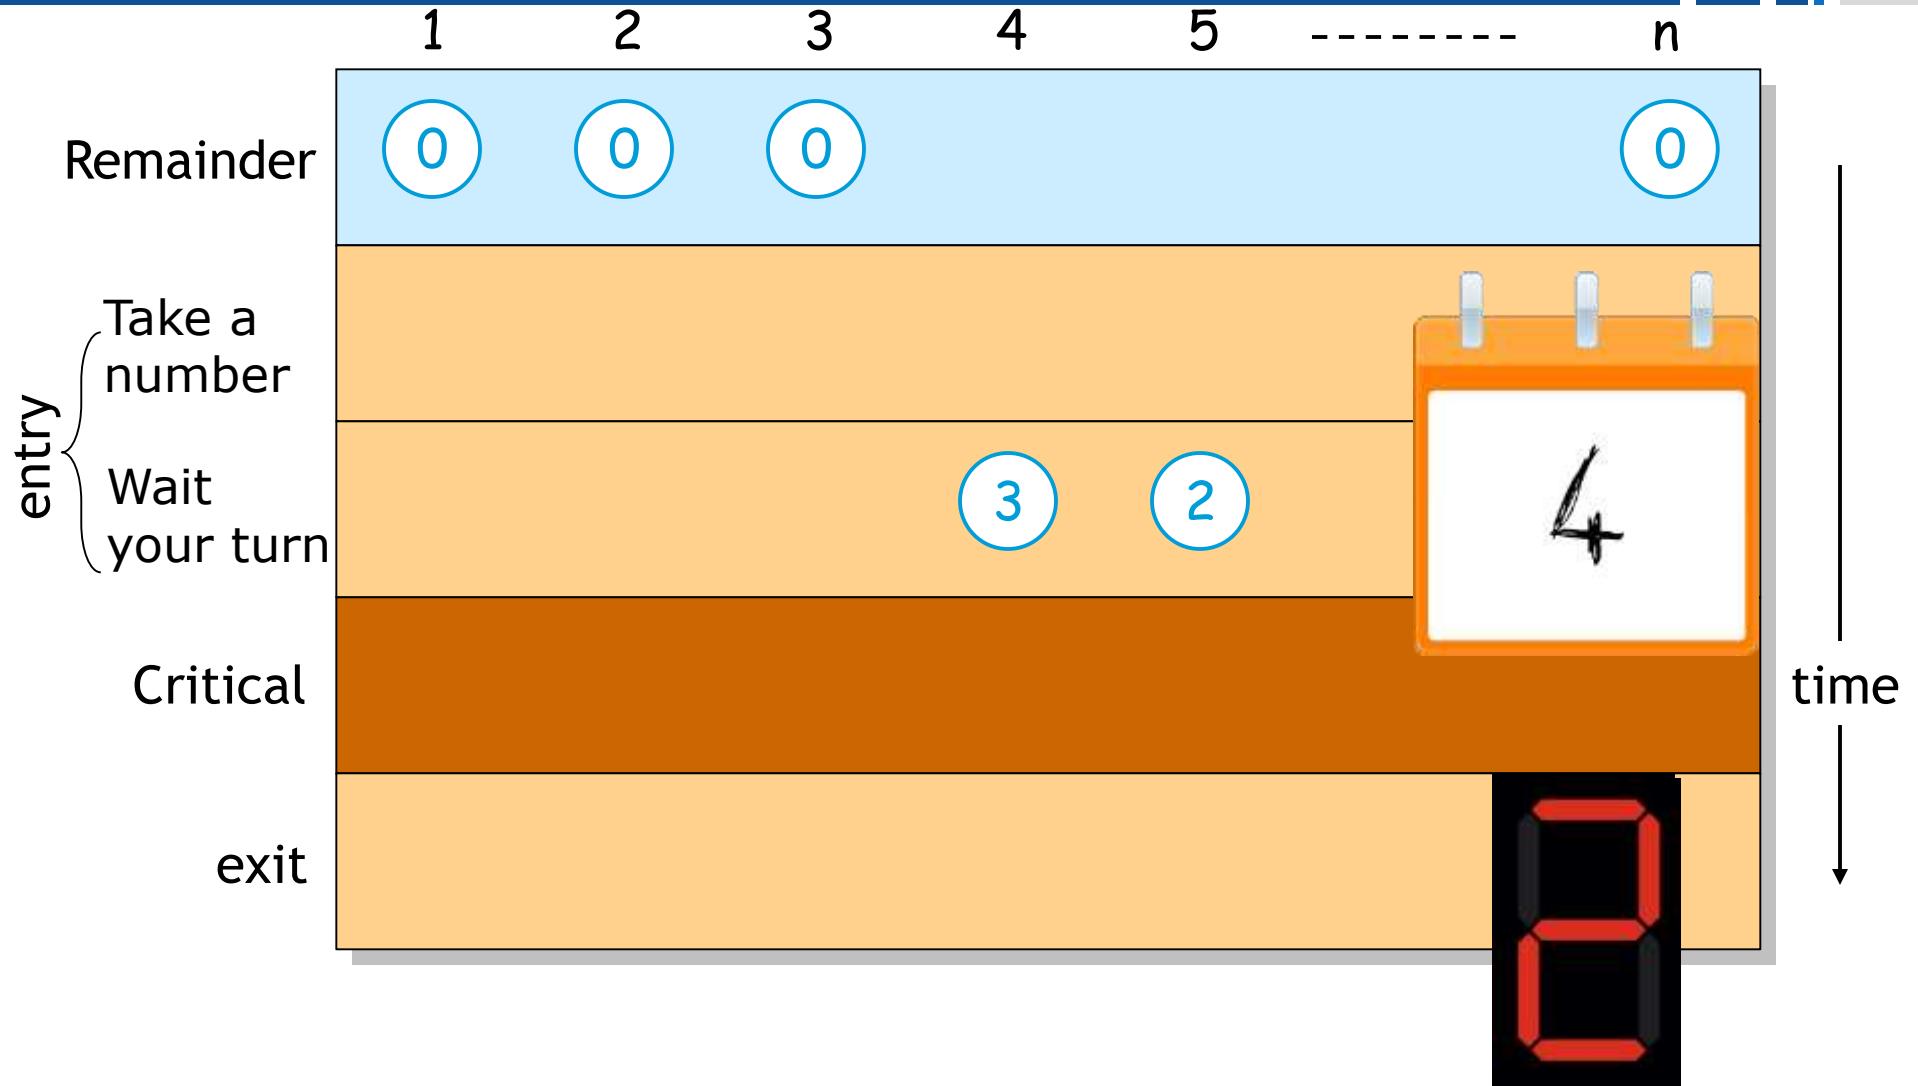
# Lamport's Bakery Algorithm - 1974



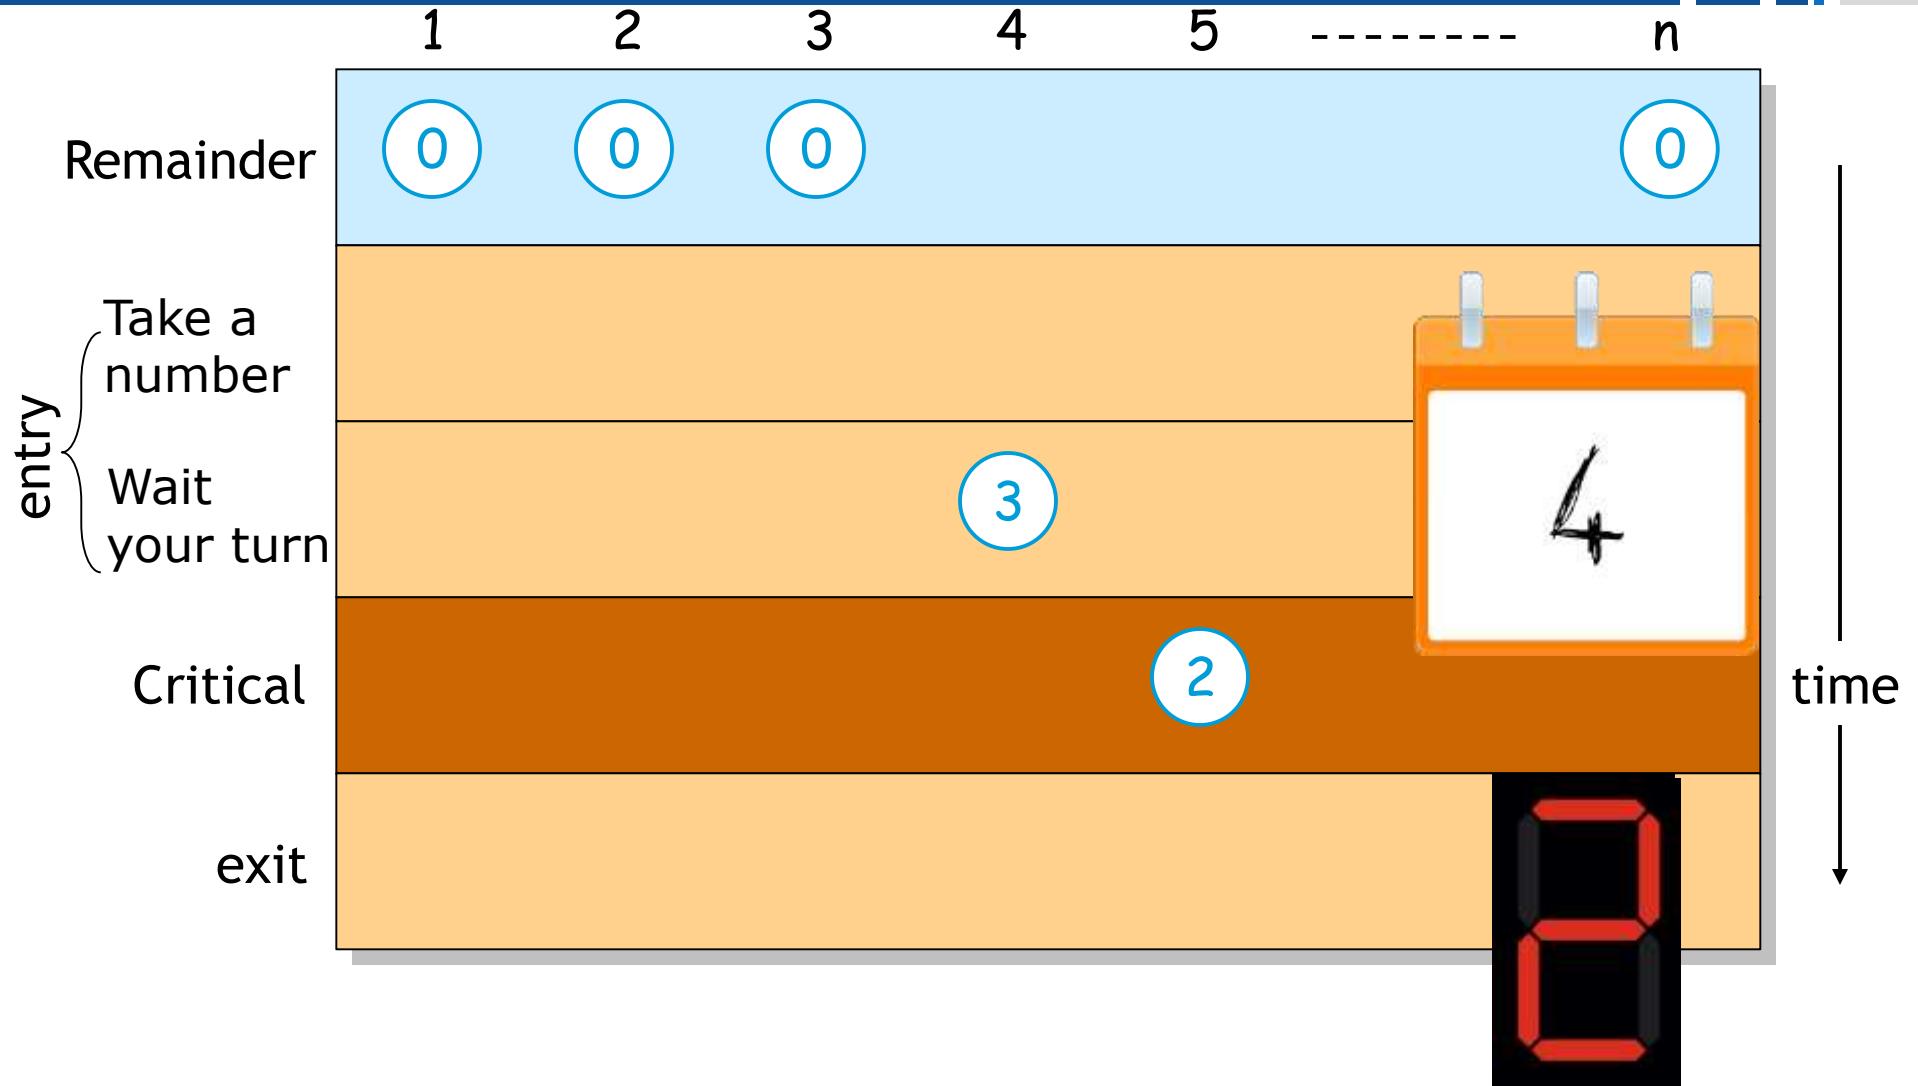
# Lamport's Bakery Algorithm - 1974



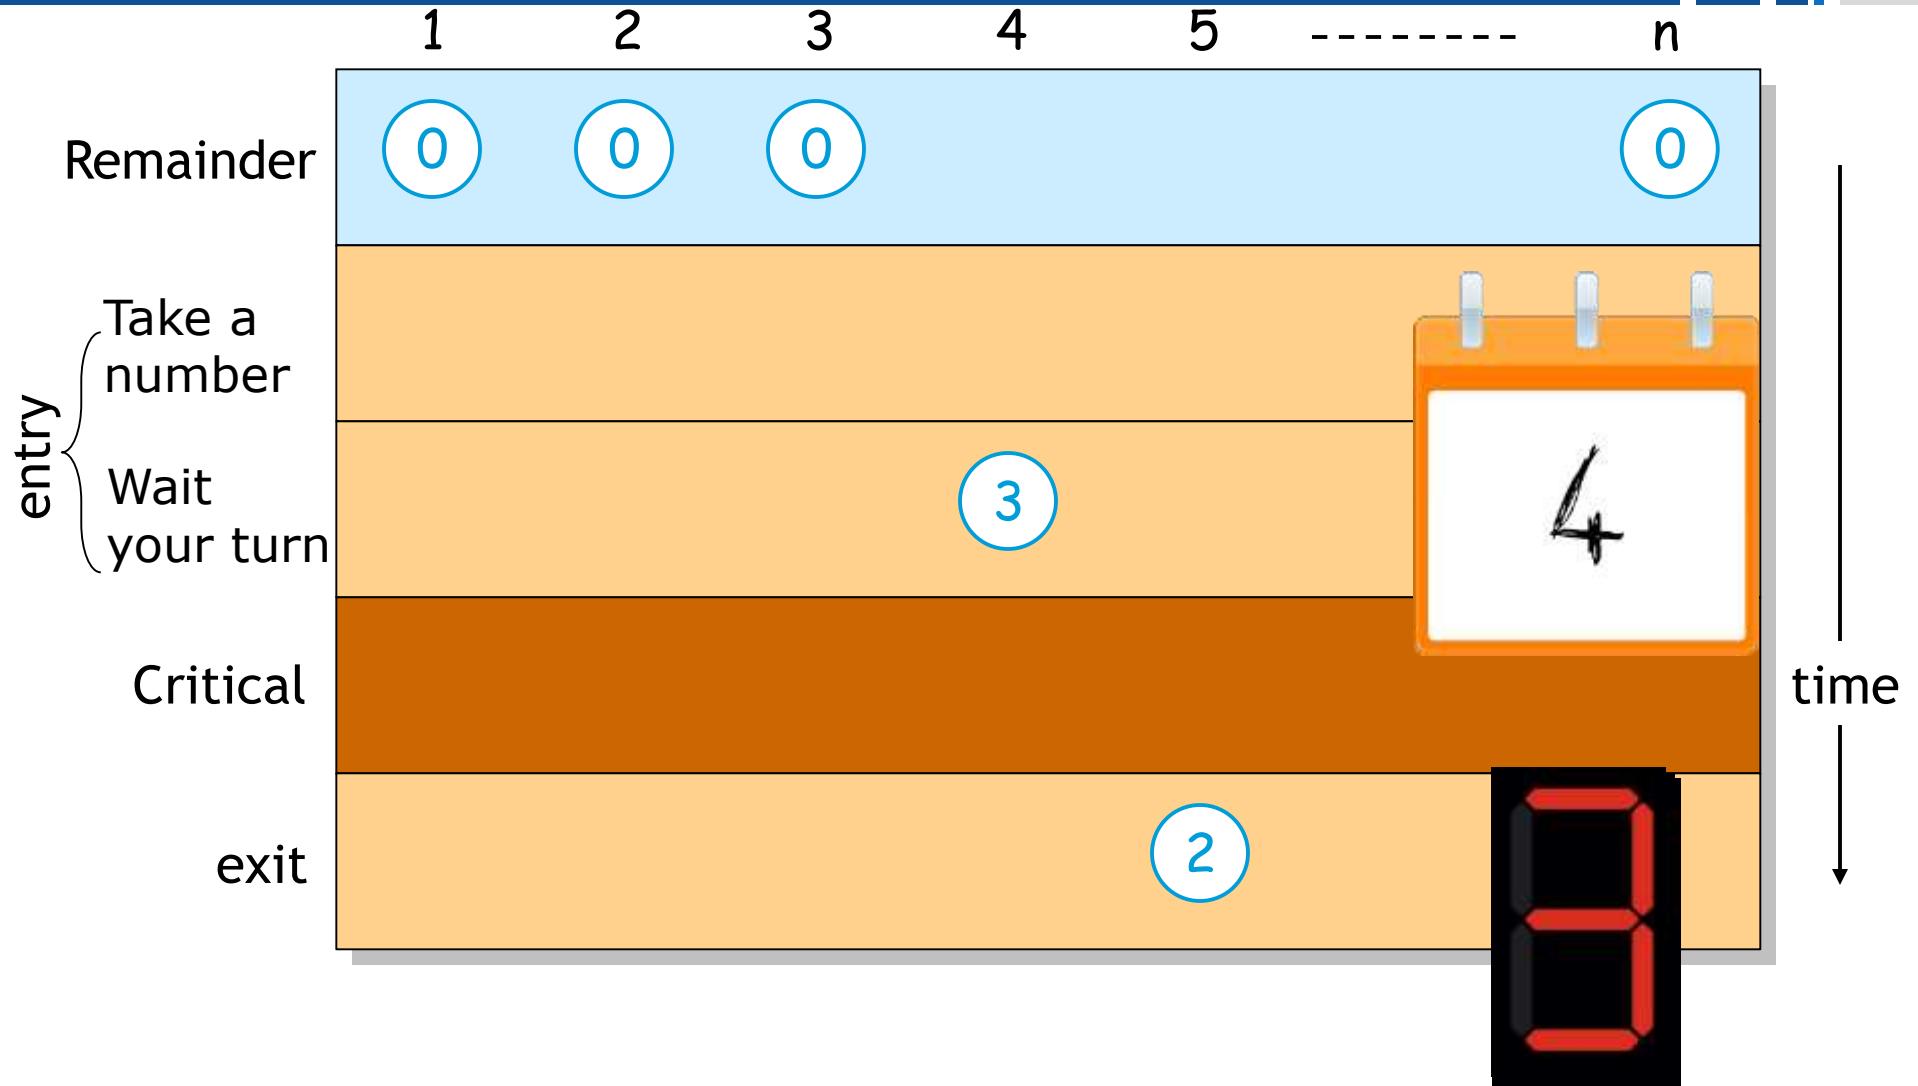
# Lamport's Bakery Algorithm - 1974



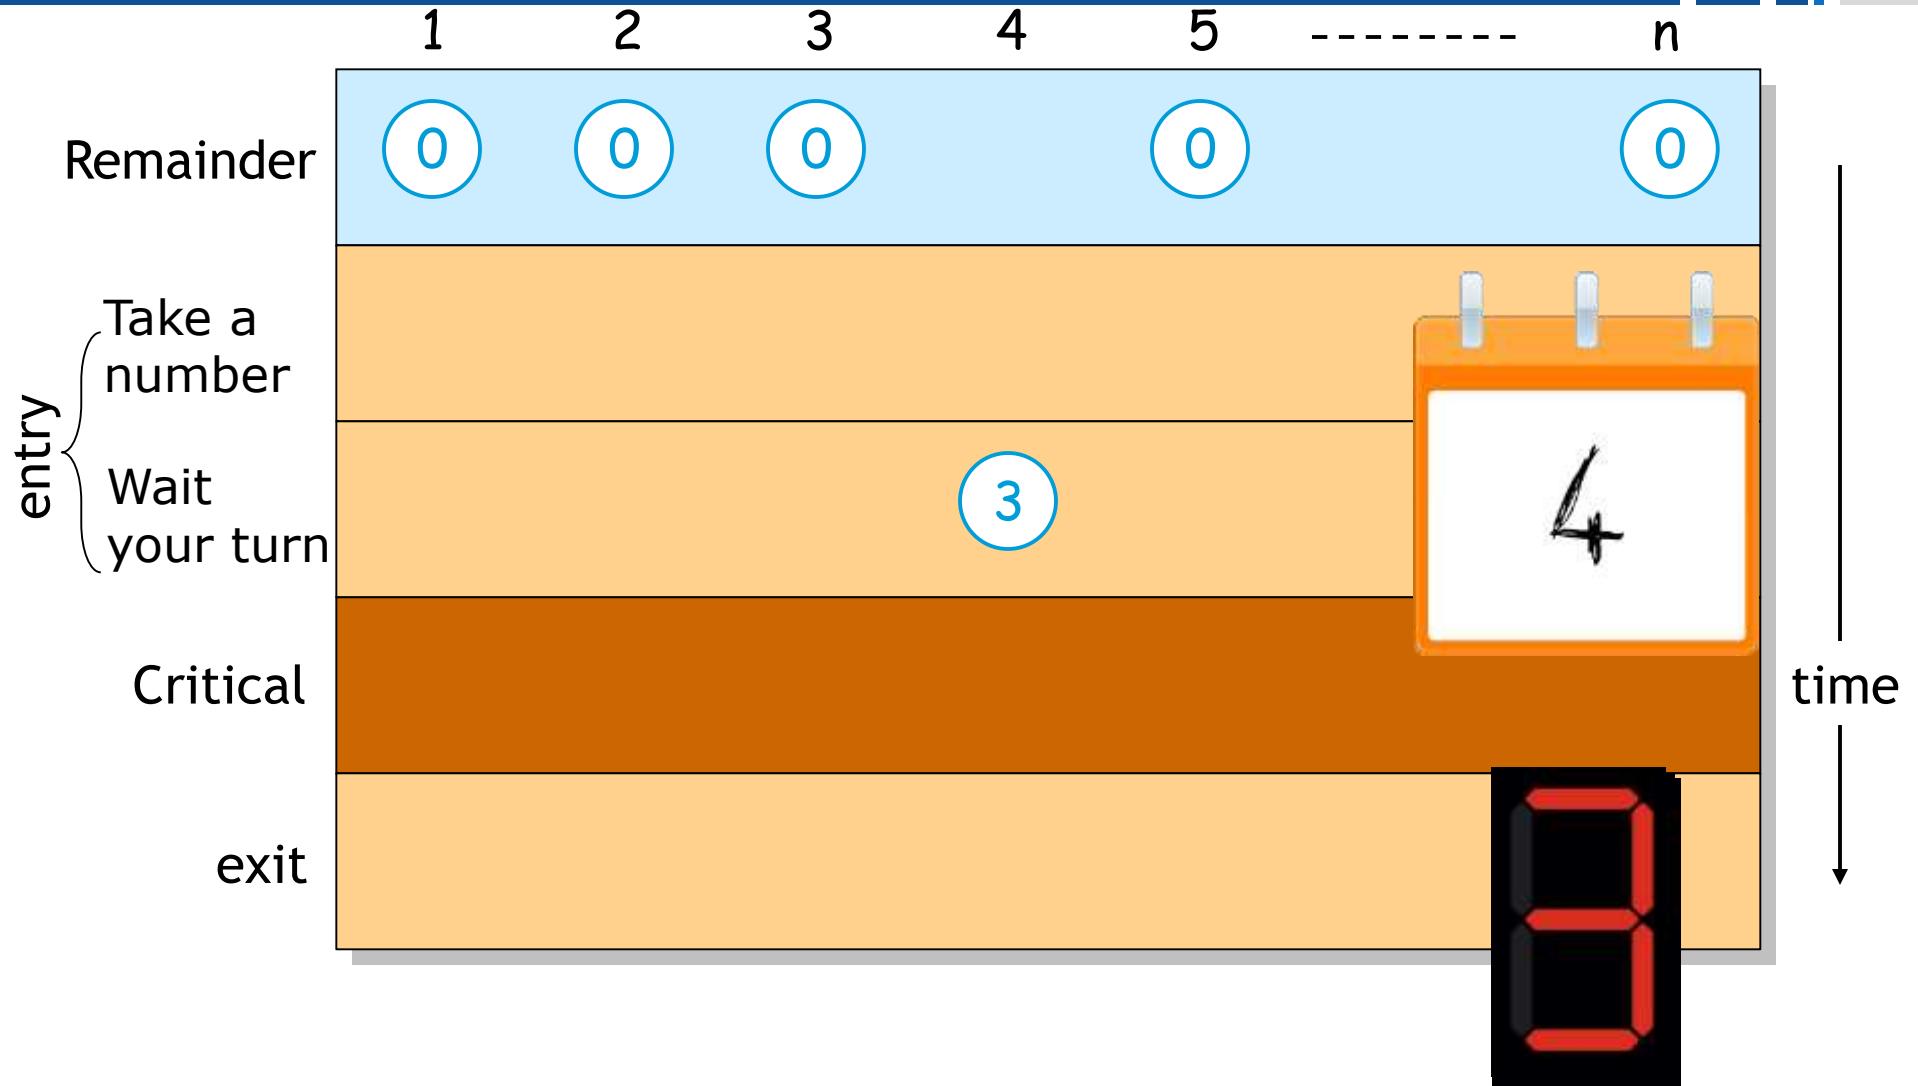
# Lamport's Bakery Algorithm - 1974



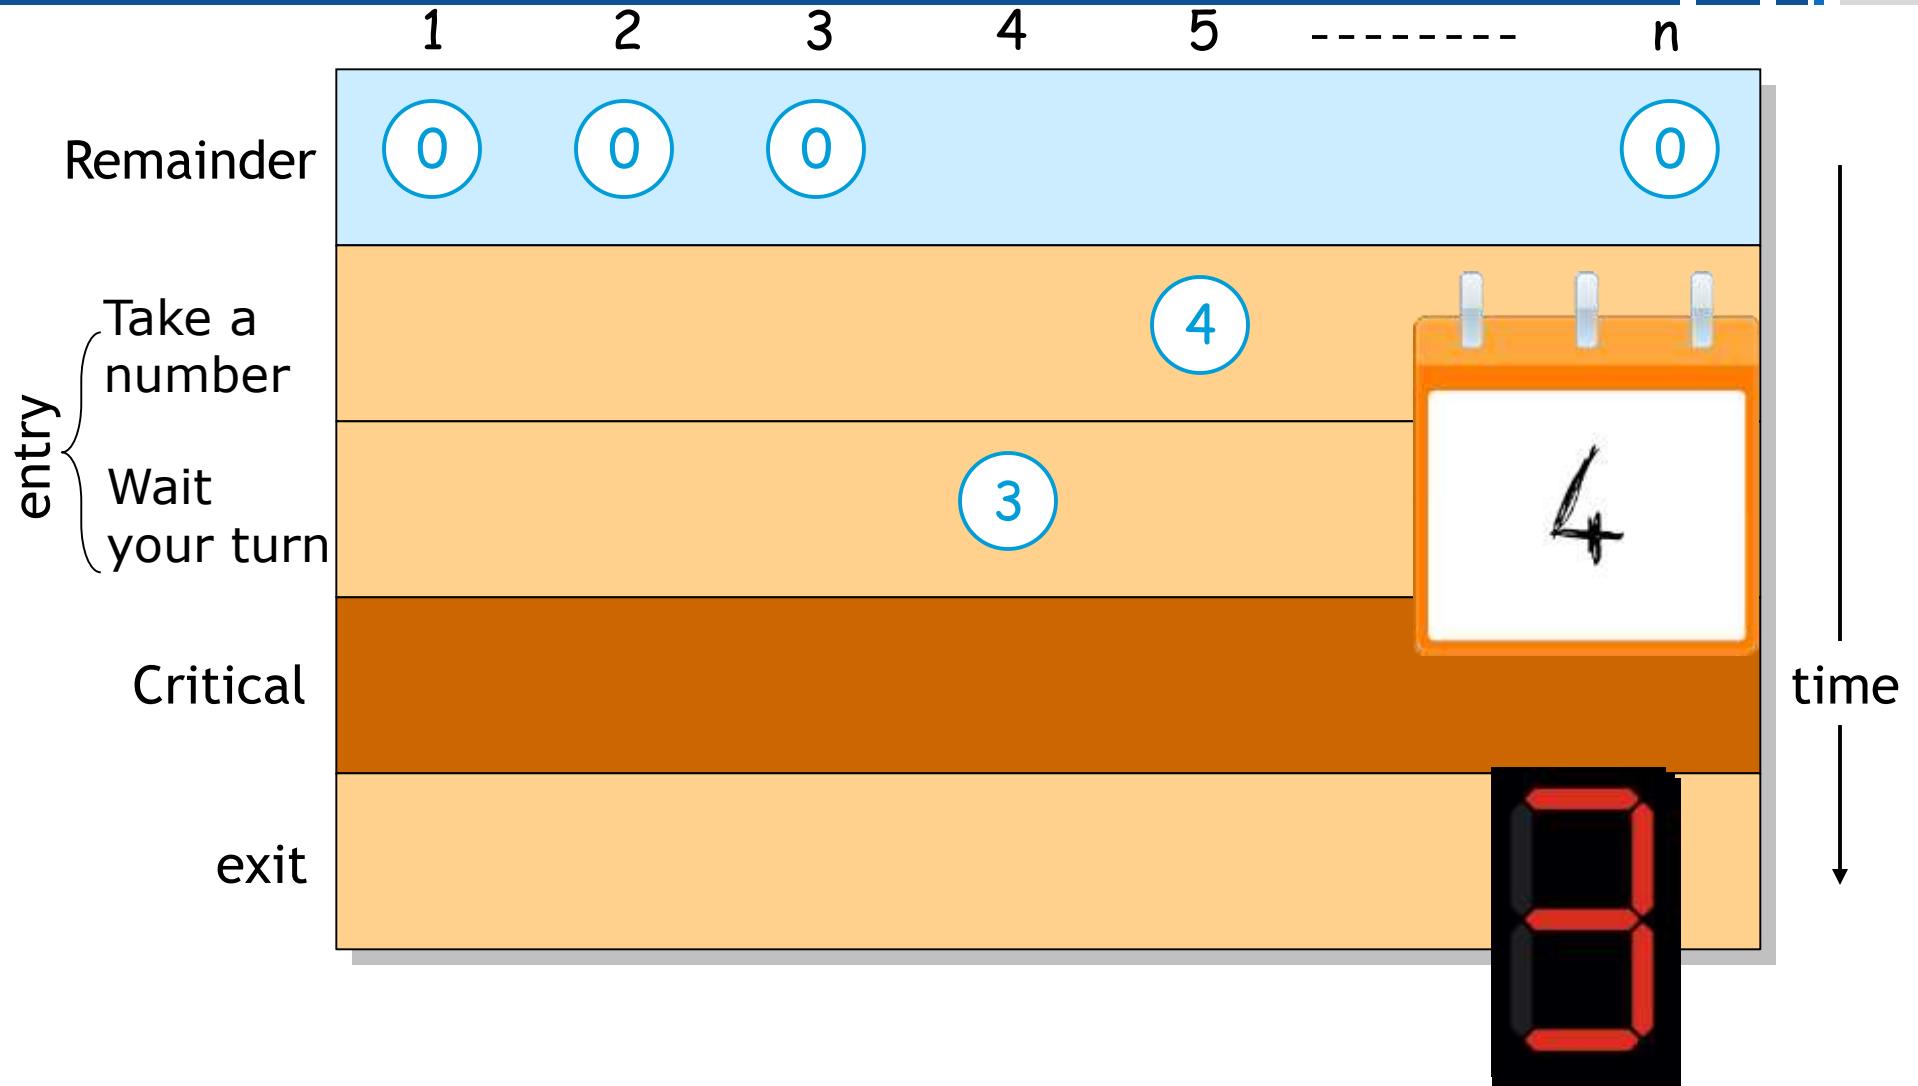
# Lamport's Bakery Algorithm - 1974



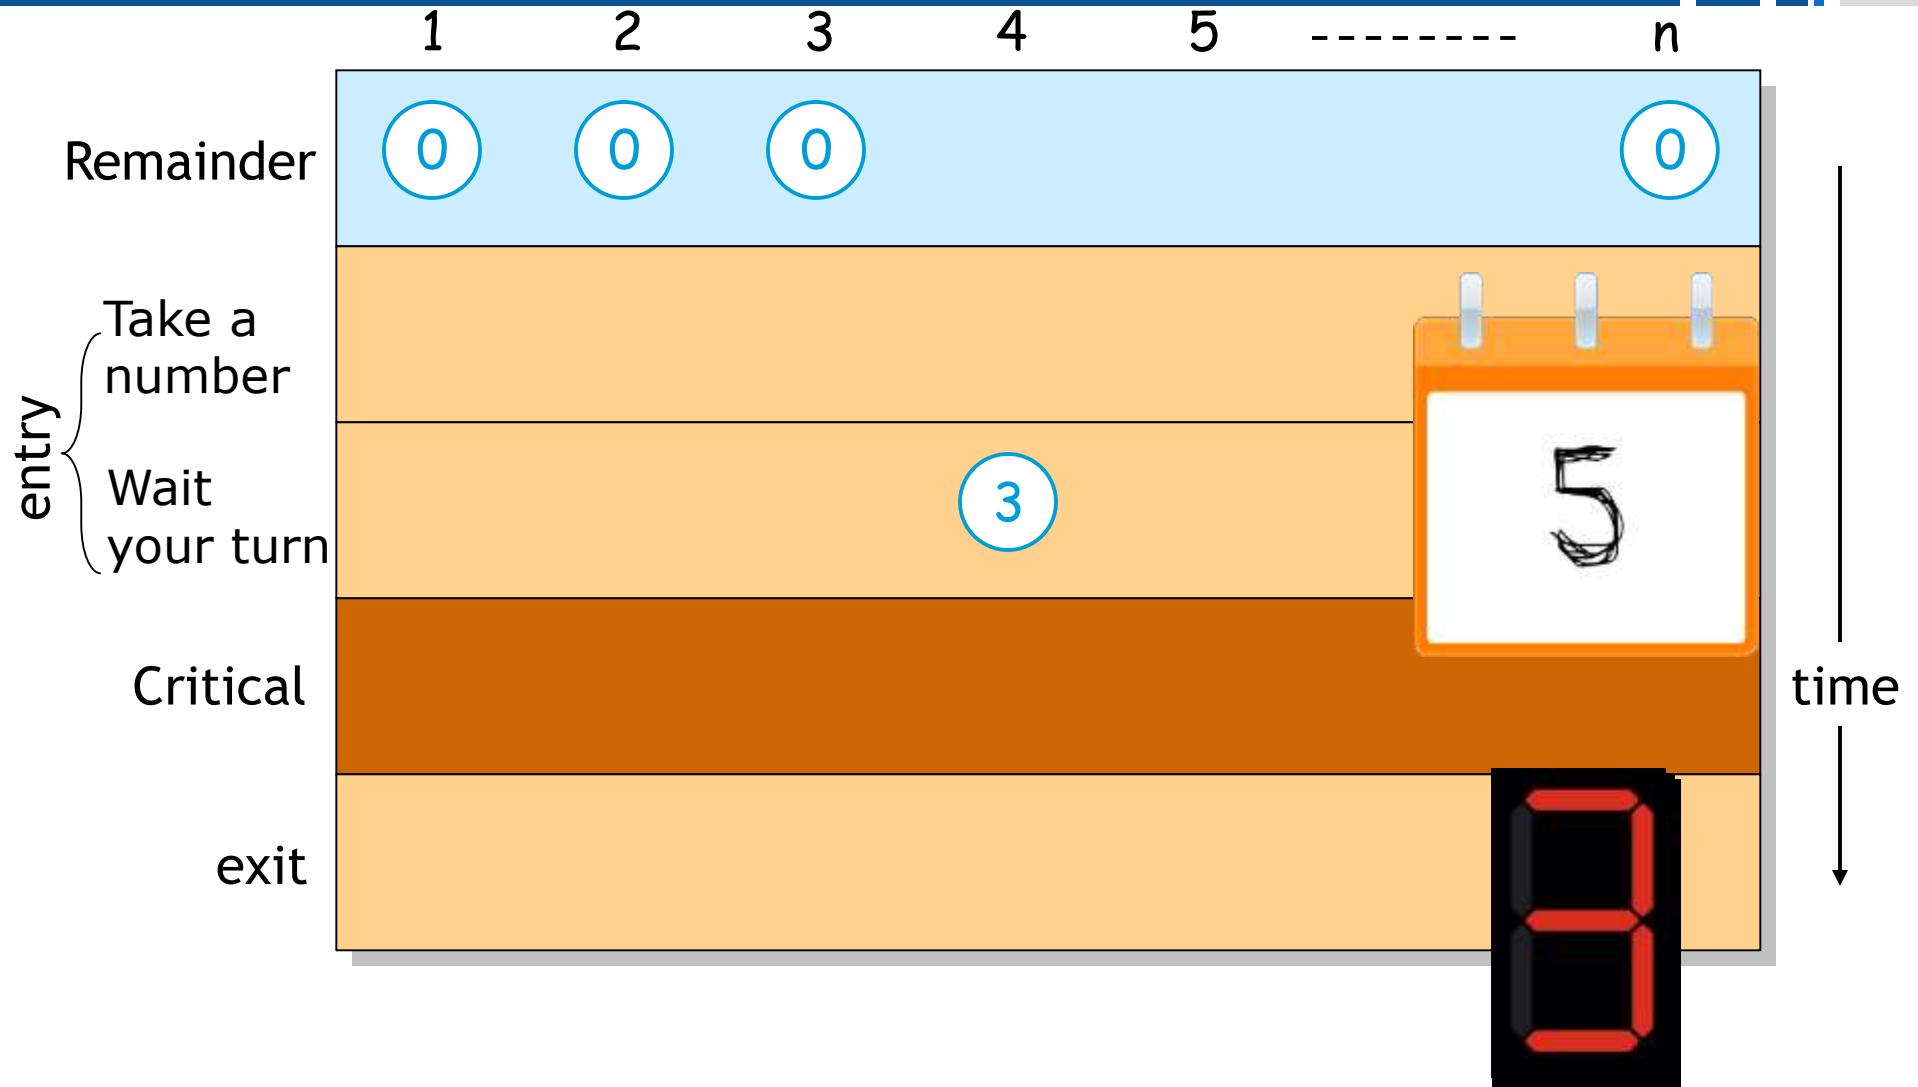
# Lamport's Bakery Algorithm - 1974



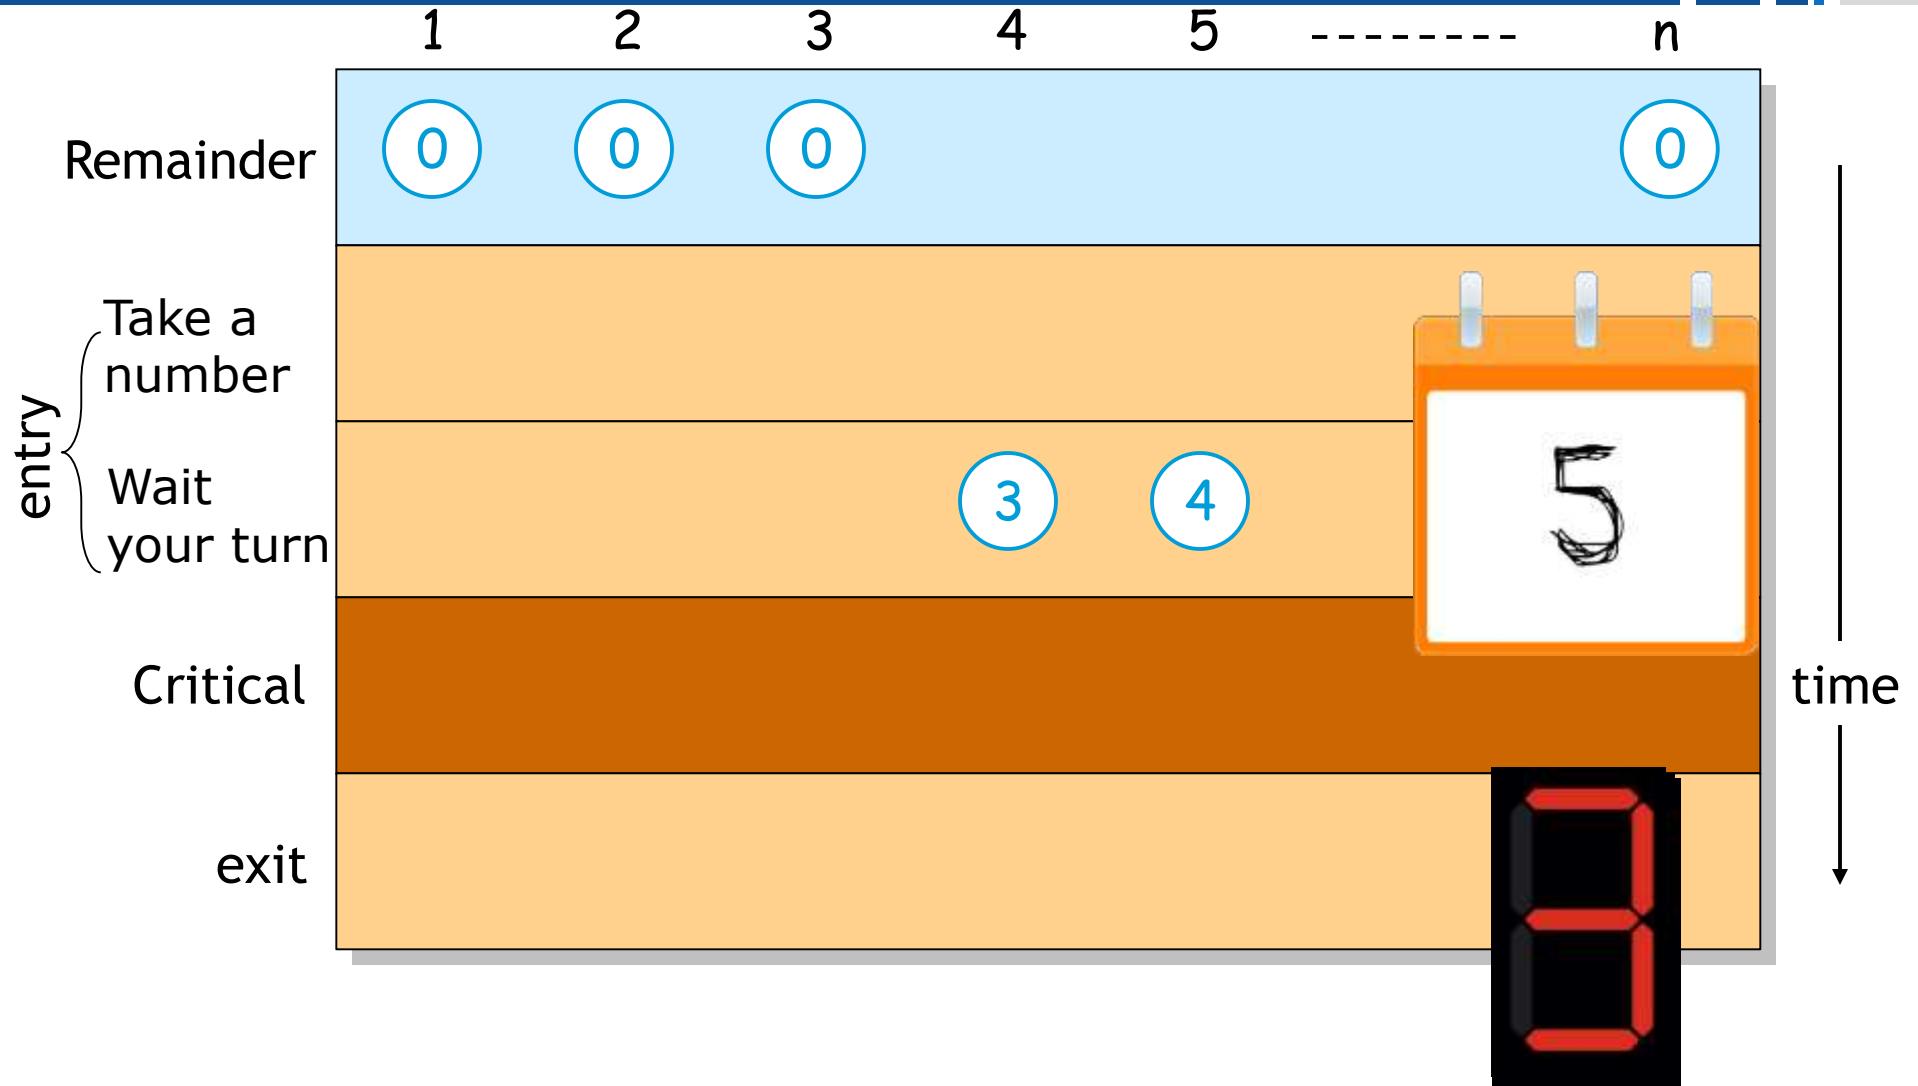
# Lamport's Bakery Algorithm - 1974



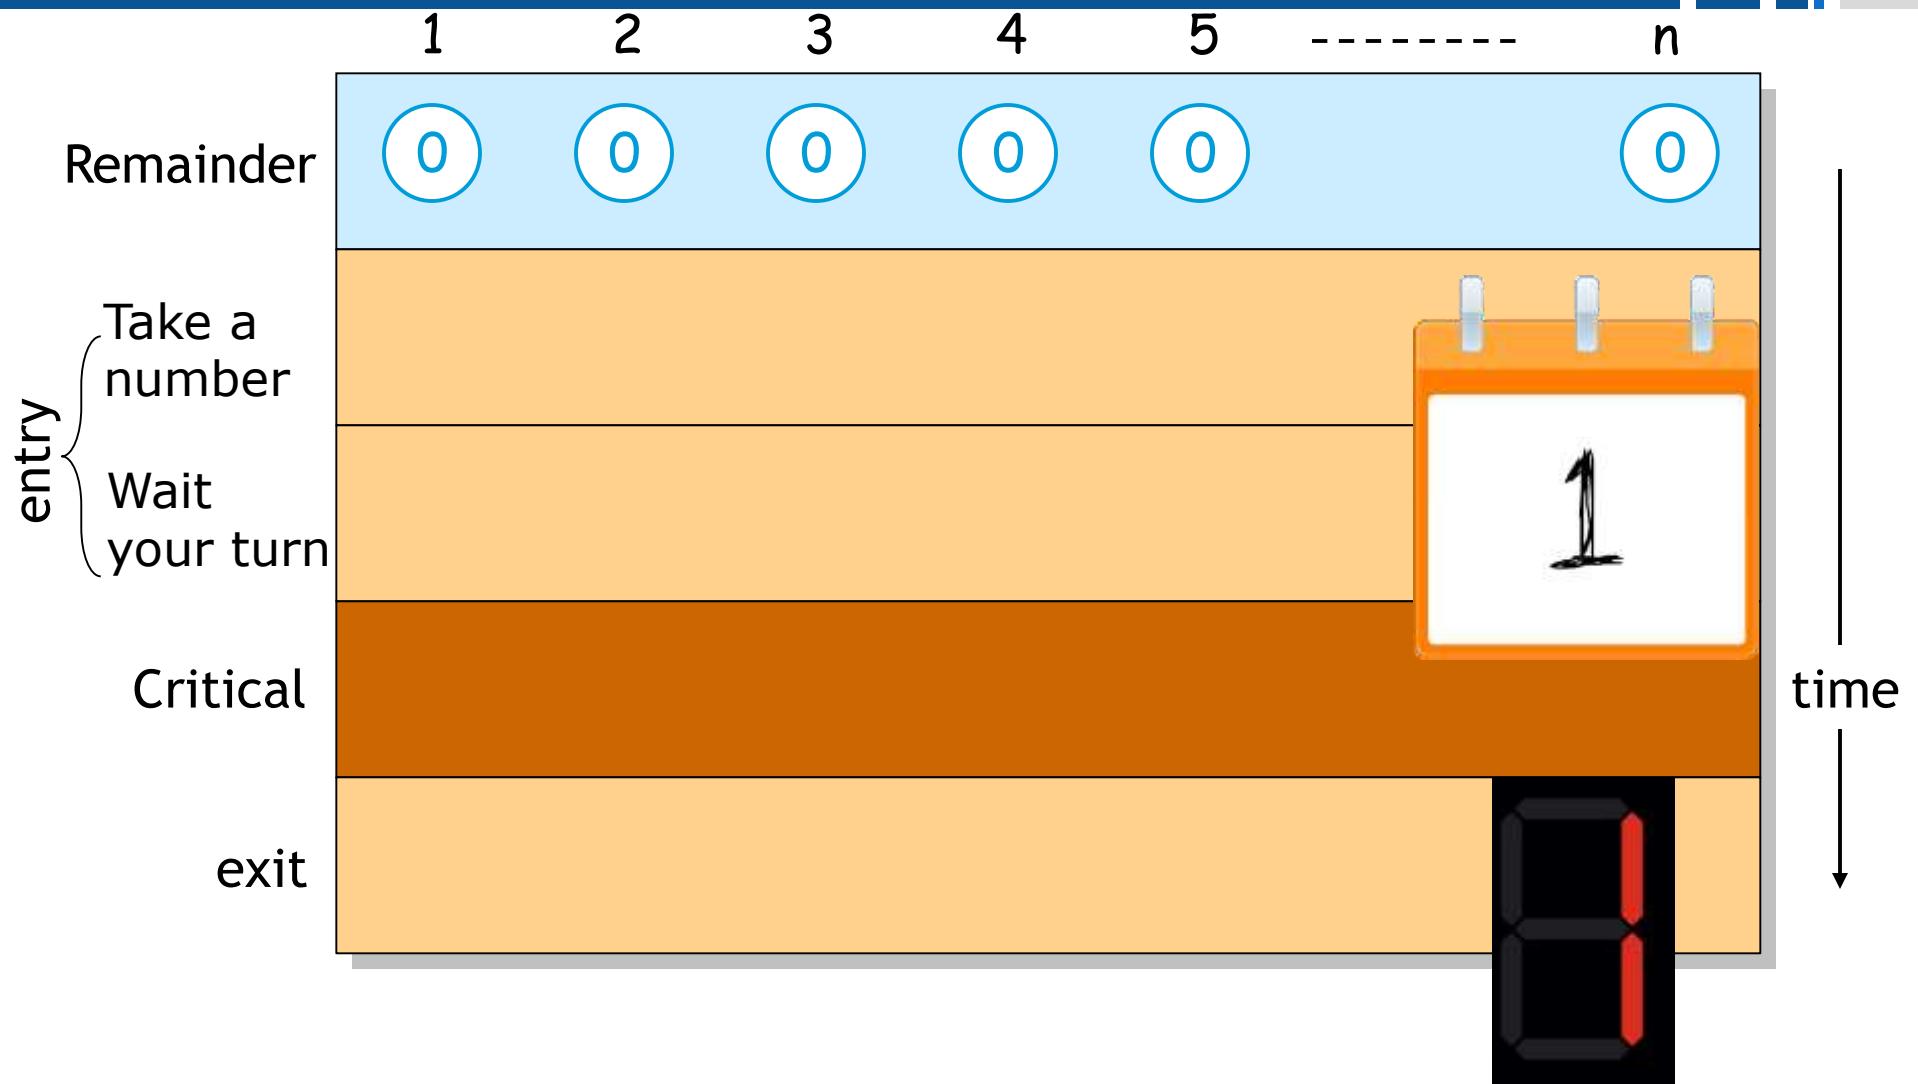
# Lamport's Bakery Algorithm - 1974



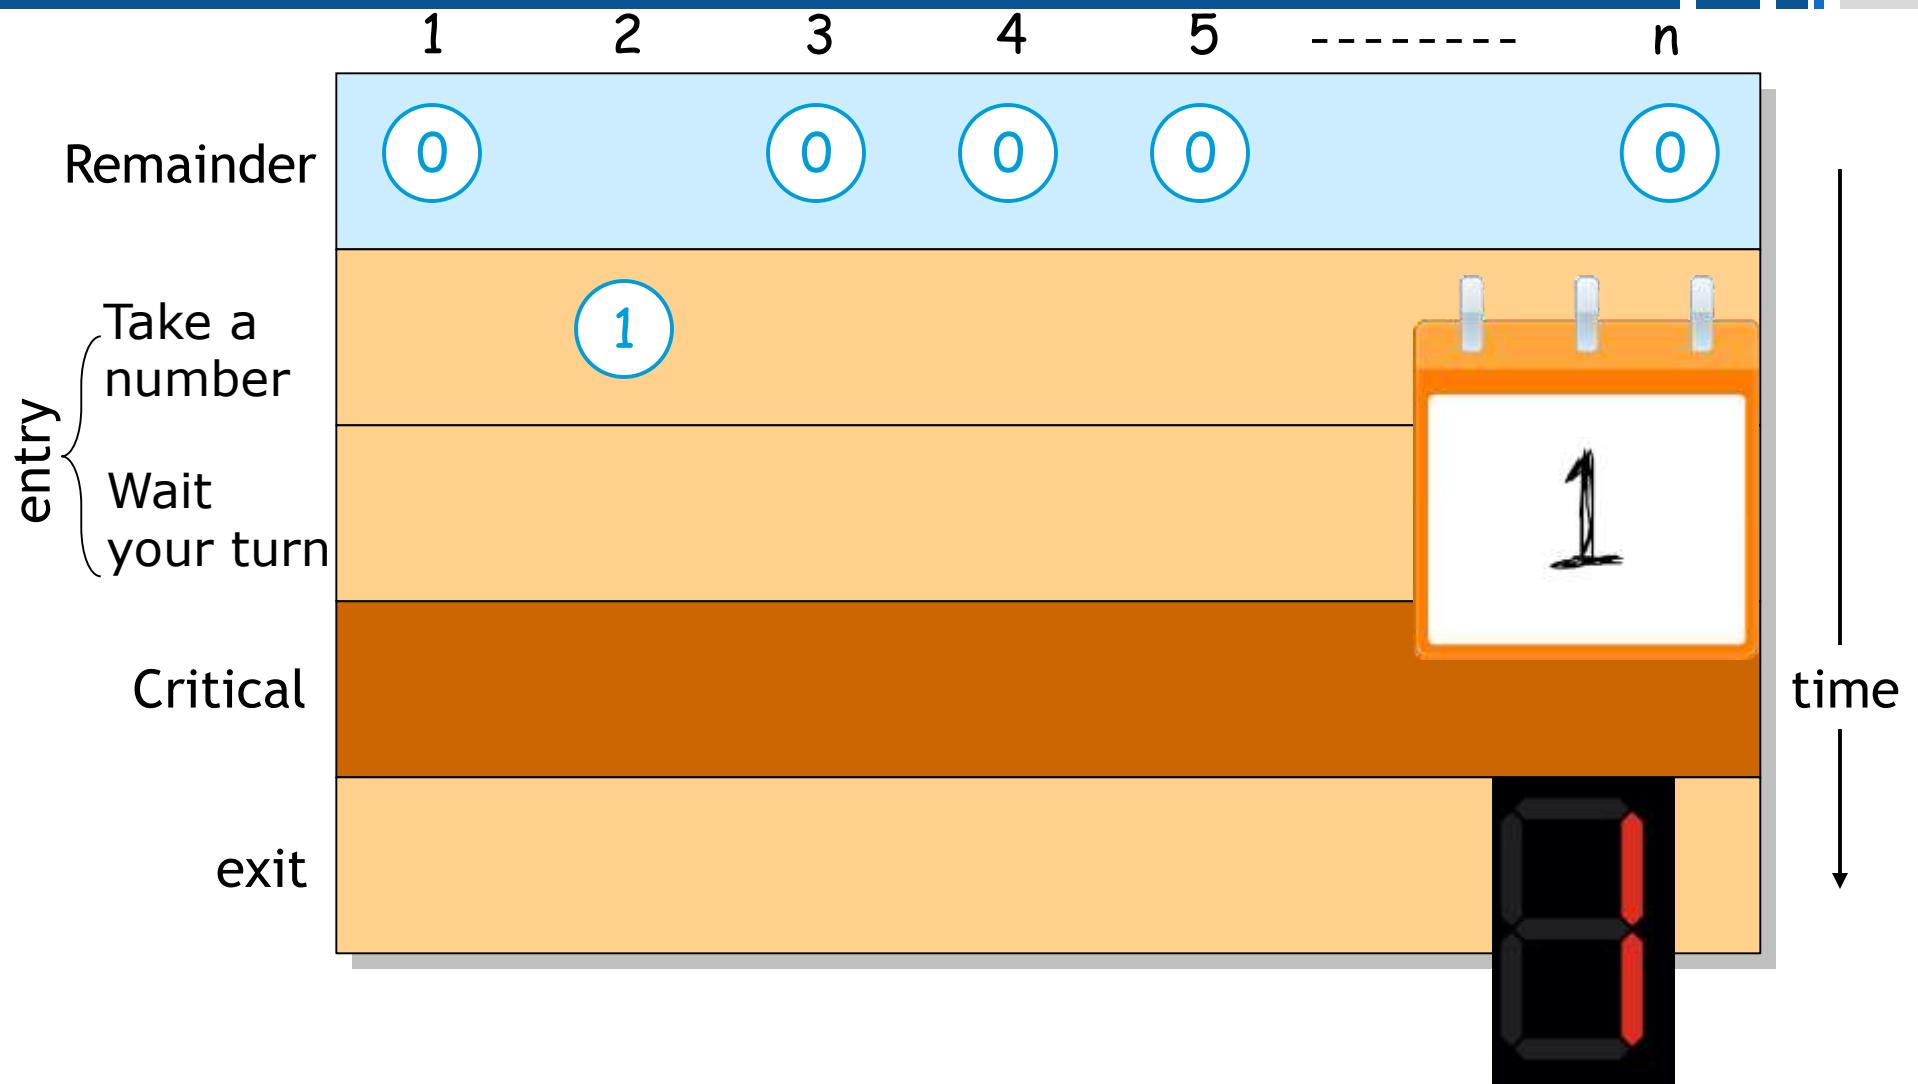
# Lamport's Bakery Algorithm - 1974



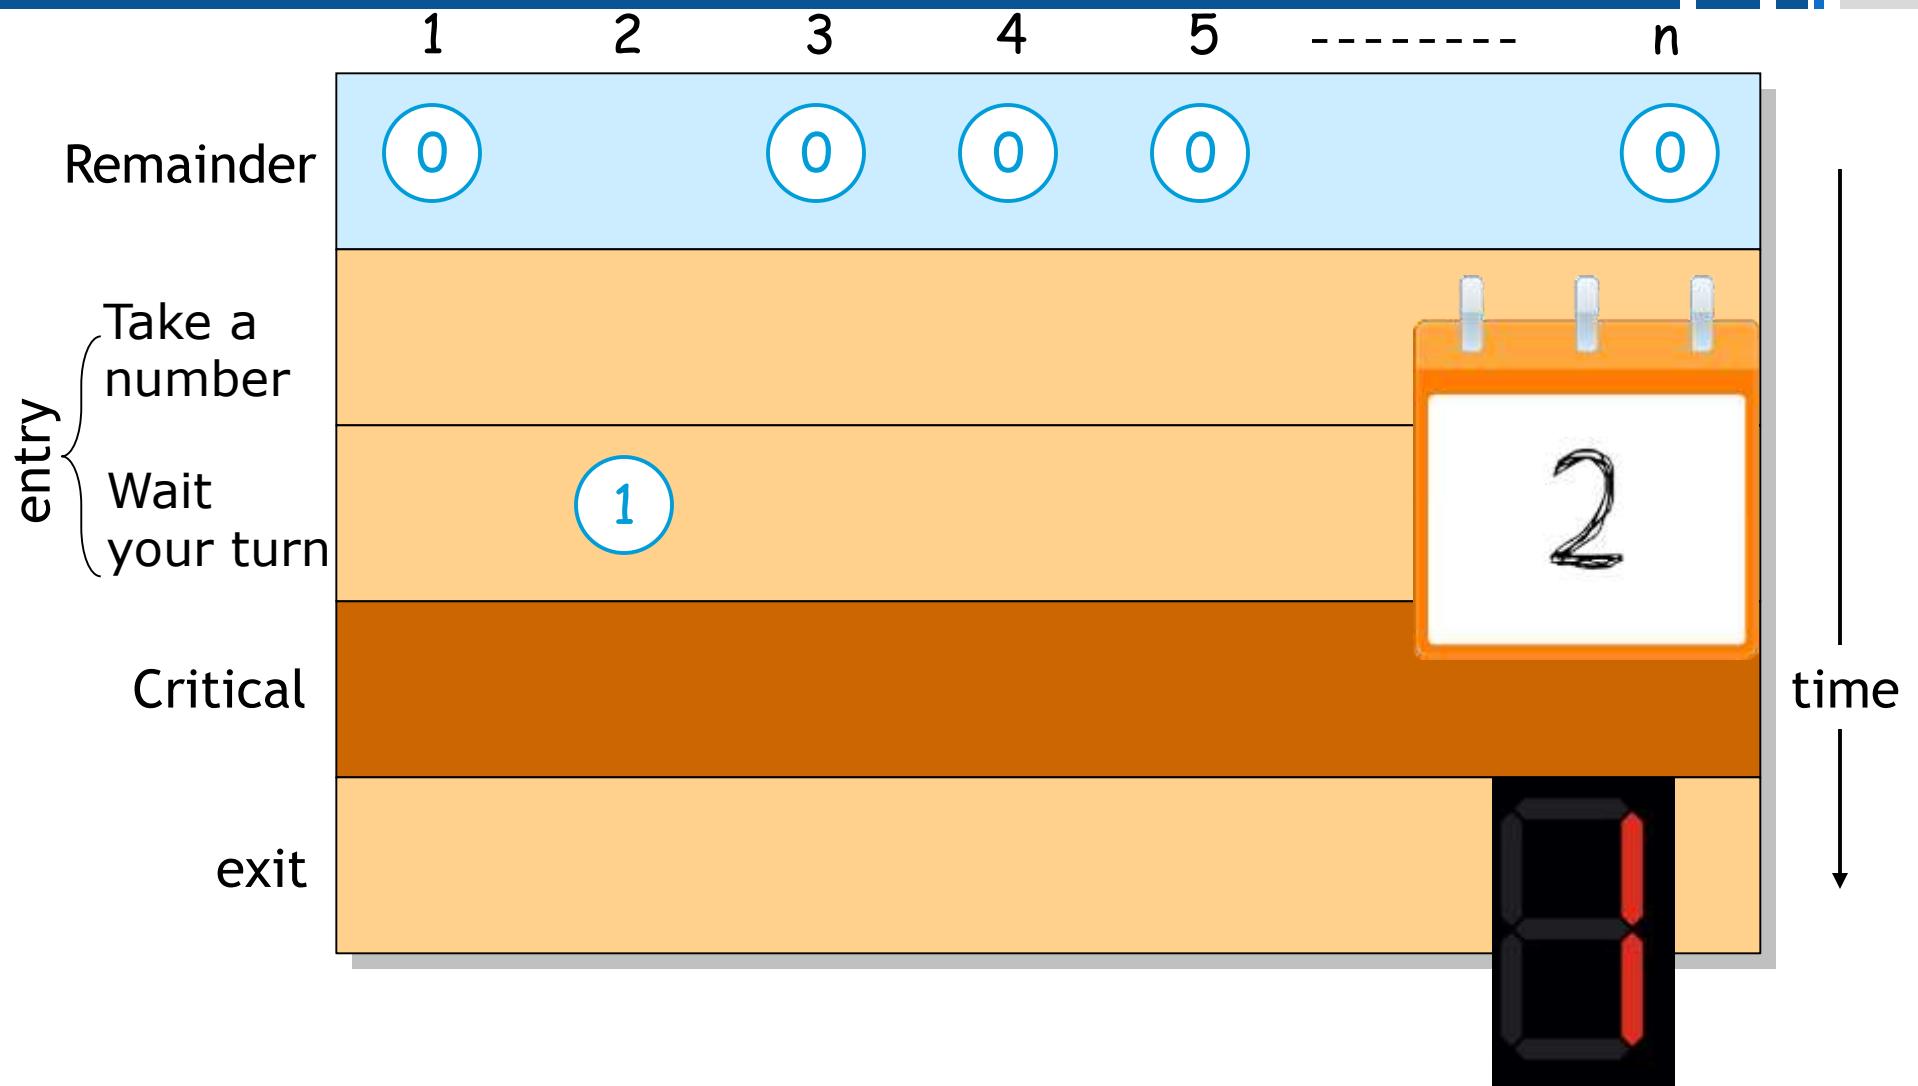
# Not quite...



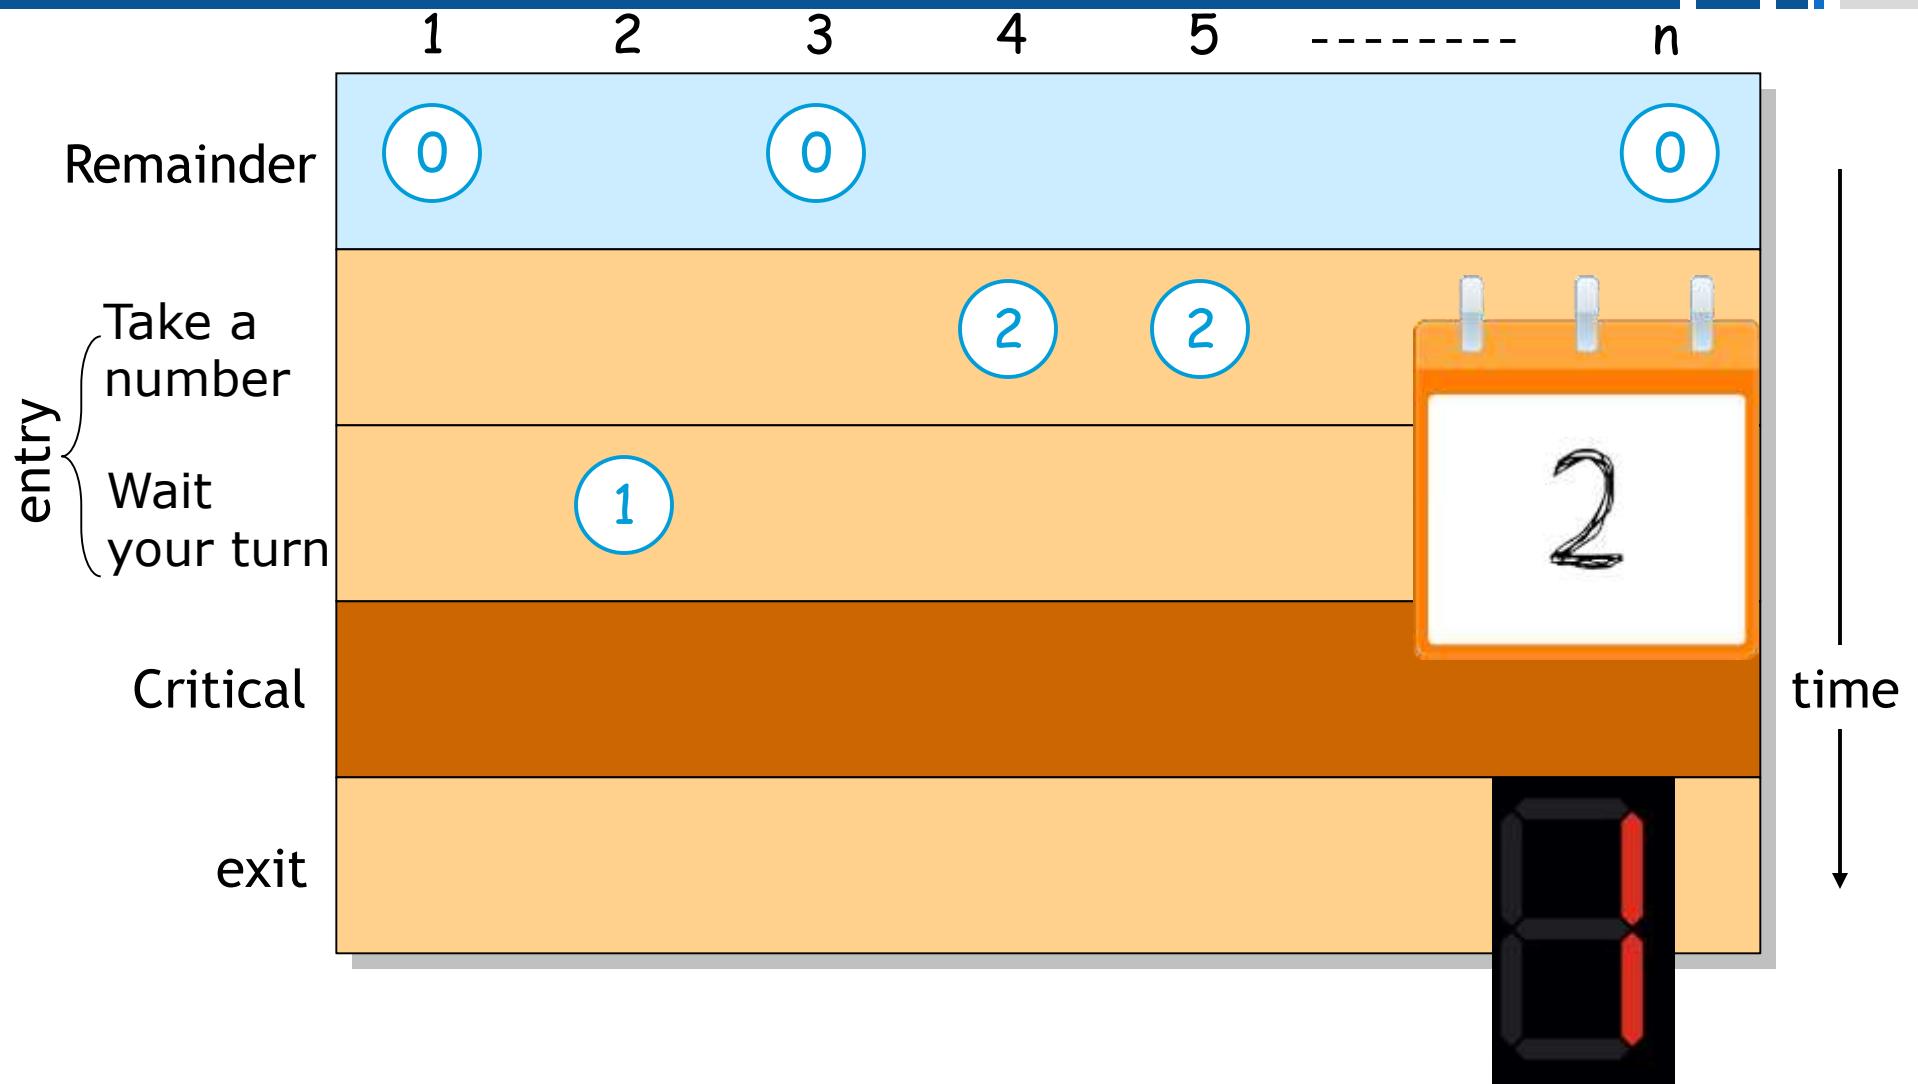
# Not quite...



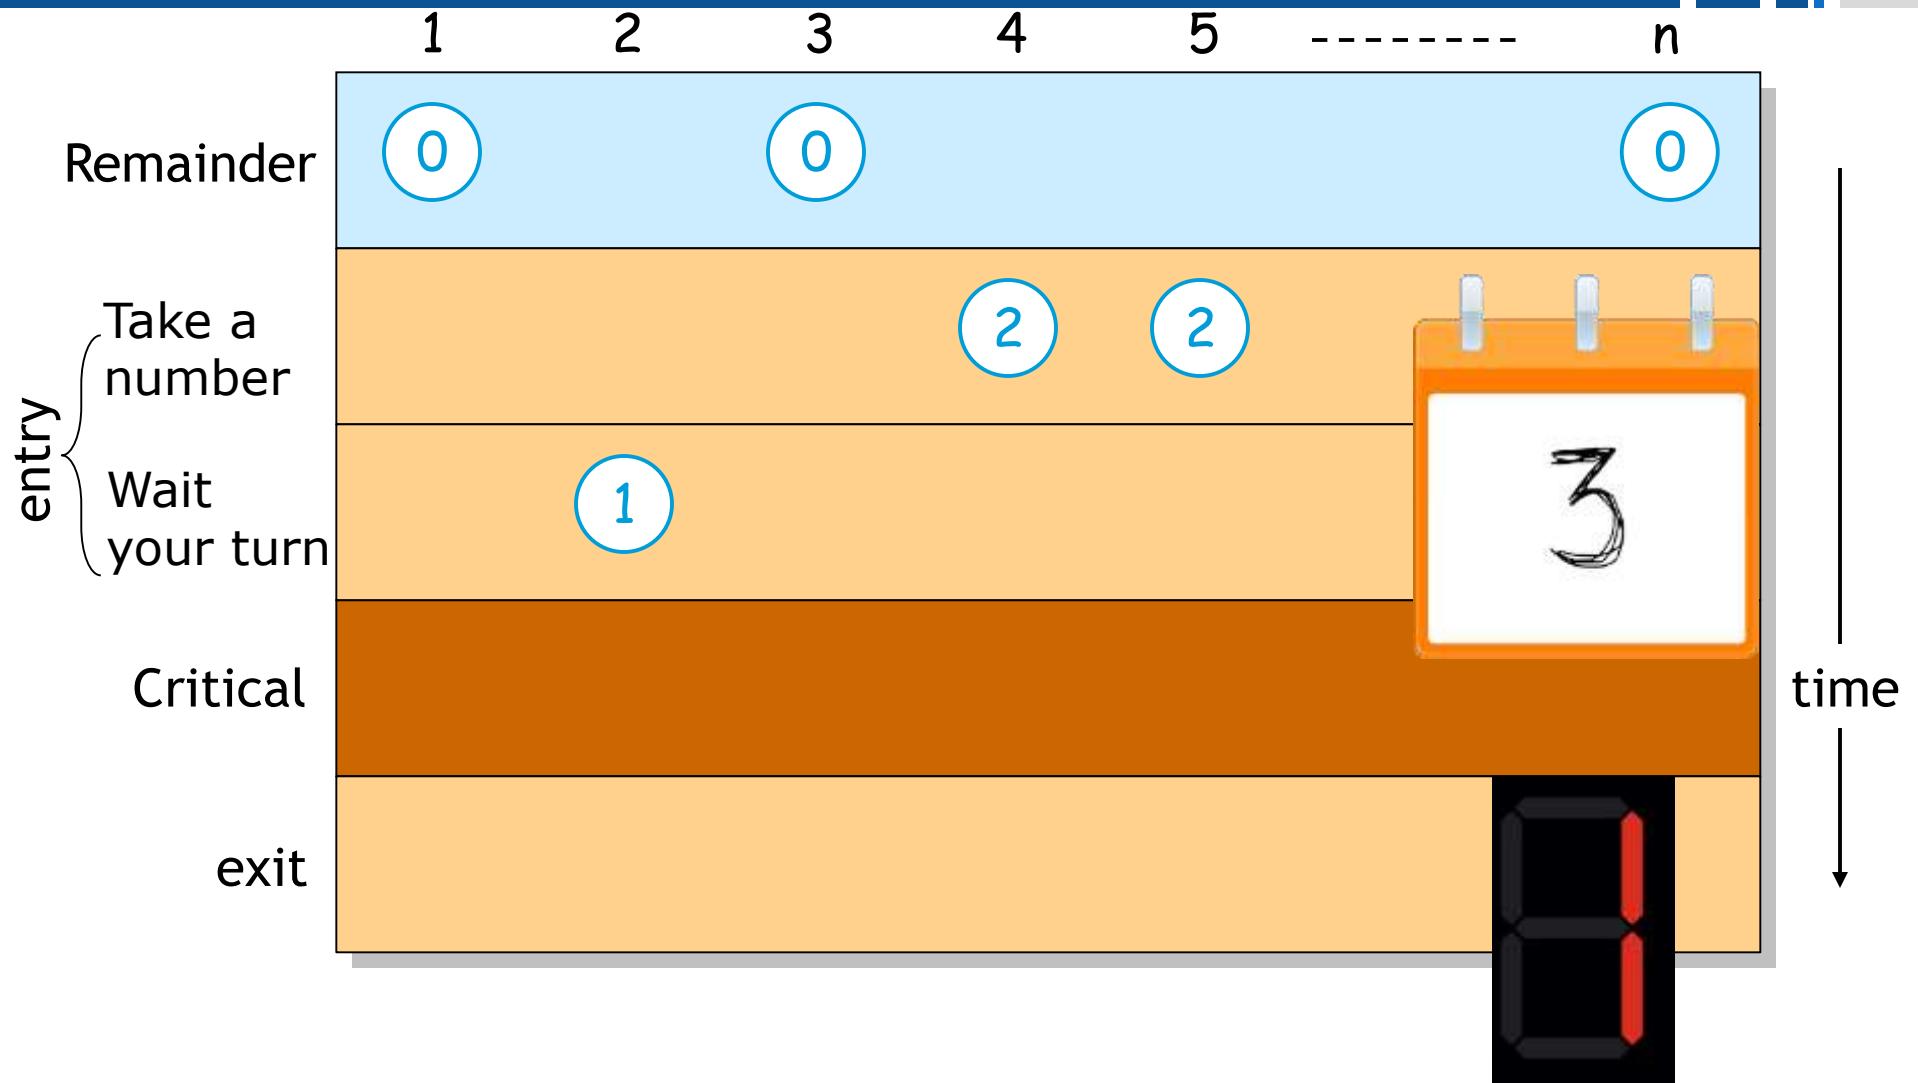
# Not quite...



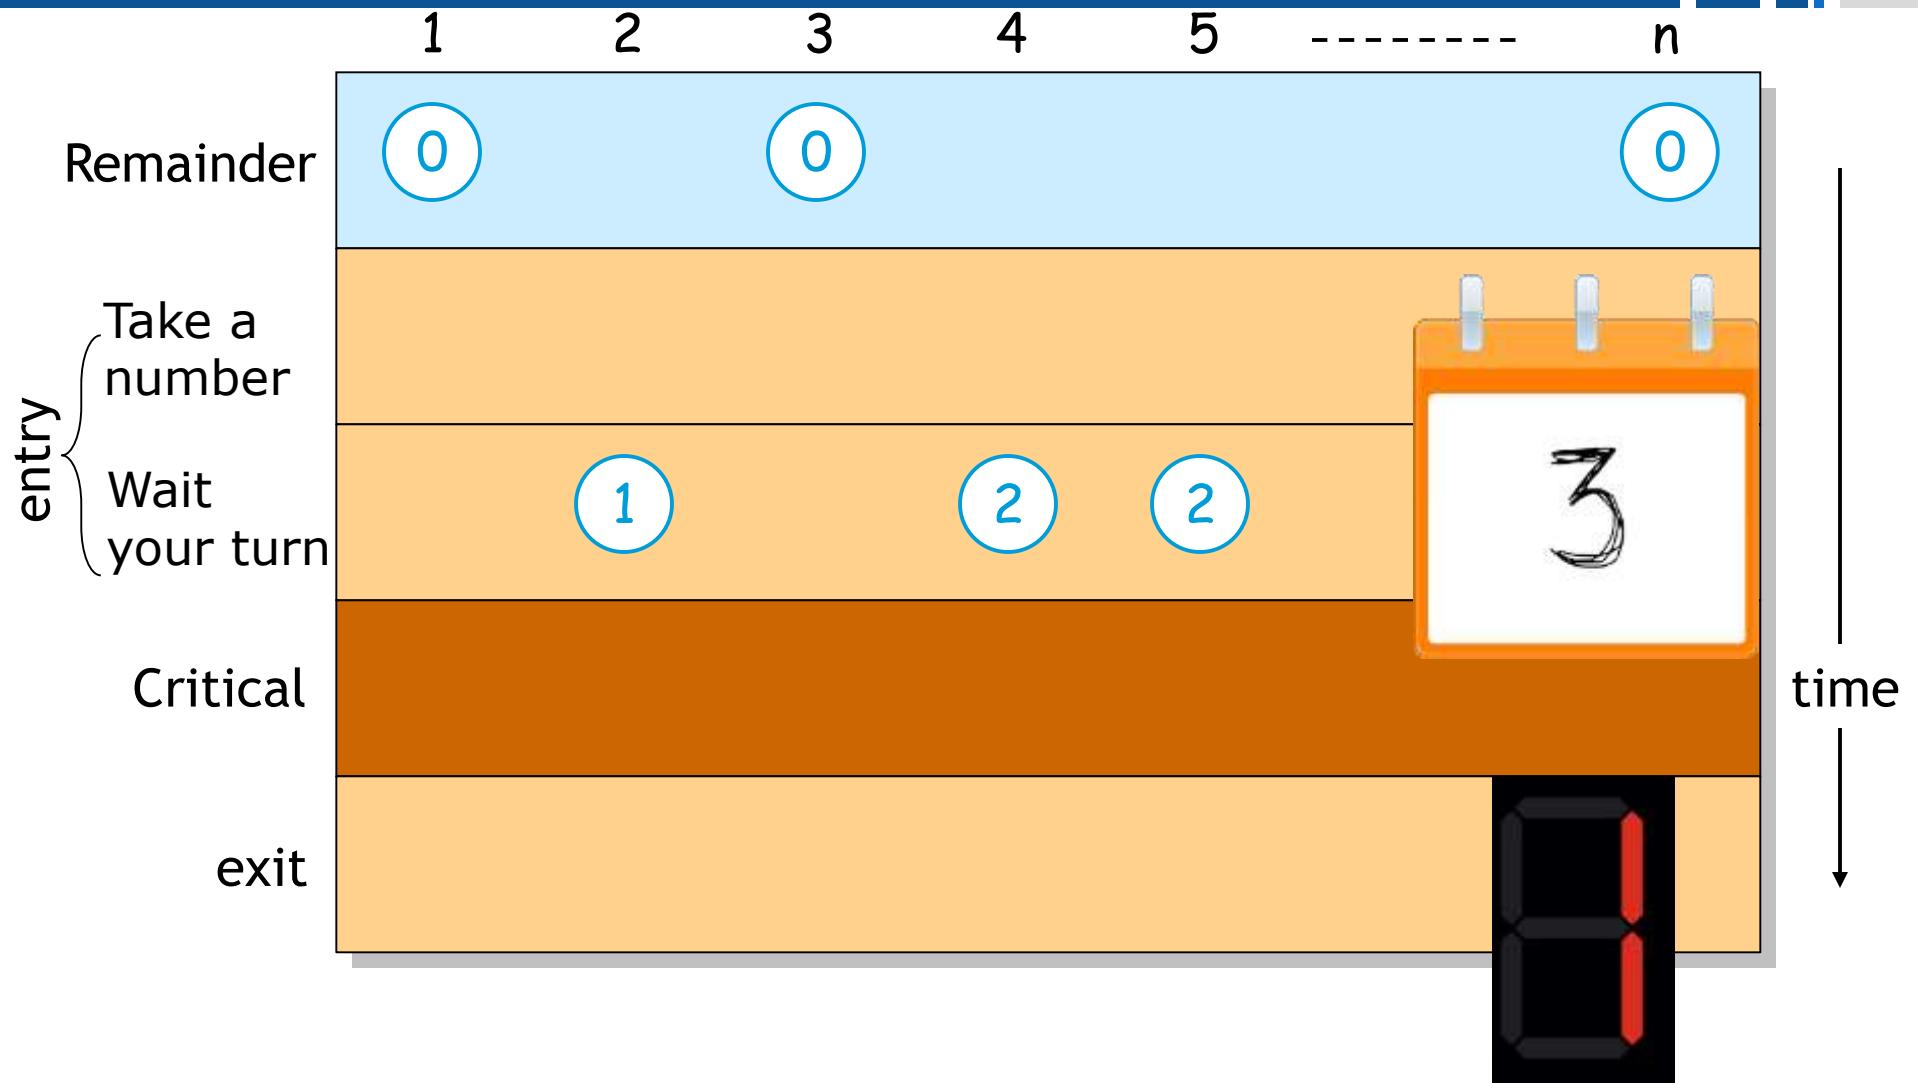
# Not quite...



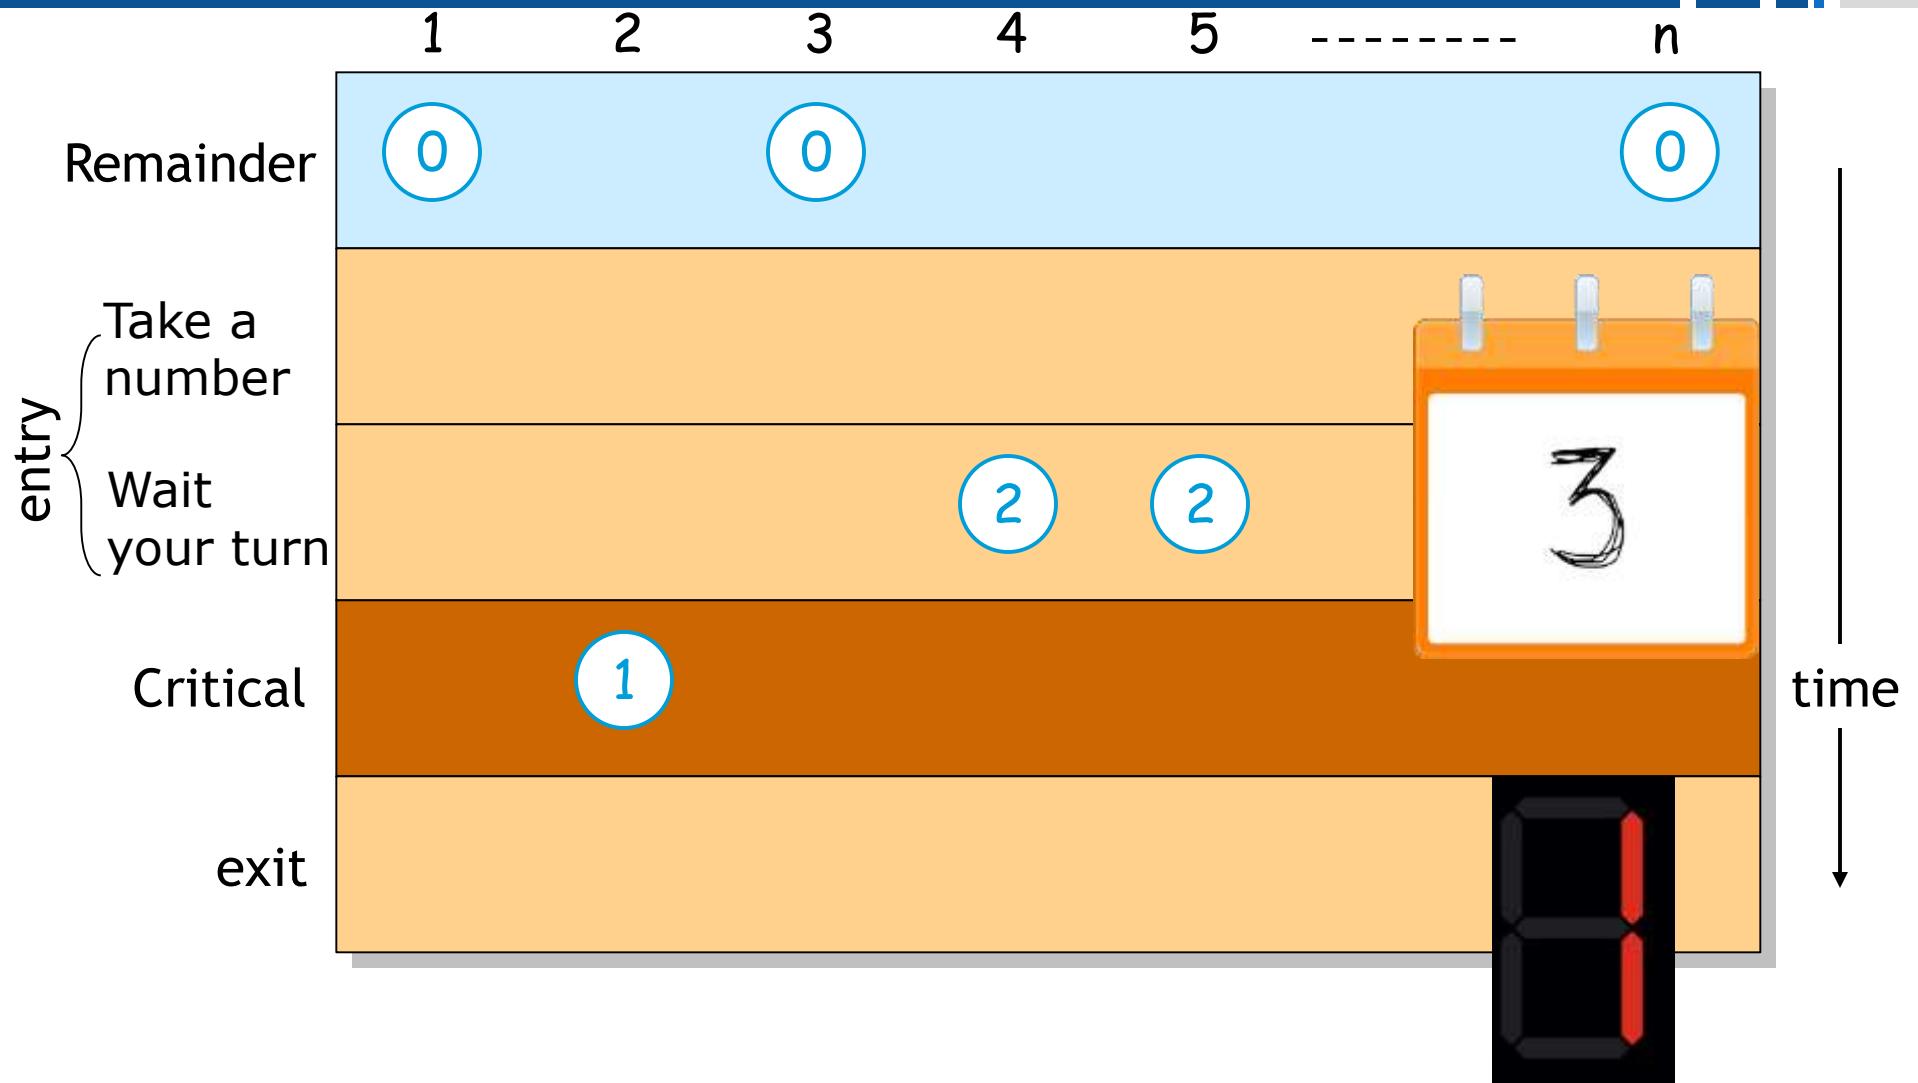
# Not quite...



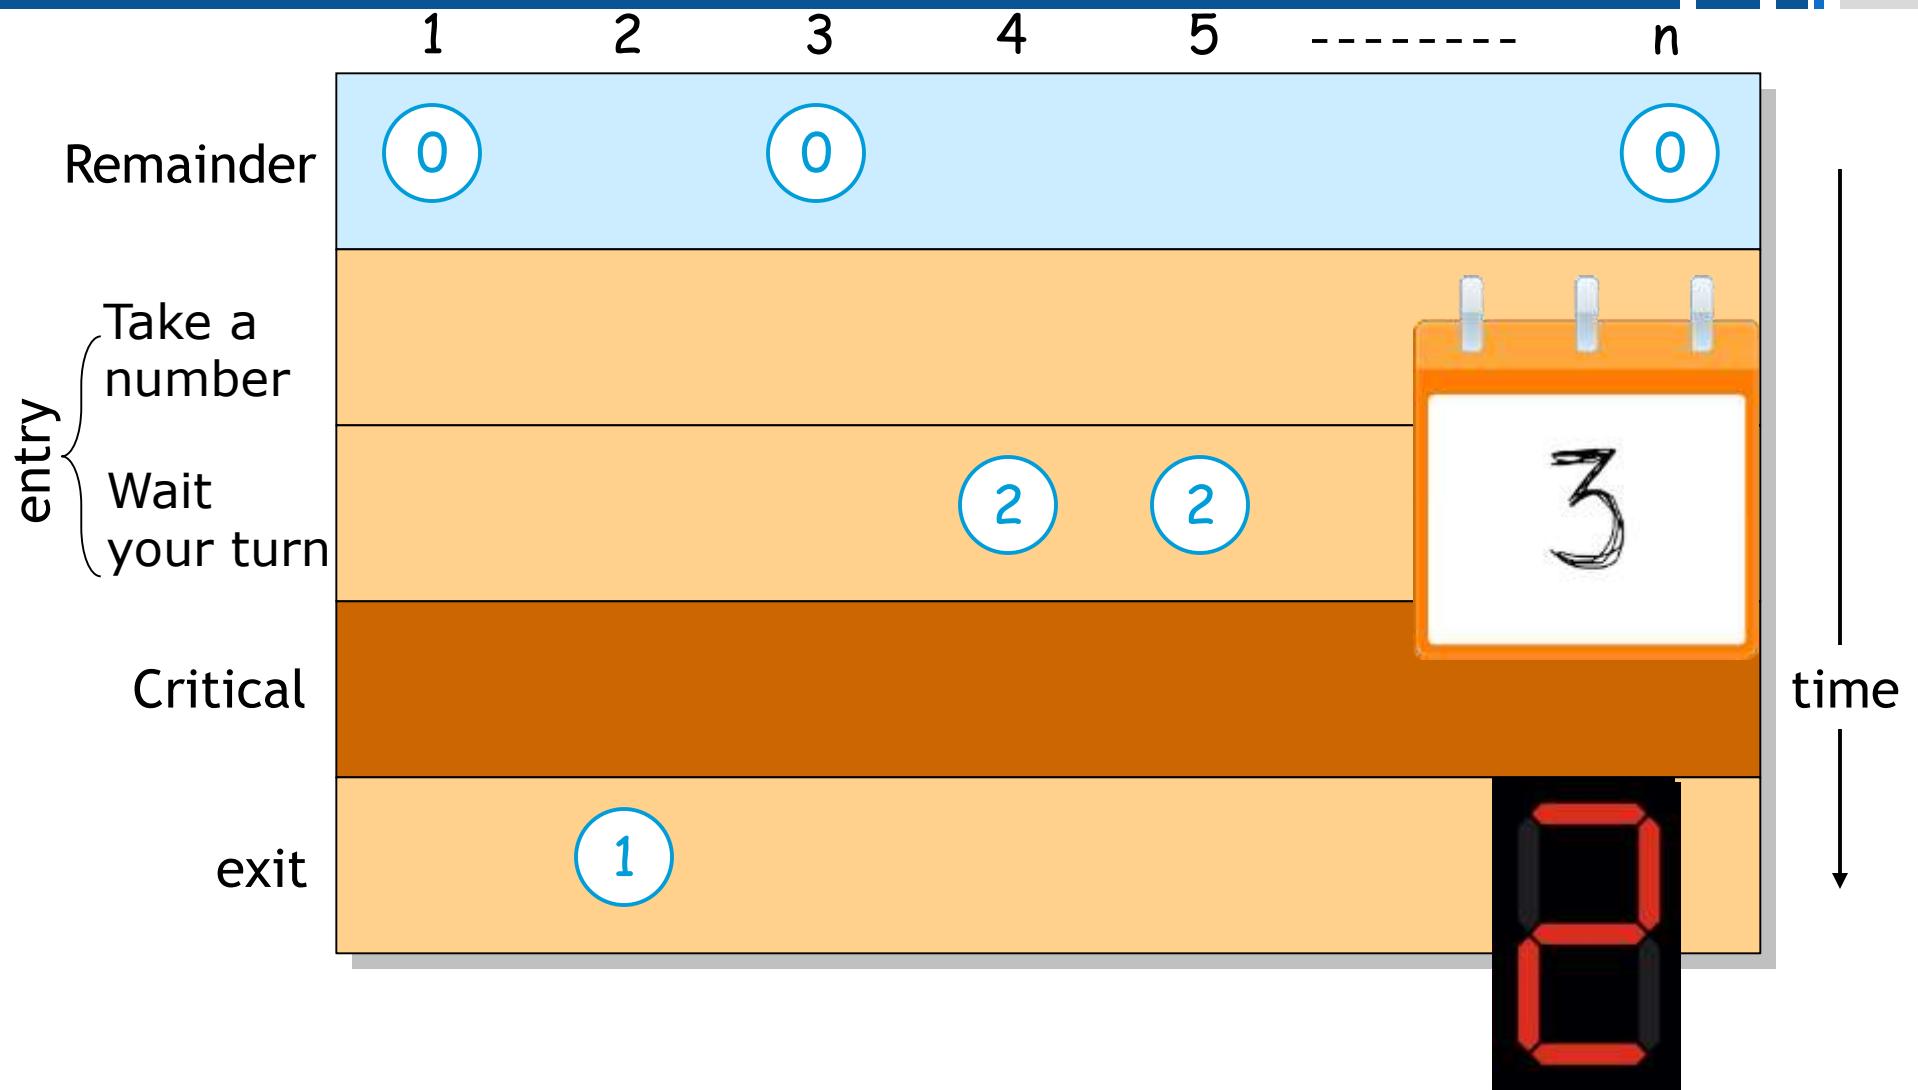
# Not quite...



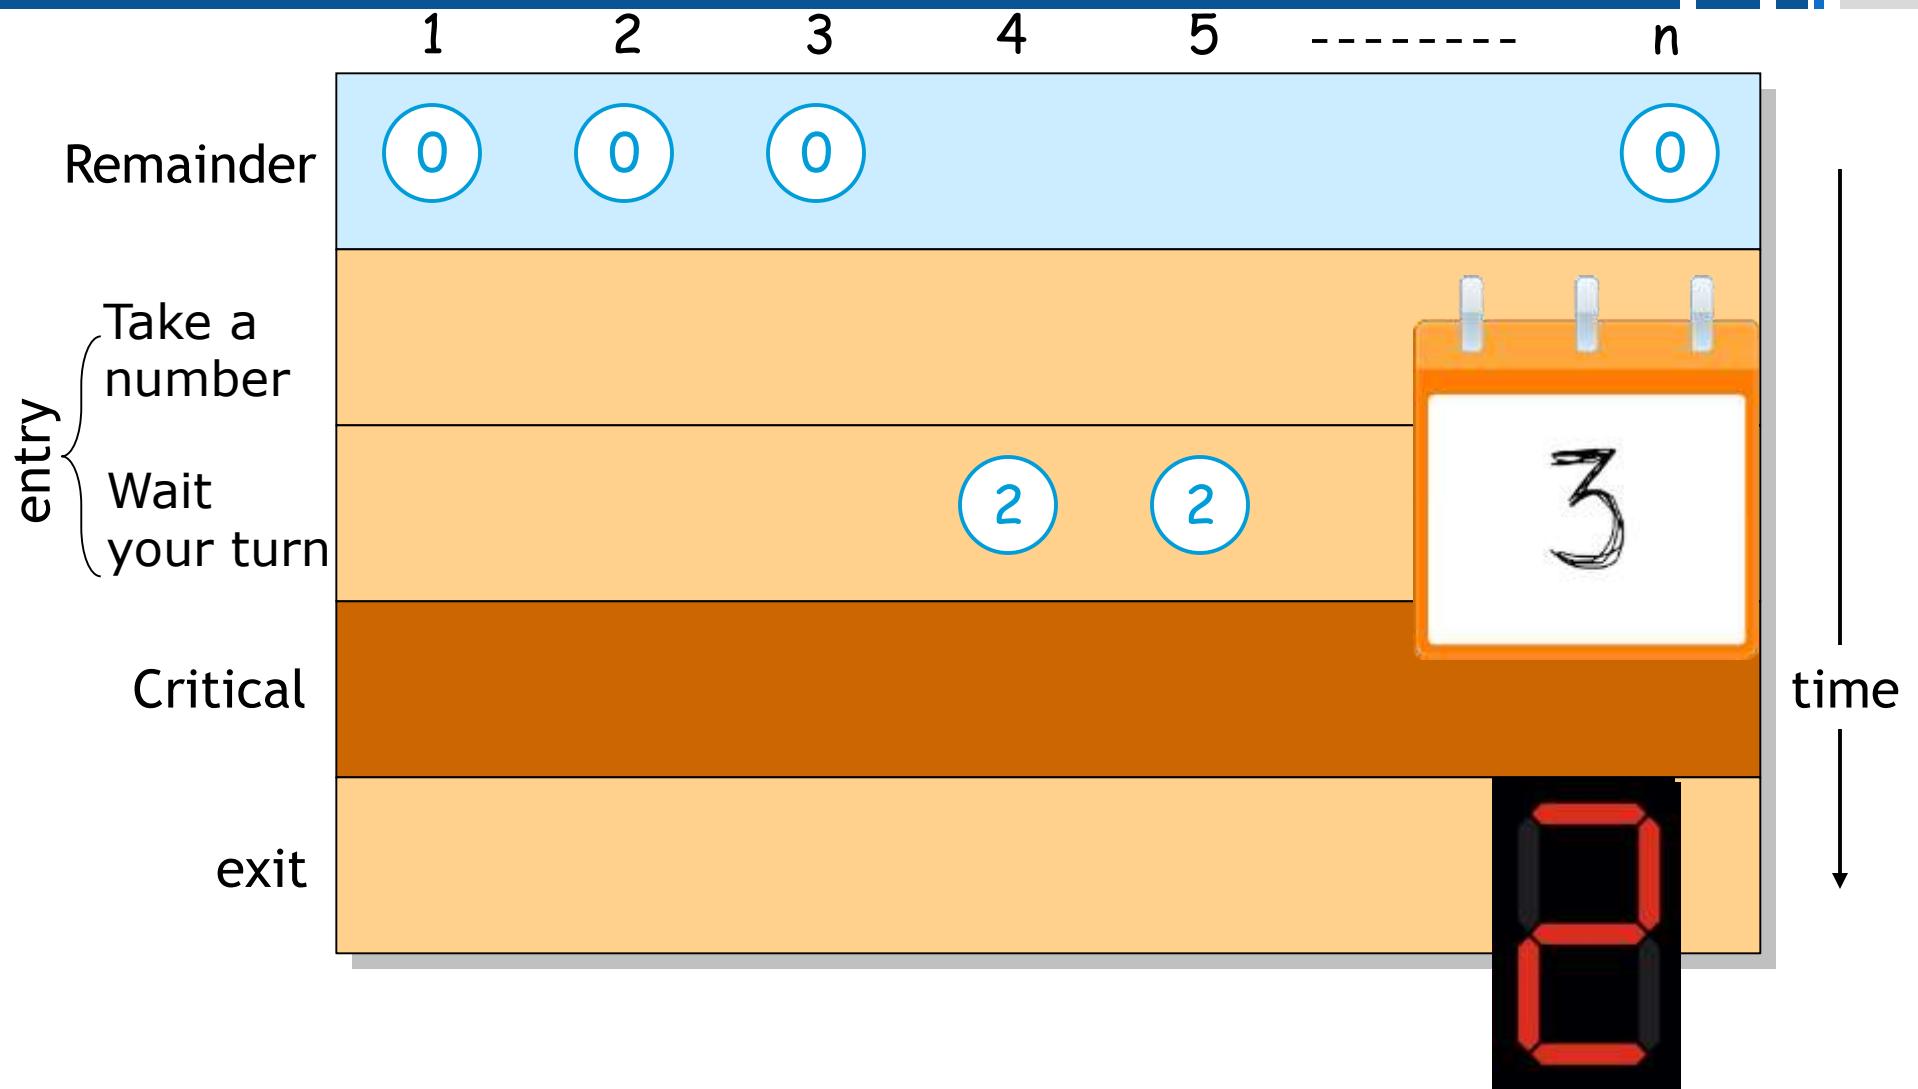
# Not quite...



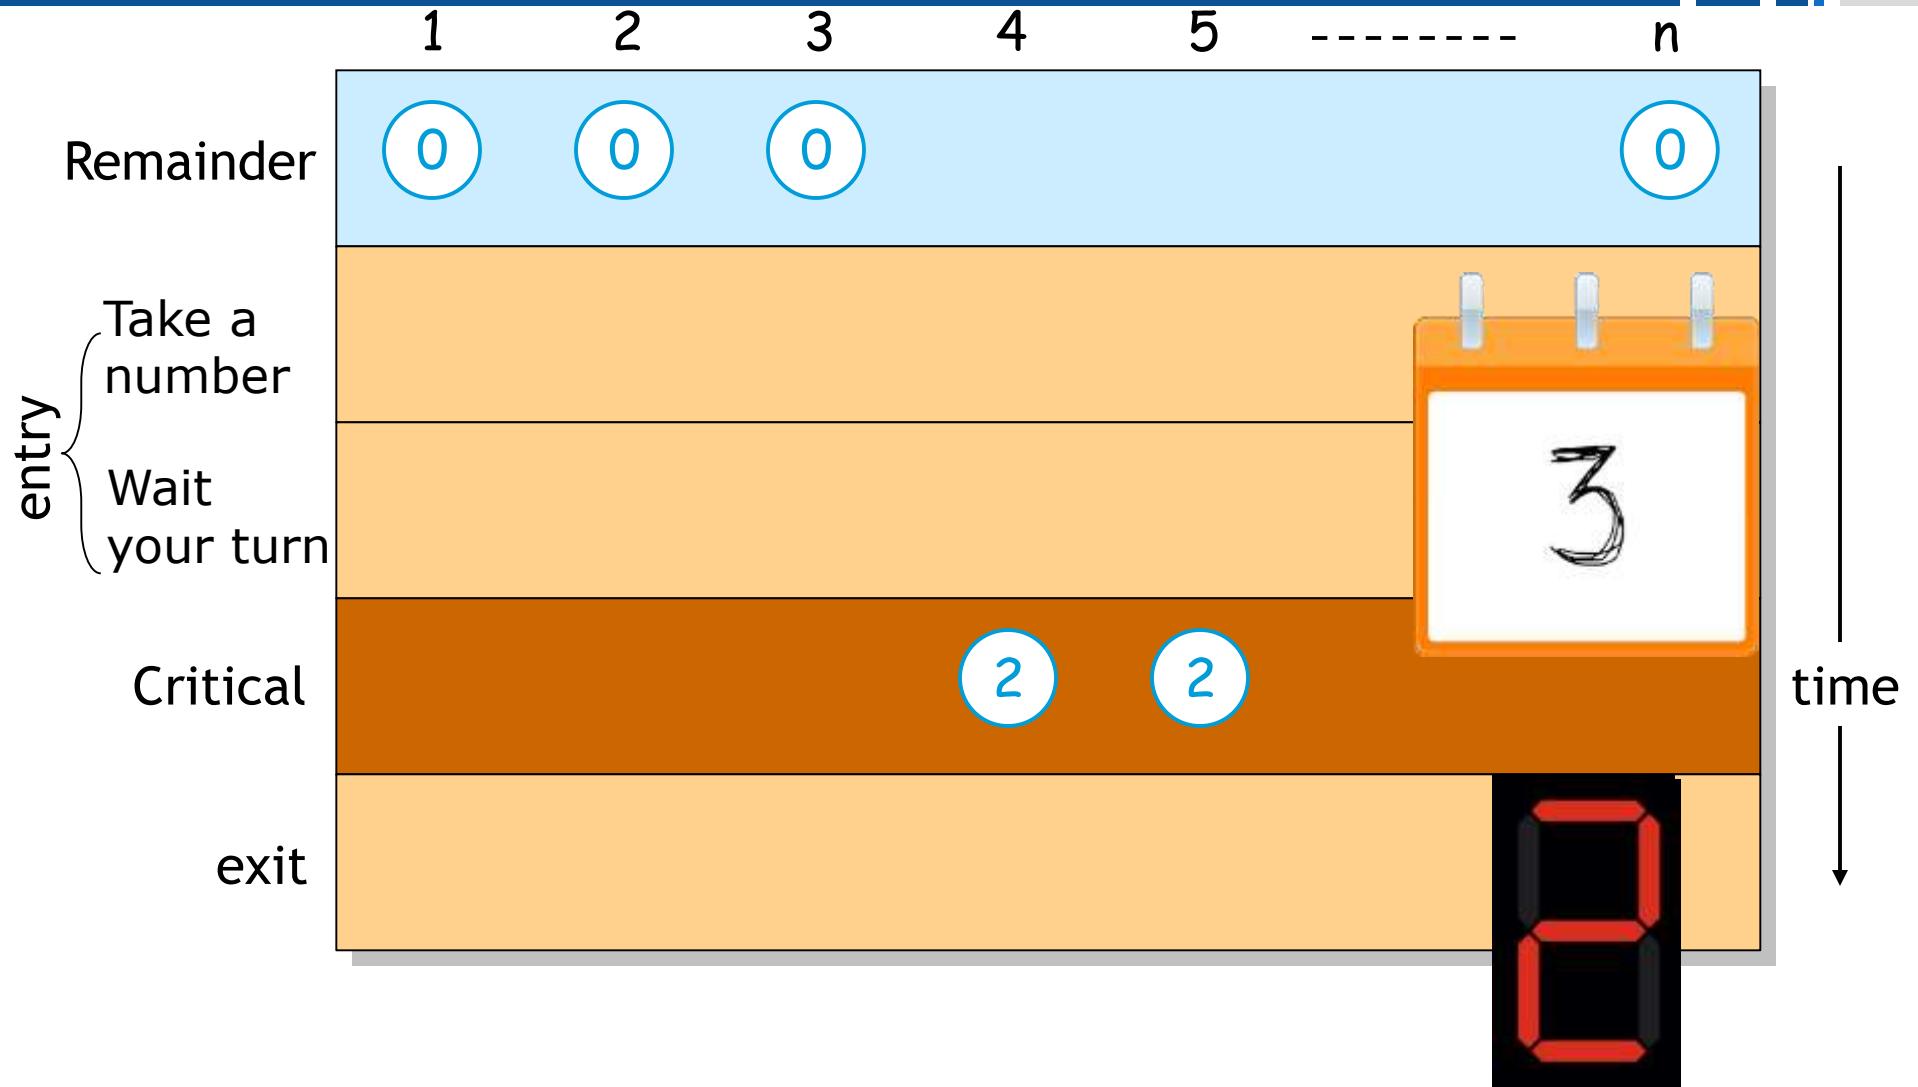
# Not quite...



# Not quite...



# Not quite...



# Implementation 1

Remainder

```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[j]>0 and  
          number[j]<number[i]);  
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

# Implementation 1

Reminder

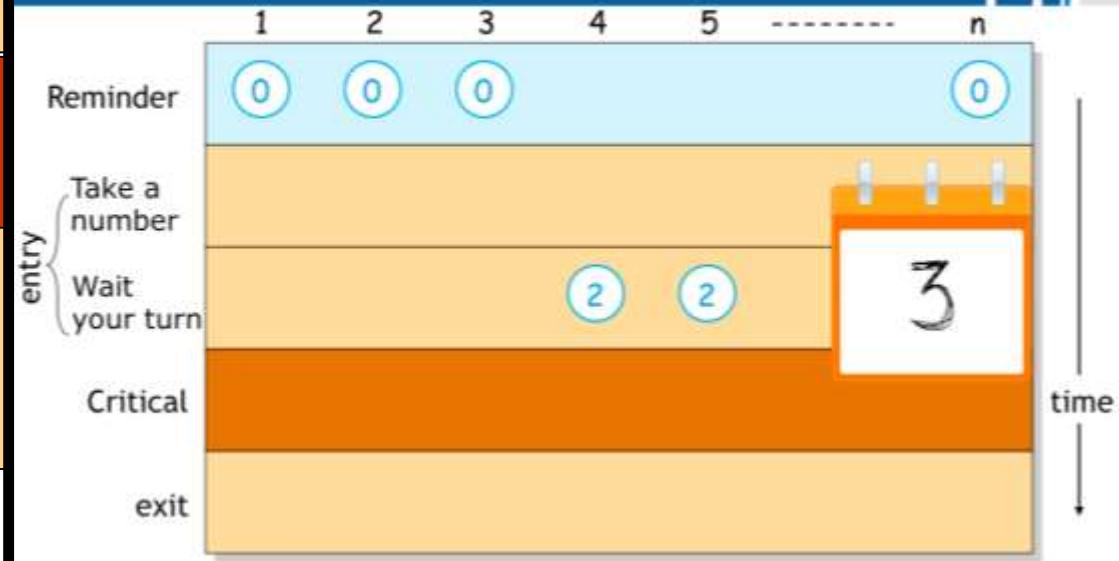
```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[j]>0 and  
          number[j]<number[i]);  
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

Lamport's Bakery Algorithm



# Implementation 1

Remainder

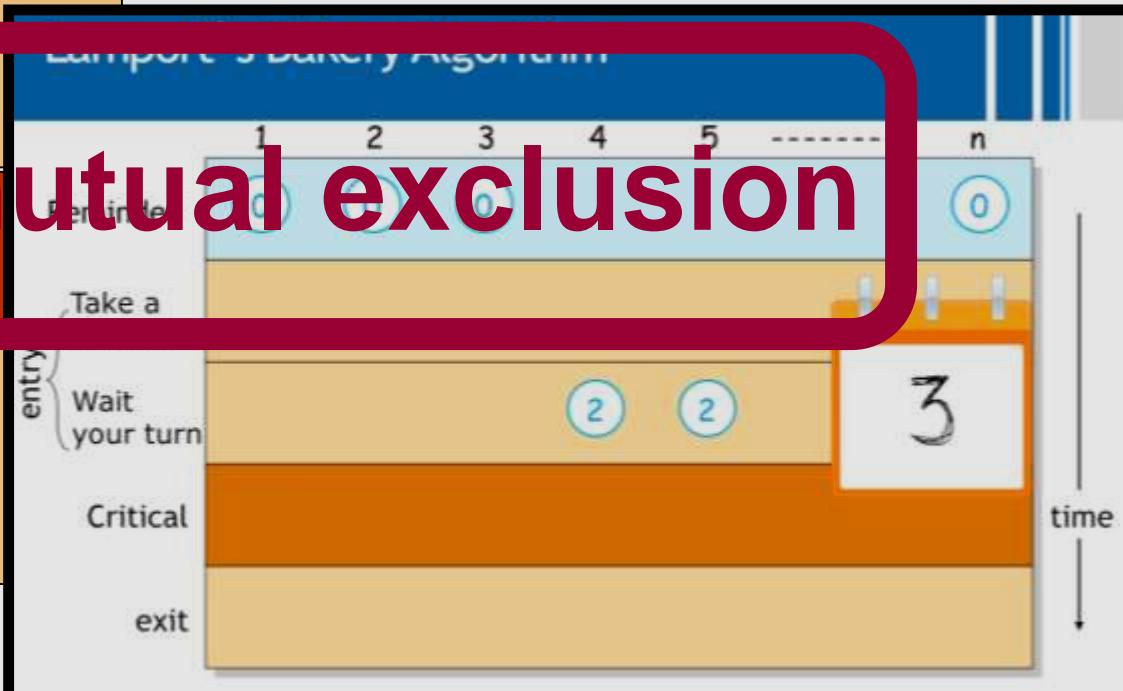
```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[i]> 0 and  
          number[j]< number[i])  
        number[j]++;  
    }  
}
```

Critical

number[i]=0;

Code for Thread *i*  
(out of *n* threads)

No mutual exclusion



# Implementation 1.1

Remainder

```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[j]>0 and  
          number[j]≤number[i]);  
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

# Implementation 1.1

Remainder

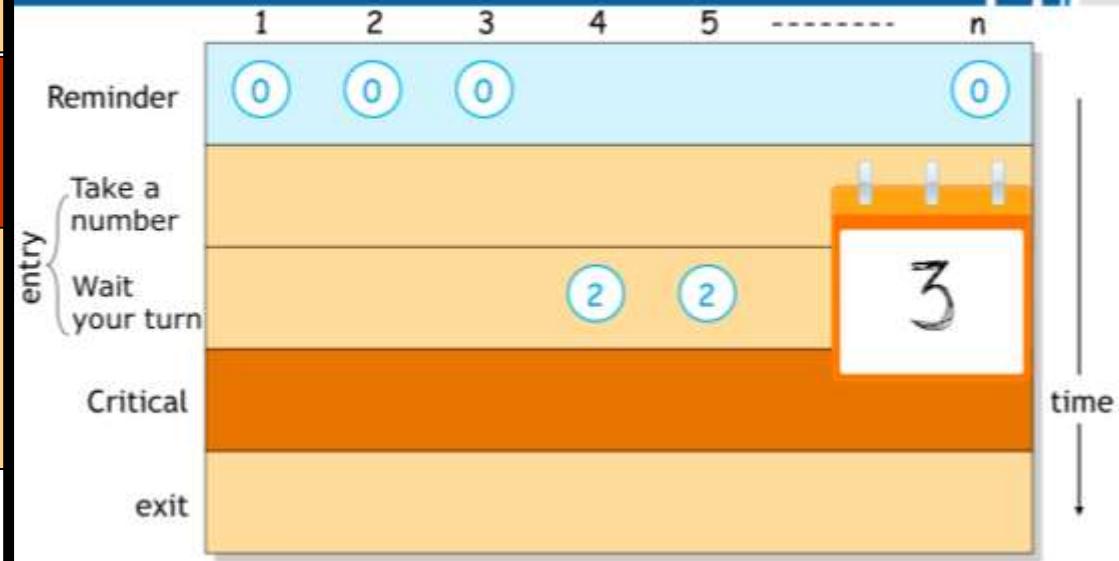
```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[j]>0 and  
          number[j]≤number[i]);  
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

Lamport's Bakery Algorithm



# Implementation 1.1

Remainder

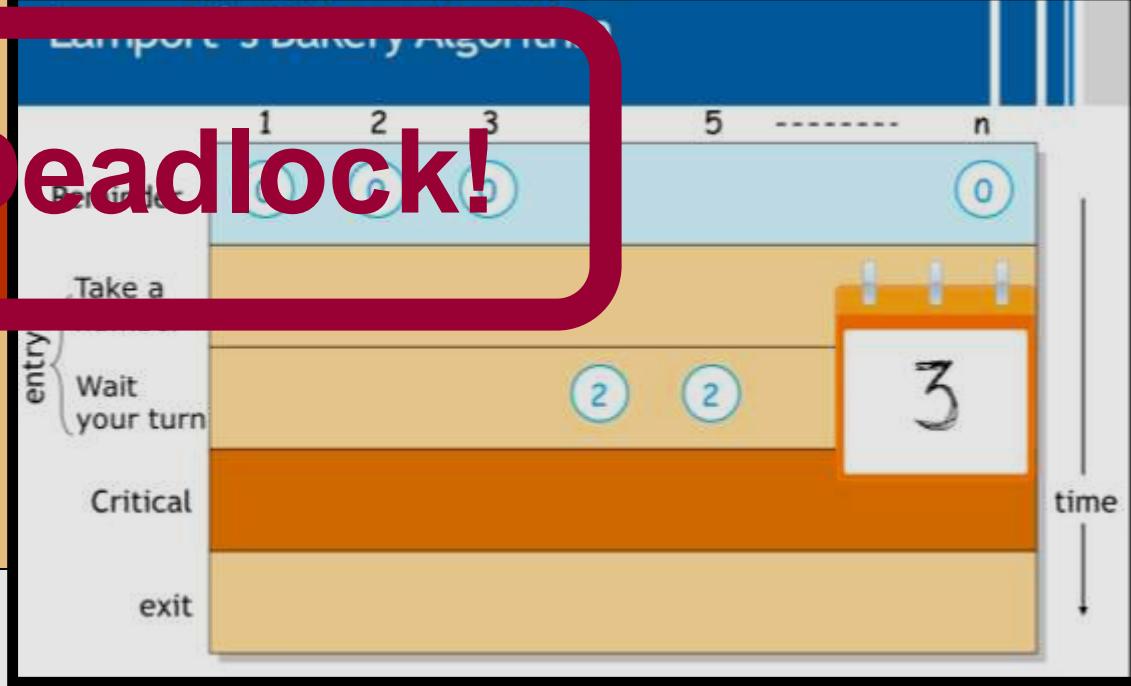
```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};  
for j=1 to n {  
    while(i≠j and number[j]>0 and  
          number[j]<=number[i]);  
}  
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

Deadlock!



# Implementation 2

Remainder

```
number[i] = 1+maxj ∈ {1,...,n}{number[j}};

for j=1 to n {
    while(i≠j and number[j]>0 and
          (number[j],j)<(number[i],i));
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

# Implementation 2

Remainder

```
number[i] = 1+maxj ∈ {1,...,n}{number[j}};

for j=1 to n {
    while(i≠j and number[j]>0 and
        (number[j],j)<(number[i],i));
}
```

Lexicographic  
order

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

# Implementation 2

Remainder

- Read all values
- Compute maximum
- Store value

```
number[i] = 1+maxj ∈ {1,...,n}{number[j]};
```

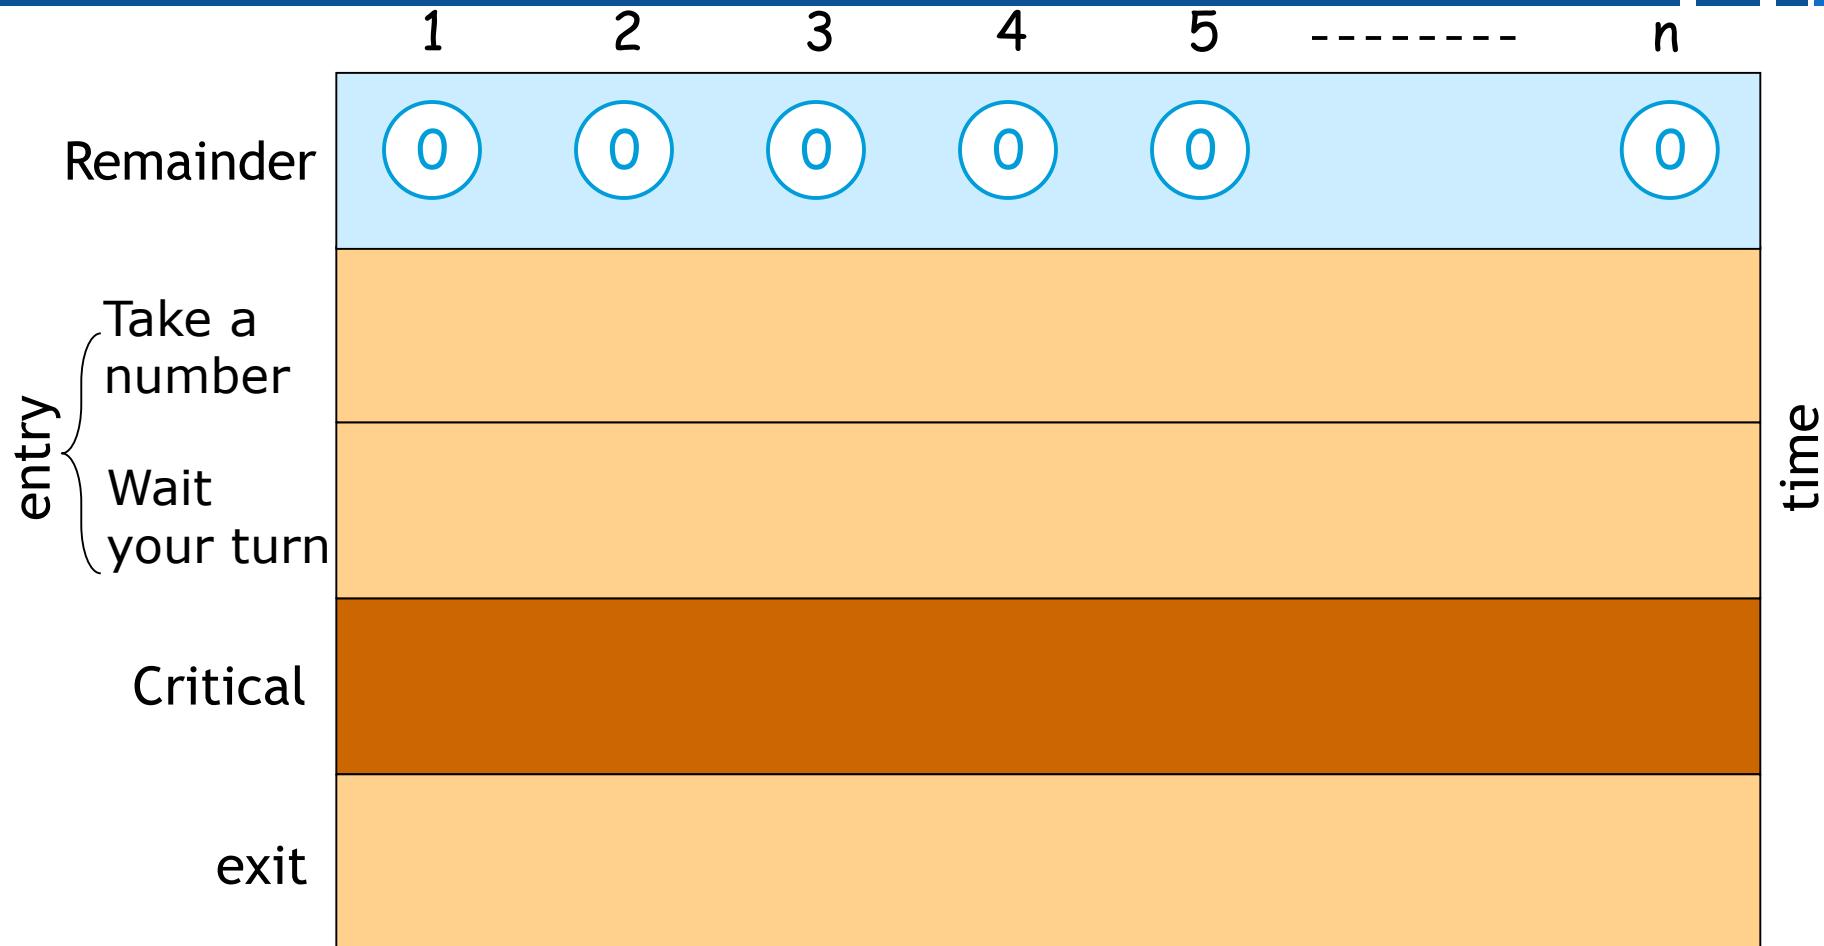
```
for j=1 to n {  
    while(i≠j and number[j]>0 and  
        (number[j],j)<(number[i],i));  
}
```

Lexicographic  
order

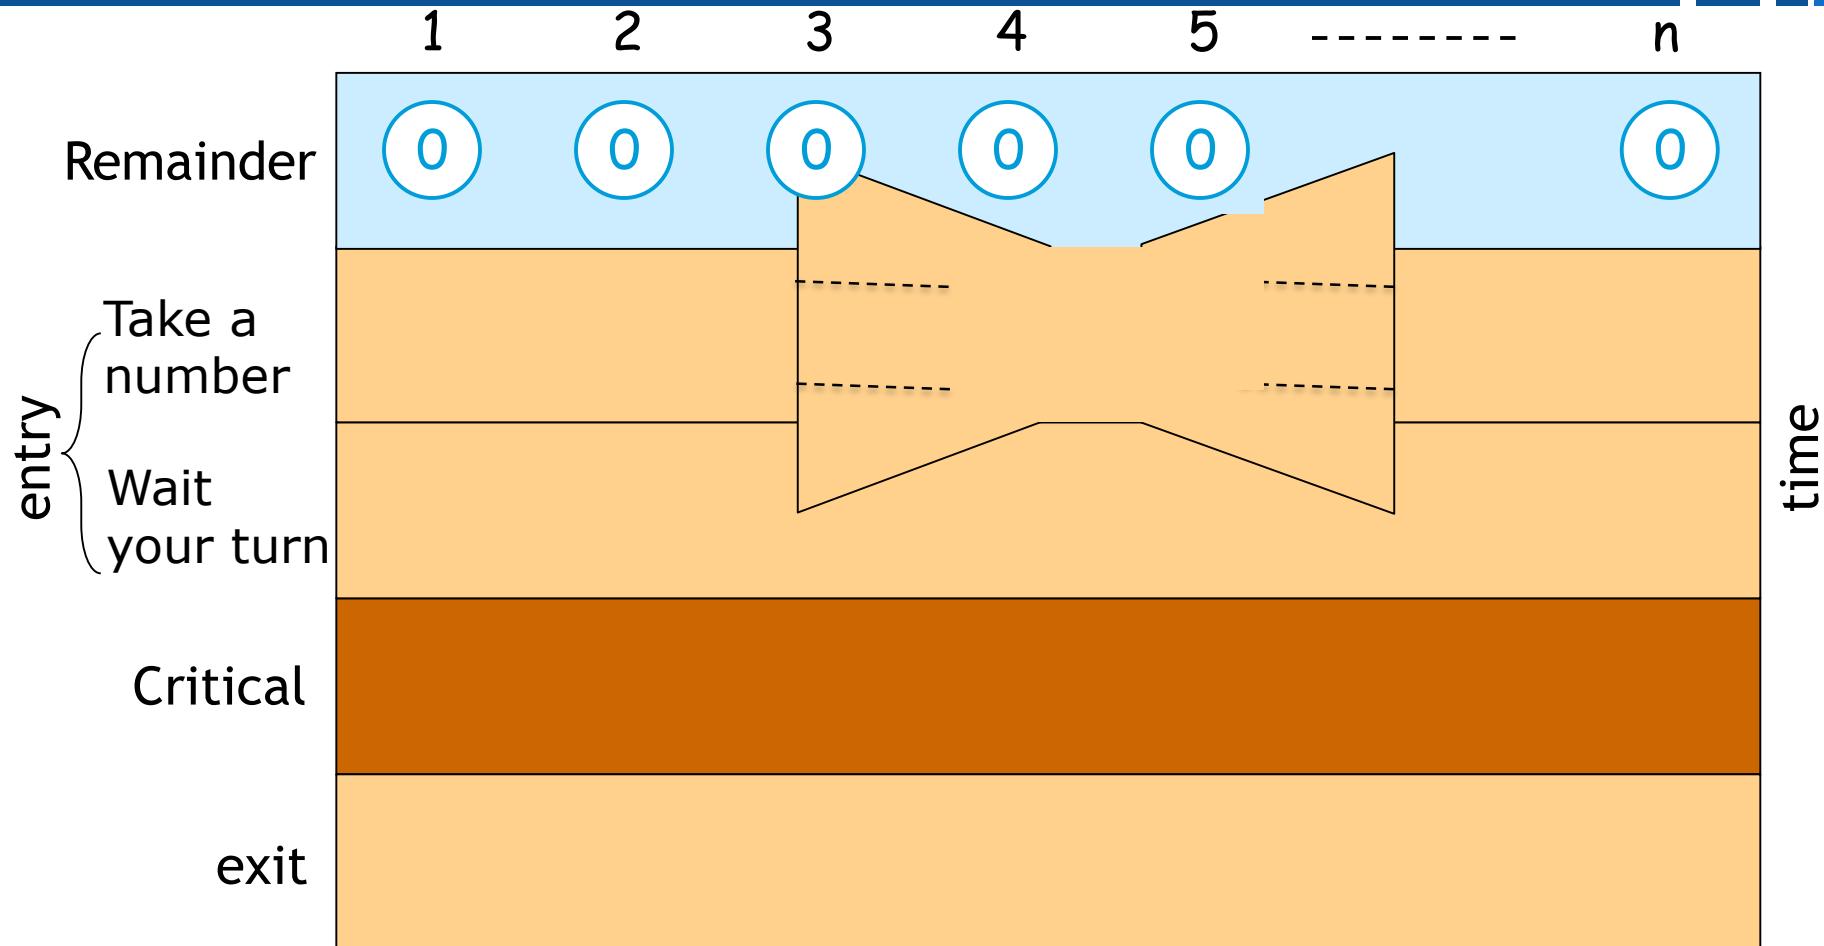
```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

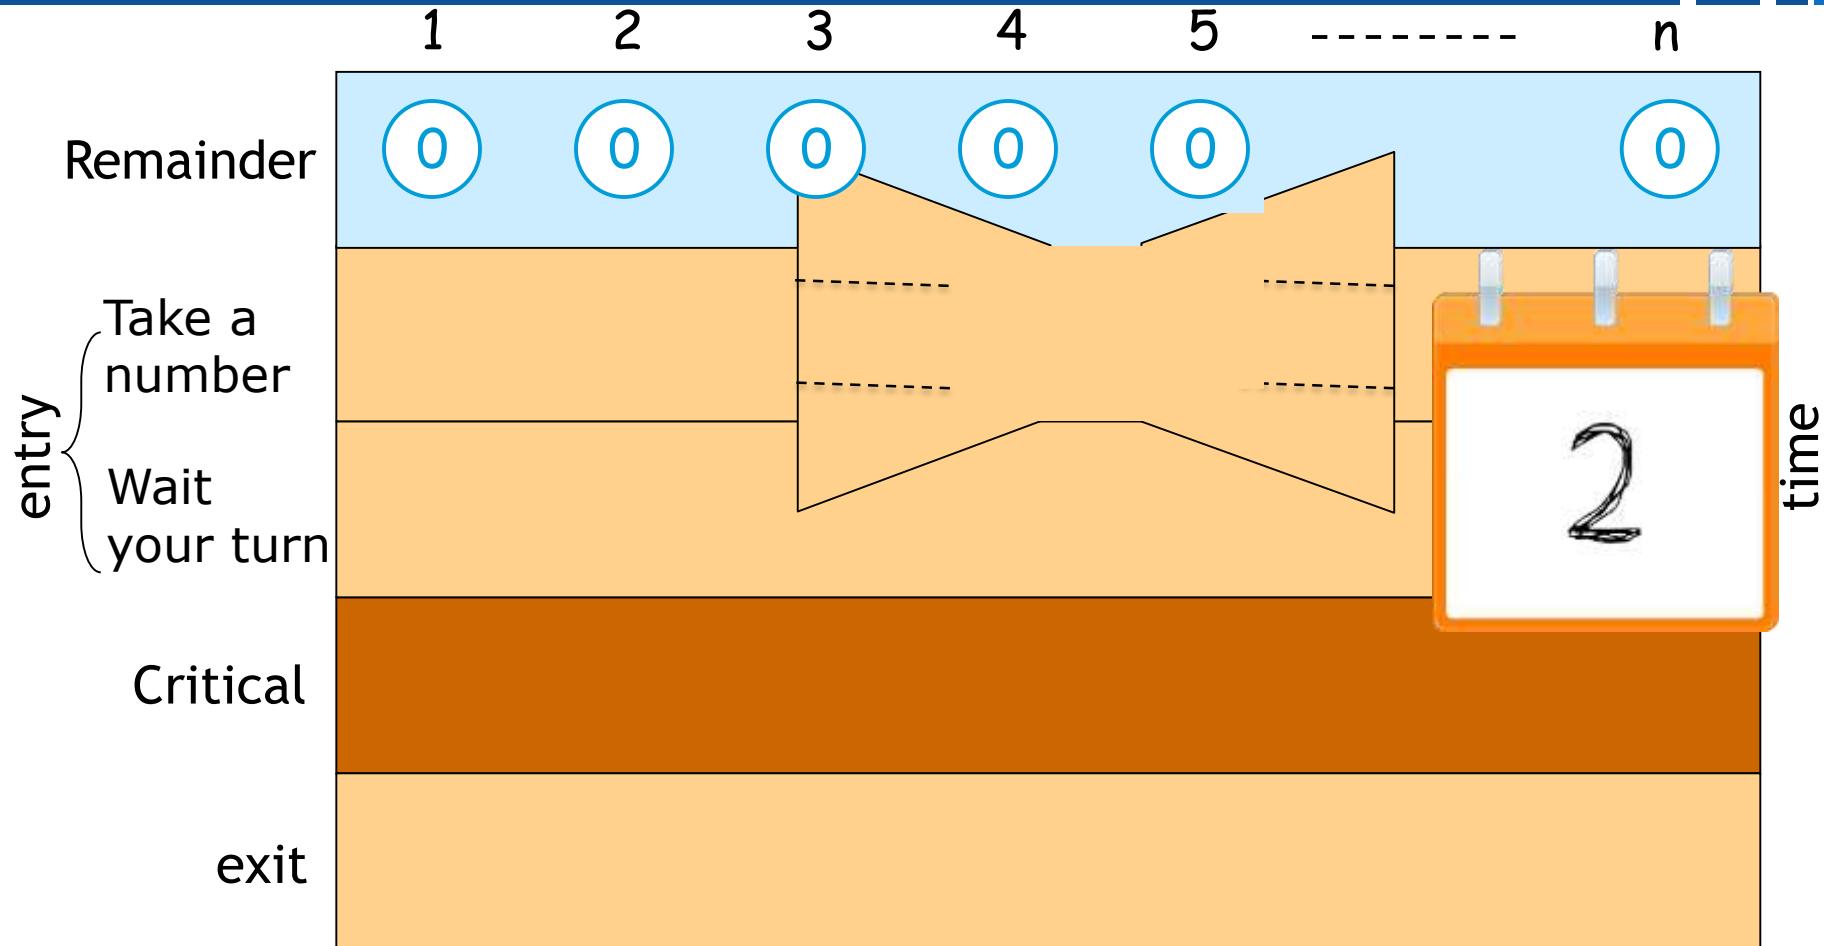
# Bakery Algorithm - Implementation 2



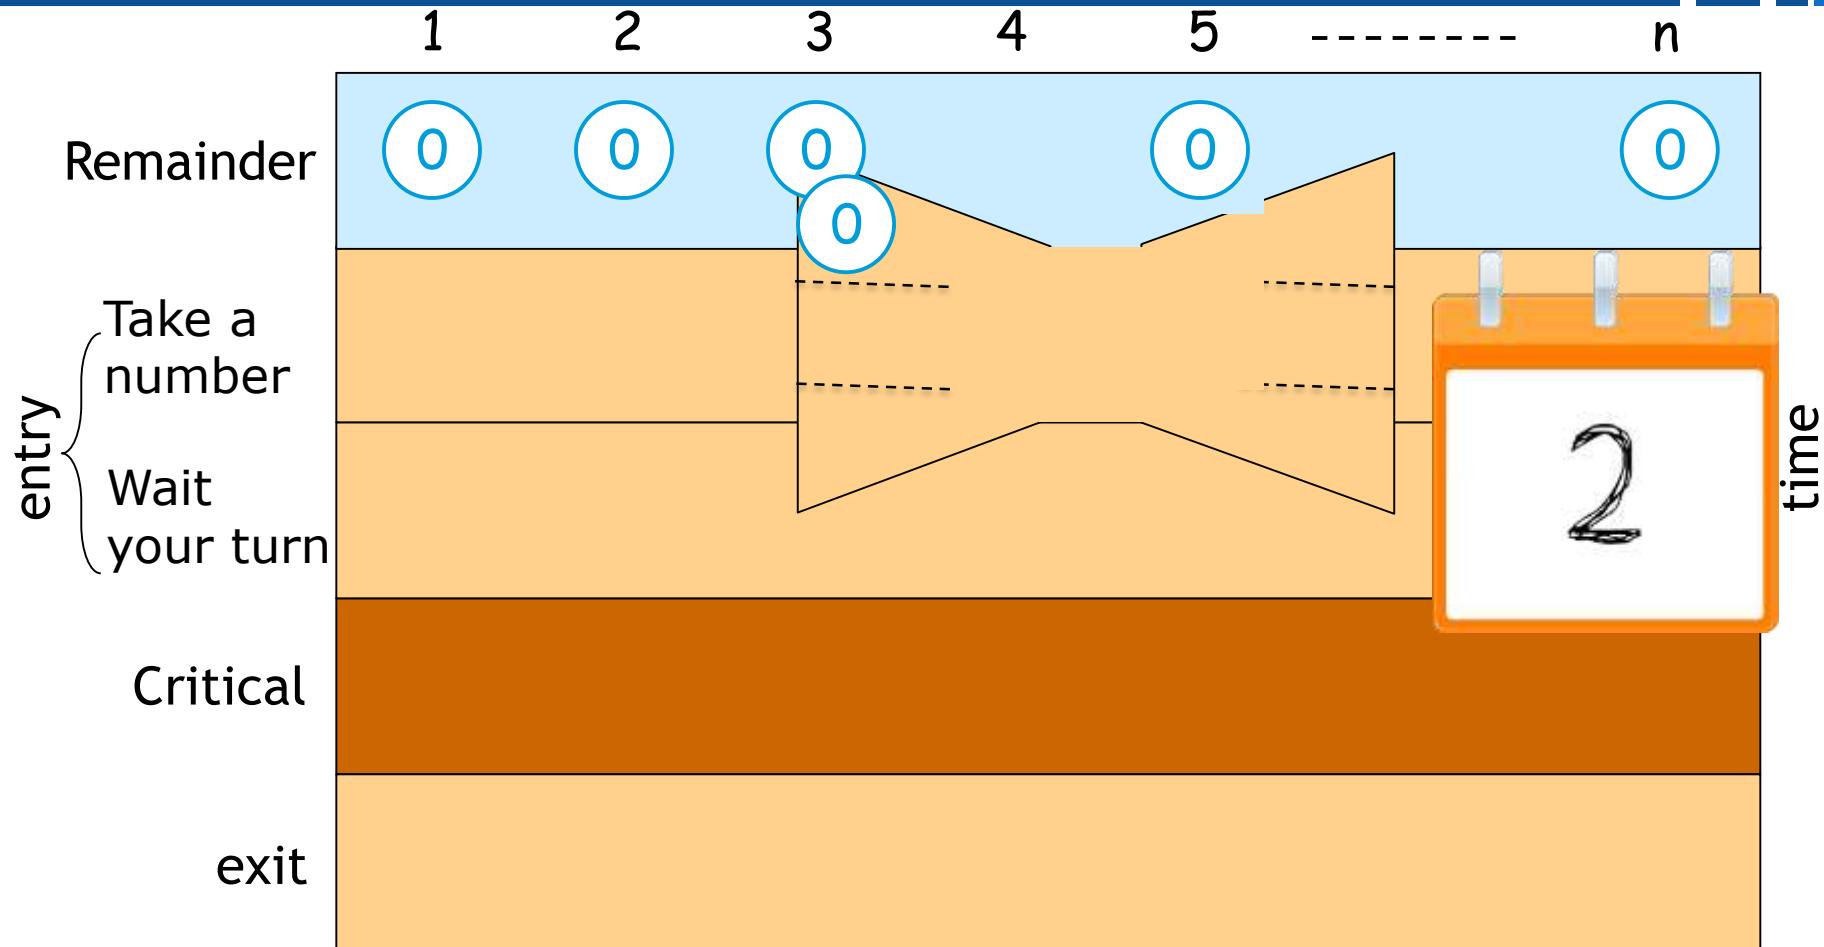
# Bakery Algorithm - Implementation 2



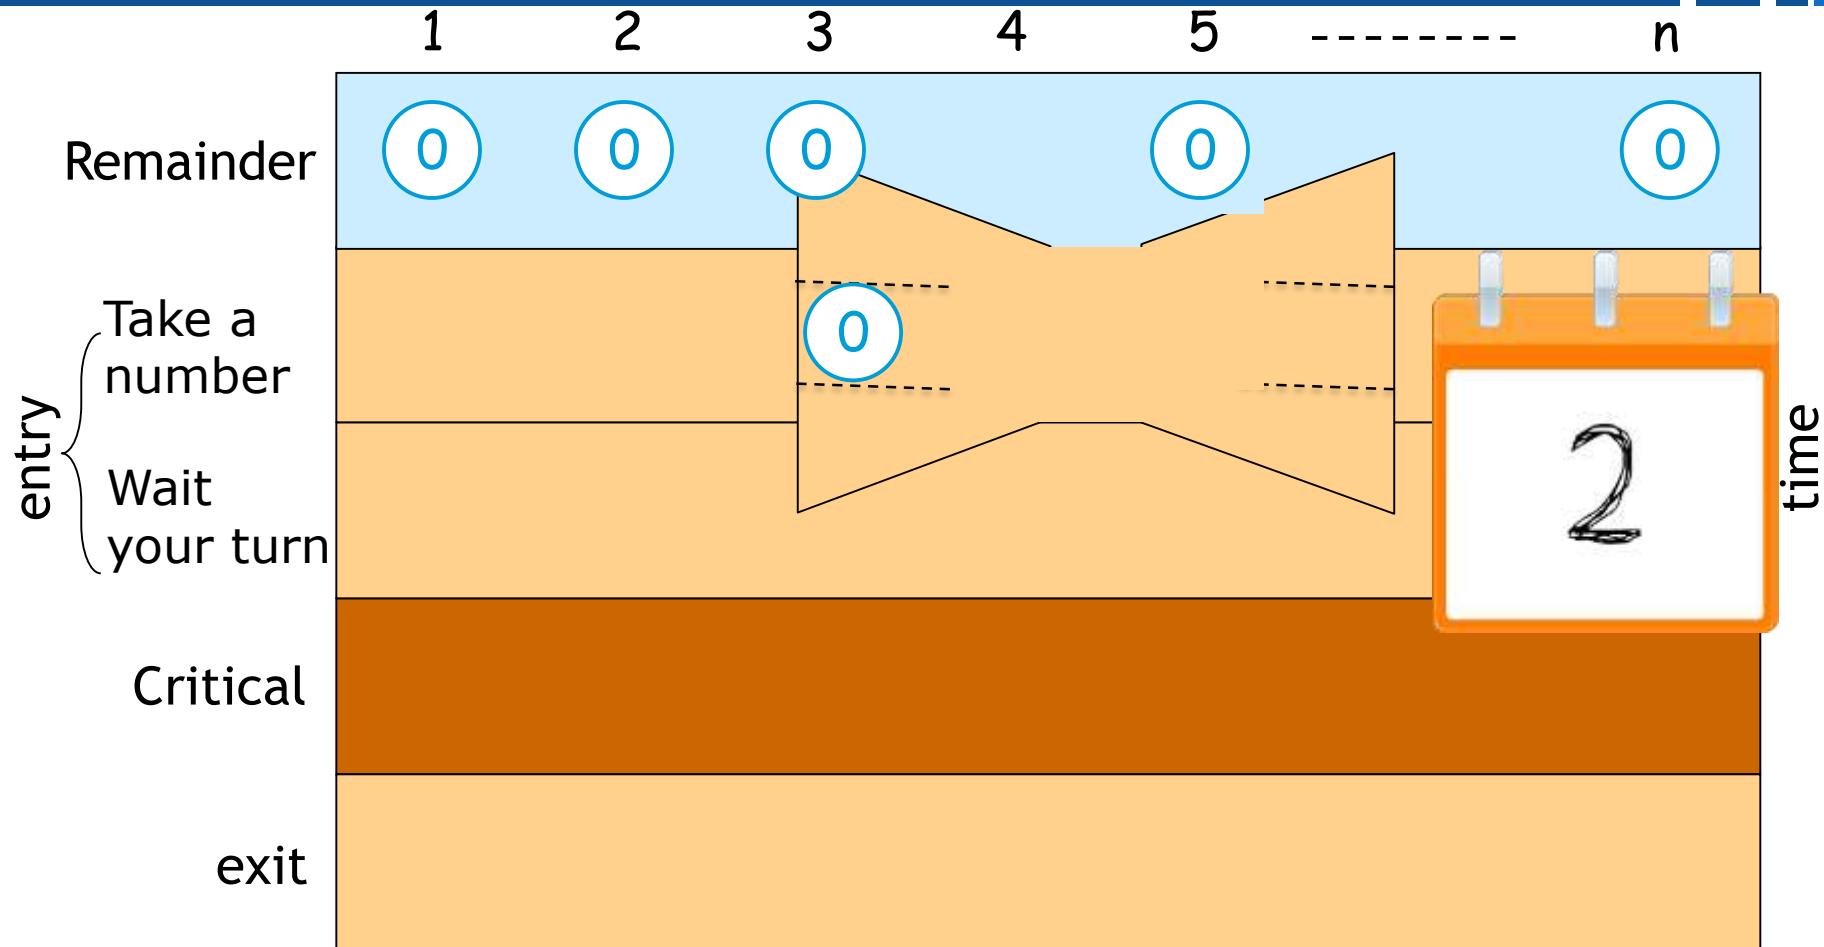
# Bakery Algorithm - Implementation 2



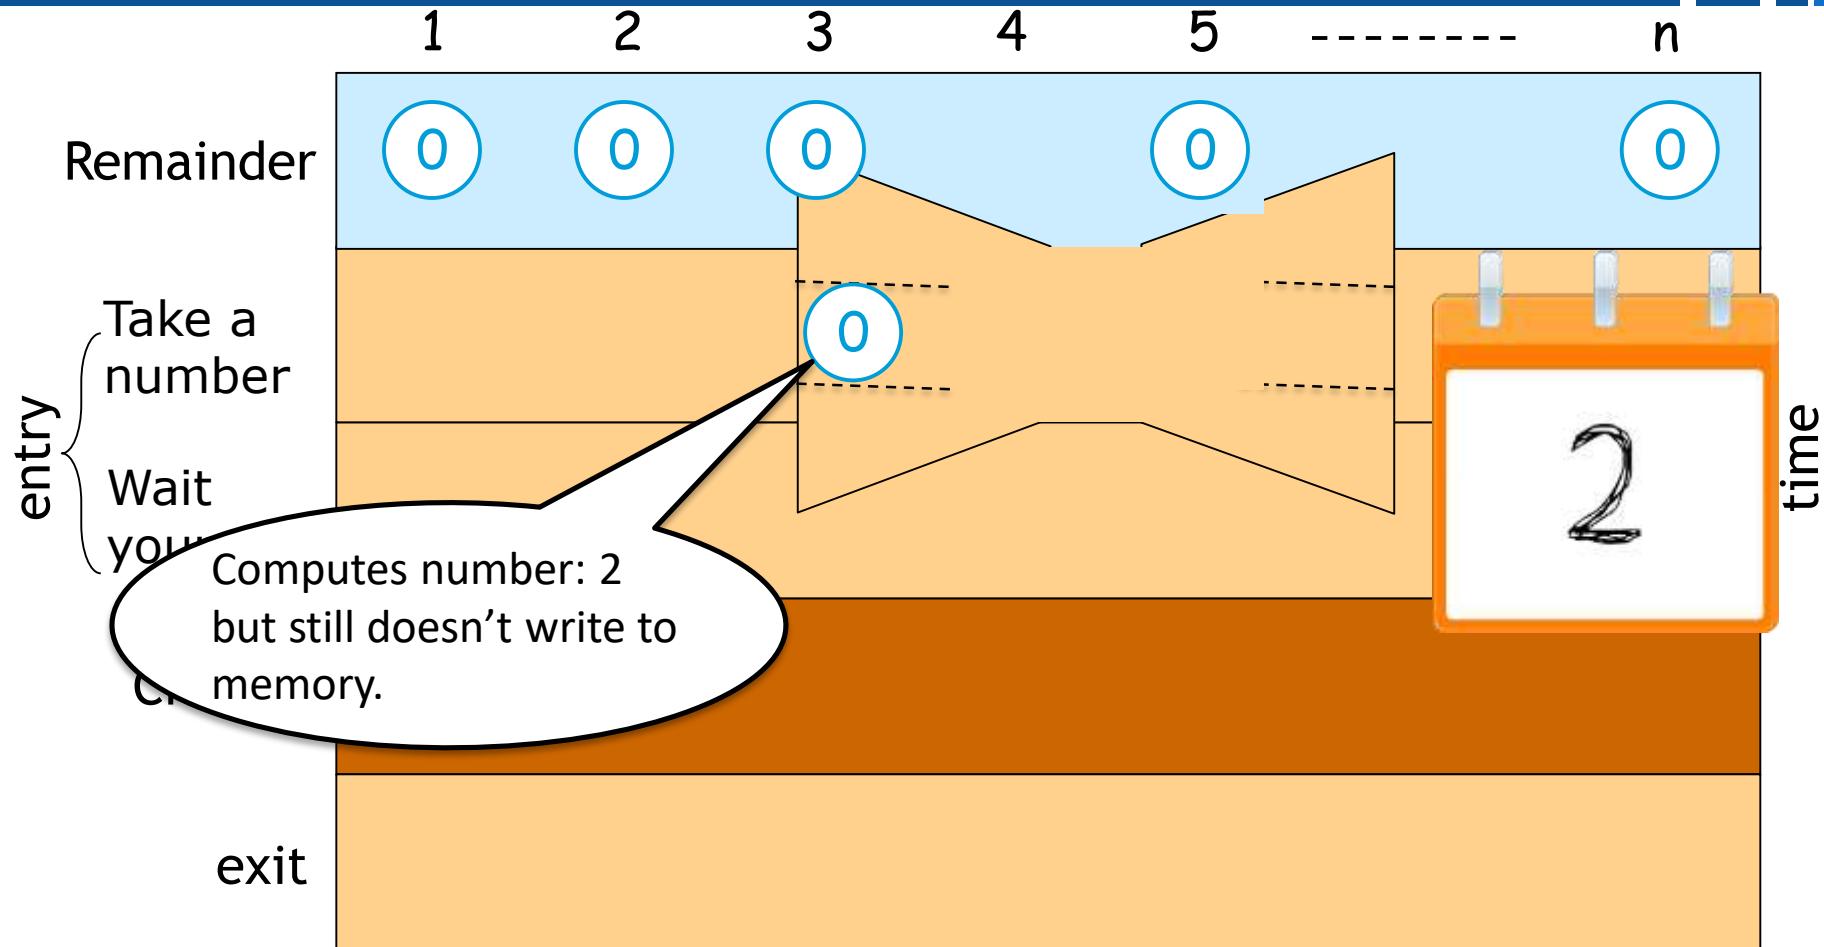
# Bakery Algorithm - Implementation 2



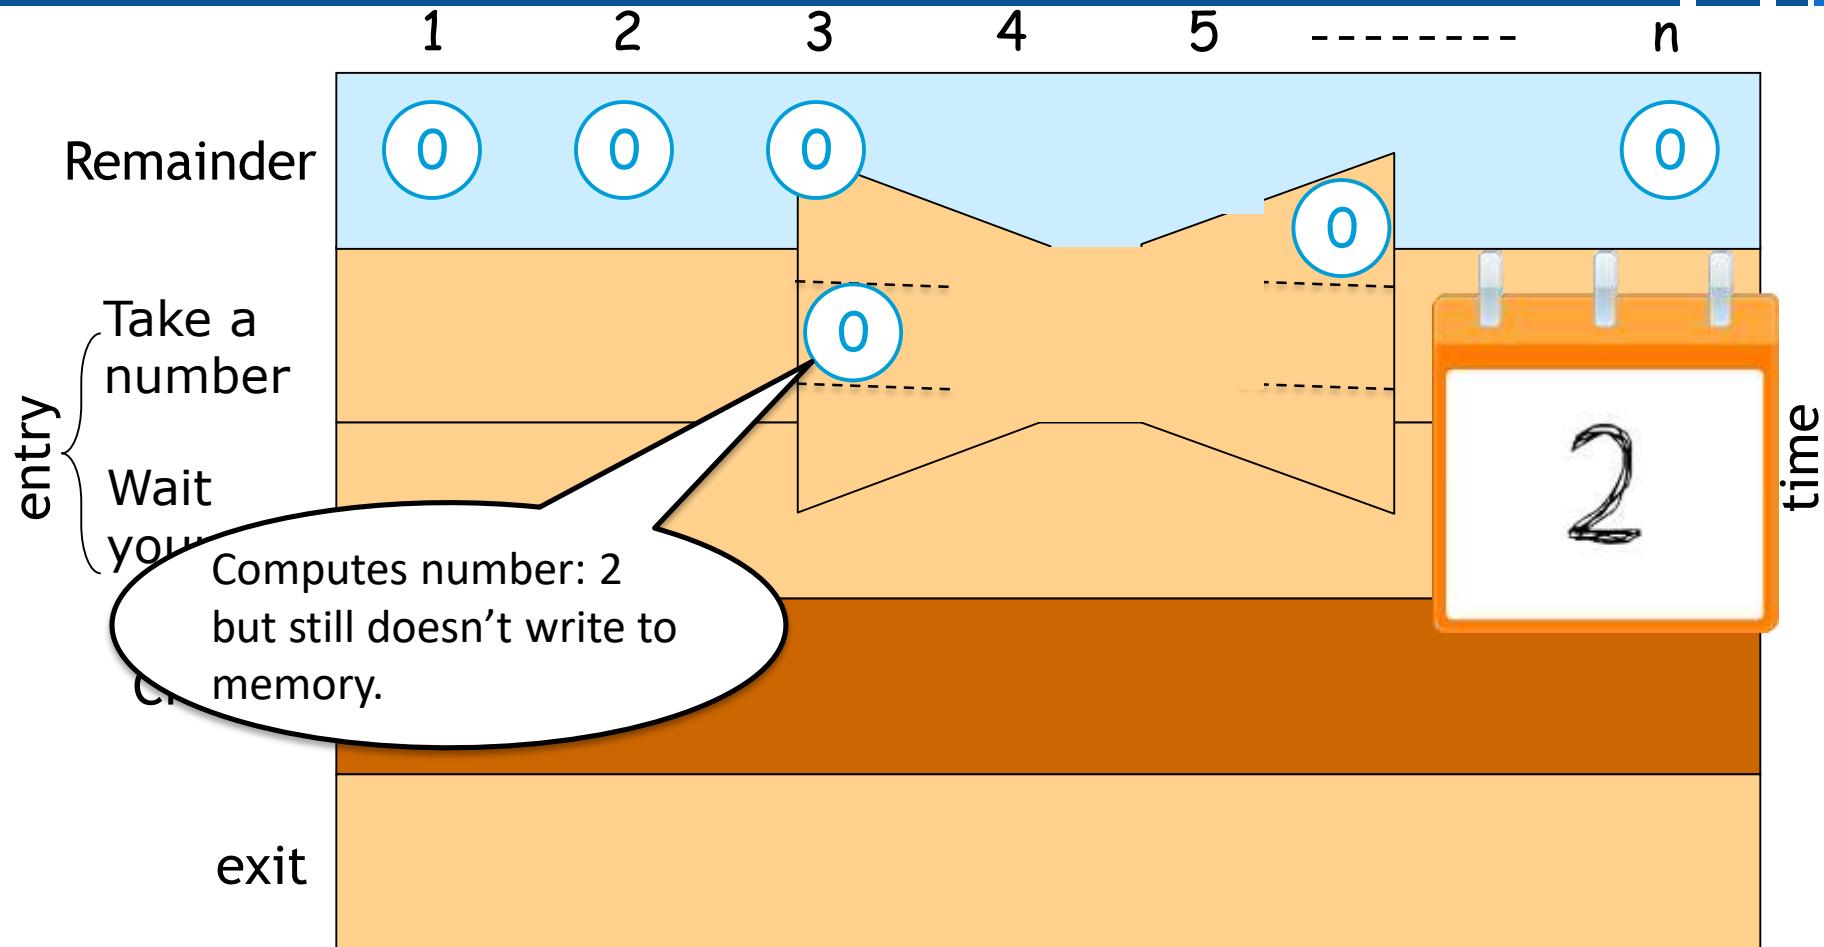
# Bakery Algorithm - Implementation 2



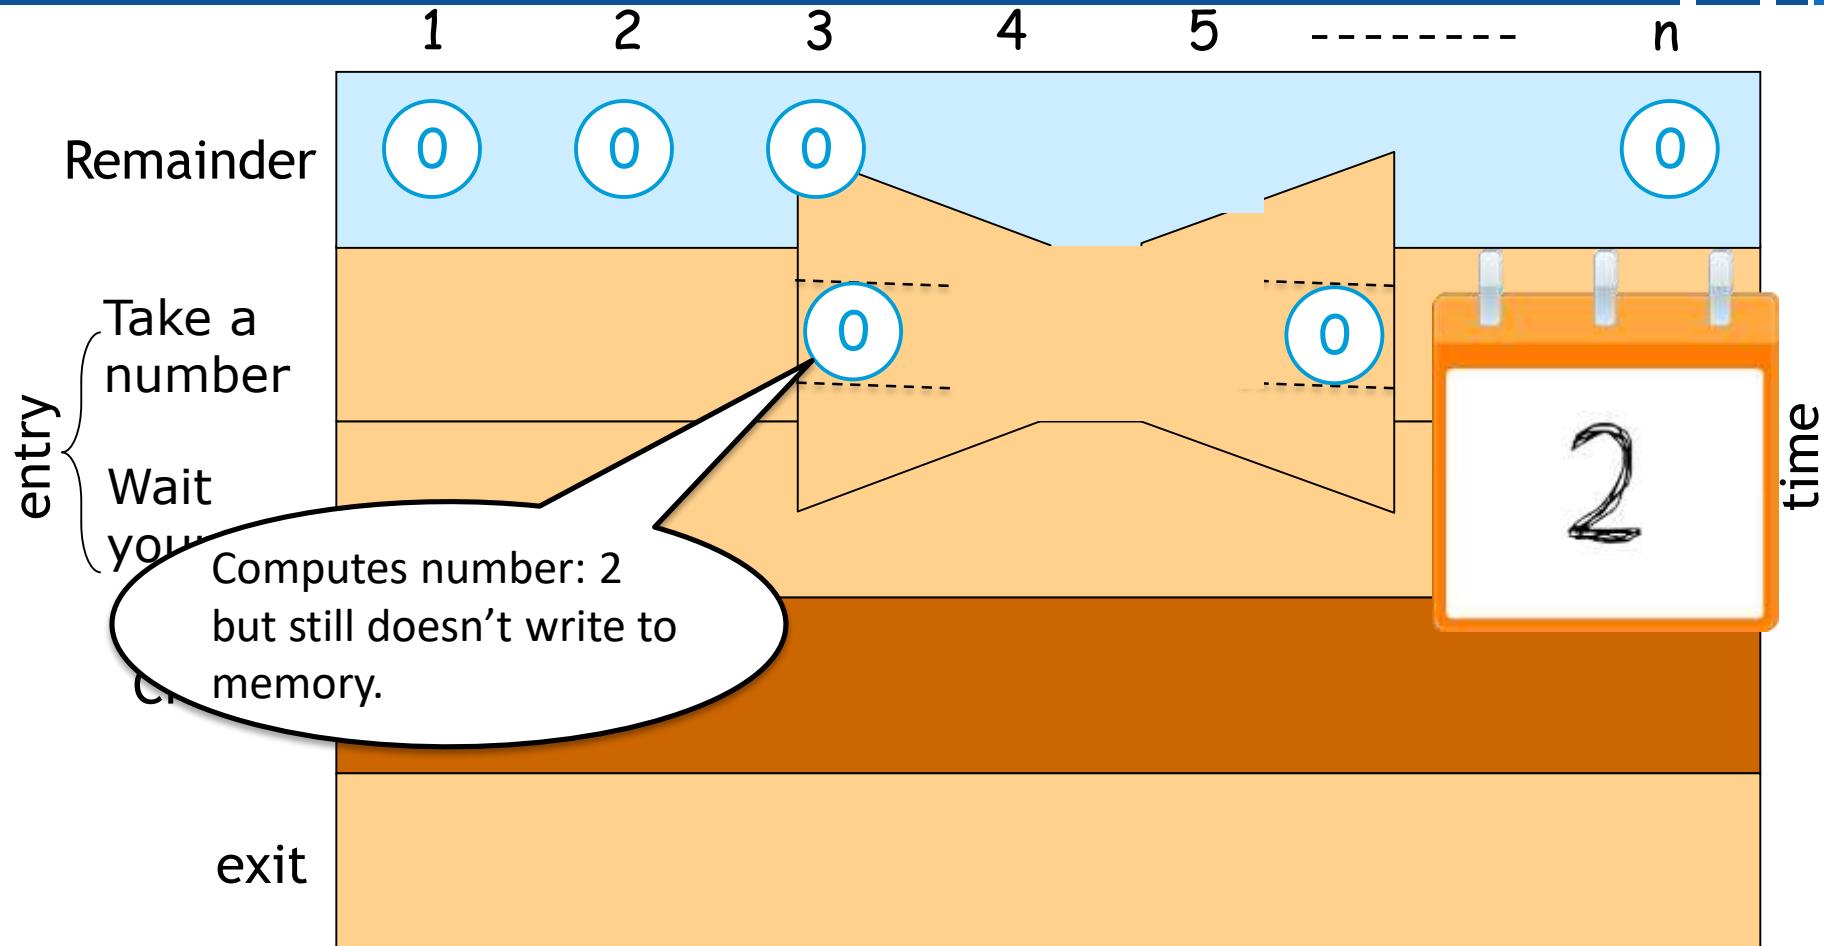
# Bakery Algorithm - Implementation 2



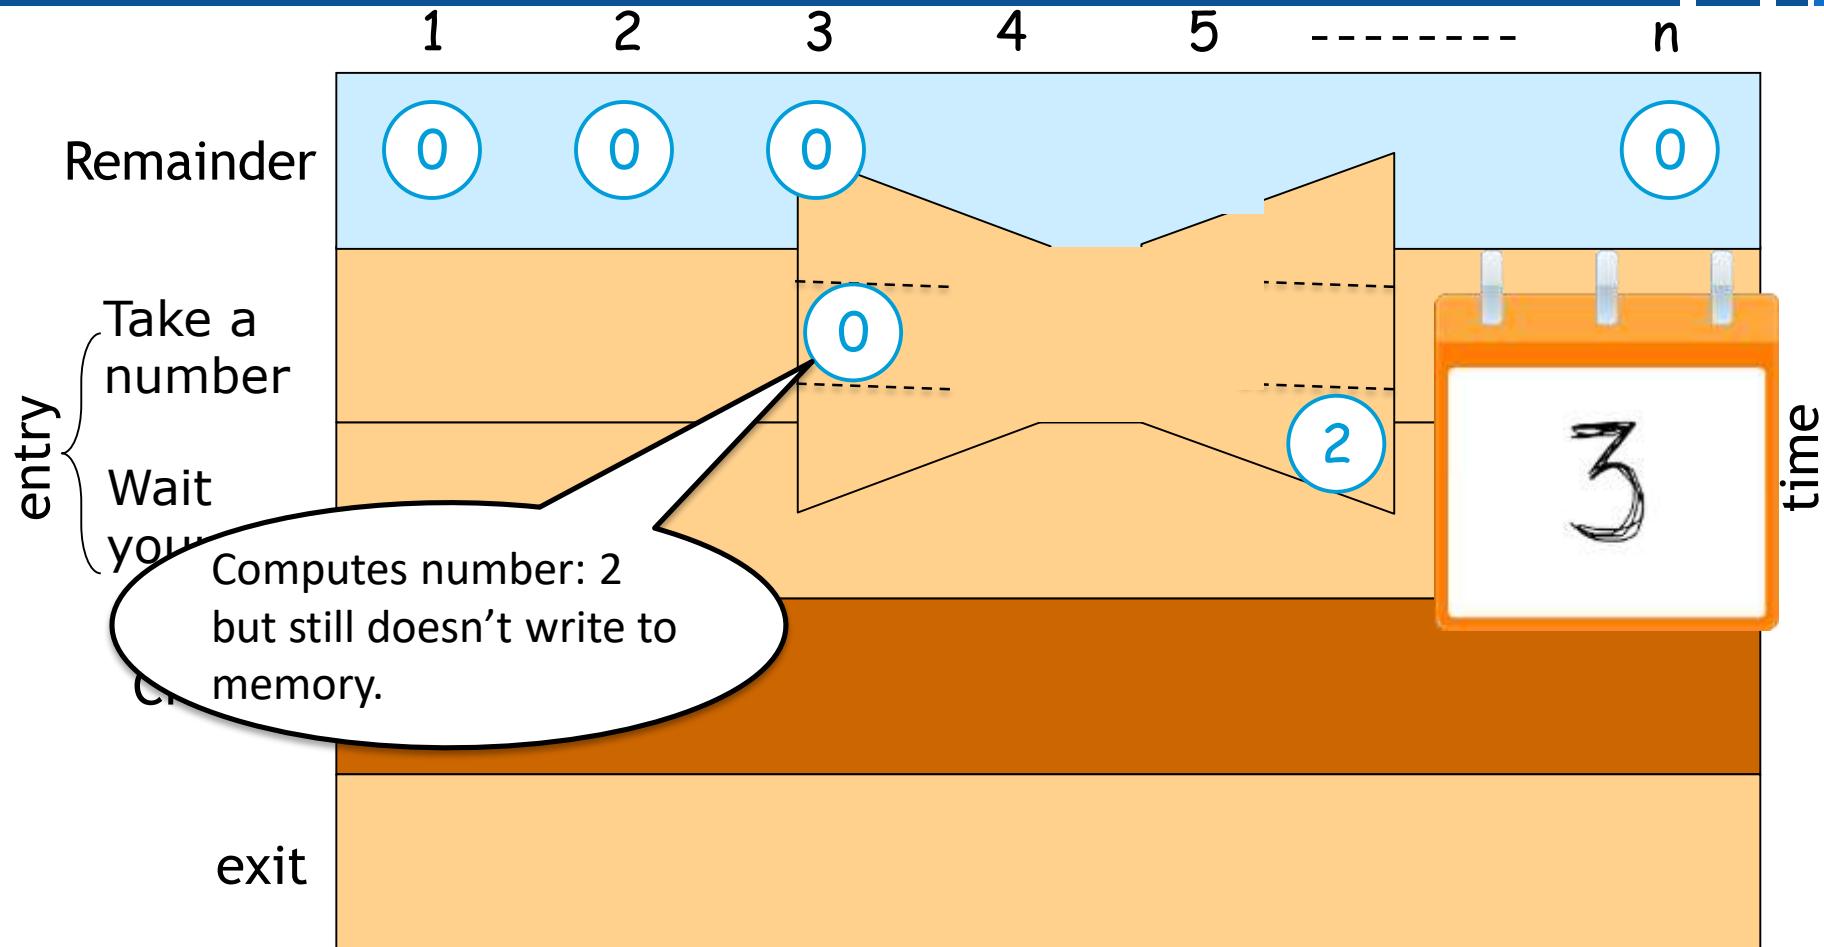
# Bakery Algorithm - Implementation 2



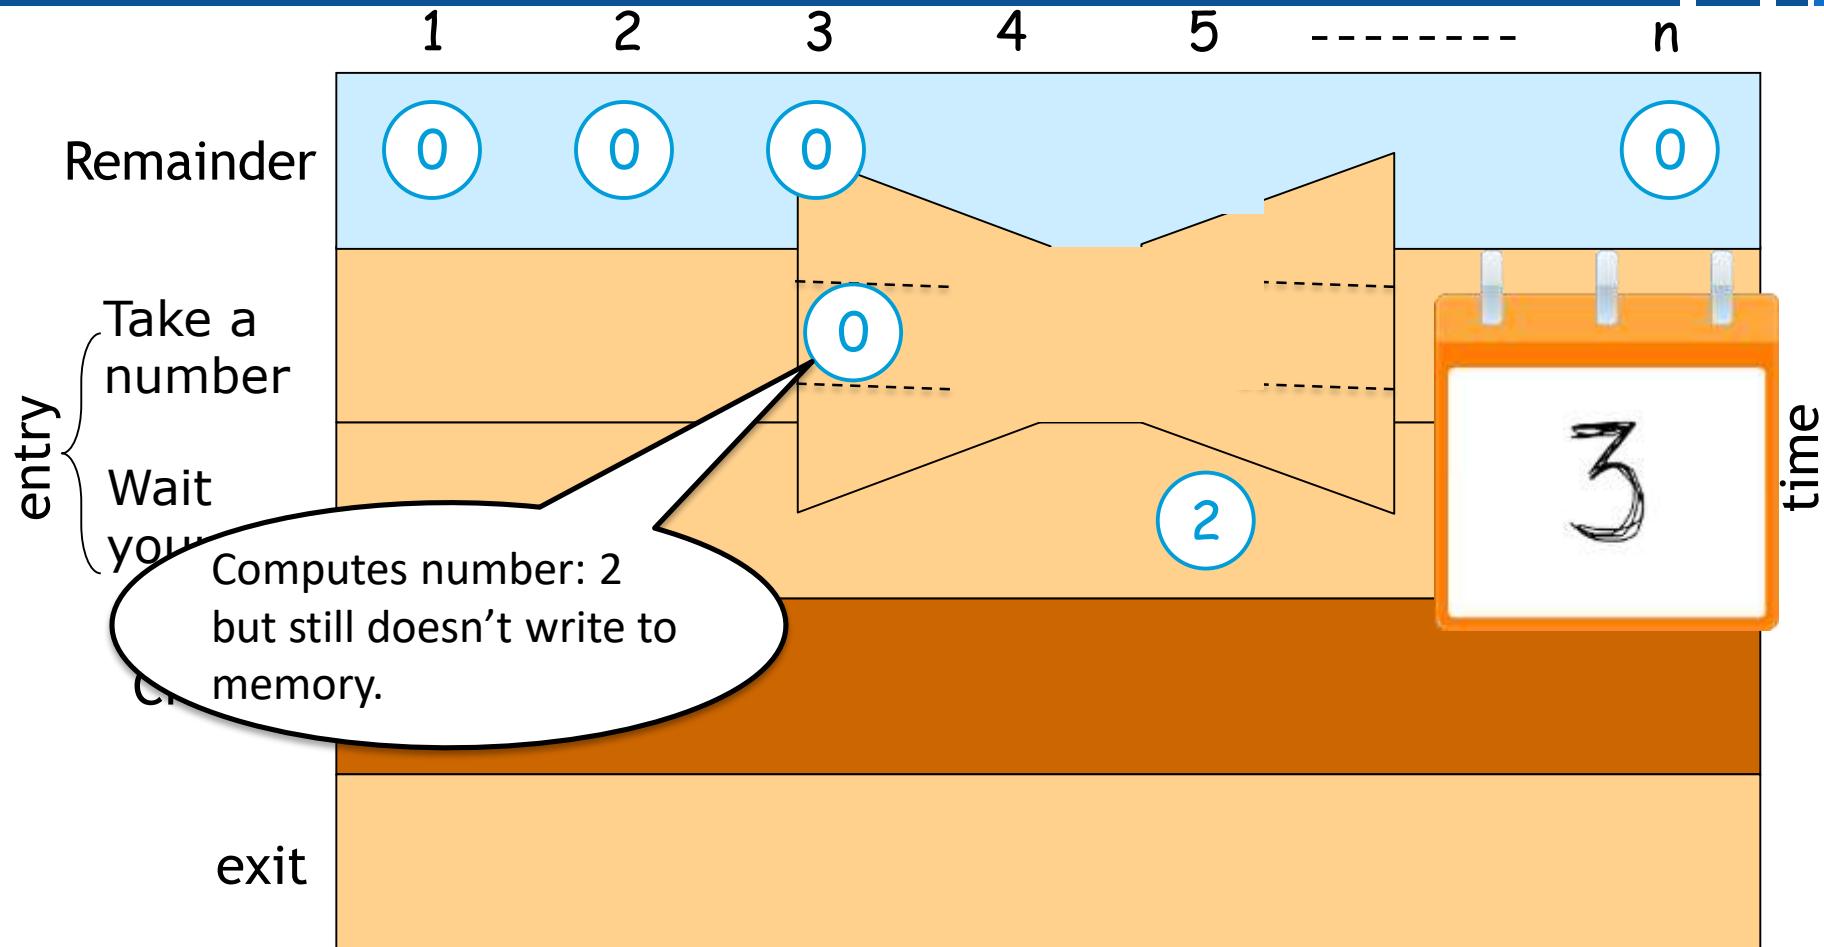
# Bakery Algorithm - Implementation 2



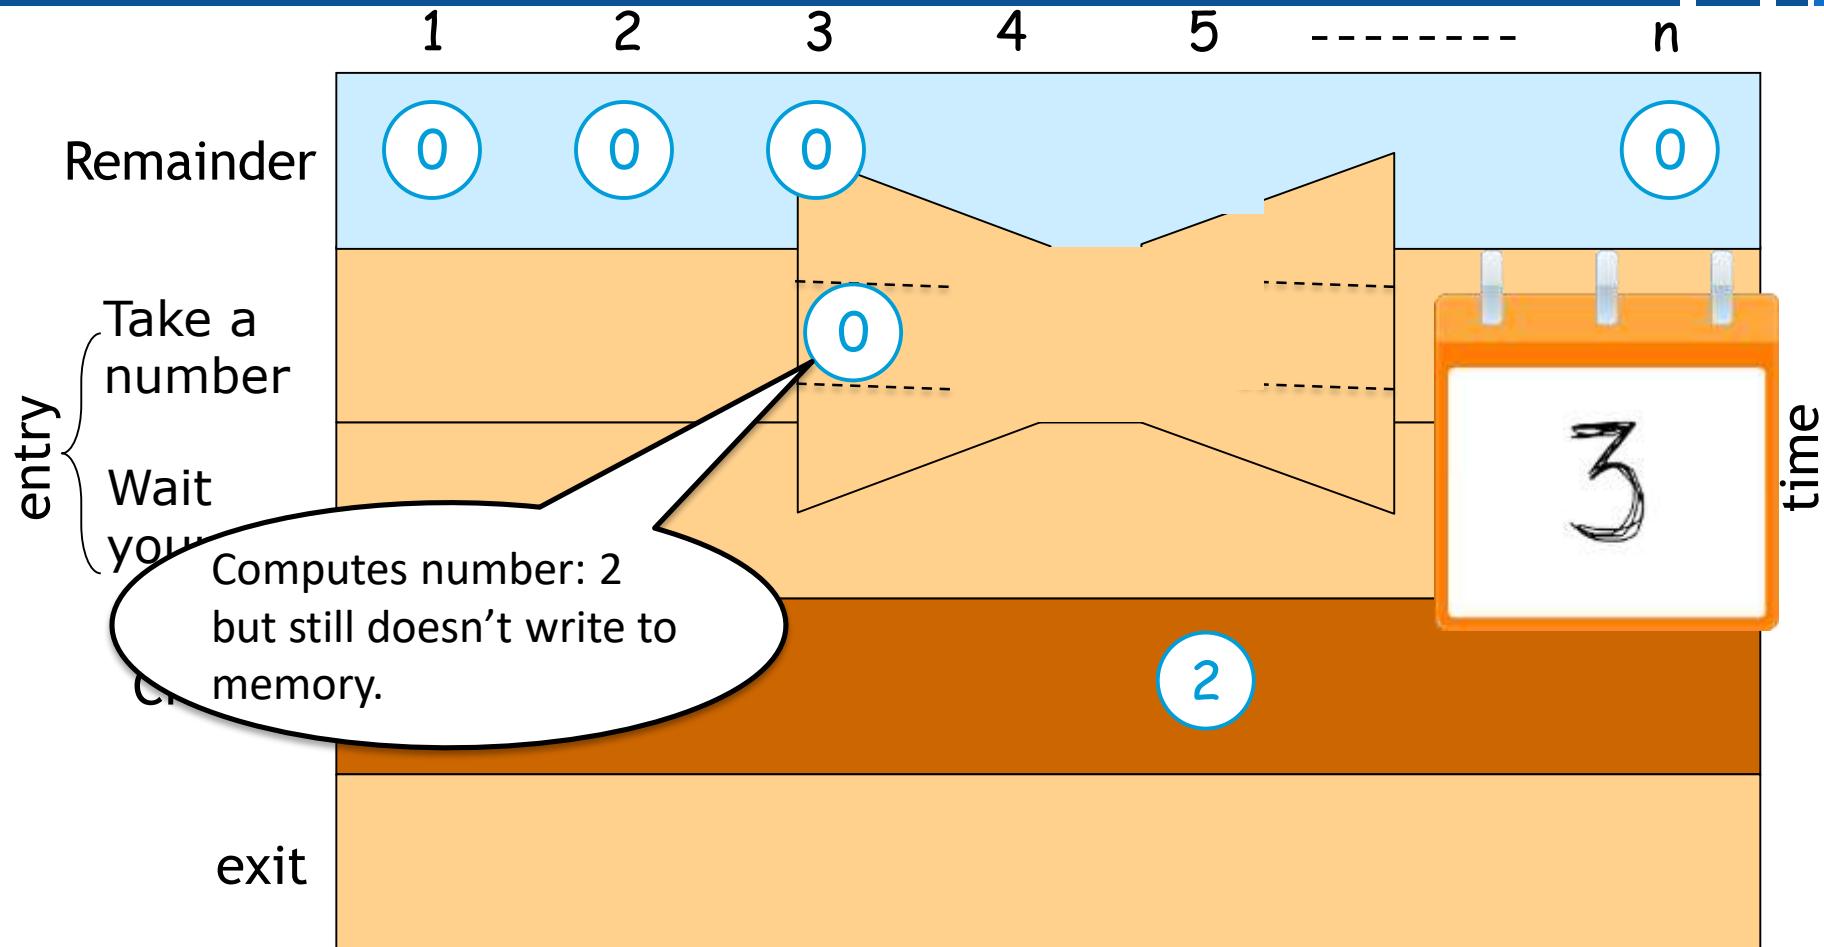
# Bakery Algorithm - Implementation 2



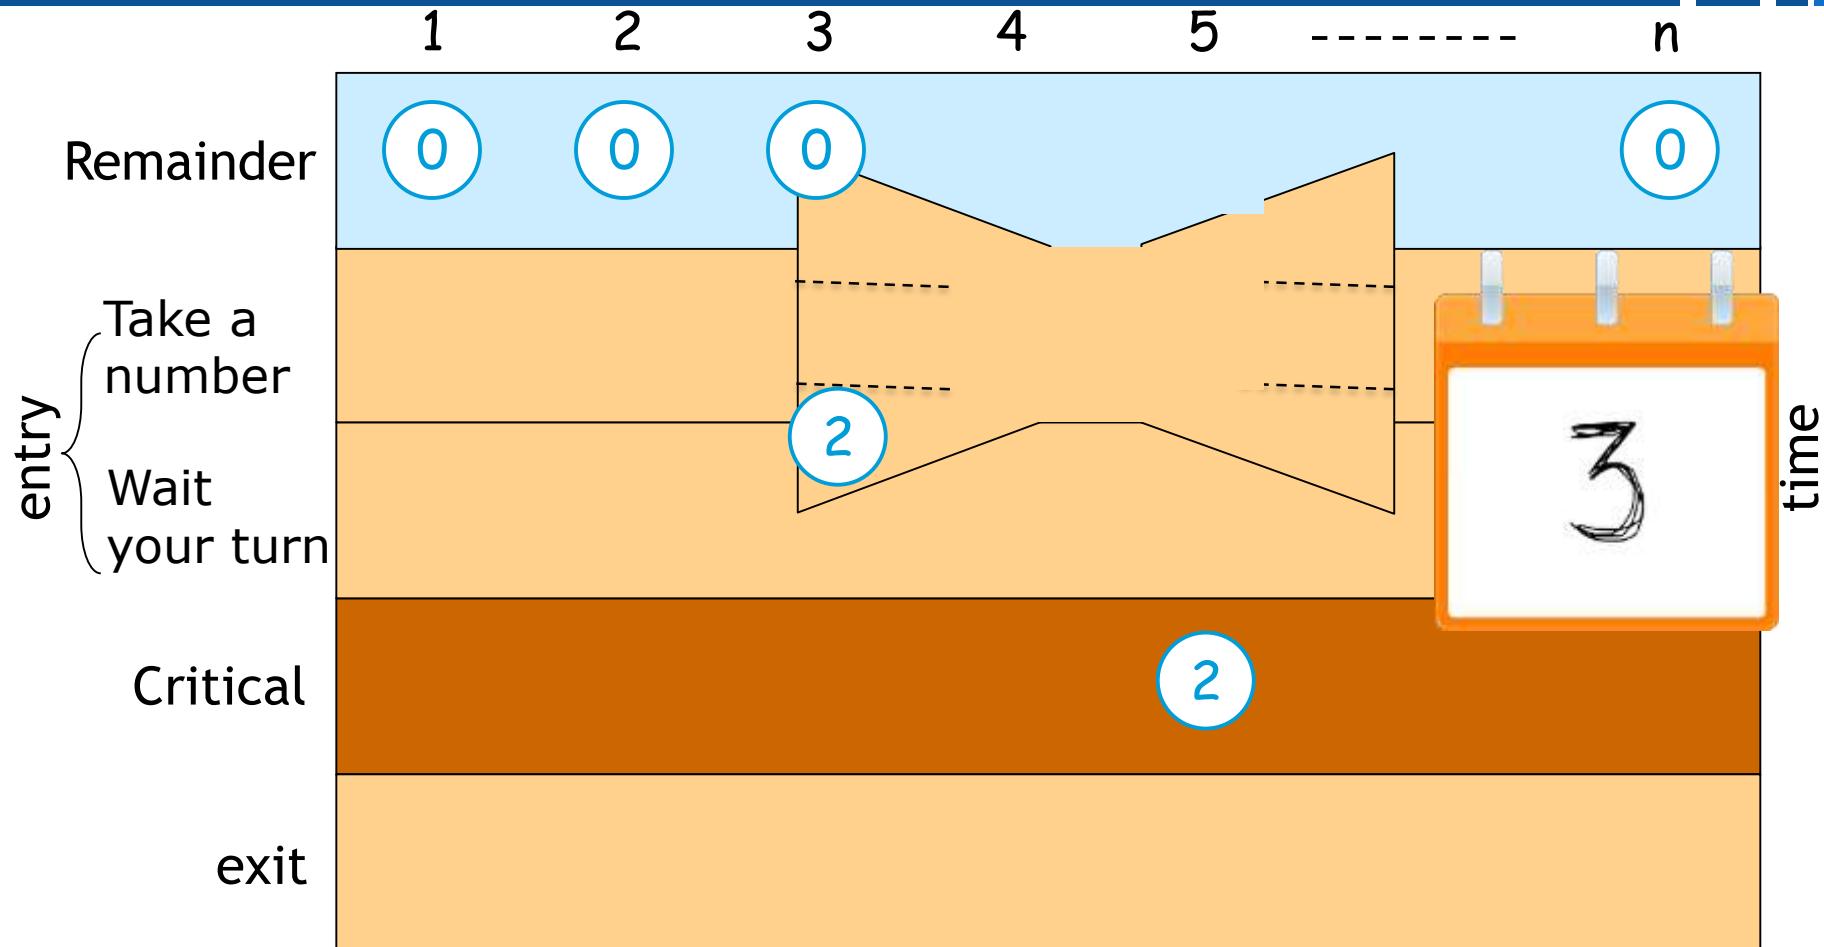
# Bakery Algorithm - Implementation 2



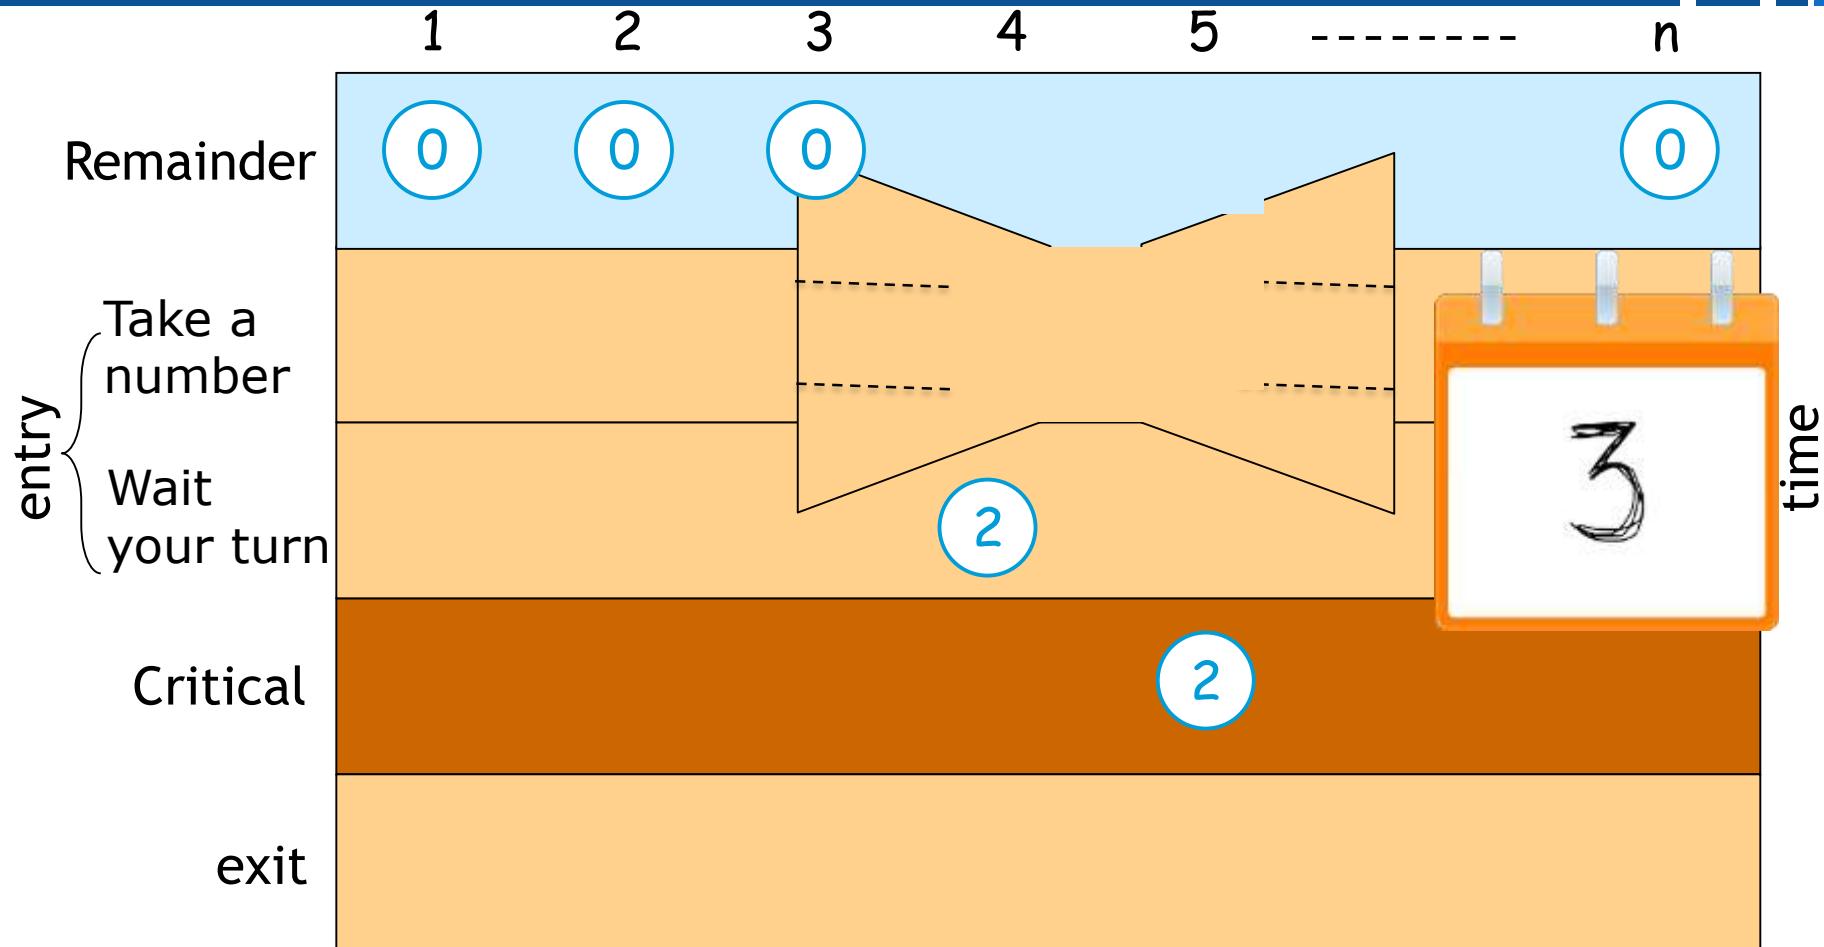
# Bakery Algorithm - Implementation 2



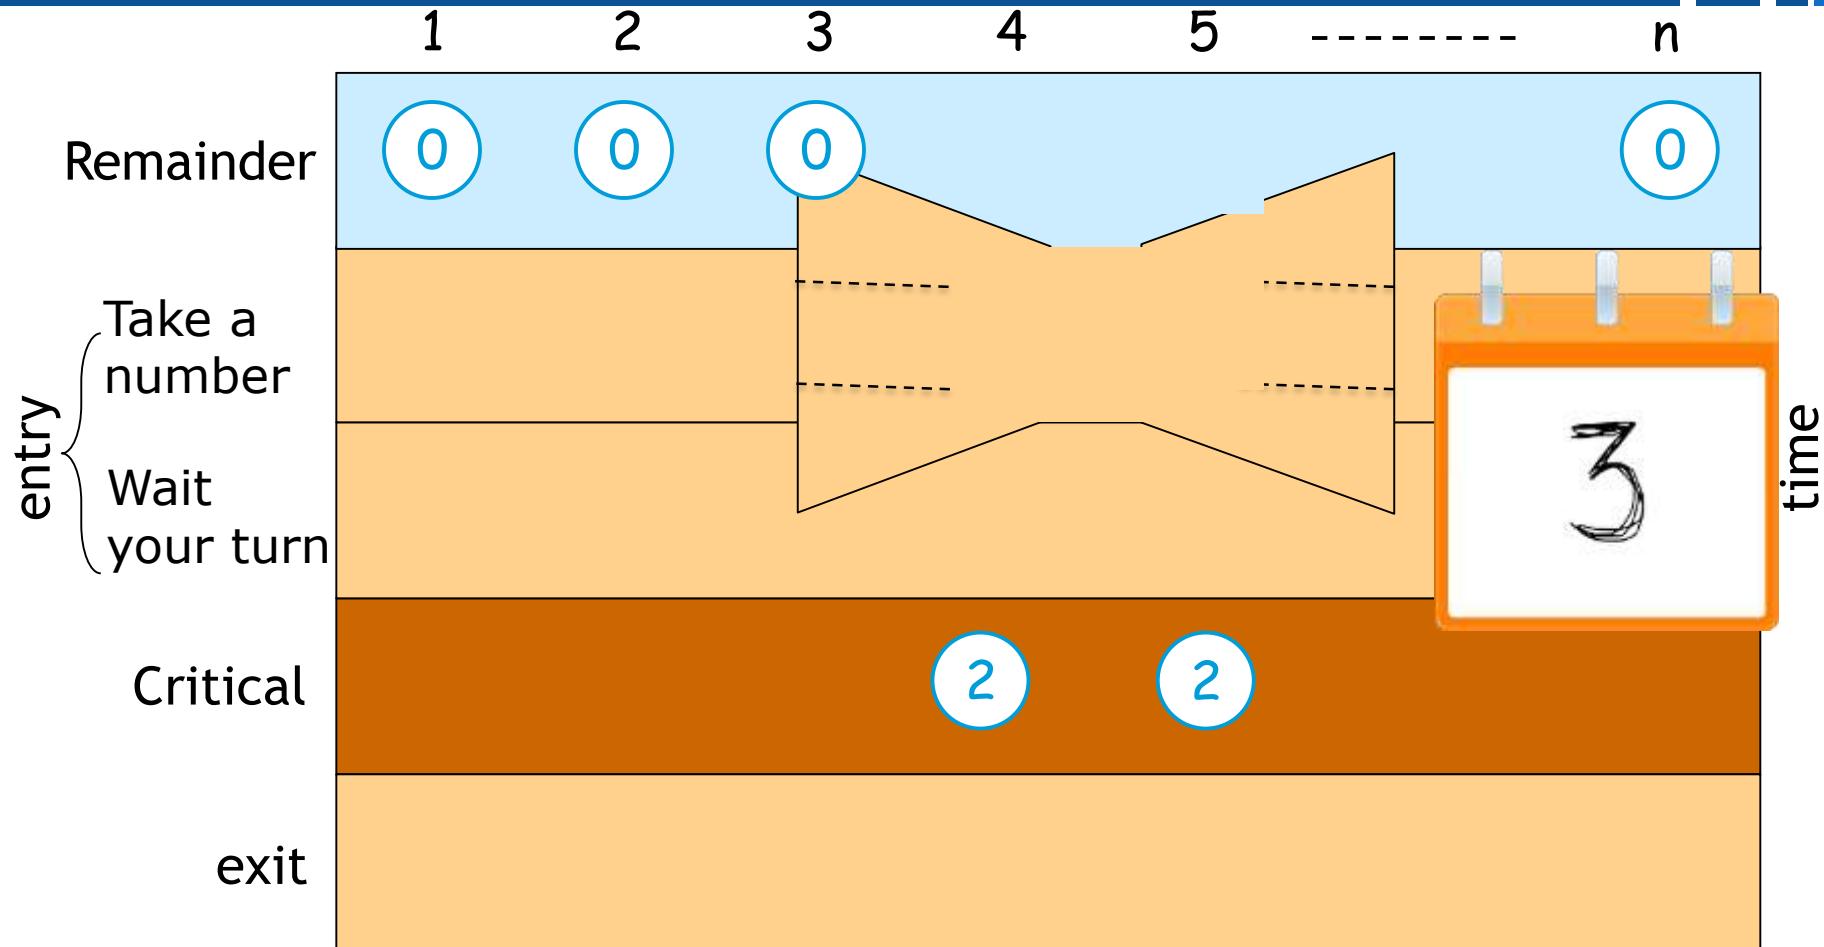
# Bakery Algorithm - Implementation 2



# Bakery Algorithm - Implementation 2



# Bakery Algorithm - Implementation 2



# Correct Implementation (1974)

Remainder

```
Choosing[i]=true;
Number[i] = 1+maxj ∈ {1,...,n}{number[j]};
Choosing[i]=false;
for j=1 to n {
    while(Choosing[j]);
    while(i≠j and number[j]>0 and
          (number[j],j)<(number[i],i));
}
```

Critical

```
number[i]=0;
```

Code for Thread *i*  
(out of *n* threads)

# Correct Implementation (1974)

Remainder

```
Choosing[i]=true;
Number[i] = 1+maxj ∈ {1,...,n}{number[j]};
Choosing[i]=false;
for j=1 to n {
    while(Choosing[j]);
    while(i≠j and number[j]>0 and
        (number[j],j)<(number[i],i));
}
```

Critical

```
number[i]=0;
```

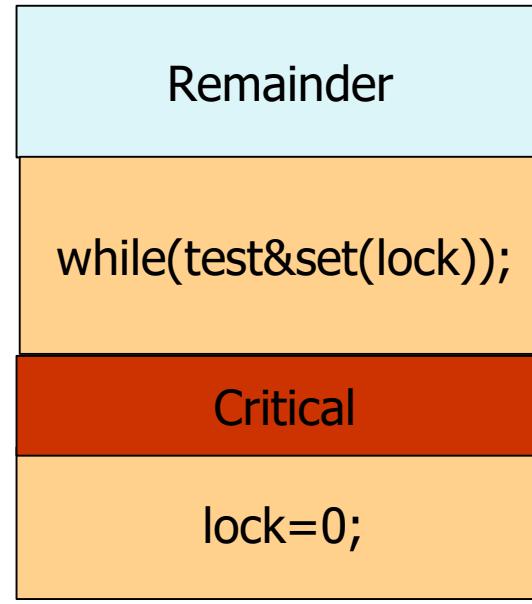
Code for Thread *i*  
(out of *n* threads)

Also ensures a notion of **First-In-First-Out**:  
If a process *i* is waiting (somewhere in the for loop) and thread *j* has not yet started the entry, then *j* will not enter the CS before *i*.

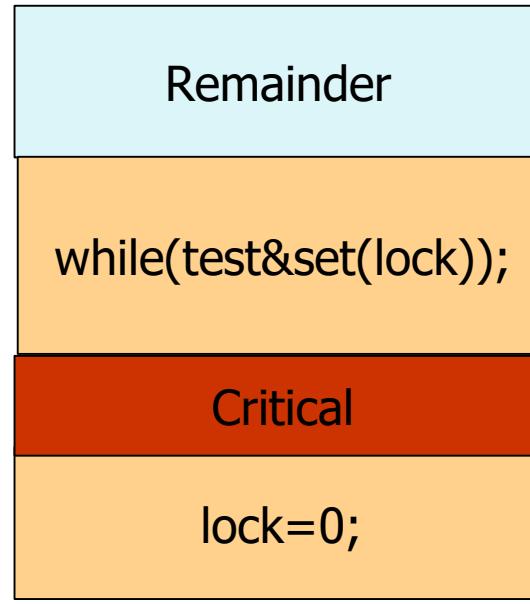
# Read-Modify-Write Instructions

- Up until now we assume that read (load) and write (store) are atomic instructions
  - There might be a context switch between a read and its corresponding write
- Modern CPUs support in hardware some RMW instructions:
  - **Test&Set(&lock):**  
 $\{i = *lock; *lock = 1; \text{return } i;\}$
  - **Fetch&Add(&p,inc):**  
 $\{val = *p; *p = val + inc; \text{return } val\}$
  - **Compare&Swap(&p,old,new):**  
 $\{\text{if}(*p \neq old) \text{return false}; *p = new; \text{return true;}\}$

# Mutual Exclusion with Test&Set



# Mutual Exclusion with Test&Set



Code for Thread i

**Not so fast: Starvation Freedom** condition does not hold!

# Burn's Algorithm

Remainder

```
waiting[i]=true;  
key[i]=1;  
while(waiting[i] and key[i]) {  
    key[i]=test&set(lock) ;  
    waiting[i]=false;
```

Critical

```
j=i+1 mod n;  
while(j≠i and not waiting[j]) {  
    j=j+1 mod n };  
if(j≠i) { waiting[j]=false;}  
else { lock=0;}
```

Code for Thread i out of n

# Up until now...

- We saw several solutions for the **mutual exclusion** problem in different settings
- All the solutions perform busy-waiting:
  - Waste of CPU resources
- Mutual exclusion is only one synchronization problem. There are many more
  - Contention → coordination
  - Some were and will be covered in the tutorials

# Next: Synchronization Primitives

- We will define **abstract data types** used to provide synchronization

# Next: Synchronization Primitives

- We will define **abstract data types** used to provide synchronization
- Clear **interface** to the data type

# Next: Synchronization Primitives

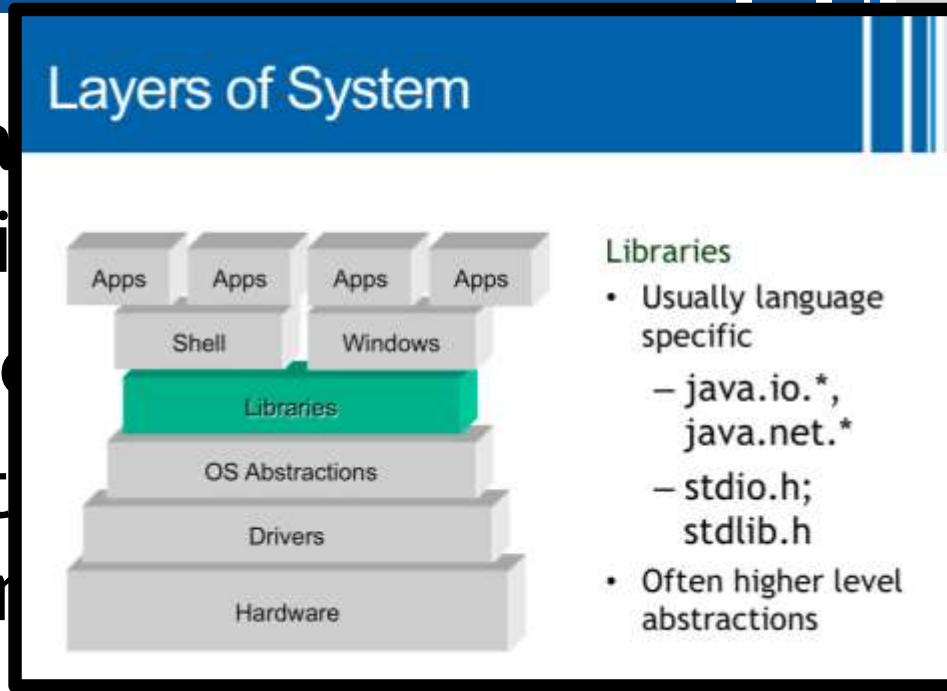
- We will define **abstract data types** used to provide synchronization
- Clear **interface** to the data type
- The data type (and its inner variables) cannot be accessed in any other way

# Next: Synchronization Primitives

- We will define **abstract data types** used to provide synchronization
- Clear **interface** to the data type
- The data type (and its inner variables) cannot be accessed in any other way
- Usually provided by the programming language with some support of the OS

# Next: Synchronization Primitives

- We will define **abstractions** that provide synchronization
- Clear **interface** to the abstraction
- The data type (and its operations) cannot be accessed in isolation
- Usually provided by the programming language with some support of the OS



# Next: Synchronization Primitives

- We will define **abstract data types** used to provide synchronization
- Clear **interface** to the data type
- The data type (and its inner variables) cannot be accessed in any other way
- Usually provided by the programming language with some support of the OS
- The abstract data type may have many implementations

# First Data Type: Semaphore

- Record with two fields:
  - value
  - List (L)



# First Data Type: Semaphore

- Record with two fields:
  - value
  - List (L)
- Two operations:



## Down(S)

```
S.value= S.value - 1  
if S.value < 0 then  
{ add this thread to S.L;  
sleep();}
```

## Up(S)

```
S.value= S.value + 1  
if S.value≤0 then  
{remove a thread T from  
S.L; Wakeup(T);}
```

## Init(S,v)

```
S.value= v
```

# First Data Type: Semaphore

- Record with two fields:
  - value
  - List (L)
- Two operations:



## Down(S)

```
S.value= S.value - 1  
if S.value < 0 then  
{ add this thread to S.L;  
sleep();}
```

## Up(S)

```
S.value= S.value + 1  
if S.value≤0 then  
{remove a thread T from  
S.L; Wakeup(T);}
```

## Init(S,v)

```
S.value= v
```

- The operations are executed **atomically**
  - How? Implementation is orthogonal to the definition

# First Data Type: Semaphore

- Record with two fields:

In literature, the operations are often called P(s) and V(s).  
In the book, wait(s), signal(s)

**Down(S)**

```
S.value= S.value - 1  
if S.value < 0 then  
{ add this thread to S.L;  
sleep();}
```

**Up(S)**

```
S.value= S.value + 1  
if S.value≤0 then  
{remove a thread T from  
S.L; Wakeup(T);}
```

**Init(S,v)**

```
S.value= v
```

- The operations are executed **atomically**
  - How? Implementation is orthogonal to the definition



# First Data Type: Semaphore

- Record with two fields:
  - value
  - List (L)
- Two operations:



Some variations bound the maximum and minimum value (e.g., binary semaphores)

**Down(S)**

```
S.value= S.value - 1  
if S.value < 0 then  
{ add this thread to S.L;  
sleep();}
```

**p(S)**

```
S.value= S.value + 1  
if S.value≤0 then  
{remove a thread T from  
S.L; Wakeup(T);}
```

**Init(S,v)**

```
S.value= v
```

- The operations are executed **atomically**
  - How? Implementation is orthogonal to the definition

# Mutual Exclusion w/ Semaphore

Remainder	
down(lock)	<pre>lock.value= lock.value - 1 if lock.value &lt; 0 then { add this thread to S.L; sleep();}</pre>
Critical	
up(lock)	<pre>lock.value= lock.value + 1 If lock.value≤0 then {remove a thread T from S.L; Wakeup(T);}</pre>
Code for Thread i	

# Mutual Exclusion w/ Semaphore

Remainder		<ul style="list-style-type: none"><li>• Assume <b>lock</b> is a shared semaphore<ul style="list-style-type: none"><li>– Initially <b>lock</b> is 1</li></ul></li></ul>
down(lock)	lock.value= lock.value - 1 if lock.value < 0 then { add this thread to S.L; sleep();}	
Critical		
up(lock)	lock.value= lock.value + 1 If lock.value≤0 then {remove a thread T from S.L; Wakeup(T);}	
Code for Thread i		

# Mutual Exclusion w/ Semaphore

Remainder	
down(lock)	<pre>lock.value= lock.value - 1 if lock.value &lt; 0 then { add this thread to S.L; sleep();}</pre>
Critical	
up(lock)	<pre>lock.value= lock.value + 1 If lock.value≤0 then {remove a thread T from S.L; Wakeup(T);}</pre>

Code for Thread i

- Assume **lock** is a shared semaphore
  - Initially **lock** is 1
- Value:
  - Cannot be more than 1
  - 0: one process in Critical Section
  - (-x): x processes are waiting

# Mutual Exclusion w/ Semaphore

Remainder		<ul style="list-style-type: none"><li>• Assume <b>lock</b> is a shared semaphore<ul style="list-style-type: none"><li>– Initially <b>lock</b> is 1</li></ul></li></ul>
down(lock)	<pre>lock.value= lock.value - 1 if lock.value &lt; 0 then { add this thread to S.L; sleep();}</pre>	<ul style="list-style-type: none"><li>• Value:<ul style="list-style-type: none"><li>– Cannot be more than 1</li><li>– 0: one process in Critical Section</li><li>– (-x): x processes are waiting</li></ul></li></ul>
Critical		
up(lock)	<pre>lock.value= lock.value + 1 If lock.value≤0 then {remove a thread T from S.L; Wakeup(T);}</pre>	<ul style="list-style-type: none"><li>• All properties hold<ul style="list-style-type: none"><li>– If L is a FIFO queue</li></ul></li></ul>
Code for Thread i		

# Mutual Exclusion w/ Semaphore

Remainder	
down(lock)	<pre>lock.value= lock.value - 1 if lock.value &lt; 0 then { add this thread to S.L; sleep();}</pre>
Critical	
up(lock)	<pre>lock.value= lock.value + 1 If lock.value≤0 then {remove a thread T from S.L; Wakeup(T);}</pre>

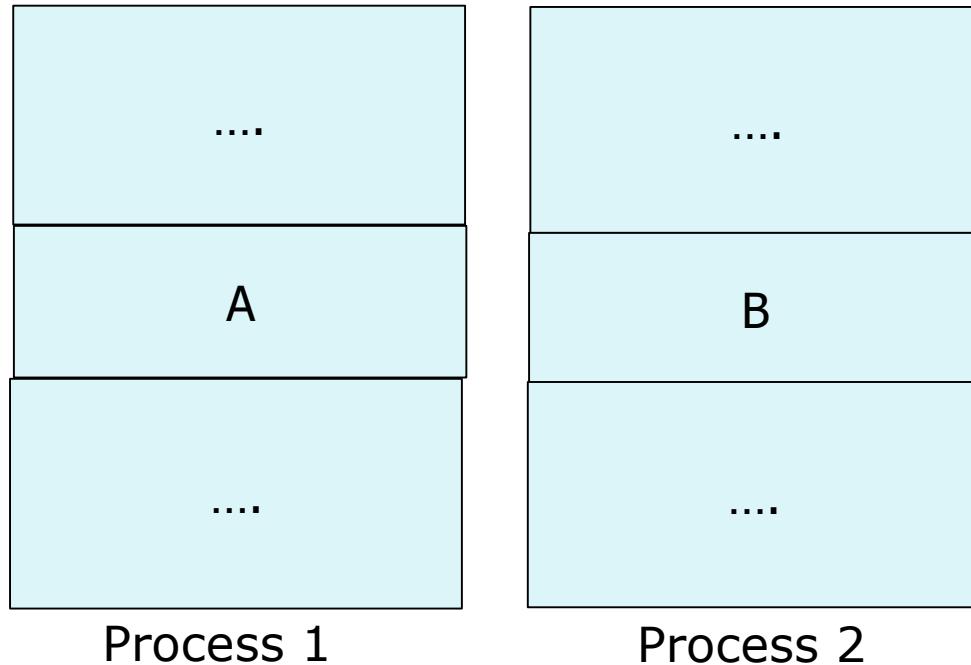
Code for Thread i

- Assume **lock** is a shared semaphore
  - Initially **lock** is 1
- Value:
  - Cannot be more than 1
  - 0: one process in Critical Section
  - (-x): x processes are waiting
- All properties hold
  - If L is a FIFO queue

0/1 (binary) semaphore suffices → Binary semaphores data type are often called Mutex objects

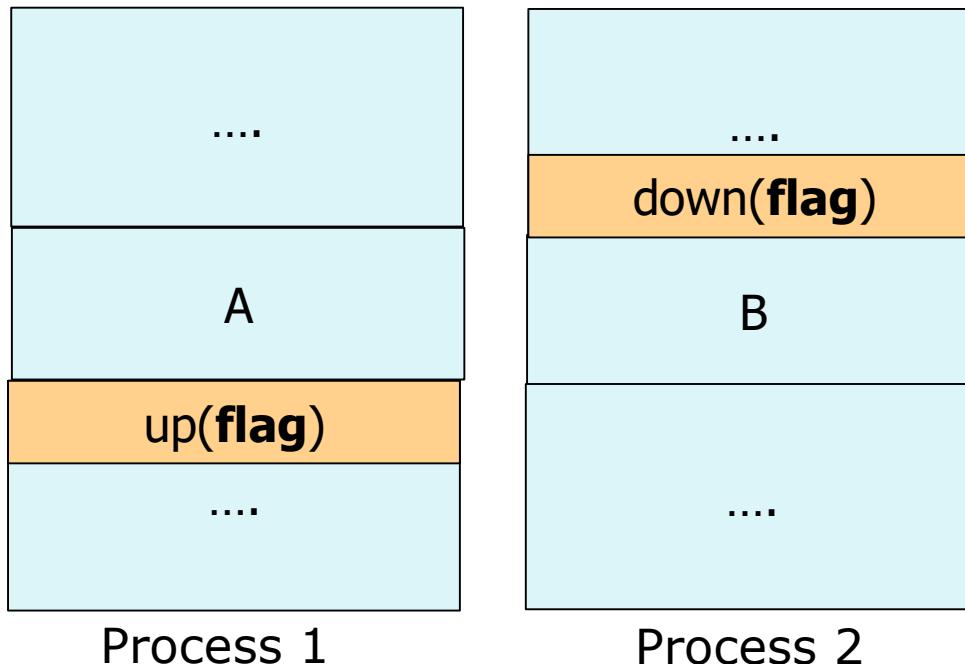
# “Execute B after A”

- A different synchronization problem
  - Coordination rather than contention
- One process needs to execute code A, before another process executes code B



# “Execute B after A”

- A different synchronization problem
  - Coordination rather than contention
- One process needs to execute code A, before another process executes code B



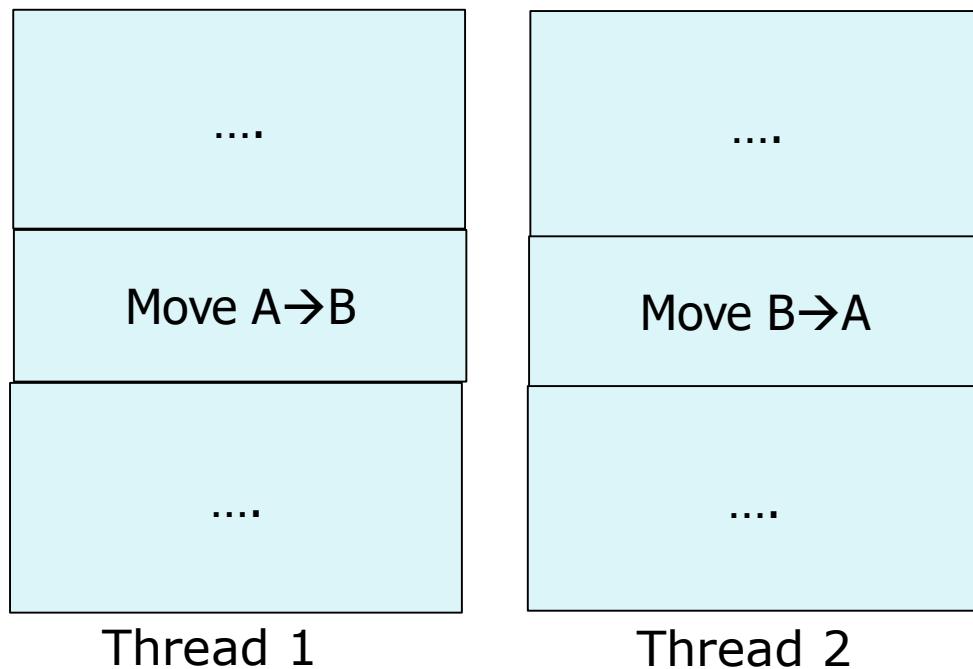
Semaphore **flag**, initialized to 0

# Another Problem: Moving Money Between Accounts

Thread 1 transfers money from account A to B, Thread 2 transfers money from B to A.  
If executed simultaneously, some errors might occur (e.g.,  
Thread 1 executes A--, while Thread 2 executes A++)

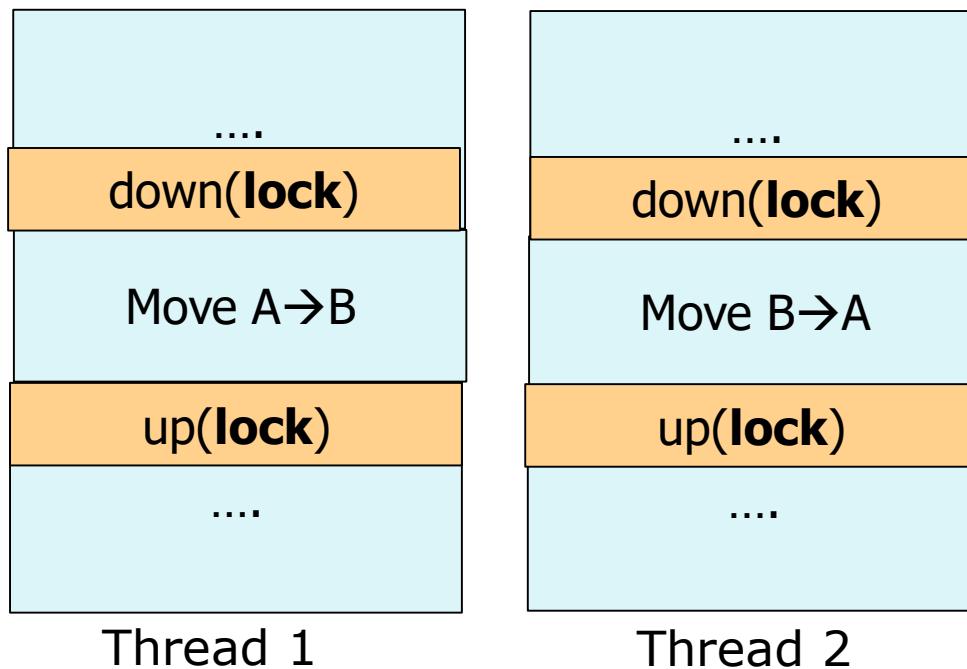
# Another Problem: Moving Money Between Accounts

Thread 1 transfers money from account A to B, Thread 2 transfers money from B to A.  
If executed simultaneously, some errors might occur (e.g.,  
Thread 1 executes A--, while Thread 2 executes A++)



# Another Problem: Moving Money Between Accounts

Thread 1 transfers money from account A to B, Thread 2 transfers money from B to A.  
If executed simultaneously, some errors might occur (e.g., Thread 1 executes A--, while Thread 2 executes A++)

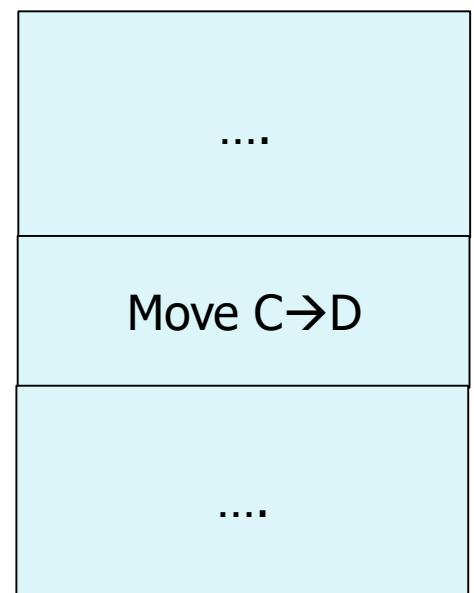
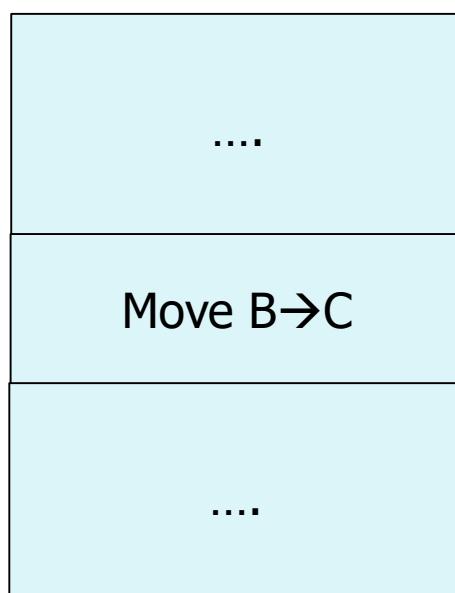
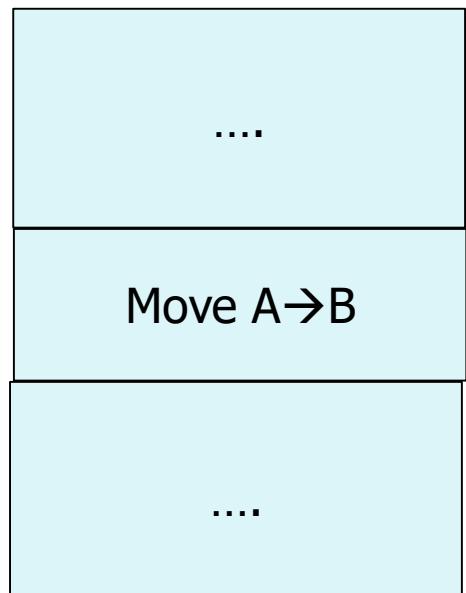


Semaphore **lock**, initialized to 1

*Mutual exclusion, where the Move is the critical section*

# Another Problem: Moving Money Between Accounts

- Three Threads



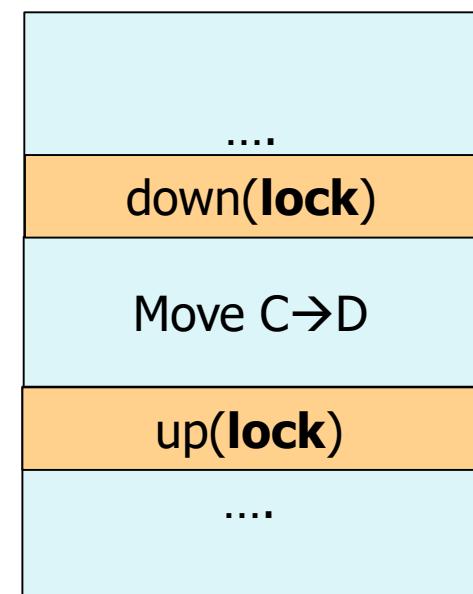
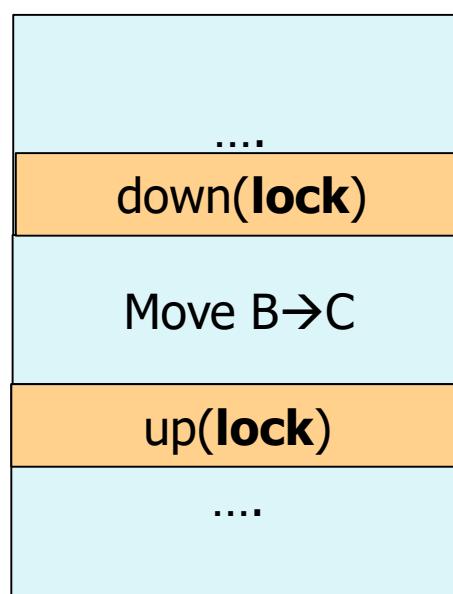
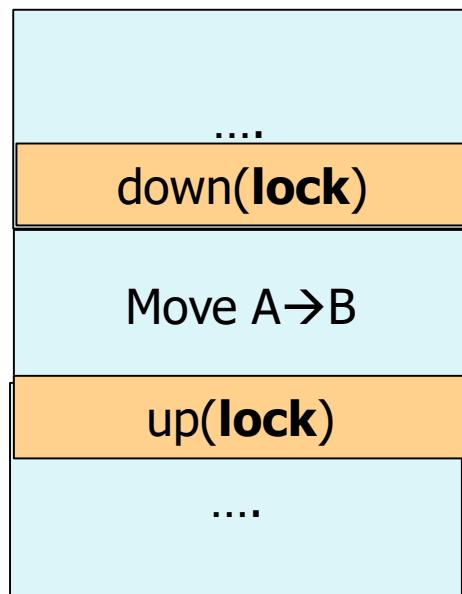
Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Three Threads
- Solution 1: mutual exclusion



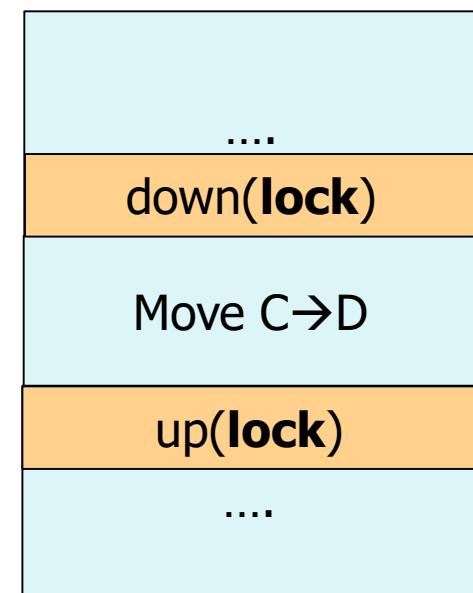
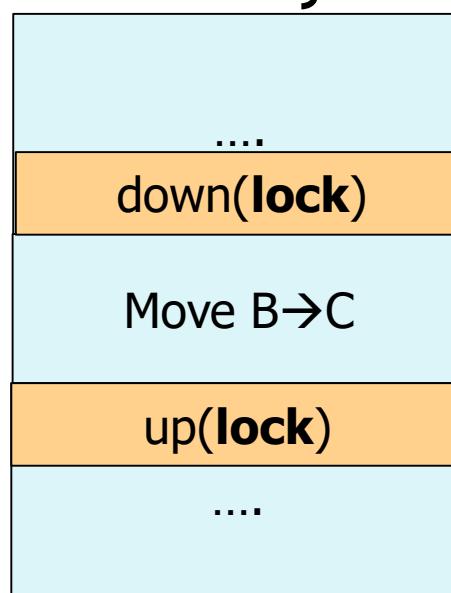
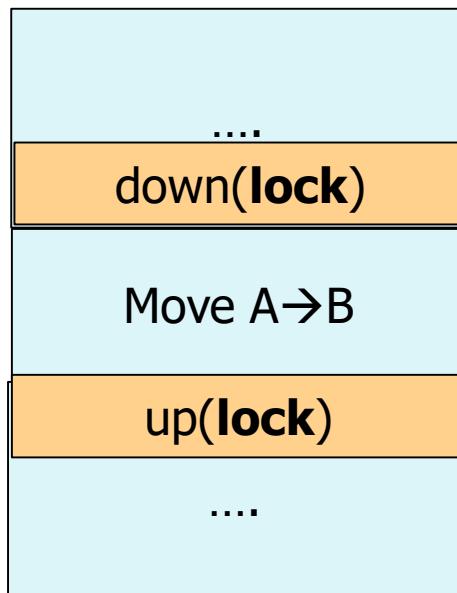
Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Three Threads
- Solution 1: mutual exclusion
  - But why Thread 1 and Thread 3 cannot be executed concurrently?



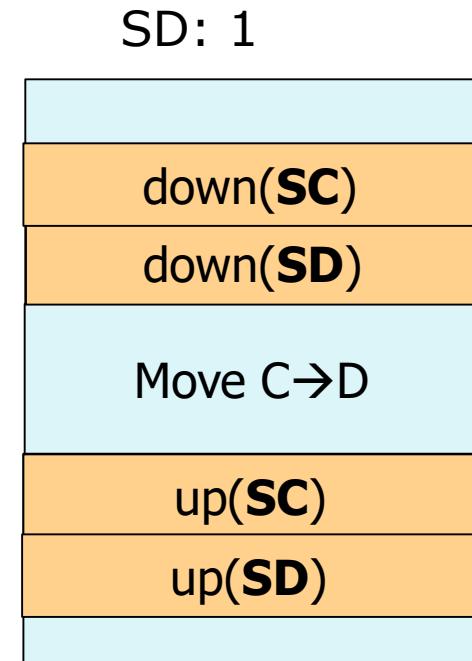
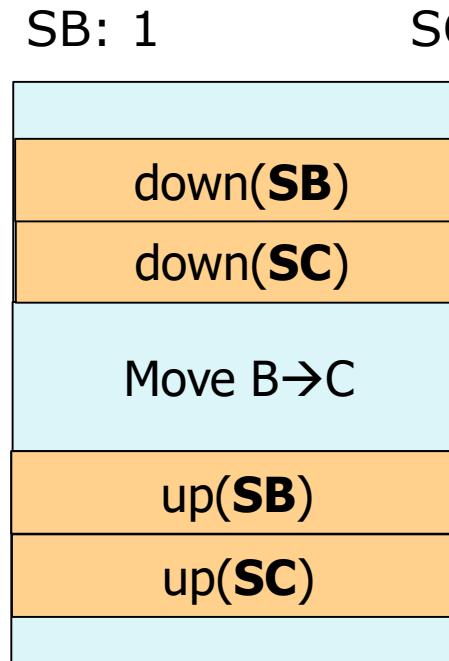
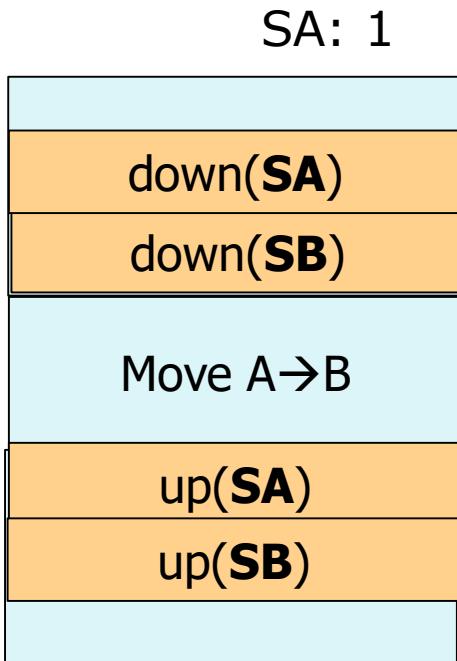
Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



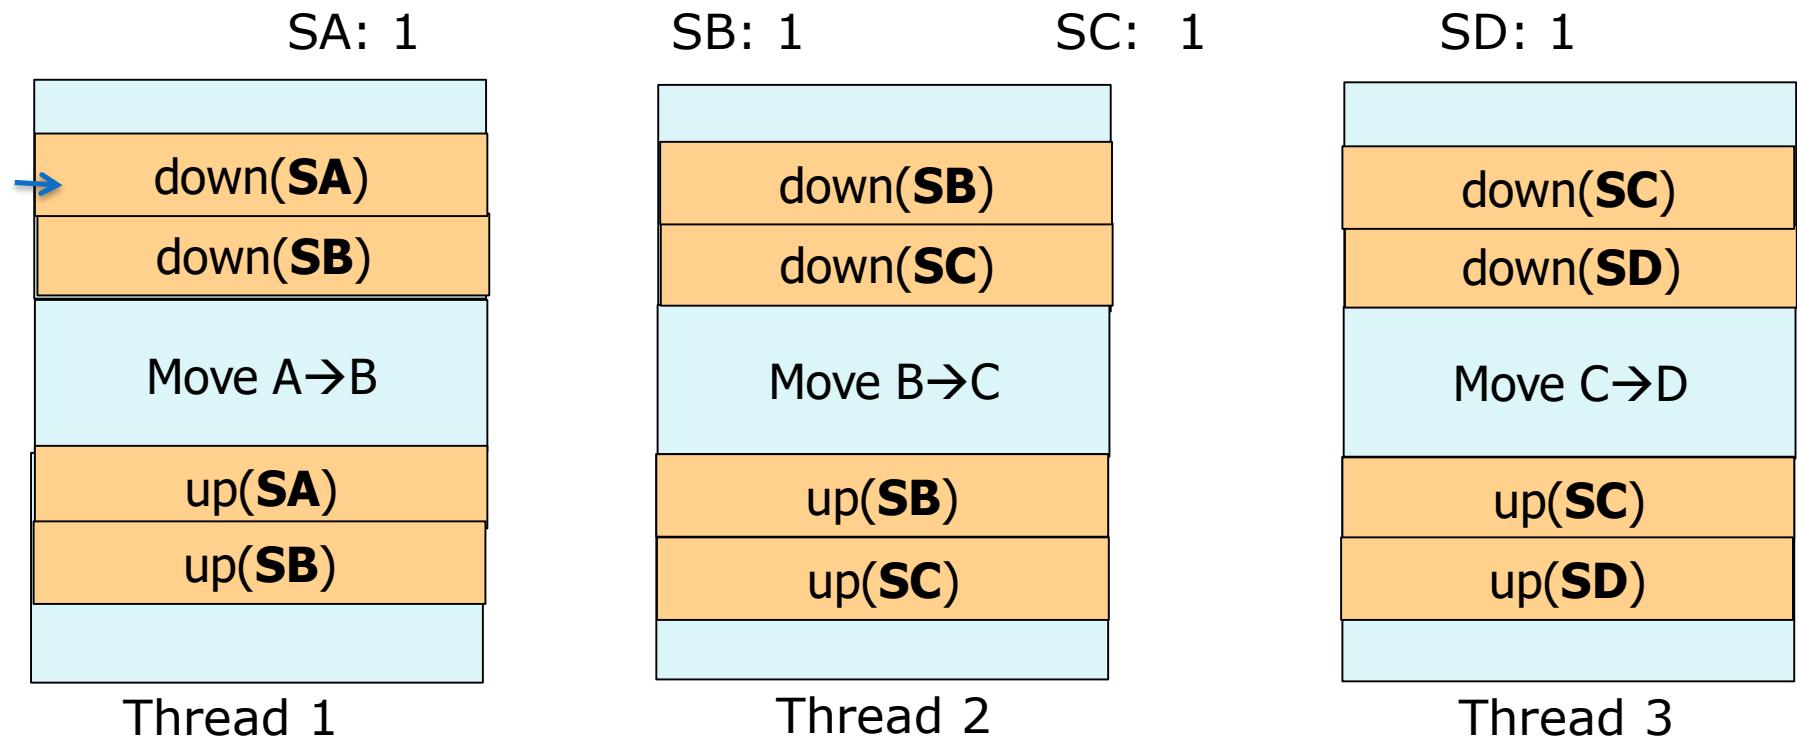
Thread 1

Thread 2

Thread 3

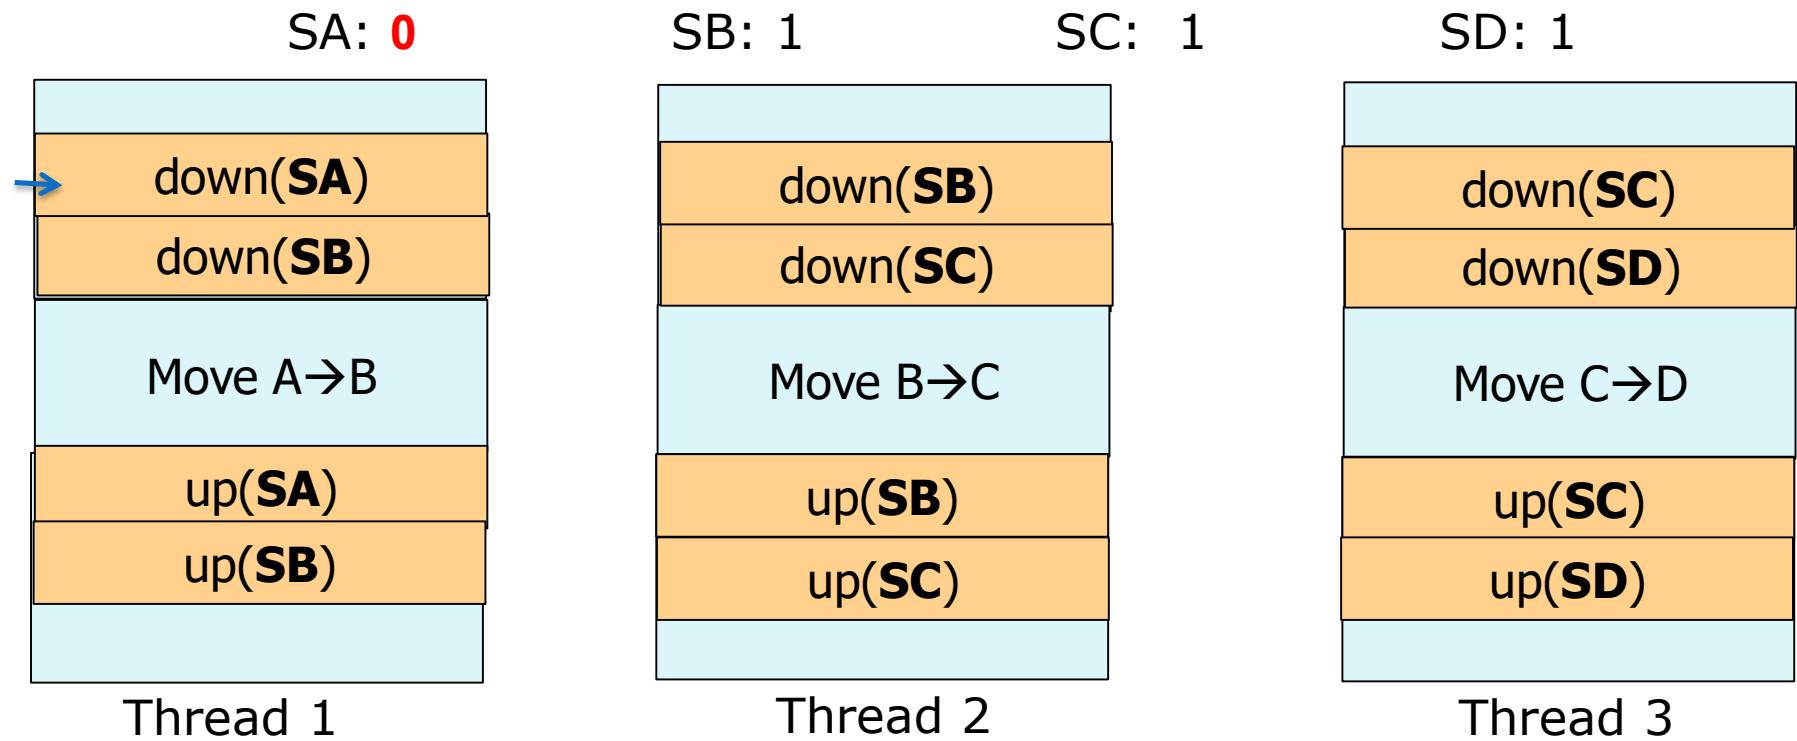
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



# Another Problem: Moving Money Between Accounts

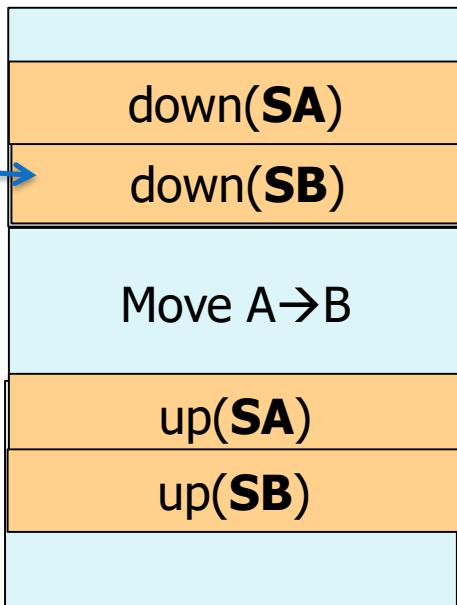
- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



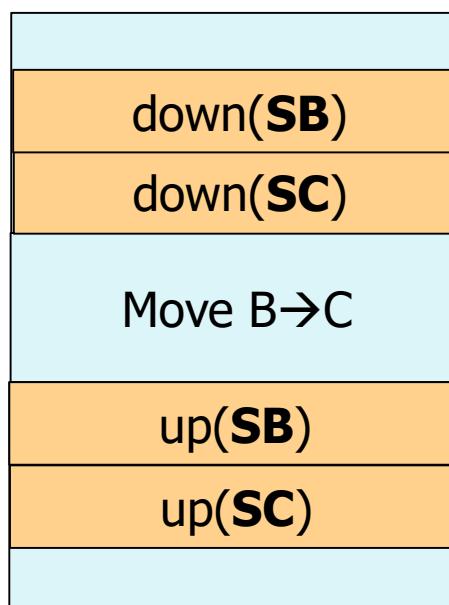
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.

SA: **0**

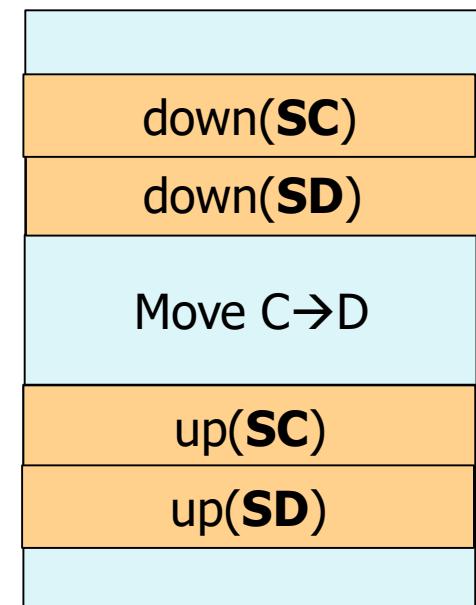


SB: 1



SC: 1

SD: 1



Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.

SA: **0**

SB: **0**

SC: 1

SD: 1

down(**SA**)

down(**SB**)

Move A→B

up(**SA**)

up(**SB**)

Thread 1

down(**SB**)

down(**SC**)

Move B→C

up(**SB**)

up(**SC**)

Thread 2

down(**SC**)

down(**SD**)

Move C→D

up(**SC**)

up(**SD**)

Thread 3

# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.

SA: **0**

SB: **0**

SC: 1

SD: 1

down(**SA**)

down(**SB**)

Move A→B

up(**SA**)

up(**SB**)

down(**SB**)

down(**SC**)

Move B→C

up(**SB**)

up(**SC**)

down(**SC**)

down(**SD**)

Move C→D

up(**SC**)

up(**SD**)

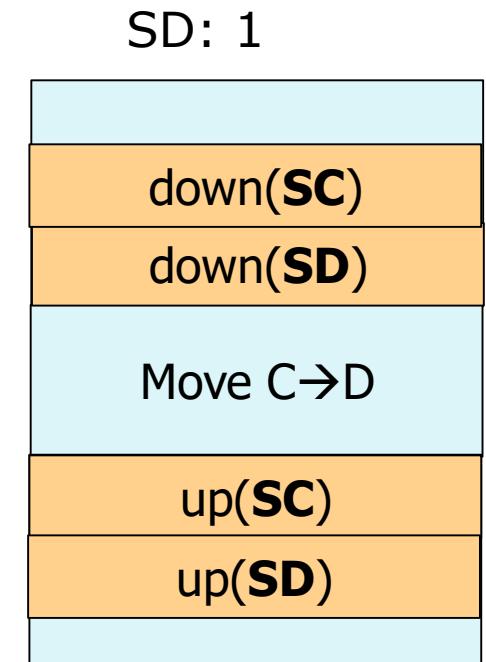
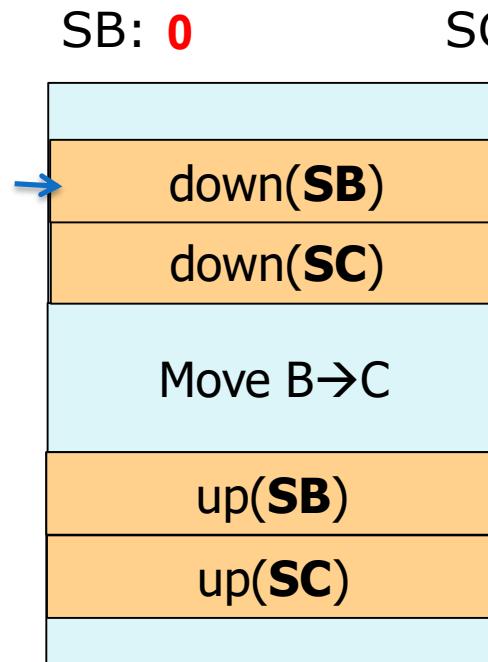
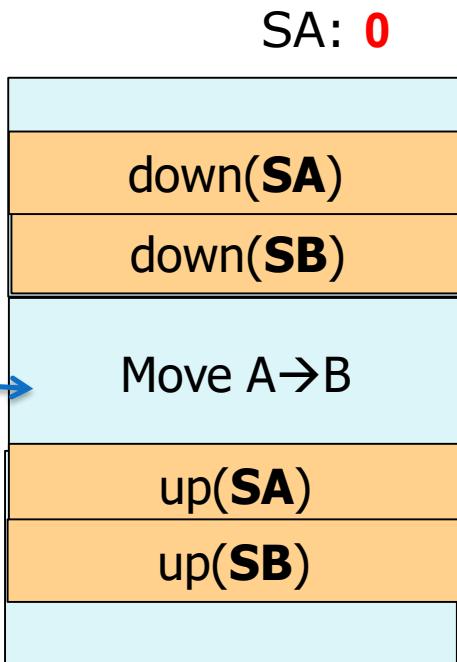
Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



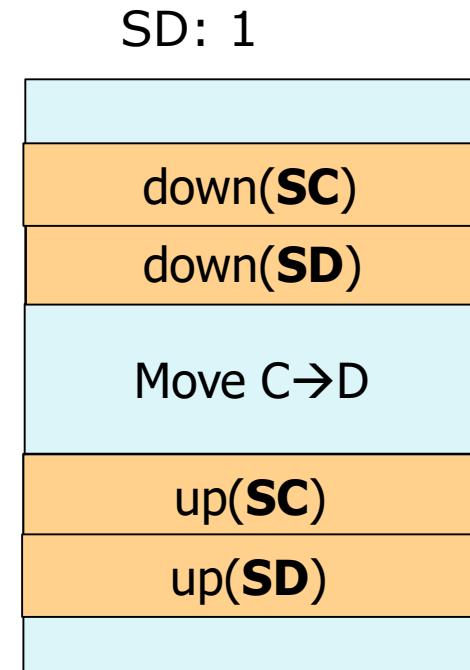
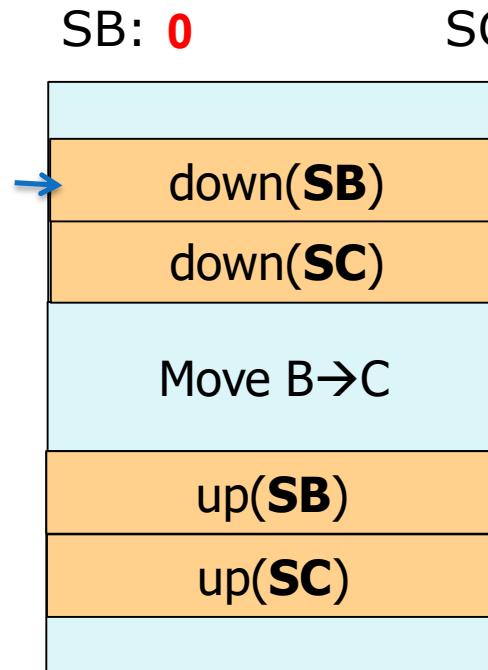
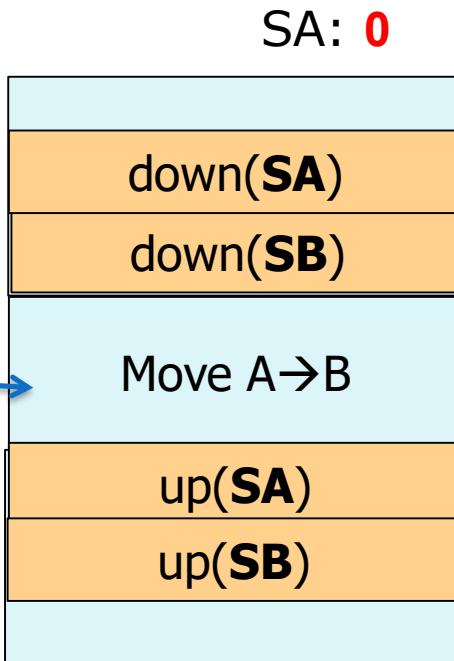
Thread 1

Thread 2

Thread 3

# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



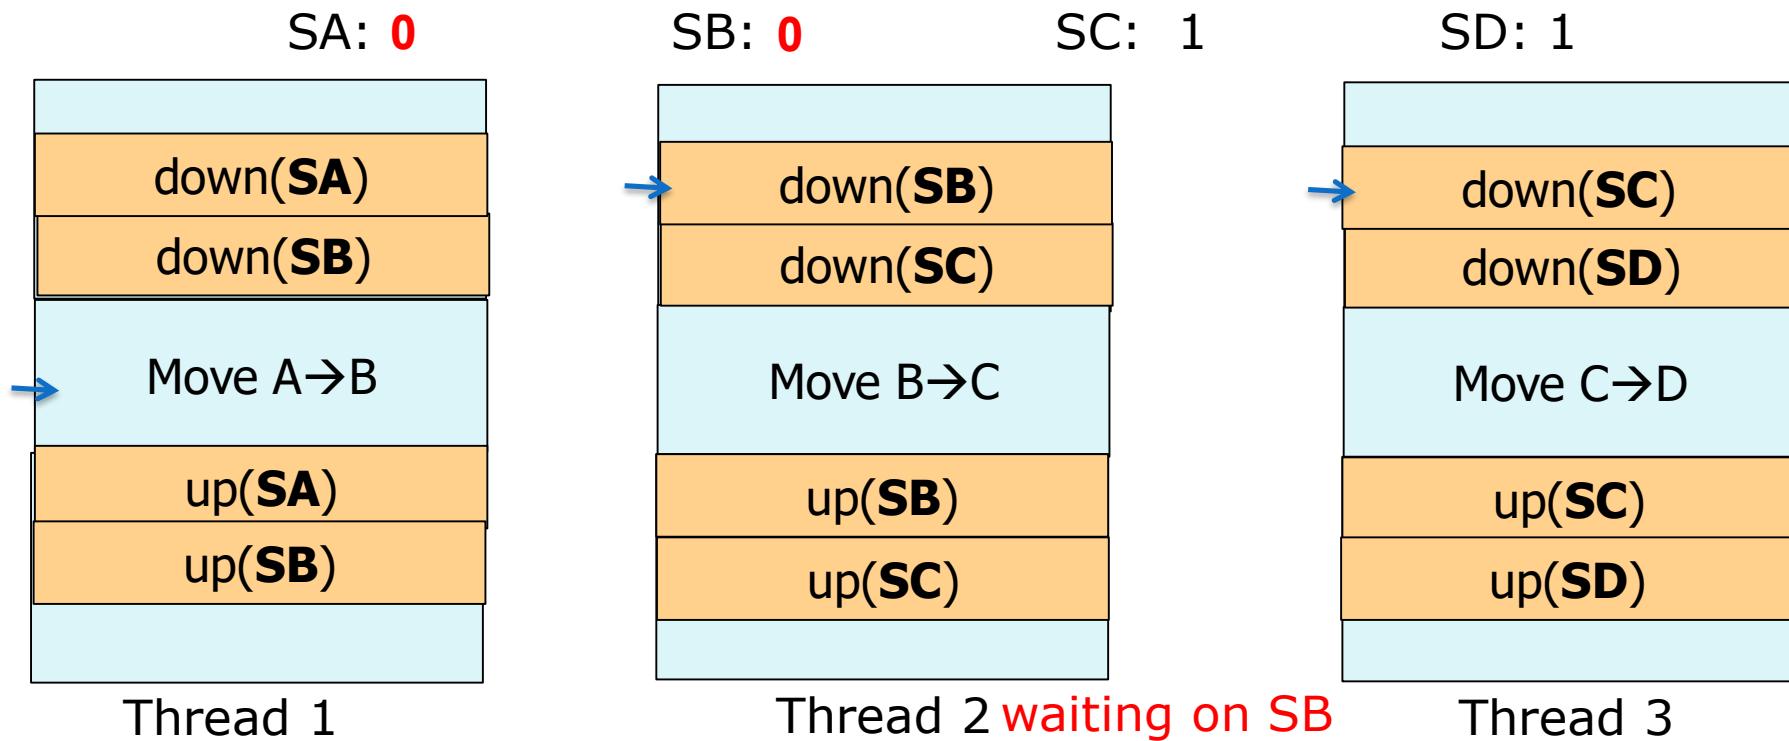
Thread 1

Thread 2 **waiting on SB**

Thread 3

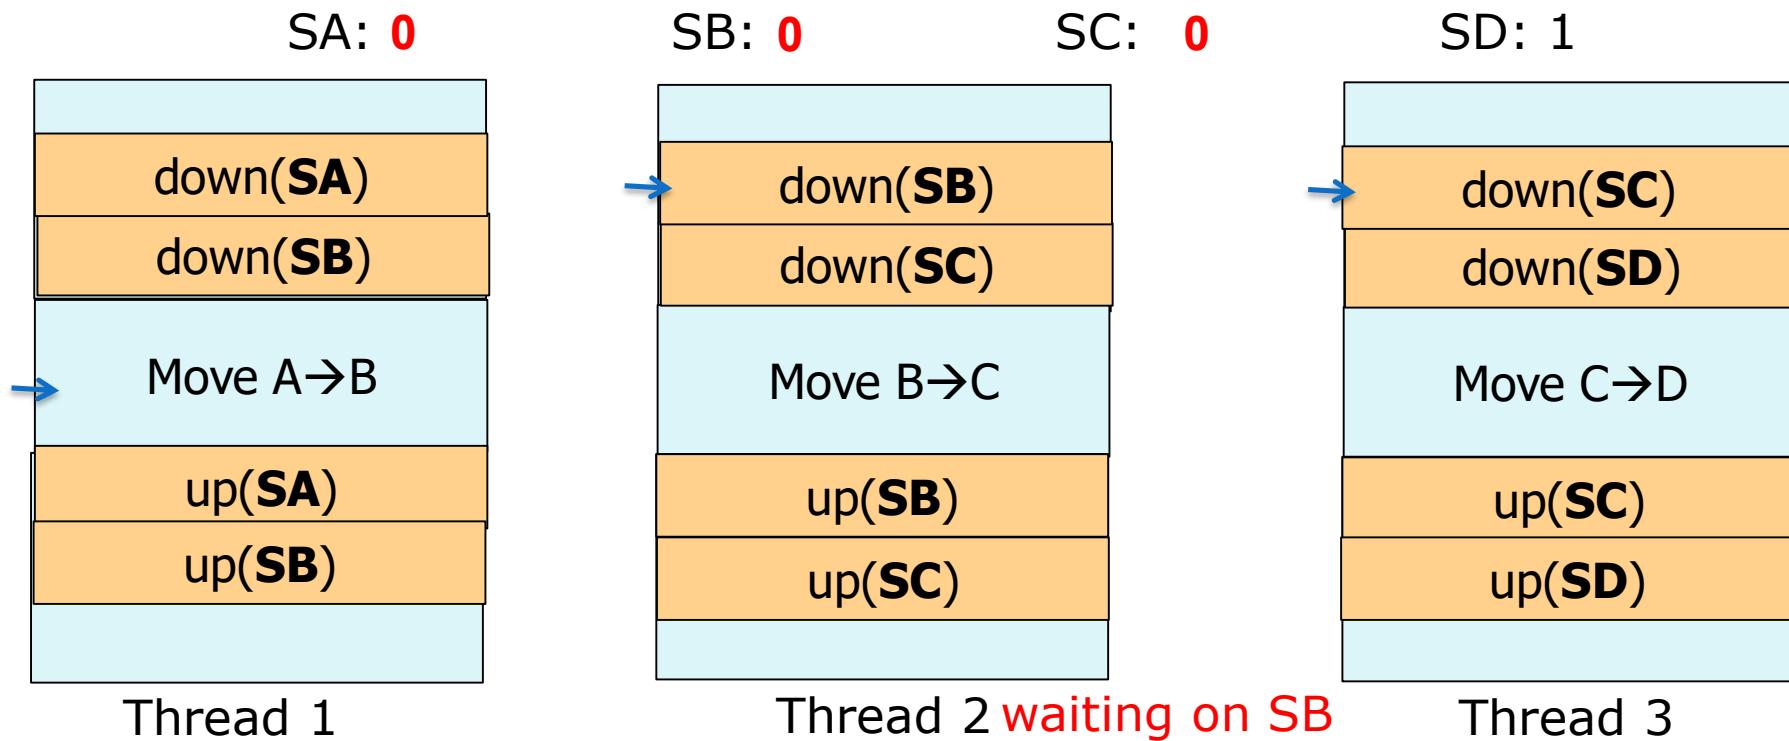
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



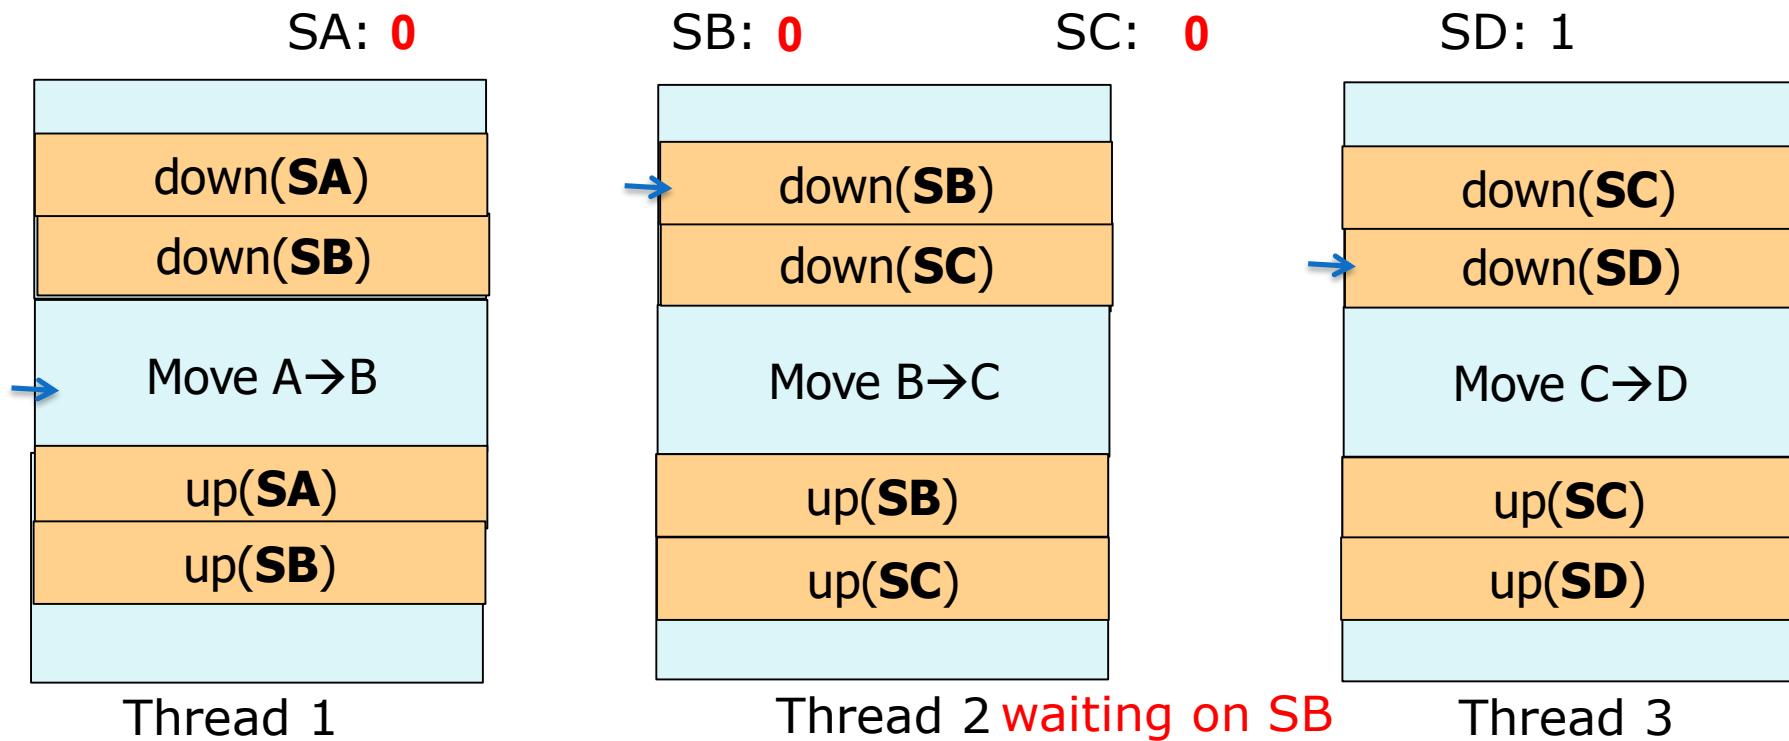
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



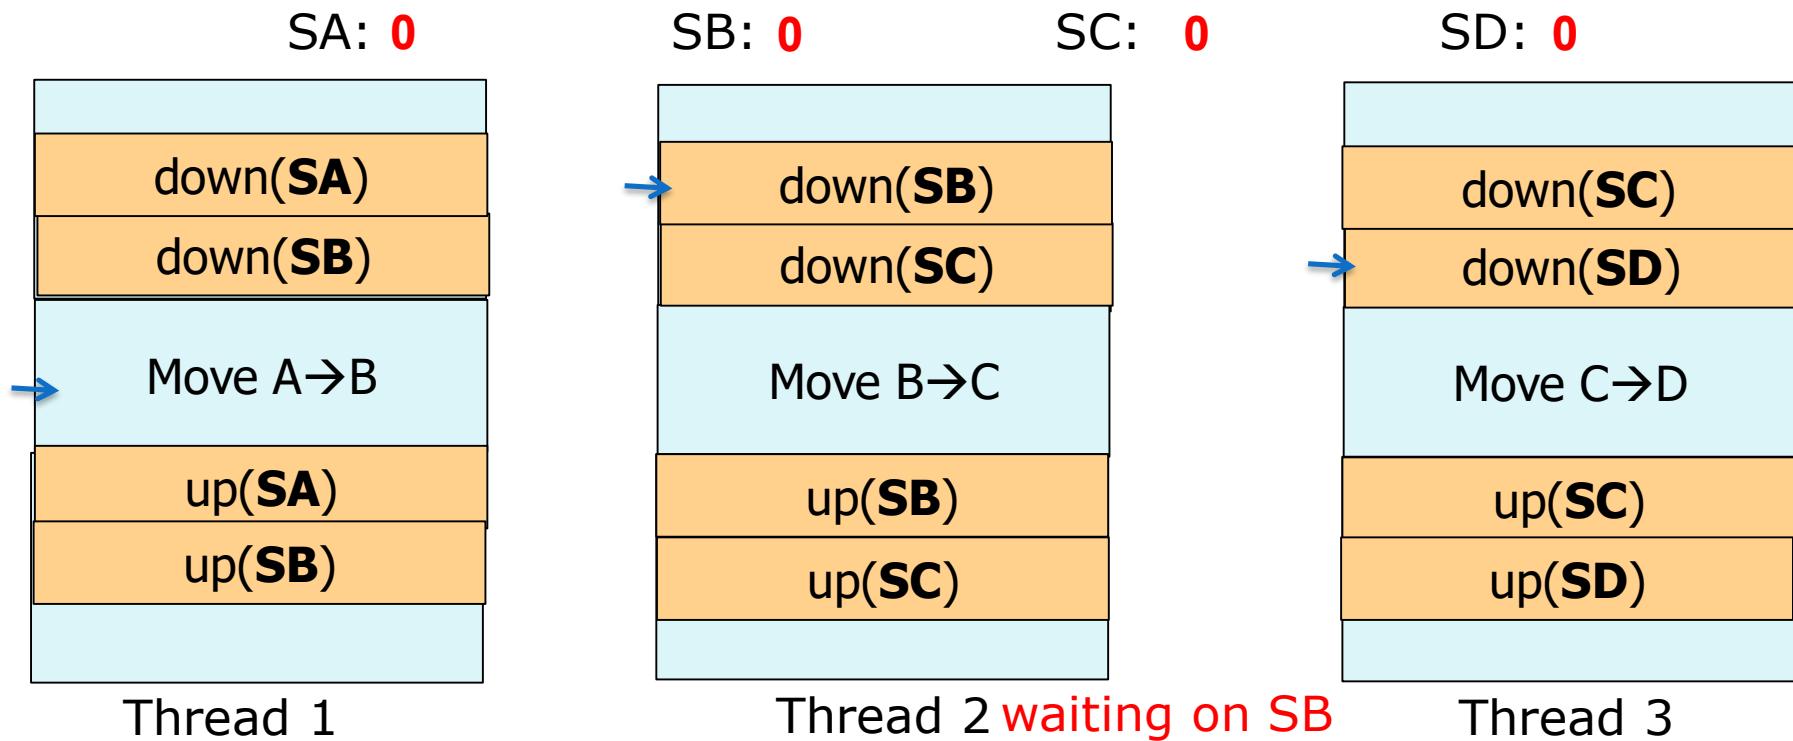
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



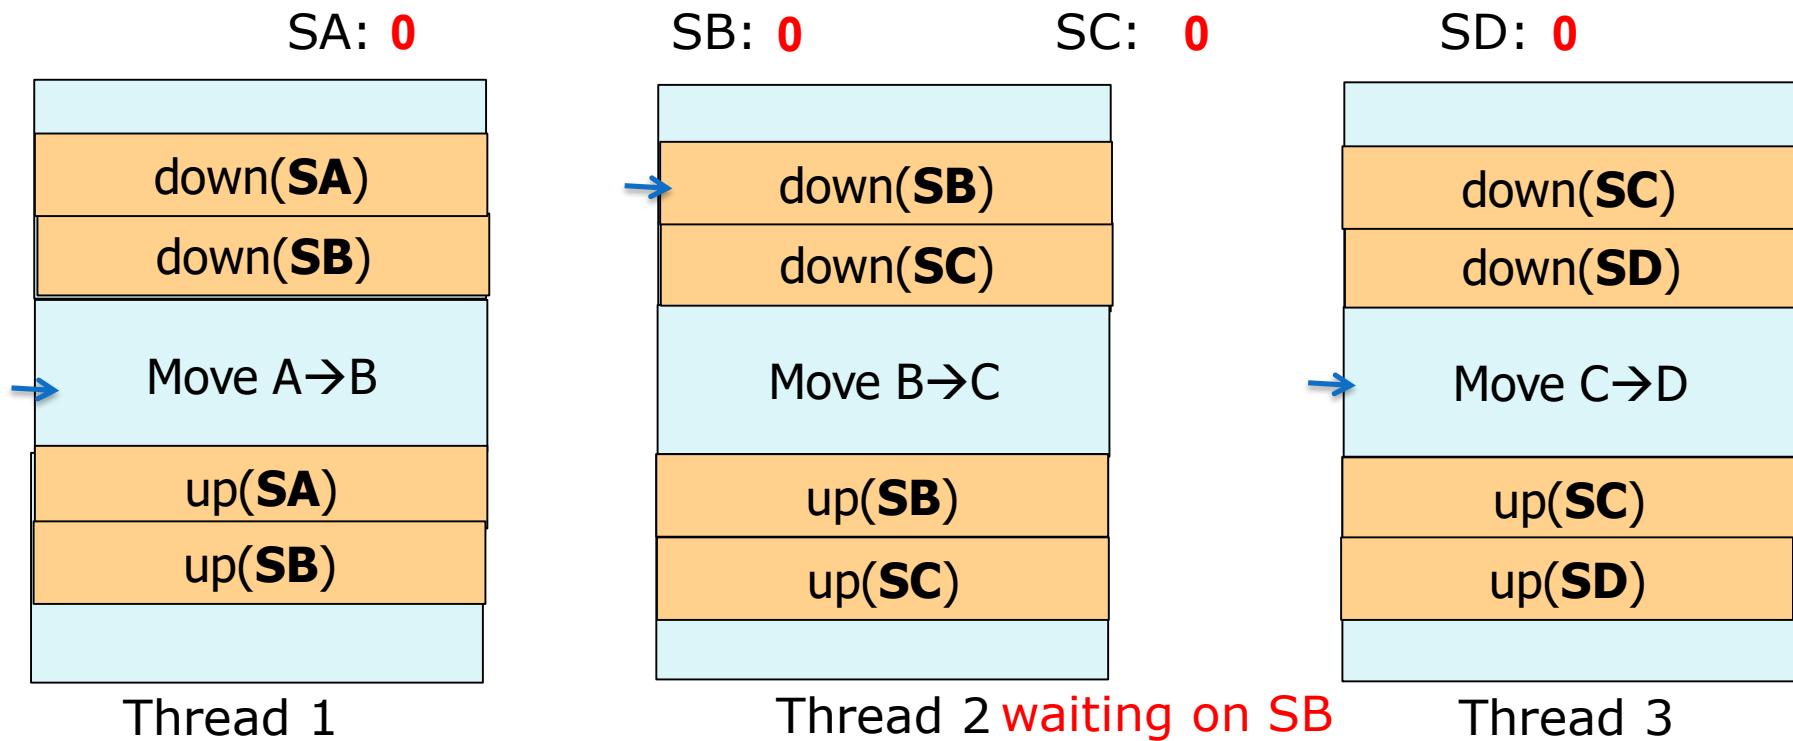
# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



# Another Problem: Moving Money Between Accounts

- Solution 2: Lock each account separately
  - Semaphore for each account: SA, SB, SC, SD. All initialized to 1.



# Semaphores are not Silver Bullets

SA: 1

SB: 1

SC: 1

SD: 1

down(**SA**)

down(**SB**)

Move A→B

up(**SA**)

up(**SB**)

down(**SB**)

down(**SC**)

Move B→C

up(**SB**)

up(**SC**)

down(**SC**)

down(**SD**)

Move C→D

up(**SC**)

up(**SD**)

down(**SD**)

down(**SA**)

Move D→A

up(**SD**)

up(**SA**)

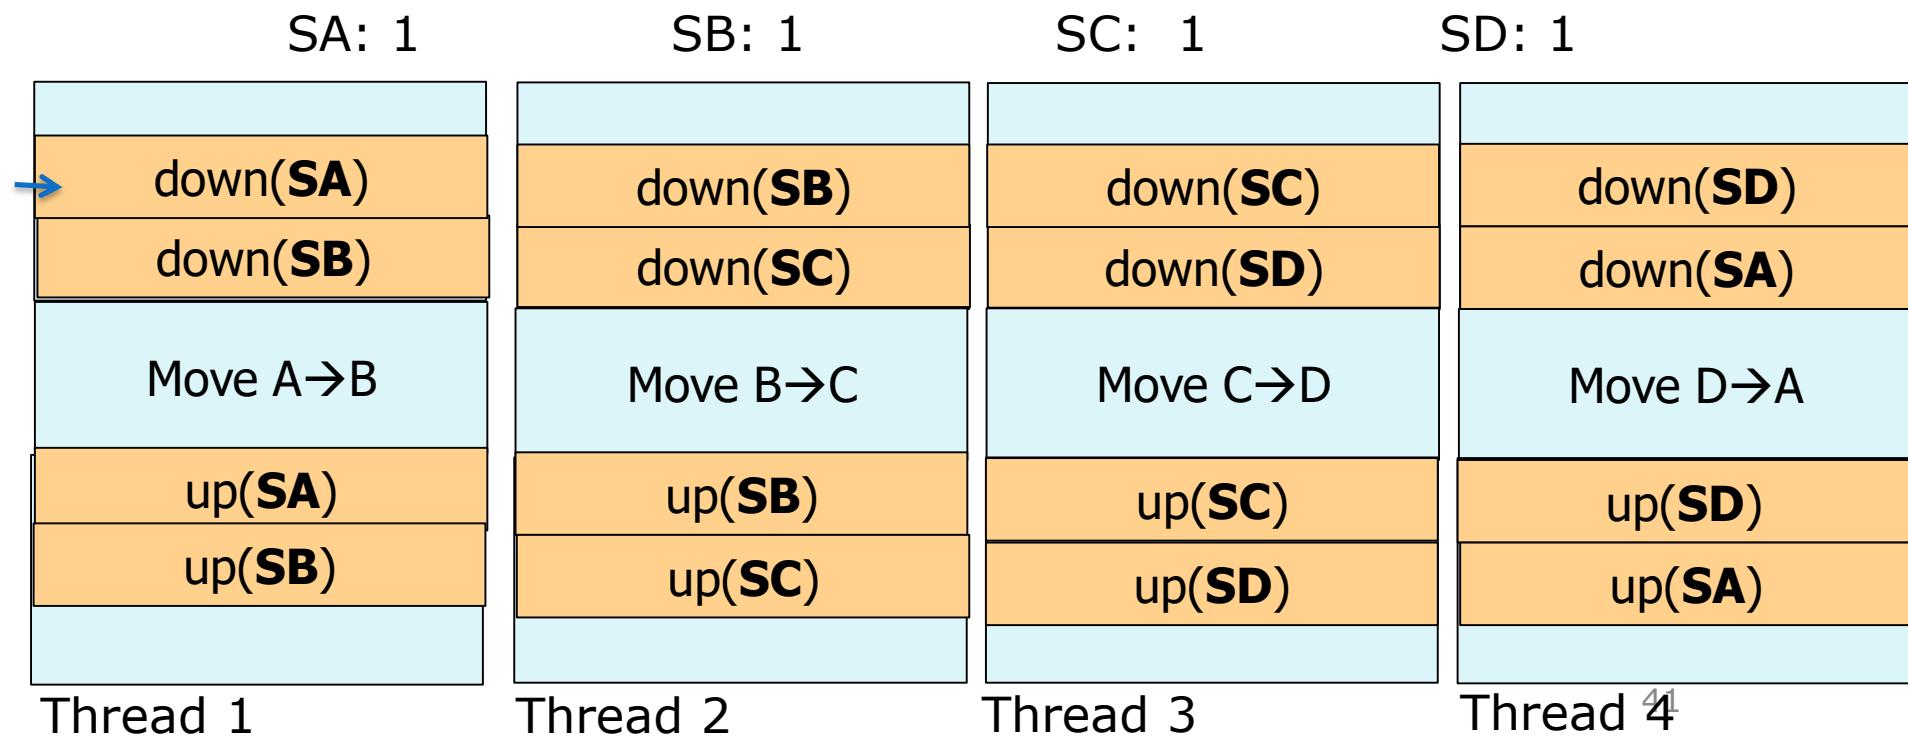
Thread 1

Thread 2

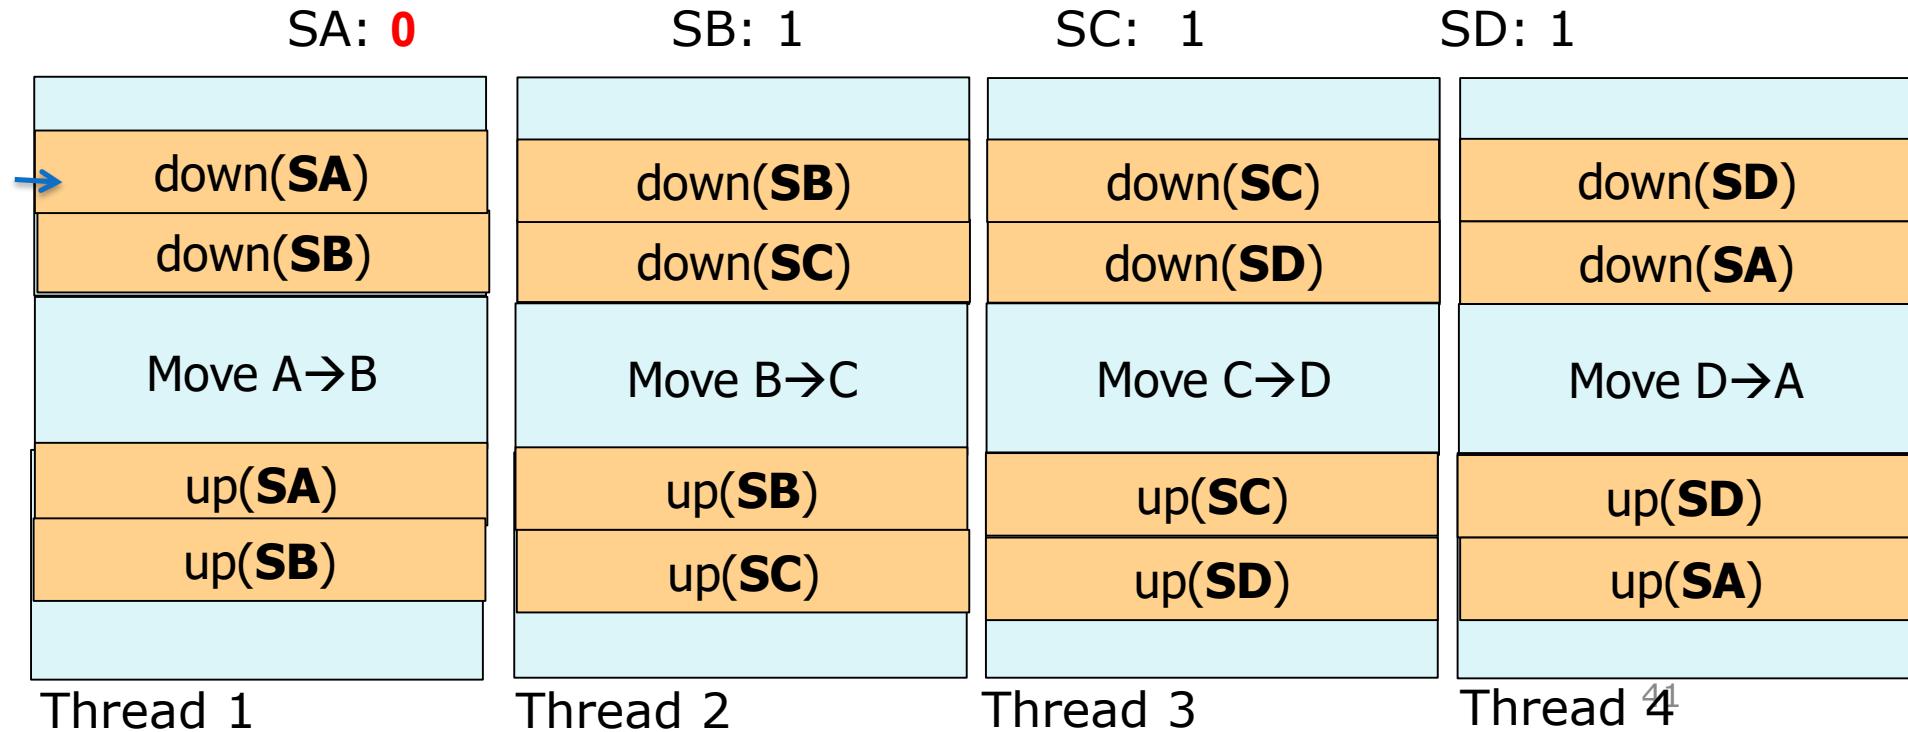
Thread 3

Thread 4<sup>41</sup>

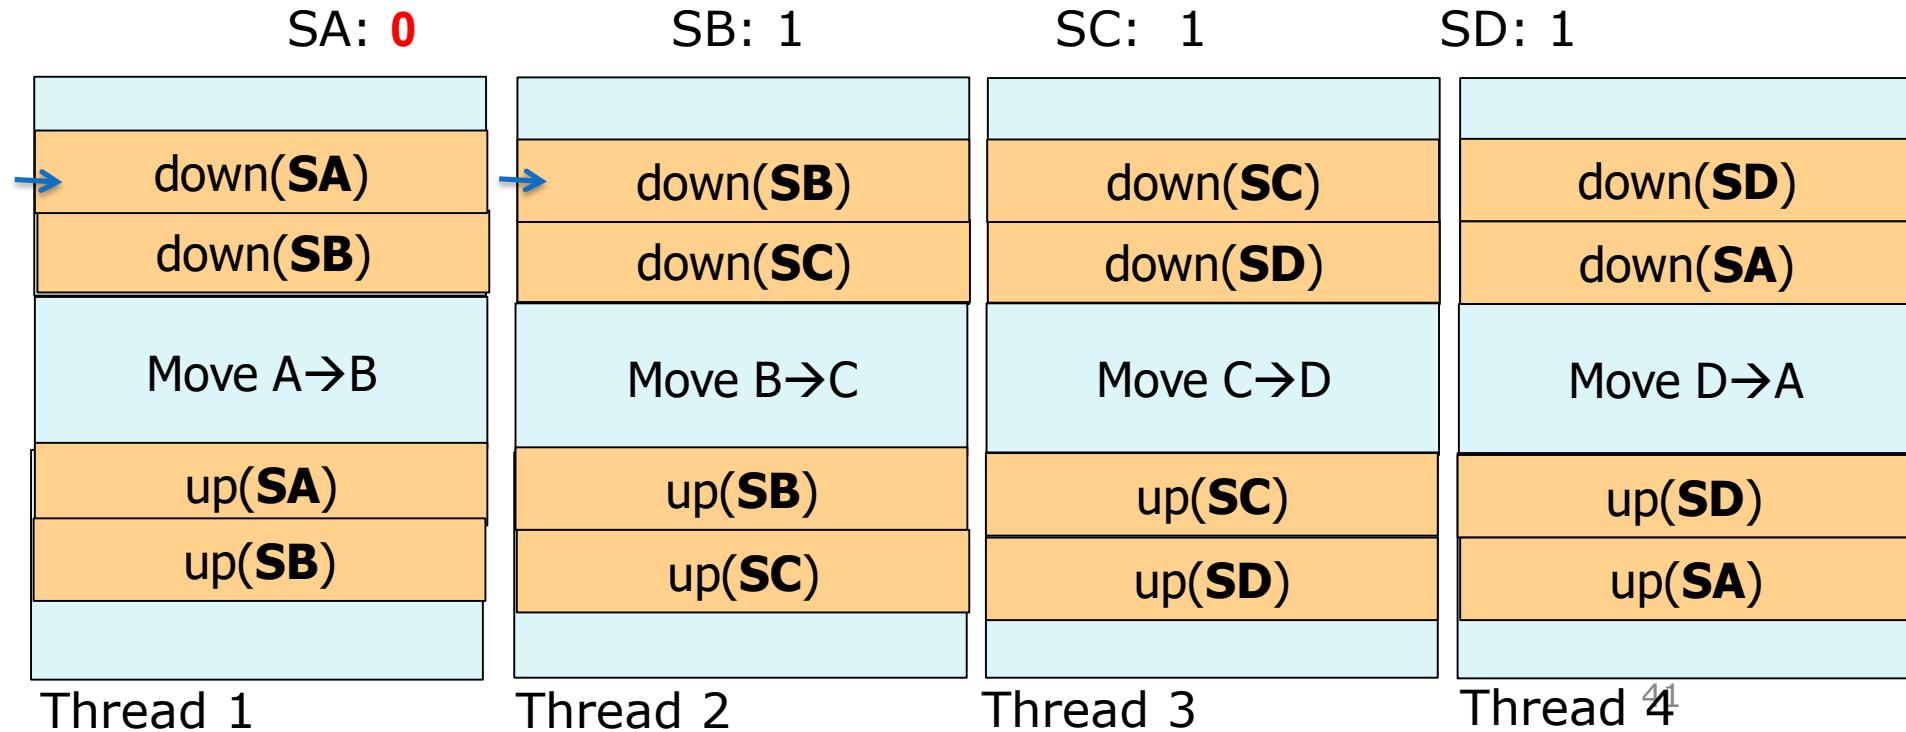
# Semaphores are not Silver Bullets



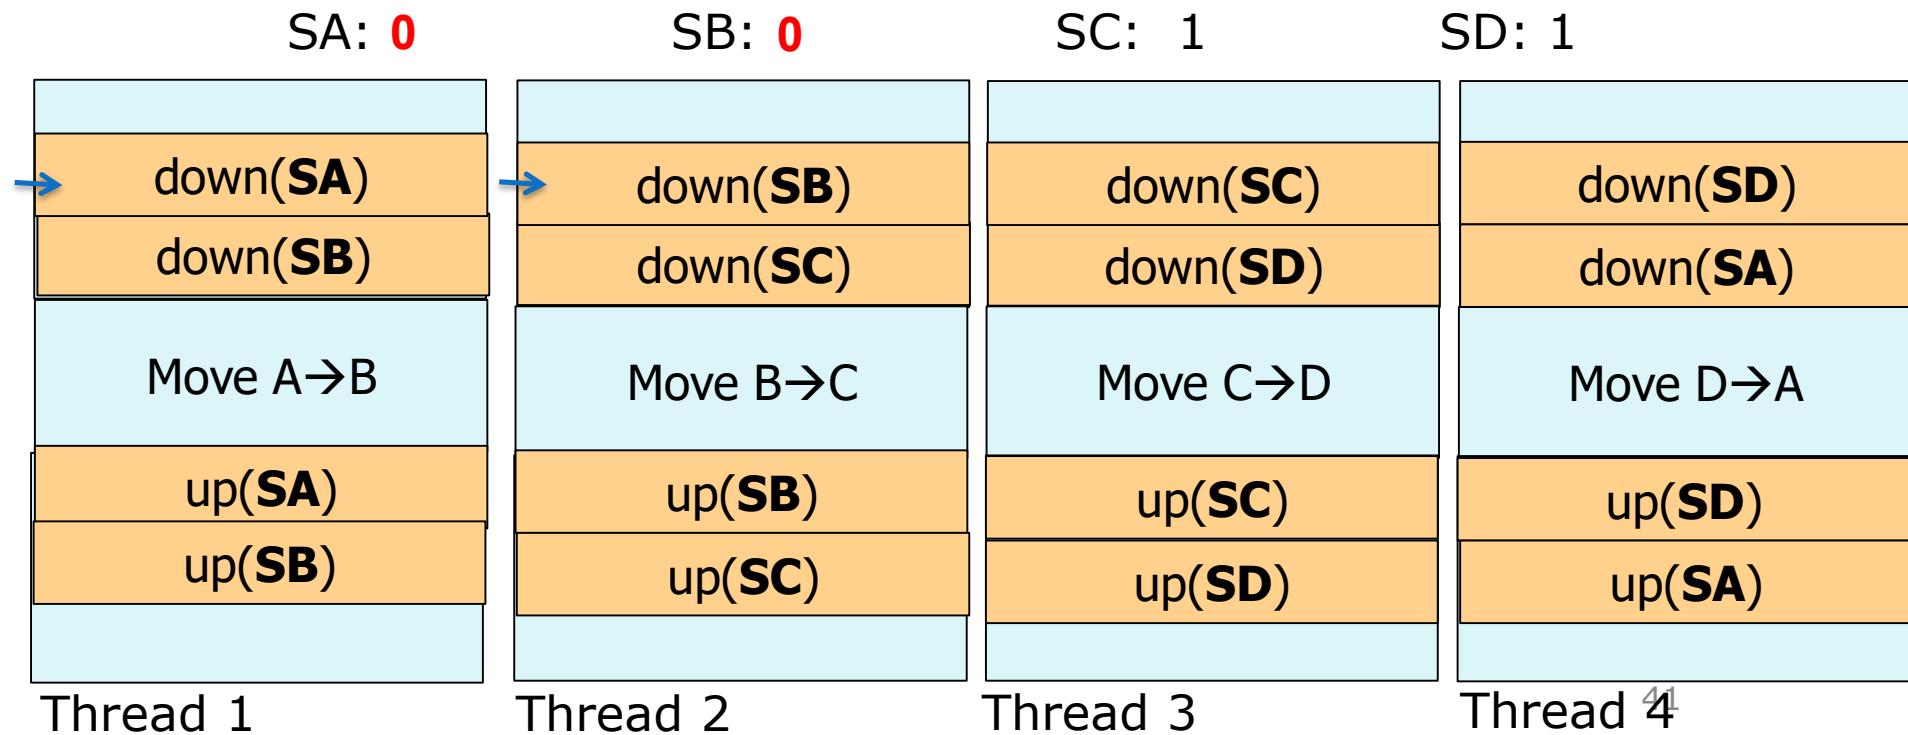
# Semaphores are not Silver Bullets



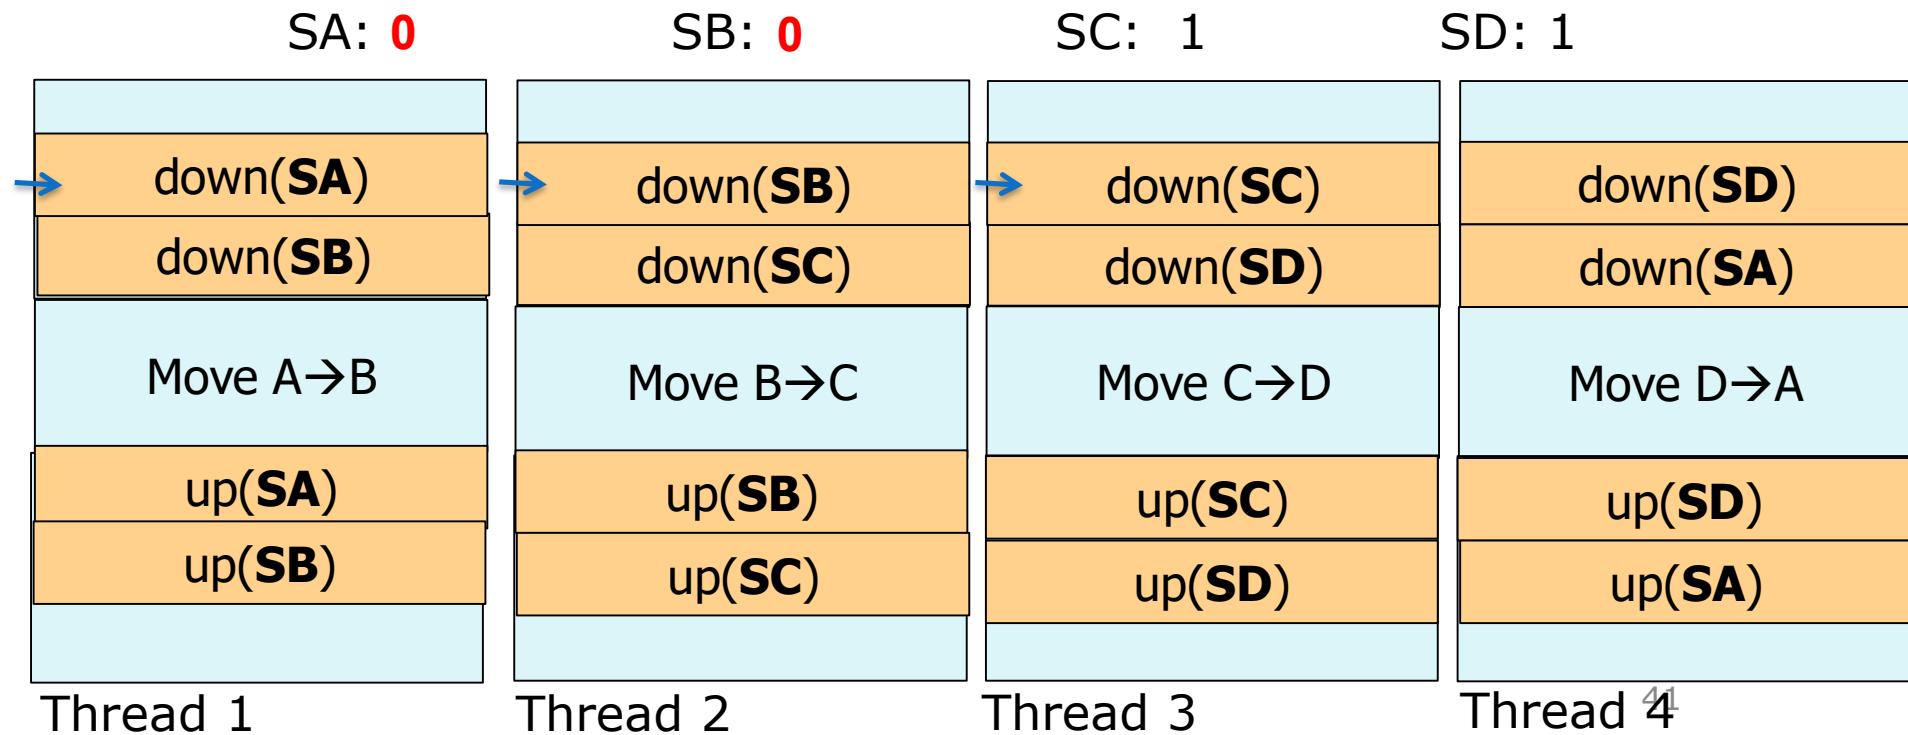
# Semaphores are not Silver Bullets



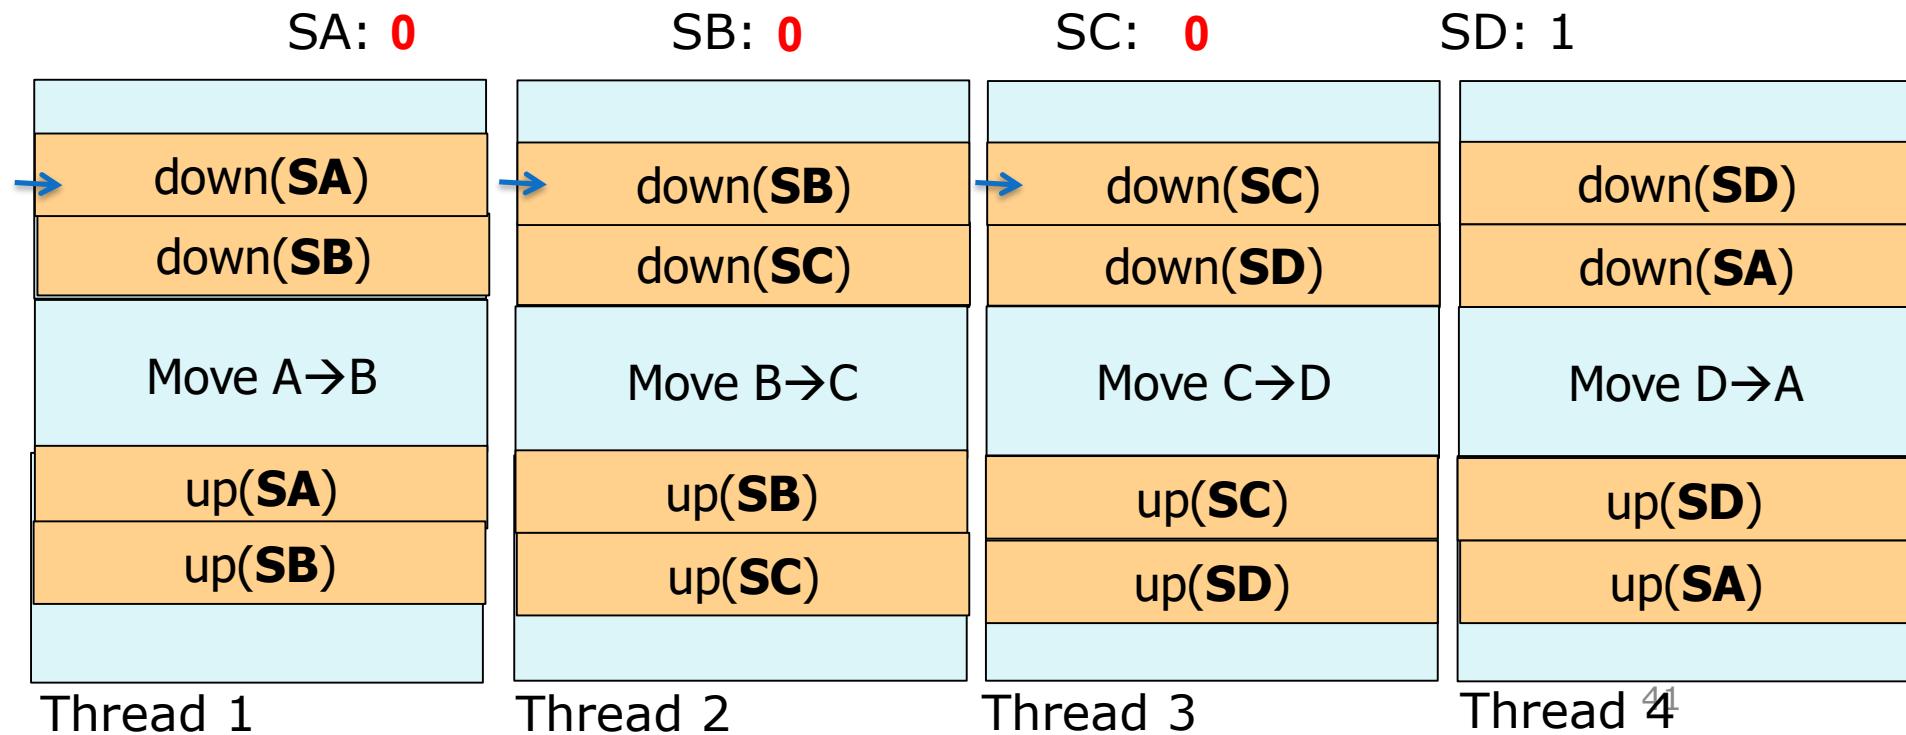
# Semaphores are not Silver Bullets



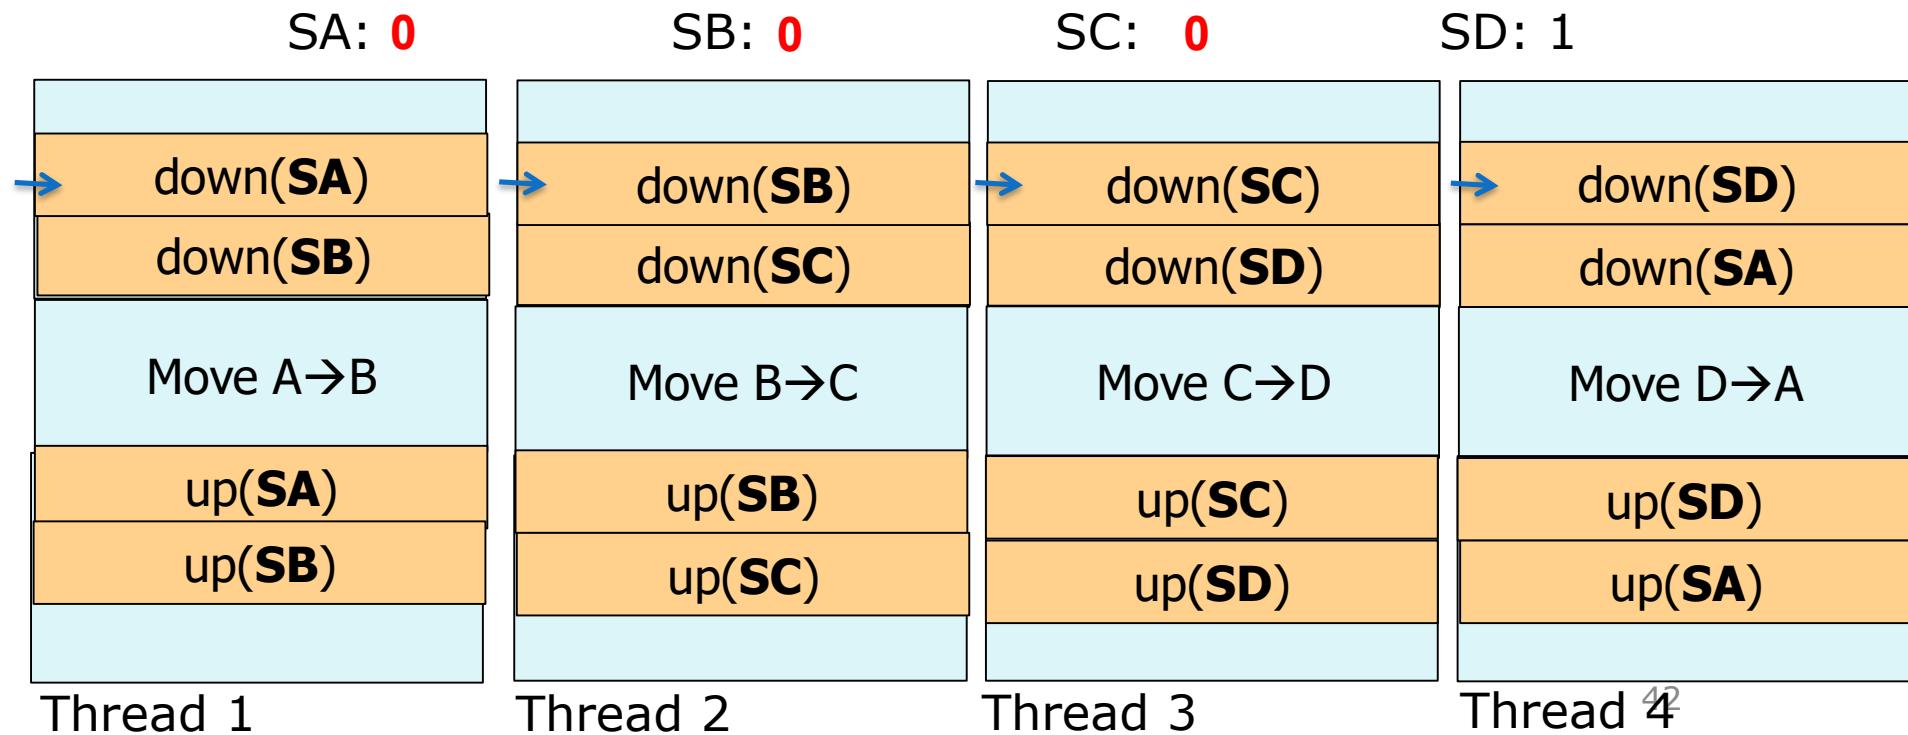
# Semaphores are not Silver Bullets



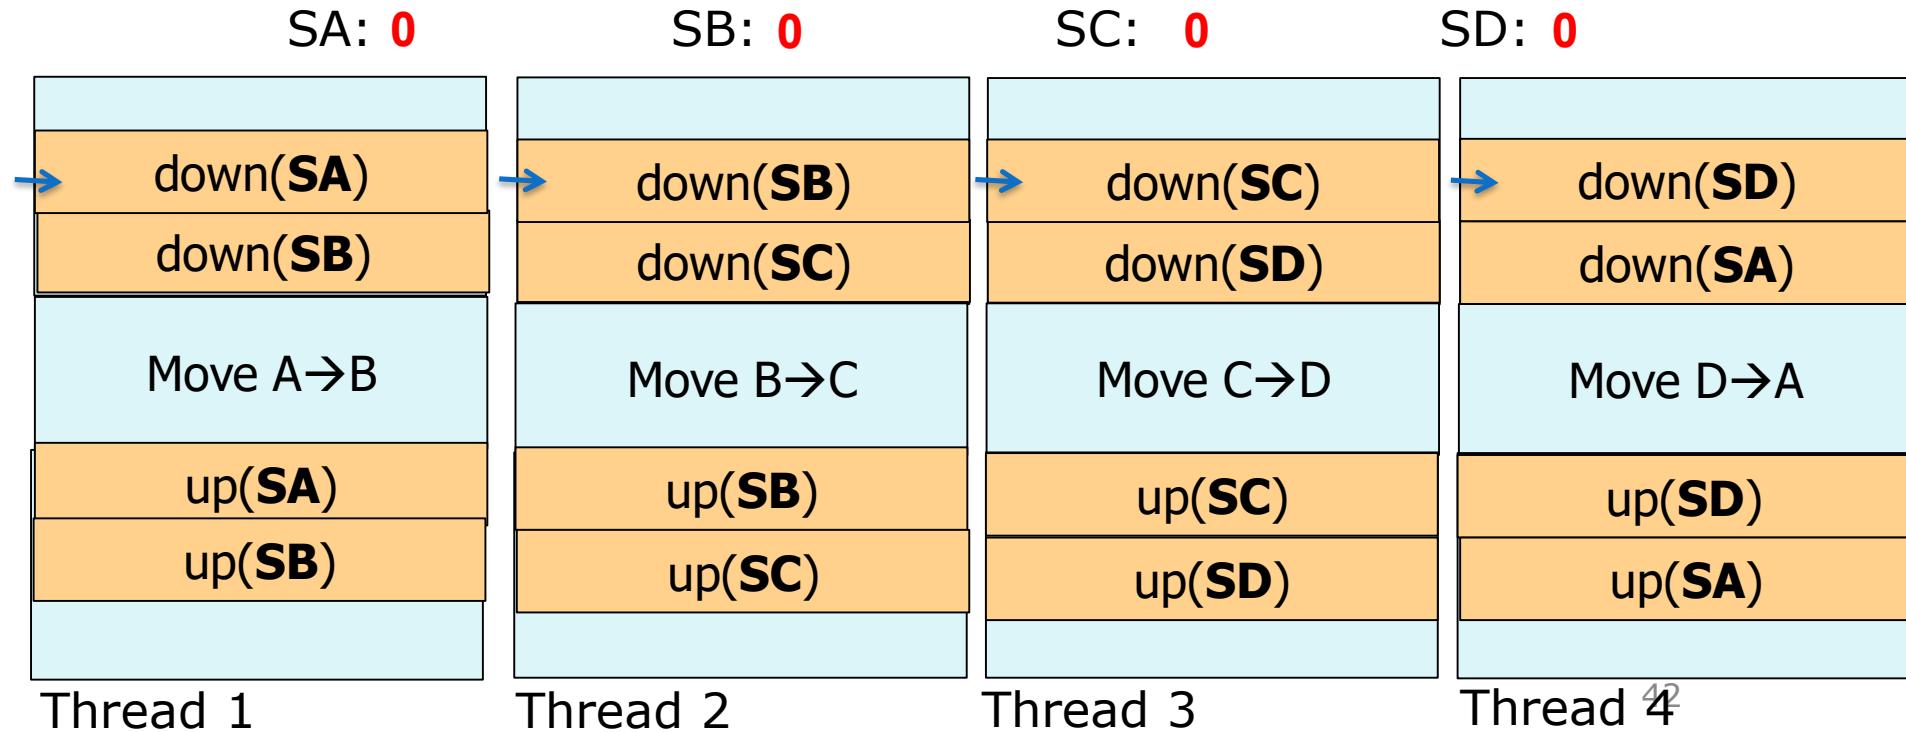
# Semaphores are not Silver Bullets



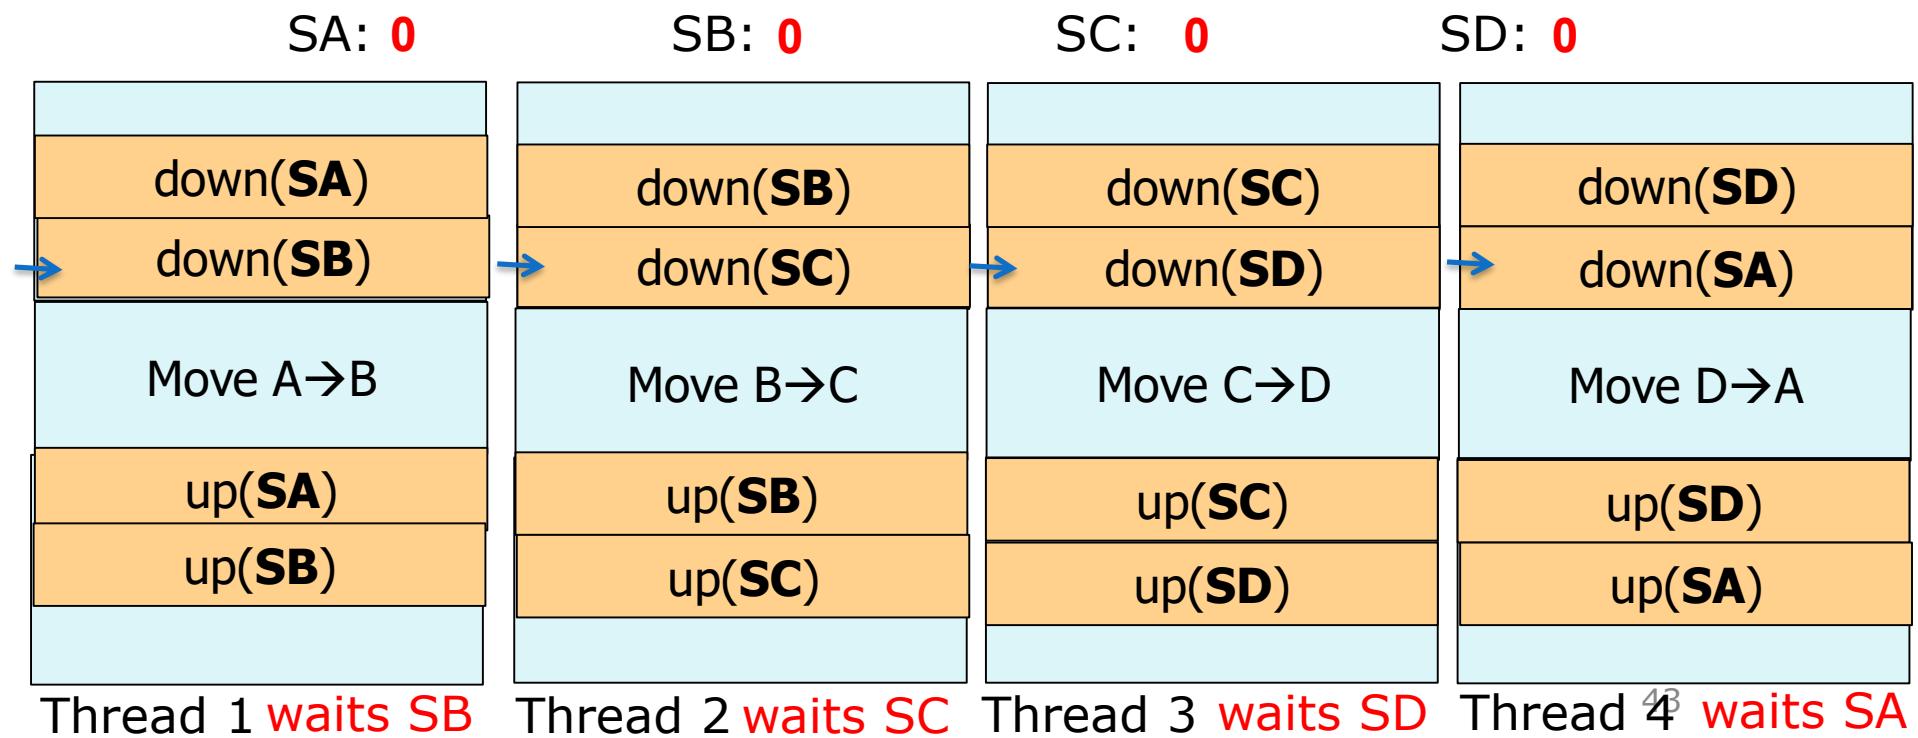
# Semaphores are not Silver Bullets



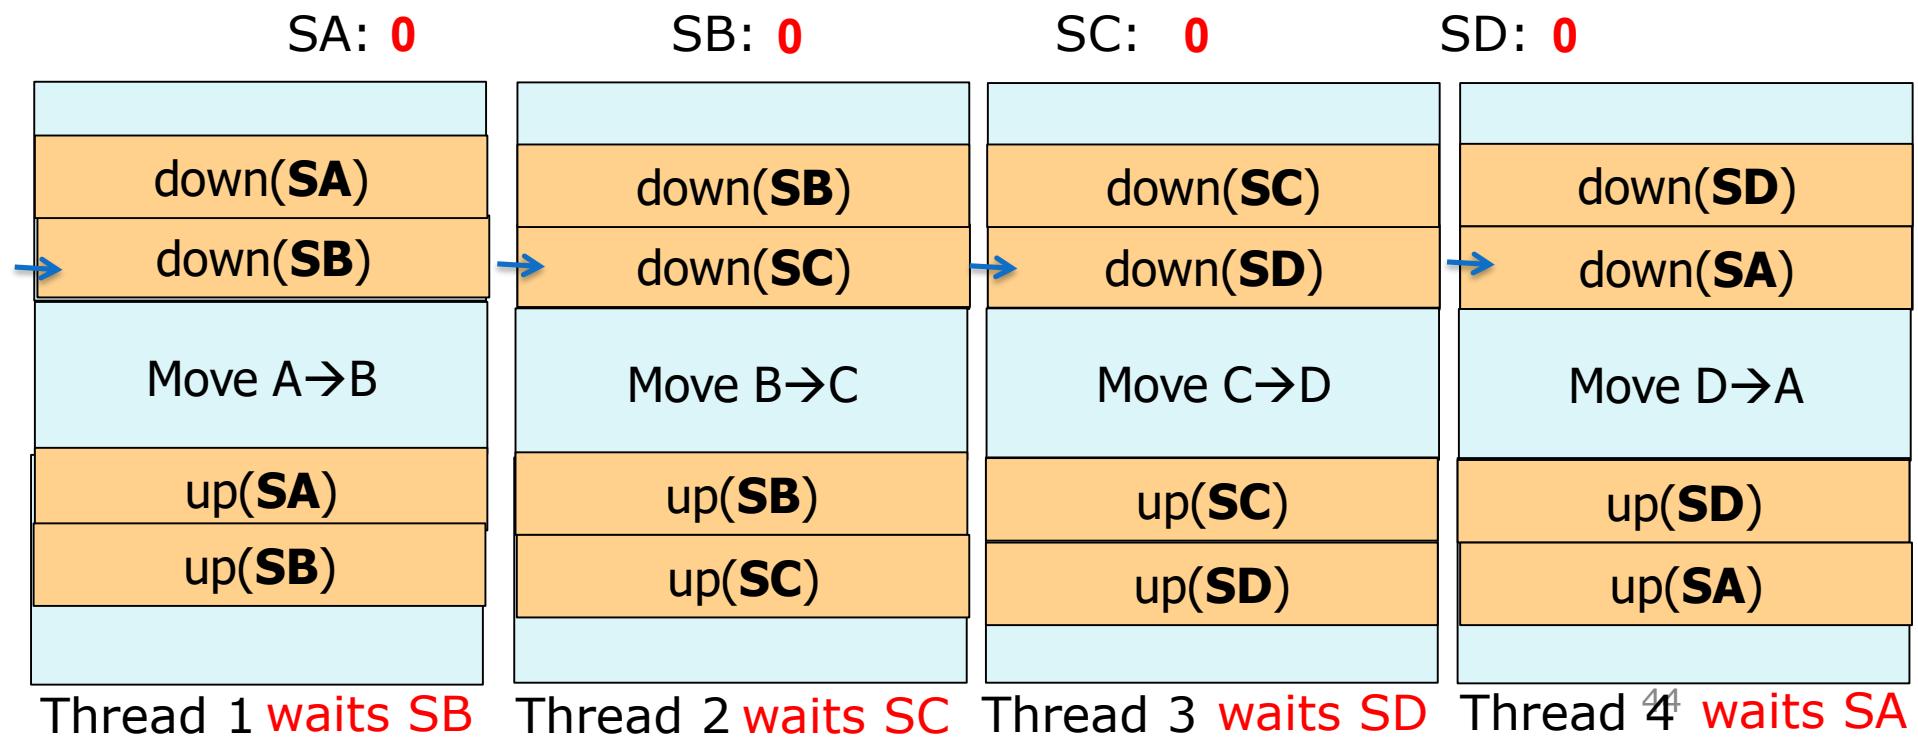
# Semaphores are not Silver Bullets



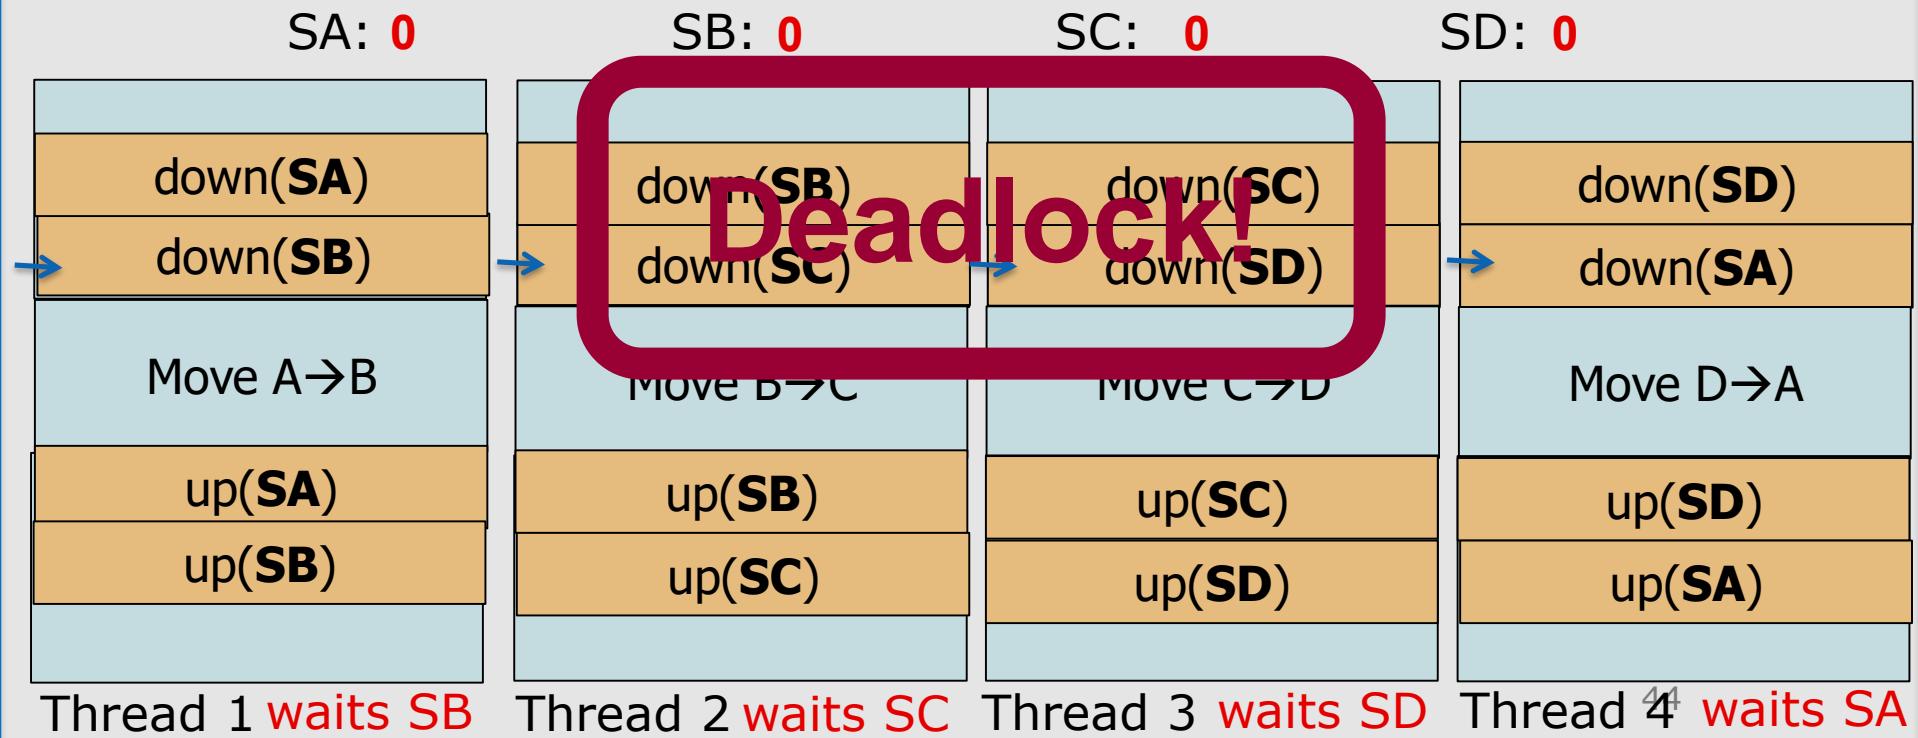
# Semaphores are not Silver Bullets



# Semaphores are not Silver Bullets



# Semaphores are not Silver Bullets



# Semaphores are not Silver Bullets

Although it is much easier to work with semaphores than read/write operations, it is still possible to cause problems (e.g., deadlocks, starvation, incorrectness)

