





Operating Systems Memory Management

David Hay
Dror Feitelson


Memory So Far




The Memory Hierarchy




The Memory Hierarchy




The Memory Hierarchy




The Memory Hierarchy




The Memory Hierarchy




The Memory Hierarchy



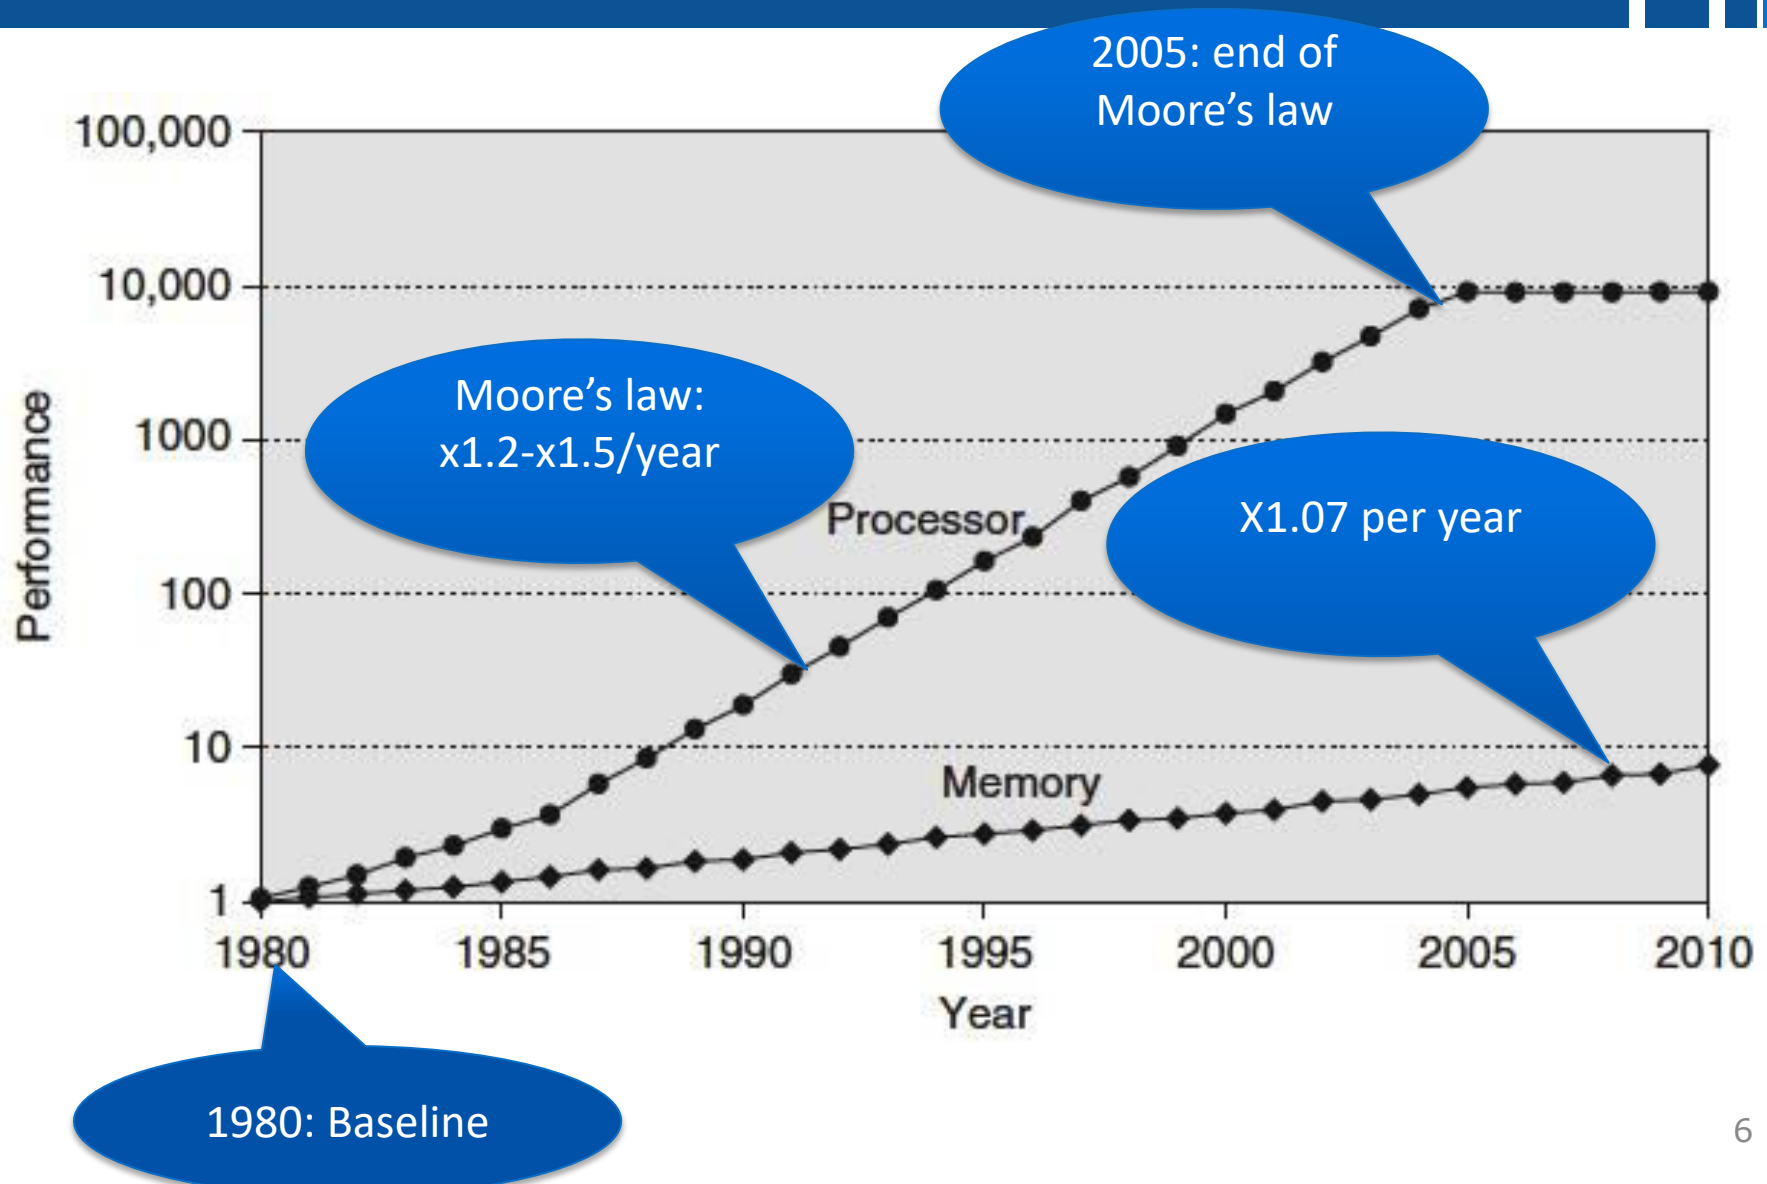
The Memory Hierarchy



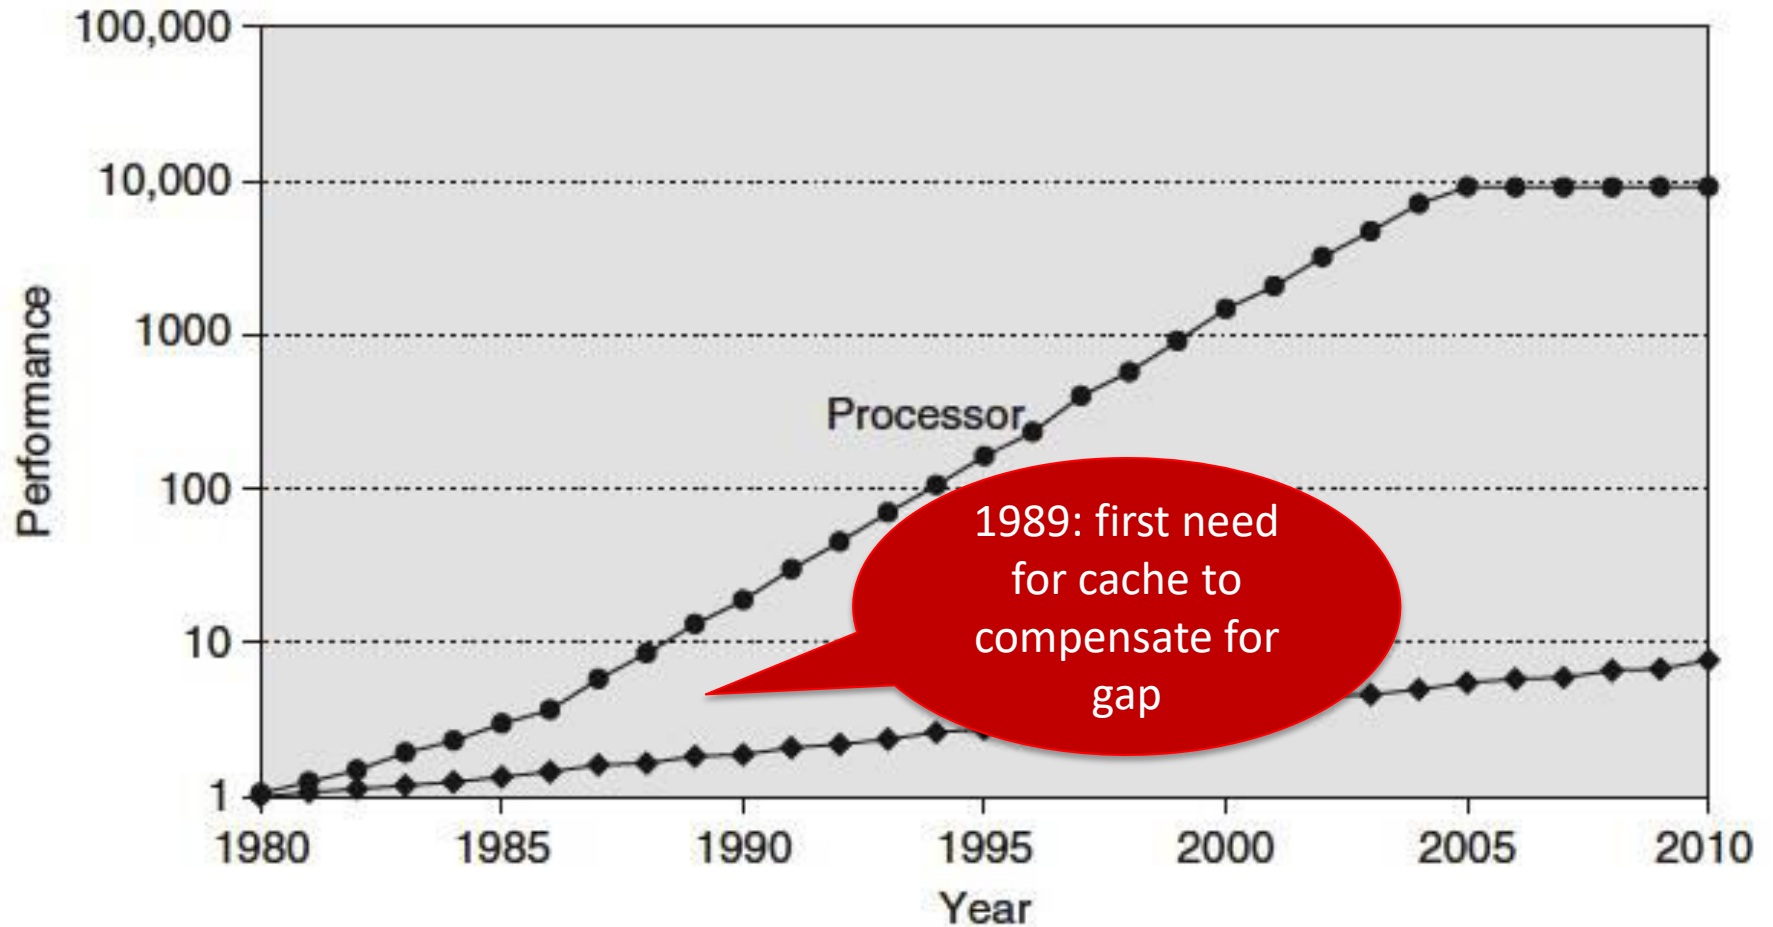
CPU-DRAM Gap 1980-2010



CPU-DRAM Gap 1980-2010



CPU-DRAM Gap 1980-2010



Caching

Small fast memory serves as cache for large slow memory

- How to find things in the cache?
- Which items should be stored in the cache?
- Which items should be evicted from the cache?

Similar considerations and solutions in HW caching and in OS memory management

Principle of Locality

Temporal locality:

If we accessed a certain address, the chances are high to access it again shortly.

- Data: updating
- Instructions: loops

Principle of Locality

Temporal locality:

If we accessed a certain address, the chances are high to access it again shortly.

- Data: updating
- Instructions: loops

Spatial locality:

If we accessed a certain address, the chances are high to access its neighbors.

- Data: arrays
- Instructions: sequential execution

Principle of Locality

Temporal locality:

If we accessed a certain address, the chances are high to access it again shortly.

- Data: arrays

- Instructions


Idea: Keep data and instructions the CPU is **most likely to need next** in fast memory close to the CPU

Spatial


If we accessed a certain address, the chances are high to access its neighbors.

- Data: arrays
- Instructions: sequential execution


Caching in the Memory Hierarchy




Caching in the Memory Hierarchy




Caching in the Memory Hierarchy




Caching in the Memory Hierarchy




Caching in the Memory Hierarchy




Caching in the Memory Hierarchy




Caching in the Memory Hierarchy



Caching in the Memory Hierarchy



Technology




Random-Access Memory (RAM)

- Static RAM (**SRAM**)
 - Retains value indefinitely, as long as it is kept powered
 - Relatively insensitive to disturbances such as electrical noise
 - Faster and more expensive than DRAM
 - Used for registers and caches
- Dynamic RAM (**DRAM**)
 - Value must be refreshed every 10-100 ms
 - Sensitive to disturbances
 - Slower and cheaper than SRAM
 - Used for main memory

Memory	Single-chip Capacity	\$/chip	\$/MByte	Access speed (ns)	Watts/chip	Watts/MByte
DRAM	128MB	\$10-\$20	\$0.08-\$0.16	40-80	1-2	0.008-0.016
SRAM	9MB	\$50-\$70	\$5.5-\$7.8	3-5	1.5-3	0.17-0.33

Disk Operation (Single-Platter View)


The disk spins
at a fixed
rotational rate



- Disk access time is **much** slower than DRAM
- **Fast** disks have access time of 5-10 **milliseconds**
 - Seek time and rotation delay are roughly equal

Disk Operation (Single-Platter View)


The disk spins
at a fixed
rotational rate



- Disk access time is **much** slower than DRAM
- **Fast** disks have access time of 5-10 **milliseconds**
 - Seek time and rotation delay are roughly equal

Disk Operation (Single-Platter View)

The disk spins at a fixed rotational rate




Read/write *head* is attached to end of the *arm* and flies over disk surface on thin cushion of air

By moving radially, arm can position read/write head over any track

- Disk access time is **much** slower than DRAM
- **Fast** disks have access time of 5-10 **milliseconds**
 - Seek time and rotation delay are roughly equal

Disk Operation (Single-Platter View)

The disk spins at a fixed rotational rate




Read/write *head* is attached to end of the *arm* and flies over disk surface on thin cushion of air

By moving radially, arm can position read/write head over any track

- Disk access time is **much** slower than DRAM
- **Fast** disks have access time of 5-10 **milliseconds**
 - Seek time and rotation delay are roughly equal

Disk Operation (Single-Platter View)

The disk spins at a fixed rotational rate



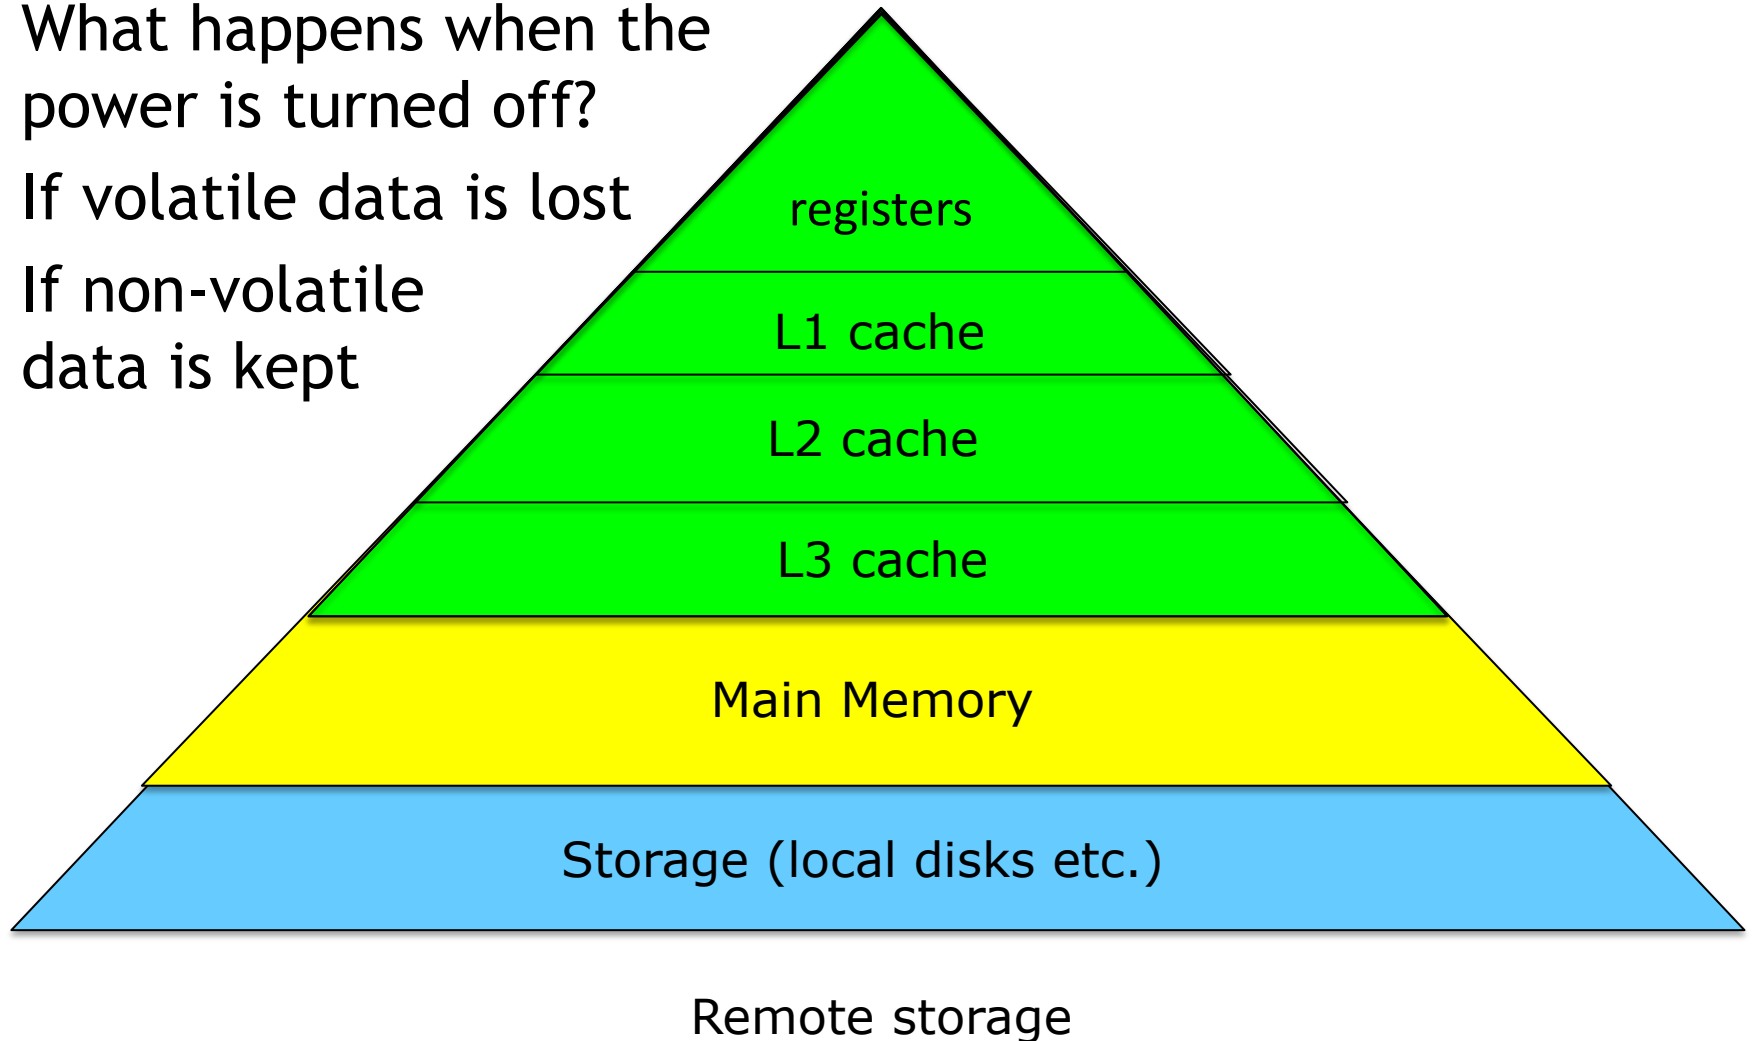
Read/write *head* is attached to end of the *arm* and flies over disk surface on thin cushion of air

By moving radially, arm can position read/write head over any track

- Disk access time is **much** slower than DRAM
- **Fast** disks have access time of 5-10 **milliseconds**
 - Seek time and rotation delay are roughly equal

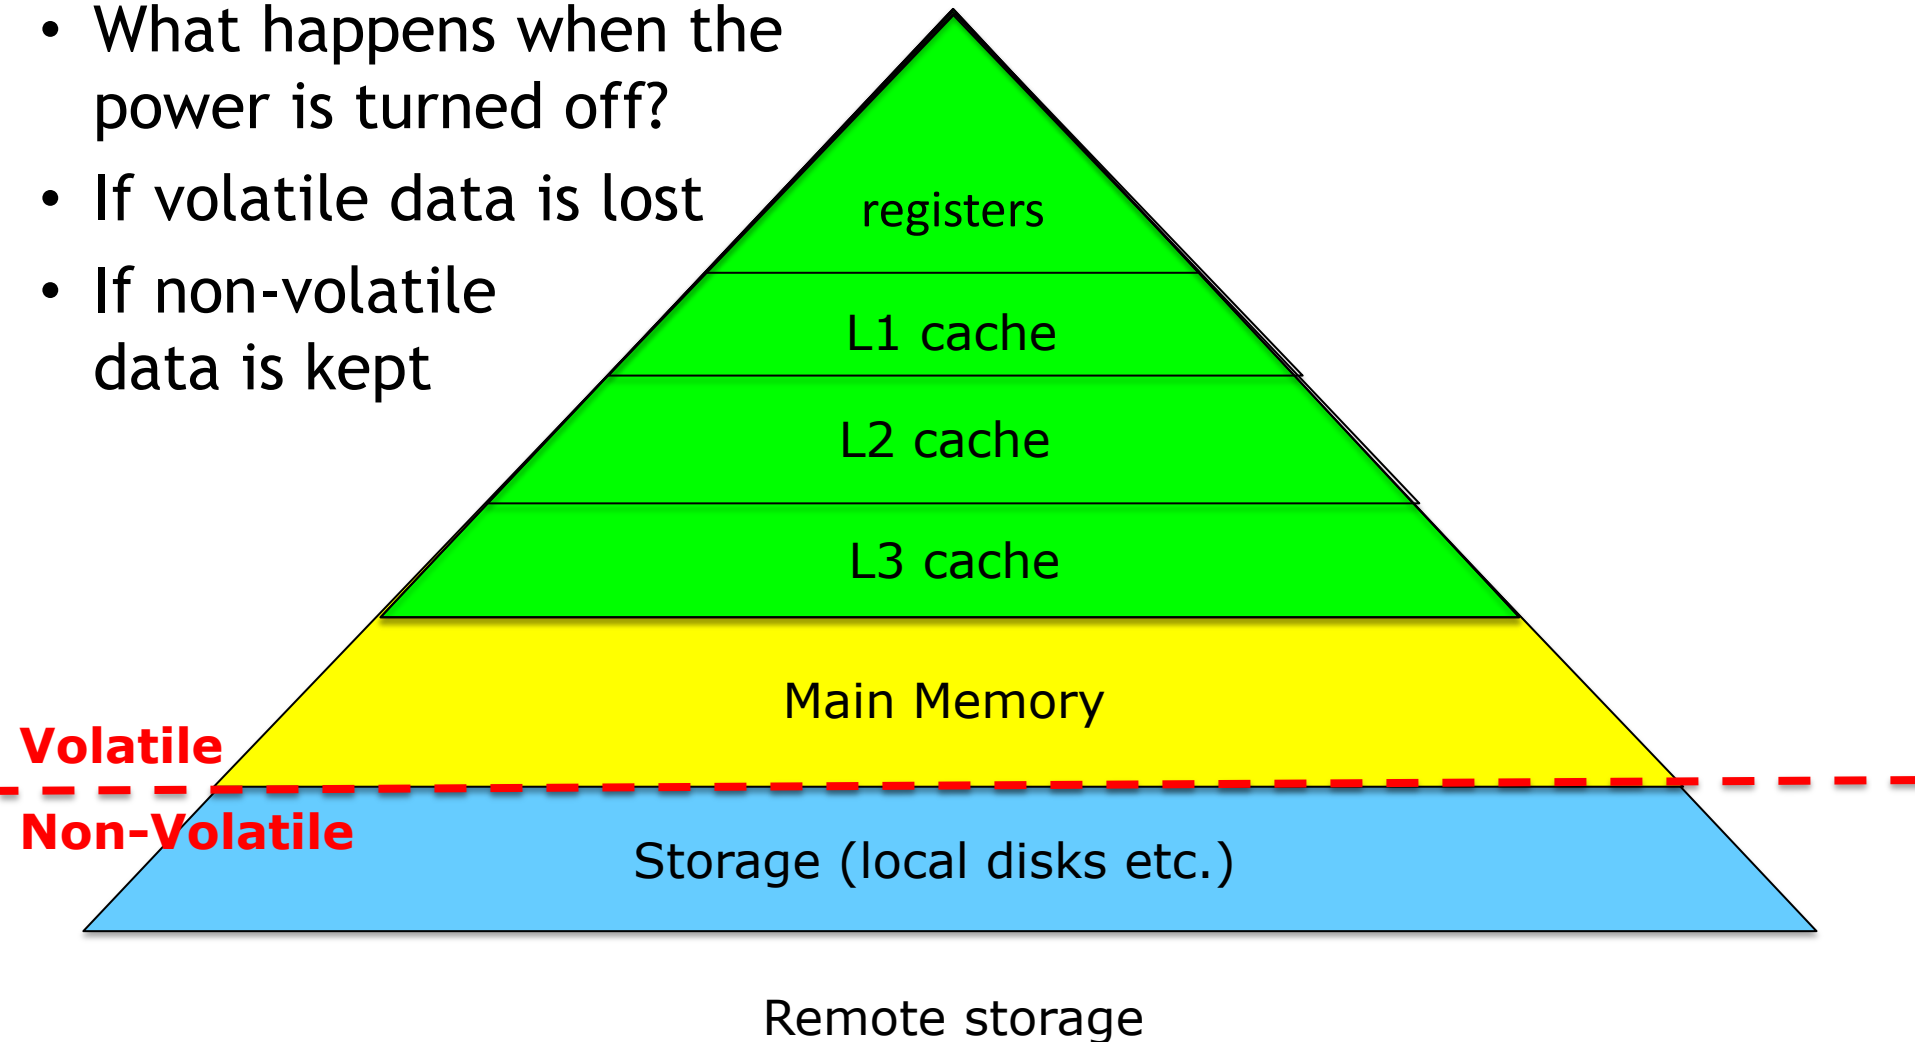
Volatility (volatile = נדיף)

- What happens when the power is turned off?
- If volatile data is lost
- If non-volatile data is kept




Volatility (volatile = נדיף)

- What happens when the power is turned off?
- If volatile data is lost
- If non-volatile data is kept




Modern Architectures

What happens with
multi-core architectures?
Multiple processors?




Modern Architectures

What happens with
multi-core architectures?
Multiple processors?




Modern Architectures

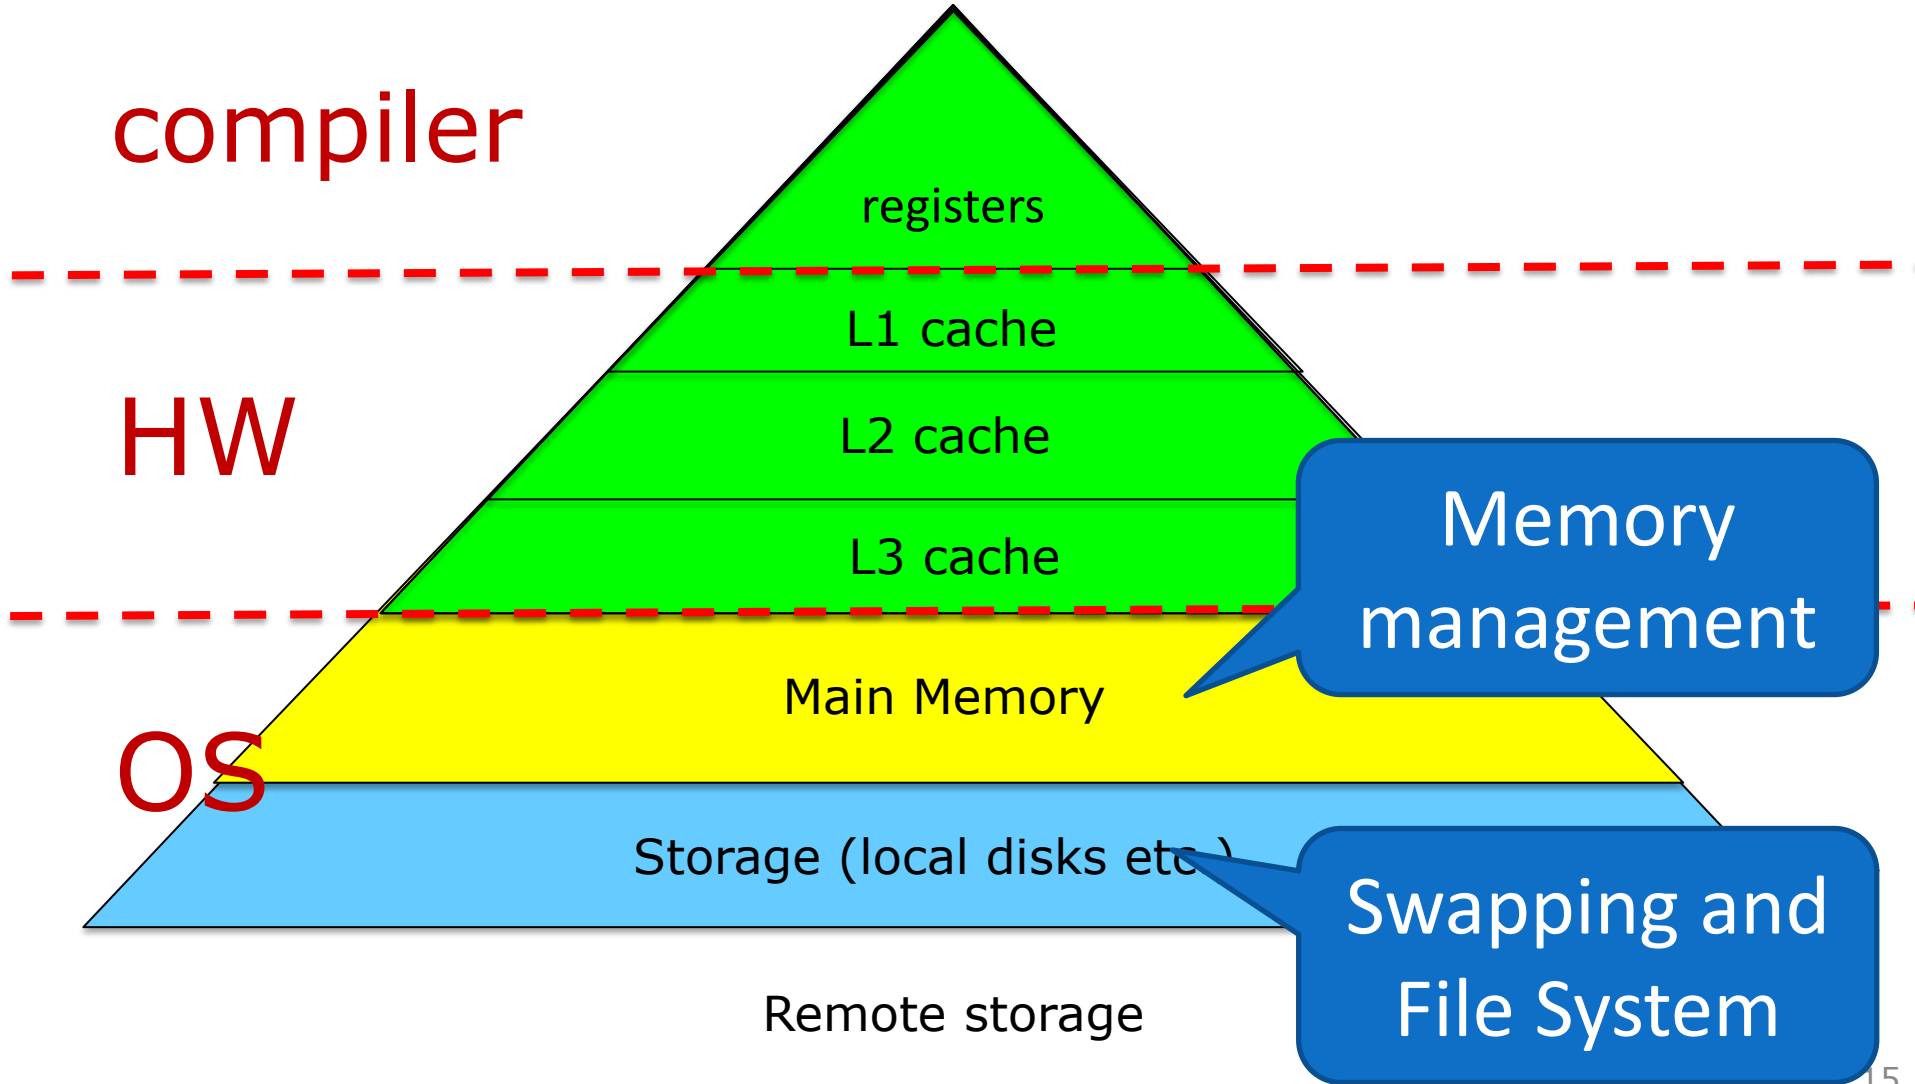
What happens with
multi-core architectures?
Multiple processors?



Responsibility



Responsibility



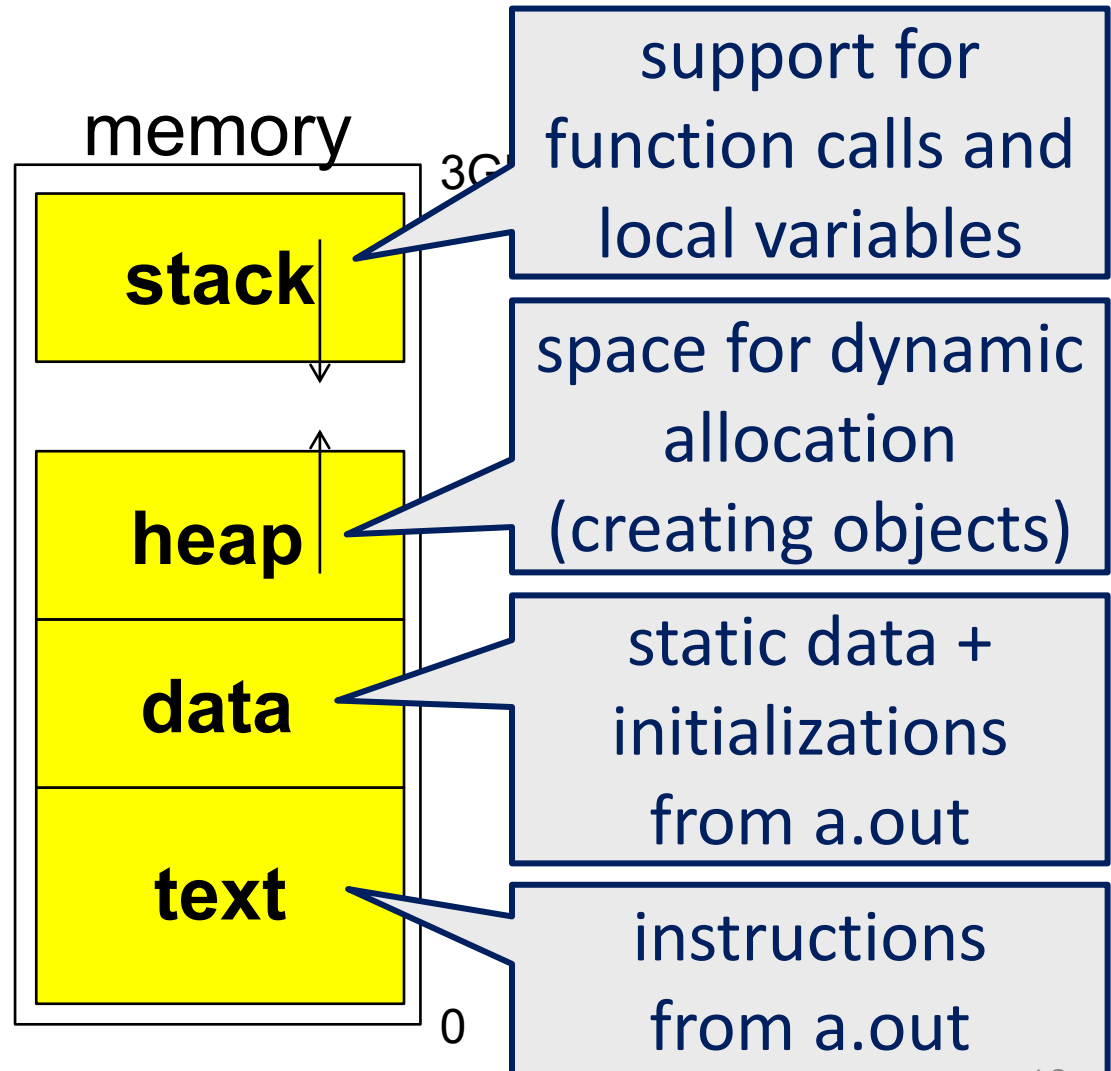
THE ADDRESS SPACE

“Address Space”


- All memory addressable by the program
- Depends on the architecture
- 32-bit architecture → 4GB memory
- Some set aside for OS
 - Windows left 2GB for user process
 - Linux left 3GB for user process
- **This is for each and every process**

Memory Segments

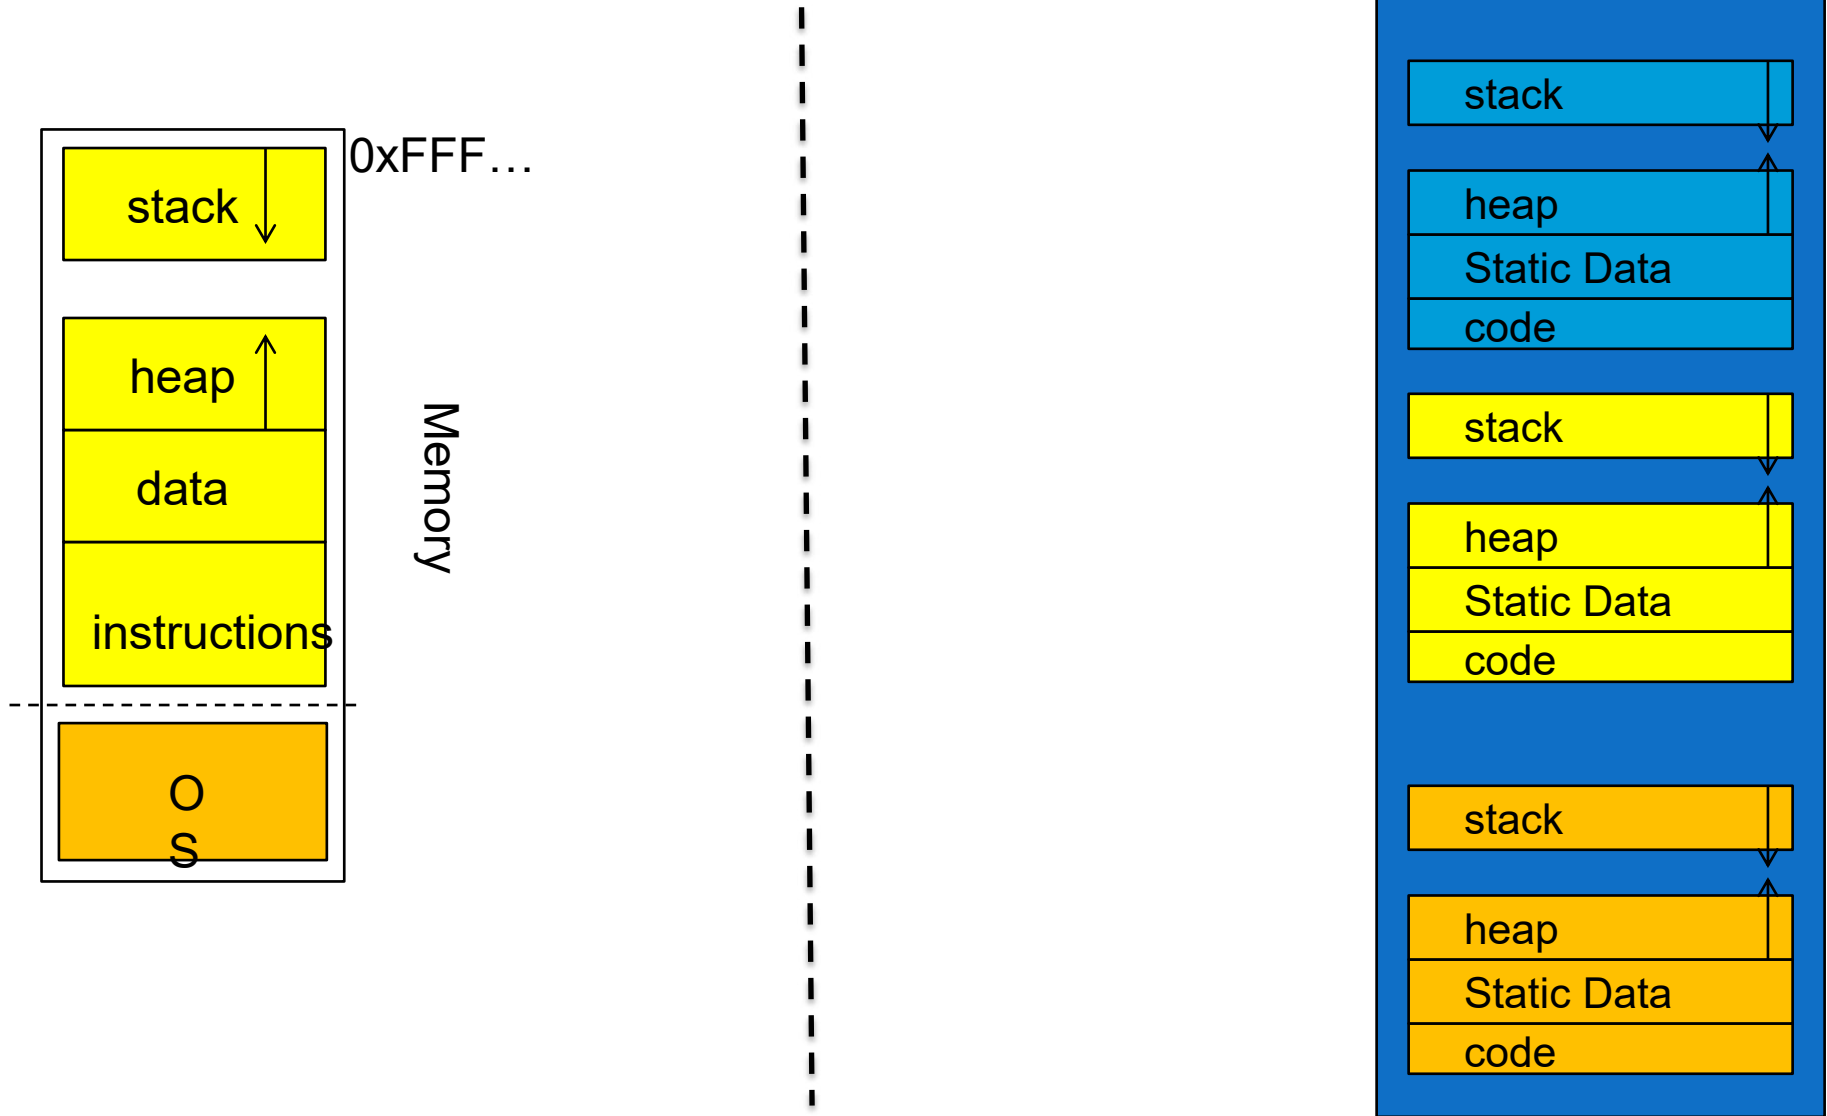
- The address space is divided into 4 segments with different uses
- Stack and heap placed to leave space for growth



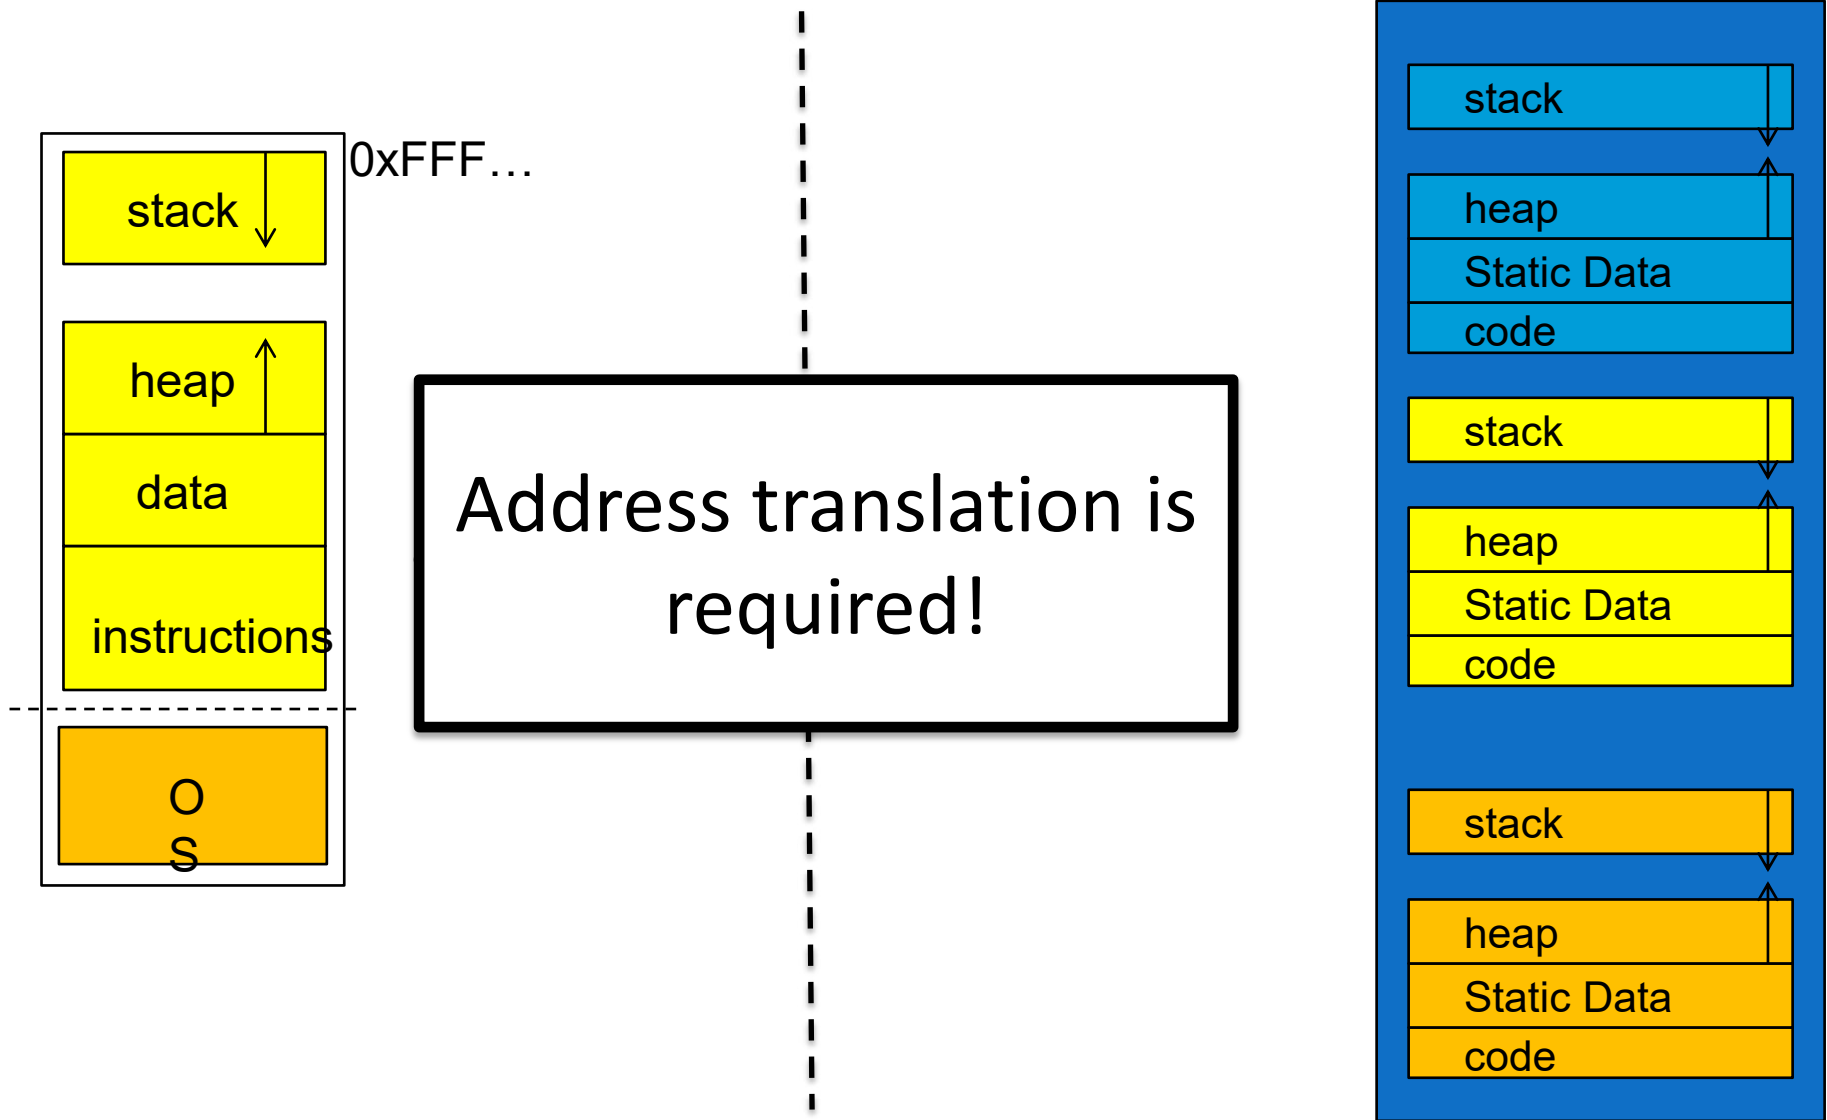
Process View Vs. Reality



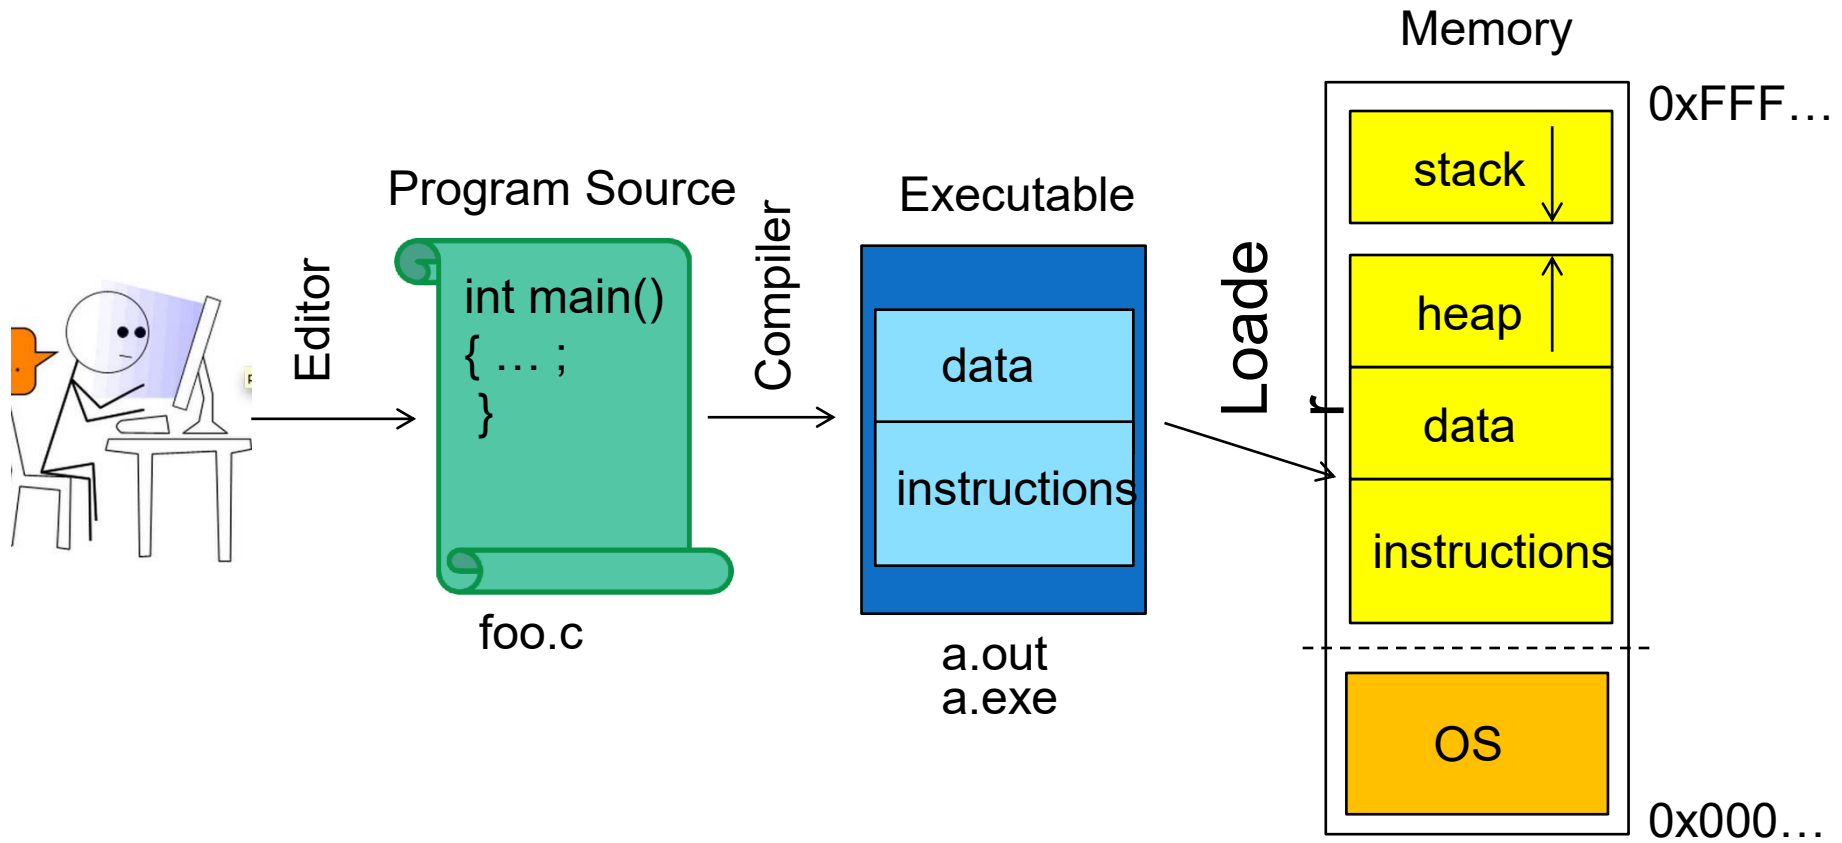
Process View Vs. Reality



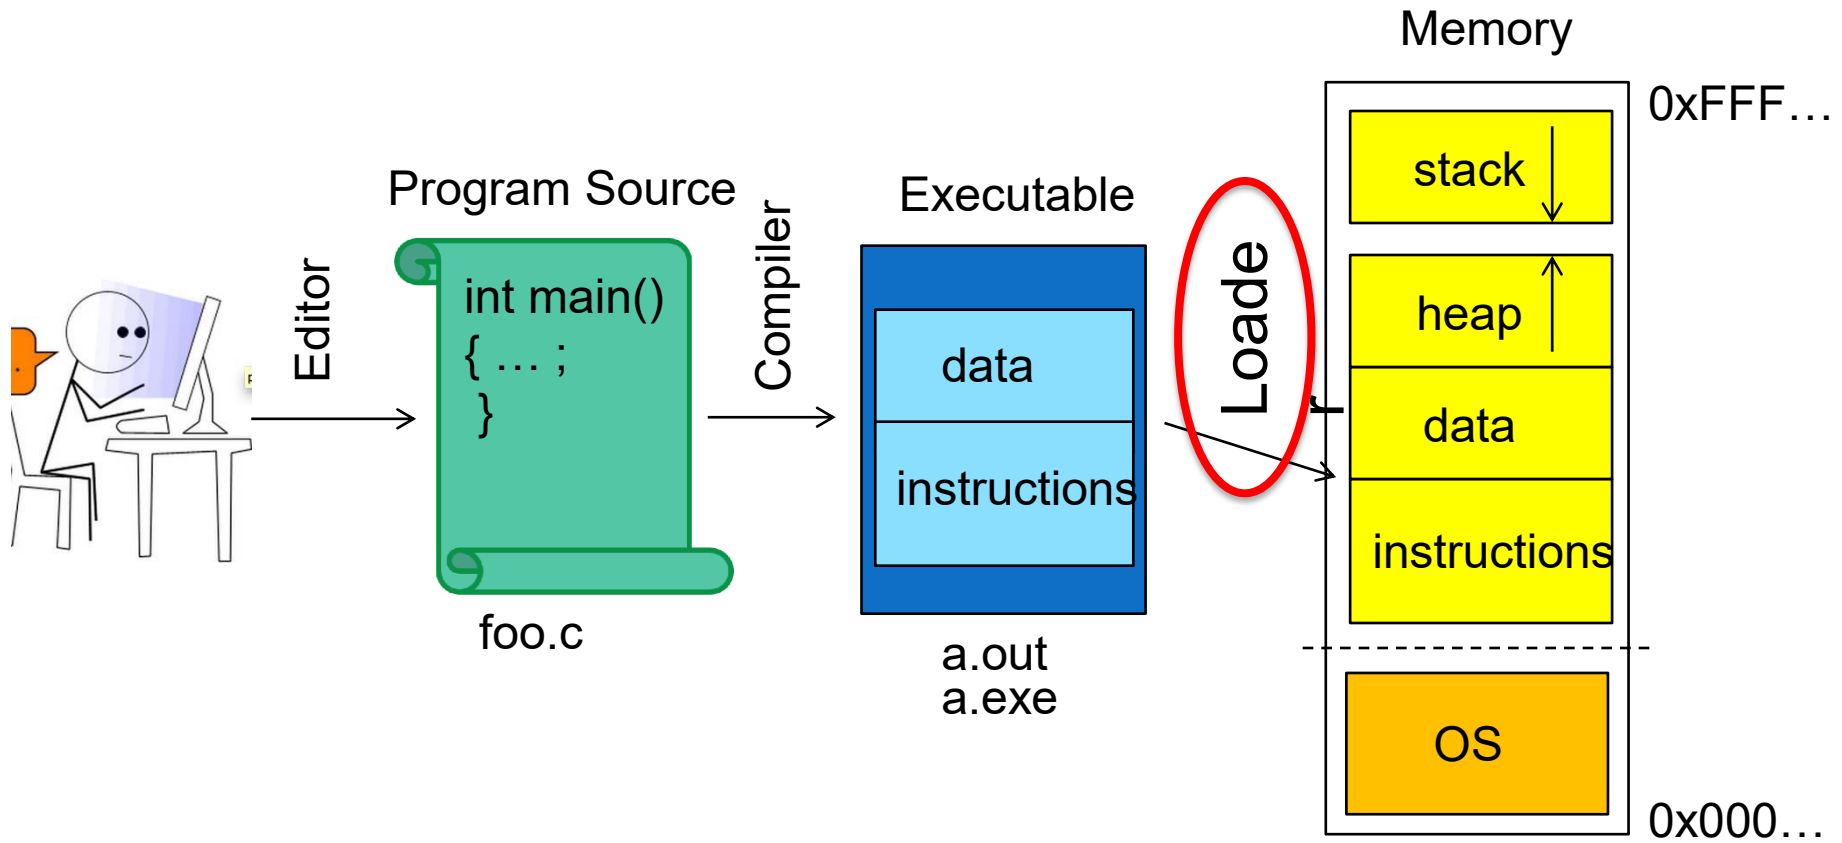
Process View Vs. Reality



When are Addresses Created?

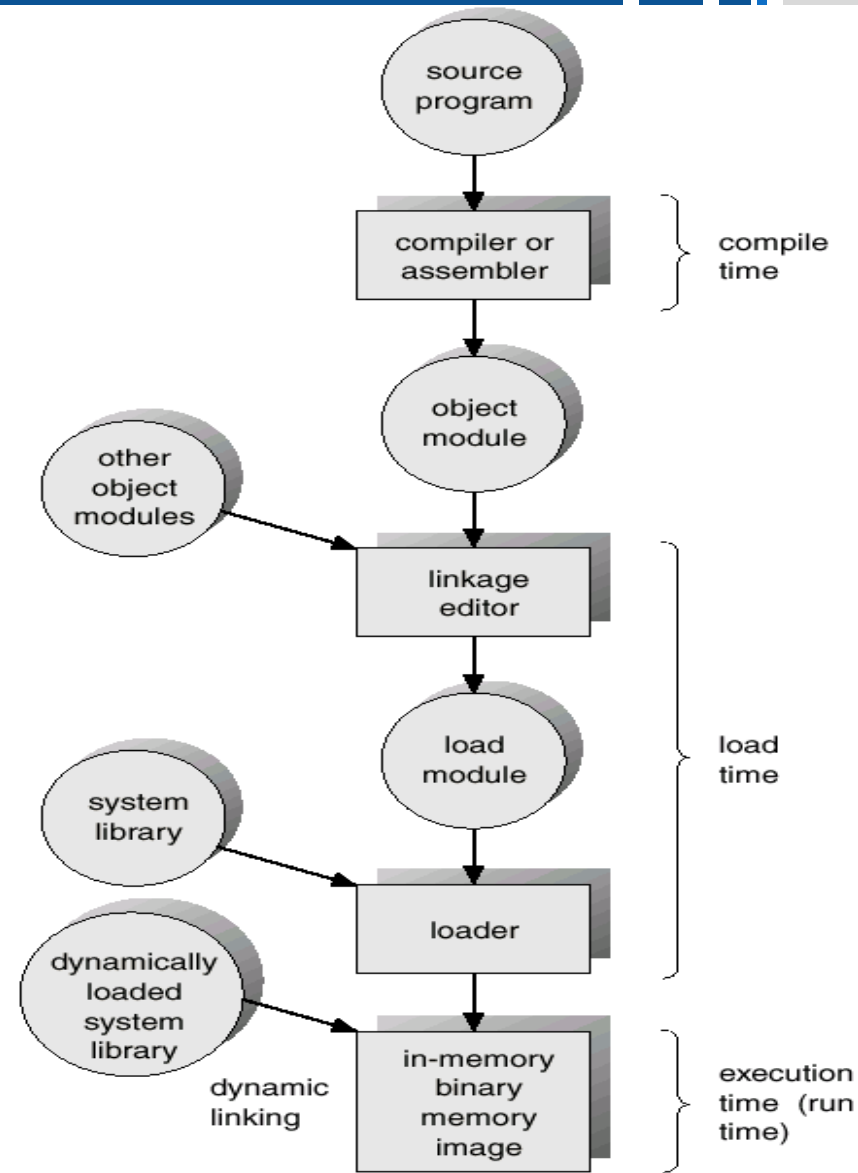


When are Addresses Created?

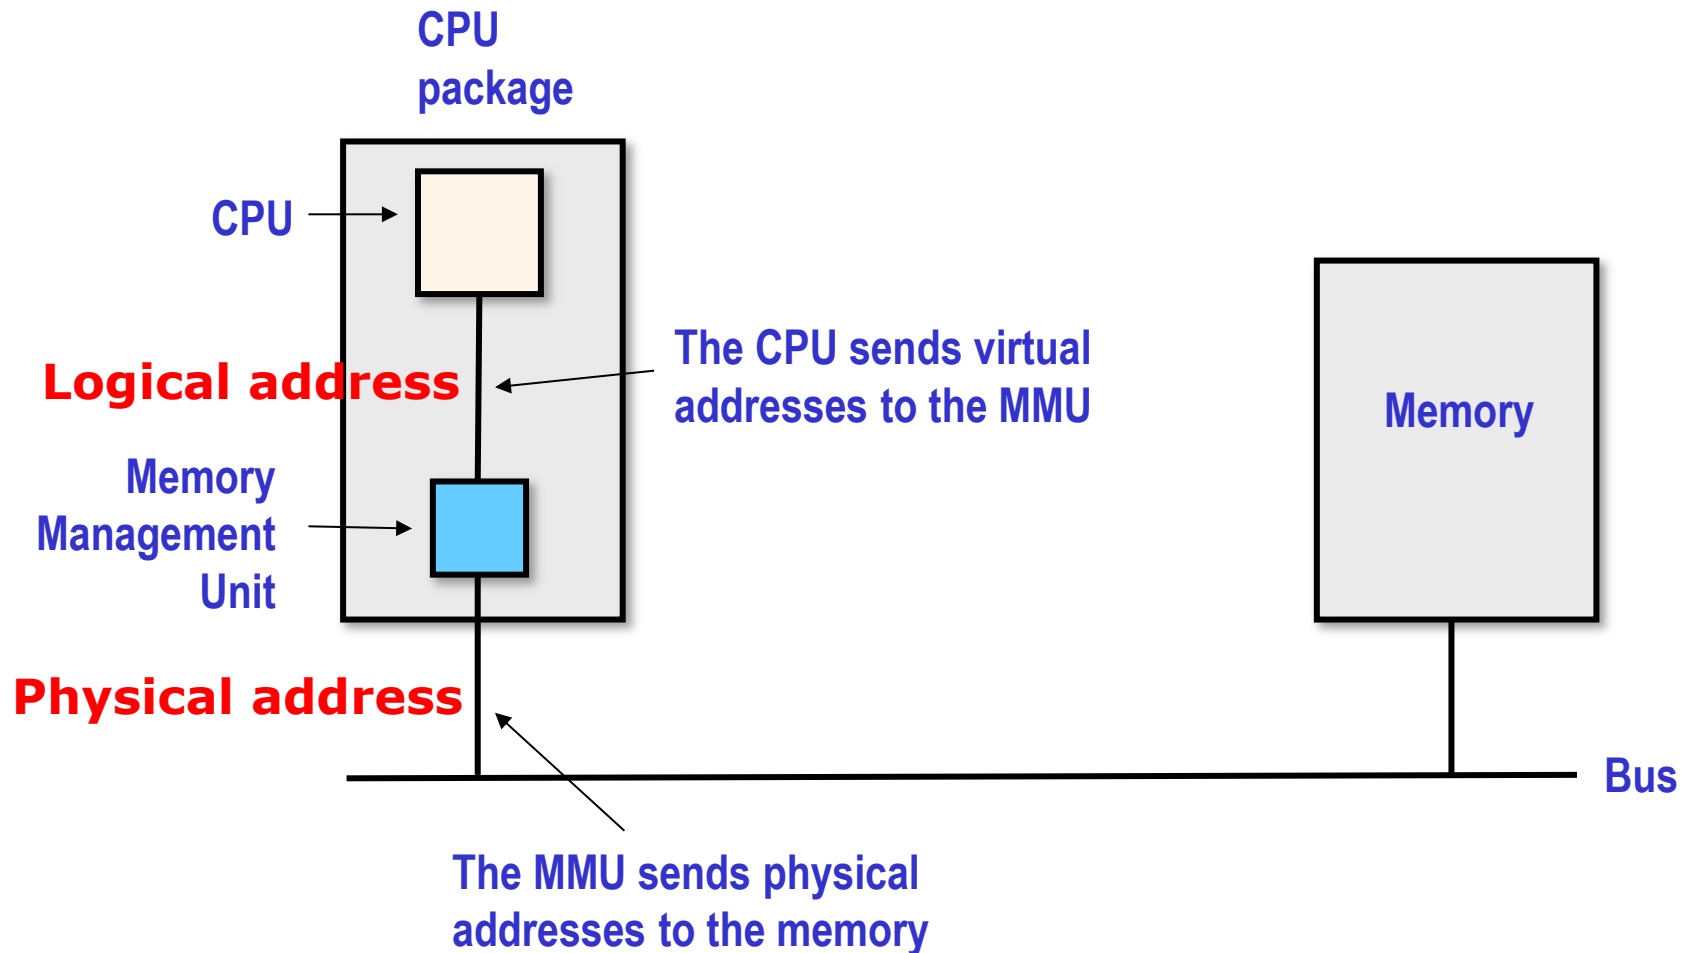


From Source to Execution...

- Programs need to be loaded to main memory before being executed
- Most addresses are known only after the load module (e.g., EXE file) is loaded into the main memory
 - Load module contains **relocatable** addresses



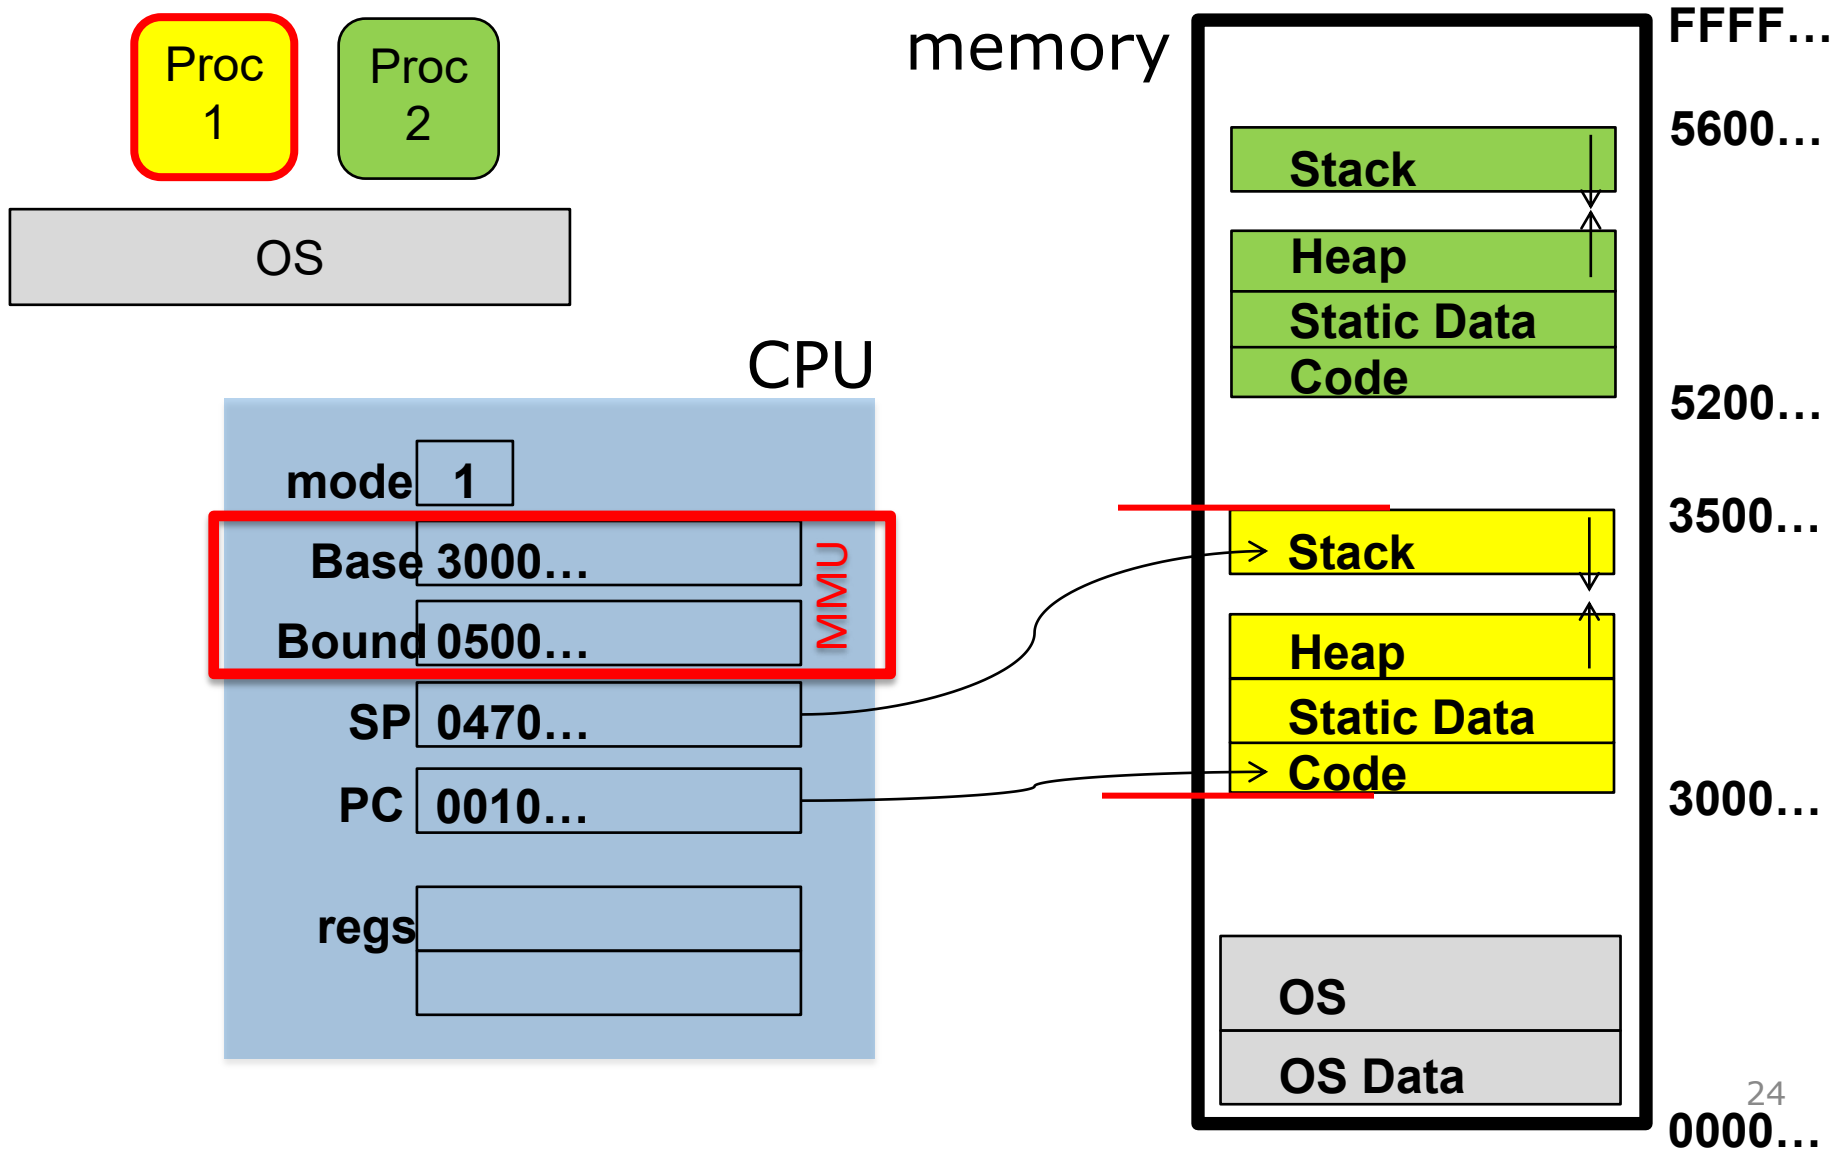
So, who is in charge of the translation?



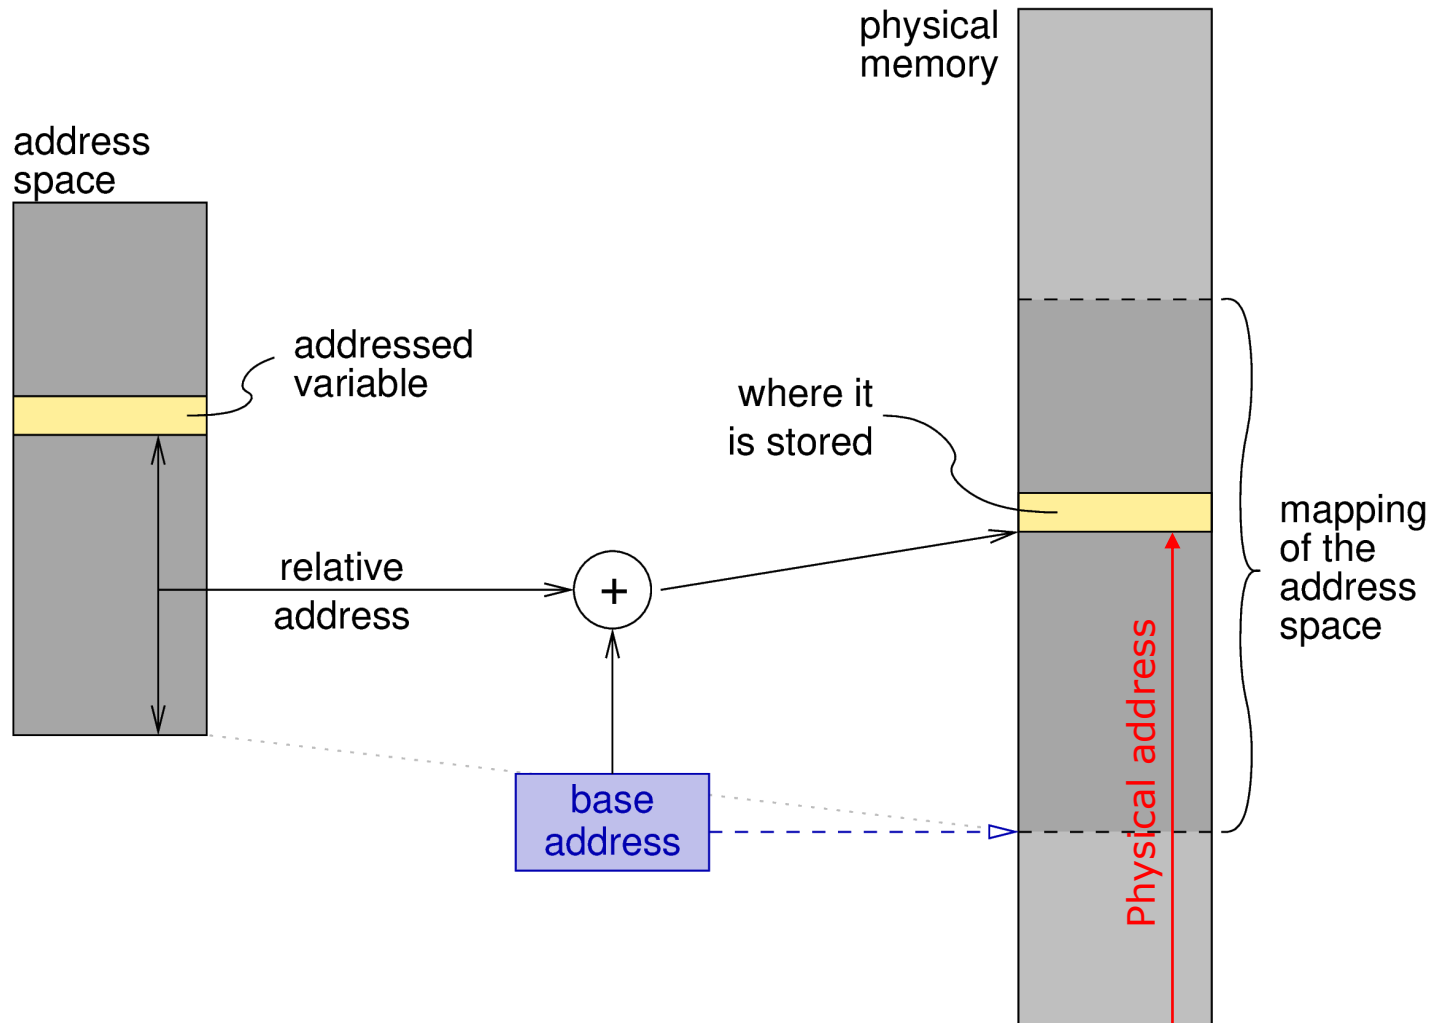
Contiguous Allocation

- Each process address space appears as a block in physical memory
- Address translation:
 - Each process gets a *base* address
 - **Logical address** x in the process is mapped to **physical address** $\text{base} + x$ in the memory
- Very simple MMU implementation
 - Need to allocate “space to grow” → internal fragmentation

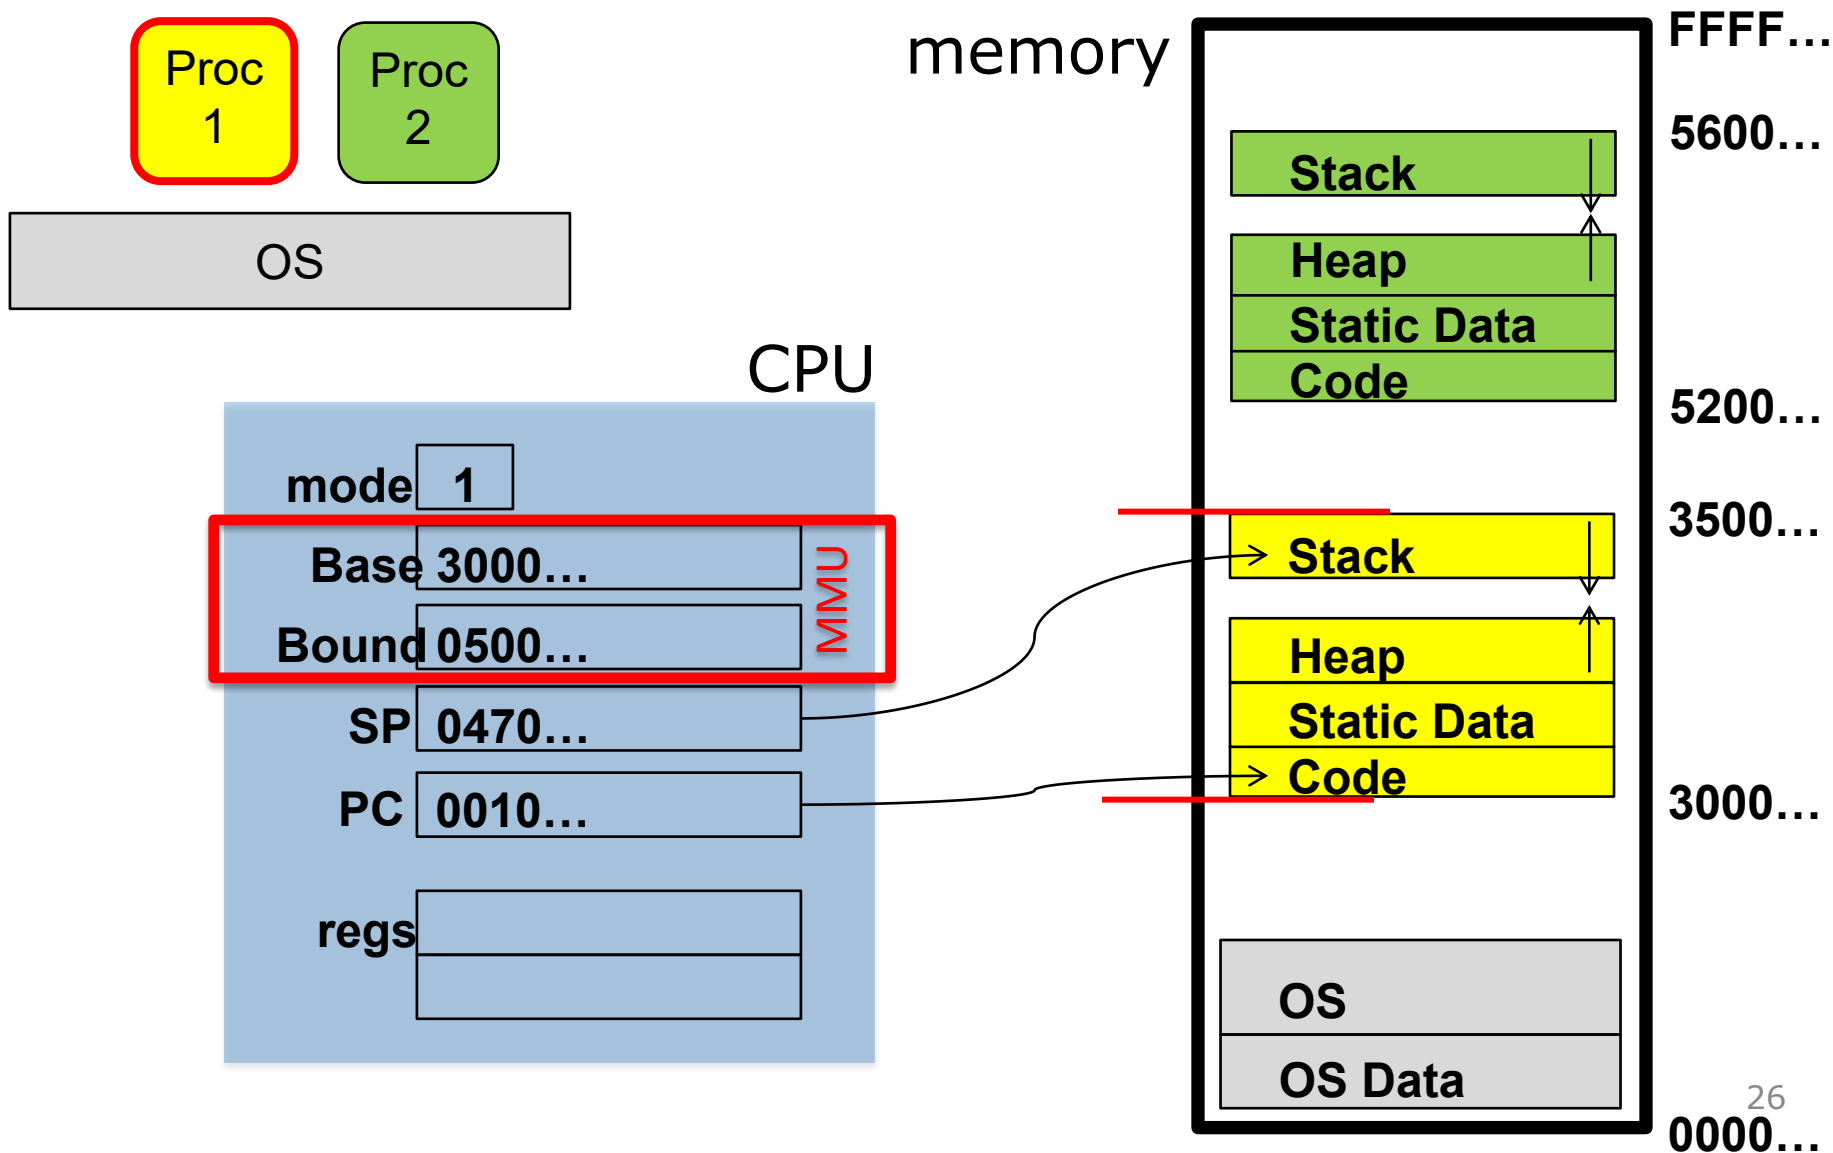
Simple Mapping Example



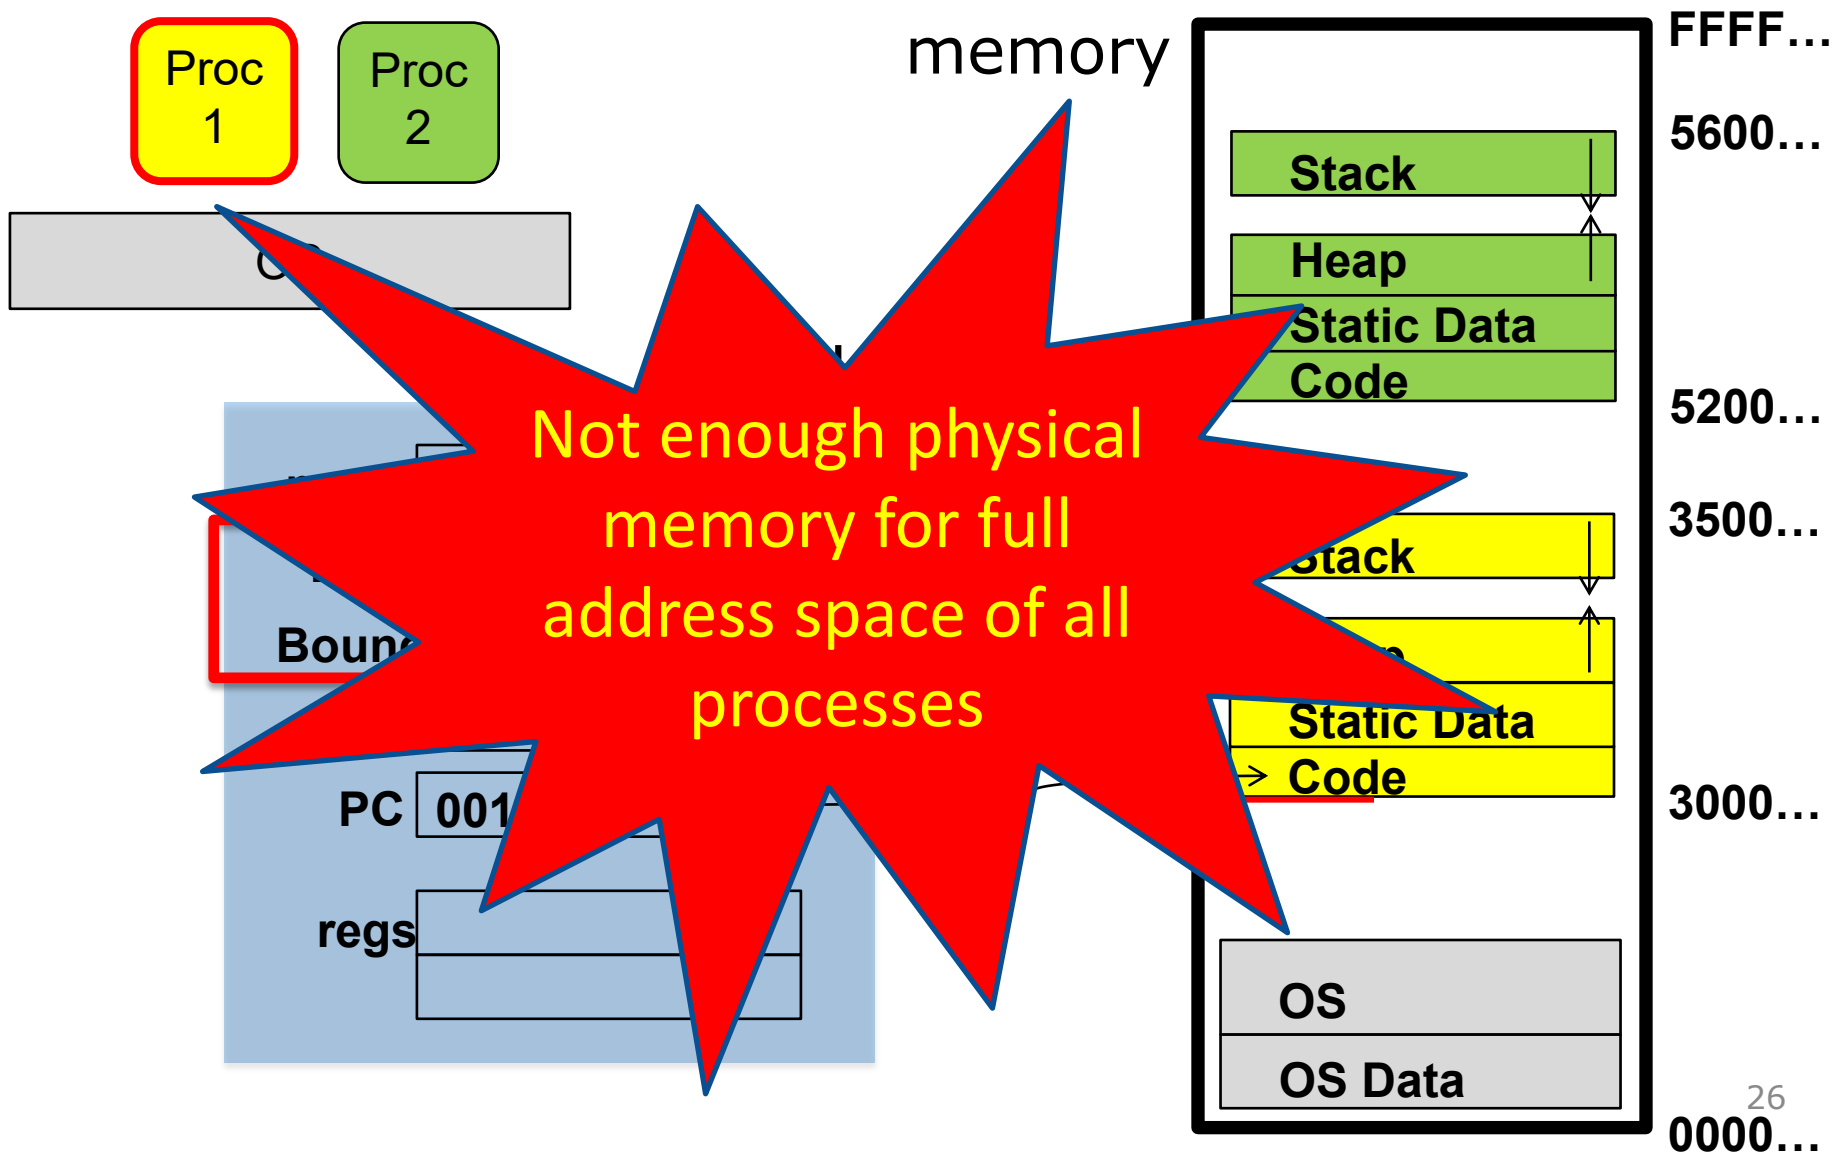
Address Translation Implementation



Problem




Problem




Fragmentation

Allocation Dynamics


Allocation Dynamics



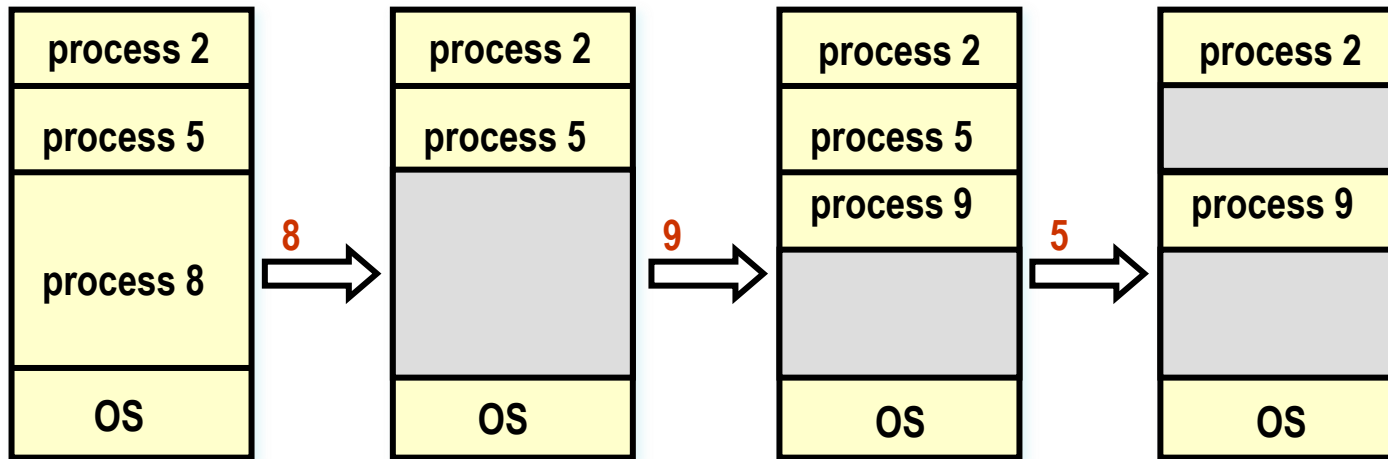
Allocation Dynamics




Allocation Dynamics




Allocation Dynamics



Allocation Dynamics



Allocation Dynamics



Result: External Fragmentation

Fragmentation

- Situation in which we have enough memory to allocate to a process, but not contiguously, as the holes are scattered all over memory
 - Internal Fragmentation: free memory inside process' allocation
 - External Fragmentation: free memory between processes' allocations
- We try to solve only external fragmentation!

Allocation Algorithms

- Input:
 - List of current free ranges
 - Request for new allocation
- Output:
 - Decision which free range to use
 - Can use only part

Allocation Algorithms

- **First fit:**

Scan the list and select the first range that fits

- Linear time, constant < 1
- May cause excessive fragmentation at low addresses

Allocation Algorithms

- **Next fit:**

Scan the list **starting from where you left off last time** and select the first range that fits

- Linear time, constant < 1
- Spreads fragmentation more evenly

Allocation Algorithms

- **Best fit:**
Scan the and select the range that provides the snuggest fit
- Linear time
- Maintains large ranges
- Causes small fragments


Fragmentation Solution: Compaction

Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes


Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes




Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes




Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes




Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes




Fragmentation Solution: Compaction

- Move around data to unify holes and create large enough holes to accommodate future processes
- **Problem:** Very costly operation. Sometime not enough space to move data



Better Solution: Using Pages


- Divide process memory space into fixed size **pages**
- Map each page independently



Better Solution: Using Pages


- Divide process memory space into fixed size **pages**
- Map each page independently

Segmentation




Better Solution: Using Pages

- Divide process memory space into fixed size **pages**
- Map each page independently




Better Solution: Using Pages

- Divide process memory space into fixed size **pages**
- Map each page independently



Better Solution: Using Pages

- Divide process memory space into fixed size **pages**
- Map each page independently



PAGING



The Concept of Paging

- The address space is divided into pages
- All pages have a fixed size
 - Typically 4KB (4096 bytes)
- The physical memory is (conceptually) divided into frames
 - The same size as pages
- Any page can be mapped to any frame
 - The mapping is done by the OS ...
 - ... and used by the MMU to access memory

Pages and Frames


Pages and Frames

Logical
memory



Pages and Frames

Logical
memory



Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Pages and Frames

Logical
memory


0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Pages and Frames


Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p


Physical
memory



Pages and Frames



Pages and Frames



Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Physical
memory

i
j
k
l
e
f
g
h
m
n
o
p
a
b
c

Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

Physical
memory

i
j
k
l
e
f
g
h
m
n
o
p
a
b
c

Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Physical
memory

i
j
k
l
e
f
g
h
m
n
o
p
a
b
c

Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

i
j
k
l
e
f
g
h
m
n
o
p
a
b
c

Pages and Frames

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

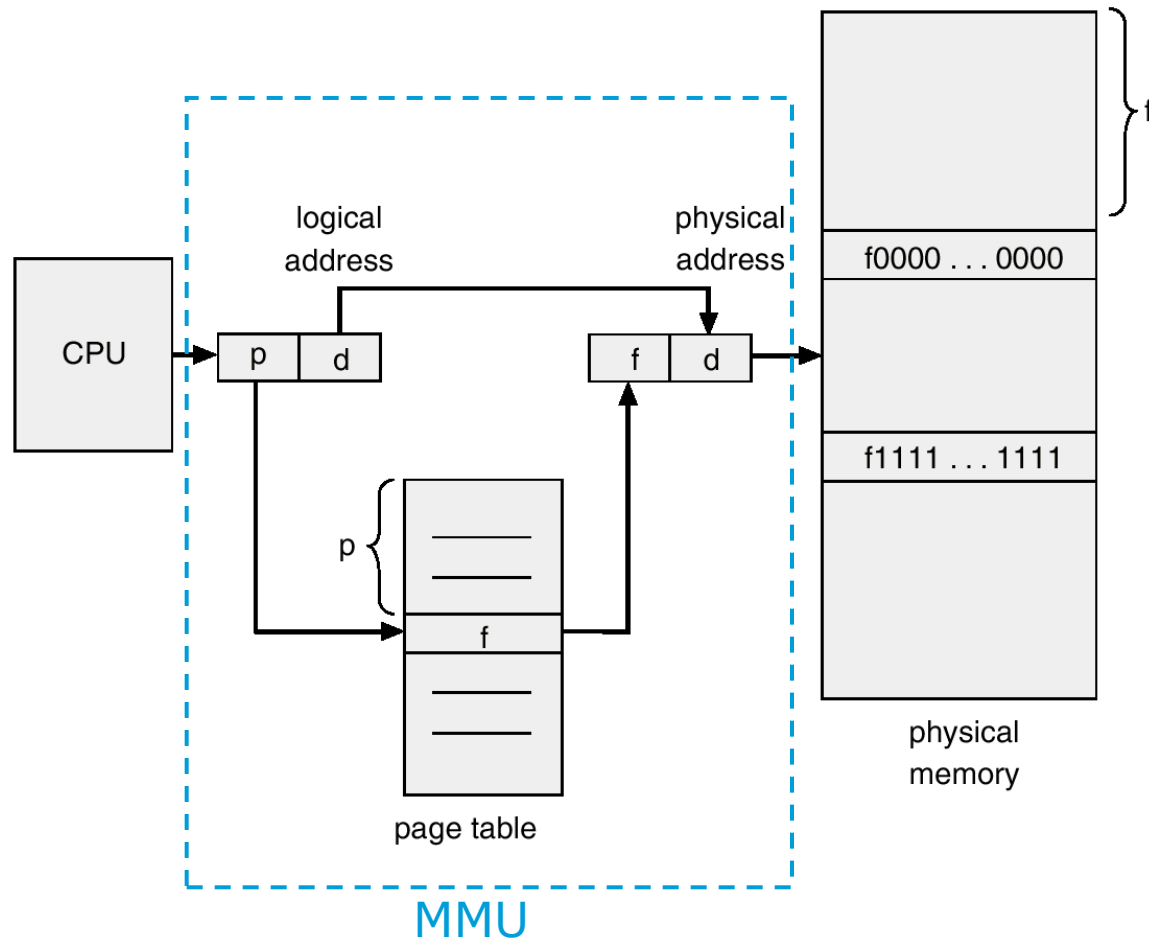
(per process)

Used by the MMU

Physical
memory

i
j
k
l
e
f
g
h
m
n
o
p
a
b
c

Address Translation Architecture



Working with the page table

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table

Address 6
00110

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table

Address 6
00110 → Page 001
Offset 10

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table

Address 6
00110 → Page 001
Offset 10

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table

Address 6
00110 → Page 001
Offset 10

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

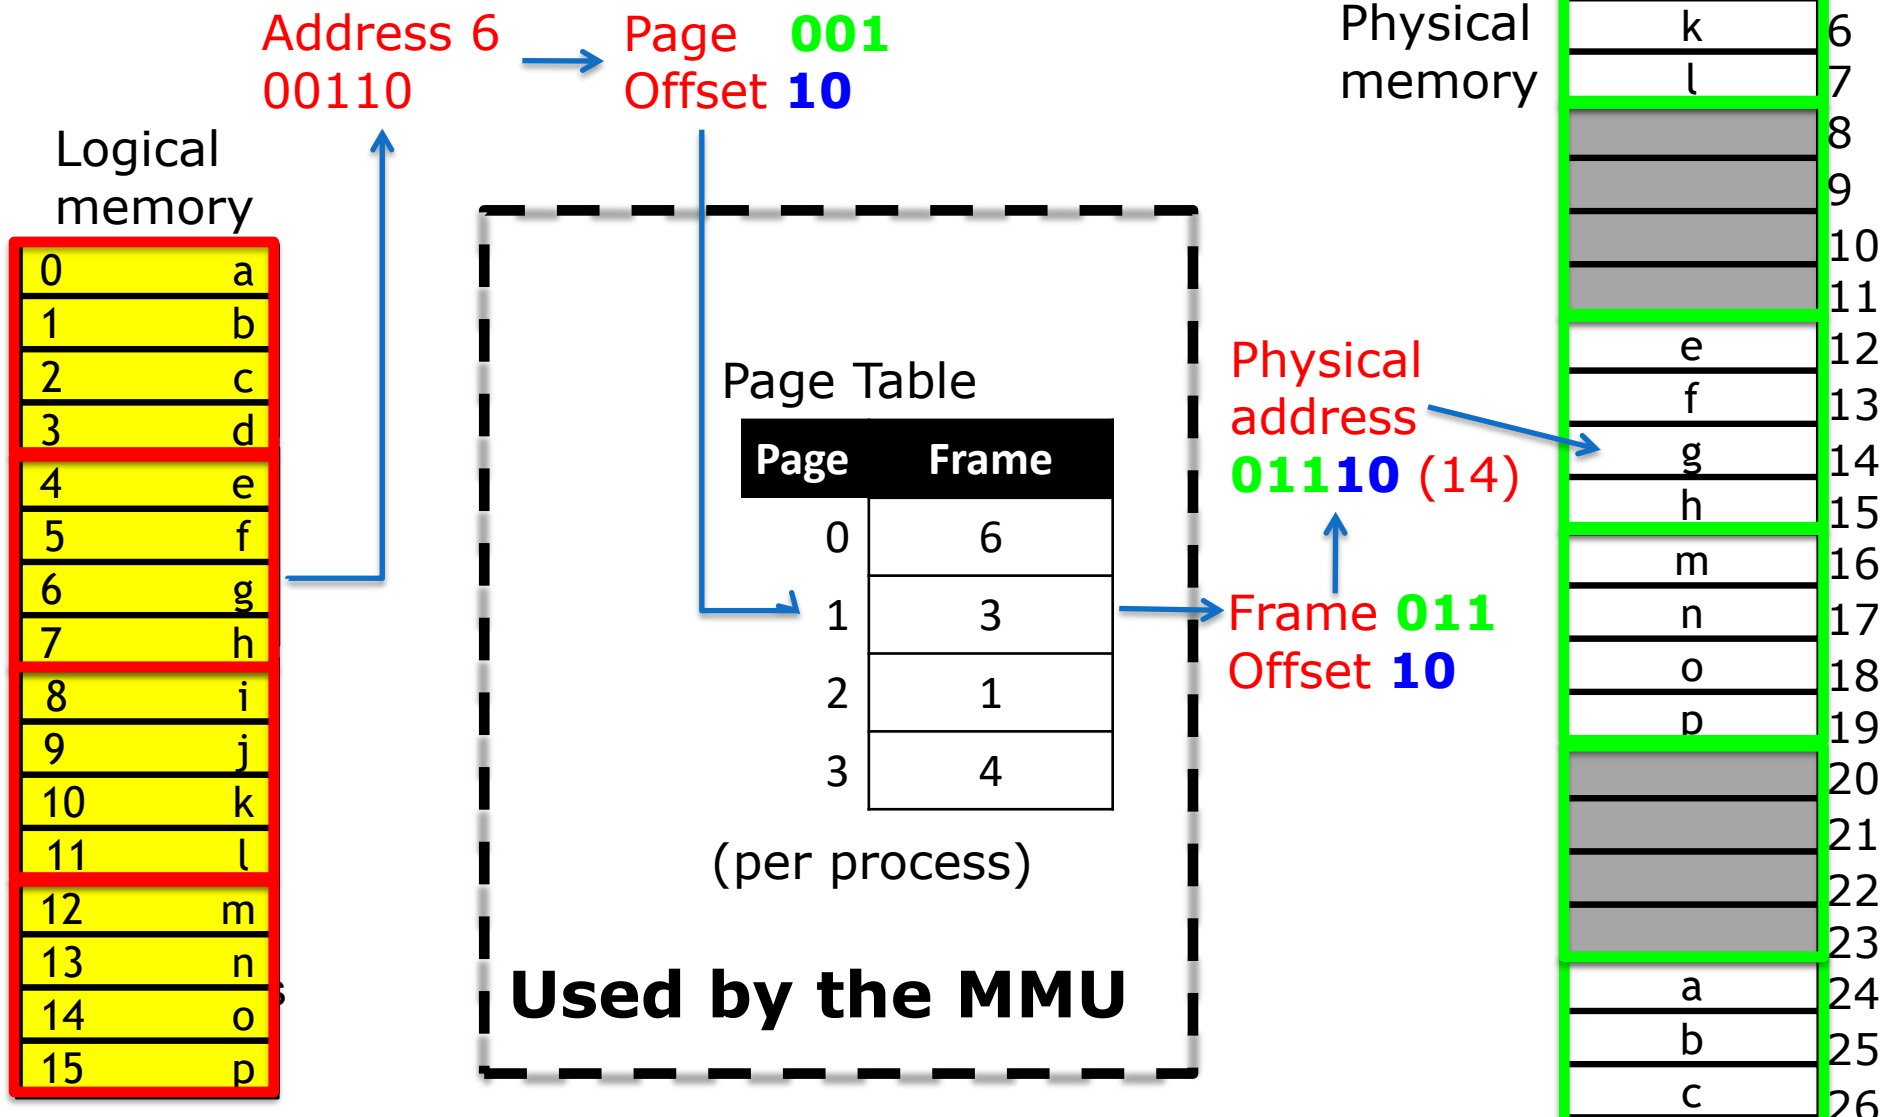
(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table



Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables

Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables
- p is the page size
s size of the process
e size of the entry in the page table
→ **Overhead = $(se/p) + p/2$**

Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables
- p is the page size
 s size of the process
 e size of the entry in the page table
→ **Overhead** = $(se/p) + p/2$

Page table size: s/p
pages → s/p entries,
each of size e

Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables
- p is the page size
 s size of the process
 e size of the entry in the page table
→ **Overhead = $(se/p) + p/2$**

Page table size: s/p
pages → s/p entries,
each of size e

Internal fragmentation

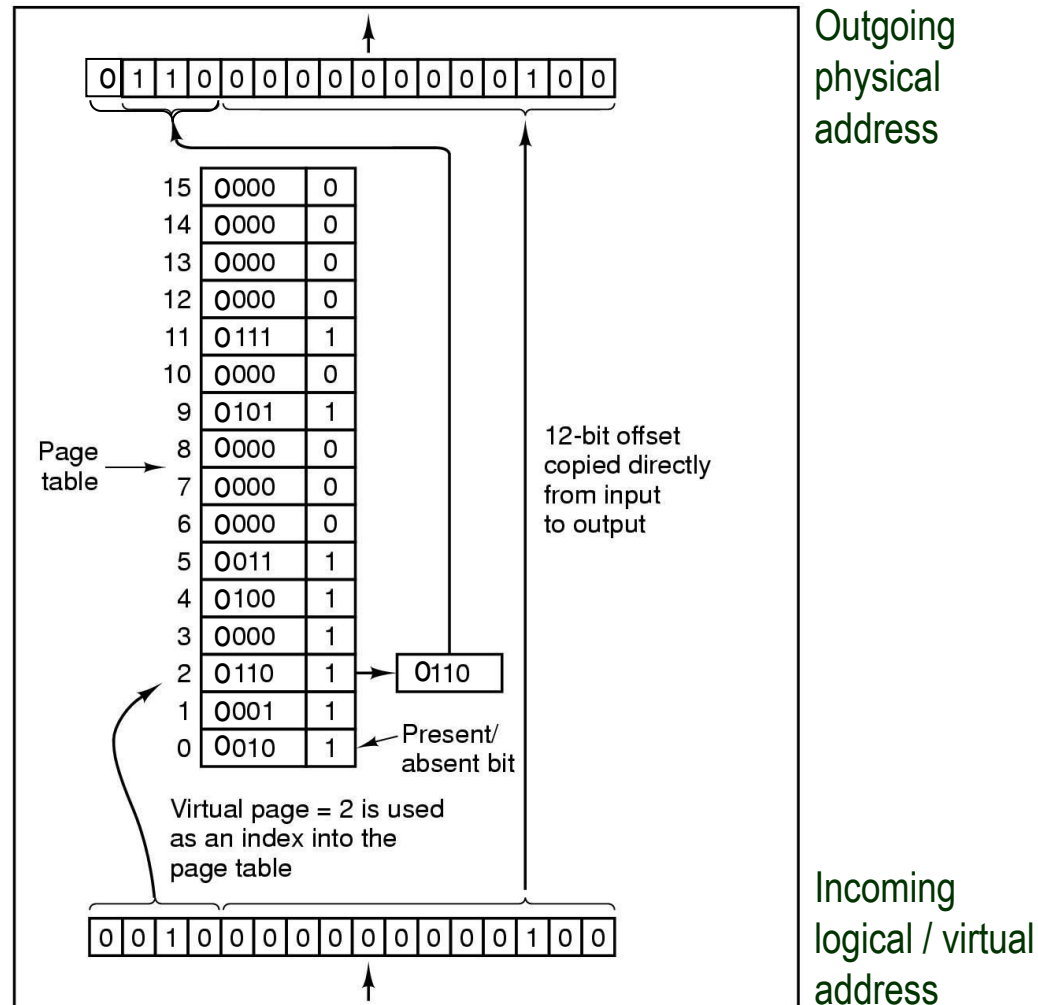
Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables
- p is the page size
 s size of the process
 e size of the entry in the page table
 - **Overhead = $(se/p) + p/2$**
 - **Optimal page size is $\sqrt{2se}$**

Optimal page size

- Small page size:
 - Less internal fragmentation
 - Larger page tables
- p is the page size
 s size of the process
 e size of the entry in the page table
→ **Overhead = $(se/p) + p/2$**
→ **Optimal page size is $\sqrt{2se}$**
- For $s=1\text{MB}$, $e=64\text{ bit}$ → $p=4\text{KB}$


Internal operation of MMU with 16 4KB pages



Address Translation

- Divide virtual address into two parts:
 - Page number (top 20 bits)
 - Offset into page (bottom 12 bits)
- Use page number as index into page table
- Get frame number from page table (20 bit)
- Combine frame number and offset to create physical address

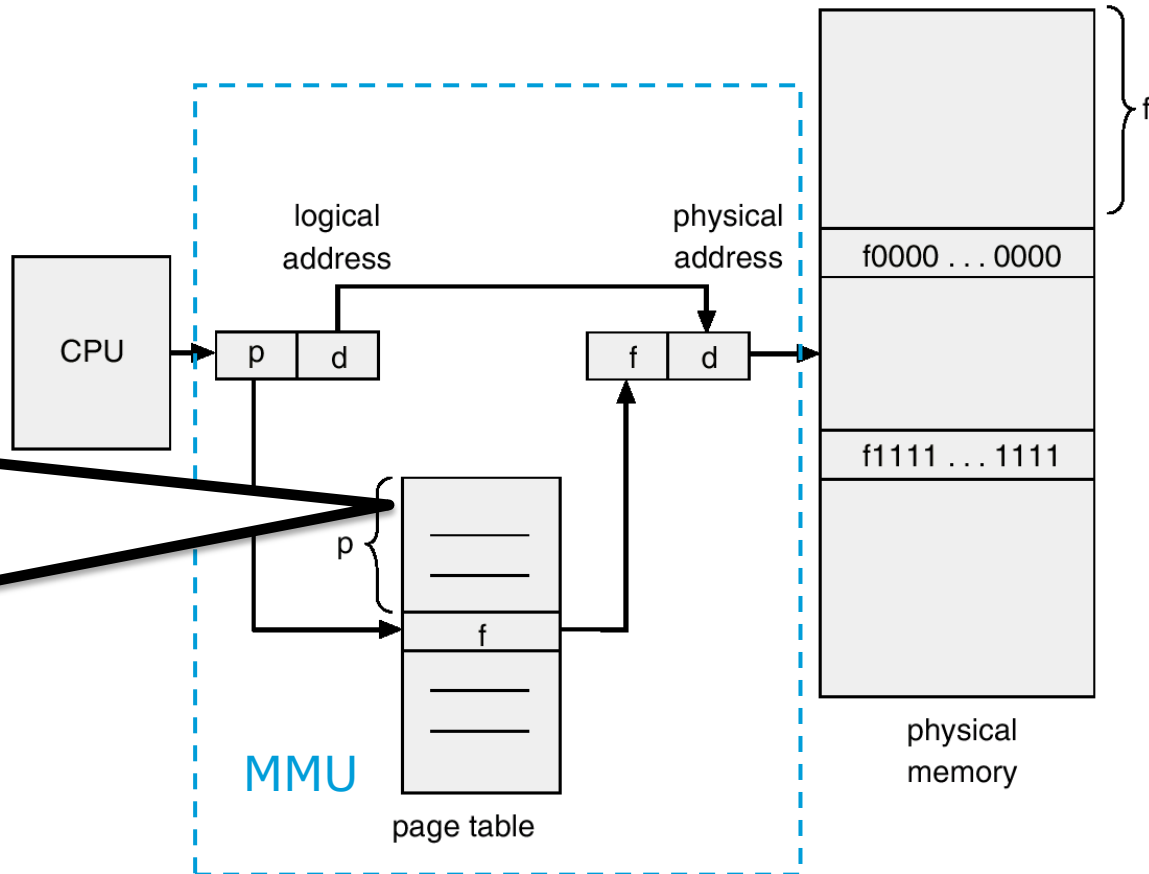
Address Translation Architecture




Address Translation Architecture

Page table size: se/p .
1MB process \rightarrow 2KB
per process.

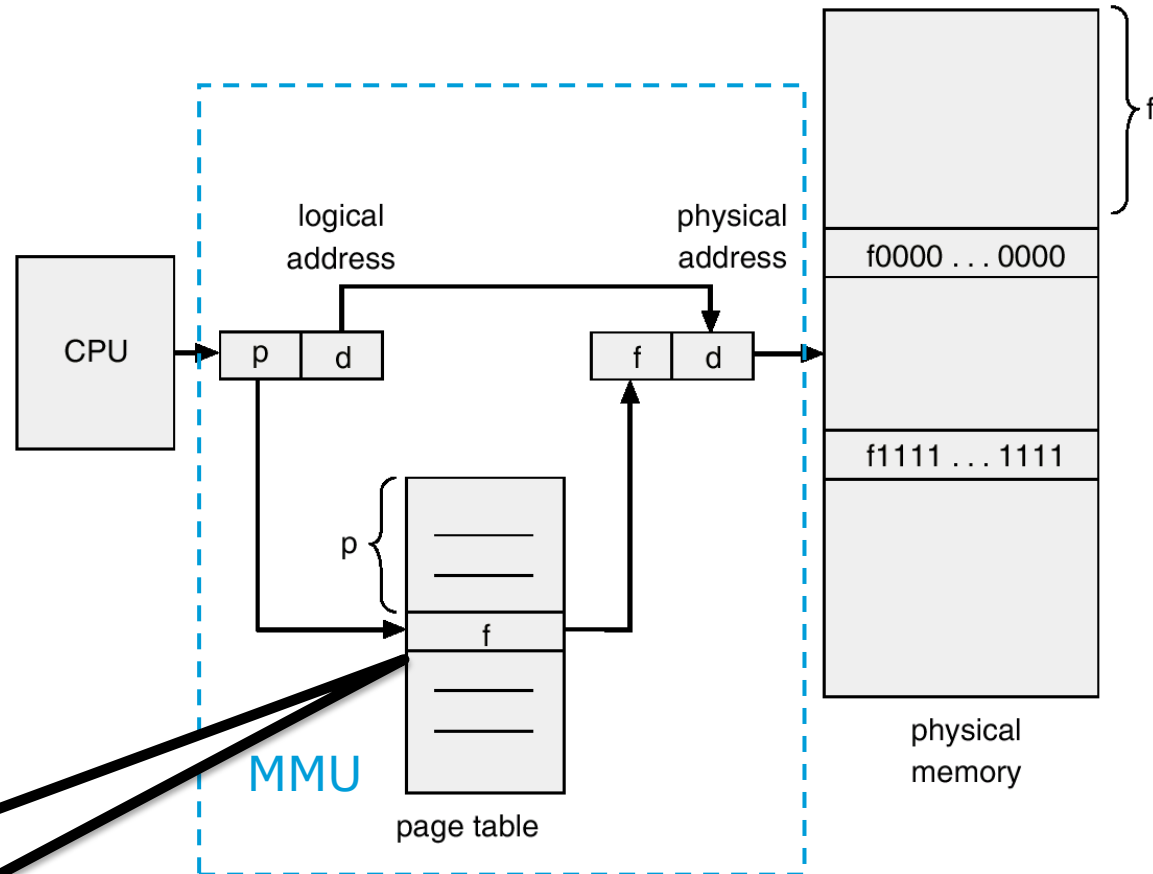
Stored in the physical
memory!



Address Translation Architecture

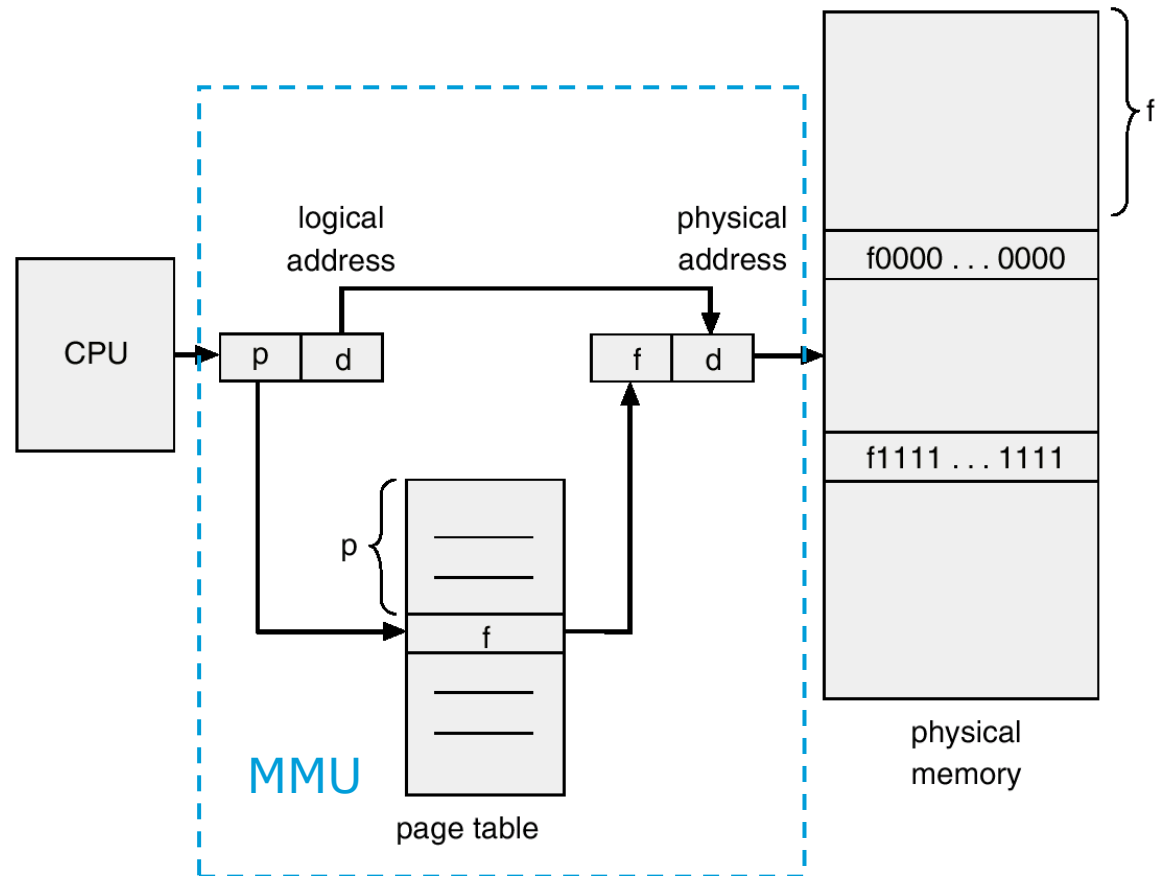


Address Translation Architecture



**Extra memory access
for each memory
access!!**

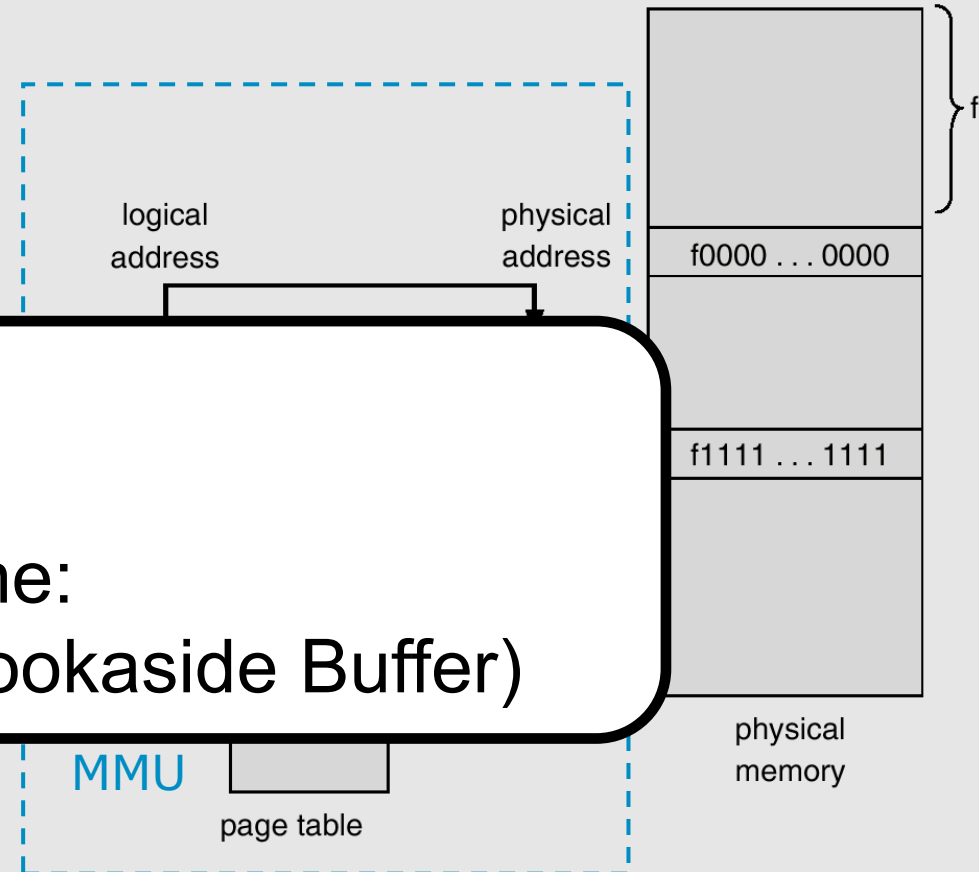
Address Translation Architecture




Address Translation Architecture

Solution: Caching

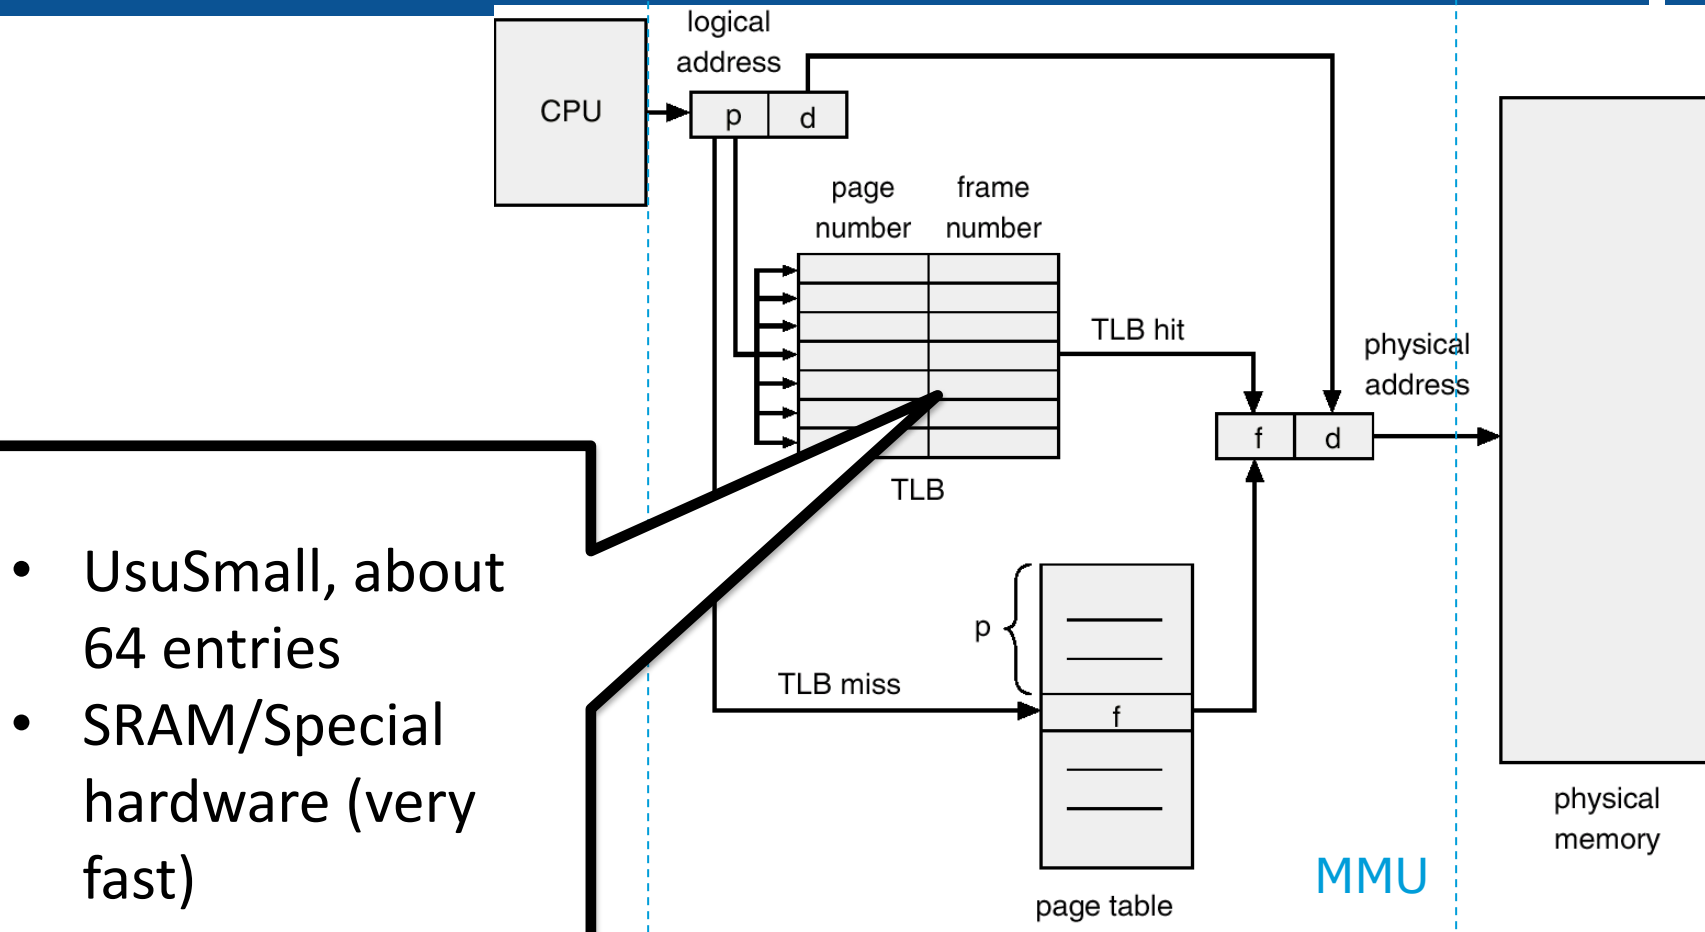
(with a special name:
TLB: Translation Lookaside Buffer)



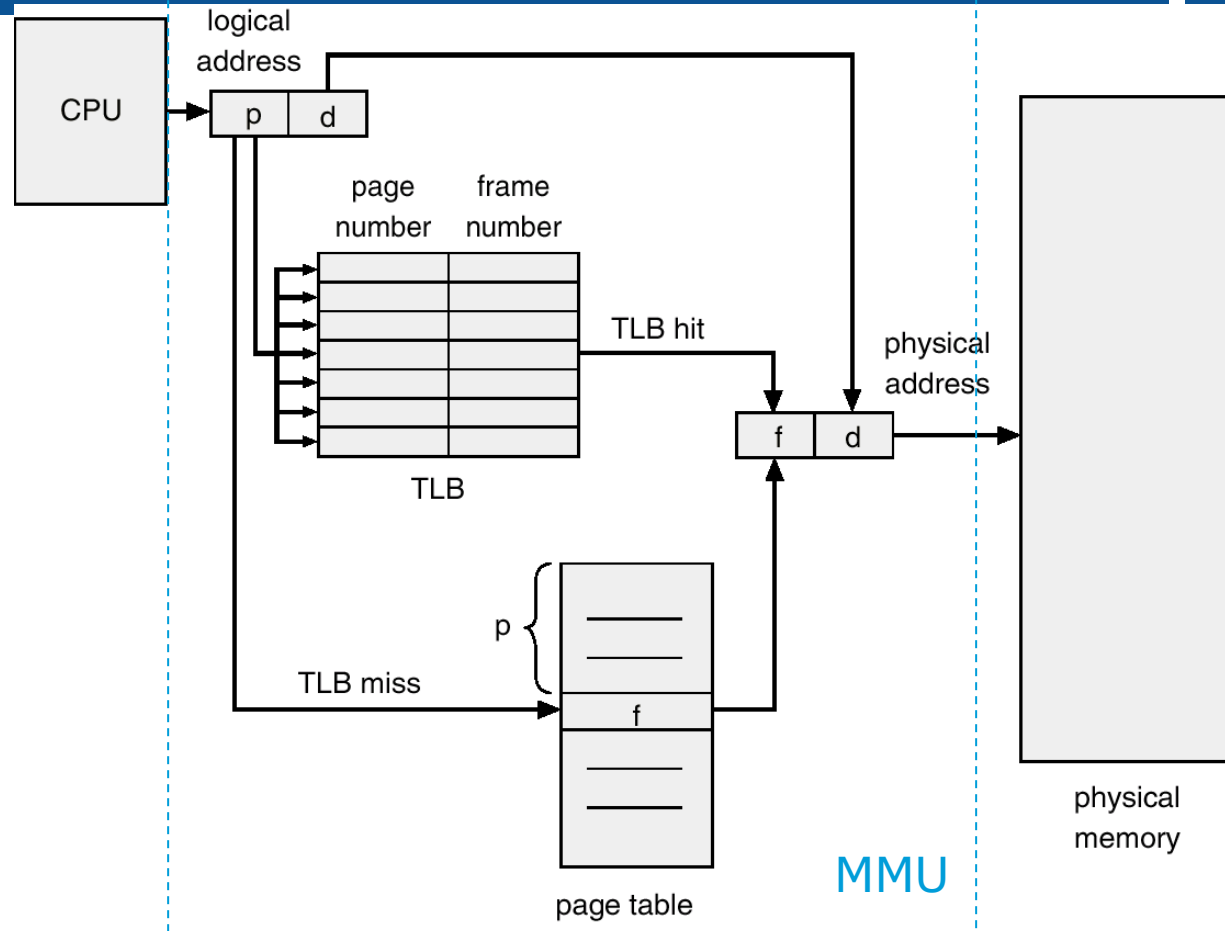
Paging Hardware With TLB



Paging Hardware With TLB



Paging Hardware With TLB



Paging

- OS keeps track of free frames
- When a process requires n pages, allocate any n free frames
 - Remove the frames from the list of free frames
 - Copy data to physical memory
 - Create page table for the process


Virtual Memory

Logical vs. Physical Memory

- Which is larger?


Logical vs. Physical Memory

- Which is larger?
- **Each process** thinks it runs alone
 - May access the entire physical memory




Logical vs. Physical Memory

- Which is larger?
- **Each process** thinks it runs alone
 - May access the entire physical memory
- The number of processes, in general, is not limited



Logical vs. Physical Memory

- Which is larger?
- **Each process** thinks it runs alone
 - May access the entire physical memory
- The number of processes, in general, is not limited



→ The required logical memory (sum of logical memory of all processes) is much larger than the actual physical memory

Virtual Memory

Virtual Memory

- Only part of the program & data needs to be in memory for execution.
 - Logical address space can therefore be much larger than physical address space.
 - Allows for more process to coexist in the main memory.
 - Allows for more efficient process creation.

Virtual Memory

- Only part of the program & data needs to be in memory for execution.
 - Logical address space can therefore be much larger than physical address space.
 - Allows for more process to coexist in the main memory.
 - Allows for more efficient process creation.
- Not currently used pages can be stored in the disk