




# Operating Systems

## Paging and Virtual Memory

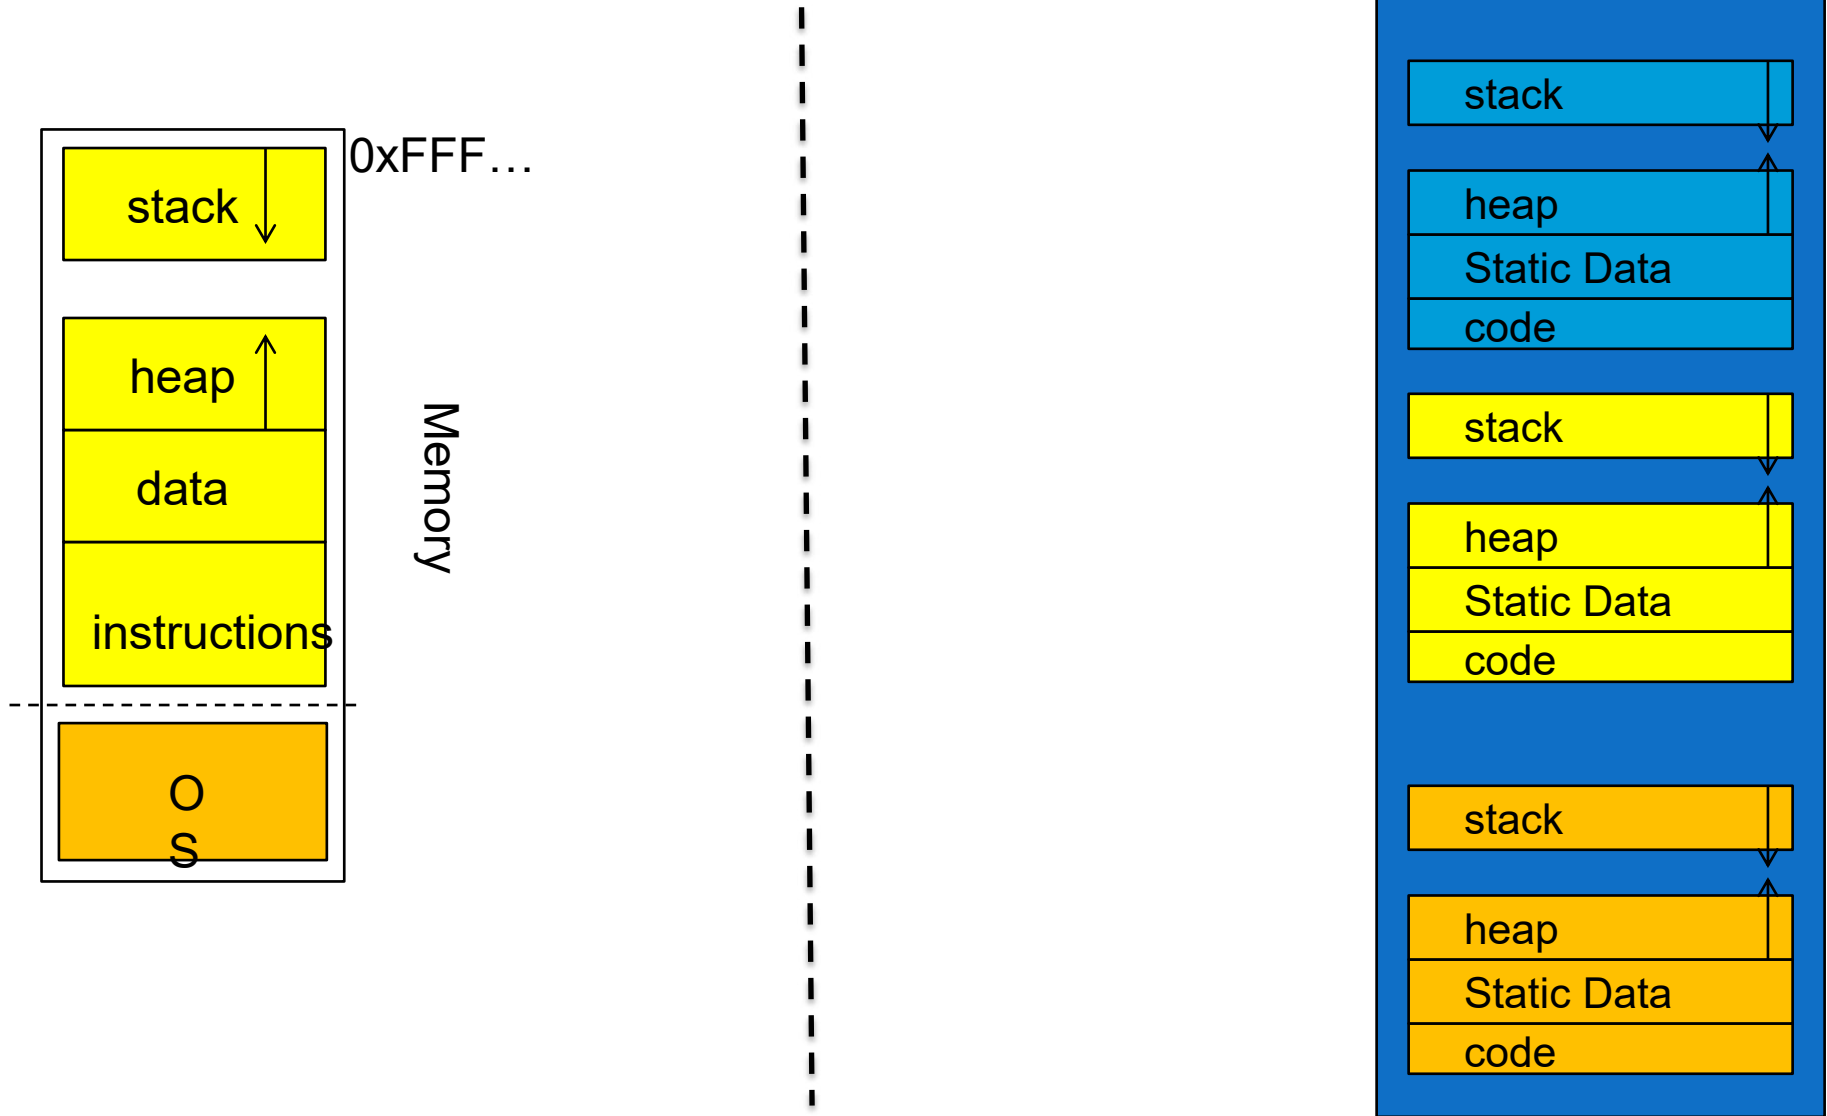
David Hay

Dror Feitelson

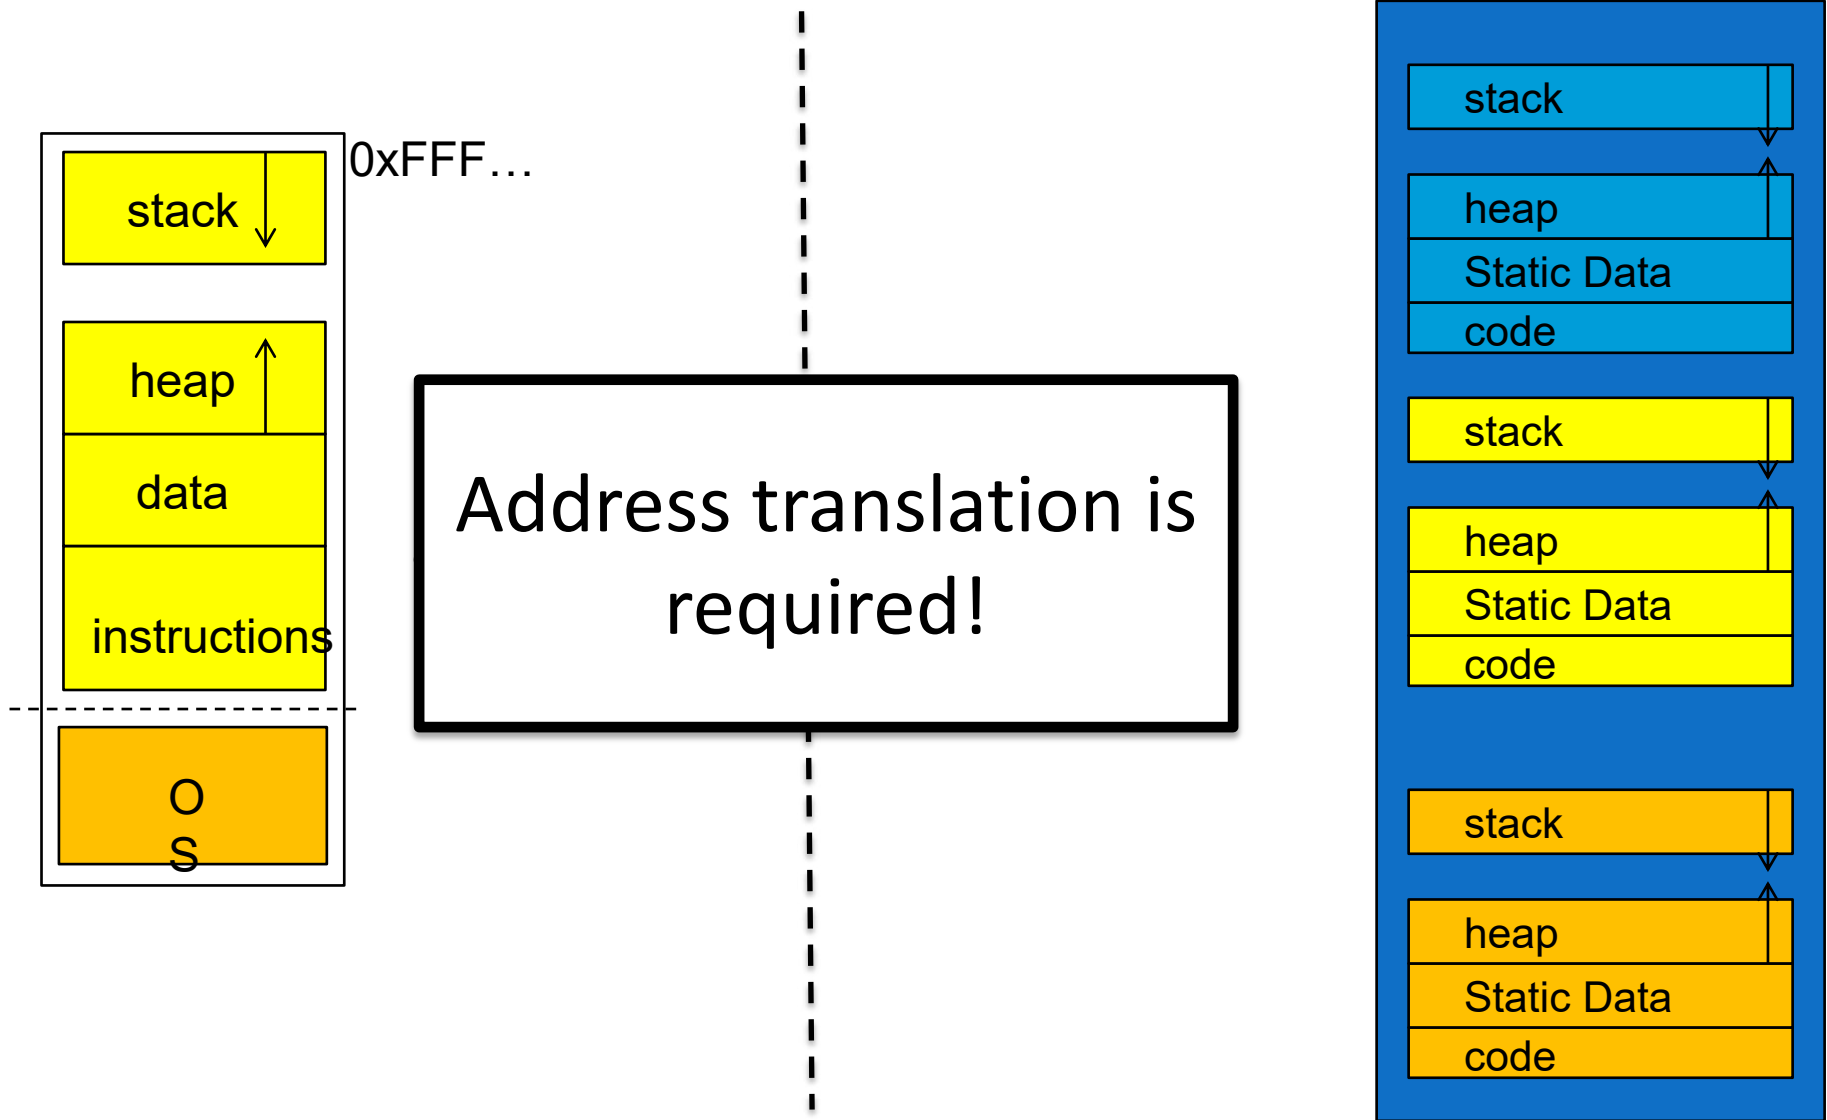
# Process View Vs. Reality




# Process View Vs. Reality



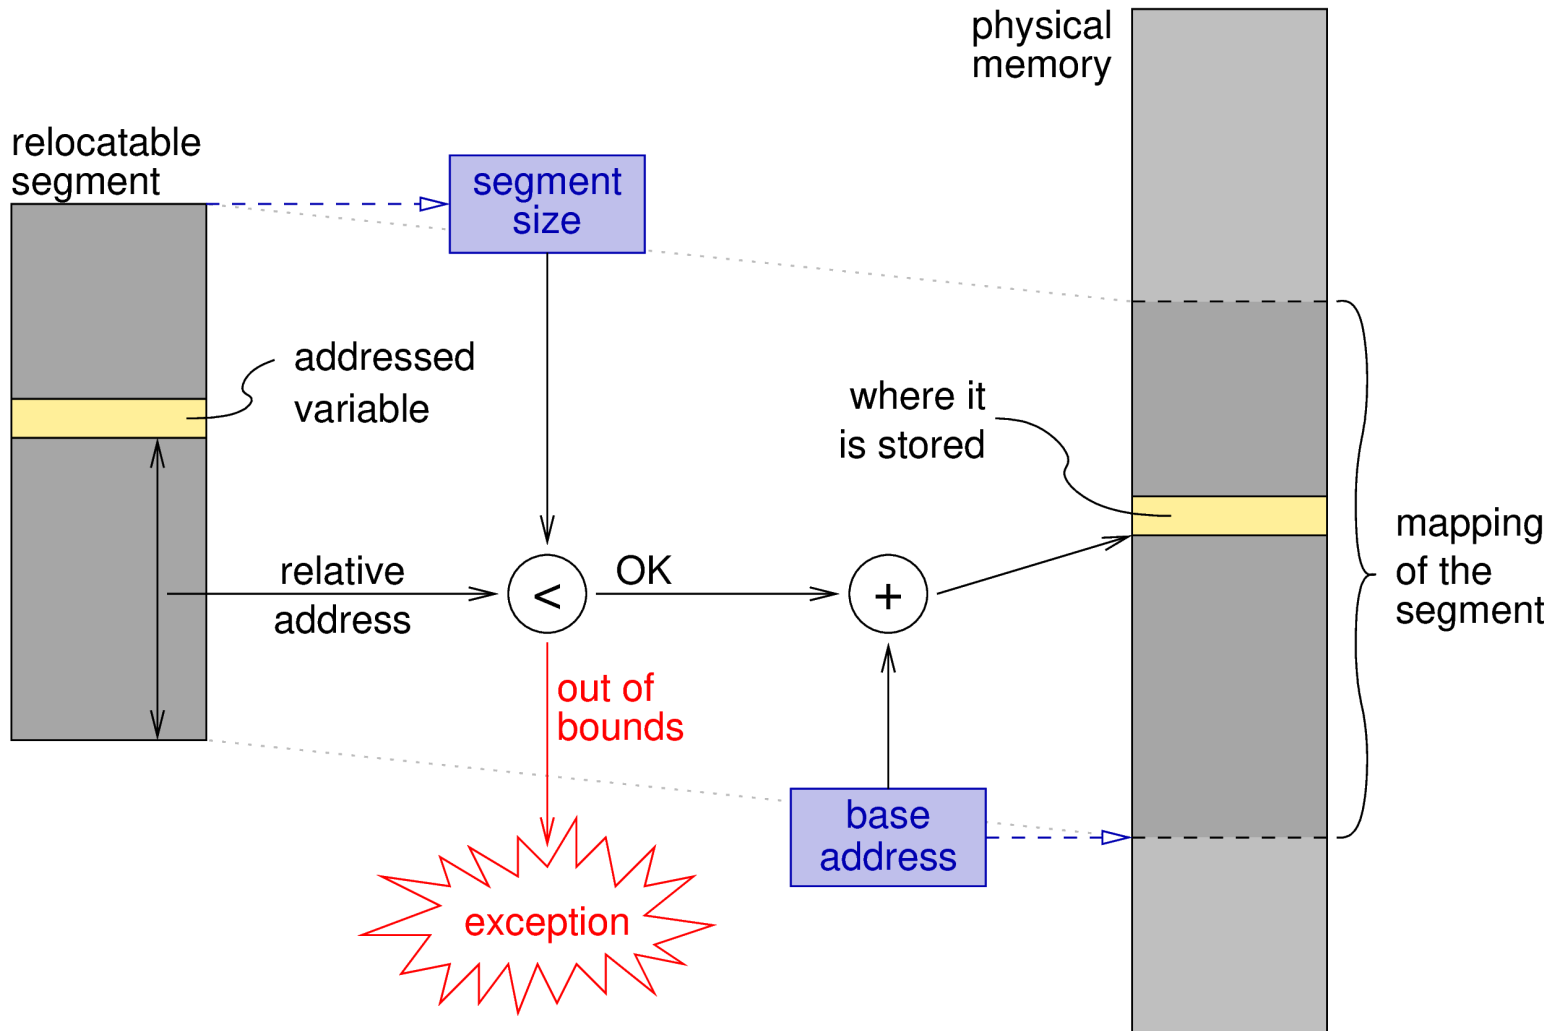
# Process View Vs. Reality



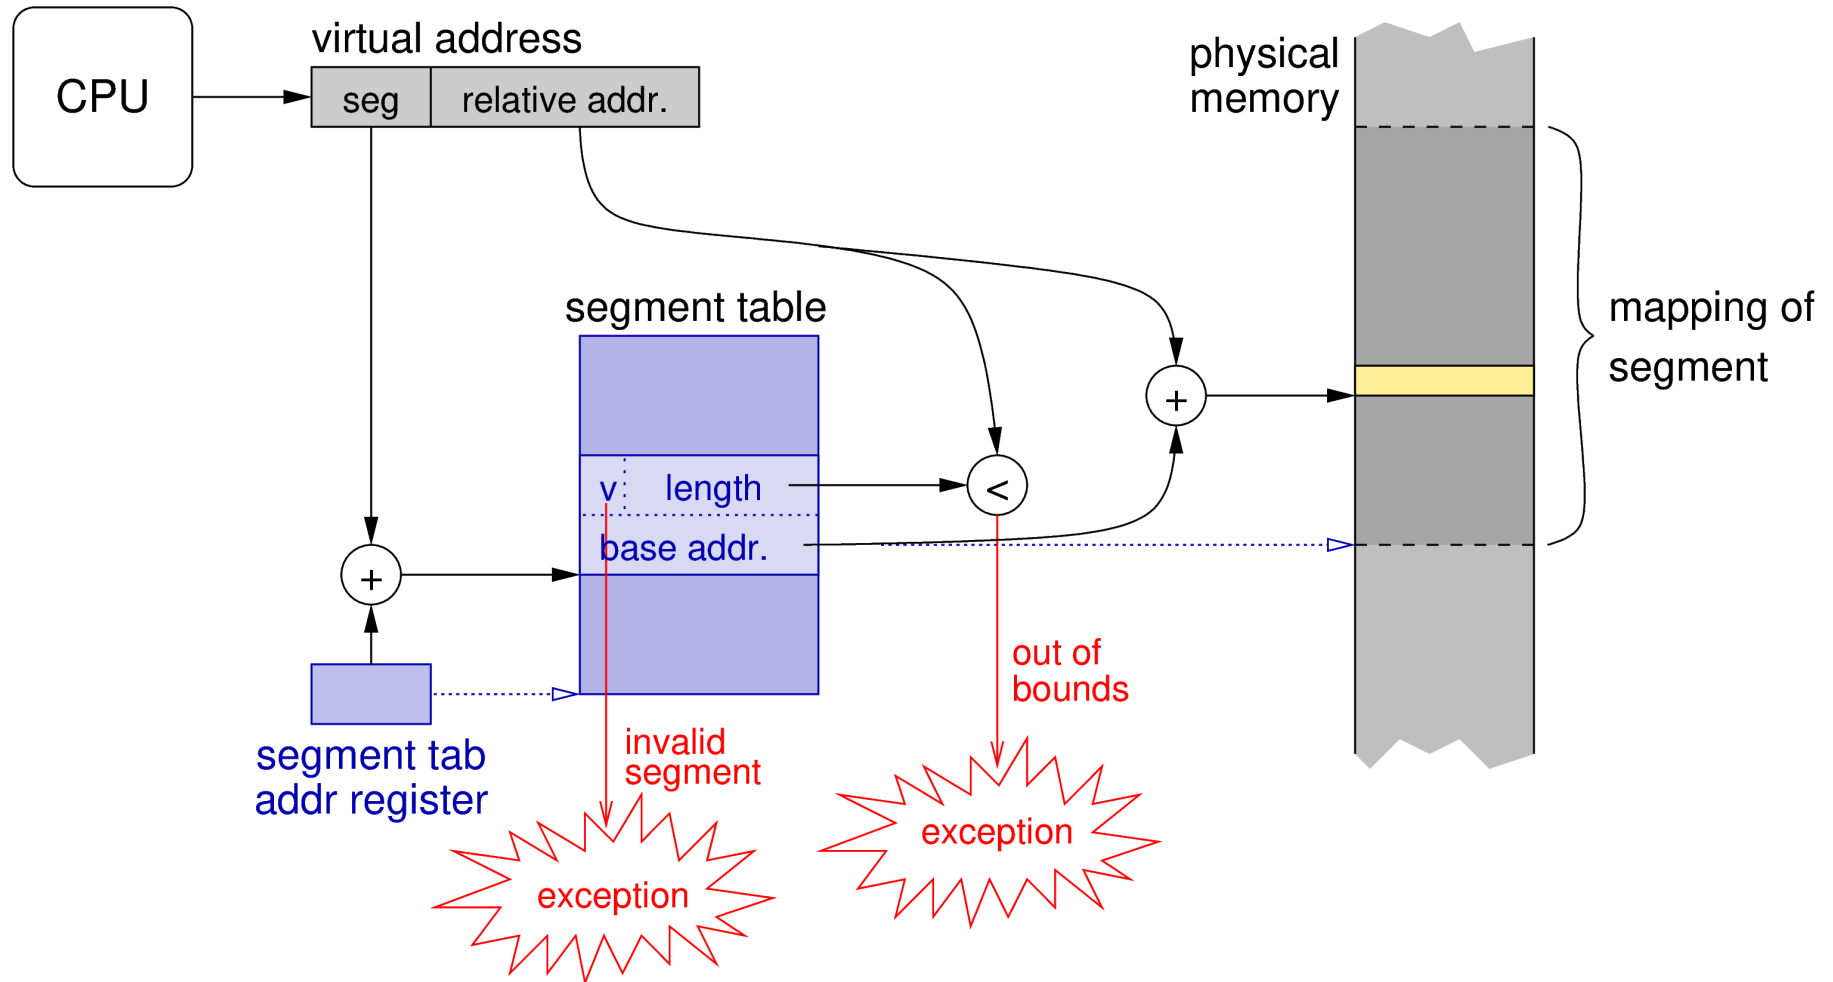
# Memory Addressing Architecture



# Segment Address Translation



# Using a Segment Table



# Fragmentation

- Situation in which we have enough memory to allocate to a process, but not contiguously, as the holes are scattered all over memory
  - Internal Fragmentation: free memory inside process' allocation
  - External Fragmentation: free memory between processes' allocations
- Contiguous allocation suffers from external fragmentation



# The Solution: Paging

- Divide process address space into fixed-size **pages** (usually 4KB)
- Divide physical memory into **frames** of the same size
- Any page can be mapped to any frame
  - No external fragmentation!
- Mapping stored in page table
- Do not need to map a whole segment
  - Only the parts we are using

# The Page Table

- Maps pages to frames
- Separate page table for each process
  - Or each segment
  - Switch tables as part of context switch
- Populated by the operating system
  - Reflects decisions what to map where
- Used by the MMU
  - To perform memory access at hardware speed

# Address Translation with the Page Table

Logical  
addr. space

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

**Used by MMU**

Physical  
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

# Address Translation with the Page Table

Address  
6=00110

Logical  
addr. space

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

**Used by MMU**

Physical  
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

# Address Translation with the Page Table

Address  
6 = 00110

→

Page | Offset  
001 | 10

Logical  
addr. space

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

**Used by MMU**

Physical  
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

# Address Translation with the Page Table

Address  
6 = 00110

Page | Offset  
001 | 10

Logical  
addr. space

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

**Used by MMU**

Physical  
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

# Address Translation with the Page Table

Address  
6=00110

PageOffset  
00110

Logical  
addr. space

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page Table

Page	Frame
0	6
1	3
2	1
3	4

(per process)

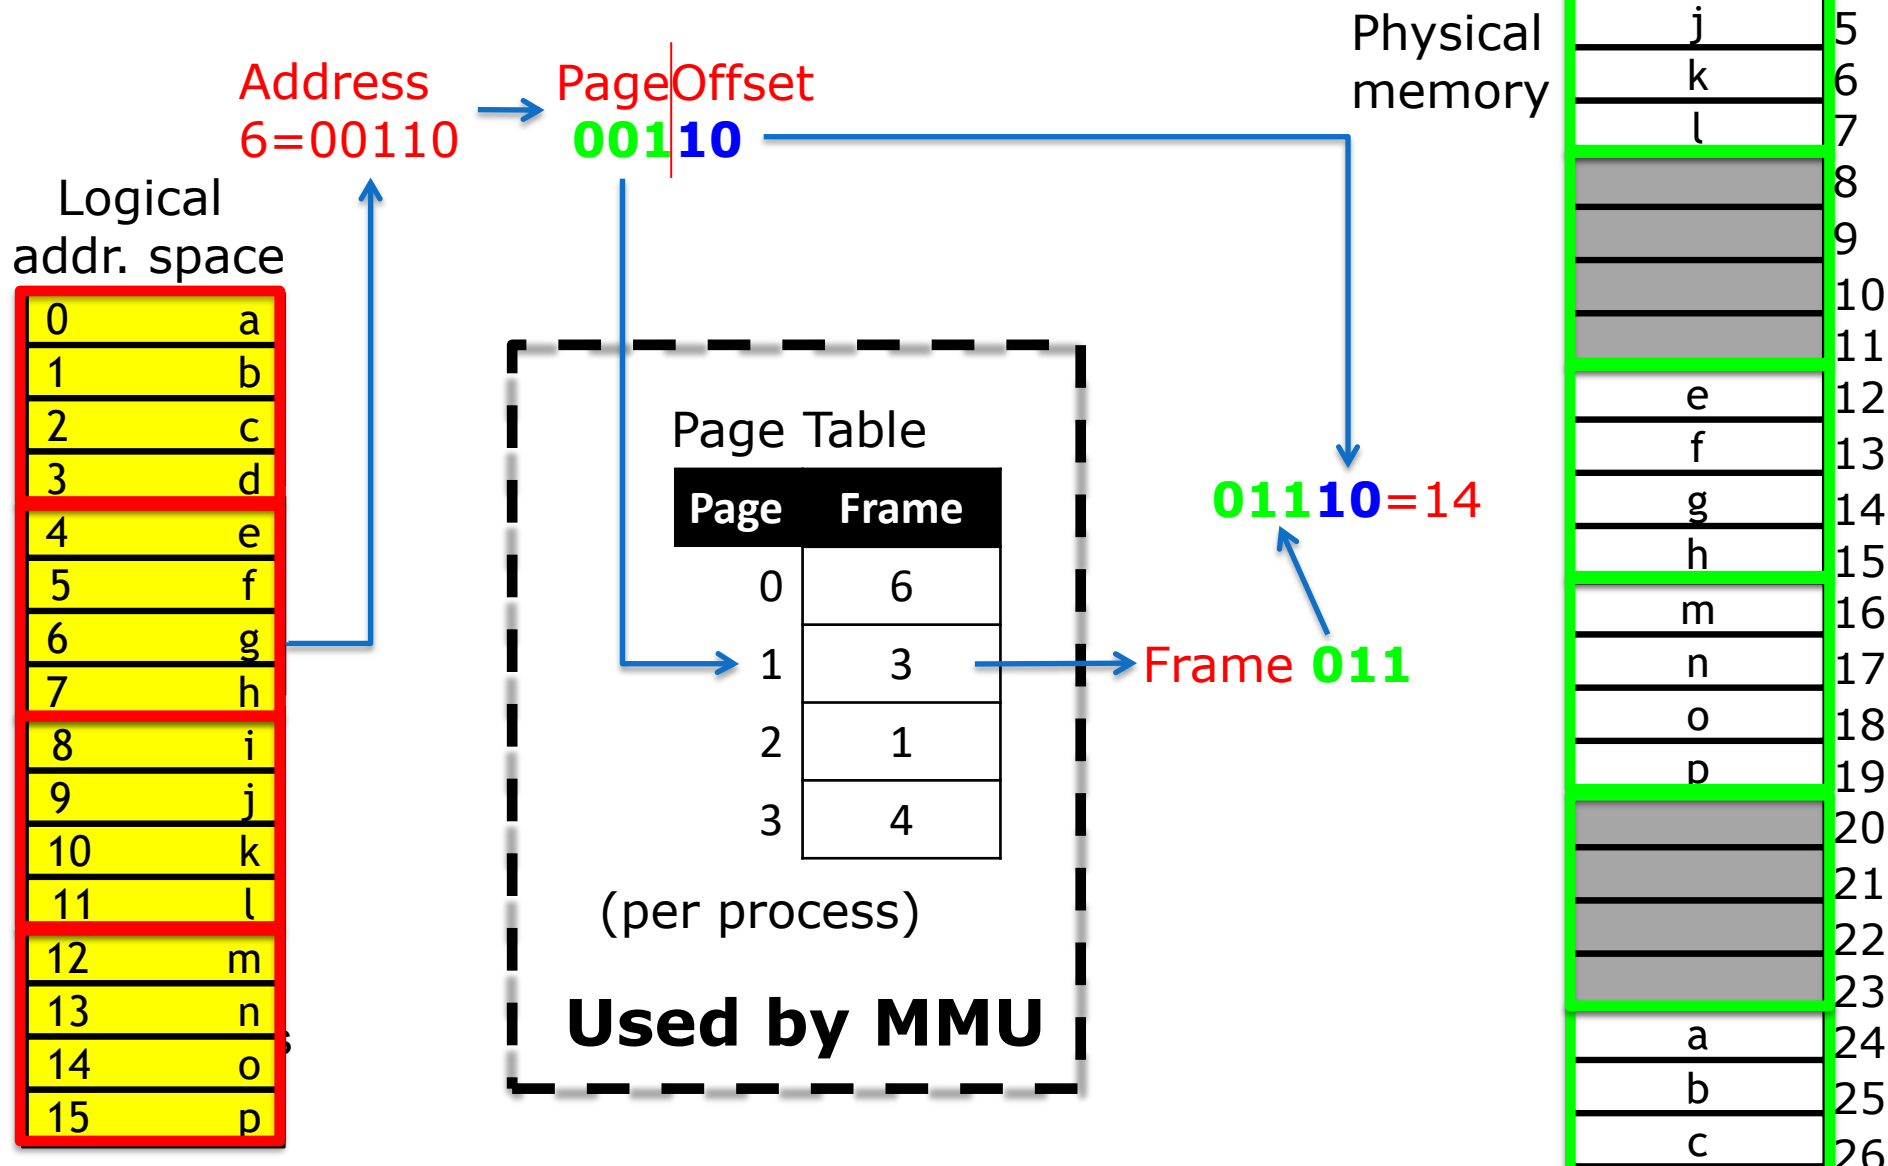
**Used by MMU**

Frame 011

Physical  
memory

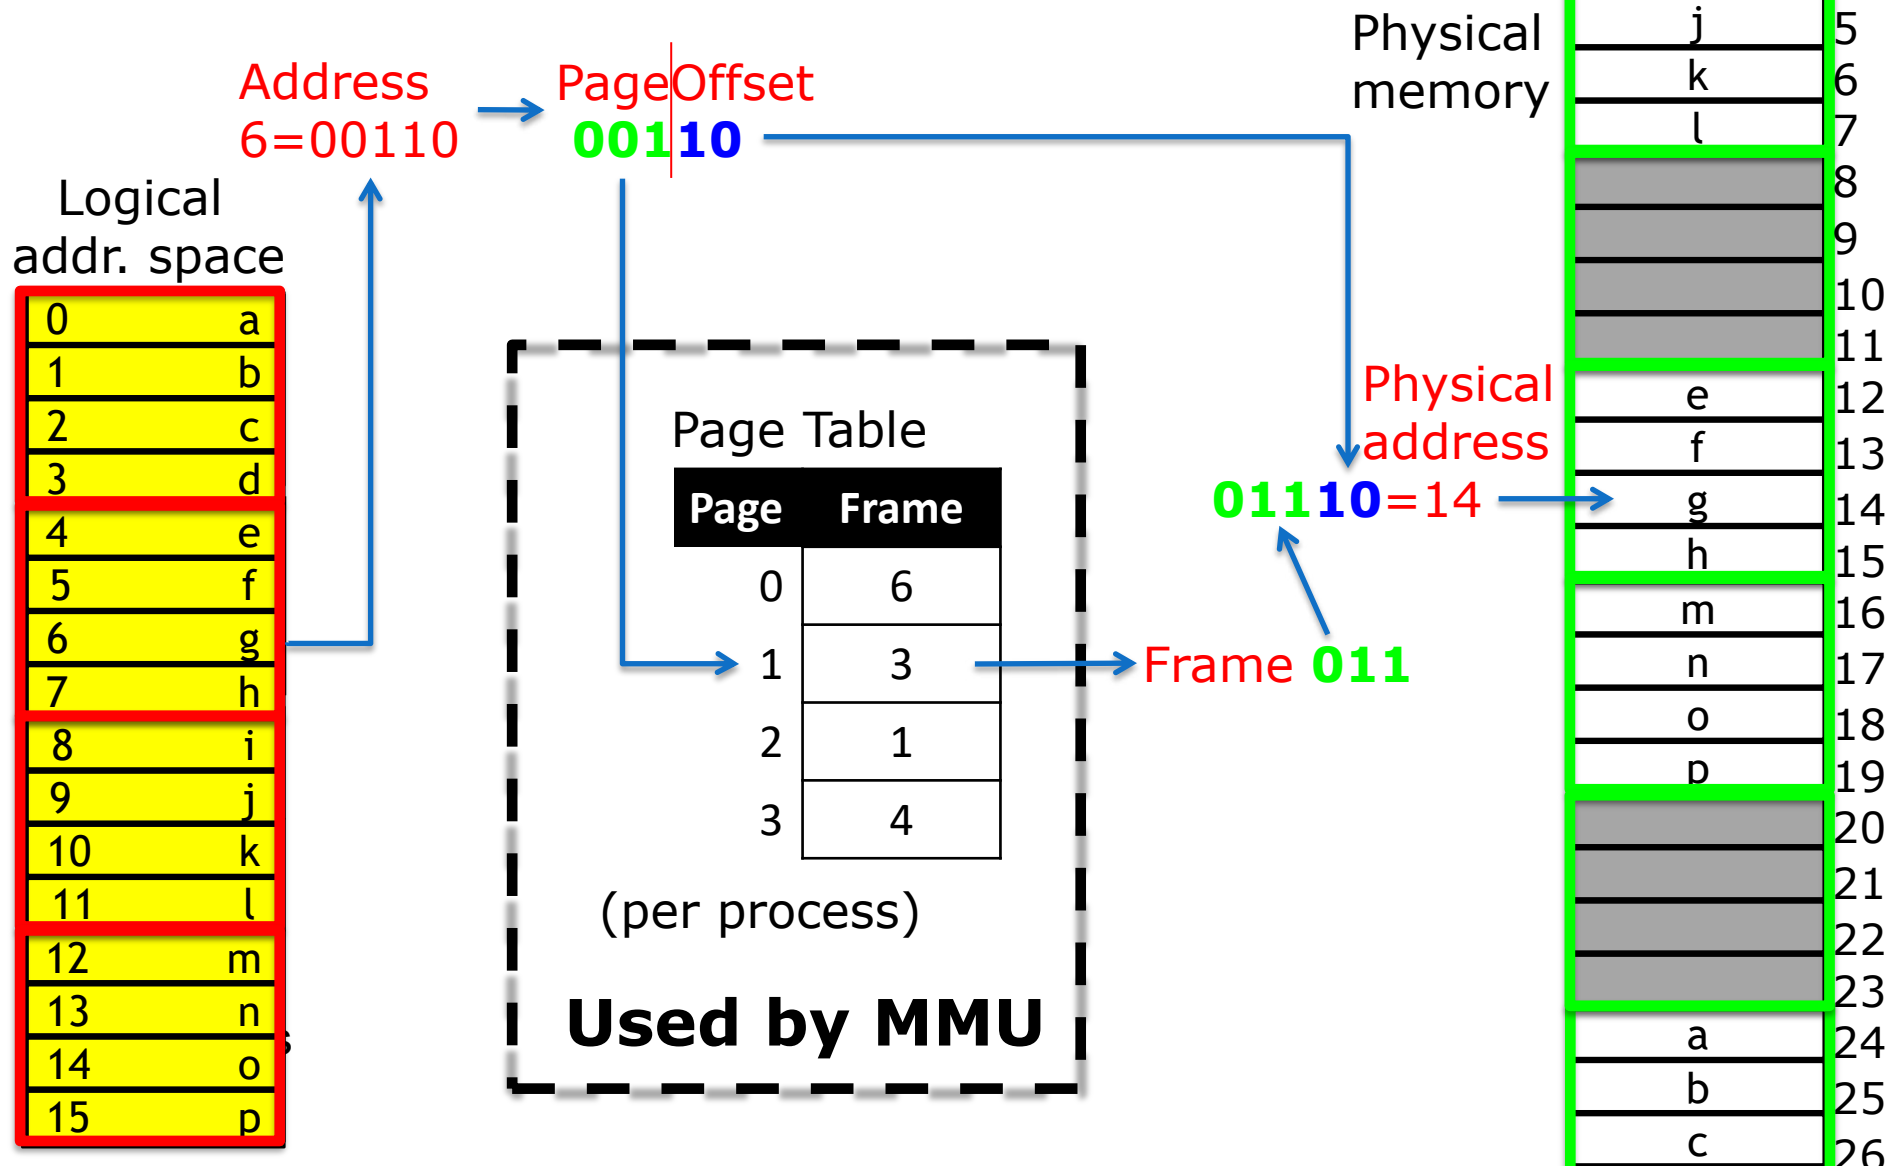
	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

# Address Translation with the Page Table





# Address Translation with the Page Table



# Overheads

- Page tables take up a lot of space
  - Need to be stored in memory
  - Improvement: optimize page size
  - Use sophisticated page table structures [later]
- Accessing the page table takes an additional memory access!
  - This is slow
  - Solution: cache recently used translations in the TLB

# Optimal page size

- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables

# Optimal page size

- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables
- $p$  = page size  
 $s$  = size of the process  
 $e$  = size of the entry in the page table  
→ **Overhead =  $(se/p) + p/2$**

# Optimal page size

- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables
- $p$  = page size  
 $s$  = size of the process  
 $e$  = size of the entry in the page table  
→ **Overhead =  $(se/p) + p/2$**

Page table size:  $s/p$   
pages →  $s/p$  entries,  
each of size  $e$

# Optimal page size

- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables
- $p$  = page size  
 $s$  = size of the process  
 $e$  = size of the entry in the page table  
→ **Overhead =  $(se/p) + p/2$**

Page table size:  $s/p$   
pages →  $s/p$  entries,  
each of size  $e$

Internal fragmentation

# Optimal page size

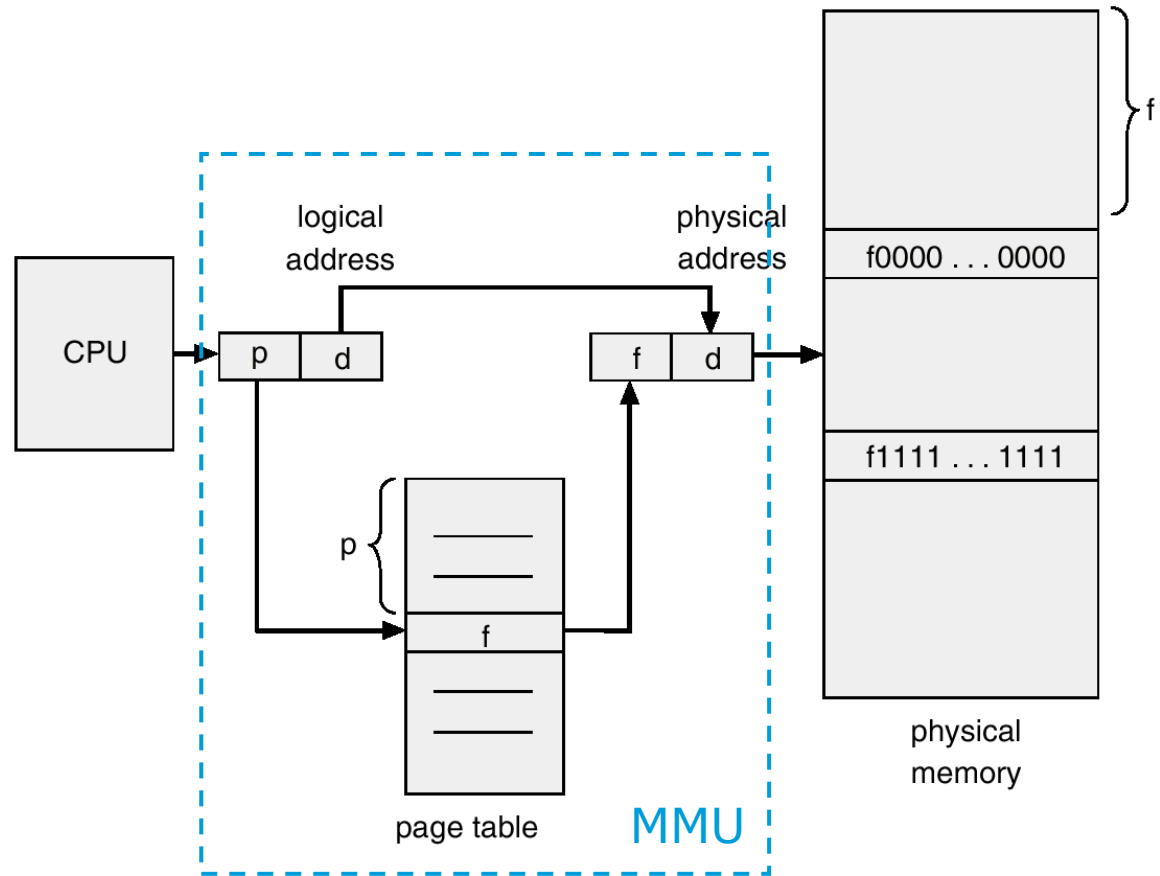
- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables
- $p$  = page size  
 $s$  = size of the process  
 $e$  = size of the entry in the page table
  - **Overhead =  $(se/p) + p/2$**
  - **Optimal page size is  $\sqrt{2se}$**

# Optimal page size

- The page size tradeoff:
  - Small pages have less internal fragmentation
  - Small pages require larger page tables
- $p$  = page size  
 $s$  = size of the process  
 $e$  = size of the entry in the page table  
→ **Overhead =  $(se/p) + p/2$**   
→ **Optimal page size is  $\sqrt{2se}$**
- For  $s=1\text{MB}$ ,  $e=64\text{ bit}$  →  $p=4\text{KB}$



# Address Translation Architecture



# Address Translation Architecture

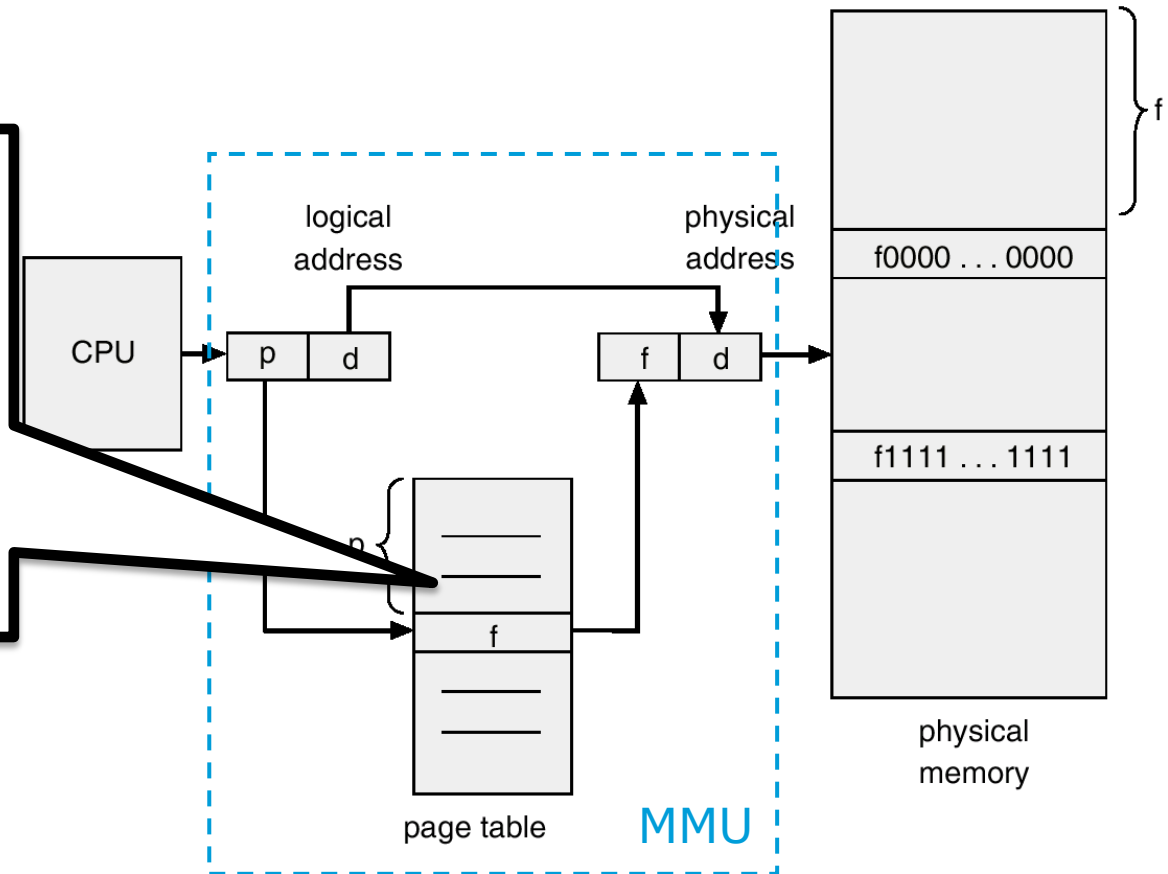
**Page table size:**

se/p

1MB process →

2KB *per process*

Need to stored it  
in memory!



# Address Translation Architecture

**Page table size:**

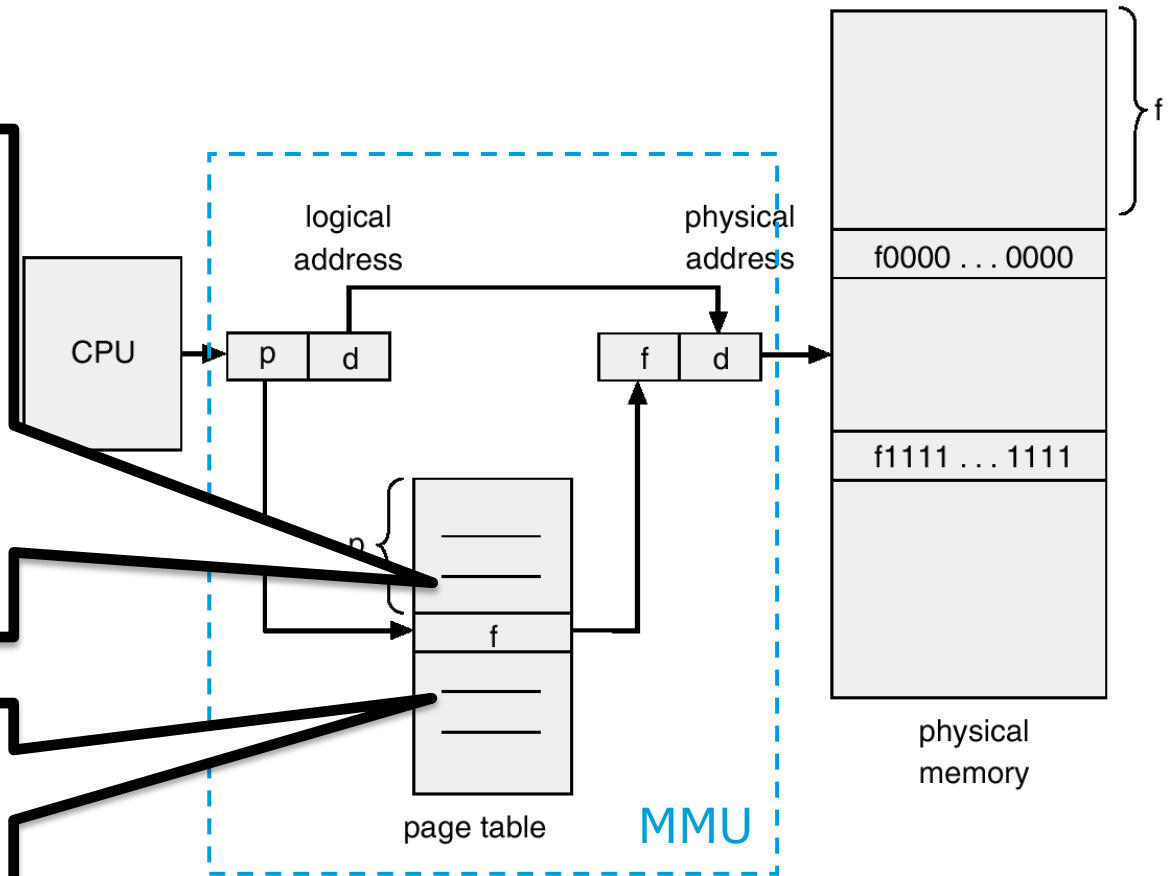
se/p

1MB process  $\rightarrow$

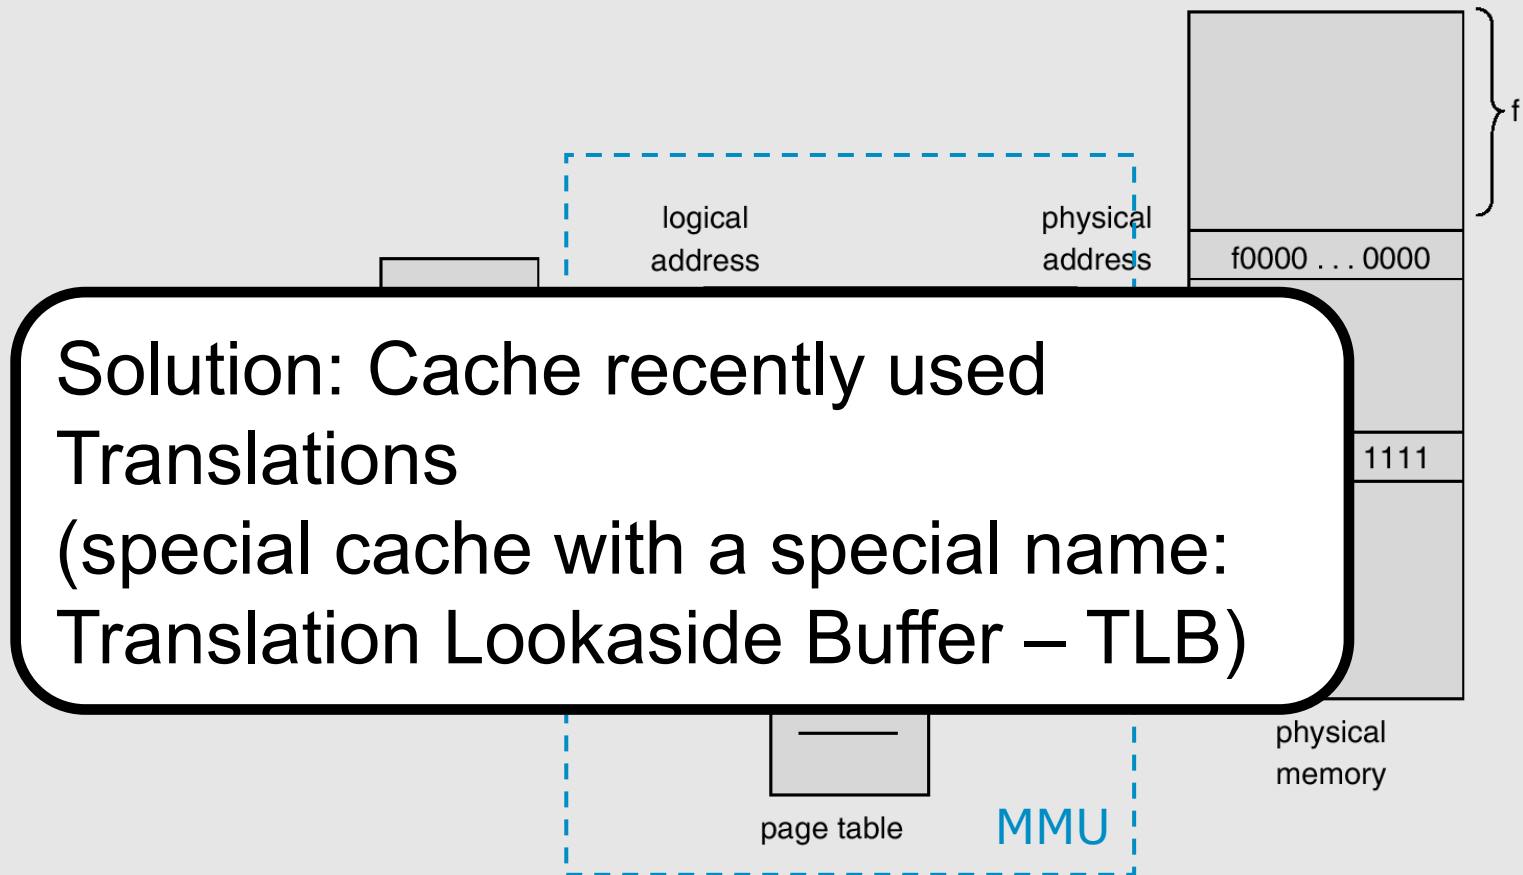
2KB *per process*

Need to stored it  
in memory!

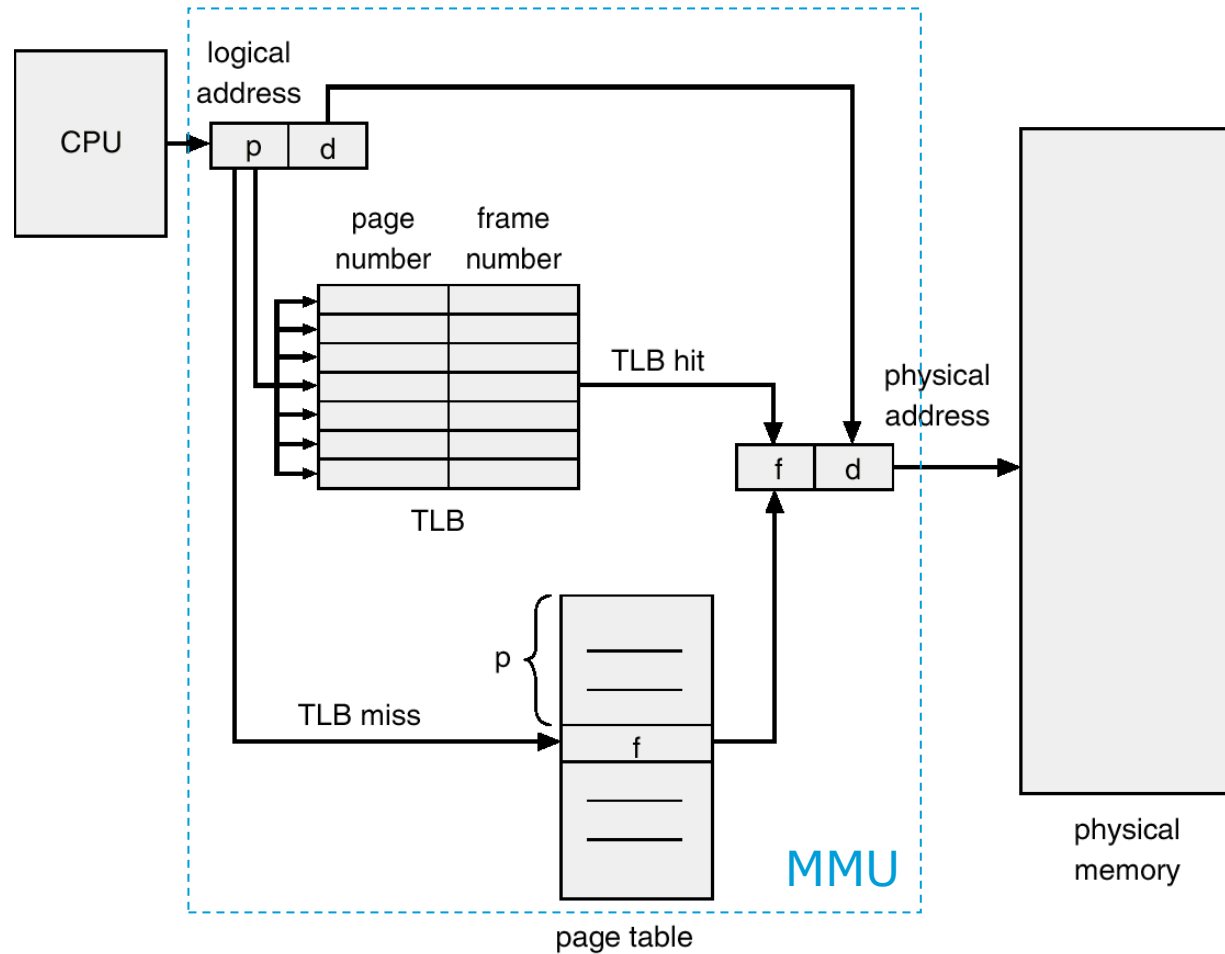
**Extra memory  
access for each  
memory access!!**



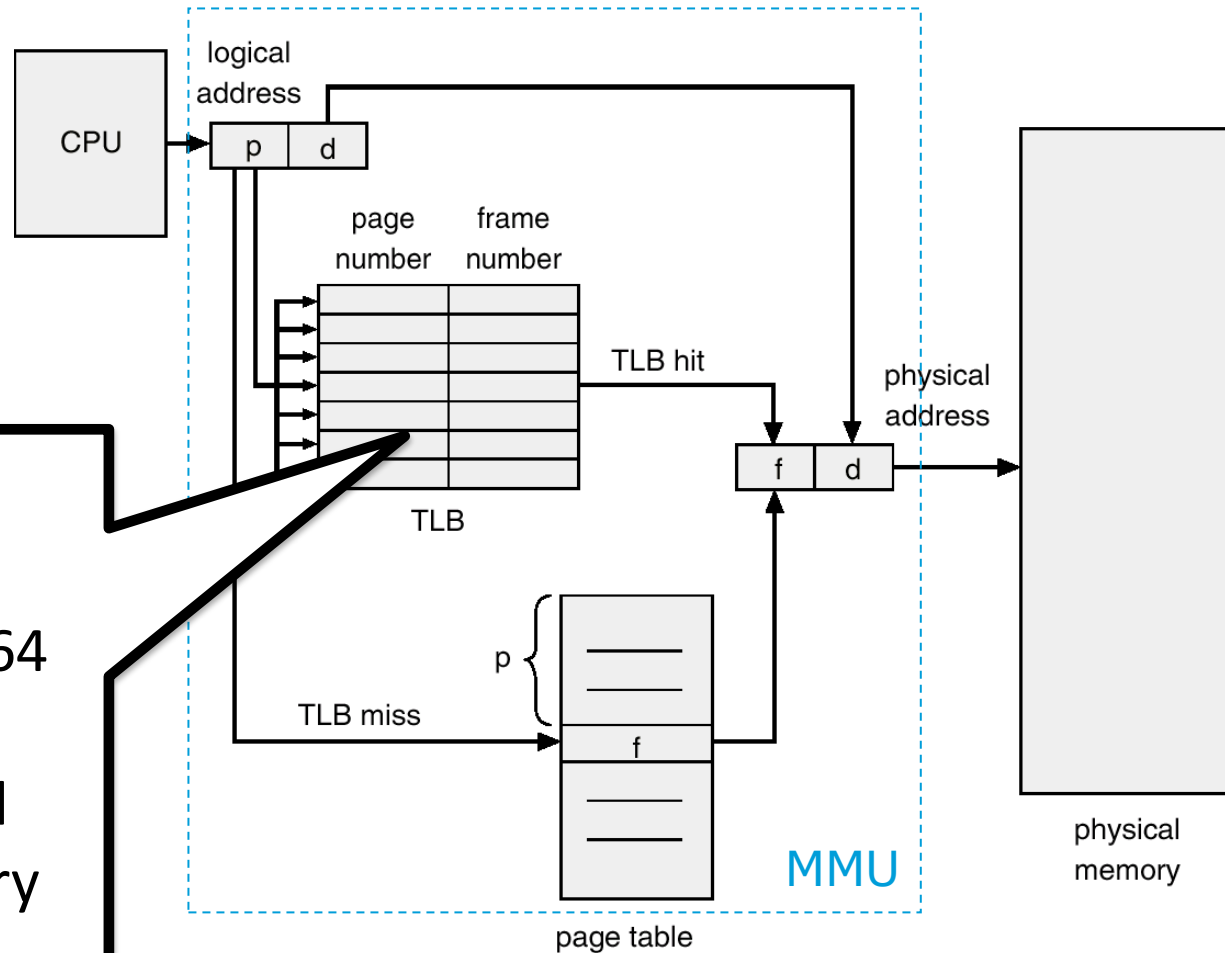
# Address Translation Architecture



# Address Translation with a TLB



# Address Translation with a TLB



- Usually fully associative
- Small, about 64 entries
- SRAM/Special hardware (very fast)

# VIRTUAL MEMORY


# Logical vs. Physical Memory

- Which is larger?




# Logical vs. Physical Memory

- Which is larger?
- **Each process thinks it runs alone**
  - May access the entire physical memory




# Logical vs. Physical Memory

- Which is larger?
- **Each process** thinks it runs alone
  - May access the entire physical memory
- The number of processes, in general, is not limited



# Logical vs. Physical Memory

- Which is larger?
- **Each process** thinks it runs alone
  - May access the entire physical memory
- The number of processes, in general, is not limited



→ The required logical memory (sum of logical memory of all processes) is much larger than the actual physical memory


# Important Observation

- Programs don't use all their address space all the time
  - Program phases: don't need initialization code after you finish it
  - May not need error handling code ever
  - Use one data structure then another
- Unused parts don't need to be mapped to memory
  - Can be stored on disk until needed
  - Reduces memory pressure
  - Allows more efficient process creation


# Virtual Memory

- The idea: **VIRTUALIZATION**
  - Disconnect from the limitations of our physical budget
  - Make it look as if we have all the memory we want
- The implementation: **DEMAND PAGING**
  - Bring pages to memory when we need them
  - Store them on disk when we don't


# Dynamics of Demand Paging



virtual  
memory




physical  
memory




Secondary  
storage  
(hard-disk)

# Dynamics of Demand Paging




# Dynamics of Demand Paging





# Dynamics of Demand Paging



# Add Bit to Page Table

Virtual  
memory

Page Table		
Page	Frame	V
0	2	1
1	3	1
2	X	0
3	X	0
4	1	1
5	0	1
6	X	0

(per process)

**Used by MMU**

0	u
1	v
2	w
3	x
4	q
5	r
6	s
7	t
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Physical  
memory

# Add Bit to Page Table

Virtual  
memory

Page Table

Page	Frame	V
0	2	1
1	3	1
2	x	0
3	X	0
4	1	1
5	0	1
6	X	0

(per process)

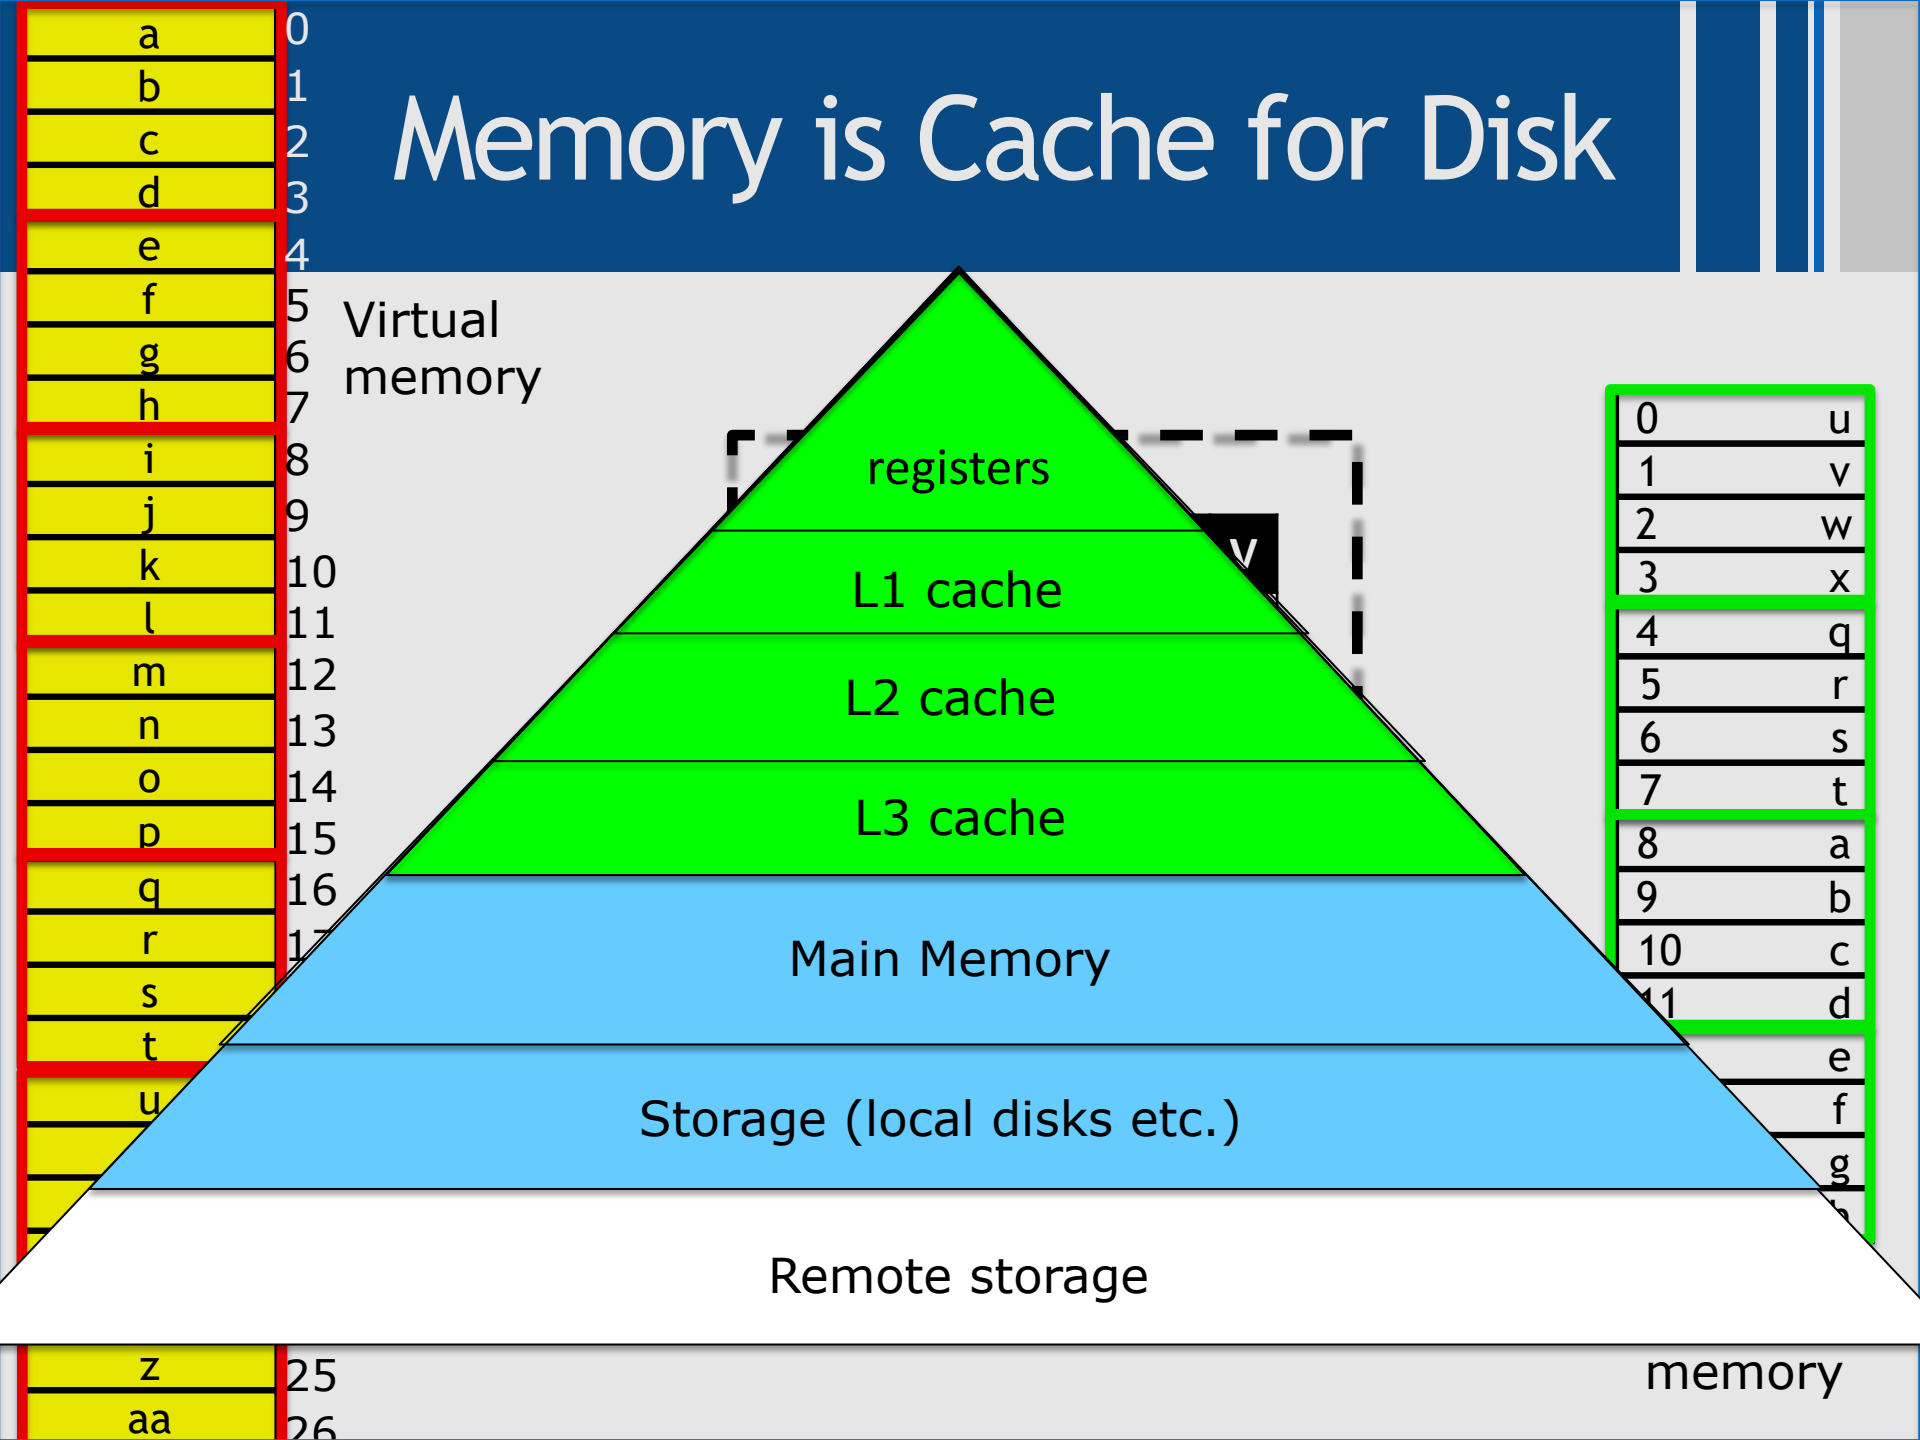
**Used by MMU**

Valid/invalid bit:  
1: Page in memory  
0: Page not in  
memory (the value  
in frame column is  
invalid)

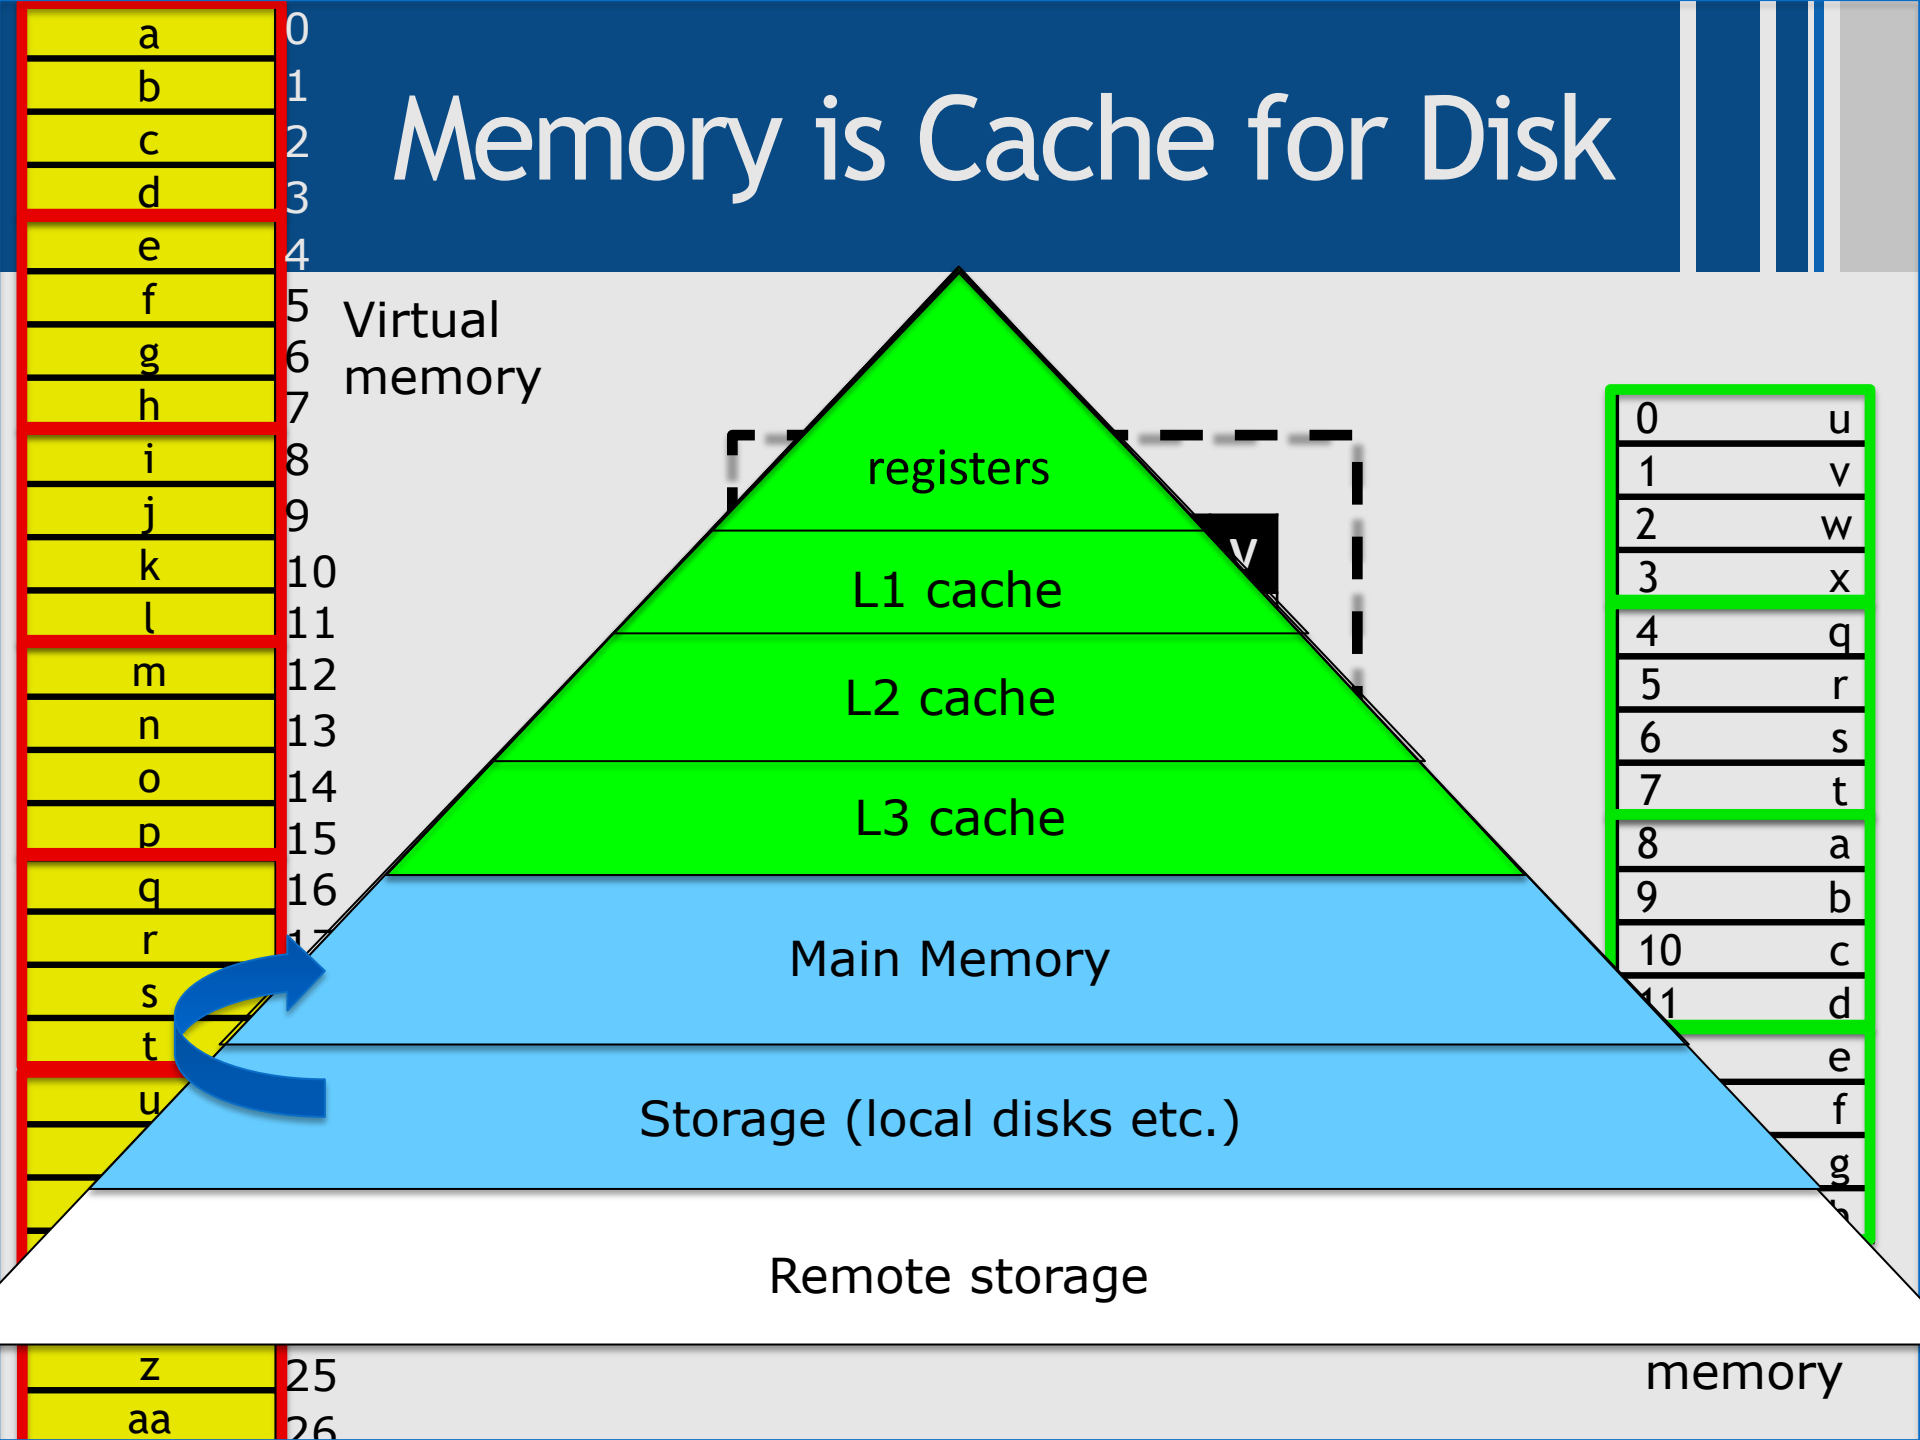
0	u
1	v
2	w
3	x
4	q
5	r
6	s
7	t
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Physical  
memory

# Memory is Cache for Disk



# Memory is Cache for Disk



# Demand Paging

- OS loads a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users

# Demand Paging

- OS loads a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- **Another option: Pre-paging**
  - OS guesses in advance which pages the process will need and pre-loads them into memory
  - Save time if the OS guesses correctly
  - More overhead if the OS is wrong

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page



# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- The data cannot be accessed!

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- The data cannot be accessed!
- MMU creates a PAGE FAULT exception

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- The data cannot be accessed!
- MMU creates a PAGE FAULT exception
- OS exception handler is invoked

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- **The data cannot be accessed!**
- MMU creates a PAGE FAULT exception
- OS exception handler is invoked
- Initiate disk operation to get required data

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- **The data cannot be accessed!**
- MMU creates a PAGE FAULT exception
- OS exception handler is invoked
- Initiate disk operation to get required data
- Process put to sleep until it arrives

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- **The data cannot be accessed!**
- MMU creates a PAGE FAULT exception
- OS exception handler is invoked
- Initiate disk operation to get required data
- Process put to sleep until it arrives
- Run other processes in this time

# What happens when the accessed page is not in memory?

- CPU issues a virtual address that is in an unmapped page
- **The data cannot be accessed!**
- MMU creates a PAGE FAULT exception
- OS exception handler is invoked
- Initiate disk operation to get required data
- Process put to sleep until it arrives
- Run other processes in this time
- When data arrives, awaken process and re-issue **the same instruction**

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26



# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

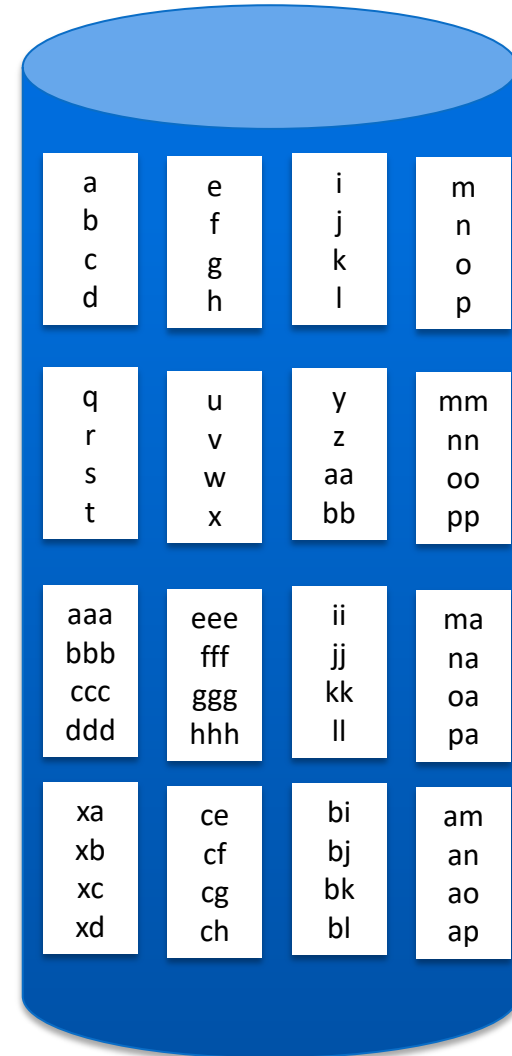

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h



# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

lw 12, R1

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0

exception

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Operating System

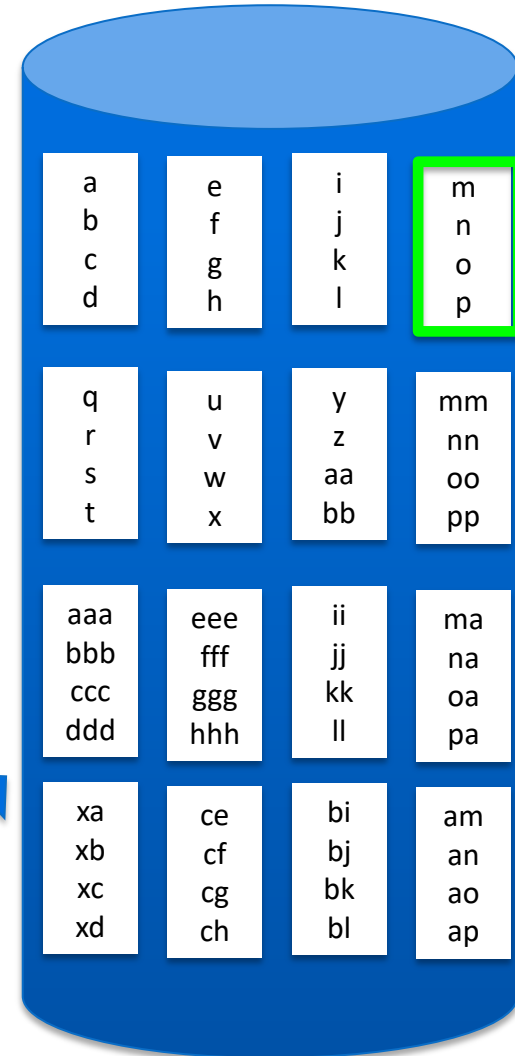

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h



# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

DMA

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

Operating System

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
4	m
5	n
6	o
7	p
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

Operating System



# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	0
5	0	1
6	0	0

0	u
1	v
2	w
3	x
4	m
5	n
6	o
7	p
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

interrupt

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	3	0
4	1	
5	0	
6	0	

0	u
1	v
2	w
3	x
4	m
5	n
6	o
7	p
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

interrupt

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

Page	Frame	
0	2	1
1	3	1
2	2	0
3	1	1
4	1	
5	0	
6	0	

0	u
1	v
2	w
3	x
4	m
5	n
6	o
7	p
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

0	u
1	v
2	w
3	x
4	m

Page	Frame	
0	2	
1	3	
2	2	
3	1	
4	1	
5	0	1
6	0	0

Change process state  
to ready and  
**restart instruction**

13	f
14	g
15	h

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# Handling Page Faults

a	0
b	1
c	2
d	3
e	4
f	5
g	6
h	7
i	8
j	9
k	10
l	11
m	12
n	13
o	14
p	15
q	16
r	17
s	18
t	19
u	20
v	21
w	22
x	23
y	24
z	25
aa	26

lw 12, R1

Page	Frame	
0	2	1
1	3	1
2	2	
3	1	1
4	1	0
5	0	1
6	0	0

write m  
to R1

0	u
1	v
2	w
3	x
4	m
5	n
6	o
7	p
8	a
9	b
10	c
11	d
12	e
13	f
14	g
15	h

Operating System

a b c d	e f g h	i j k l	m n o p
q r s t	u v w x	y z aa bb	mm nn oo pp
aaa bbb ccc ddd	eee fff ggg hhh	ii jj kk ll	ma na oa pa
xa xb xc xd	ce cf cg ch	bi bj bk bl	am an ao ap

# But There's A Slight Problem

# But There's A Slight Problem

What if there is no  
free frame to map  
the page to?

# Page Eviction

- Need to select some mapped page and evict it
  - Copy its data to disk
  - Use the frame for the new page
- Which page should it be?
  - Called “the victim”



# Replacement Algorithms

- How do we choose the *victim*?

# Replacement Algorithms

- How do we choose the *victim*?
  - We can just choose at random...

# Replacement Algorithms

- How do we choose the *victim*?
  - We can just choose at random...
  - But better not to choose often used pages (will probably need to be brought back in soon)

# Replacement Algorithms

- How do we choose the *victim*?
  - We can just choose at random...
  - But better not to choose often used pages (will probably need to be brought back in soon)
- Many policies are possible
  - Optimal
  - Random
  - FIFO (first-in-first-out), second chance FIFO
  - NRU (not recently used)
  - LRU (least recently used), pseudo-LRU
  - LFU (least frequently used)
  - Etc

# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time

# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5






# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5


5

1
2
3
4

# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5






# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm



- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm


- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5




- **Infeasible**: need to know the future

# (Infeasible) Optimal Algorithm

a.k.a. Bélády's Algorithm, clairvoyant algorithm

- Replace the page that **will not be used** for the longest period of time
- Page numbers: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



- **Infeasible**: need to know the future
- Used only for comparison
  - Given an algorithm, how close is it to optimal?

# Random Replacement

- Just evict any page randomly
- The other extreme from optimal
  - Optimal uses full knowledge of everything, including the future
  - Random uses no knowledge of anything, including the past
- Also used for comparison
  - Given an algorithm, is it better than random?

# First-In-First-Out (FIFO)

- What information can the OS use?



# First-In-First-Out (FIFO)

- What information can the OS use?
- Sort the pages in the order they were loaded to memory
  - E.g., maintain the order as a link list

# First-In-First-Out (FIFO)

- What information can the OS use?
- Sort the pages in the order they were loaded to memory
  - E.g., maintain the order as a link list
- The victim is the first page (in this order)
  - The “oldest” page

# First-In-First-Out (FIFO)

- What information can the OS use?
- Sort the pages in the order they were loaded to memory
  - E.g., maintain the order as a link list
- The victim is the first page (in this order)
  - The “oldest” page
- Disadvantage: page in the memory the longest may actually be used often

# First-In-First-Out (FIFO)

- What information can the OS use?
- Sort the pages in the order they were loaded to memory
  - E.g., maintain the order as a link list
- The victim is the first page (in this order)
  - The “oldest” page
- Disadvantage: page in the memory the longest may actually be used often
- Used in Windows NT
  - Independent of any hardware support

# Belady's Anomaly

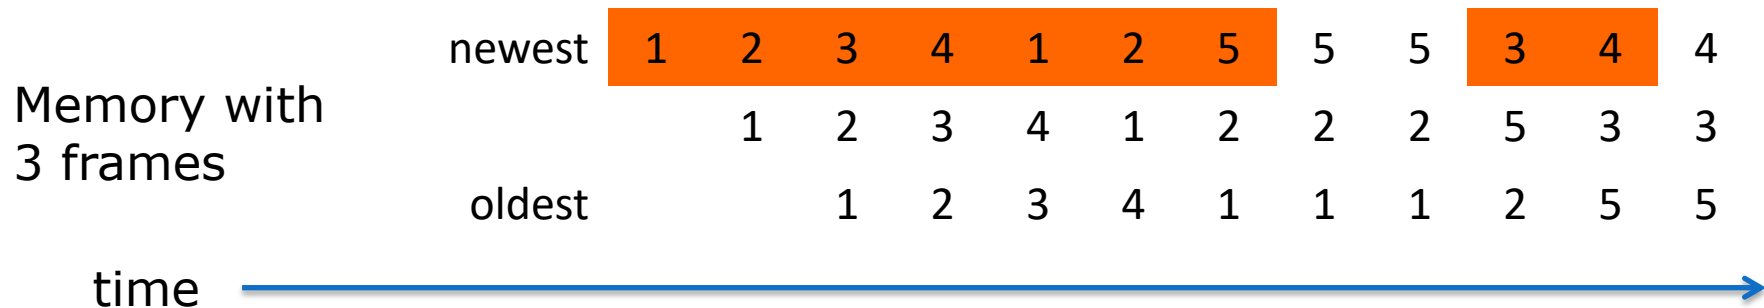
- Excepted behavior: Larger memory  $\rightarrow$  less page faults

# Belady's Anomaly

- Excepted behavior: Larger memory → less page faults
- Consider: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5


# Belady's Anomaly

- Excepted behavior: Larger memory → less page faults
- Consider: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5



# Belady's Anomaly

- Excepted behavior: Larger memory → less page faults
- Consider: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

	newest	1	2	3	4	1	2	5	5	5	3	4	4
Memory with 3 frames			1	2	3	4	1	2	2	2	5	3	3
	oldest			1	2	3	4	1	1	1	2	5	5
	time												
	newest	1	2	3	4	4	4	5	1	2	3	4	5
Memory with 4 frames			1	2	3	3	3	4	5	1	2	3	4
				1	2	2	2	3	4	5	1	2	3
	oldest				1	1	1	2	3	4	5	1	2



# Stop and Think: What Do We Really Want?

# Stop and Think: What Do We Really Want?

- Ideally, to know the future

# Stop and Think: What Do We Really Want?

- Ideally, to know the future
  - But why? What information do we want?

# Stop and Think:

## What Do We Really Want?

- Ideally, to know the future
  - But why? What information do we want?
- We actually want to identify the process's **WORKING SET**
- This is the set of pages it is currently using
  - Pages in the working set should be retained
  - Pages not in the working set can be evicted

# Stop and Think:

## What Do We Really Want?

- Ideally, to know the future
  - But why? What information do we want?
- We actually want to identify the process's **WORKING SET**
- This is the set of pages it is currently using
  - Pages in the working set should be retained
  - Pages not in the working set can be evicted
- Based on the principle of locality

# The Working Set

- So how can we identify the working set?

# The Working Set

- So how can we identify the working set?
- How is it even defined formally?

# The Working Set

- So how can we identify the working set?
- How is it even defined formally?
- Parametric definition: the set of pages accessed in the last  $k$  memory accesses
  - With random access, working set size  $\approx k$
  - With locality, working set size  $\ll k$



# The Working Set

- So how can we identify the working set?
- How is it even defined formally?
- Parametric definition: the set of pages accessed in the last  $k$  memory accesses
  - With random access, working set size  $\approx k$
  - With locality, working set size  $\ll k$
- In theory: start with  $k=1$ , increase  $k$  until working set stabilizes

# The Working Set

- So how can we identify the working set?
- How is it even defined formally?
- Parametric definition: the set of pages accessed in the last  $k$  memory accesses
  - With random access, working set size  $\approx k$
  - With locality, working set size  $\ll k$
- In theory: start with  $k=1$ , increase  $k$  until working set stabilizes
- In practice: (for large  $k$ ) which pages have been used recently?

# Hardware Support

- Memory access is done at clock speed
- So need hardware support to track it
- But need to also limit overhead

Reference bit:

a.k.a. used bit

Turned on when a page is accessed

Dirty bit:

Turned on when a page is modified

- Done by MMU at each access
- Supported on Intel processors

# Not Recently Used (NRU)

- We cannot know the future, so we try to estimate according to the past

# Not Recently Used (NRU)

- We cannot know the future, so we try to estimate according to the past
- Use the pages' reference bits

# Not Recently Used (NRU)

- We cannot know the future, so we try to estimate according to the past
- Use the pages' **reference bits**
- Periodically (on clock interrupt), all reference bits are cleared
  - Another version: cleared if all blocks are referenced

# Not Recently Used (NRU)

- We cannot know the future, so we try to estimate according to the past
- Use the pages' **reference bits**
- Periodically (on clock interrupt), all reference bits are cleared
  - Another version: cleared if all blocks are referenced
- When need to choose a victim, choose randomly among the pages with zero reference bit
  - Logic: have not been accessed since last clearing

# Not Recently Used (NRU)

- We cannot know the future, so we try to estimate according to the past
- Use the pages' **reference bits**
- Periodically (on clock interrupt), all reference bits are cleared
  - Another version: cleared if all blocks are referenced
- When need to choose a victim, choose randomly among the pages with zero reference bit
  - Logic: have not been accessed since last clearing
- Can be implemented, but very crude



# Least Recently Used (LRU)

# Least Recently Used (LRU)

- Temporal locality → pages that were recently used are likely to be used again

# Least Recently Used (LRU)

- Temporal locality → pages that were recently used are likely to be used again
- So evict the one that was least recently used
  - Gives good approximation of working set
  - Same logic as NRU, but better accuracy

# Least Recently Used (LRU)

- Temporal locality → pages that were recently used are likely to be used again
- So evict the one that was least recently used
  - Gives good approximation of working set
  - Same logic as NRU, but better accuracy
- LRU implementation is difficult/expensive

# Least Recently Used (LRU)

- Temporal locality → pages that were recently used are likely to be used again
- So evict the one that was least recently used
  - Gives good approximation of working set
  - Same logic as NRU, but better accuracy
- LRU implementation is difficult/expensive
  - Timestamps? How many bits?
  - Must find minimal timestamp on each eviction


# Least Recently Used (LRU)

- Temporal locality → pages that were recently used are likely to be used again
- So evict the one that was least recently used
  - Gives good approximation of working set
  - Same logic as NRU, but better accuracy
- LRU implementation is difficult/expensive
  - Timestamps? How many bits?
  - Must find minimal timestamp on each eviction
  - Sorted list? Re-sort on every access?
  - List overhead:  $\log_2(n)$  bits / page

# The Clock Algorithm

a.k.a. Second Chance Algorithm


- Think of all the frames as a circular list
- With a hand pointing at one of them
- Each page has a reference bit
- When you need to evict a page:
  1. As long as the pointed page has been referenced, clear the bit and move on
  2. Evict the first non-referenced page found



# The Clock Algorithm

a.k.a. Second Chance Algorithm


- Think of all the frames as a circular list
- With a hand pointing at one of them
- Each page has a reference bit
- When you need to evict a page:
  1. As long as the pointed page has been referenced, clear the bit and move on
  2. Evict the first non-referenced page found




Give it a second chance




# The Clock Algorithm




# The Clock Algorithm




# The Clock Algorithm




# The Clock Algorithm




# The Clock Algorithm




# The Clock Algorithm



# The Clock Algorithm




# The Clock Algorithm





# The Clock Algorithm




# From References to milli-secs

- Clock algorithm is rather crude
- No real discrimination between pages with different activity patterns
- Improvement: try to track time since last reference
  - Not only that it existed
- Method: maintain crude timestamp per page


# Clock+Time Algorithm

- Maintain all pages of a process in a circular list
- Maintain virtual time variable
- Every time a page is referenced, set R=1
- Upon a clock tick:
  - For each page with R=1, set time of last use to virtual time
  - Advance virtual time, reset R for all pages




# Clock+Time Algorithm

- Upon a page fault, if eviction is needed




# Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow




# Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow
  2. If  $R=1$ , advance arrow, go to Line 1




# Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow
  2. If  $R=1$ , advance arrow, go to Line 1
  3. If  $R=0$  AND  $(\text{current virtual time}) - (\text{time of last use}) < k$  advance arrow, go to Line 1



# Clock+Time Algorithm


- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow
  2. If  $R=1$ , advance arrow, go to Line 1
  3. If  $R=0$  AND  $(\text{current virtual time}) - (\text{time of last use}) < k$  advance arrow, go to Line 1
  4. Otherwise, evict page





# Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow
  2. If  $R=1$ , advance arrow, go to Line 1
  3. If  $R=0$  AND  $(\text{current virtual time}) - (\text{time of last use}) < k$  advance arrow, go to Line 1
  4. Otherwise, evict page



# Evicting dirty and clean pages

# Evicting dirty and clean pages

- Dirty page: A page that was modified since the last time it was written to the disk

# Evicting dirty and clean pages

- Dirty page: A page that was modified since the last time it was written to the disk
  - The copy in the disk is different than the copy in the main memory (write back)

# Evicting dirty and clean pages

- Dirty page: A page that was modified since the last time it was written to the disk
  - The copy in the disk is different than the copy in the main memory (write back)
  - Evicting dirty page → waiting to write the page to the disk and then waiting for the new page to the memory

# Evicting dirty and clean pages


- Dirty page: A page that was modified since the last time it was written to the disk
  - The copy in the disk is different than the copy in the main memory (write back)
  - Evicting dirty page → waiting to write the page to the disk and then waiting for the new page to the memory
  - Evicting clean page → just waiting for the new page to the memory (safe to override)

# Evicting dirty and clean pages

- Dirty page: A page that was modified since the last time it was written to the disk
  - The copy in the disk is different than the copy in the main memory (write back)
  - Evicting dirty page → waiting to write the page to the disk and then waiting for the new page to the memory
  - Evicting clean page → just waiting for the new page to the memory (safe to override)
- Prefer to replace clean pages

# Improved Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  1. Look at the page with the arrow
  2. If  $R=1$ , advance arrow, go to Line 1
  3. If  $R=0$  AND  $(\text{current virtual time}) - (\text{time of last use}) < k$  advance arrow, go to Line 1
  4. If dirty, advance arrow, go to Line 1
  5. Otherwise, evict page






# Improved Clock+Time Algorithm

- Upon a page fault, if eviction is needed
  - Look at the page with the arrow
  - If  $R=1$ , advance arrow, go to Line 4

- If  $R=0$ , (current page is clean)
  - If no candidate is found:  
Evict the oldest page (clean or dirty)


4. If dirty, advance arrow, go to Line 1

5. Otherwise, evict page



# Improved Clock+Time Algorithm

- Upon a page fault, if eviction is required:
  - If dirty but old (more than k):**
    1. Look at schedule for eviction; the page is written to the disk in parallel to the process actions
    2. If  $R=1$ , parallel to the process actions to Line 1 (using DMA)
    3. If  $R=0$  AND  $(\text{current virtual time} - (\text{time of last use})) < k$   
advance arrow, go to Line 1
    4. **If dirty, advance arrow, go to Line 1**
    5. Otherwise, evict page



# Global vs. Local Paging

- When process P1 causes a page fault, should the OS choose a victim from one of P1 pages or from any process?
- Some operating systems have local paging (only from P1 pages) and some global (“best” candidate to evict)