

Scheduling

1

**OPERATING SYSTEMS COURSE
THE HEBREW UNIVERSITY
SPRING 2023**

Outline

2

- **CPU Scheduling**

- First come first serve (FCFS)
- Shortest time first (SJF)
- Shortest remaining time first (SRTF)
- Priority scheduling (PS)
- Round robin (RR)
- Multi-level queue
- Multi-level feedback queue

- **Parallel System Scheduler**

Scheduling Criteria

3

- CPU utilization (max)
 - keep the CPU as busy as possible
- Throughput (max)
 - number of processes that complete their execution per time unit
- Waiting time (min)
 - amount of time a process has been waiting in the ready queue
- Turnaround time (min)
 - amount of time to execute a particular process

Scheduling Criteria

4

First-Come First-Served Scheduling

5

First-Come First-Served Scheduling

6

- The process that requests the CPU first is allocated the CPU first.
 - The **simplest** CPU-scheduling algorithm
 - **Implementation** with a FIFO queue
 - Add to the tail of the queue
 - Remove from the head to the queue
 - The **code** is simple to write and understand

First-Come First-Served Scheduling

7

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
|----------------|-------------------|

| | |
|-------|----|
| P_1 | 10 |
|-------|----|

| | |
|-------|---|
| P_2 | 1 |
|-------|---|

| | |
|-------|---|
| P_3 | 1 |
|-------|---|

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The **Gantt Chart** for the schedule is:

- Waiting time for $P_1 = 0$; $P_2 = 10 + cs$; $P_3 = 11 + 2cs$
- Average waiting time: $(0 + 10 + 11 + 3cs)/3 = 7 + cs$

First-Come First-Served Scheduling

7

| Process | Burst Time |
|---------|------------|
| P_1 | 10 |
| P_2 | 1 |
| P_3 | 1 |

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The **Gantt Chart** for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 10 + cs$; $P_3 = 11 + 2cs$
- Average waiting time: $(0 + 10 + 11 + 3cs)/3 = 7 + cs$

First-Come First-Served Scheduling

8

- Processes:

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :

- Performance:

| Metric (Avg) | FCFS |
|-------------------------|---|
| <i>CPU Utilization</i> | $26/(26+3\mathbf{cs})$ |
| <i>Turn around time</i> | $((8)+(12+cs)+(21+2cs)+(26+3cs))/4 = 16.75 + 1.5cs$ |
| <i>Waiting</i> | $((0)+(8+cs)+(12+2cs)+(21+3\mathbf{cs}))/4 = 10.25 + 1.5cs$ |
| <i>Throughput</i> | $4/(26 + 3\mathbf{cs})$ |


First-Come First-Served Scheduling

8

- Processes:

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :



- Performance:

| Metric (Avg) | FCFS |
|-------------------------|---|
| <i>CPU Utilization</i> | $26/(26+3cs)$ |
| <i>Turn around time</i> | $((8)+(12+cs)+(21+2cs)+(26+3cs))/4 = 16.75 + 1.5cs$ |
| <i>Waiting</i> | $((0)+(8+cs)+(12+2cs)+(21+3cs))/4 = 10.25 + 1.5cs$ |
| <i>Throughput</i> | $4/(26 + 3cs)$ |

First-Come First-Served Scheduling

9

- Fair scheduler
- Average waiting time is often quite long.
 - The problem: short processes wait after long processes
 - Provide good performance when the variance between the jobs is small
- FCFS is non-preemptive
 - Once the CPU has been allocated to a process, that process keeps the CPU until it releases the CPU.

Shortest Job First (SJF) Scheduling

10

Shortest-Job-First Scheduling

11

- When the CPU is available, it is assigned to the process that has the **smallest next CPU burst**.
 - Another term – **Shortest Job Next (SJN)**
 - FCFS scheduling is used to break the tie
- **Non-preemptive** scheduler
- Provably optimal
 - **Minimum average waiting time** for a given set of processes [if all the jobs arrive at the same time, “offline” scheduling]
- Problems:
 - Jobs execution time is often **unknown**
 - **Fairness** (including starvation)

Shortest-Job-First Scheduling

12

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
|----------------|-------------------|

| | |
|-------|---|
| P_1 | 8 |
|-------|---|

| | |
|-------|---|
| P_2 | 4 |
|-------|---|

| | |
|-------|---|
| P_3 | 9 |
|-------|---|

| | |
|-------|---|
| P_4 | 5 |
|-------|---|

The Gantt Chart for the schedule is:

- Waiting time for $P_2 = 0$; $P_4 = 4$; $P_1 = 9$; $P_3 = 17$;
- Average waiting time: $(0 + 4 + 9 + 17) / 4 = 7.5$
- Using FCFS scheme : $(0 + 8 + 12 + 21) / 4 = 10.25$

Ignoring
context
switch
cost

Shortest-Job-First Scheduling

12

| Process | Burst Time |
|---------|------------|
|---------|------------|


| | |
|-------|---|
| P_1 | 8 |
|-------|---|

| | |
|-------|---|
| P_2 | 4 |
|-------|---|

| | |
|-------|---|
| P_3 | 9 |
|-------|---|

| | |
|-------|---|
| P_4 | 5 |
|-------|---|

The Gantt Chart for the schedule is:



- Waiting time for $P_2 = 0$; $P_4 = 4$; $P_1 = 9$; $P_3 = 17$;
- Average waiting time: $(0 + 4 + 9 + 17) / 4 = 7.5$
- Using FCFS scheme : $(0 + 8 + 12 + 21) / 4 = 10.25$

Ignoring
context
switch
cost

Shortest-Job-First Scheduling

13

- Processes:

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart:

- Performance:

| Metric (Avg) | SJF |
|------------------|---|
| Utilization | $26/(26+3CS)$ |
| Turn around time | $(4+9+CS+17+2CS+26+3CS)/4 = 14 + 1.5cs$ |
| Waiting | $(0+4+CS+9+2CS+17+3CS)/4 = 7.5 + 1.5cs$ |
| Throughput | $4/(26 + 3CS)$ |


Shortest-Job-First Scheduling

13

- Processes:

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart:



- Performance:

| Metric (Avg) | SJF |
|------------------|---|
| Utilization | $26/(26+3CS)$ |
| Turn around time | $(4+9+CS+17+2CS+26+3CS)/4 = 14 + 1.5cs$ |
| Waiting | $(0+4+CS+9+2CS+17+3CS)/4 = 7.5 + 1.5cs$ |
| Throughput | $4/(26 + 3CS)$ |

Shortest Remaining Time First (SRTF)

14

Shortest Remaining Time First (SRTF)

15

- Often called Shortest Remaining Time.
- A **preemptive** variant of SJF
 - If a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
 - Short processes are handled very quickly
- Same problems as SJF.
 - Jobs execution time is often **unknown**
 - **Fairness** (including starvation)

Example of SRTF

16

| <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|----------------|---------------------|-------------------|
| P_1 | 0 | 8 |
| P_2 | 1 | 4 |
| P_3 | 2 | 9 |
| P_4 | 3 | 5 |

- Gantt Chart:

Ignoring
context
switch
cost


- Average waiting time = $((10-1) + (1-1) + (17-2) + (5-3)) / 4 = 6.5$

Example of SRTF

16

| <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|----------------|---------------------|-------------------|
| P_1 | 0 | 8 |
| P_2 | 1 | 4 |
| P_3 | 2 | 9 |
| P_4 | 3 | 5 |

- Gantt Chart:



Ignoring
context
switch
cost

- Average waiting time = $((10-1) + (1-1) + (17-2) + (5-3)) / 4 = 6.5$

Priority Scheduling

17

Priority Scheduling

18

- The CPU is **allocated** to the process with the **highest priority**.
 - A priority is associated with each **process**
 - Fixed range of number, such as 0 to 7
 - We use **low** numbers to represent high priority
 - Equal-priority processes are scheduled in **FCFS** order
 - **SJF** is a special case of the general priority-scheduling algorithm
 - The priority is the predicted next CPU burst

Example of Priority Scheduling

19

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| P_1 | 10 | 3 |
| P_2 | 1 | 1 |
| P_3 | 2 | 4 |
| P_4 | 1 | 5 |
| P_5 | 5 | 2 |

- Gantt Chart:

- Average waiting time = $(6 + 0 + 16 + 18 + 1) / 5 = 8.2$

Ignoring
context
switch
cost

Example of Priority Scheduling

19

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| P_1 | 10 | 3 |
| P_2 | 1 | 1 |
| P_3 | 2 | 4 |
| P_4 | 1 | 5 |
| P_5 | 5 | 2 |

- Gantt Chart:



- Average waiting time = $(6 + 0 + 16 + 18 + 1) / 5 = 8.2$

Ignoring
context
switch
cost

Priority Scheduling

20

- Priorities can be defined
 - Internally
 - Use some measurable quantity
 - Time limits, memory requirements...
 - Externally
 - Set by criteria external to OS
 - importance, political factors
- Priority scheduling can be
 - Preemptive
 - Non-preemptive
- Major problem
 - Indefinite blocking or starvation
 - Solution: aging
 - Gradually increase the priority of processes that wait for a long time

Round-Robin Scheduling

21

Round-Robin Scheduling

22

- A small unit of time, called a **time quantum** is defined.
 - Generally from 10 to 100 milliseconds
- The ready queue is treated as a **circular**, FIFO queue.
- The CPU scheduler goes around the ready queue
 - Allocate the CPU to each process for a time interval of up to 1 time quantum.
- Two cases
 - CPU burst less than 1 time quantum
 - The process released the CPU voluntarily
 - CPU burst longer than 1 time quantum
 - Context switch will be executed

Example of Round-Robin Scheduling

23

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
|----------------|-------------------|

| | |
|-------|----|
| P_1 | 24 |
|-------|----|

| | |
|-------|---|
| P_2 | 3 |
|-------|---|

| | |
|-------|---|
| P_3 | 3 |
|-------|---|

- Round-robin scheduling
 - A time-quantum of 4 milliseconds
- Average waiting time = $(6 + 4 + 7) / 3 = 5.66$
- RR scheduling is preemptive.

Example of Round-Robin Scheduling

23


| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
|----------------|-------------------|

| | |
|-------|----|
| P_1 | 24 |
|-------|----|

| | |
|-------|---|
| P_2 | 3 |
|-------|---|

| | |
|-------|---|
| P_3 | 3 |
|-------|---|

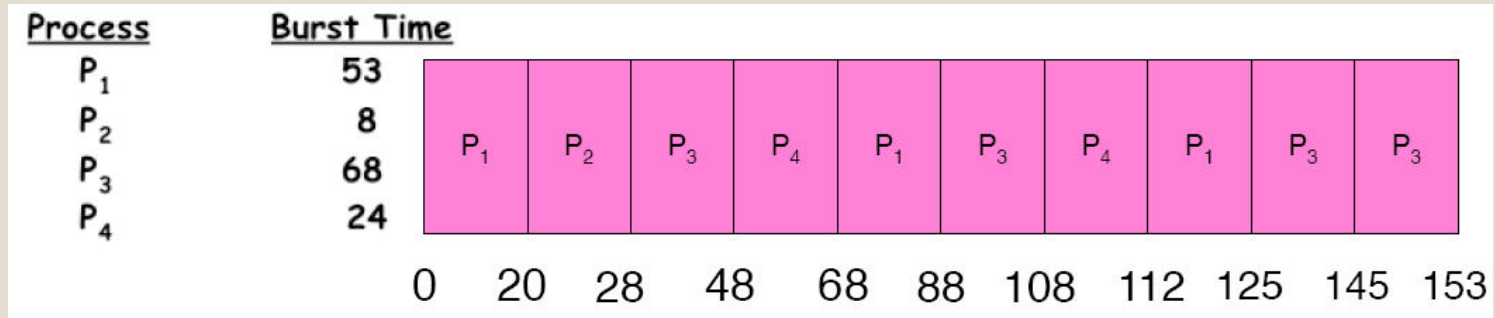
- Round-robin scheduling
 - A time-quantum of 4 milliseconds



- Average waiting time = $(6 + 4 + 7) / 3 = 5.66$
- RR scheduling is **preemptive**.

Example of RR with Time Quantum = 20

24



- Waiting Time:
 - P1: $0 + (68 - 20) + (112 - 88) = 72$
 - P2: $(20 - 0) = 20$
 - P3: $(28 - 0) + (88 - 48) + (125 - 108) = 85$
 - P4: $(48 - 0) + (108 - 68) = 88$
- Completion Time (Turn Around Time):
 - P1: 125
 - P2: 28
 - P3: 153
 - P4: 112
- Average Waiting Time: $(72 + 20 + 85 + 88) / 4 = 66.25$
- Average Completion Time: $(125 + 28 + 153 + 112) / 4 = 104.5$

Example of Round-Robin Scheduling

(25)

- Processes:
- Quantum=3

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :

- Performance:

| Metric | RR |
|-----------------------------|---|
| <i>Utilization</i> | $26/(26+9CS)$ |
| <i>Avg Turn around time</i> | $(23+16+26+21+29cs)/4 = 21.5$ ignoring CS |
| <i>Avg Waiting</i> | $(15+12+17+16+29cs)/4 = 15$ ignoring CS |
| <i>Throughput</i> | $4/(26 + 9CS)$ |


Example of Round-Robin Scheduling

(25)

- Processes:
- Quantum=3

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :



- Performance:

| Metric | RR |
|----------------------|---|
| Utilization | $26/(26+9CS)$ |
| Avg Turn around time | $(23+16+26+21+29cs)/4 = 21.5$ ignoring CS |
| Avg Waiting | $(15+12+17+16+29cs)/4 = 15$ ignoring CS |
| Throughput | $4/(26 + 9CS)$ |

Example of Round-Robin Scheduling

(26)

- Processes:
- Quantum=6

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :


Example of Round-Robin Scheduling

(26)

- Processes:
- Quantum=6

| Process | Burst Time |
|---------|------------|
| A | 8 |
| B | 4 |
| C | 9 |
| D | 5 |

- Gantt Chart :



Comparing FCFS and RR

27

- Assuming zero-cost context switching time, is RR always better than FCFS?
- Assume 10 jobs, all start at the same time, and each require 100 seconds of CPU time
- RR scheduler quantum of 1 second
- Completion Times (CT)
 - Both FCFS and RR finish at the same time
 - But average response time is much worse under RR!

| Job # | FCFS CT | RR CT |
|-------|---------|-------|
| 1 | 100 | 991 |
| 2 | 200 | 992 |
| ... | ... | ... |
| 9 | 900 | 999 |
| 10 | 1000 | 1000 |

Bad when all jobs are same length


RR performance

28

- Run time distribution:
 - RR is **poor** if the jobs **variance** is **small**
 - RR is good for “real life”
- Context switch **hurts** the performance!
 - Context switch's **time**
 - **Cache** state must be shared between all jobs with RR
 - Total time for RR longer even for zero-cost context switch!

Comparing FCFS and RR

29



| | Quantum | P ₁ | P ₂ | P ₃ | P ₄ | Average |
|-----------------|------------|----------------|----------------|----------------|----------------|------------------|
| Wait Time | Best FCFS | 32 | 0 | 85 | 8 | $31\frac{1}{4}$ |
| | Q = 1 | 84 | 22 | 85 | 57 | 62 |
| | Q = 5 | 82 | 20 | 85 | 58 | $61\frac{1}{4}$ |
| | Q = 8 | 80 | 8 | 85 | 56 | $57\frac{1}{4}$ |
| | Q = 10 | 82 | 10 | 85 | 68 | $61\frac{1}{4}$ |
| | Q = 20 | 72 | 20 | 85 | 88 | $66\frac{1}{4}$ |
| | Worst FCFS | 68 | 145 | 0 | 121 | $83\frac{1}{2}$ |
| Completion Time | Best FCFS | 85 | 8 | 153 | 32 | $69\frac{1}{2}$ |
| | Q = 1 | 137 | 30 | 153 | 81 | $100\frac{1}{2}$ |
| | Q = 5 | 135 | 28 | 153 | 82 | $99\frac{1}{2}$ |
| | Q = 8 | 133 | 16 | 153 | 80 | $95\frac{1}{2}$ |
| | Q = 10 | 135 | 18 | 153 | 92 | $99\frac{1}{2}$ |
| | Q = 20 | 125 | 28 | 153 | 112 | $104\frac{1}{2}$ |
| | Worst FCFS | 121 | 153 | 68 | 145 | $121\frac{3}{4}$ |

Multilevel Queue Scheduling

30

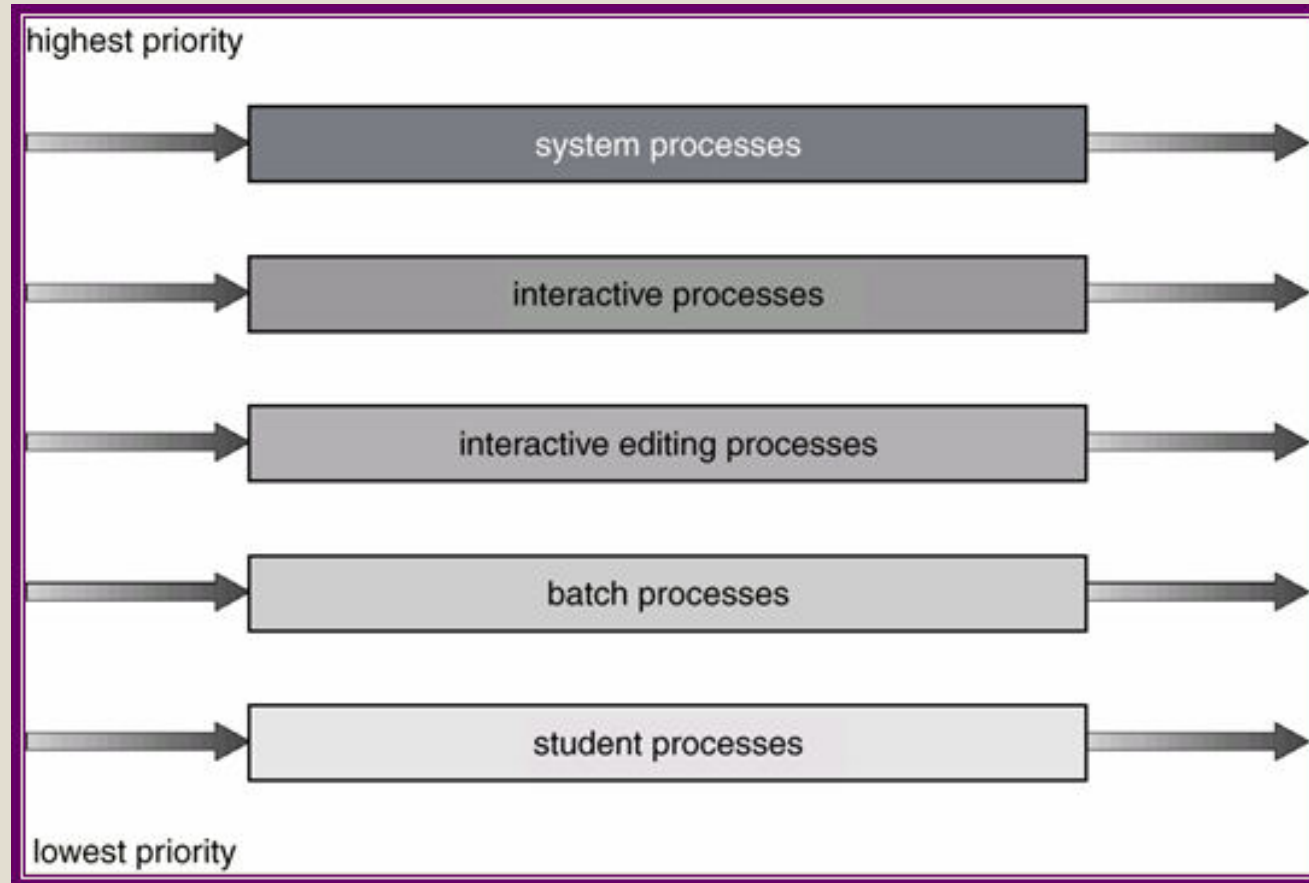
Multilevel Queue Scheduling

31

- Multilevel queue-scheduling algorithm
 - Partition the ready queue into several **separate groups**
 - Each group has its own scheduling **algorithm**
 - Scheduling among the queues
 - Fixed-priority preemptive scheduling
 - Possibility of starvation
 - Time-slice between the queues
 - A certain portion of the CPU time
- Processes are classified into different groups
 - Foreground (or interactive) processes
 - Background (or batch) processes

Multilevel Queue Scheduling

32



Multilevel Feedback Queue

33

- A **process** can **move** between the various queues
- For example:
 - Use too much CPU time: move to a lower-priority queue
 - Wait too long: move to a higher-priority queue
 - This form of aging prevents starvation
- Multilevel-feedback-queue scheduler defined by:
 - Number of queues
 - Scheduling algorithms for each queue
 - Method used to determine which queue a process will enter when that process needs service
 - Method used to determine when to upgrade a process
 - Method used to determine when to demote a process

Example of Multilevel Feedback Queue


34

- Three queues:
 - Q_0 – time quantum 8 milliseconds
 - Q_1 – time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel feedback queues

35

- The most general scheme



Outline

36

- CPU Scheduling
- **Parallel Systems Scheduling**
 - Schedulers
 - Performance Evaluation

Parallel System

37




Parallel System

37



Parallel System

37



Supercomputers

38

- Provide extremely high speed of calculation (HPC)
- Each computer has a basic OS.
- The computers are connected with high speed network
- Regularly run parallel jobs
- Each job contain the required number of processors
- No preemption.

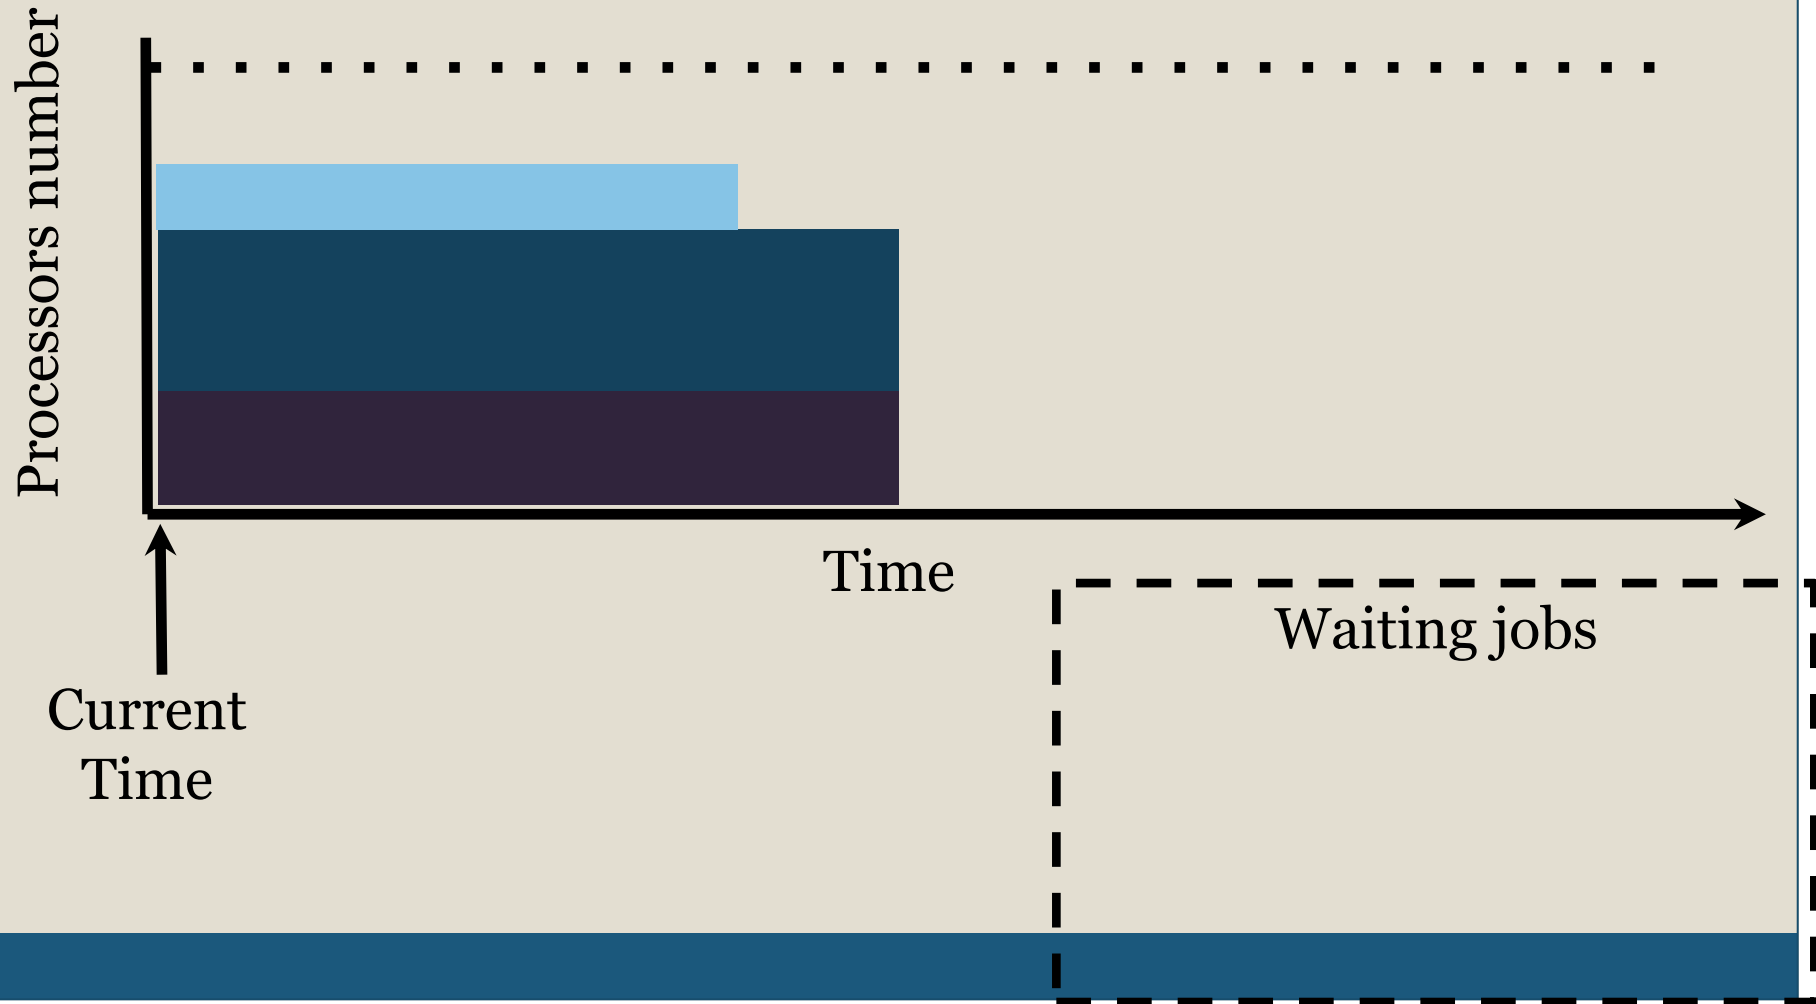
Supercomputers' Scheduler

39

- The scheduler's purpose is to determine which **processors** will be **allocated** to each **job**.
- Example of schedulers: FCFS, SJF, Backfilling schedulers.

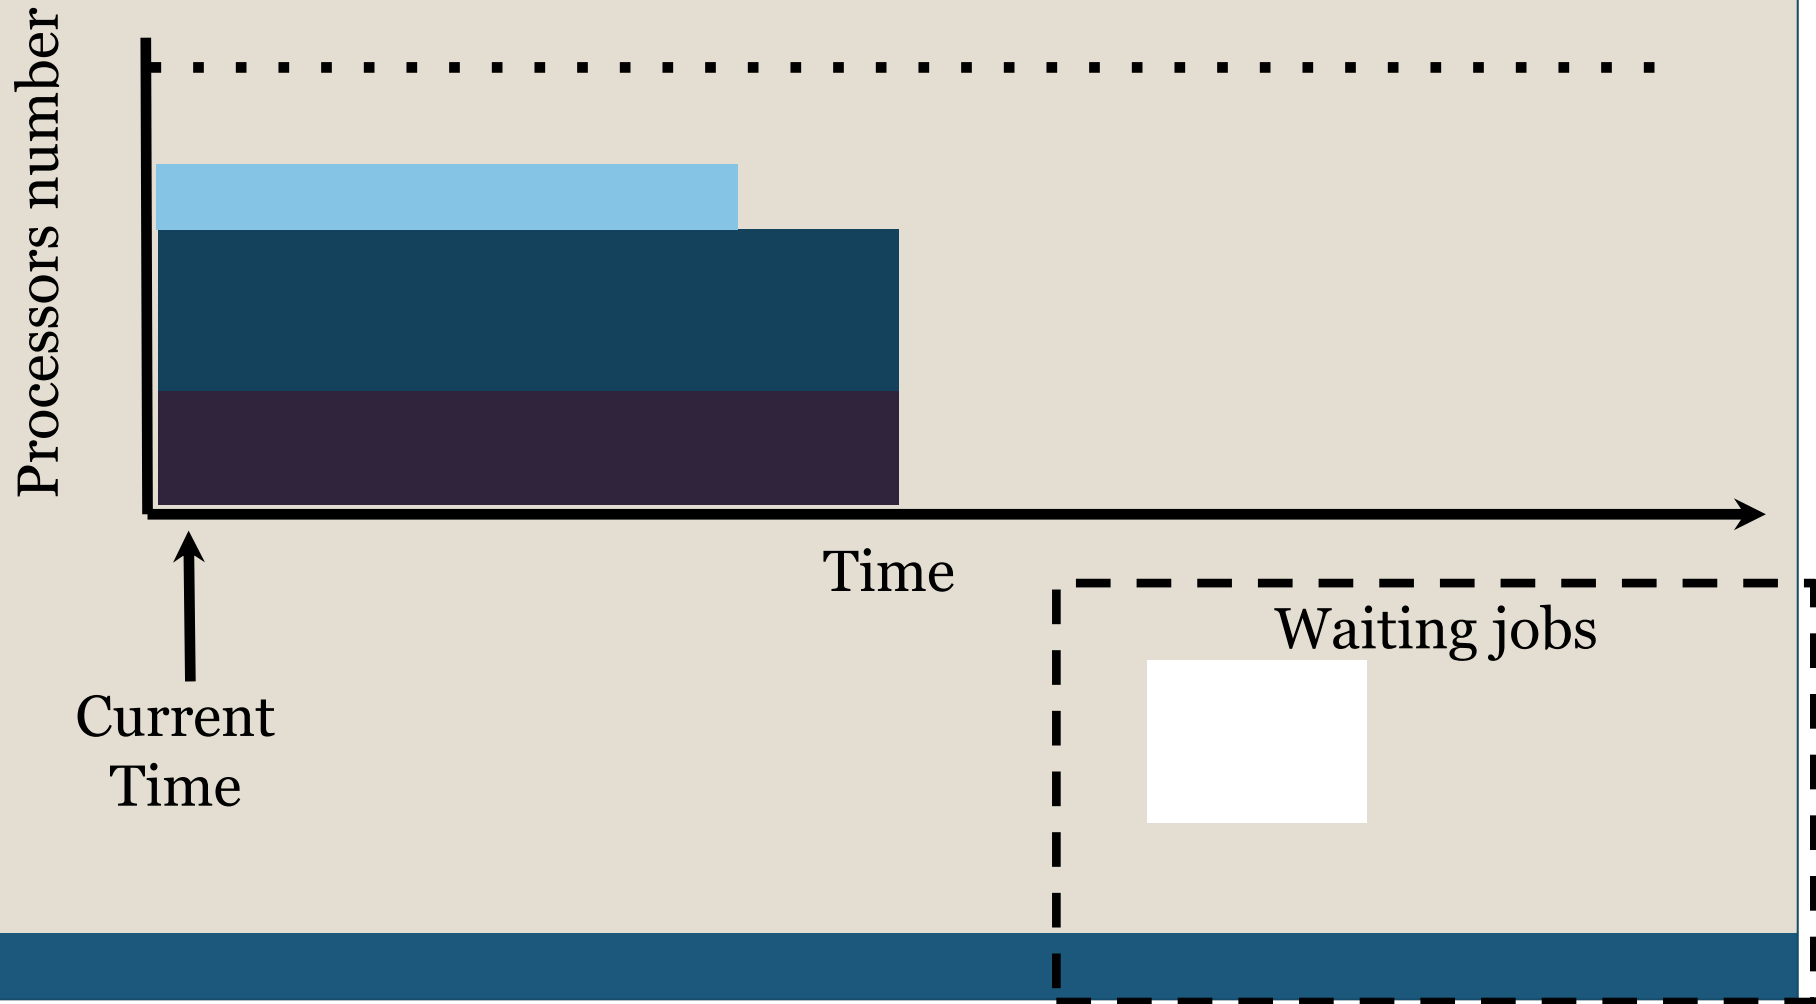
FCFS Scheduler

40



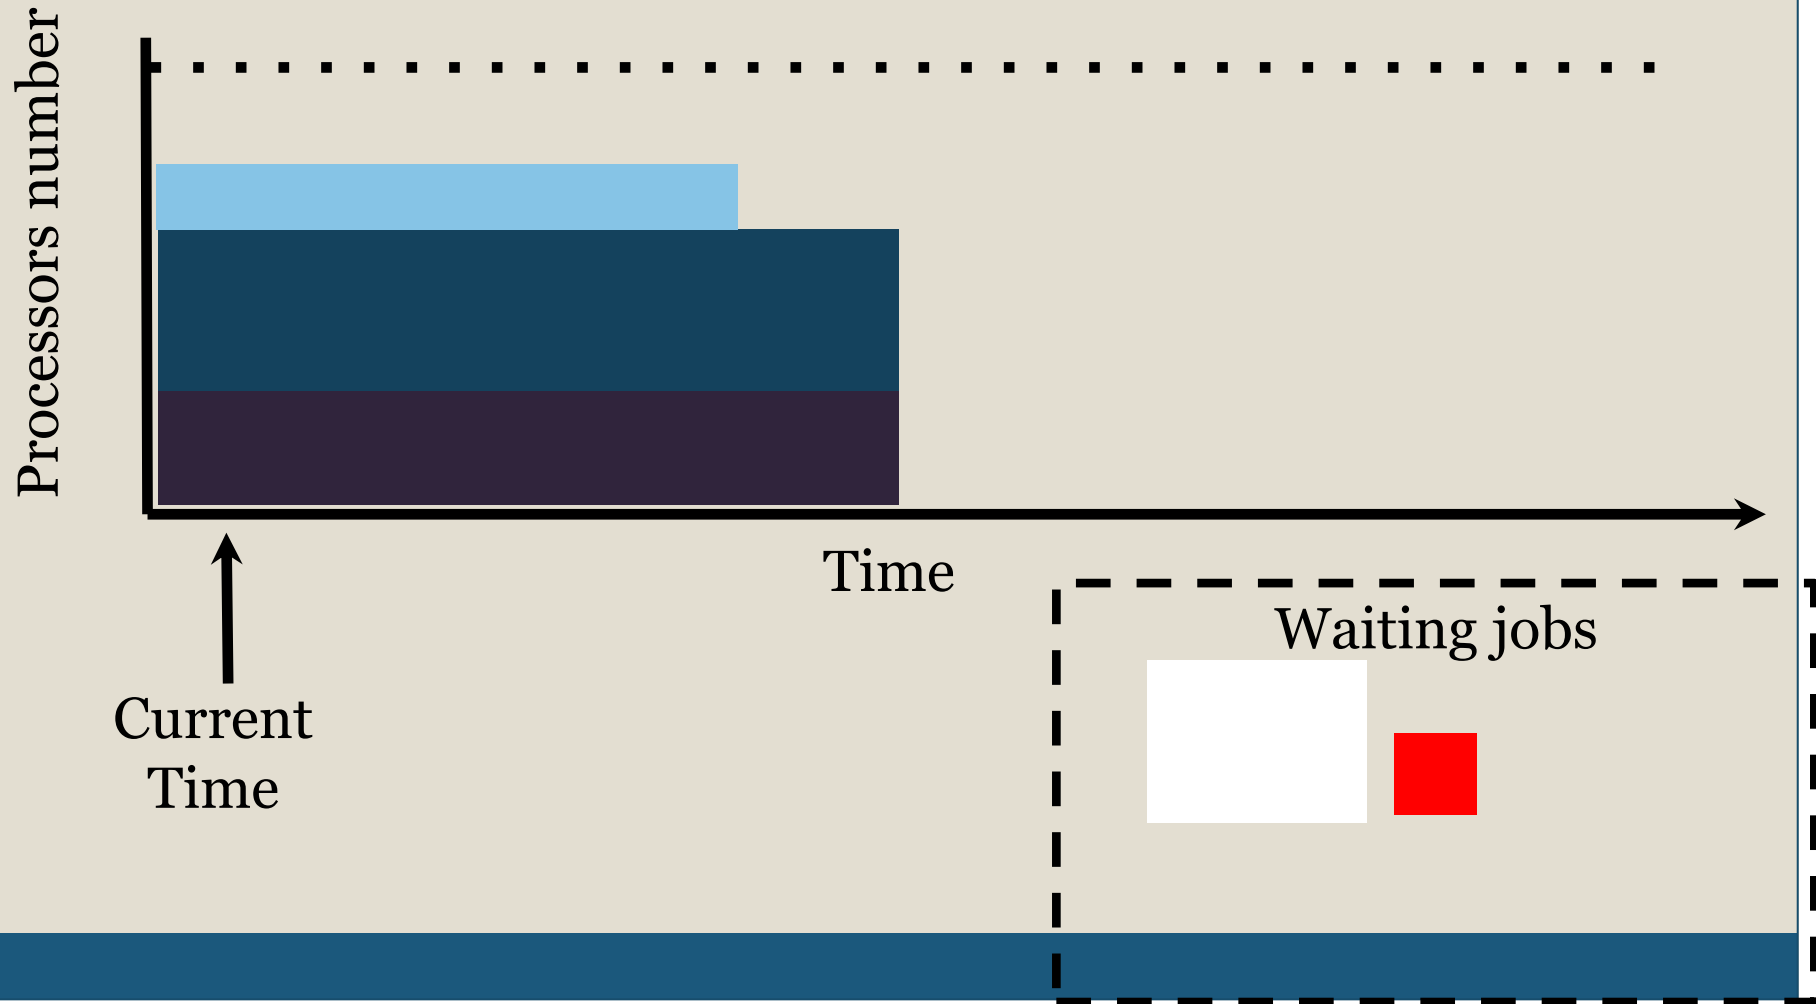
FCFS Scheduler

40



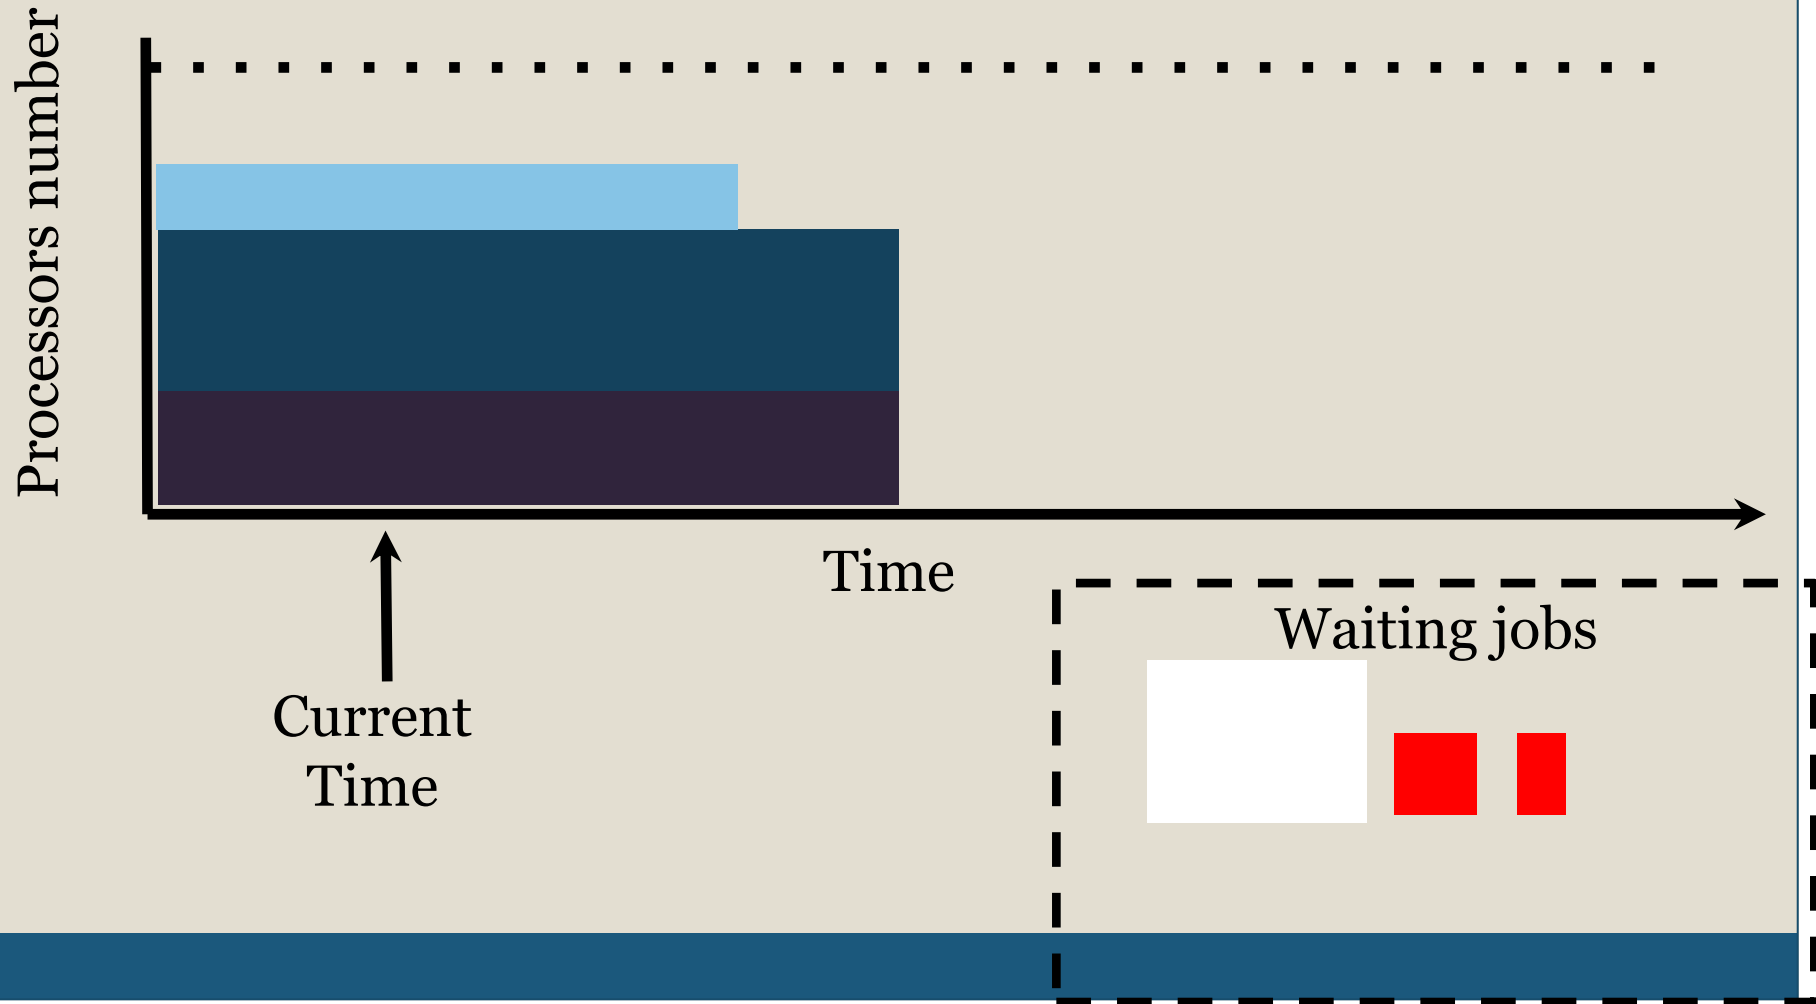
FCFS Scheduler

40



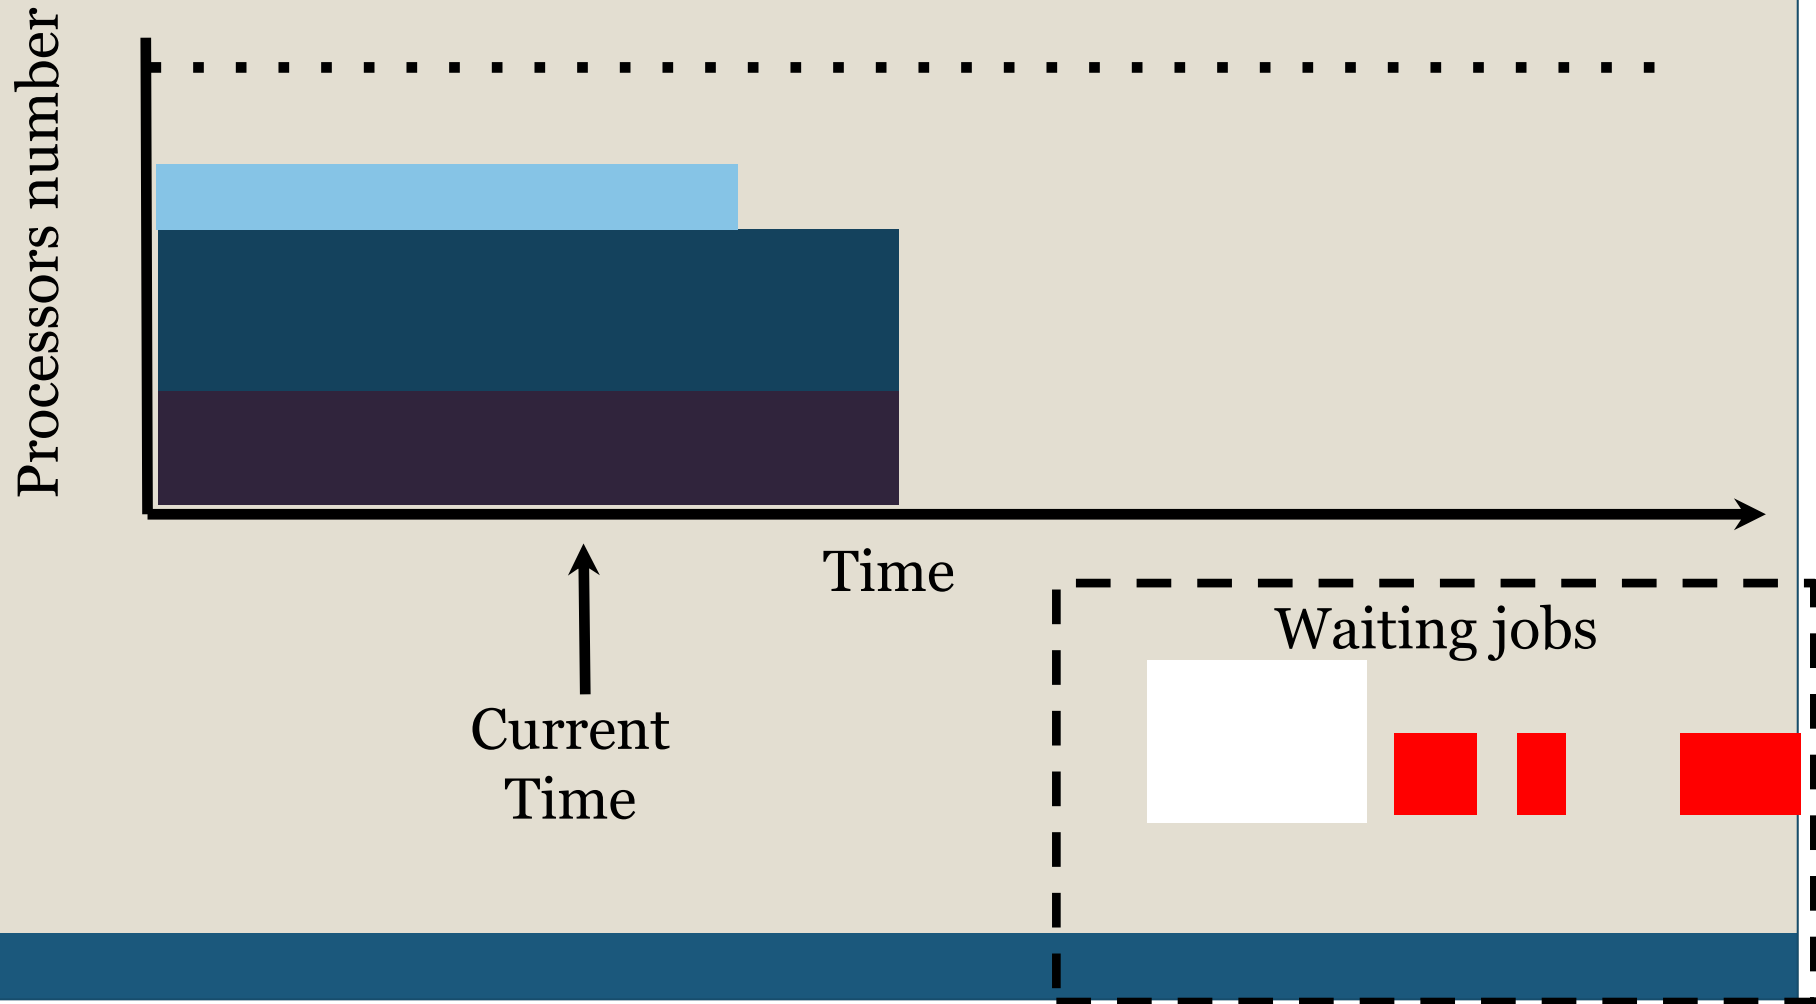
FCFS Scheduler

40



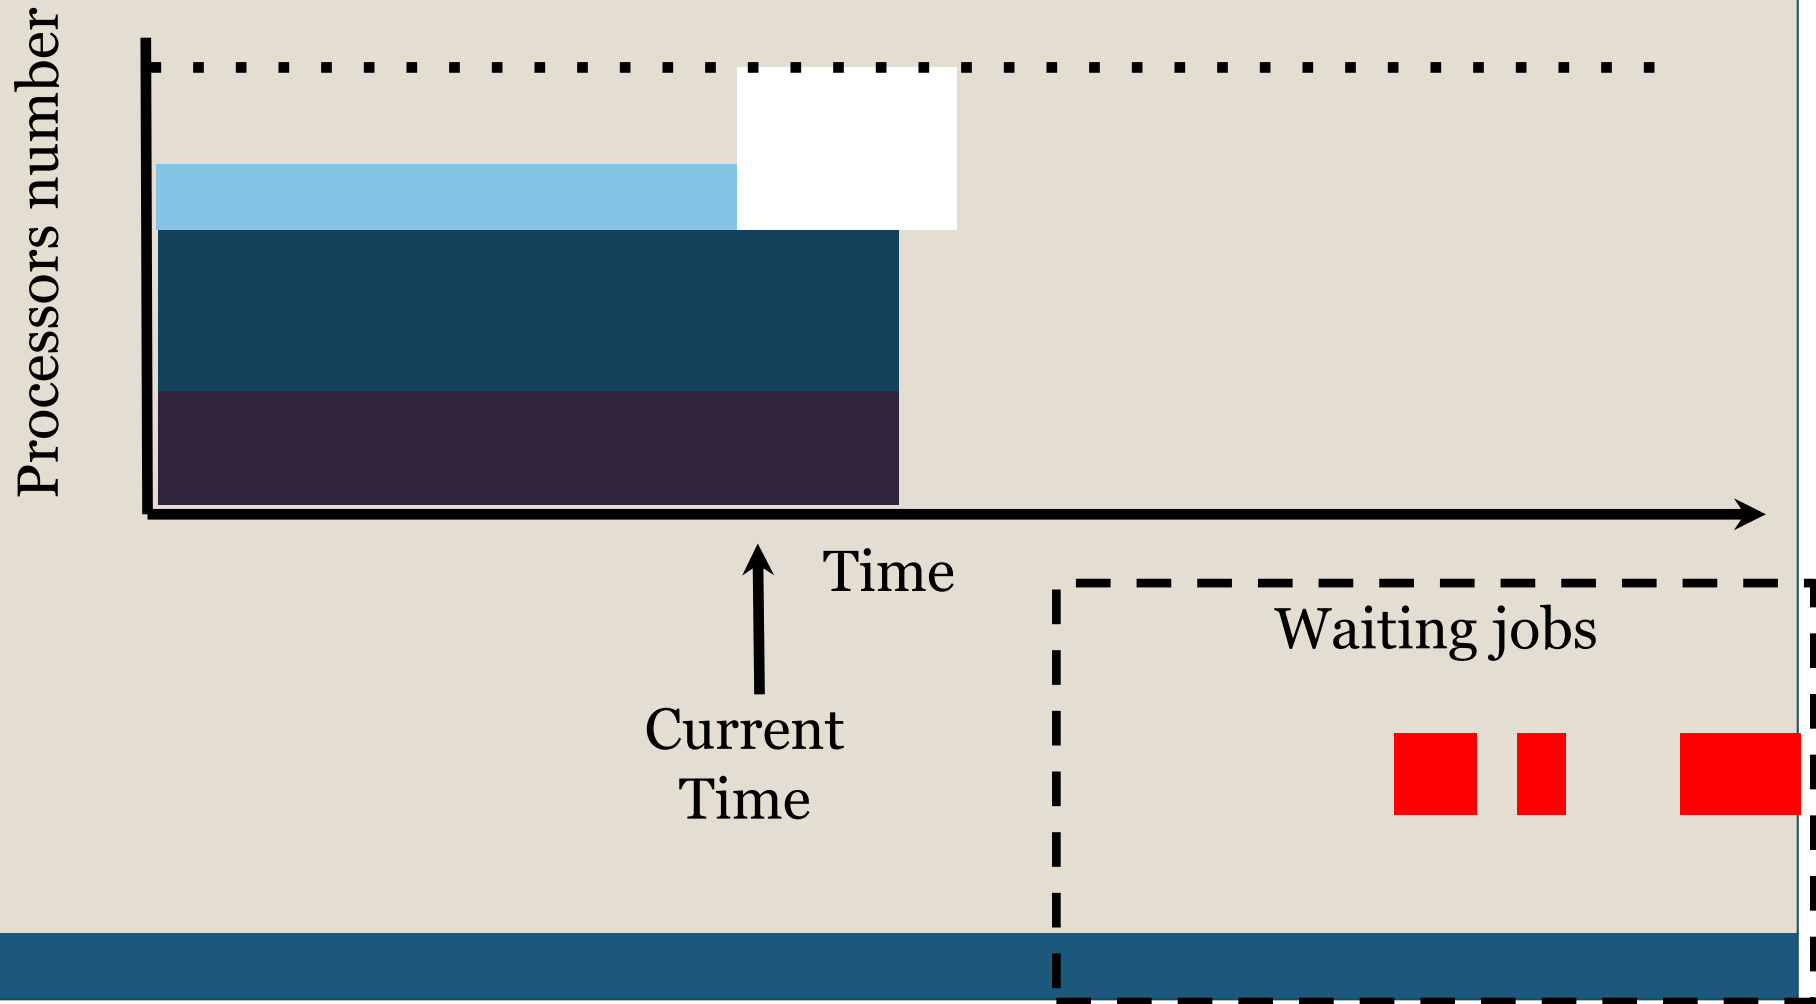
FCFS Scheduler

40




FCFS Scheduler

40



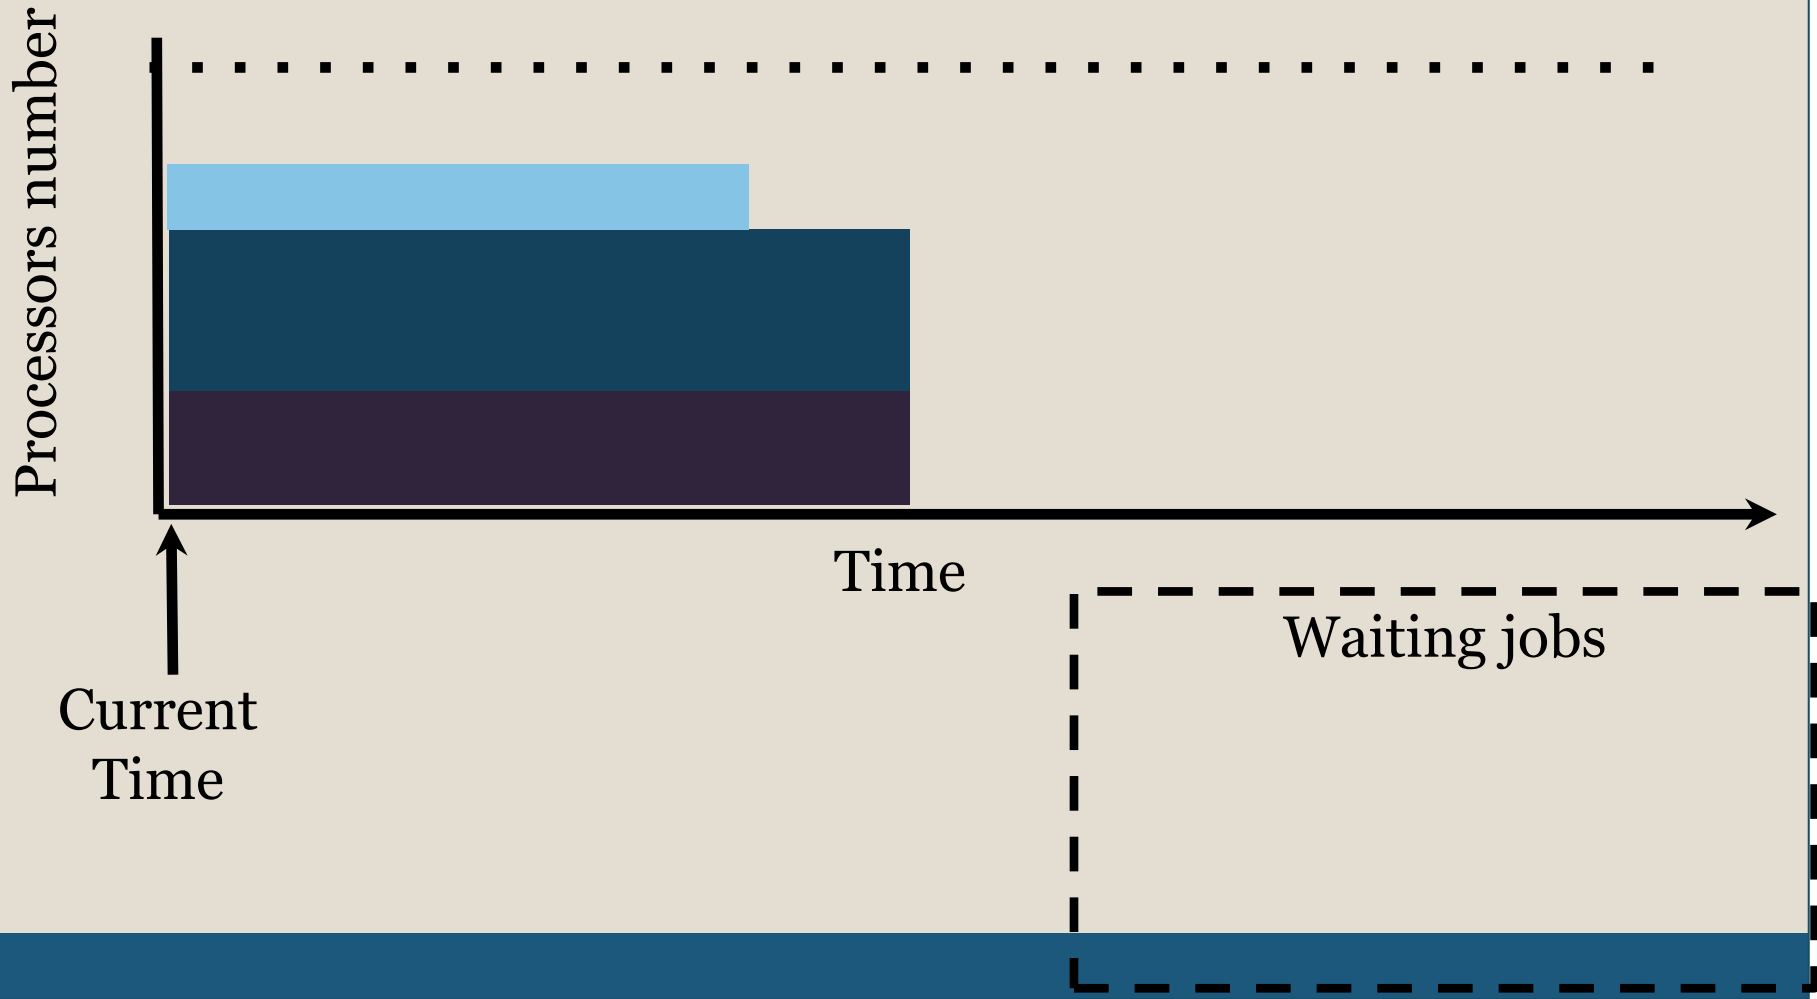
FCFS Scheduler

40




Scheduler with Backfilling

41




Scheduler with Backfilling

41



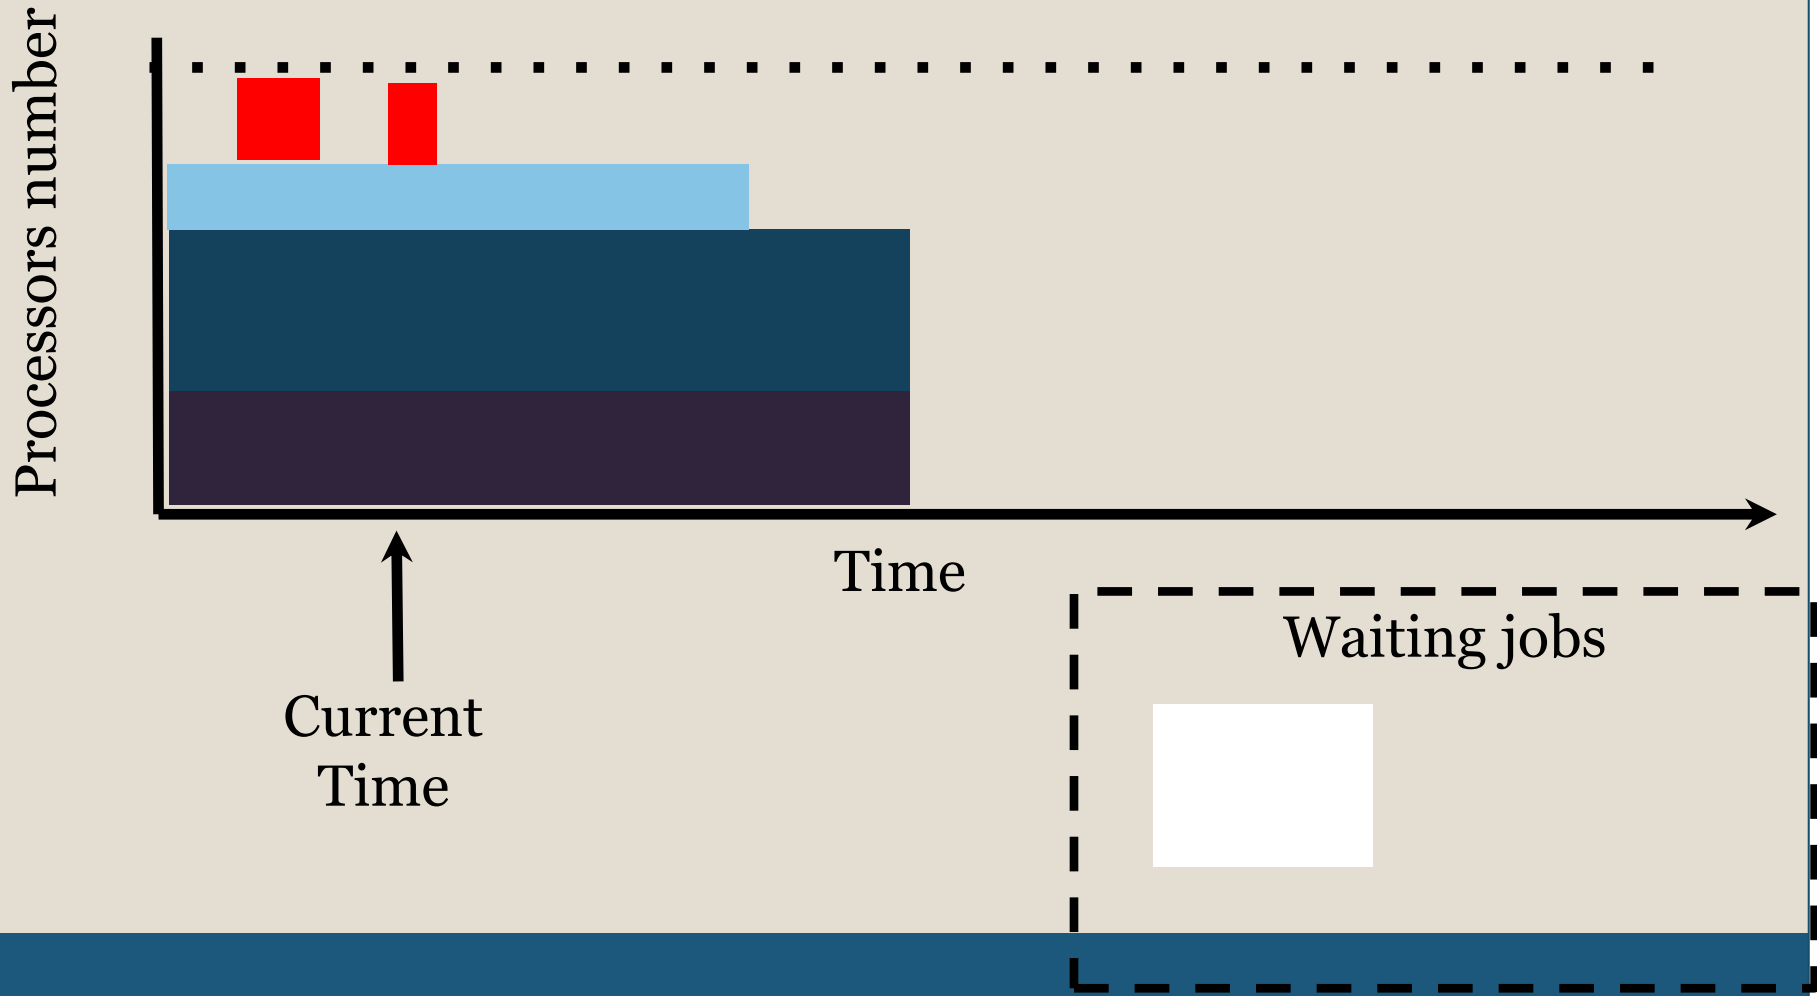
Scheduler with Backfilling

41




Scheduler with Backfilling

41



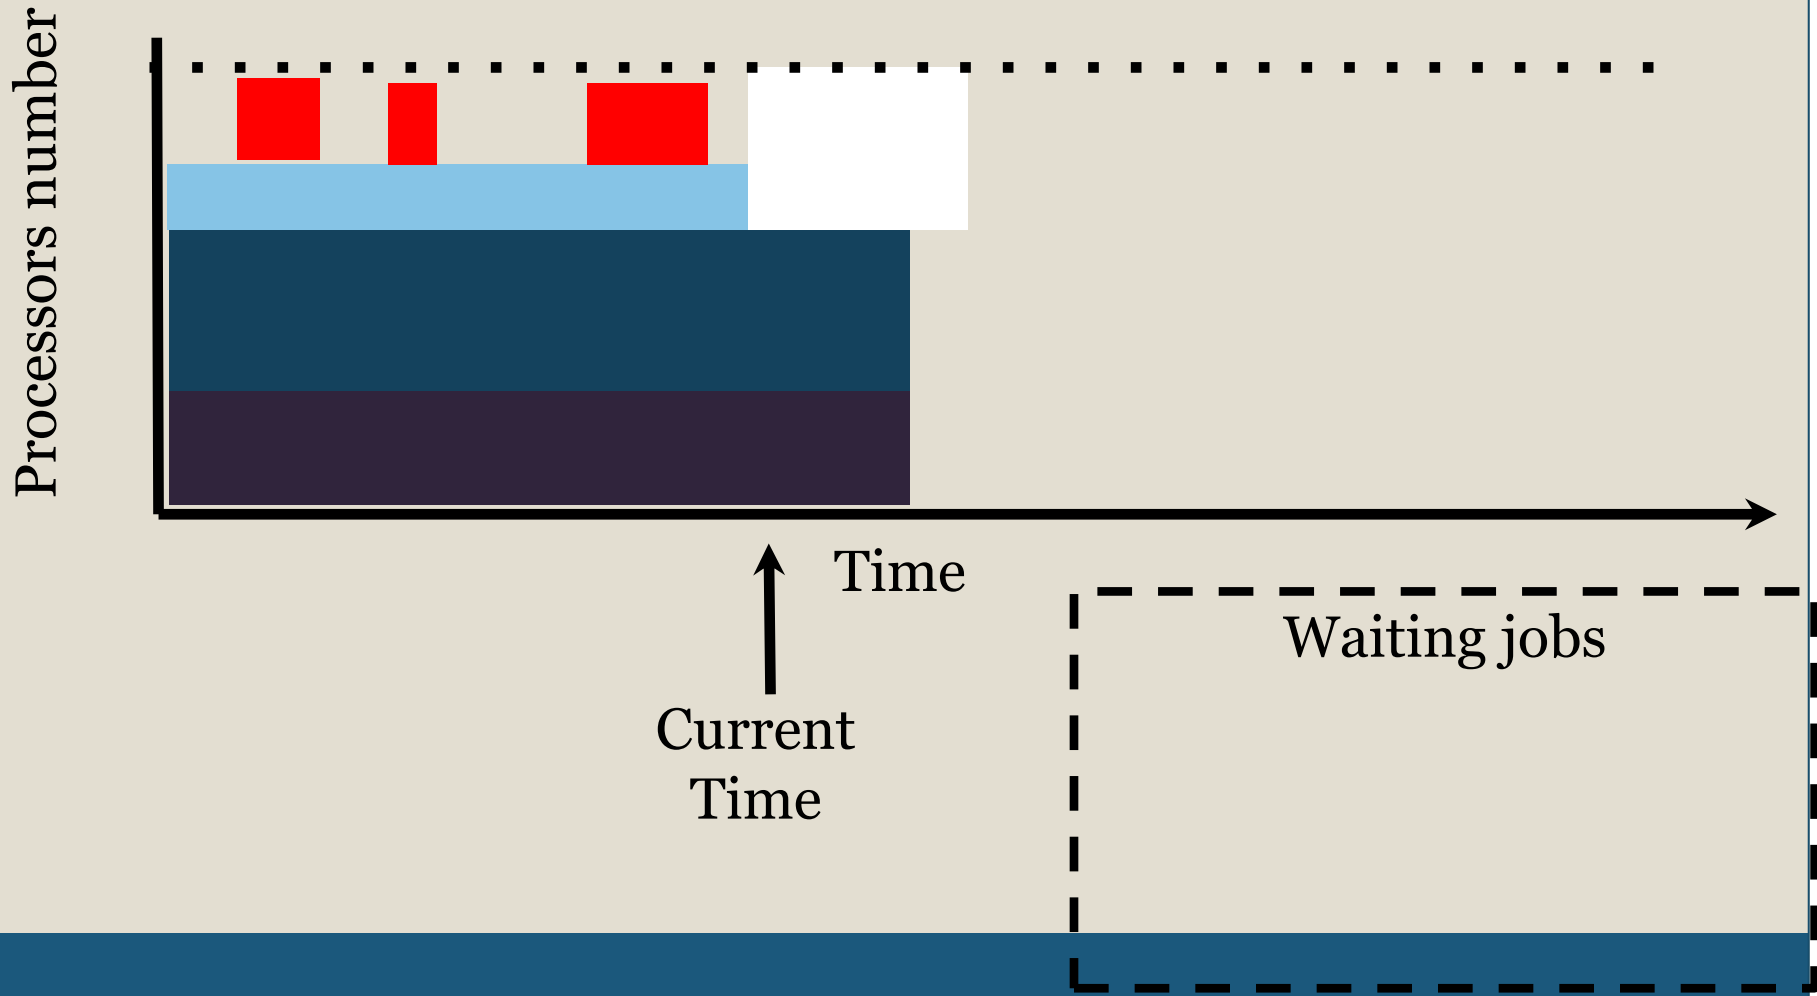
Scheduler with Backfilling

41




Scheduler with Backfilling

41



Scheduler with Backfilling

41



The EASY Scheduler


42

- Very common and simple algorithm
- Relatively fair
- Uses backfilling based on arrival time.

EASY Data Structures

43


- List of running jobs
 - Number of processors they use
 - Expected termination
- List of queued jobs
 - How many processors they need
 - How long they are expected to run
 - Sorted in order of arrival



EASY Operation

44


1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

44

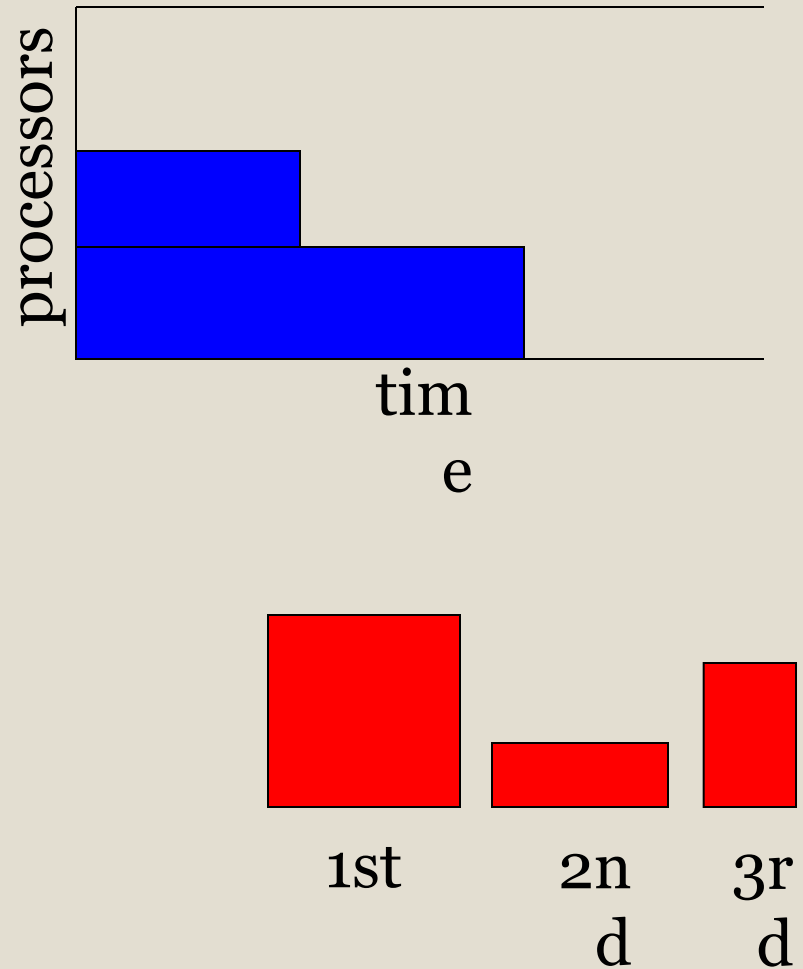
1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

45



1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

45



1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

46


1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

46


1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



EASY Operation

46

1. Schedule jobs on available processors in FCFS order
2. Make reservation for first job that cannot run
3. Schedule additional jobs provided they do not conflict with this reservation



This is called
“backfilling”

2n
d

3r
d



Backfilling Conditions

47

1. Backfill job will terminate before reservation time

OR

2. Backfill job uses only “extra” processors



Runtime Estimates

48

- When users submit jobs, they provide
 1. The number of processors to use
 2. An estimate of the **job runtime**
- **Estimates** are used to predict when processors will become **free** for reservation
- Also used to verify that backfill job will terminate before reservation
- If it does not, it will be killed

Questions?

