

Memory Management

1

**OPERATING SYSTEMS COURSE
THE HEBREW UNIVERSITY
SPRING 2022**

Basic Numbers

2

- Basic units:

- KB = 2^{10} bytes
- MB = 2^{20} bytes
- GB = 2^{30} bytes

- Basic calculations:

- $4\text{GB} / 8\text{KB} = (2^2 / 2^3) * (2^{30} / 2^{10}) = 2^{(20 - 1)} = 2^{19}$
- How many numbers can 8 digit number present?
 - ✦ 2^8 numbers (values between 0 to $2^8 - 1$)
- What is the decimal value of 1101?
 - ✦ $1 * 2^0 + 0 * 2^1 + 1 * 2^2 + 1 * 2^3 = 13$

CPU

3

- Can only use data that is stored in registers and memory.
- Executes a set of instructions:
 - Data handling – set a register, store (register in memory), load (from memory to register)
 - Arithmetic operations on registers - Bitwise operations, compare, and basic mathematics operations.
 - Control flow using registers– branch, conditional branch
- Therefore, programs must be brought (from disk) into memory for them to be run

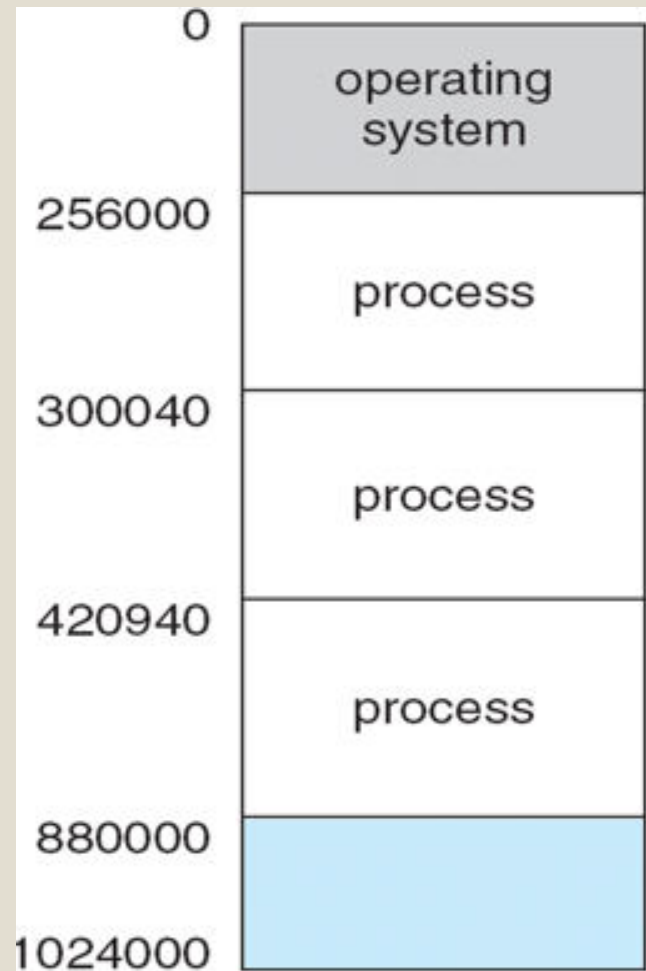
Physical Address

4

- The address in the particular storage cell of main memory.
- The physical address space depends on the memory size
 - If we have 8GB of RAM, the address length is 33 [$\log_2(8\text{GB})$].
- Also called real address, or binary address

Physical Address Basic Example

5



Virtual and Physical address spaces

6

- OS can reserve addresses on the memory, and use physical address to access them
- For “regular” processes we want to choose the exact address dynamically, during run time.
- Therefore, processes don’t know their physical addresses during compile time
- Solution:
 - Programs use “virtual memory” to describe their addresses. Virtual memory starts from 0 for each process
 - Hardware (MMU) & the OS translate virtual memory to physical memory.
- New problem: how to do this translation.


Virtual Address Space (VAS)

7

- The set of ranges of virtual addresses that an operating system makes available to a process
- Also called **logical address space**.
- In 32-bit operating system, the virtual address space size is $2^{32} = 4\text{GB}$, in 64-bit it's 2^{64} (= a lot)
- For example, a pointer to a function or variable is actually a virtual address.

Virtual Address Space

Data contains static variables and globals.



MMU

9

- The user program deals with *logical* addresses - it never sees the *real* physical addresses
- Memory-Management Unit (MMU) is the hardware device that maps virtual to physical address


Segmentation

10

- Memory segmentation is the division of a computer's primary memory into segments
- A program is a collection of segments. A segment is a logical unit such as:
 - main program
 - function
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays

User's View of a Program

11



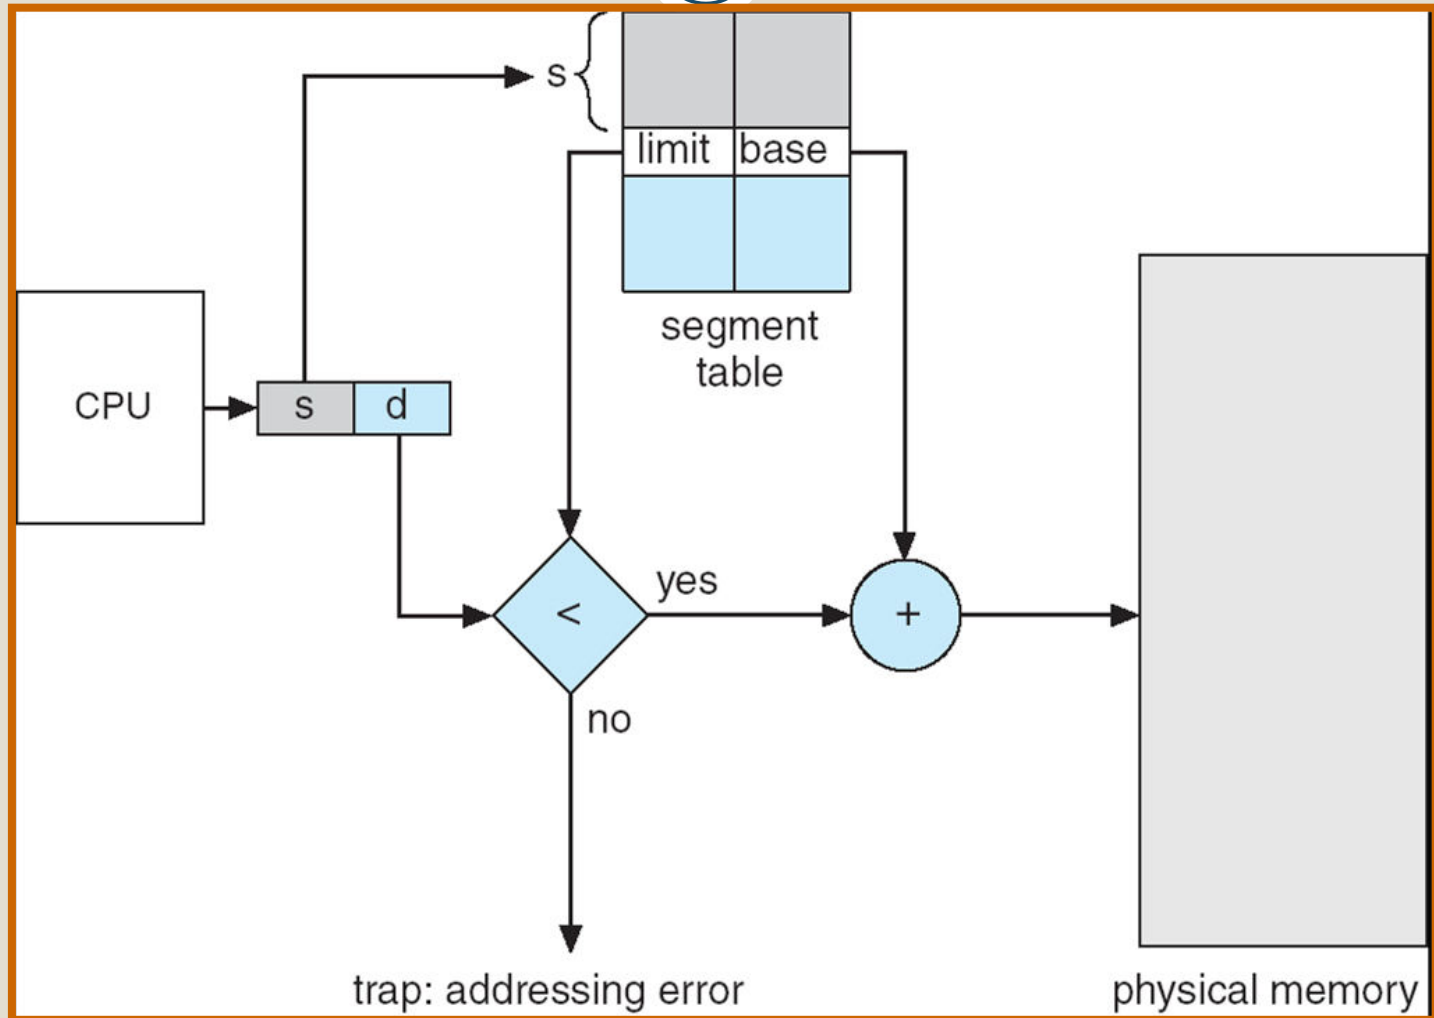
Segmentation Architecture

12

- A Logical address consists of a two tuple:
 <segment-number, offset>
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segment resides in memory
 - **limit** – specifies the length of the segment.
 Offset o is valid if $o < \text{limit}$
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program; segment number s is legal if $s < \text{STLR}$

Segmentation Hardware

13

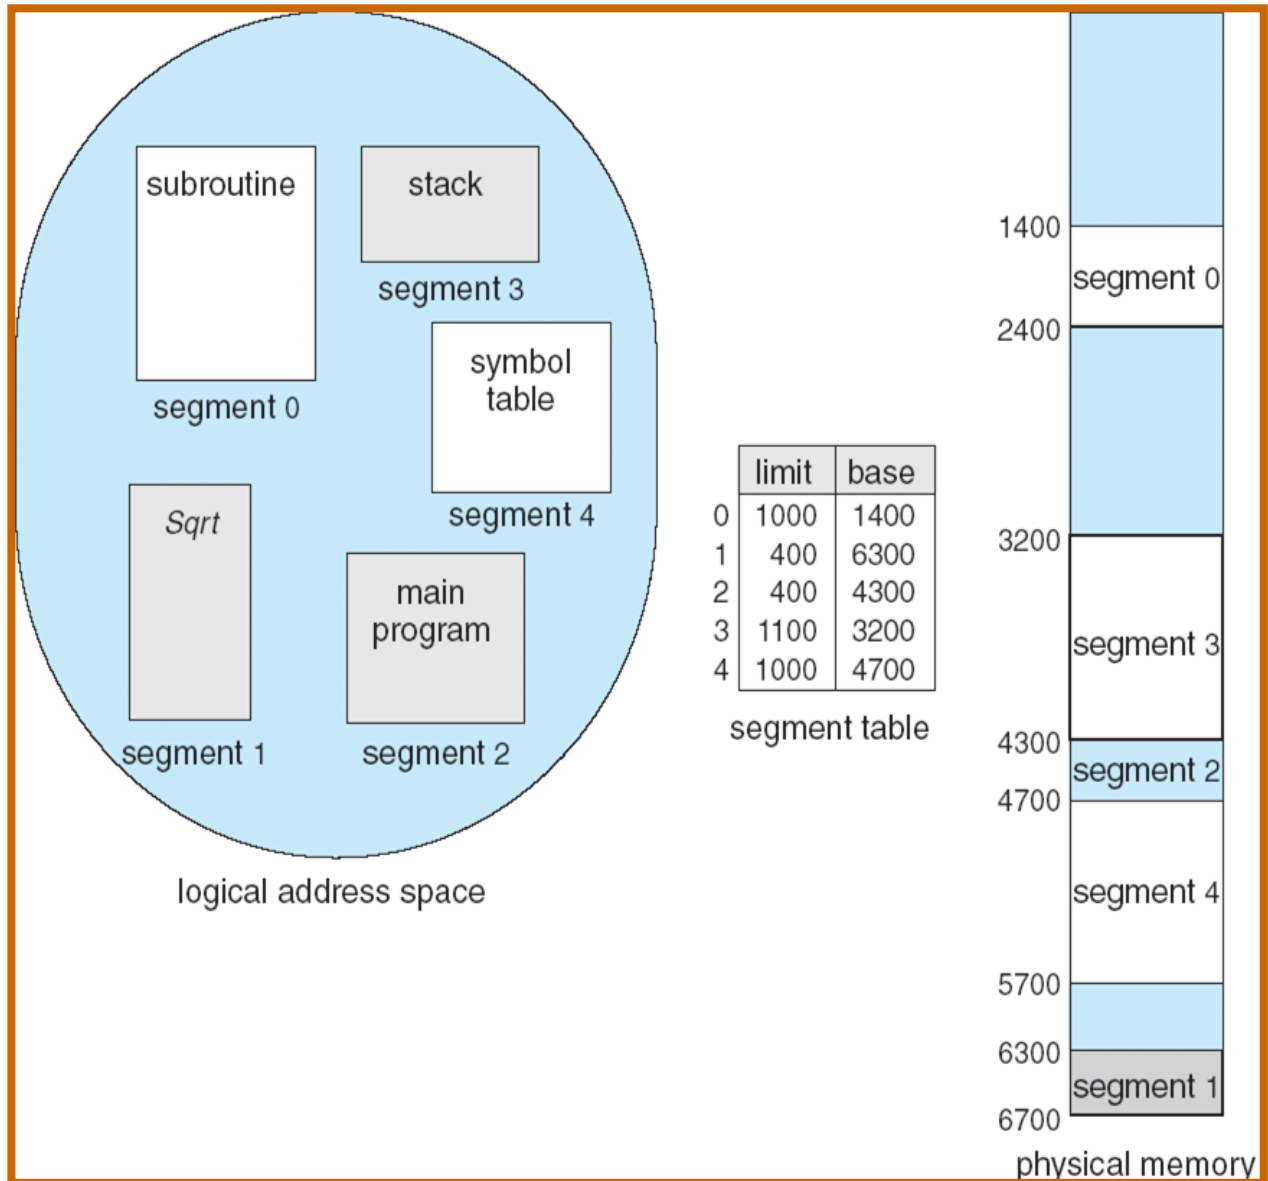


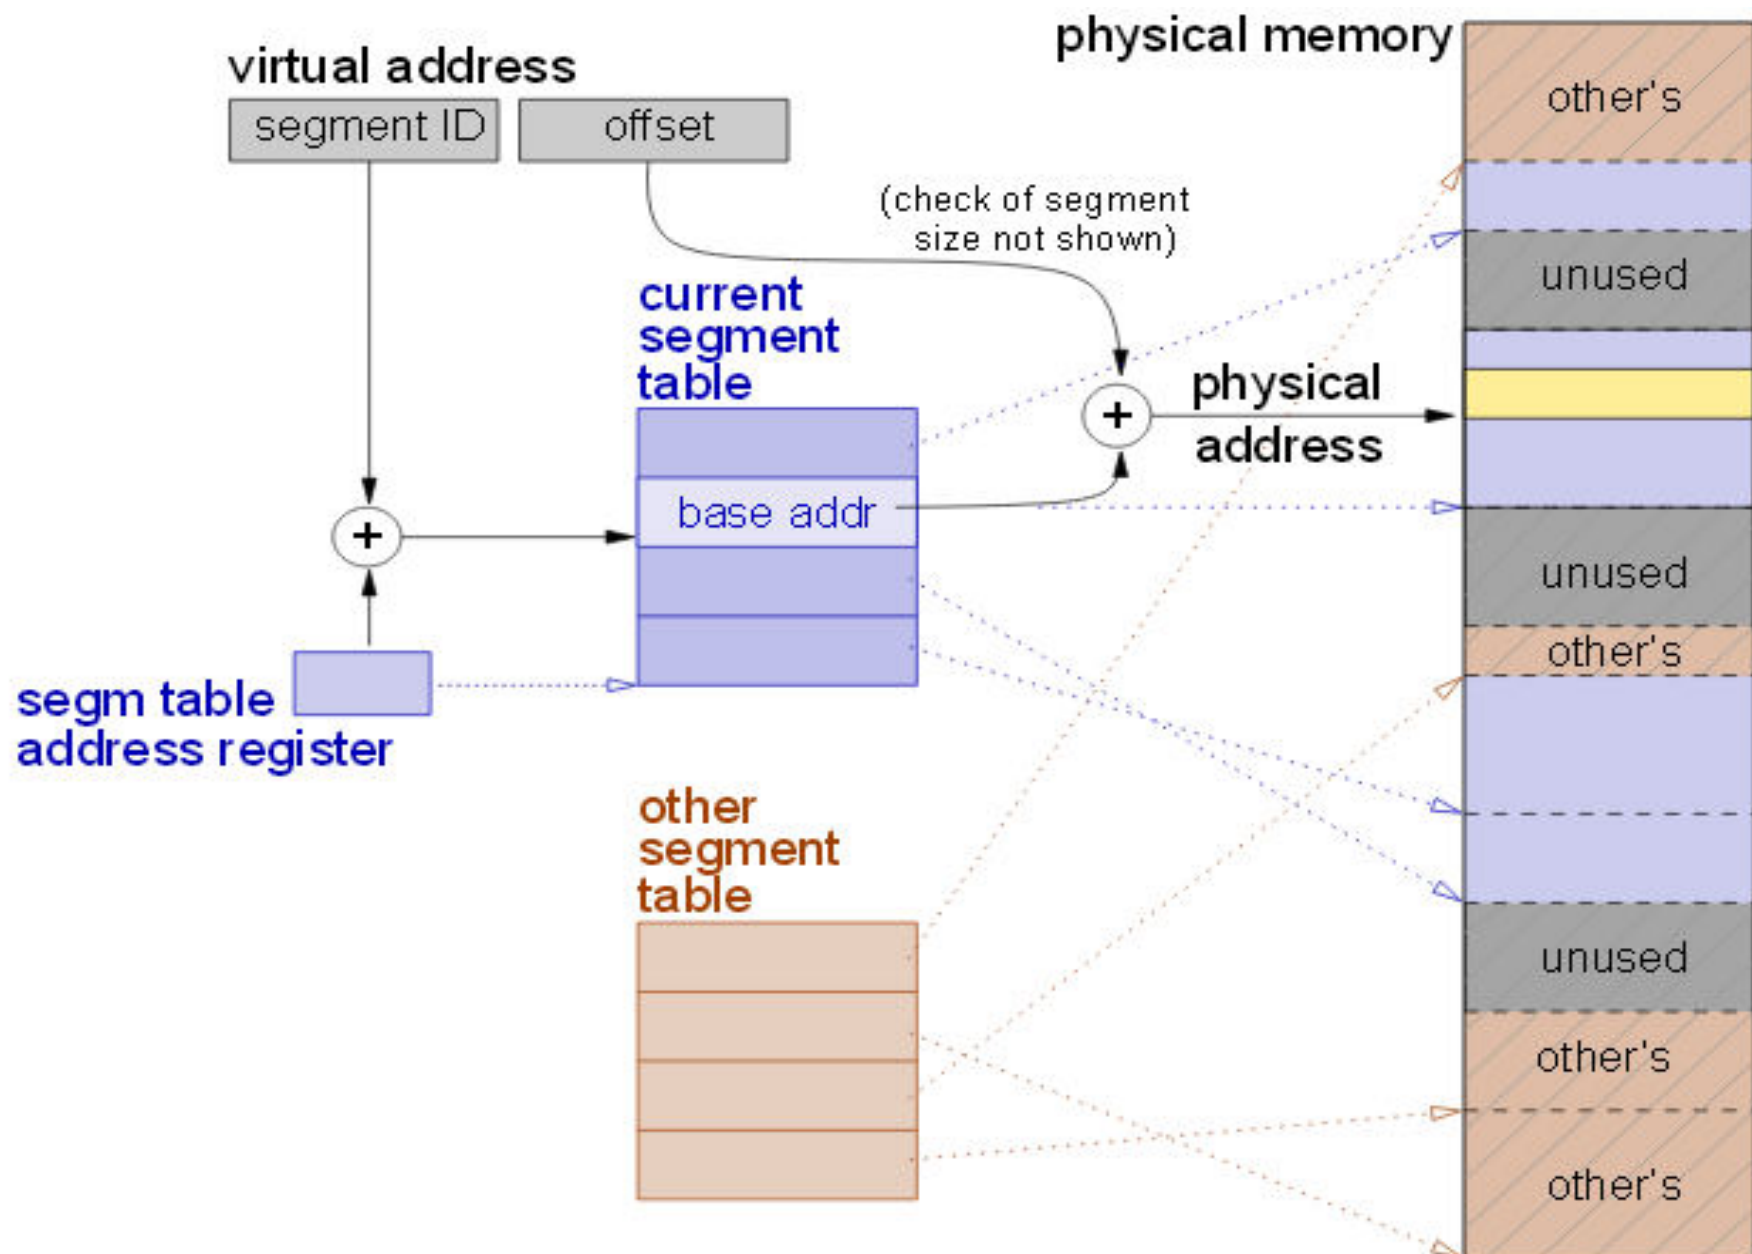
Segmentation Architecture - Cont.

14

- Additional bits in each entry in the segment table:
 - validation bit (also called present-bit)
 - read/write/execute privileges
- Using the validation bit, when there is no space in the RAM, we can store segments in the disk and mark the validation bit 0.

Example of Segmentation





Segmentations: Summary

17

- Divide the program into segments such as: code, heap, stack, data-structures.
- The addresses of each segment are sequential on the physical memory
- Given a virtual address number X in segment Y , we need to:
 - Get the base address of segment Y in the physical memory ($Y_PHYSICAL$)
 - Check that the segment's size is bigger than X (otherwise, throw segmentation fault)
 - Access physical address $Y_PHYSICAL + X$.
- This is done practically by:
 - Given virtual address with n bits, the first bits are the segment number, and the rest of the bits to the offset within the segment
 - There is a table, called Segment table, that maps for each segment its base address and size
- Problem: external fragmentation


<i>SEG</i>	<i>OFFSET</i>
------------	---------------

Paging

18

- The logical address space of process can be noncontiguous; process is allocated physical memory whenever the latter is available.
- Physical memory is partitioned into fixed-sized blocks called **frames**.
- logical memory is partitioned into blocks of the same size called **pages**.
- Keep track of all free frames.
- To run a program of size n pages, need to find n free frames and load program.
- Set up a page table to translate logical to physical addresses.
- No external fragmentation but there is bounded internal fragmentation.

Paging Model of Logical and Physical Memory




logical
memory

0	1
1	4
2	3
3	7


page table

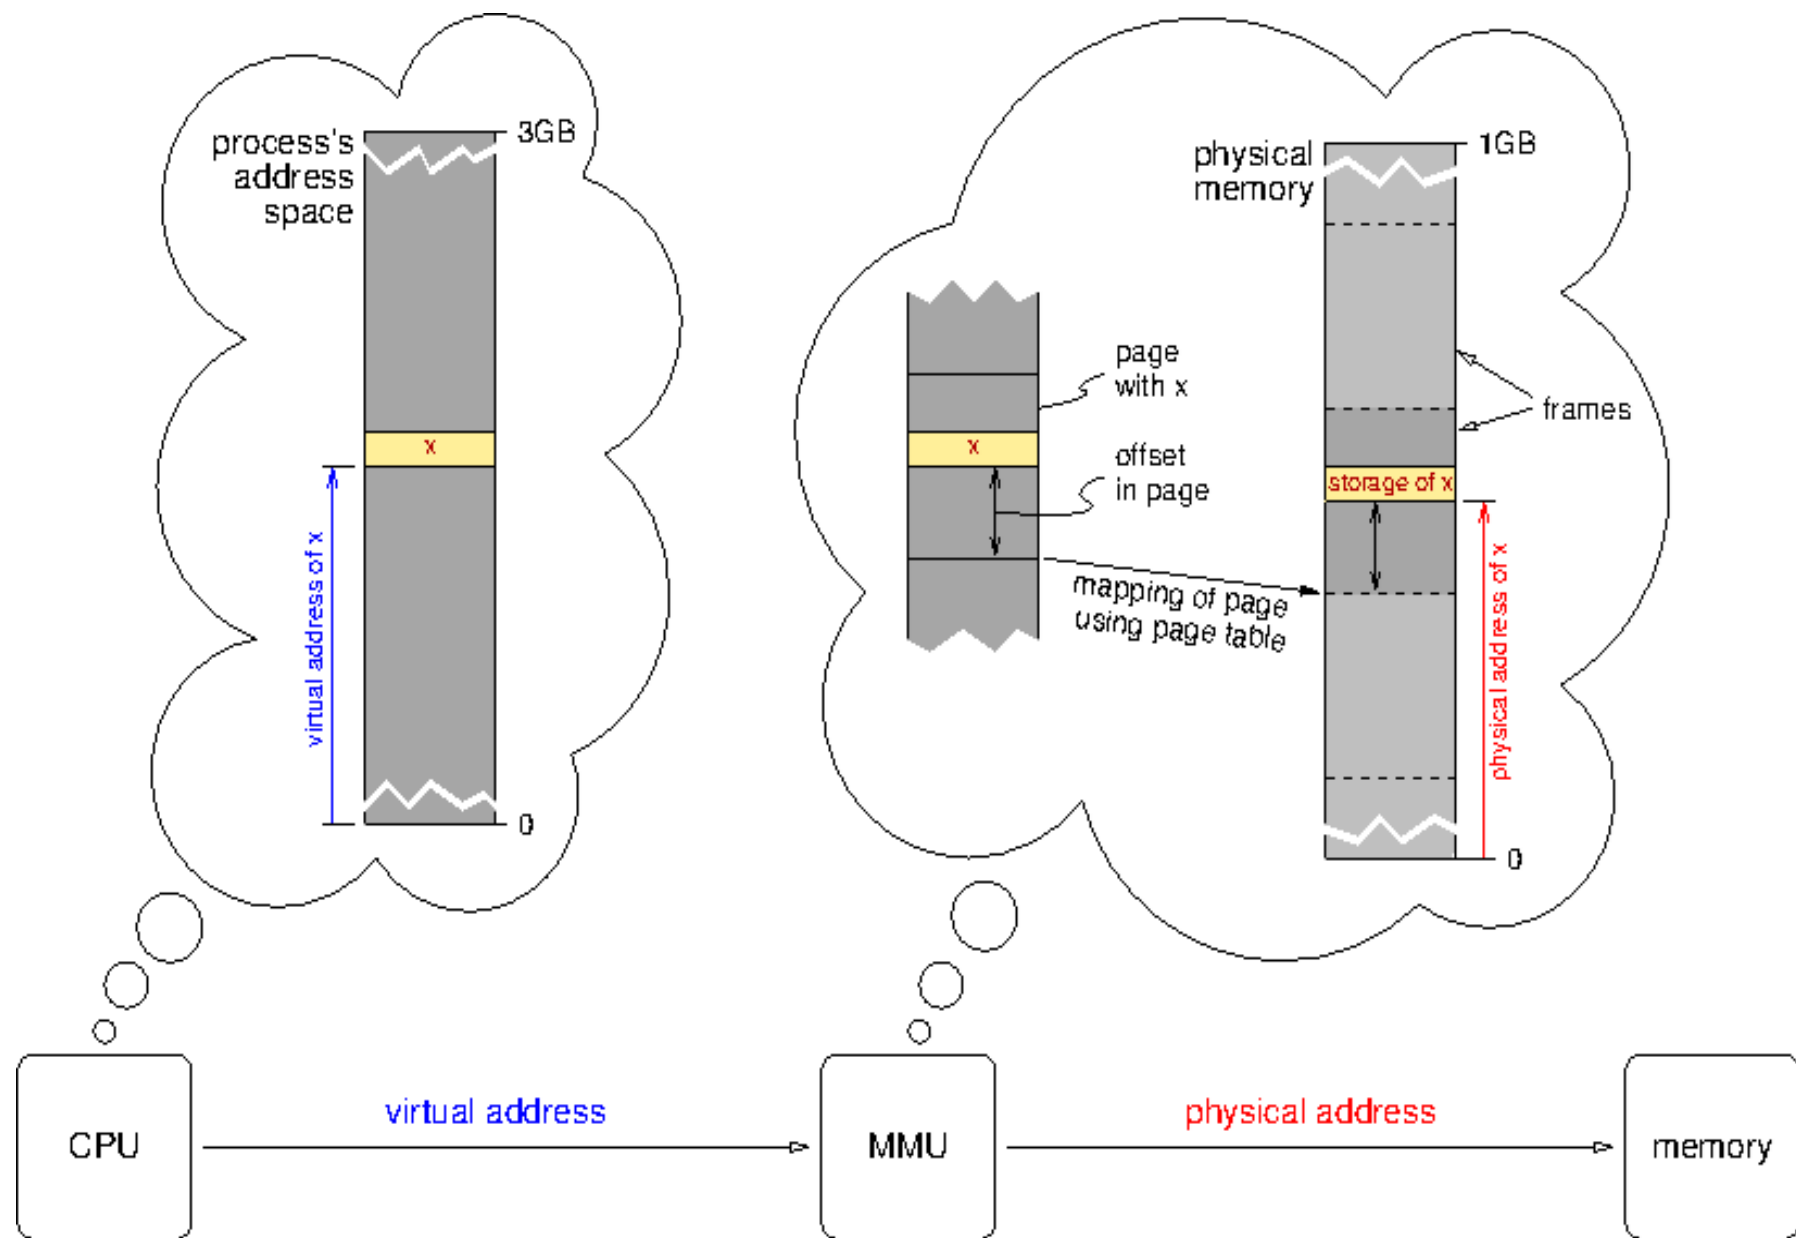
frame
number



physical
memory

Frames Allocation Example





Address Translation Scheme

22

- Address generated by CPU is divided into:
- **Page number (p)** – used as an index into a **page table** which contains base address of each page in physical memory
- **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit
- For a given logical address space of size 2^m words and **page size** 2^n words:

page number	page offset
p	d
$m - n$	n

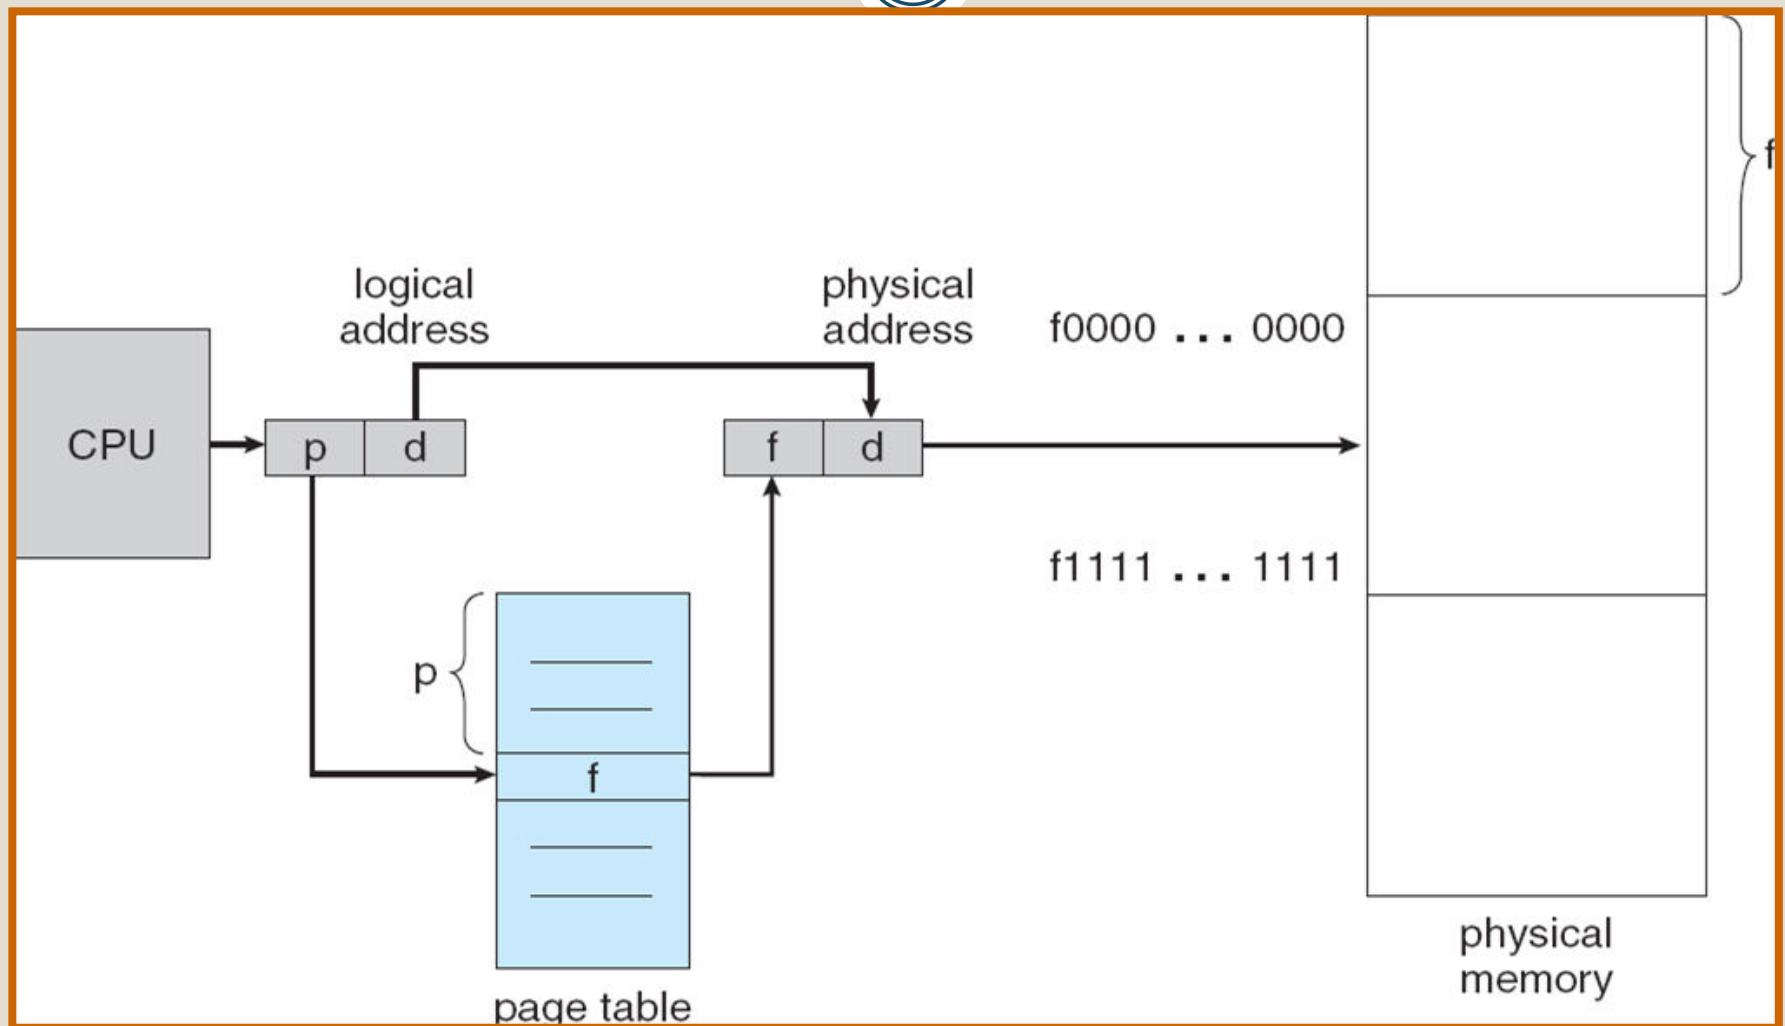
Implementation of Page Table

23

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PRLR)** indicates size of the page table

Paging Hardware

24



Working with the page table

25

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
------	-------

0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Slide from
lecture

Working with the page table

26

Address 6
0110

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Slide from
lecture

Working with the page table

27

Address 6
0110

→

Page 01
Offset 10

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Slide from
lecture

Working with the page table

28

Address 6
0110

Page 01
Offset 10

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
0	6
1	3
2	1
3	4

(per process)

Used by the MMU

Physical
memory

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Slide from
lecture

Working with the page table

29

Address 6
0110

Page 01
Offset 10

Physical
memory

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
0	6
1	3
2	1
3	4

Frame 011
Offset 10

(per process)

Used by the MMU

Slide from
lecture

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Working with the page table

30

Address 6
0110

Page 01
Offset 10

Physical
memory

Logical
memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Page	Frame
0	6
1	3
2	1
3	4

Physical
address
01110 (14)

Frame 011
Offset 10

(per process)

Used by the MMU

Slide from
lecture

	0
	1
	2
	3
i	4
j	5
k	6
l	7
	8
	9
	10
	11
e	12
f	13
g	14
h	15
m	16
n	17
o	18
p	19
	20
	21
	22
	23
a	24
b	25
c	26

Page Table Additional Info

26

- **Valid bit** - indicates whether the page is assigned to a frame.
- **Modified - Dirty bit** - indicates whether the page was modified.
- **Used bit** - when was the last accesses to page, i.e. timestamp.
 - E.g. for the clock algorithm for page replacement
- **Access permissions** - read-only, read-write

Page Replacement

27

- If the **valid bit** is off, the page is not mapped to a frame
- This caused an exception named **page fault**
- The OS searches for an available frame
 - If there is one – the page is loaded from the disk to this frame
 - Otherwise, the OS pages out (swaps out) a page to the disk, brings the requested frame to this location in the memory, and updates the page table accordingly.
- To choose the evicted page, the OS uses Page Replacement Algorithms, such as the clock algorithm (using **used bit**)

Structure of the Page Table


28

- As the number of processes increases, the percentage of memory devoted to page tables also increases.
- The following structures solved this problem:
 - Hierarchical Paging
 - Inverted Page Tables


Hierarchical Page Tables

29

- Also called **multilevel page table**
- Break up the logical address space into multiple page tables
 - We don't need to preserve all the page tables, but only those that we use.
 - This works well because many times programs don't use the middle addresses ->
- A simple technique is a two-level page table (the page table is paged).



Two-Level Page-Table Scheme



Two-Level Paging Example

31

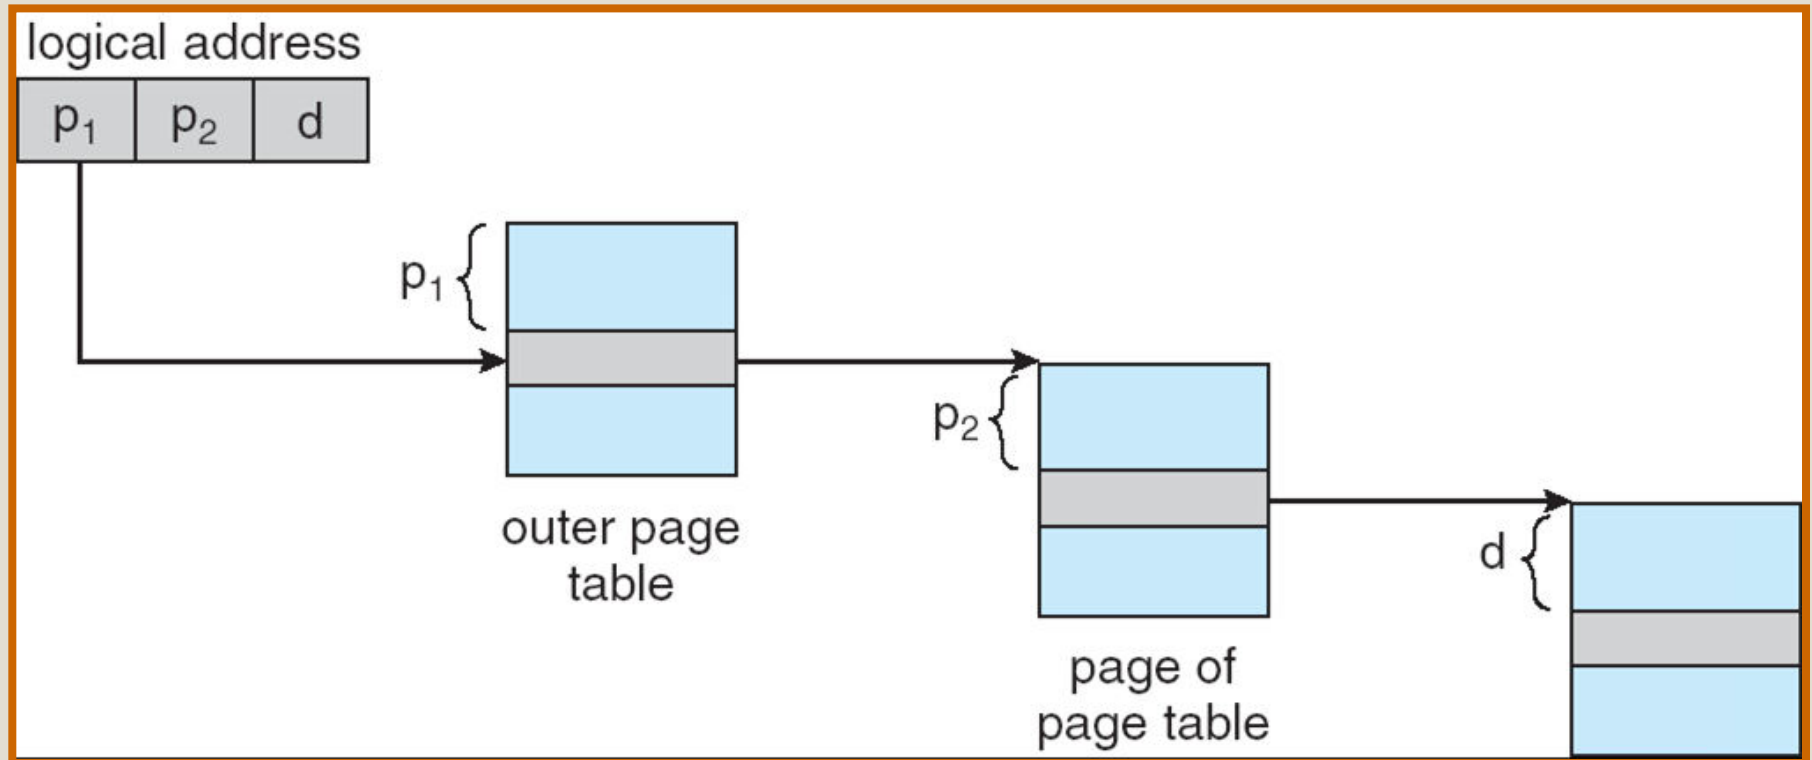
- On a 32-bit machine, with 1K page size
- A logical is divided into
 - a page number consisting of 22 bits ($p_1 + p_2$)
 - a page offset consisting of 10 bits (d)
- Since the page table is paged, the page number is divided into:
 - a 12-bit page number (p_1)
 - a 10-bit page offset (p_2)
- Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d
12	10	10

- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table.

Address-Translation Scheme

32



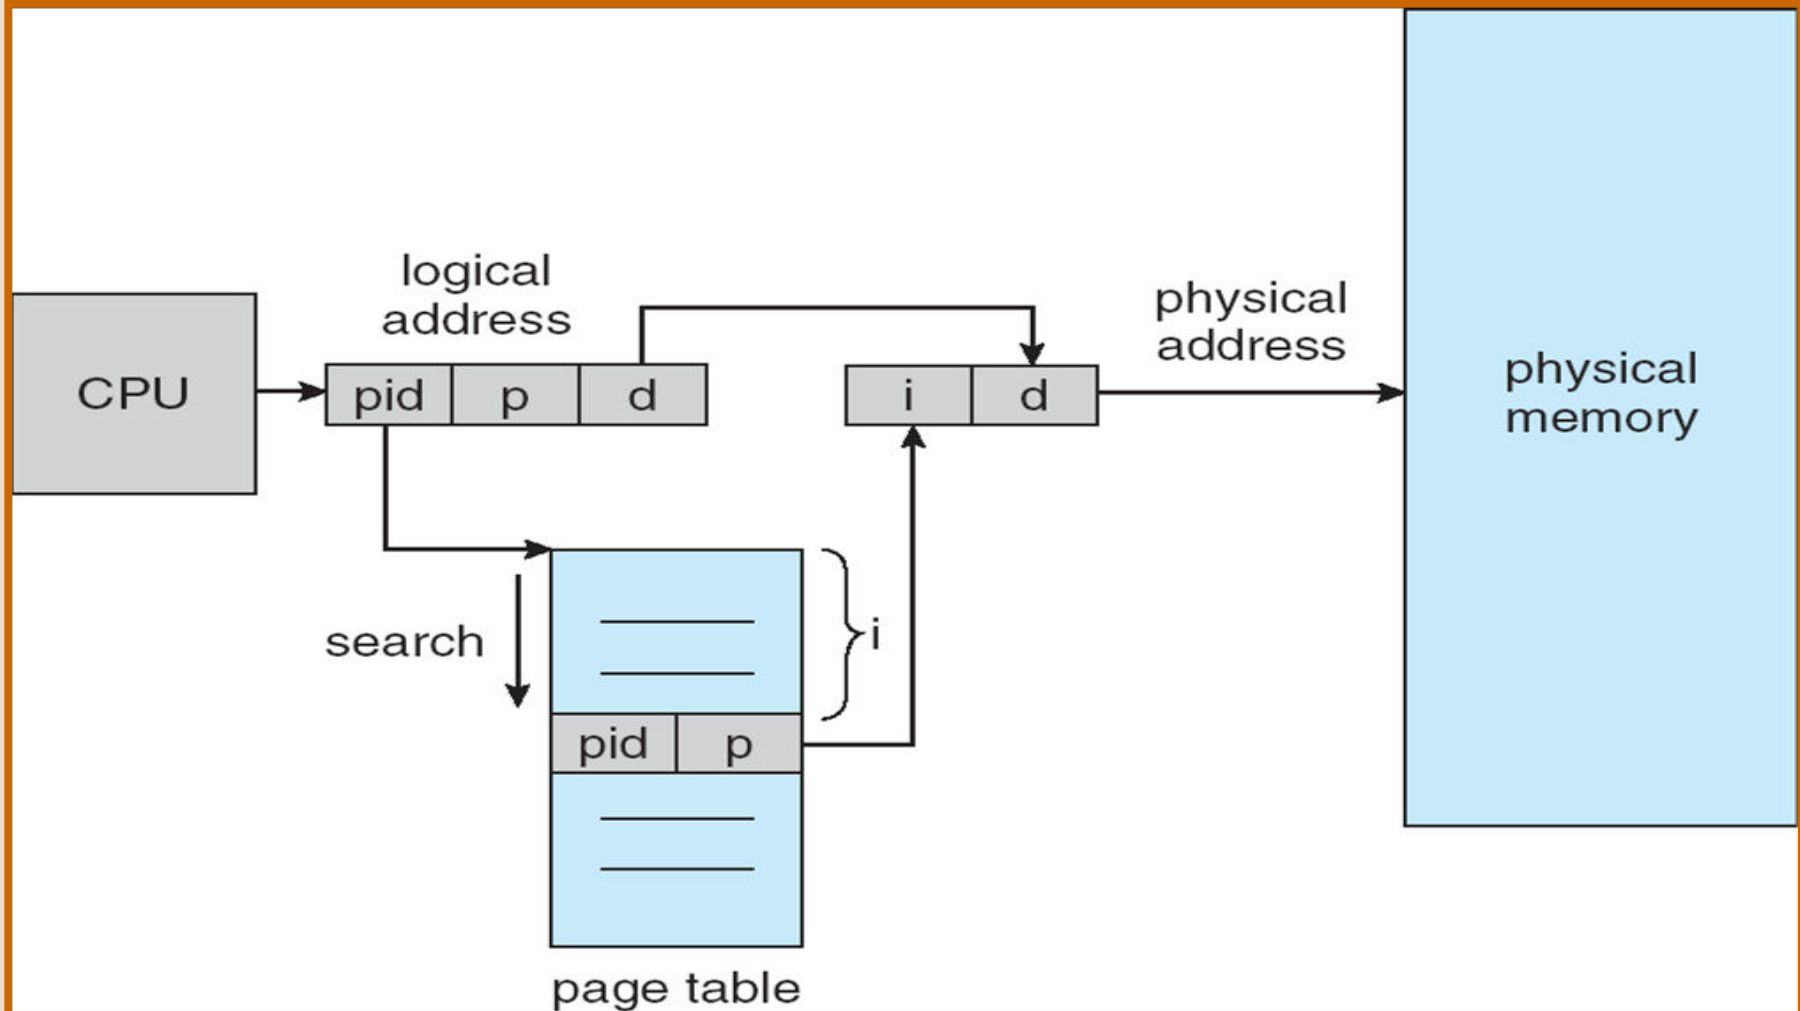
Inverted Page Table

33

- There is one table with the size of the physical memory shared for all processes.
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page.
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.


Inverted Page Table Architecture

34



Page Table - Limitations

35

- Mapping virtual addresses to physical memory adds overhead to every memory access.
 - Even in 1-level paging, every data/ instruction access requires two memory accesses: one for the page table and one for the data/instruction.
 - Usually it's more than two accesses (multi-level paging).
- 
- The CPU uses a cache of recently used mappings from the operating system's page table.
 - This cache is called the **Translation look-aside buffers (TLBs)** which is an associative cache.

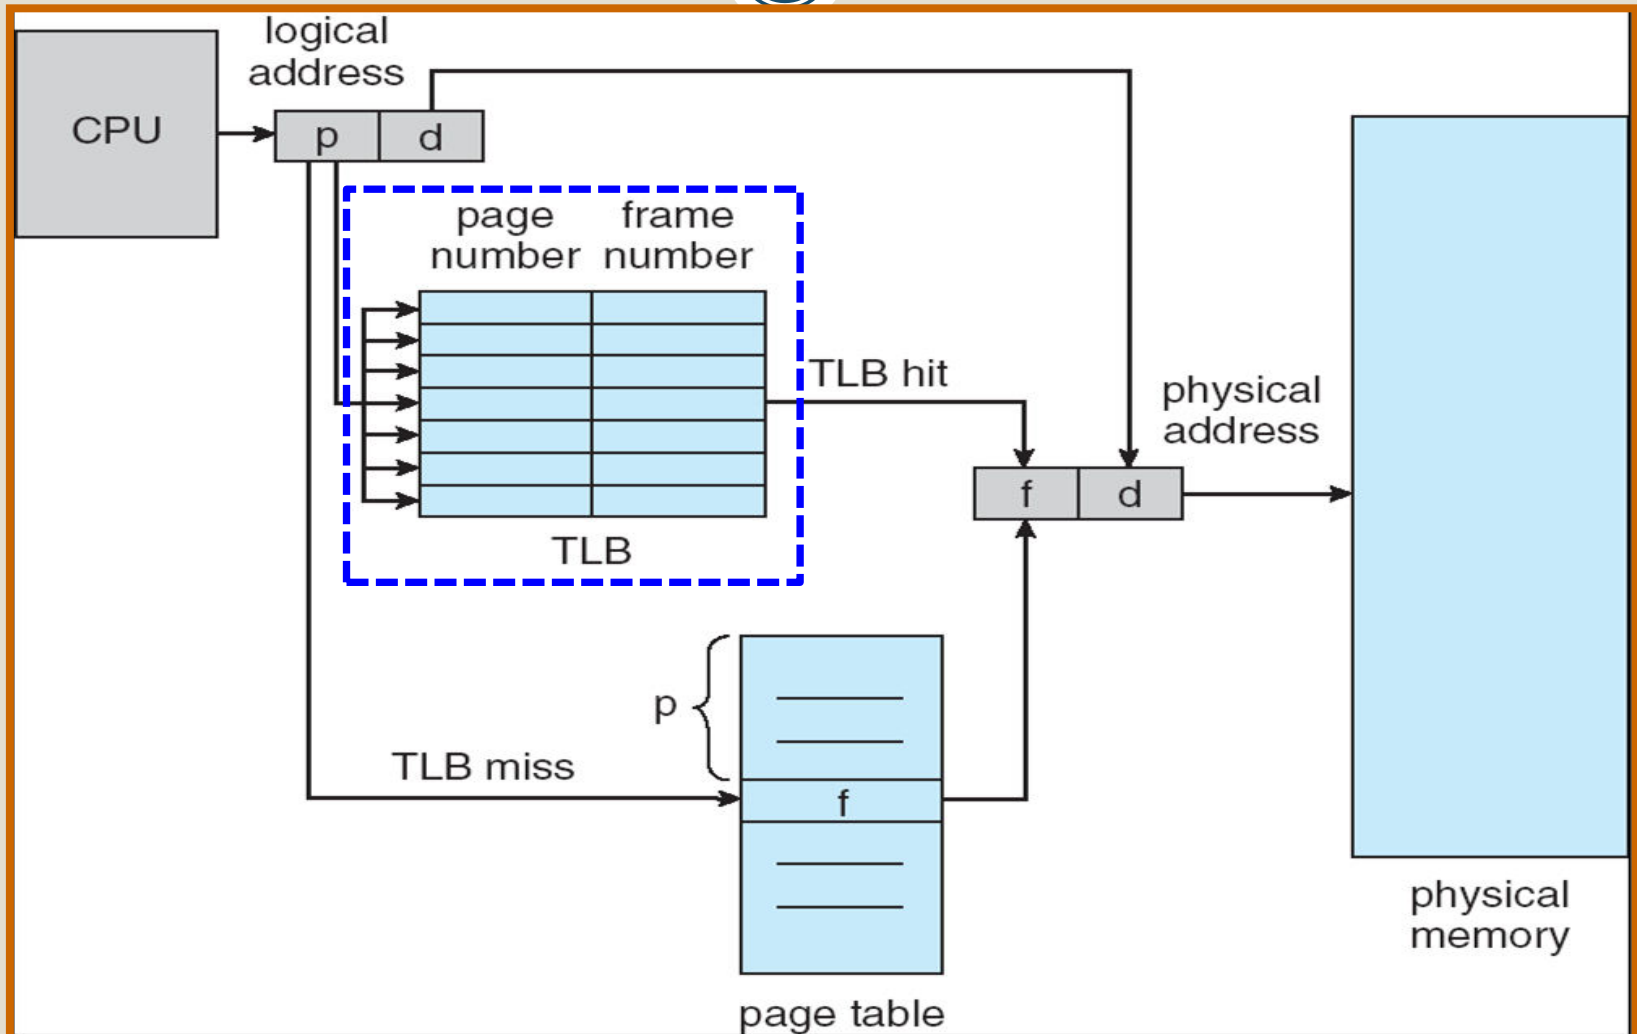
TLB

36

- Usually it's fully associative with 64 entries
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process.

Paging Hardware With TLB

37



Putting Everything Together: Address Translation

38

Virtual Address:

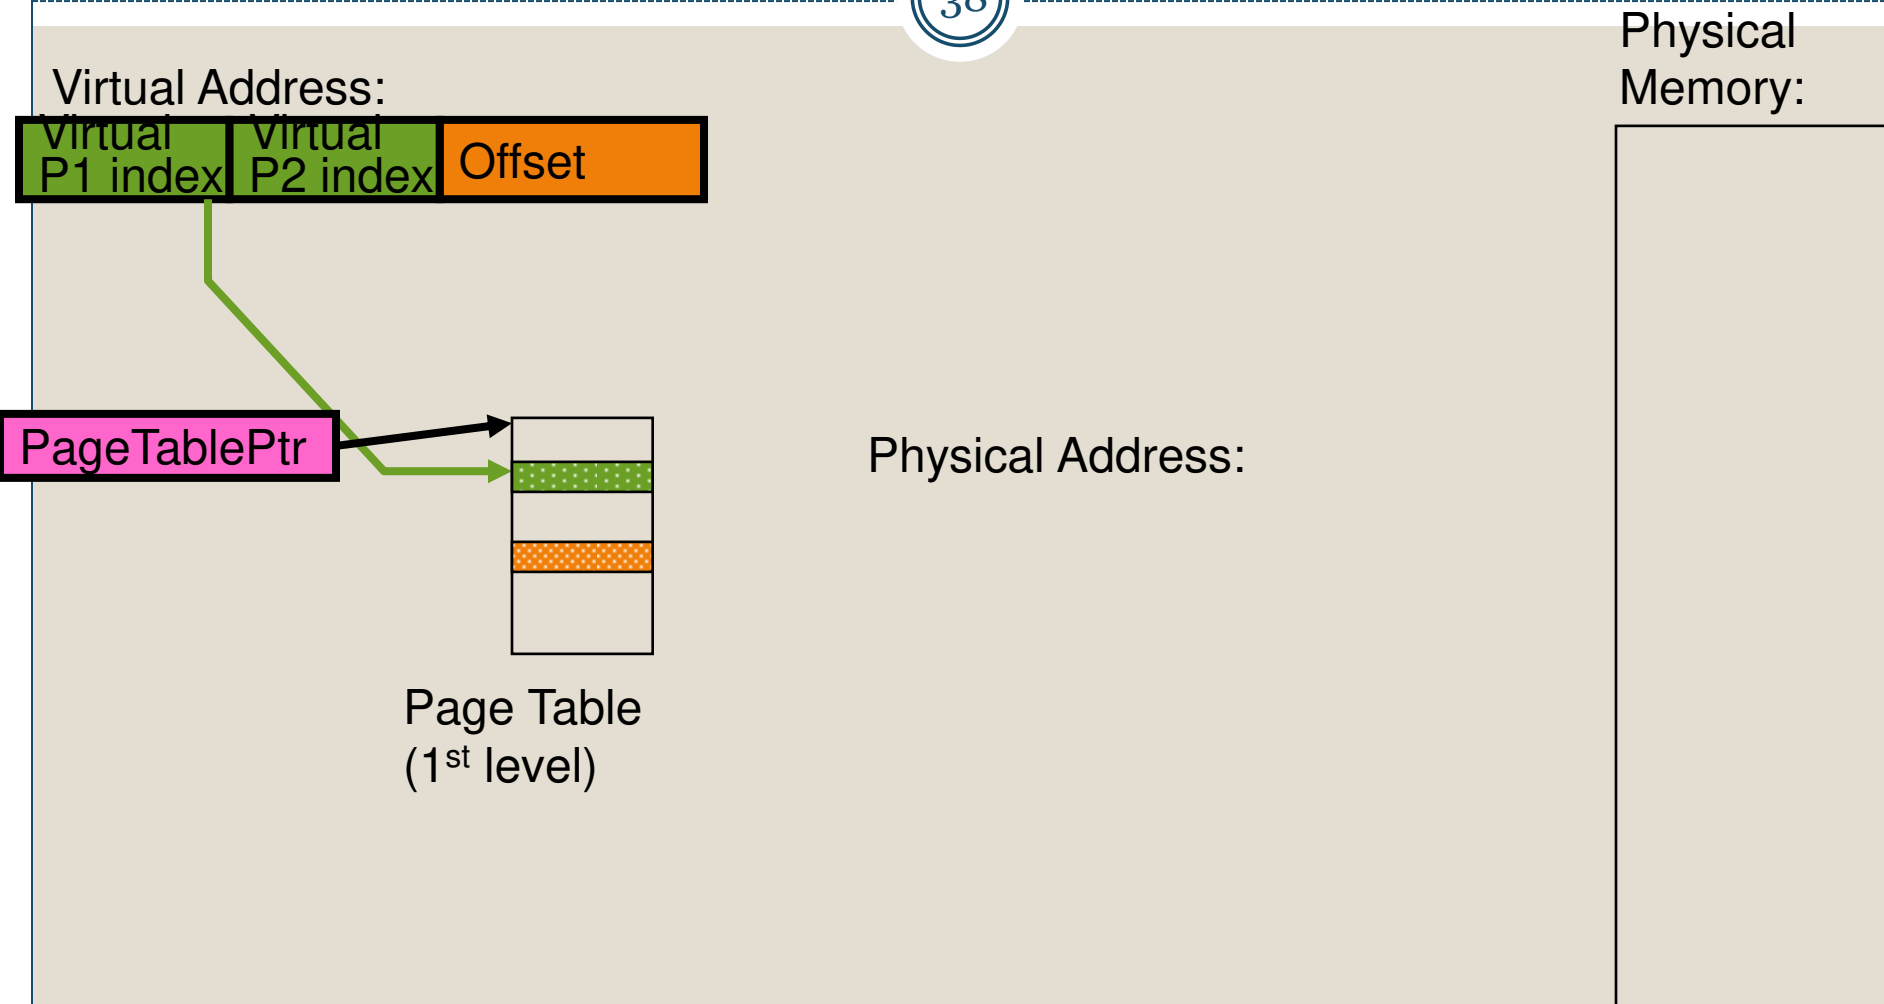
Virtual P1 index	Virtual P2 index	Offset
---------------------	---------------------	--------

Physical
Memory:

Physical Address:

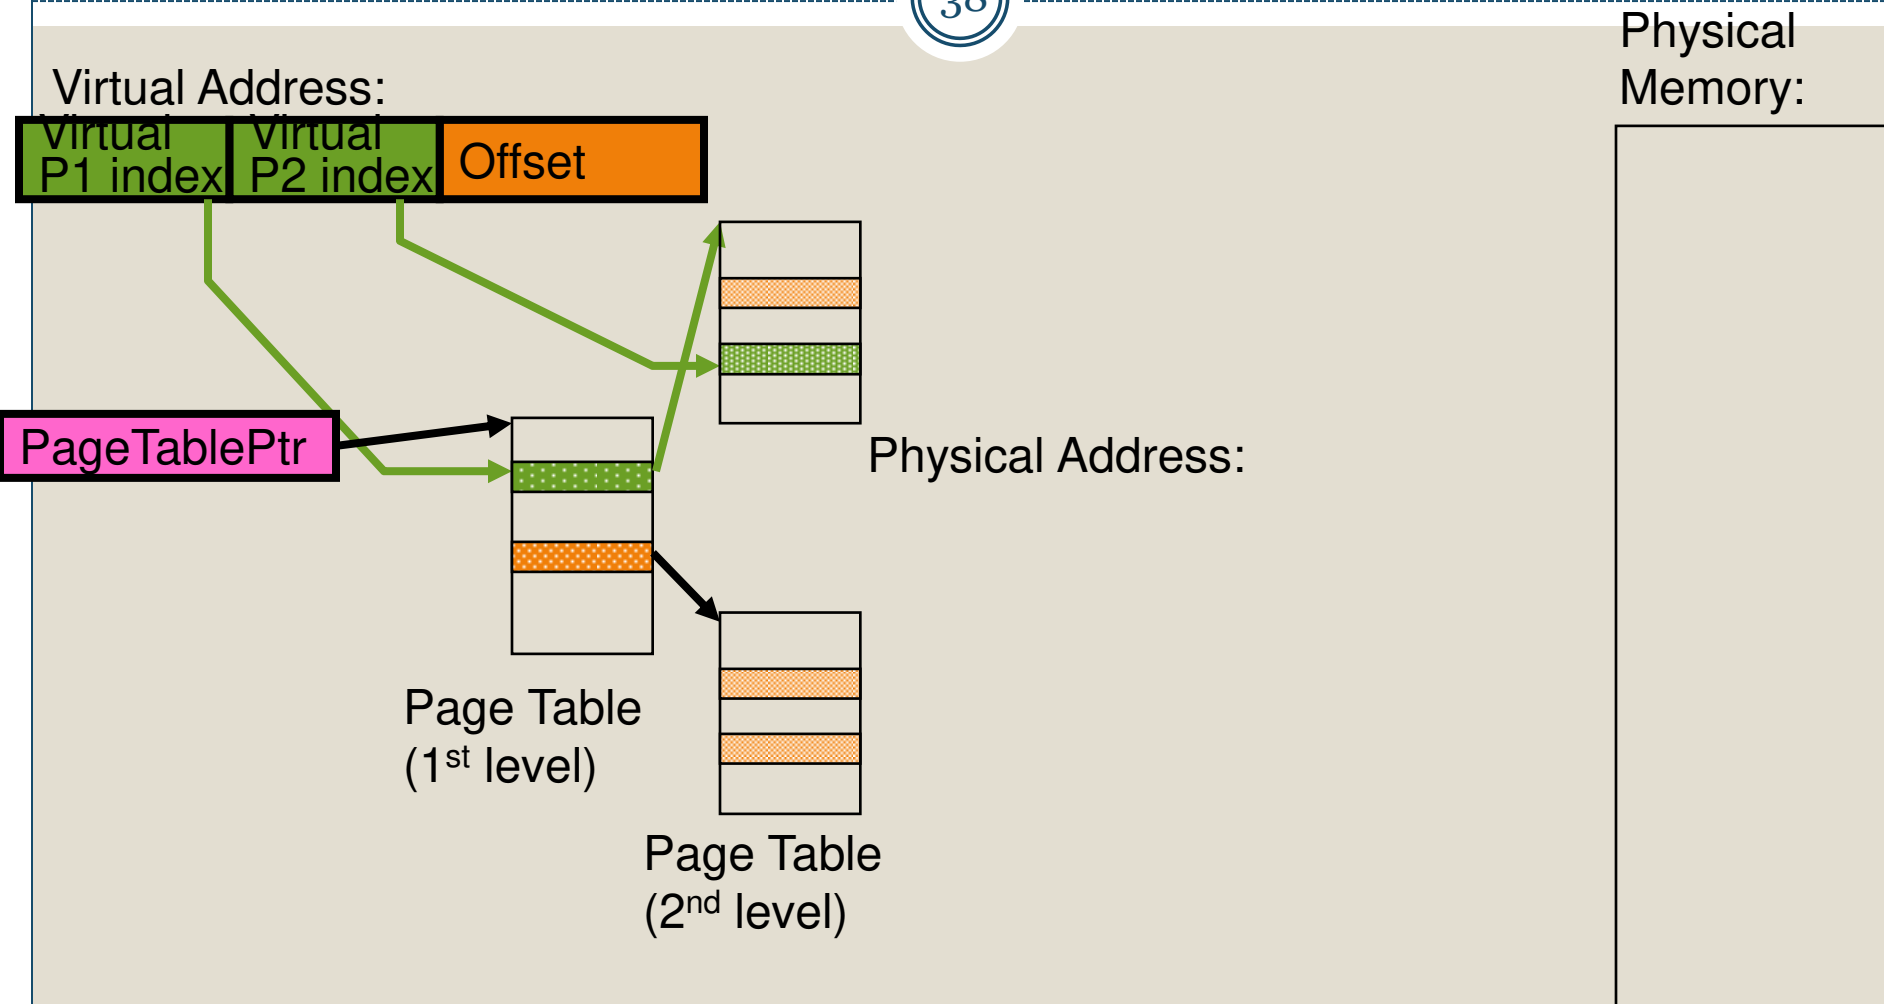
Putting Everything Together: Address Translation

38



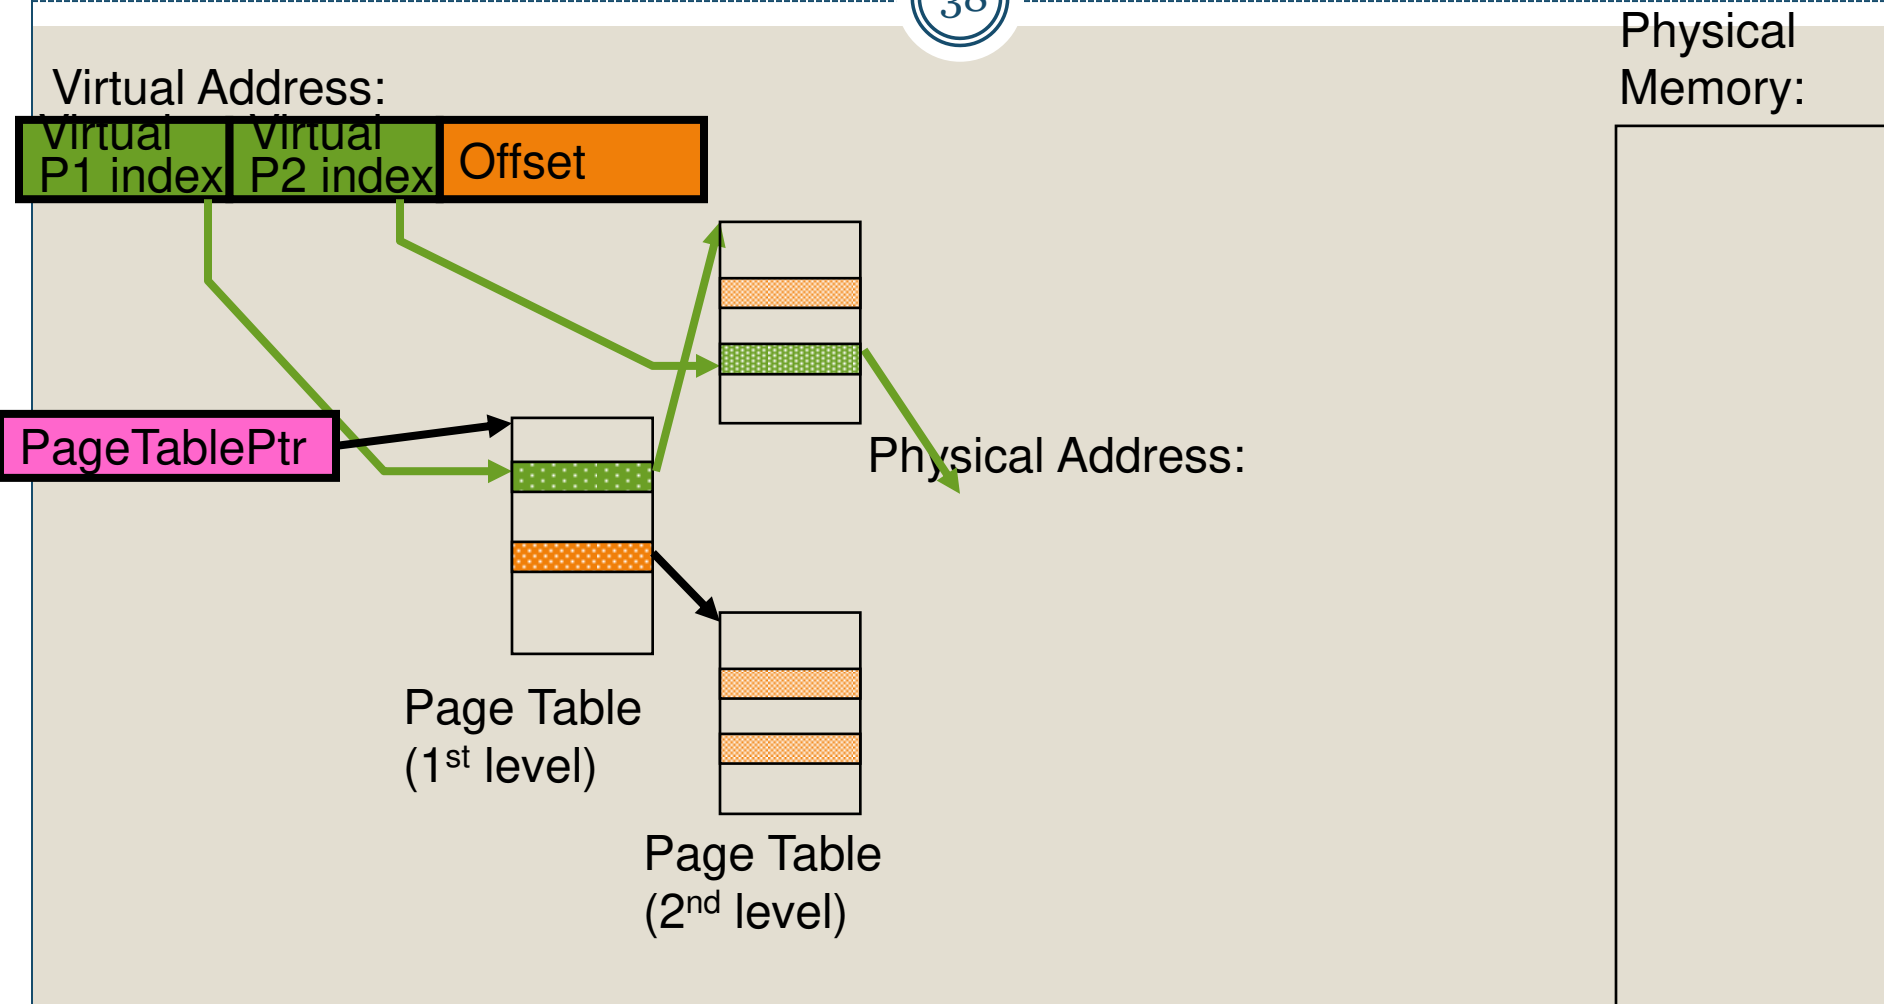
Putting Everything Together: Address Translation

38



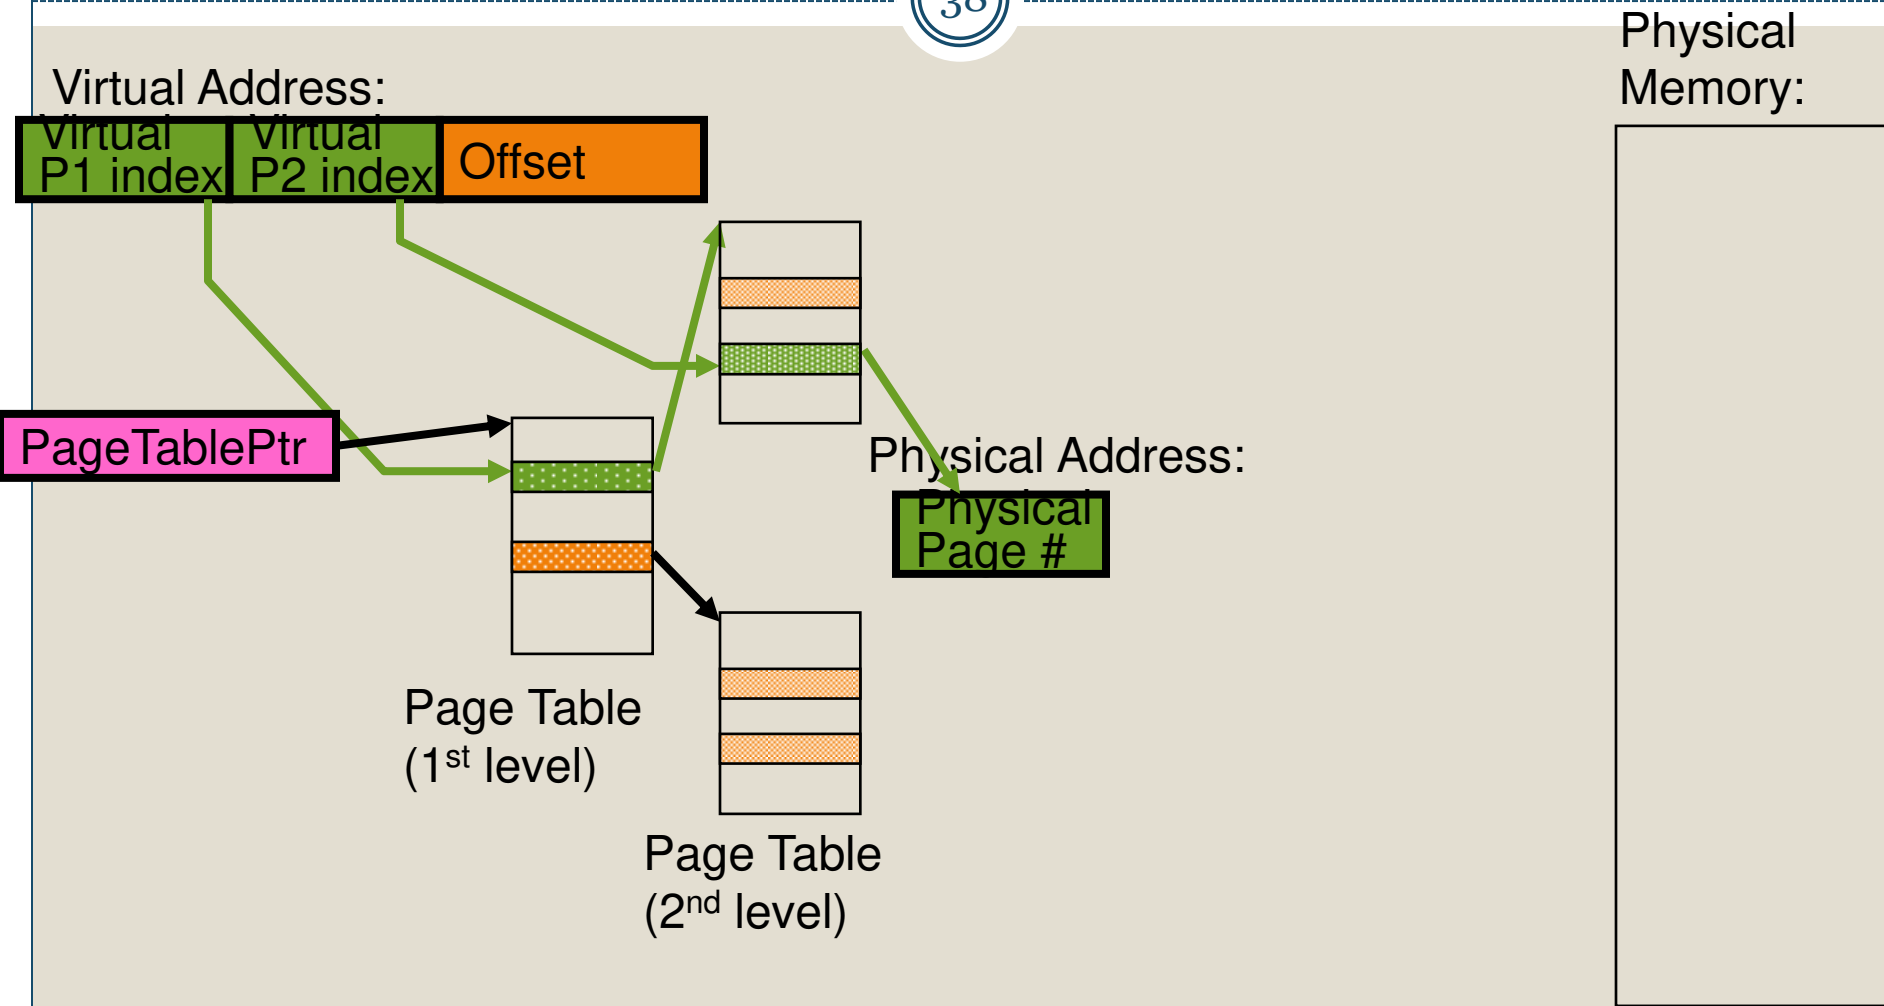
Putting Everything Together: Address Translation

38



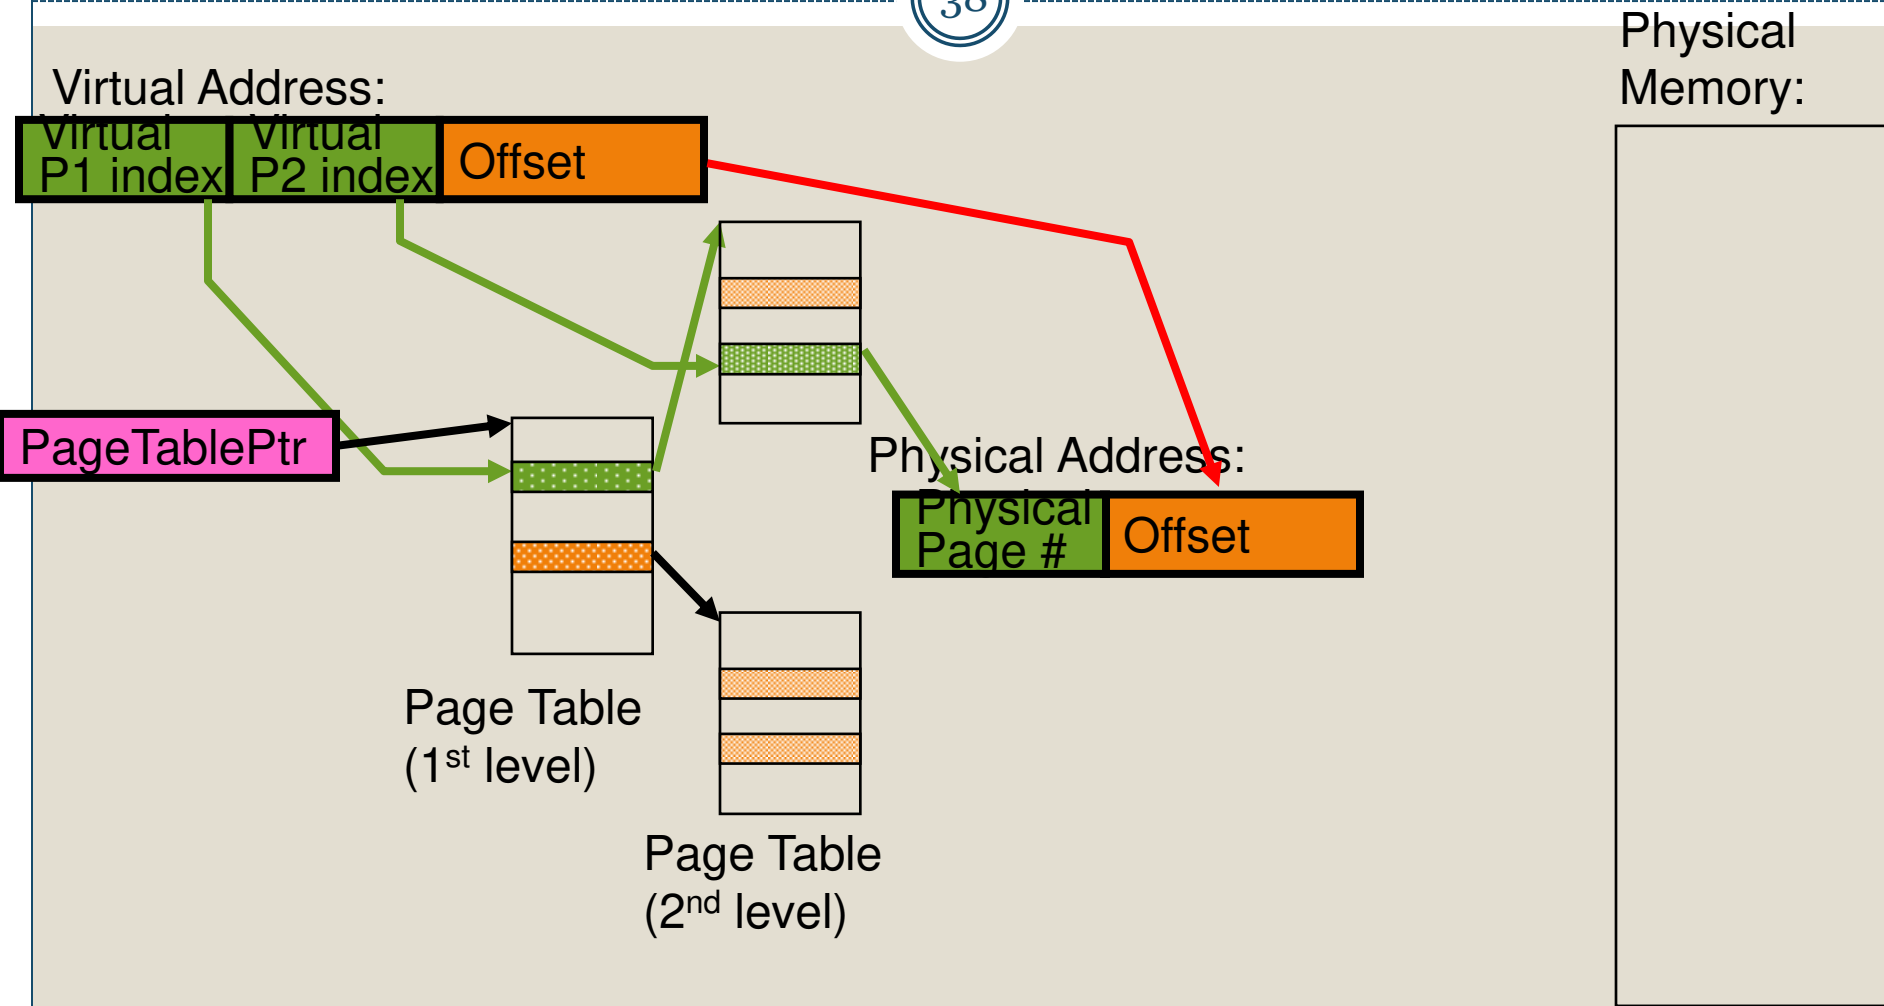
Putting Everything Together: Address Translation

38



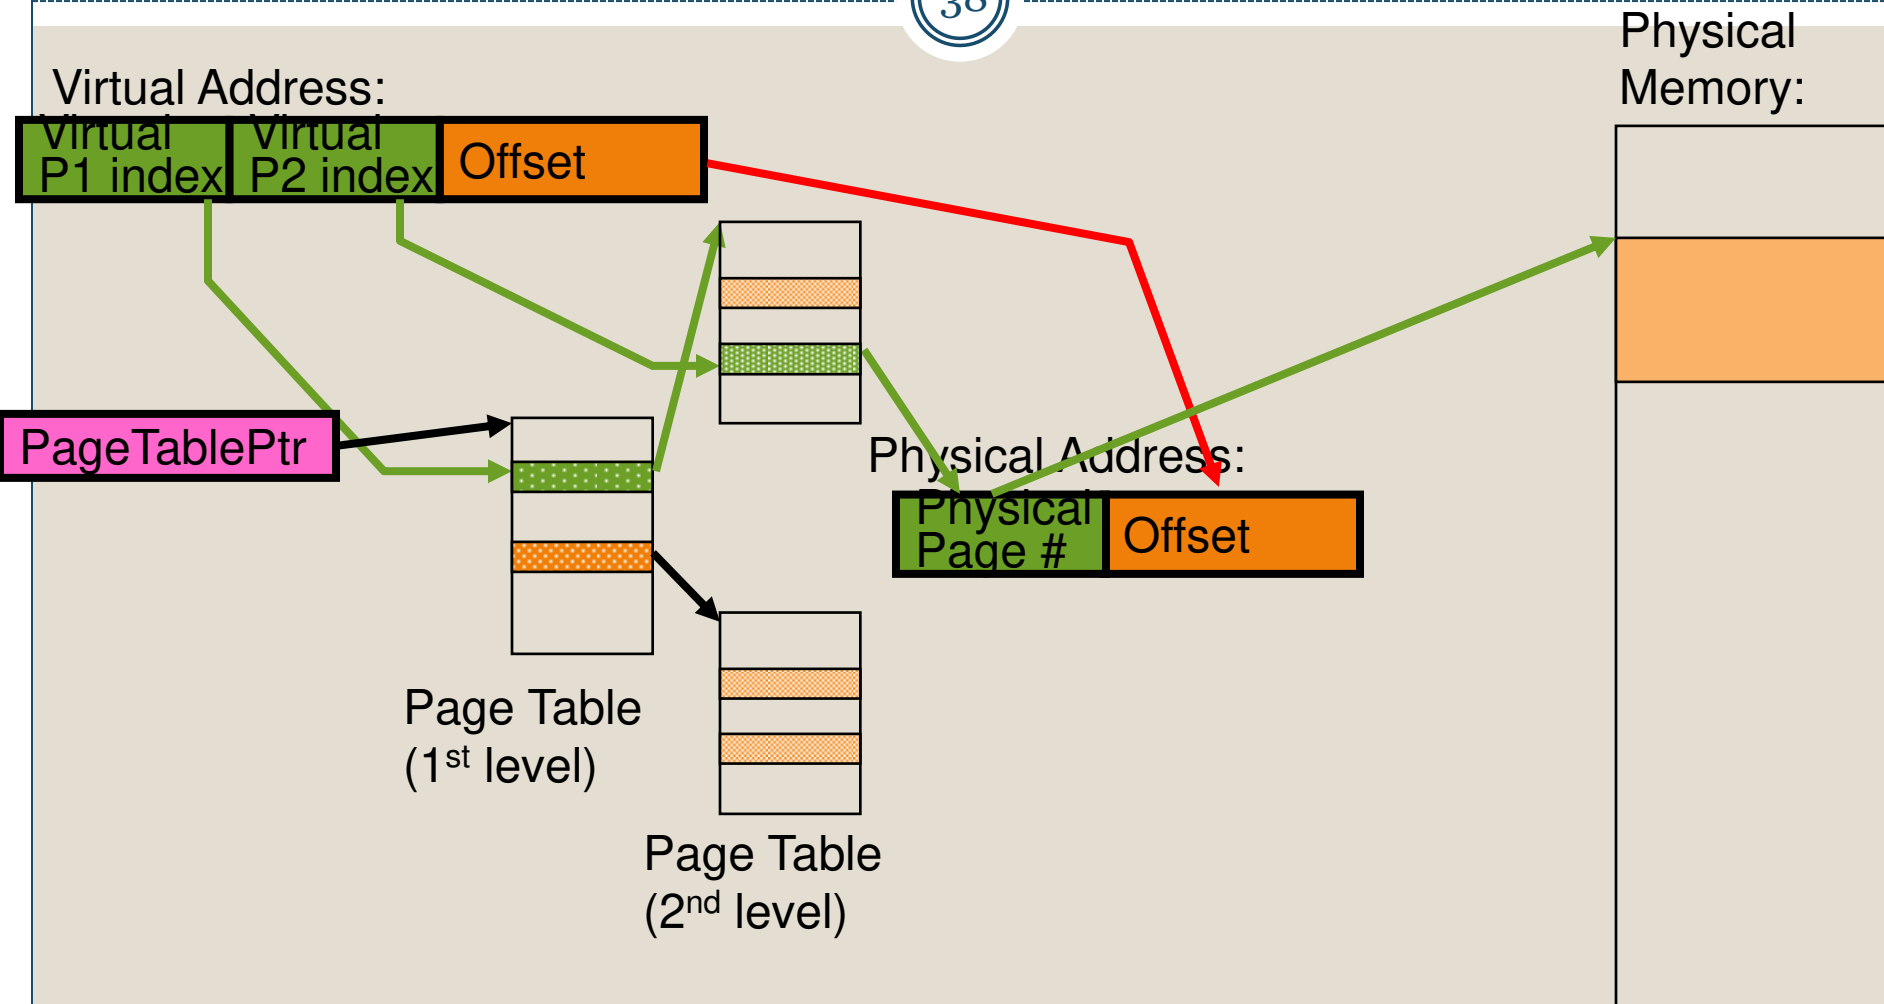
Putting Everything Together: Address Translation

38



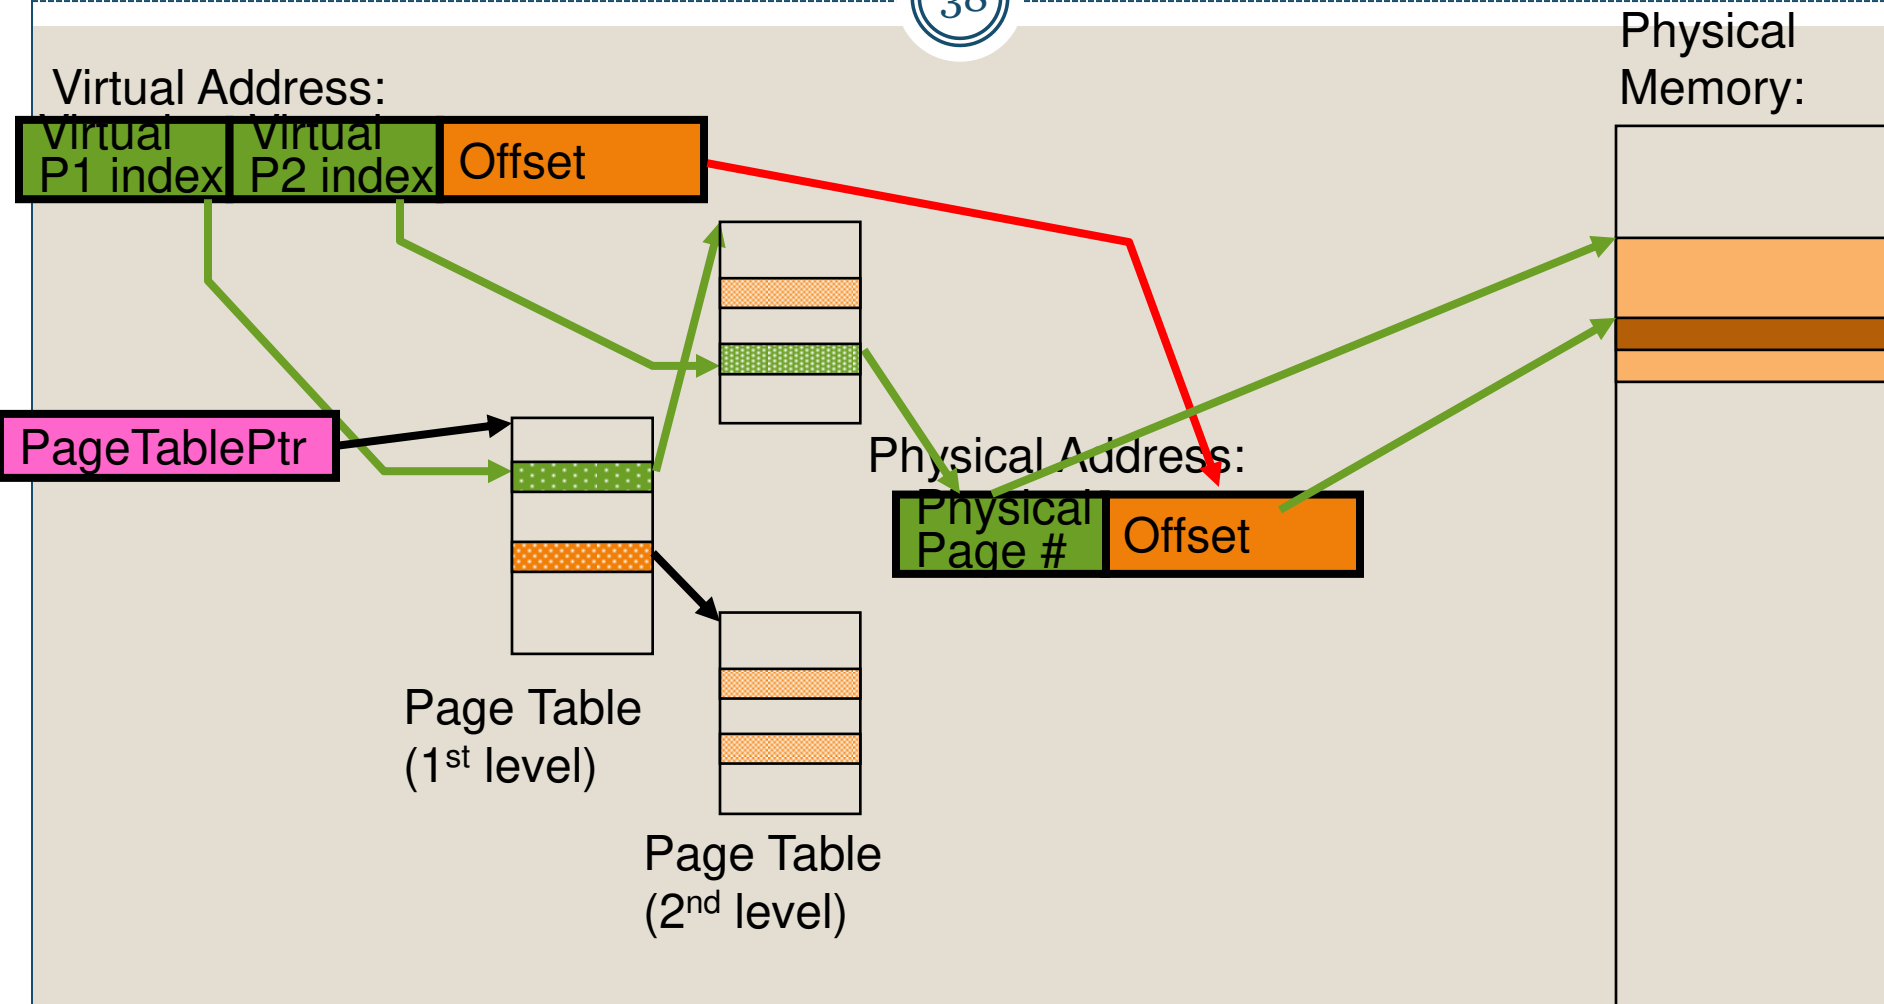
Putting Everything Together: Address Translation

38



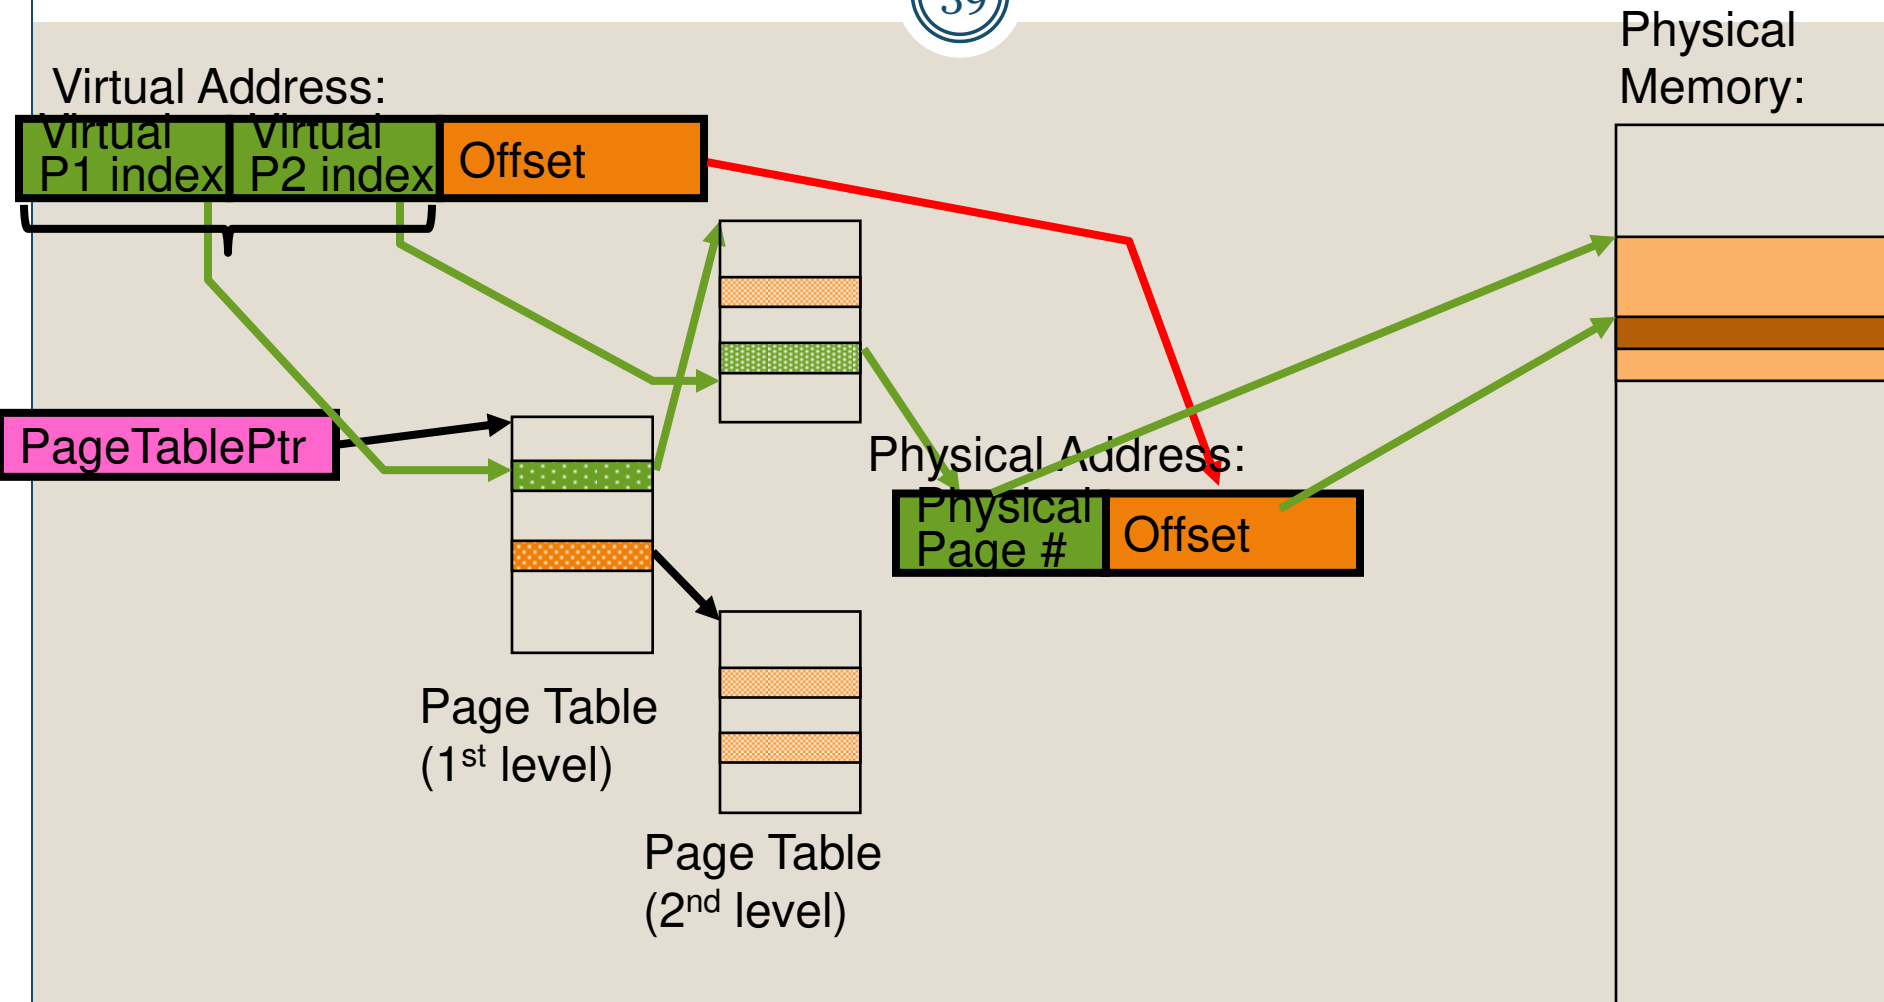
Putting Everything Together: Address Translation

38



Putting Everything Together: TLB

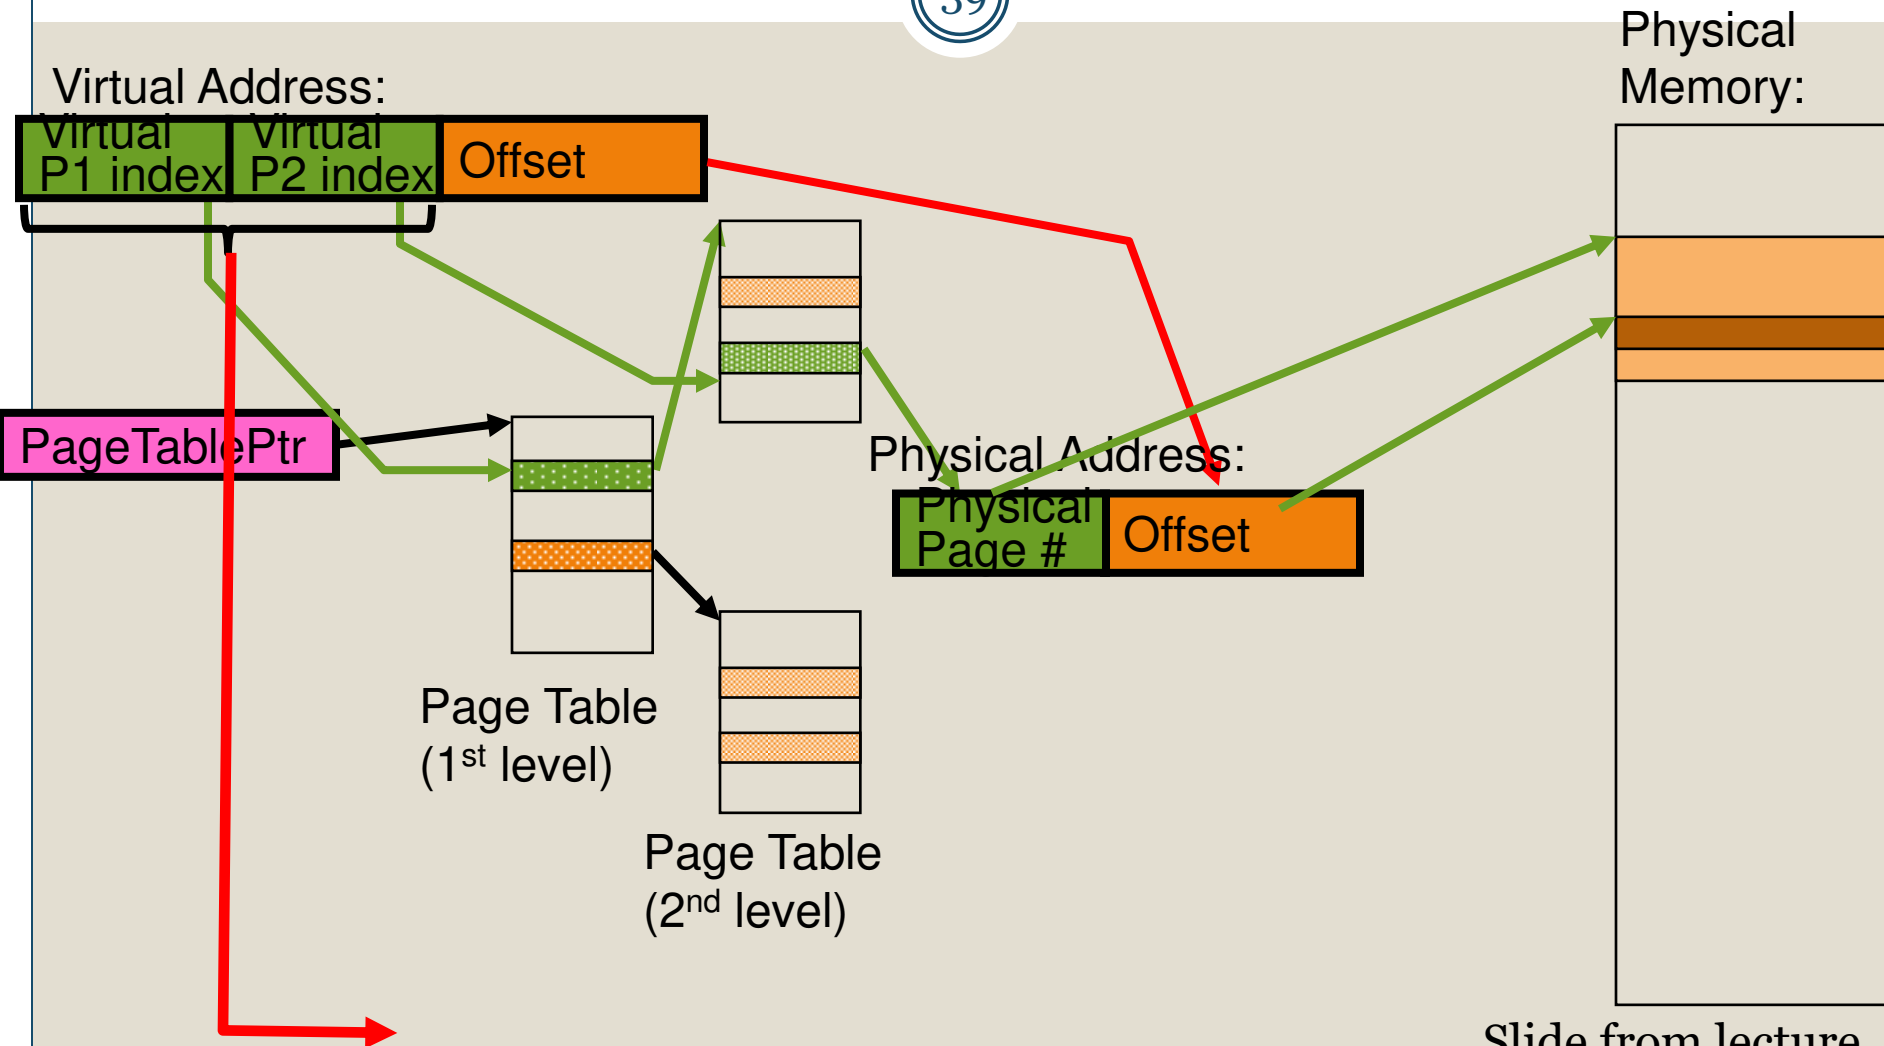
39



Slide from lecture

Putting Everything Together: TLB

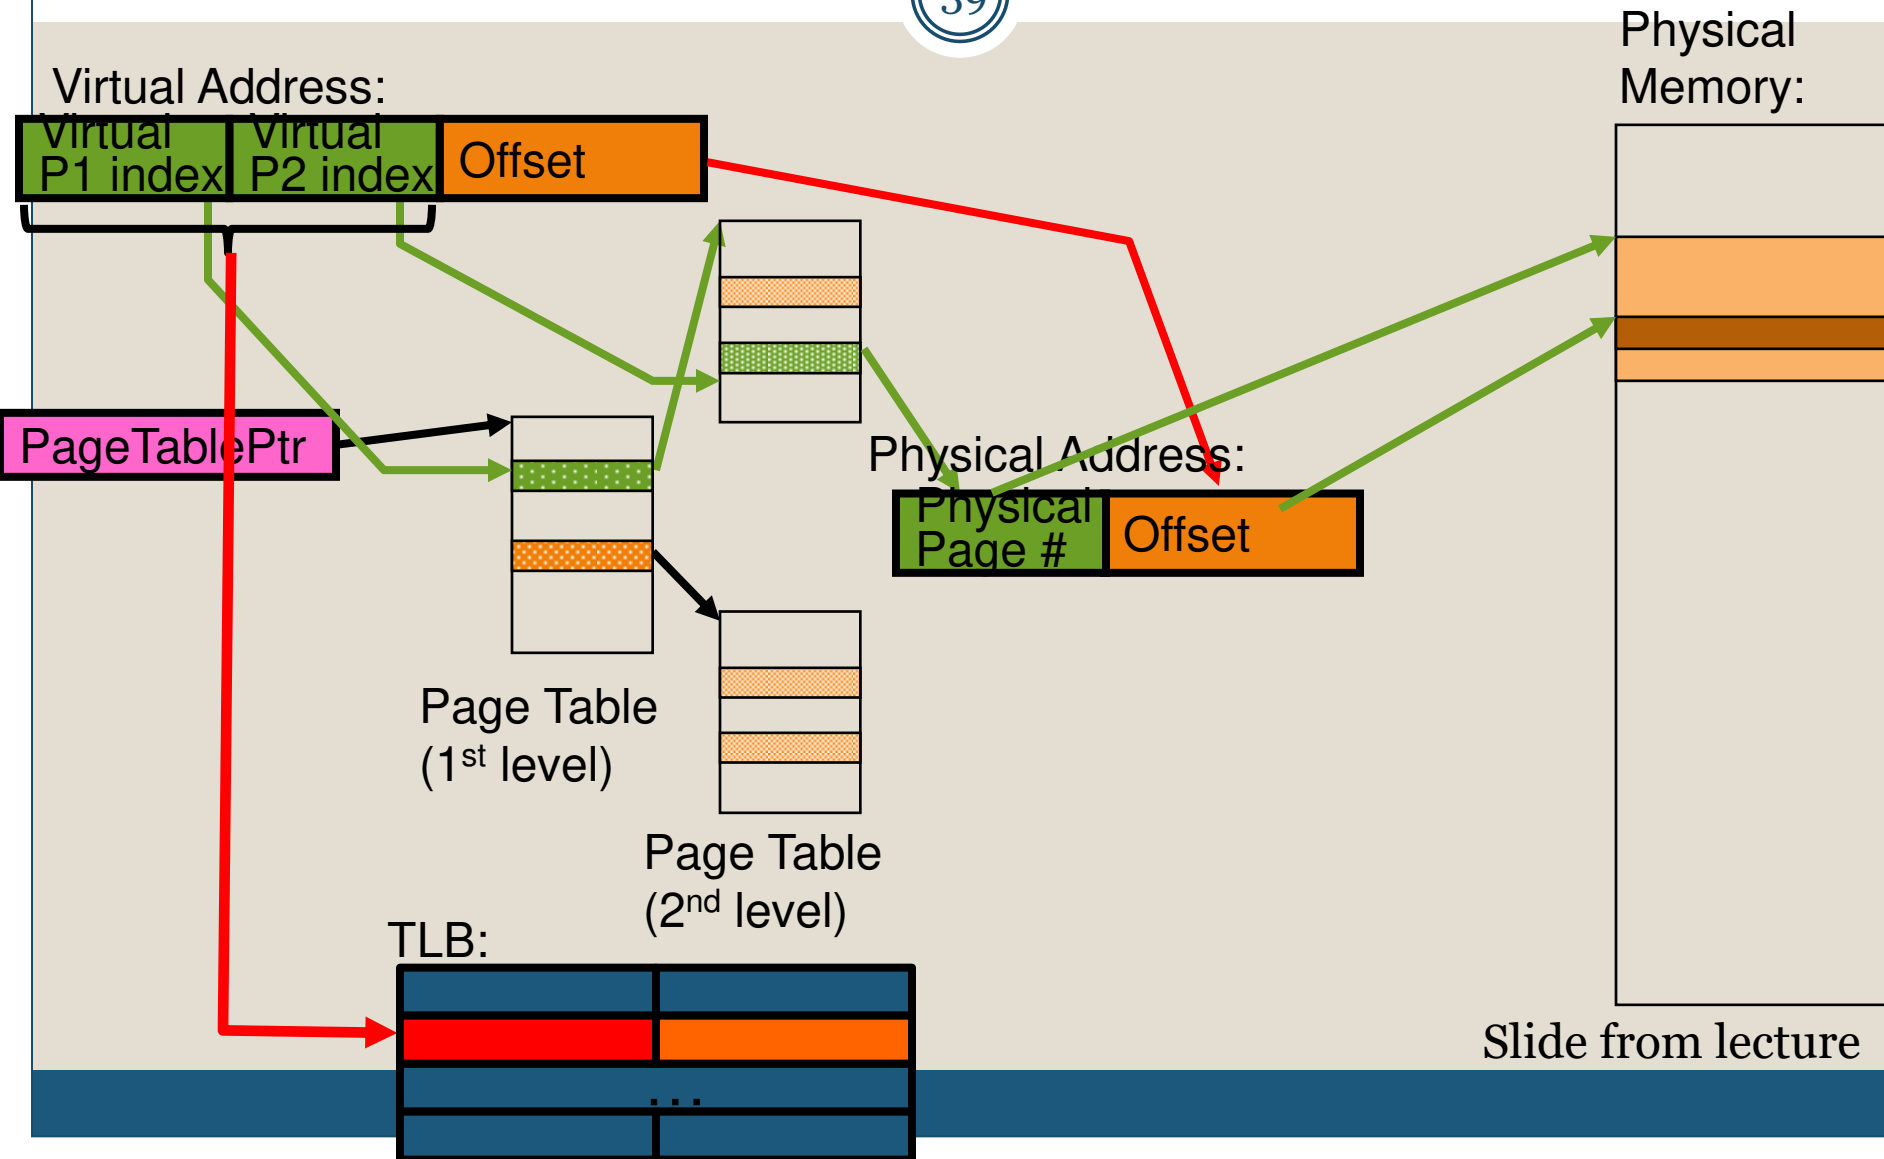
39



Slide from lecture

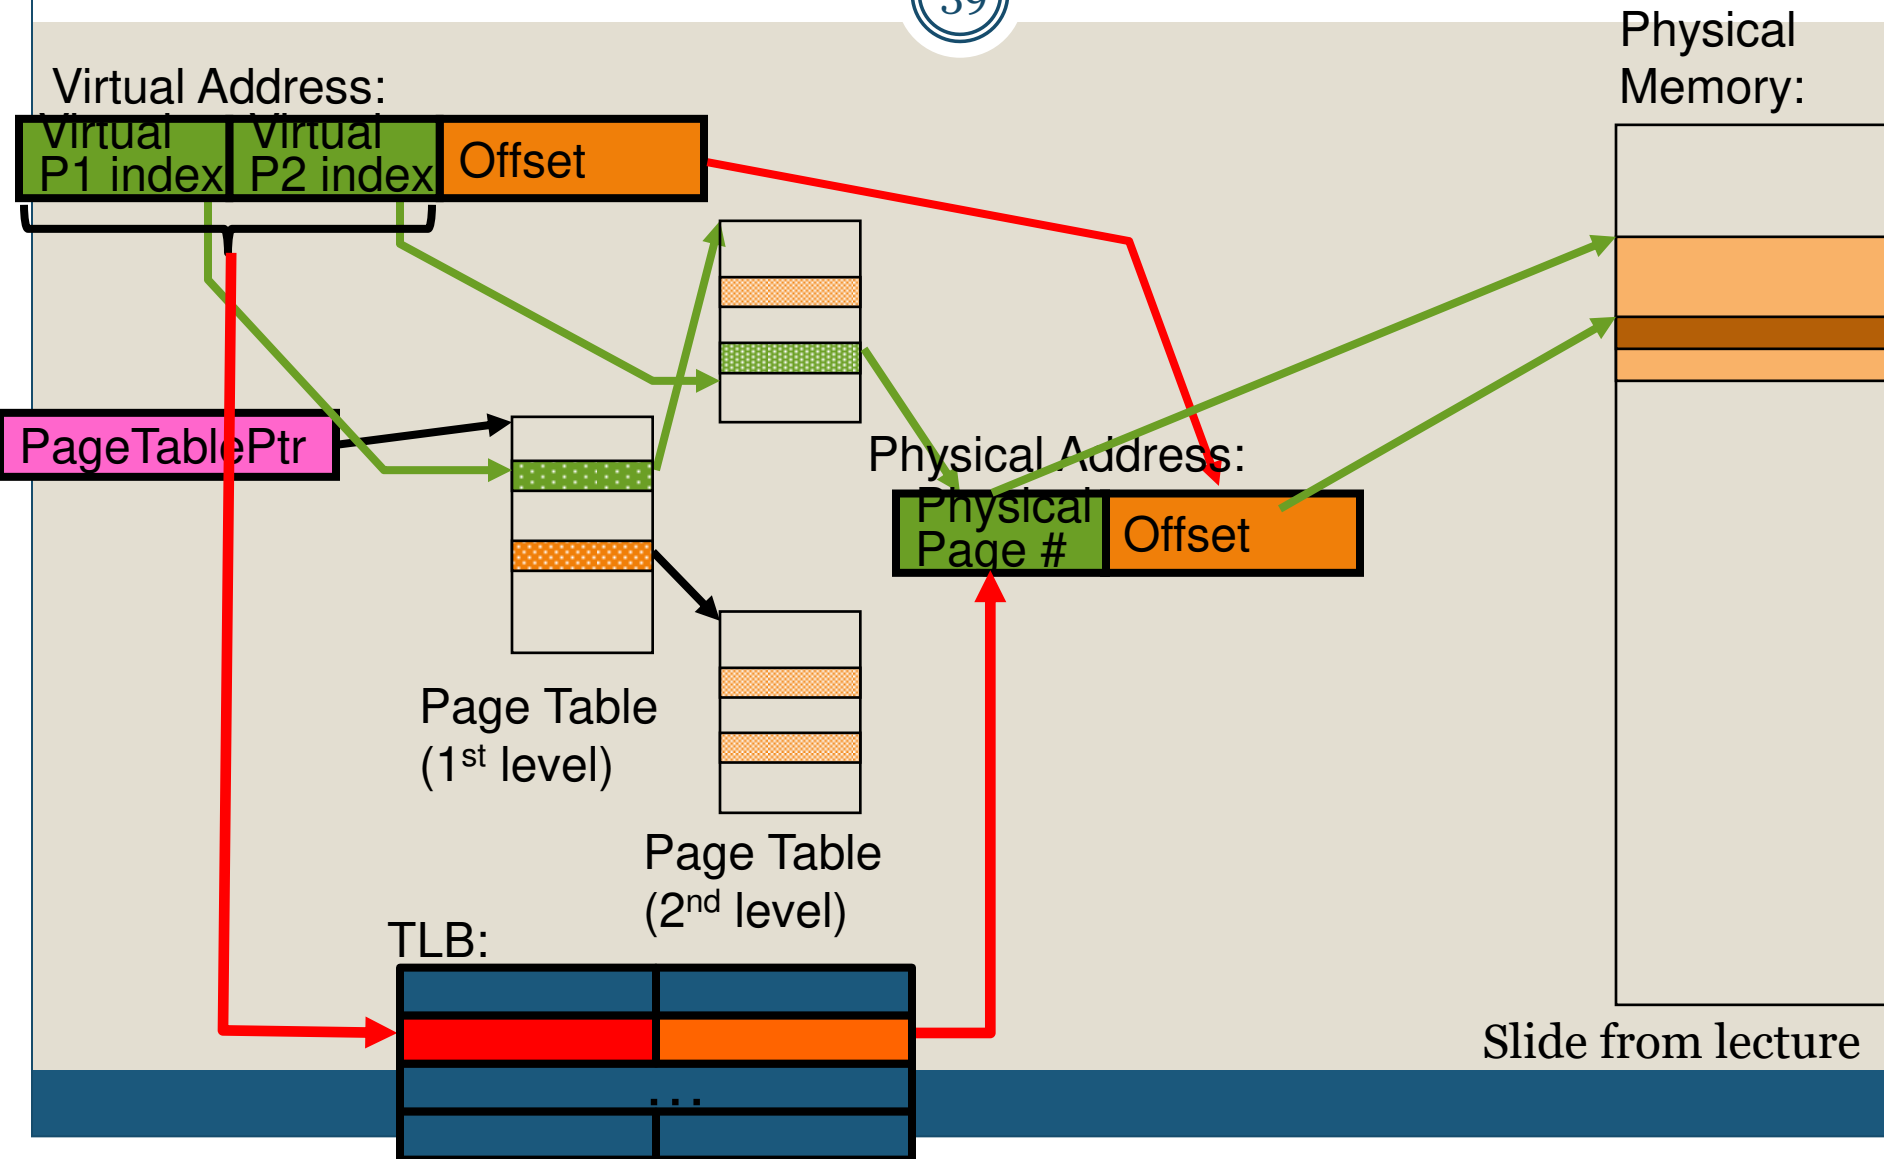
Putting Everything Together: TLB

39



Putting Everything Together: TLB

39



Ex 4

41

- In EX 4 you will implement a virtual memory interface using hierarchical page tables of arbitrary depth using simulated physical memory.
- You'll have to implement 2 functions:
 - VMread(*virtualAddress*, **value*) – reads the word from the virtual address *virtualAddress* into **value*. Returns 1 on success and 0 on failure.
 - VMwrite(*virtualAddress*, *value*) – writes the word *value* into the virtual address *virtualAddress*. Returns 1 on success and 0 on failure.

Q & A

42

